

20090118_Twitter_CommunityGuidelines	4
20090918_Twitter_CommunityGuidelines	6
20091025_Twitter_CommunityGuidelines	9
20100301_Twitter_CommunityGuidelines	12
20100323_Twitter_CommunityGuidelines	22
20101004_Twitter_CommunityGuidelines	31
20101020_Twitter_CommunityGuidelines	40
20110124_Twitter_CommunityGuidelines	48
20110412_Twitter_CommunityGuidelines	56
20110902_Twitter_CommunityGuidelines	64
20110903_Twitter_CommunityGuidelines	73
20120510_Twitter_CommunityGuidelines	82
20120919_Twitter_CommunityGuidelines	91
20121029_Twitter_CommunityGuidelines	101
20130530_Twitter_CommunityGuidelines	111
20130813_Twitter_CommunityGuidelines	128
20130821_Twitter_CommunityGuidelines	147
20130929_Twitter_CommunityGuidelines	166
20131209_Twitter_CommunityGuidelines	185
20140729_Twitter_CommunityGuidelines	204
20141031_Twitter_CommunityGuidelines	223
20150314_Twitter_CommunityGuidelines	242
20150419_Twitter_CommunityGuidelines	261
20150527_Twitter_CommunityGuidelines	280
20150531_Twitter_CommunityGuidelines	299
20150713_Twitter_CommunityGuidelines	318
20151230_Twitter_CommunityGuidelines	338
20160117_Twitter_CommunityGuidelines	358
20160206_Twitter_CommunityGuidelines	376
20160504_Twitter_CommunityGuidelines	395

20160630_Twitter_CommunityGuidelines	414
20170714_Twitter_CommunityGuidelines	431
20171103_Twitter_CommunityGuidelines	449
20171222_Twitter_CommunityGuidelines	479
20171223_Twitter_CommunityGuidelines	508
20171231_Twitter_CommunityGuidelines	537
20180103_Twitter_CommunityGuidelines	565
20180228_Twitter_CommunityGuidelines	593
20180414_Twitter_CommunityGuidelines	621
20180424_Twitter_CommunityGuidelines	649
20180507_Twitter_CommunityGuidelines	677
20180512_Twitter_CommunityGuidelines	705
20180513_Twitter_CommunityGuidelines	734
20180601_Twitter_CommunityGuidelines	762
20180722_Twitter_CommunityGuidelines	790
20181024_Twitter_CommunityGuidelines	818
20181028_Twitter_CommunityGuidelines	846
20181113_Twitter_CommunityGuidelines	877
20181128_Twitter_CommunityGuidelines	908
20181206_Twitter_CommunityGuidelines	941
20190130_Twitter_CommunityGuidelines	974
20190428_Twitter_CommunityGuidelines	1007
20190607_Twitter_CommunityGuidelines	1041
20190727_Twitter_CommunityGuidelines	1075
20190829_Twitter_CommunityGuidelines	1113
20190925_Twitter_CommunityGuidelines	1151
20191001_Twitter_CommunityGuidelines	1189
20191204_Twitter_CommunityGuidelines	1227
20191217_Twitter_CommunityGuidelines	1265
20200113_Twitter_CommunityGuidelines	1303

20200114_Twitter_CommunityGuidelines	1341
20200204_Twitter_CommunityGuidelines	1379
20200331_Twitter_CommunityGuidelines	1418
20200605_Twitter_CommunityGuidelines	1458
20200805_Twitter_CommunityGuidelines	1498
20200917_Twitter_CommunityGuidelines	1539
20200918_Twitter_CommunityGuidelines	1580
20201012_Twitter_CommunityGuidelines	1622
20201026_Twitter_CommunityGuidelines	1663
20201028_Twitter_CommunityGuidelines	1704
20201203_Twitter_CommunityGuidelines	1746
20210120_Twitter_CommunityGuidelines	1787
20210127_Twitter_CommunityGuidelines	1828
20210527_Twitter_CommunityGuidelines	1869
20210626_Twitter_CommunityGuidelines	1911
20210914_Twitter_CommunityGuidelines	1955
20210922_Twitter_CommunityGuidelines	2003
20211006_Twitter_CommunityGuidelines	2051
20211010_Twitter_CommunityGuidelines	2102
20211207_Twitter_CommunityGuidelines	2156
20211130_Twitter_CommunityGuidelines	2216
20211017_Twitter_CommunityGuidelines	2274
20211221_Twitter_CommunityGuidelines	2330

The Twitter Rules

Submitted Jan 14 by crystal

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor user's content and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

***Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others ***Privacy:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National

Identity numbers, without their express authorization and permission.

***Violence and Threats:** You may not publish or post direct, specific threats of violence against others.

***Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.

***Unlawful Use:** You may not use our service for any unlawful purposes or for promotion of illegal activities.

***Serial Accounts:** You may not create serial accounts for disruptive or abusive purposes.

*** Name Squatting:** You may not engage in name squatting Accounts that are inactive for more than 6 months may be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be name squatting are:

the number of accounts created
creating accounts for the purpose of preventing others from using those account names
creating accounts for the purpose of selling those accounts
using feeds of third-party content to update and maintain accounts under the names of those third parties

***Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.

***Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:

If you have followed a large amount of users in a short amount of time;
If you have a small number of followers compared to the amount of people you are following; If your updates consist mainly of links, and not personal updates;
If a large number of people are blocking you;
The number of spam complaints that have been filed against you;
If you post duplicate content over multiple accounts;
If you repost other user's content without attribution.

***Pornography:** You may not use obscene or pornographic images in either your profile picture or user background.

Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the Terms of Service.

The Twitter Rules

Submitted Jan 14 by crystal

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor user's content and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

***Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others

*** Trademark:** We reserve the right to reclaim user names on behalf of businesses or individuals that hold legal claim or trademark on those user names. Accounts

using business names and/or logos to mislead others will be permanently suspended.

***Privacy:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission.

***Violence and Threats:** You may not publish or post direct, specific threats of violence against others.

***Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.

***Unlawful Use:** You may not use our service for any unlawful purposes or for promotion of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.

***Verified accounts:** You may not use the Verified Account badge unless it is provided by Twitter. Accounts using the badge as part of profile pictures, background images, or in any way implying false verification will be permanently suspended.

Spam and Abuse

Twitter strives to protect its users from spam and abuse. Technical abuse and user abuse is not tolerated on Twitter.com, and will result in permanent suspension. Any

accounts engaging in the activities specified below are subject to permanent suspension.

* **Serial Accounts:** You may not create serial accounts for disruptive or abusive purposes. Mass account creation will result in suspension of all accounts.

* **Name Squatting:** You may not engage in name squatting. Accounts that are inactive for more than 6 months may also be removed without further notice. Some of

the factors that we take into account when determining what conduct is considered to be name squatting are:

the number of accounts created

creating accounts for the purpose of preventing others from using those account names

creating accounts for the purpose of selling those accounts

using feeds of third-party content to update and maintain accounts under the names of those third parties

***Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations

***Selling user names:** Unless you have been specifically permitted to do so in a separate agreement with Twitter, you agree that you will not reproduce, duplicate,

copy, sell, trade or resell the Services for any purpose.

Where "Services" is defined as follows: Your use of Twitter's products, services and web sites (referred to collectively as the "Services" in this document and excluding any services provided to you by Twitter under a separate written agreement) is subject to the terms of a legal agreement between you and Twitter.

***Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.

***Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:

If you have followed a large amount of users in a short amount of time;

If you have followed and unfollowed people in a short time period, particularly by automated means (aggressive follower churn); If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;

If you have a small number of followers compared to the amount of people you are following;

If your updates consist mainly of links, and not personal updates;

If a large number of people are blocking you;

The number of spam complaints that have been filed against you;

If you post duplicate content over multiple accounts or multiple duplicate updates on one account

If you post multiple unrelated updates to a topic using #

If you post multiple unrelated updates to a trending or popular topic
If you send large numbers of duplicate @replies
If you send large numbers of unsolicited @replies in an attempt to spam a service or link
If you repost other user's content without attribution.
If you have attempted to "sell" followers, particularly through tactics considered aggressive following or follower churn.

Your account may be suspended for Terms of Service violations if any of the above is true; accounts created to replace suspended accounts will be permanently suspended.

***Pornography:** You may not use obscene or pornographic images in either your profile picture or user background.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the Terms of Service.

The Twitter Rules

Submitted Jan 14 by crystal

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor user's content and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

***Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others

*** Trademark:** We reserve the right to reclaim user names on behalf of businesses or individuals that hold legal claim or trademark on those user names. Accounts using business names and/or logos to mislead others will be permanently suspended.

***Privacy:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission.

***Violence and Threats:** You may not publish or post direct, specific threats of violence against others. ***Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright

procedures are set forth in the Terms of Service.

***Unlawful Use:** You may not use our service for any unlawful purposes or for promotion of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.

***Verified accounts:** You may not use the Verified Account badge unless it is provided by Twitter. Accounts using the badge as part of profile pictures, background images, or in any way implying false verification will be permanently suspended.

Spam and Abuse

Twitter strives to protect its users from spam and abuse. Technical abuse and user abuse is not tolerated on Twitter.com, and will result in permanent suspension. Any

accounts engaging in the activities specified below are subject to permanent suspension.

* **Serial Accounts:** You may not create serial accounts for disruptive or abusive purposes. Mass account creation will result in suspension of all accounts.

* **Name Squatting:** You may not engage in name squatting. Accounts that are inactive for more than 6 months may also be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be name squatting are:

the number of accounts created

creating accounts for the purpose of preventing others from using those account names

creating accounts for the purpose of selling those accounts

using feeds of third-party content to update and maintain accounts under the names of those third parties

***Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations ***Selling user names:** Unless you have been specifically permitted to do so in a separate agreement with Twitter, you

agree that you will not reproduce, duplicate, copy, sell, trade or resell the Services for any purpose.

Where "Services" is defined as follows: Your use of Twitter's products, services and web sites (referred to collectively as the "Services" in this document and excluding any services provided to you by Twitter under a separate written agreement) is subject to the terms of a legal agreement between you and Twitter.

***Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.

***Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:

If you have followed a large amount of users in a short amount of time;

If you have followed and unfollowed people in a short time period, particularly by automated means (aggressive follower churn);

If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;

If you have a small number of followers compared to the amount of people you are following; If your updates consist mainly of links, and not personal updates;
If a large number of people are blocking you;
The number of spam complaints that have been filed against you;

If you post duplicate content over multiple accounts or multiple duplicate updates on one account
If you post multiple unrelated updates to a topic using #
If you post multiple unrelated updates to a trending or popular topic
If you send large numbers of duplicate @replies

If you send large numbers of unsolicited @replies in an attempt to spam a service or link
If you repost other user's content without attribution.

If you have attempted to "sell" followers, particularly through tactics considered aggressive following or follower churn.

Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising "more followers fast," or any other site that offers to automatically add followers to your account).

Your account may be suspended for Terms of Service violations if any of the above is true; accounts created to replace suspended accounts will be permanently suspended.

***Pornography:** You may not use obscene or pornographic images in either your profile picture or user background.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the Terms of Service.

If you post multiple unrelated updates to a trending or popular topic
If you send large numbers of duplicate @replies
If you send large numbers of unsolicited @replies in an attempt to spam a service or link
If you repost other user's content without attribution.
If you have attempted to "sell" followers, particularly through tactics considered aggressive following or follower churn.

Your account may be suspended for Terms of Service violations if any of the above is true; accounts created to replace suspended accounts will be permanently suspended.

***Pornography:** You may not use obscene or pornographic images in either your profile picture or user background.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the Terms of Service.

The Twitter Rules (20100301)

Submitted Jan 14, 2009 by **crystal**

The Twitter Rules

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor user's content and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

***Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others

***Trademark:** We reserve the right to reclaim user names on behalf of businesses or individuals that hold legal claim or trademark on those user names. Accounts using business names and/or logos to mislead others will be permanently suspended.

***Privacy:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission.

***Violence and Threats:** You may not publish or post direct, specific threats of violence against others.

***Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.

***Unlawful Use:** You may not use our service for any unlawful purposes or for promotion of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.

***Verified accounts:** You may not use the Verified Account badge unless it is provided by Twitter. Accounts using the badge as part of profile pictures,

background images, or in any way implying false verification will be permanently suspended.

Spam and Abuse

Twitter strives to protect its users from spam and abuse. Technical abuse and user abuse is not tolerated on Twitter.com, and will result in permanent suspension. Any accounts engaging in the activities specified below are subject to permanent suspension.

* **Serial Accounts:** You may not create serial accounts for disruptive or abusive purposes, or with overlapping use cases. Mass account creation may result in suspension of all related accounts. Please note that any violation of the Twitter Rules is cause for permanent suspension of all accounts.

* **Name Squatting:** You may not engage in name squatting. Accounts that are inactive for more than 6 months may also be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be name squatting are:

- the number of accounts created
- creating accounts for the purpose of preventing others from using those account names
- creating accounts for the purpose of selling those accounts
- using feeds of third-party content to update and maintain accounts under the names of those third parties

***Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations

***Selling user names:** Unless you have been specifically permitted to do so in a separate agreement with Twitter, you agree that you will not reproduce, duplicate, copy, sell, trade or resell the Services for any purpose.

Where "Services" is defined as follows: Your use of Twitter's products, services and web sites (referred to collectively as the "Services" in this document and excluding any services provided to you by Twitter under a separate written agreement) is subject to the terms of a legal agreement between you and Twitter.

***Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.

***Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes “spamming” will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:

- If you have followed a large amount of users in a short amount of time;
- If you have followed and unfollowed people in a short time period, particularly by automated means (aggressive follower churn);
- If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
- If you have a small number of followers compared to the amount of people you are following;
- If your updates consist mainly of links, and not personal updates;
- If you post misleading links;
- If a large number of people are blocking you;
- The number of spam complaints that have been filed against you;
- If you post duplicate content over multiple accounts or multiple duplicate updates on one account;
- If you post multiple unrelated updates to a topic using #;
- If you post multiple unrelated updates to a trending or popular topic;
- If you send large numbers of duplicate @replies;
- If you send large numbers of unsolicited @replies in an attempt to spam a service or link;
- If you add a large number of unrelated users to lists in an attempt to spam a service or link;
- If you repeatedly post other users' Tweets as your own;
- If you have attempted to "sell" followers, particularly through tactics considered aggressive following or follower churn;
- Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising "more followers fast," or any other site that offers to automatically add followers to your account).

Your account may be suspended for Terms of Service violations if any of the above is true. Please see our help pages on [Following Best Practices](#) and [Automation Rules and Best Practices](#) for a more detailed discussion of how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

***Pornography:** You may not use obscene or pornographic images in either your profile picture or user background.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account

without further notice in the event that, in its judgment, you violate these Rules or the Terms of Service.

Topic is closed for comments

Following Limits and Best Practices (20091014)

Submitted Oct 14 by ginger

What is following?

Following is a way to receive information, in the form of tweets, from a person, company, or organization. If an account seems interesting, you can follow that account and see their updates in your timeline. This page discusses some of the limits and best practices regarding following on Twitter.

What are the limits?

We don't limit the number of followers you can have. However, we do monitor how aggressively users follow other users. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

What is aggressive following?

Following users is a way to see their updates in your timeline. If an account seems interesting, feel free to follow it! Many Twitter users receive email notifications when someone follows their account or check out the profiles of new followers to see if they share common interests. If some accounts are aggressively or indiscriminately following hundreds of accounts just to garner attention, it makes Twitter a less-nice place to hang out.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, that's fine! You can just visit the person's profile page and then un-follow them. Aggressive follow churn is when an account repeatedly follows and un-follows large numbers of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio.

These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Additional limits if you are following 2000 or more people:

The rules about aggressive following and follow churn still apply. In addition, every user can follow 2000 people total. Once you've followed 2000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following. When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—basically, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or

total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improves site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including high profile and API accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter monitors following behaviors and limits some behavior in order to make using Twitter a more positive experience for everyone.

Twitter facilitates social networking, but it's not a social networking website. In fact, the way Twitter works is quite different from social networks: when you accept friend requests on social networks, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different, because instead of indicating a mutual relationship, following is a one-way action that means you want to receive information, in the form of tweets, from someone. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring mutual following.

If you follow too many people, there's no way you can keep up with everyone's updates in your timeline. If you need to communicate with someone but don't need to see their updates everyday, you don't have to follow them. Send them an @reply when you need to; it doesn't require following and your update will appear in the person's replies tab, so they can reply back.

In addition, limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter, so monitoring for abuse is one way to reduce spam on Twitter.

Using Third Party Applications to “Get More Followers Fast!”

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. [The Twitter Rules](#) addresses the use of these programs by disallowing:

- Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. Please do not give your username and password out to any third-party application that you have not thoroughly researched.

Following limits? But I'm whitelisted!

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users, and degrades the Twitter experience for everyone.

Automation Rules & Best Practices (201001)

Submitted Nov 08, 2009 by [ginger](#)

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password.

Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically

sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

Updates posted using the general API *without Twitter OAuth* will instead say "from API" as the source. We do not allow advertising posted through general API applications. If you find your account posting ads such as these, please change your password and revoke unwanted connections. You can find more information, and a technical walk-through for allowing and revoking API access, the help page on [Connecting to Third-Party Applications](#).

The Twitter Rules (20100323)

crystal Jan 14, 2009

The Twitter Rules

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor user's content and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

***Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others

***Trademark:** We reserve the right to reclaim user names on behalf of businesses or individuals that hold legal claim or trademark on those user names. Accounts using business names and/or logos to mislead others will be permanently suspended.

***Privacy:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission.

***Violence and Threats:** You may not publish or post direct, specific threats of violence against others.

***Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.

***Unlawful Use:** You may not use our service for any unlawful purposes or for promotion of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.

***Verified accounts:** You may not use the Verified Account badge unless it is provided by Twitter. Accounts using the badge as part of profile pictures, background images, or in any way implying false verification will be permanently suspended.

Spam and Abuse

Twitter strives to protect its users from spam and abuse. Technical abuse and user abuse is not tolerated on Twitter.com, and will result in permanent suspension. Any accounts engaging in the activities specified below are subject to permanent suspension.

*** Serial Accounts:** You may not create serial accounts for disruptive or abusive purposes, or with overlapping use cases. Mass account creation may result in suspension of all related accounts. Please note that any violation of the Twitter Rules is cause for permanent suspension of all accounts.

*** Name Squatting:** You may not engage in name squatting. Accounts that are inactive for more than 6 months may also be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be name squatting are:

- the number of accounts created
- creating accounts for the purpose of preventing others from using those account names
- creating accounts for the purpose of selling those accounts
- using feeds of third-party content to update and maintain accounts under the names of those third parties

***Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations

***Selling user names:** Unless you have been specifically permitted to do so in a separate agreement with Twitter, you agree that you will not reproduce, duplicate, copy, sell, trade or resell the Services for any purpose.

Where "Services" is defined as follows: Your use of Twitter's products, services and web sites (referred to collectively as the "Services" in this document and excluding any services provided to you by Twitter under a separate written agreement) is subject to the terms of a legal agreement between you and Twitter.

***Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.

***Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:

- If you have followed a large amount of users in a short amount of time;
- If you have followed and unfollowed people in a short time period, particularly by automated means (aggressive follower churn);
- If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
- If you have a small number of followers compared to the amount of people you are following;
- If your updates consist mainly of links, and not personal updates;
- If you post misleading links;
- If a large number of people are blocking you;
- The number of spam complaints that have been filed against you;
- If you post duplicate content over multiple accounts or multiple duplicate updates on one account;
- If you post multiple unrelated updates to a topic using #;
- If you post multiple unrelated updates to a trending or popular topic;
- If you send large numbers of duplicate @replies;
- If you send large numbers of unsolicited @replies in an attempt to spam a service or link;
- If you add a large number of unrelated users to lists in an attempt to spam a service or link;
- If you repeatedly post other users' Tweets as your own;
- If you have attempted to "sell" followers, particularly through tactics considered aggressive following or follower churn;
- Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising "more followers fast," or any other site that offers to automatically add followers to your account).

Your account may be suspended for Terms of Service violations if any of the above is true. Please see our help pages on [Following Best Practices](#) and [Automation Rules and Best Practices](#) for a more detailed discussion of how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

***Pornography:** You may not use obscene or pornographic images in either your profile picture or user background.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the Terms of Service.

Following Rules and Best Practices **(201005*)**

What is following?

Following is a way to receive information, in the form of tweets, from a person, company, or organization. If an account seems interesting, you can follow that account and see their updates in your timeline. This page discusses some of the limits and best practices regarding following on Twitter.

What are the limits?

We don't limit the number of followers you can have. However, we do monitor how aggressively users follow other users. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine. Please note that the only automated following behavior that Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. Please review our [Automation Rules & Best Practices](#) for more information on automating your account.

What is aggressive following?

Following users is a way to see their updates in your timeline. If an account seems interesting, feel free to follow it! Many Twitter users receive email notifications when someone follows their account or check out the profiles of new followers to see if they share common interests. If some accounts are aggressively or indiscriminately following hundreds of accounts just to garner attention, it makes Twitter a less-nice place to hang out.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, that's fine! You can just visit the person's profile page and then un-follow them. Aggressive follow churn is when an account repeatedly follows and un-follows large numbers of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio.

These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Additional limits if you are following 2000 or more people:

The rules about aggressive following and follow churn still apply. In addition, every user can follow 2000 people

total. Once you've followed 2000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following. When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—basically, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter monitors following behaviors and limits some behavior in order to make using Twitter a more positive experience for everyone.

Twitter facilitates social networking, but it's not a social networking website. In fact, the way Twitter works is quite different from social networks: when you accept friend requests on social networks, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different, because instead of indicating a mutual relationship, following is a one-way action that means you want to receive information, in the form of tweets, from someone. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring mutual following.

If you follow too many people, there's no way you can keep up with everyone's updates in your timeline. If you need to communicate with someone but don't need to see their updates everyday, you don't have to follow them. Send them an @reply when you need to; it doesn't require following and your update will appear in the person's replies tab, so they can reply back.

In addition, limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter, so monitoring for abuse is one way to reduce spam on Twitter.

Using Third Party Applications to “Get More Followers Fast!”

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. [The Twitter Rules](#) addresses the use of these programs by disallowing:

- Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. Please do not give your username and password out to any third-party application that you have not thoroughly researched.

Following limits? But I'm whitelisted!

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users, and degrades the Twitter experience for everyone.

Automation Rules & Best Practices (201001)

Submitted Nov 08, 2009 by ginger

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted

automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#).

Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

Updates posted using the general API *without Twitter OAuth* will instead say "from API" as the source. We do not allow advertising posted through general API applications. If you find your account posting ads such as these, please change your password and revoke unwanted connections. You can find more information, and a technical walk-through for allowing and revoking API access, the help page on [Connecting to Third-Party Applications](#).

The Twitter Rules (20101004)

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor user's content and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

- **Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others
- **Trademark:** We reserve the right to reclaim user names on behalf of businesses or individuals that hold legal claim or trademark on those user names. Accounts using business names and/or logos to mislead others will be permanently suspended.
- **Privacy:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission.
- **Violence and Threats:** You may not publish or post direct, specific threats of violence against others.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Unlawful Use:** You may not use our service for any unlawful purposes or for promotion of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Verified accounts:** You may not use the Verified Account badge unless it is provided by Twitter. Accounts using the badge as part of profile pictures, background images, or in a way that falsely implies verification will be suspended.

Spam and Abuse

Twitter strives to protect its users from spam and abuse. Technical abuse and user abuse is not tolerated on Twitter.com, and will result in permanent suspension. Any accounts engaging in the activities specified below are subject to permanent suspension.

- **Serial Accounts:** You may not create serial accounts for disruptive or abusive purposes, or with overlapping use cases. Mass account creation may result in suspension of all related accounts. Please note that any violation of the Twitter Rules is cause for permanent suspension of all accounts.
- **Username Squatting:** You may not engage in username squatting. Accounts that are inactive for more than 6 months may also be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be username squatting are:
 - the number of accounts created
 - creating accounts for the purpose of preventing others from using those account names
 - creating accounts for the purpose of selling those accounts
 - using feeds of third-party content to update and maintain accounts under the names of those third parties
- **Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations.
- **Selling user names:** Unless you have been specifically permitted to do so in a separate agreement with Twitter, you agree that you will not reproduce, duplicate, copy, sell, trade or resell the Services for any purpose, where "Services" is defined as follows: Your use of Twitter's products, services and web sites (referred to collectively as the "Services" in this document and excluding any services provided to

you by Twitter under a separate written agreement) is subject to the terms of a legal agreement between you and Twitter.

- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - If you have followed a large amount of users in a short amount of time;
 - If you have followed and unfollowed people in a short time period, particularly by automated means (aggressive follower churn);
 - If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - If you have a small number of followers compared to the amount of people you are following;
 - If your updates consist mainly of links, and not personal updates;
 - If you post misleading links;
 - If a large number of people are blocking you;
 - The number of spam complaints that have been filed against you;
 - If you post duplicate content over multiple accounts or multiple duplicate updates on one account;
 - If you post multiple unrelated updates to a topic using #;
 - If you post multiple unrelated updates to a trending or popular topic;
 - If you send large numbers of duplicate @replies;
 - If you send large numbers of unsolicited @replies in an attempt to spam a service or link;
 - If you add a large number of unrelated users to lists in an attempt to spam a service or link;
 - If you repeatedly post other users' Tweets as your own;
 - If you have attempted to "sell" followers, particularly through tactics considered aggressive following or follower churn;
 - Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising "more followers fast," or any other site that offers to automatically add followers to your account);
 - If you create false or misleading Points of Interest;
 - If you create Points of Interest to namesquat or spam.
- **Pornography:** You may not use obscene or pornographic images in either your profile picture or user background

Your account may be suspended for Terms of Service violations if any of the above is true. Please see our help pages on [Following Best Practices](#) and [Automation Rules and Best Practices](#) for a more detailed discussion of how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the Terms of Service.

Following Rules and Best Practices **(201005*)**

What is following?

Following is a way to receive information, in the form of tweets, from a person, company, or organization. If an account seems interesting, you can follow that account and see their updates in your timeline. This page discusses some of the limits and best practices regarding following on Twitter.

What are the limits?

We don't limit the number of followers you can have. However, we do monitor how aggressively users follow other users. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine. Please note that the only automated following behavior that Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. Please review our [Automation Rules & Best Practices](#) for more information on automating your account.

What is aggressive following?

Following users is a way to see their updates in your timeline. If an account seems interesting, feel free to follow it! Many Twitter users receive email notifications when someone follows their account or check out the profiles of new followers to see if they share common interests. If some accounts are aggressively or indiscriminately following hundreds of accounts just to garner attention, it makes Twitter a less-nice place to hang out.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, that's fine! You can just visit the person's profile page and then un-follow them. Aggressive follow churn is when an account repeatedly follows and un-follows large numbers of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio.

These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Additional limits if you are following 2000 or more people:

The rules about aggressive following and follow churn still apply. In addition, every user can follow 2000 people total. Once you've followed 2000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following. When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—basically, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter monitors following behaviors and limits some behavior in order to make using Twitter a more positive experience for everyone.

Twitter facilitates social networking, but it's not a social networking website. In fact, the way Twitter works is quite different from social networks: when you accept friend requests on social networks, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different, because instead of indicating a mutual relationship, following is a one-way action that means you want to receive information, in the form of tweets, from someone. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring mutual following.

If you follow too many people, there's no way you can keep up with everyone's updates in your timeline. If you need to communicate with someone but don't need to see their updates everyday, you don't have to follow them. Send them an @reply when you need to; it doesn't require following and your update will appear in the person's replies tab, so they can reply back.

In addition, limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter, so monitoring for abuse is one way to reduce spam on Twitter.

Using Third Party Applications to “Get More Followers Fast!”

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. [The Twitter Rules](#) addresses the use of these programs by disallowing:

- Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. Please do not give your username and password out to any third-party application that you have not thoroughly researched.

Following limits? But I'm whitelisted!

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users, and degrades the Twitter experience for everyone.

Automation Rules & Best Practices (2010/01)

Submitted Nov 08, 2009 by **ginger**

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.

3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of

users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

Updates posted using the general API *without Twitter OAuth* will instead say "from API" as the source. We do not allow advertising posted through general API applications. If you find your account posting ads such as these,

please change your password and revoke unwanted connections. You can find more information, and a technical walk-through for allowing and revoking API access, the help page on [Connecting to Third-Party Applications](#).

The Twitter Rules (20101020)

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor user's content and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

- Impersonation: You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others
- Trademark: We reserve the right to reclaim user names on behalf of businesses or individuals that hold legal claim or trademark on those user names. Accounts using business names and/or logos to mislead others will be permanently suspended.
- Privacy: You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission.
- Violence and Threats: You may not publish or post direct, specific threats of violence against others.
- Copyright: We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- Unlawful Use: You may not use our service for any unlawful purposes or for promotion of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- Verified accounts: You may not use the Verified Account badge unless it is provided by Twitter. Accounts using the badge as part of profile pictures, background images, or in a way that falsely implies verification will be suspended.

Spam and Abuse

Twitter strives to protect its users from spam and abuse. Technical abuse and user abuse is not tolerated on Twitter.com, and will result in permanent suspension. Any accounts engaging in the activities specified below are subject to permanent suspension.

- Serial Accounts: You may not create serial accounts for disruptive or abusive purposes, or with overlapping use cases. Mass account creation may result in suspension of all related accounts. Please note that any violation of the Twitter Rules is cause for permanent suspension of all accounts.
- Username Squatting: You may not engage in username squatting. Accounts that are inactive for more than 6 months may also be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be username squatting are:
 - the number of accounts created
 - creating accounts for the purpose of preventing others from using those account names
 - creating accounts for the purpose of selling those accounts
 - using feeds of third-party content to update and maintain accounts under the names of those third parties
- Invitation spam: You may not use Twitter.com's address book contact import to send repeat, mass invitations.
- Selling user names: Unless you have been specifically permitted to do so in a separate agreement with Twitter, you agree that you will not reproduce, duplicate, copy, sell, trade or resell the Services for any purpose, where "Services" is defined as follows: Your use of Twitter's products, services and web sites (referred to collectively as the "Services" in this document and excluding any services provided to you by Twitter under a separate written agreement) is subject to the terms of a legal agreement between you and Twitter.
- Malware/Phishing: You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.
- Spam: You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:

- If you have followed a large amount of users in a short amount of time;
- If you have followed and unfollowed people in a short time period, particularly by automated means (aggressive follower churn);
- If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
- If you have a small number of followers compared to the amount of people you are following;
- If your updates consist mainly of links, and not personal updates;
- If you post misleading links;
- If a large number of people are blocking you;
- The number of spam complaints that have been filed against you;
- If you post duplicate content over multiple accounts or multiple duplicate updates on one account;
- If you post multiple unrelated updates to a topic using #;
- If you post multiple unrelated updates to a trending or popular topic;
- If you send large numbers of duplicate @replies;
- If you send large numbers of unsolicited @replies in an attempt to spam a service or link;
- If you add a large number of unrelated users to lists in an attempt to spam a service or link;
- If you repeatedly post other users' Tweets as your own;
- If you have attempted to "sell" followers, particularly through tactics considered aggressive following or follower churn;
- Creating or purchasing accounts in order to gain followers;
- Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising "more followers fast," or any other site that offers to automatically add followers to your account);
- If you create false or misleading Points of Interest;
- If you create Points of Interest to namesquat or spam.
- Pornography: You may not use obscene or pornographic images in either your profile picture or user background

Your account may be suspended for Terms of Service violations if any of the above is true. Please see our help pages on [Following Best Practices](#) and [Automation Rules and Best Practices](#) for a more detailed discussion of how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the Terms of Service.

Following Rules and Best Practices (201005*)

What is following?

Following is a way to receive information, in the form of tweets, from a person, company, or organization. If an account seems interesting, you can follow that account and see their updates in your timeline. This page discusses some of the limits and best practices regarding following on Twitter.

What are the limits?

We don't limit the number of followers you can have. However, we do monitor how aggressively users follow other users. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine. Please note that the only automated following behavior that Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. Please review our [Automation Rules & Best Practices](#) for more information on automating your account.

What is aggressive following?

Following users is a way to see their updates in your timeline. If an account seems interesting, feel free to follow it! Many Twitter users receive email notifications when someone follows their account or check out the profiles of new followers to see if they share common interests. If some accounts are aggressively or indiscriminately following hundreds of accounts just to garner attention, it makes Twitter a less-nice place to hang out.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, that's fine! You can just visit the person's profile page and then un-follow them. Aggressive follow churn is when an account repeatedly follows and unfollows large numbers of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio.

These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Additional limits if you are following 2000 or more people:

The rules about aggressive following and follow churn still apply. In addition, every user can follow 2000 people total. Once you've followed 2000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following. When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—basically, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter monitors following behaviors and limits some behavior in order to make using Twitter a more positive experience for everyone.

Twitter facilitates social networking, but it's not a social networking website. In fact, the way Twitter works is quite different from social networks: when you accept friend requests on social networks, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different, because instead of indicating a mutual relationship, following is a one-way action that means you want to receive information, in the form of tweets, from someone. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring mutual following.

If you follow too many people, there's no way you can keep up with everyone's updates in your timeline. If you

need to communicate with someone but don't need to see their updates everyday, you don't have to follow them. Send them an @reply when you need to; it doesn't require following and your update will appear in the person's replies tab, so they can reply back.

In addition, limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter, so monitoring for abuse is one way to reduce spam on Twitter.

Using Third Party Applications to “Get More Followers Fast!”

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. [The Twitter Rules](#) addresses the use of these programs by disallowing:

- Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. Please do not give your username and password out to any third-party application that you have not thoroughly researched.

Following limits? But I'm whitelisted!

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users, and degrades the Twitter experience for everyone.

Automation Rules & Best Practices (2010/01)

Submitted Nov 08, 2009 by [ginger](#)

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts

that help people stay updated on their [city's emergency services](#), connect to national health alerts, and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#).

Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated “thanks for following” message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don’t like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

The Twitter Rules (20110124)

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor user's content and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

- **Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others
- **Trademark:** We reserve the right to reclaim user names on behalf of businesses or individuals that hold legal claim or trademark on those user names. Accounts using business names and/or logos to mislead others will be permanently suspended.
- **Privacy:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission.
- **Violence and Threats:** You may not publish or post direct, specific threats of violence against others.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Unlawful Use:** You may not use our service for any unlawful purposes or for promotion of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter Badges:** You may not use a Verified Account badge or Promoted Products badge unless it is provided by Twitter. Accounts using these badges as part of profile pictures, background images, or in a way that falsely implies affiliation with Twitter will be suspended.

Spam and Abuse

Twitter strives to protect its users from spam and abuse. Technical abuse and user abuse is not tolerated on Twitter.com, and will result in permanent suspension. Any accounts engaging in the activities specified below are subject to permanent suspension.

- **Serial Accounts:** You may not create serial accounts for disruptive or abusive purposes, or with overlapping use cases. Mass account creation may result in suspension of all related accounts. Please note that any violation of the Twitter Rules is cause for permanent suspension of all accounts.
- **Username Squatting:** You may not engage in username squatting. Accounts that are inactive for more than 6 months may also be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be username squatting are:
 - the number of accounts created
 - creating accounts for the purpose of preventing others from using those account names
 - creating accounts for the purpose of selling those accounts
 - using feeds of third-party content to update and maintain accounts under the names of those third parties
- **Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations.
- **Selling user names:** Unless you have been specifically permitted to do so in a separate agreement with Twitter, you agree that you will not reproduce, duplicate, copy, sell, trade or resell the Services for any purpose, where "Services" is defined as follows: Your use of Twitter's products, services and web sites (referred to collectively as the "Services" in this document and excluding any services provided to you by Twitter under a separate written agreement) is subject to the terms of a legal agreement between you and Twitter.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.

- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - If you have followed a large amount of users in a short amount of time;
 - If you have followed and unfollowed people in a short time period, particularly by automated means (aggressive follower churn);
 - If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - If you have a small number of followers compared to the amount of people you are following;
 - If your updates consist mainly of links, and not personal updates;
 - If you post misleading links;
 - If a large number of people are blocking you;
 - The number of spam complaints that have been filed against you;
 - If you post duplicate content over multiple accounts or multiple duplicate updates on one account;
 - If you post multiple unrelated updates to a topic using #;
 - If you post multiple unrelated updates to a trending or popular topic;
 - If you send large numbers of duplicate @replies or mentions;
 - If you send large numbers of unsolicited @replies or mentions in an attempt to spam a service or link;
 - If you add a large number of unrelated users to lists in an attempt to spam a service or link;
 - If you repeatedly post other users' Tweets as your own;
 - If you have attempted to "sell" followers, particularly through tactics considered aggressive following or follower churn;
 - Creating or purchasing accounts in order to gain followers;
 - Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising "more followers fast," or any other site that offers to automatically add followers to your account);
 - If you create false or misleading Points of Interest;
 - If you create Points of Interest to namesquat or spam.
- **Pornography:** You may not use obscene or pornographic images in either your profile picture or user background

Your account may be suspended for Terms of Service violations if any of the above is true. Please see our help pages on [Following Best Practices](#) and [Automation Rules and Best Practices](#) for a more detailed discussion of how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the Terms of Service.

Following Rules and Best Practices **(201005*)**

What is following?

Following is a way to receive information, in the form of tweets, from a person, company, or organization. If an account seems interesting, you can follow that account and see their updates in your timeline. This page discusses some of the limits and best practices regarding following on Twitter.

What are the limits?

We don't limit the number of followers you can have. However, we do monitor how aggressively users follow other users. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine. Please note that the only automated following behavior that Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. Please review our [Automation Rules & Best Practices](#) for more information on automating your account.

What is aggressive following?

Following users is a way to see their updates in your timeline. If an account seems interesting, feel free to follow it! Many Twitter users receive email notifications when someone follows their account or check out the profiles of new followers to see if they share common interests. If some accounts are aggressively or indiscriminately following hundreds of accounts just to garner attention, it makes Twitter a less-nice place to hang out.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, that's fine! You can just visit the person's profile page and then un-follow them. Aggressive follow churn is when an account repeatedly follows and unfollows large numbers of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio.

These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Additional limits if you are following 2000 or more people:

The rules about aggressive following and follow churn still apply. In addition, every user can follow 2000 people total. Once you've followed 2000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following. When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—basically, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter monitors following behaviors and limits some behavior in order to make using Twitter a more positive experience for everyone.

Twitter facilitates social networking, but it's not a social networking website. In fact, the way Twitter works is quite different from social networks: when you accept friend requests on social networks, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different, because instead of indicating a mutual relationship, following is a one-way action that means you want to receive information, in the form of tweets, from someone. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring mutual following.

If you follow too many people, there's no way you can keep up with everyone's updates in your timeline. If you need to communicate with someone but don't need to see their updates everyday, you don't have to follow them. Send them an @reply when you need to; it doesn't require following and your update will appear in the person's replies tab, so they can reply back.

In addition, limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter, so monitoring for abuse is one way to reduce spam on Twitter.

Using Third Party Applications to “Get More Followers Fast!”

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. [The Twitter Rules](#) addresses the use of these programs by disallowing:

- Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. Please do not give your username and password out to any third-party application that you have not thoroughly researched.

Following limits? But I'm whitelisted!

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users, and degrades the Twitter experience for everyone.

Automation Rules & Best Practices (2010/01)

Submitted Nov 08, 2009 by ginger

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password.

Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

The Twitter Rules (20110412*)

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor user's content and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

- **Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others
- **Trademark:** We reserve the right to reclaim user names on behalf of businesses or individuals that hold legal claim or trademark on those user names. Accounts using business names and/or logos to mislead others will be permanently suspended.
- **Privacy:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission.
- **Violence and Threats:** You may not publish or post direct, specific threats of violence against others.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Unlawful Use:** You may not use our service for any unlawful purposes or for promotion of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter Badges:** You may not use a Verified Account badge or Promoted Products badge unless it is provided by Twitter. Accounts using these badges as part of profile pictures, background images, or in a way that falsely implies affiliation with Twitter will be suspended.

Spam and Abuse

Twitter strives to protect its users from spam and abuse. Technical abuse and user abuse is not tolerated on Twitter.com, and will result in permanent suspension. Any accounts engaging in the activities specified below are subject to permanent suspension.

- **Serial Accounts:** You may not create serial accounts for disruptive or abusive purposes, or with overlapping use cases. Mass account creation may result in suspension of all related accounts. Please note that any violation of the Twitter Rules is cause for permanent suspension of all accounts.
- **Username Squatting:** You may not engage in username squatting. Accounts that are inactive for more than 6 months may also be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be username squatting are:
 - the number of accounts created
 - creating accounts for the purpose of preventing others from using those account names
 - creating accounts for the purpose of selling those accounts
 - using feeds of third-party content to update and maintain accounts under the names of those third parties
- **Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations.
- **Selling user names:** Unless you have been specifically permitted to do so in a separate agreement with Twitter, you agree that you will not reproduce, duplicate, copy, sell, trade or resell the Services for any purpose, where "Services" is defined as follows: Your use of Twitter's products, services and web sites (referred to collectively as the "Services" in this document and excluding any services provided to you by Twitter under a separate written agreement) is subject to the terms of a legal agreement between you and Twitter.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.

- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - If you have followed a large amount of users in a short amount of time;
 - If you have followed and unfollowed people in a short time period, particularly by automated means (aggressive follower churn);
 - If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - If you have a small number of followers compared to the amount of people you are following;
 - If your updates consist mainly of links, and not personal updates;
 - If you post misleading links;
 - If a large number of people are blocking you;
 - The number of spam complaints that have been filed against you;
 - If you post duplicate content over multiple accounts or multiple duplicate updates on one account;
 - If you post multiple unrelated updates to a topic using #;
 - If you post multiple unrelated updates to a trending or popular topic;
 - If you send large numbers of duplicate @replies or mentions;
 - If you send large numbers of unsolicited @replies or mentions in an attempt to spam a service or link;
 - If you add a large number of unrelated users to lists in an attempt to spam a service or link;
 - If you repeatedly post other users' Tweets as your own;
 - If you have attempted to "sell" followers, particularly through tactics considered aggressive following or follower churn;
 - Creating or purchasing accounts in order to gain followers;
 - Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising "more followers fast," or any other site that offers to automatically add followers to your account);
 - If you create false or misleading Points of Interest;
 - If you create Points of Interest to namesquat or spam.
- **Pornography:** You may not use obscene or pornographic images in either your profile picture or user background

Your account may be suspended for Terms of Service violations if any of the above is true. Please see our help pages on [Following Best Practices](#) and [Automation Rules and Best Practices](#) for a more detailed discussion of how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the Terms of Service.

Following Rules and Best Practices (20110412*)

Following is a way to receive information, in the form of tweets, from a person, company, or organization. If an account seems interesting, you can follow that account and see their updates in your timeline. This page discusses some of the limits and best practices regarding following on Twitter.

What are the limits?

We don't limit the number of followers you can have. However, we do monitor how aggressively users follow other users. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine. Please note that the only automated following behavior that Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. Please review our [Automation Rules and Best Practices](#) for more information on automating your account.

What is aggressive following?

Following users is a way to see their updates in your timeline. If an account seems interesting, feel free to follow it! Many Twitter users receive email notifications when someone follows their account or check out the profiles of new followers to see if they share common interests. If some accounts are aggressively or indiscriminately following hundreds of accounts just to garner attention, it makes Twitter a less-nice place to hang out.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, that's fine! You can just visit the person's profile page and then un-follow them. Aggressive follow churn is when an account repeatedly follows and un-follows large numbers of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio.

These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Additional limits if you are following 2000 or more people:

The rules about aggressive following and follow churn still apply. In addition, every user can follow 2000 people total. Once you've followed 2000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following. When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—basically, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter facilitates social networking, but it's not a social networking website. In fact, Twitter works quite differently from social networks: when you accept friend requests on social networks, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different, because instead of indicating a mutual relationship, following is a one-way action that means you want to receive information, in the form of tweets, from someone. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring mutual following.

If you follow too many people, there's no way you can keep up with everyone's updates in your timeline. If you need to communicate with someone but don't need to see their updates everyday, you don't have to follow them. Send them an @reply when you need to; it doesn't require following and your update will appear in the person's replies tab, so they can reply back.

In addition, limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter, so monitoring for abuse is one way to reduce spam on Twitter.

Using Third Party Applications to “Get More Followers Fast!”

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. [The Twitter Rules](#) addresses the use of these programs by disallowing:

- Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising 'more followers fast,' or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. Please do not give your username and password out to any third-party application that you have not thoroughly researched.

Revoke access for any third-party application that you don't recognize by visiting the [Connections tab](#) in Account Settings.

Following limits? But I'm whitelisted!

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users, and degrades the Twitter experience for everyone.

Automation Rules & Best Practices (2010/01)

Submitted Nov 08, 2009 by [ginger](#)

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account

when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to [unrelated](#) trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#).

Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

The Twitter Rules (20110902)

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor user's content and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

- **Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others
- **Trademark:** We reserve the right to reclaim user names on behalf of businesses or individuals that hold legal claim or trademark on those user names. Accounts using business names and/or logos to mislead others will be permanently suspended.
- **Privacy:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission.
- **Violence and Threats:** You may not publish or post direct, specific threats of violence against others.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Unlawful Use:** You may not use our service for any unlawful purposes or for promotion of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter Badges:** You may not use a Verified Account badge or Promoted Products badge unless it is provided by Twitter. Accounts using these badges as part of profile pictures, background images, or in a way that falsely implies affiliation with Twitter will be suspended.

Spam and Abuse

Twitter strives to protect its users from spam and abuse. Technical abuse and user abuse is not tolerated on Twitter.com, and will result in permanent suspension. Any accounts engaging in the activities specified below are subject to permanent suspension.

- **Serial Accounts:** You may not create serial accounts for disruptive or abusive purposes, or with overlapping use cases. Mass account creation may result in suspension of all related accounts. Please note that any violation of the Twitter Rules is cause for permanent suspension of all accounts.
- **Username Squatting:** You may not engage in username squatting. Accounts that are inactive for more than 6 months may also be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be username squatting are:
 - the number of accounts created
 - creating accounts for the purpose of preventing others from using those account names
 - creating accounts for the purpose of selling those accounts
 - using feeds of third-party content to update and maintain accounts under the names of those third parties
- **Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations.
- **Selling user names:** Unless you have been specifically permitted to do so in a separate agreement with Twitter, you agree that you will not reproduce, duplicate, copy, sell, trade or resell the Services for any purpose, where "Services" is defined as follows: Your use of Twitter's products, services and web sites (referred to collectively as the "Services" in this document and excluding any services provided to

you by Twitter under a separate written agreement) is subject to the terms of a legal agreement between you and Twitter.

- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - If you have followed a large amount of users in a short amount of time;
 - If you have followed and unfollowed people in a short time period, particularly by automated means (aggressive follower churn);
 - If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - If you have a small number of followers compared to the amount of people you are following;
 - If your updates consist mainly of links, and not personal updates;
 - If you post misleading links;
 - If a large number of people are blocking you;
 - The number of spam complaints that have been filed against you;
 - If you post duplicate content over multiple accounts or multiple duplicate updates on one account;
 - If you post multiple unrelated updates to a topic using #;
 - If you post multiple unrelated updates to a trending or popular topic;
 - If you send large numbers of duplicate @replies or mentions;
 - If you send large numbers of unsolicited @replies or mentions in an attempt to spam a service or link;
 - If you add a large number of unrelated users to lists in an attempt to spam a service or link;
 - If you repeatedly post other users' Tweets as your own;
 - If you have attempted to "sell" followers, particularly through tactics considered aggressive following or follower churn;
 - Creating or purchasing accounts in order to gain followers;
 - Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising "more followers fast," or any other site that offers to automatically add followers to your account);
 - If you create false or misleading Points of Interest;
 - If you create Points of Interest to namesquat or spam.
- **Pornography:** You may not use obscene or pornographic images in either your profile picture or user background

Your account may be suspended for Terms of Service violations if any of the above is true. Please see our help pages on [Following Best Practices](#) and [Automation Rules and Best Practices](#) for a more detailed discussion of how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the Terms of Service. Having Trouble?

Check out our complete list of articles outlining our [policies, guidelines, and best practices](#).

Following Rules and Best Practices (20110412*)

Following is a way to receive information, in the form of tweets, from a person, company, or organization. If an account seems interesting, you can follow that account and see their updates in your timeline. This page discusses some of the limits and best practices regarding following on Twitter.

What are the limits?

We don't limit the number of followers you can have. However, we do monitor how aggressively users follow other users. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine. Please note that the only automated following behavior that Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. Please review our [Automation Rules and Best Practices](#) for more information on automating your account.

What is aggressive following?

Following users is a way to see their updates in your timeline. If an account seems interesting, feel free to follow it! Many Twitter users receive email notifications when someone follows their account or check out the profiles of new followers to see if they share common interests. If some accounts are aggressively or indiscriminately following hundreds of accounts just to garner attention, it makes Twitter a less-nice place to hang out.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, that's fine! You can just visit the person's profile page and then un-follow them. Aggressive follow churn is when an account repeatedly follows and un-follows large numbers of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio.

These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Additional limits if you are following 2000 or more people:

The rules about aggressive following and follow churn still apply. In addition, every user can follow 2000 people total. Once you've followed 2000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following. When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—basically, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error

message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter facilitates social networking, but it's not a social networking website. In fact, Twitter works quite differently from social networks: when you accept friend requests on social networks, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different, because instead of indicating a mutual relationship, following is a one-way action that means you want to receive information, in the form of tweets, from someone. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring mutual following.

If you follow too many people, there's no way you can keep up with everyone's updates in your timeline. If you need to communicate with someone but don't need to see their updates everyday, you don't have to follow them. Send them an @reply when you need to; it doesn't require following and your update will appear in the person's replies tab, so they can reply back.

In addition, limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter, so monitoring for abuse is one way to reduce spam on Twitter.

Using Third Party Applications to “Get More Followers Fast!”

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. [The Twitter Rules](#) addresses the use of these programs by disallowing:

- Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. Please do not give your username and password out to any third-party application that you have not thoroughly researched.

Revoke access for any third-party application that you don't recognize by visiting the [Connections tab](#) in Account Settings.

Following limits? But I'm whitelisted!

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users, and degrades the Twitter experience for everyone.

Automation Rules & Best Practices (2010/01)

Submitted Nov 08, 2009 by [ginger](#)

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being

filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate

offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the time-stamp, like this:

The Twitter Rules (20110903)

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor user's content and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

- **Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others
- **Trademark:** We reserve the right to reclaim user names on behalf of businesses or individuals that hold legal claim or trademark on those user names. Accounts using business names and/or logos to mislead others will be permanently suspended.
- **Privacy:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission.
- **Violence and Threats:** You may not publish or post direct, specific threats of violence against others.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Unlawful Use:** You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter Badges:** You may not use a Verified Account badge or Promoted Products badge unless it is provided by Twitter. Accounts using these badges as part of profile pictures, background images, or in a way that falsely implies affiliation with Twitter will be suspended.

Spam and Abuse

Twitter strives to protect its users from spam and abuse. Technical abuse and user abuse is not tolerated on Twitter.com, and will result in permanent suspension. Any accounts engaging in the activities specified below are subject to permanent suspension.

- **Serial Accounts:** You may not create serial accounts for disruptive or abusive purposes, or with overlapping use cases. Mass account creation may result in suspension of all related accounts. Please note that any violation of the Twitter Rules is cause for permanent suspension of all accounts.
- **Username Squatting:** You may not engage in username squatting. Accounts that are inactive for more than 6 months may also be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be username squatting are:

- the number of accounts created
- creating accounts for the purpose of preventing others from using those account names
- creating accounts for the purpose of selling those accounts
- using feeds of third-party content to update and maintain accounts under the names of those third parties
- **Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations.
- **Selling user names:** Unless you have been specifically permitted to do so in a separate agreement with Twitter, you agree that you will not reproduce, duplicate, copy, sell, trade or resell the Services for any purpose, where "Services" is defined as follows: Your use of Twitter's products, services and web sites (referred to collectively as the "Services" in this document and excluding any services provided to you by Twitter under a separate written agreement) is subject to the terms of a legal agreement between you and Twitter.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - If you have followed a large amount of users in a short amount of time;
 - If you have followed and unfollowed people in a short time period, particularly by automated means (aggressive follower churn);
 - If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - If you have a small number of followers compared to the amount of people you are following;
 - If your updates consist mainly of links, and not personal updates;
 - If you post misleading links;
 - If a large number of people are blocking you;
 - The number of spam complaints that have been filed against you;
 - If you post duplicate content over multiple accounts or multiple duplicate updates on one account;
 - If you post multiple unrelated updates to a topic using #;
 - If you post multiple unrelated updates to a trending or popular topic;
 - If you send large numbers of duplicate @replies or mentions;
 - If you send large numbers of unsolicited @replies or mentions in an attempt to spam a service or link;
 - If you add a large number of unrelated users to lists in an attempt to spam a service or link;
 - If you repeatedly post other users' Tweets as your own;
 - If you have attempted to "sell" followers, particularly through tactics considered aggressive following or follower churn;
 - Creating or purchasing accounts in order to gain followers;

- Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising "more followers fast," or any other site that offers to automatically add followers to your account);
- If you create false or misleading Points of Interest;
- If you create Points of Interest to namesquat or spam.
- **Pornography:** You may not use obscene or pornographic images in either your profile picture or user background

Your account may be suspended for Terms of Service violations if any of the above is true. Please see our help pages on [Following Best Practices](#) and [Automation Rules and Best Practices](#) for a more detailed discussion of how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the Terms of Service.

Having Trouble?

Check out our complete list of articles outlining our [policies, guidelines, and best practices](#).

Following Rules and Best Practices (20110412*)

Following is a way to receive information, in the form of tweets, from a person, company, or organization. If an account seems interesting, you can follow that account and see their updates in your timeline. This page discusses some of the limits and best practices regarding following on Twitter.

What are the limits?

We don't limit the number of followers you can have. However, we do monitor how aggressively users follow other users. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine. Please note that the only automated following behavior that Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. Please review our [Automation Rules and Best Practices](#) for more information on automating your account.

What is aggressive following?

Following users is a way to see their updates in your timeline. If an account seems interesting, feel free to follow it! Many Twitter users receive email notifications when someone follows their account or check out the profiles of new followers to see if they share common interests. If some accounts are aggressively or indiscriminately following hundreds of accounts just to garner attention, it makes Twitter a less-nice place to hang out.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, that's fine! You can just visit the person's profile page and then un-follow them. Aggressive follow churn is when an account repeatedly follows and un-follows large numbers of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio.

These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Additional limits if you are following 2000 or more people:

The rules about aggressive following and follow churn still apply. In addition, every user can follow 2000 people total. Once you've followed 2000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following. When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—basically, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter facilitates social networking, but it's not a social networking website. In fact, Twitter works quite differently from social networks: when you accept friend requests on social networks, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different, because instead of indicating a mutual relationship, following is a one-way action that means you want to receive information, in the form of tweets, from someone. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring mutual following.

If you follow too many people, there's no way you can keep up with everyone's updates in your timeline. If you need to communicate with someone but don't need to see their updates everyday, you don't have to follow them. Send them an @reply when you need to; it doesn't require following and your update will appear in the person's replies tab, so they can reply back.

In addition, limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter, so monitoring for abuse is one way to reduce spam on Twitter.

Using Third Party Applications to “Get More Followers Fast!”

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. [The Twitter Rules](#) addresses the use of these programs by disallowing:

- Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. Please do not give your username and password out to any third-party application that you have not thoroughly researched.

Revoke access for any third-party application that you don't recognize by visiting the [Connections tab](#) in Account Settings.

Following limits? But I'm whitelisted!

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users, and degrades the Twitter experience for everyone.

[Automation Rules & Best Practices \(2010/01\)](#)

Submitted Nov 08, 2009 by [ginger](#)

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results. Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

The Twitter Rules (20120510)

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor user's content and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

- **Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others
- **Trademark:** We reserve the right to reclaim user names on behalf of businesses or individuals that hold legal claim or trademark on those user names. Accounts using business names and/or logos to mislead others will be permanently suspended.
- **Privacy:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission.
- **Violence and Threats:** You may not publish or post direct, specific threats of violence against others.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Unlawful Use:** You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter Badges:** You may not use a Verified Account badge or Promoted Products badge unless it is provided by Twitter. Accounts using these badges as part of profile pictures, background images, or in a way that falsely implies affiliation with Twitter will be suspended.

Spam and Abuse

Twitter strives to protect its users from spam and abuse. Technical abuse and user abuse is not tolerated on Twitter.com, and will result in permanent suspension. Any accounts engaging in the activities specified below are subject to permanent suspension.

- **Serial Accounts:** You may not create serial accounts for disruptive or abusive purposes, or with overlapping use cases. Mass account creation may result in suspension of all related accounts. Please note that any violation of the Twitter Rules is cause for permanent suspension of all accounts.
- **Username Squatting:** You may not engage in username squatting. Accounts that are inactive for more than 6 months may also be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be username squatting are:

- the number of accounts created
- creating accounts for the purpose of preventing others from using those account names
- creating accounts for the purpose of selling those accounts
- using feeds of third-party content to update and maintain accounts under the names of those third parties
- **Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations.
- **Selling user names:** You may not buy or sell Twitter usernames.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - If you have followed a large amount of users in a short amount of time;
 - If you have followed and unfollowed people in a short time period, particularly by automated means (aggressive follower churn);
 - If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - If you have a small number of followers compared to the amount of people you are following;
 - If your updates consist mainly of links, and not personal updates;
 - If you post misleading links;
 - If a large number of people are blocking you;
 - The number of spam complaints that have been filed against you;
 - If you post duplicate content over multiple accounts or multiple duplicate updates on one account;
 - If you post multiple unrelated updates to a topic using #;
 - If you post multiple unrelated updates to a trending or popular topic;
 - If you send large numbers of duplicate @replies or mentions;
 - If you send large numbers of unsolicited @replies or mentions in an attempt to spam a service or link;
 - If you add a large number of unrelated users to lists in an attempt to spam a service or link;
 - If you repeatedly post other users' Tweets as your own;
 - If you have attempted to "sell" followers, particularly through tactics considered aggressive following or follower churn;
 - Creating or purchasing accounts in order to gain followers;
 - Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising "more followers fast," or any other site that offers to automatically add followers to your account);
 - If you create false or misleading Points of Interest;
 - If you create Points of Interest to namesquat or spam.
- **Pornography:** You may not use obscene or pornographic images in either your profile picture or user background

Your account may be suspended for Terms of Service violations if any of the above is true. Please see our help pages on [Following Best Practices](#) and [Automation Rules and Best Practices](#) for a more detailed discussion of how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the Terms of Service.

Having Trouble?

Check out our complete list of articles outlining our [policies, guidelines, and best practices](#).

Following Rules and Best Practices (20110412*)

Following is a way to receive information, in the form of tweets, from a person, company, or organization. If an account seems interesting, you can follow that account and see their updates in your timeline. This page discusses some of the limits and best practices regarding following on Twitter.

What are the limits?

We don't limit the number of followers you can have. However, we do monitor how aggressively users follow other users. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine. Please note that the only automated following behavior that Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. Please review our [Automation Rules and Best Practices](#) for more information on automating your account.

What is aggressive following?

Following users is a way to see their updates in your timeline. If an account seems interesting, feel free to follow it! Many Twitter users receive email notifications when someone follows their account or check out the profiles of new followers to see if they share common interests. If some accounts are aggressively or indiscriminately following hundreds of accounts just to garner attention, it makes Twitter a less-nice place to hang out.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, that's fine! You can just visit the person's profile page and then un-follow them. Aggressive follow churn is when an account repeatedly follows and un-follows large numbers of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio.

These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Additional limits if you are following 2000 or more people:

The rules about aggressive following and follow churn still apply. In addition, every user can follow 2000 people total. Once you've followed 2000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following. When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—basically, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter facilitates social networking, but it's not a social networking website. In fact, Twitter works quite differently from social networks: when you accept friend requests on social networks, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different, because instead of indicating a mutual relationship, following is a one-way action that means you want to receive information, in the form of tweets, from someone. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring mutual following.

If you follow too many people, there's no way you can keep up with everyone's updates in your timeline. If you need to communicate with someone but don't need to see their updates everyday, you don't have to follow them. Send them an @reply when you need to; it doesn't require following and your update will appear in the person's replies tab, so they can reply back.

In addition, limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter, so monitoring for abuse is one way to reduce spam on Twitter.

Using Third Party Applications to “Get More Followers Fast!”

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. [The Twitter Rules](#) addresses the use of these programs by disallowing:

- Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. Please do not give your username and password out to any third-party application that you have not thoroughly researched.

Revoke access for any third-party application that you don't recognize by visiting the [Connections tab](#) in Account Settings.

Following limits? But I'm whitelisted!

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users, and degrades the Twitter experience for everyone.

Automation Rules & Best Practices (2010/01)

Submitted Nov 08, 2009 by [ginger](#)

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password.

Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

The Twitter Rules (20120919*)

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor user's content and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

- **Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others
- **Trademark:** We reserve the right to reclaim user names on behalf of businesses or individuals that hold legal claim or trademark on those user names. Accounts using business names and/or logos to mislead others will be permanently suspended.
- **Privacy:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission.
- **Violence and Threats:** You may not publish or post direct, specific threats of violence against others.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Unlawful Use:** You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter Badges:** You may not use a Verified Account badge or Promoted Products badge unless it is provided by Twitter. Accounts using these badges as part of profile pictures, background images, or in a way that falsely implies affiliation with Twitter will be suspended.

Spam and Abuse

Twitter strives to protect its users from spam and abuse. Technical abuse and user abuse is not tolerated on Twitter.com, and will result in permanent suspension. Any accounts engaging in the activities specified below are subject to permanent suspension.

- **Serial Accounts:** You may not create serial accounts for disruptive or abusive purposes, or with overlapping use cases. Mass account creation may result in suspension of all related accounts. Please note that any violation of the Twitter Rules is cause for permanent suspension of all accounts.
- **Username Squatting:** You may not engage in username squatting. Accounts that are inactive for more than 6 months may also be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be username squatting are:

- the number of accounts created
- creating accounts for the purpose of preventing others from using those account names
- creating accounts for the purpose of selling those accounts
- using feeds of third-party content to update and maintain accounts under the names of those third parties
- **Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations.
- **Selling user names:** You may not buy or sell Twitter usernames.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - If you have followed a large amount of users in a short amount of time;
 - If you have followed and unfollowed people in a short time period, particularly by automated means (aggressive follower churn);
 - If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - If you have a small number of followers compared to the amount of people you are following;
 - If your updates consist mainly of links, and not personal updates;
 - If you post misleading links;
 - If a large number of people are blocking you;
 - The number of spam complaints that have been filed against you;
 - If you post duplicate content over multiple accounts or multiple duplicate updates on one account;
 - If you post multiple unrelated updates to a topic using #;
 - If you post multiple unrelated updates to a trending or popular topic;
 - If you send large numbers of duplicate @replies or mentions;
 - If you send large numbers of unsolicited @replies or mentions in an attempt to spam a service or link;
 - If you add a large number of unrelated users to lists in an attempt to spam a service or link;
 - If you repeatedly post other users' Tweets as your own;
 - If you have attempted to "sell" followers, particularly through tactics considered aggressive following or follower churn;
 - Creating or purchasing accounts in order to gain followers;
 - Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising "more followers fast," or any other site that offers to automatically add followers to your account);
 - If you create false or misleading Points of Interest;
 - If you create Points of Interest to namesquat or spam.
- **Pornography:** You may not use obscene or pornographic images in either your profile picture or user background

Your account may be suspended for Terms of Service violations if any of the above is true. Please see our help pages on [Following Best Practices](#) and [Automation Rules and Best Practices](#) for a more detailed discussion of how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the Terms of Service.

Having Trouble?

Check out our complete list of articles outlining our [policies, guidelines, and best practices](#).

Following Rules and Best Practices (20110919*)

To learn more about following, please read "[FAQs About Following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that the only automated following behavior that Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. Please review our [Automation Rules and Best Practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and un-follow them. Aggressive follow churn is when an account *repeatedly* follows and then un-follows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 2000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 2000 people total**. Once you've followed 2000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows

people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

Tip: If you need to communicate with someone but don't need to see their updates every day, don't follow them. Visit their profile or send them an @reply when you need to; sending @replies doesn't require following, and your reply will appear in the person's @mentions tab so they can reply back.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

The [Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do**

not give your username and password out to any third-party application that you have not thoroughly researched.

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Having Trouble?

Check out our [Following Troubleshooting section](#) for solutions to common problems

Automation Rules & Best Practices (2010/01)

Submitted Nov 08, 2009 by [ginger](#)

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find

out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social

action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for

the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

The Twitter Rules (20121029)

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor user's content and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

- **Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others
- **Trademark:** We reserve the right to reclaim user names on behalf of businesses or individuals that hold legal claim or trademark on those user names. Accounts using business names and/or logos to mislead others will be permanently suspended.
- **Privacy:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission.
- **Violence and Threats:** You may not publish or post direct, specific threats of violence against others.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Unlawful Use:** You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter Badges:** You may not use a Verified Account badge or Promoted Products badge unless it is provided by Twitter. Accounts using these badges as part of profile photos, header photos, background images, or in a way that falsely implies affiliation with Twitter will be suspended.

Spam and Abuse

Twitter strives to protect its users from spam and abuse. Technical abuse and user abuse is not tolerated on Twitter.com, and will result in permanent suspension. Any accounts engaging in the activities specified below are subject to permanent suspension.

- **Serial Accounts:** You may not create serial accounts for disruptive or abusive purposes, or with overlapping use cases. Mass account creation may result in suspension of all related accounts. Please note that any violation of the Twitter Rules is cause for permanent suspension of all accounts.
- **Username Squatting:** You may not engage in username squatting. Accounts that are inactive for more than 6 months may also be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be username squatting are:
 - the number of accounts created

- creating accounts for the purpose of preventing others from using those account names
- creating accounts for the purpose of selling those accounts
- using feeds of third-party content to update and maintain accounts under the names of those third parties
- **Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations.
- **Selling user names:** You may not buy or sell Twitter usernames.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - If you have followed a large amount of users in a short amount of time;
 - If you have followed and unfollowed people in a short time period, particularly by automated means (aggressive follower churn);
 - If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - If you have a small number of followers compared to the amount of people you are following;
 - If your updates consist mainly of links, and not personal updates;
 - If you post misleading links;
 - If a large number of people are blocking you;
 - The number of spam complaints that have been filed against you;
 - If you post duplicate content over multiple accounts or multiple duplicate updates on one account;
 - If you post multiple unrelated updates to a topic using #;
 - If you post multiple unrelated updates to a trending or popular topic;
 - If you send large numbers of duplicate @replies or mentions;
 - If you send large numbers of unsolicited @replies or mentions in an attempt to spam a service or link;
 - If you add a large number of unrelated users to lists in an attempt to spam a service or link;
 - If you repeatedly post other users' Tweets as your own;
 - If you have attempted to "sell" followers, particularly through tactics considered aggressive following or follower churn;
 - Creating or purchasing accounts in order to gain followers;
 - Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising "more followers fast," or any other site that offers to automatically add followers to your account);
 - If you create false or misleading Points of Interest;
 - If you create Points of Interest to namesquat or spam.
- **Pornography:** You may not use obscene or pornographic images in either your profile photo, header photo, or user background.

Your account may be suspended for Terms of Service violations if any of the above is true. Please see our help pages on [Following Best Practices](#) and [Automation Rules and Best Practices](#) for a more detailed discussion of how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the Terms of Service.

Having Trouble?

Check out our complete list of articles outlining our [policies, guidelines, and best practices](#).

Following Rules and Best Practices (20110919*)

To learn more about following, please read "[FAQs About Following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that the only automated following behavior that Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. Please review our [Automation Rules and Best Practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and un-follow them. Aggressive follow churn is when an account *repeatedly* follows and then un-follows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 2000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 2000 people total**. Once you've followed 2000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows

people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

Tip: If you need to communicate with someone but don't need to see their updates every day, don't follow them. Visit their profile or send them an @reply when you need to; sending @replies doesn't require following, and your reply will appear in the person's @mentions tab so they can reply back.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

The [Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do**

not give your username and password out to any third-party application that you have not thoroughly researched.

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Having Trouble?

Check out our [Following Troubleshooting section](#) for solutions to common problems

Automation Rules & Best Practices (2010/01)

Submitted Nov 08, 2009 by [ginger](#)

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find

out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social

action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for

the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

The Twitter Rules (20130530)

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

- **Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others.
- **Trademark:** We reserve the right to reclaim usernames on behalf of businesses or individuals that hold legal claim or trademark on those usernames. Accounts using business names and/or logos to mislead others may be permanently suspended.
- **Private information:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission.
- **Violence and Threats:** You may not publish or post direct, specific threats of violence against others.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Unlawful Use:** You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter Badges:** You may not use badges, such as but not limited to the Promoted or Verified Twitter badge, unless provided by Twitter. Accounts using these badges as part of profile photos, header photos, background images, or in a way that falsely implies affiliation with Twitter may be suspended.

Spam and Abuse

Twitter strives to protect its users from spam and abuse. Technical abuse and user abuse is not tolerated on Twitter.com, and will result in permanent suspension. Any accounts engaging in the activities specified below are subject to permanent suspension.

- **Serial Accounts:** You may not create multiple accounts for disruptive or abusive purposes, or with overlapping use cases. Mass account creation may result in suspension of all related accounts. Please note that any violation of the Twitter Rules is cause for permanent suspension of all accounts.
- **Username Squatting:** You may not engage in username squatting. Accounts that are inactive for more than six months may also be removed without further notice. Some of

the factors that we take into account when determining what conduct is considered to be username squatting are:

- the number of accounts created
- creating accounts for the purpose of preventing others from using those account names
- creating accounts for the purpose of selling those accounts
- using feeds of third-party content to update and maintain accounts under the names of those third parties
- **Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations.
- **Selling usernames:** You may not buy or sell Twitter usernames.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - If you have followed and/or unfollowed large amounts of users in a short time period, particularly by automated means (aggressive following or follower churn);
 - If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - If your updates consist mainly of links, and not personal updates;
 - If a large number of people are blocking you;
 - If a large number of spam complaints have been filed against you;
 - If you post duplicate content over multiple accounts or multiple duplicate updates on one account;
 - If you post multiple unrelated updates to a topic using #, trending or popular topic, or promoted trend;
 - If you send large numbers of duplicate @replies or mentions;
 - If you send large numbers of unsolicited @replies or mentions in an aggressive attempt to bring attention to a service or link;
 - If you add a large number of unrelated users to lists in an attempt to bring attention to an account, service or link;
 - If you repeatedly create false or misleading content in an attempt to bring attention to an account, service or link;
 - Randomly or aggressively favoriting Tweets through automation in an attempt to bring attention to an account, service or link;
 - Randomly or aggressively Retweeting accounts through automation in an attempt to bring attention to an account, service or link;
 - If you repeatedly post other users' account information as your own (bio, Tweets, url, etc.);
 - If you post misleading links (e.g. affiliate links, links to malware/click jacking pages, etc.);
 - Creating multiple misleading accounts in order to gain followers;
 - Selling followers;

- Purchasing followers;
- Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising "more followers fast," or any other site that offers to automatically add followers to your account);
- **Pornography:** You may not use obscene or pornographic images in either your profile photo, header photo, or user background.

Your account may be suspended for Terms of Service violations if any of the above is true. Please see our help pages on [Following rules and best practices](#) and [Automation rules and best practices](#) for a more detailed discussion of how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the [Terms of Service](#).

Have Questions?

Check out our complete list of articles outlining our [policies, guidelines, and best practices](#). To report an account for violation of the Twitter Rules, [please use our forms](#).

Impersonation policy (201305*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter Impersonation Policy.

An account will not be removed if

- the user shares your name but has no other commonalities, or
- the profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same background or avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to [Twitter's Parody Policy](#) for more information about these types of accounts.

For Frequently Asked Questions about reporting impersonation on Twitter, [click here](#).

Trademark policy (20121003)

Updated October 3, 2012

What is a Trademark Policy Violation on Twitter?

Using a company or business name, logo, or other trademark-protected materials in a manner that may mislead or confuse others with regard to its brand or business affiliation may be considered a trademark policy violation.

How Does Twitter Respond To Reported Trademark Policy Violations?

When we receive reports of trademark policy violations from holders of federal or international trademark registrations, we review the account and may take the following actions:

- When there is a clear intent to mislead others through the unauthorized use of a trademark, Twitter will suspend the account and notify the account holder.
- When we determine that an account appears to be confusing users, but is not purposefully passing itself off as the trademarked good or service, we give the account holder an opportunity to clear up any potential confusion. We may also release a username for the trademark holder's active use.
- We are responsive to reports about confusing or misleading [Promoted Tweet](#) and [Promoted Trend](#) copy, as well as [Promoted Account](#) profile information. For more information, please visit our [Trademark Policy for Promoted Products](#).

What Is not a Trademark Policy Violation?

Using another's trademark in a way that has nothing to do with the product or service for which the trademark was granted is not a violation of Twitter's trademark policy.

- Twitter usernames are provided on a first-come, first-served basis and may not be reserved. For information on why you may not be able to select a certain username, please see our [Why Can't I Register Certain Usernames](#) help page.
- Search and interest keywords advertisers choose are subject to applicable trademark law. Although Twitter may review reports of confusing keywords, we may not disable keywords in every case. For more information about our policies specific to Promoted Products, please see our [Promoted Products Policy](#) and our [Trademark Policy for Promoted Products](#).

Guidelines for News feed, Commentary, and Fan Accounts

Twitter users are allowed to create news feed, commentary, and fan accounts. Twitter provides a platform for its users to share and receive a wide range of ideas and content, and we greatly value and respect our users' expression. Because of these principles, we do not actively monitor users' content and will not edit or remove user content, except in cases of violations of our [Terms of Service](#).

An account's profile information should make it clear that the account is not actually the company or business entity that is the subject of the news feed/commentary/fan account. Here are some suggestions for distinguishing your account:

- **Username:** The username should not be the trademarked name of the subject of the news feed, commentary, or fan account.
- **Name:** The profile name should not be the trademarked name of the company or include the trademarked name in a misleading manner.

- **Bio:** The bio should include a statement to distinguish it from the real company, such as "Unofficial Account," "Fan Account," or "Not affiliated with..."
- **Profile photo, header photo, or background image:** The account should not use another's trademark, logo or other copyright-protected image without express permission.
- **Communication with other users:** The account should not, through private or public communication with other users, try to deceive or mislead others about your identity. Users may also choose to use different language to indicate that an account is not associated with the actual brand/company/product so long as it is clear and not confusing to others, and does not mislead or deceive.

If an account is reported to be confusing, we may request that the account holder make further changes to bring the account in compliance with these best practices.

How Can I Make My Own Account's Brand or Business Affiliation Clear?

We strongly recommend that you use all of Twitter's account settings (account name, location, web, and bio) to make your account's affiliation clear.

- Please see our [How to Change Your Profile Picture or Information](#) page for instructions on customizing your account. In particular, we recommend clearly stating your location, including your website if you have one, and clearly describing your brand or business in the bio, if applicable.
- You can also consider linking your official website to your Twitter account. [Twitter Buttons](#) are an easy way to add Twitter actions like "Follow" and "Mention" directly to your website.

What Information is Required When Reporting Trademark Policy Violations?

In order to investigate trademark policy violations, please provide all of the following information:

Username of the reported account (e.g., @safety or <http://www.twitter.com/safety>):

Your company name:

Your company Twitter account (if there is one):

Company website:

Your trademarked word, symbol, etc. (e.g. Twitter):

Trademark registration number:

Trademark registration office (e.g., USPTO):

Note: A federal or international trademark registration number is required. If the name you are reporting is not a registered mark (e.g., a government agency or non-profit organization), please let us know:

Your first and last name:

Title:

Address:

Phone:

Fax:

Email (must be from company domain):

Description of confusion (e.g., passing off as your company, including specific descriptions of content or behavior):

Requested Action (e.g., removal of violating account or transfer of trademarked username to an existing company account):

How Do I Report a Trademark Policy Violation?

You do not need a Twitter account to submit a trademark report. Holders of registered trademarks can report possible violations to Twitter's Policy Team through our [support forms](#).

Please submit trademark-related requests from your company email address and follow the format above to help expedite our response. Also, be sure to clearly describe to us why the account or Tweets may cause confusion with your mark.

Private information posted on Twitter (20130523)

Posting another person's private and confidential information is a violation of the [Twitter Rules](#).

Some examples of private and confidential information are:

- credit card information
- social security or other national identity numbers
- addresses or locations that are considered and treated as private
- non-public, personal phone numbers
- non-public, personal email addresses

Keep in mind that although you may consider certain information to be private, not all postings of such information may be a violation of this policy. If information was previously posted or displayed elsewhere on the Internet prior to being put on Twitter, it is not a violation of this policy.

If the private information you're reporting is not your own, please note that you must have documentation that you're authorized to act on behalf of the person whose confidential information is posted.

For frequently asked questions about reporting private information posted on Twitter, [click here](#). To learn more about protecting your private information on Twitter and other websites, [click here](#).

Abusive behavior policy (20130524*)

User disputes and false statements

Twitter is a platform that provides a global communication service which encompasses a variety of users with different voices, ideas and perspectives. As a policy, we do not mediate content or intervene in disputes between users. However, targeted abuse or harassment may constitute a violation of the [Twitter Rules](#) and [Terms of Service](#).

Offensive content

Users are allowed to post content, including potentially inflammatory content, provided they do not violate the [Twitter Terms of Service and Rules](#). Twitter does not screen content and we do not remove potentially offensive content unless such content is a violation of our [Terms of Service](#).

If you believe the content or behavior you are reporting is prohibited in your local jurisdiction, please contact your local authorities so they can accurately assess the content or behavior for possible violations of local law. If Twitter is contacted directly by law enforcement, we can work with them and provide assistance for their investigation as well as guidance around possible options. You can point local law enforcement to our [Law Enforcement Guidelines](#).

For Frequently Asked Questions about reporting abusive behavior on Twitter, [click here](#).

To learn more about what you can do when you encounter abusive behavior on Twitter and other websites, [click here](#).

Copyright and DMCA policy (20130430*)

Topics Covered:

- [What is the DMCA ?](#)
- [What Types of Copyright Complaints Does Twitter Respond To?](#)
- [Am I a Copyright Holder? How Do I Know?](#)
- [What Information Do You Need To Process a Copyright Complaint?](#)
- [How Do I File a Copyright Complaint?](#)
- [How Are Claims Processed?](#)
- [What Happens Next?](#)
- [What Happens If I Receive a DMCA Notification?](#)
- [When Should I File a Counter-Notice?](#)
- [What Information Do You Need to Process a Counter-notice?](#)
- [What Happens After I Submit a Counter-notice?](#)
- [Filing a DMCA Notification or Counter-notice Is Serious Business!](#)

What is the DMCA?

[Section 512](#) of the Digital Millennium Copyright Act (“DMCA”) outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

What Types of Copyright Complaints Does Twitter Respond To?

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile photo, header photo, or background, allegations concerning the unauthorized use of a copyrighted image uploaded through our photo hosting service,

or Tweets containing links to allegedly infringing materials.

Tip: If you are concerned about the use of your brand or entity's name, please review Twitter's [Trademark Policy](#). If you are concerned about the use of a fictional character, please see our [Parody, Commentary, and Fan Accounts Policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <http://chillingeffects.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

Tip: If you are reporting the content of a Tweet, please give us a direct link to that Tweet following these instructions: <https://support.twitter.com/articles/80586>. Or please specify if the alleged infringement is in the background, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to Twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate DMCA notifications may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with access to the reporter's complaint along with instructions on how to file a counter-notice, and forward a copy of the complaint to [Chilling Effects](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a copy of the takedown notice, along with instructions for filing a counter-notification.

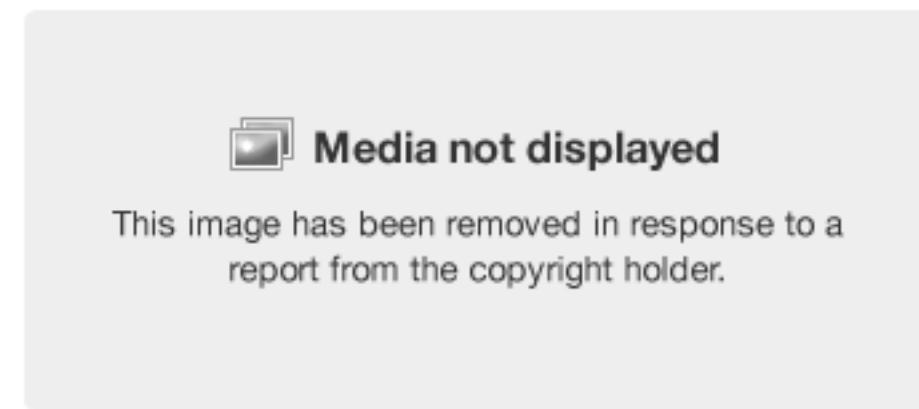
Tip: If you've not yet received a copy of the DMCA notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a copy of each DMCA notification and counter-notice that we process to [Chilling Effects](#), where they are posted to a public-facing website (with your personal information removed).

Withheld Tweet:



Withheld Media:



Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a DMCA Notification?

If you receive a DMCA notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to

read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

If you believe that the material reported in the DMCA notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Tip: Re-posting material removed in response to a DMCA notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the DMCA notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

Tip: We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a DMCA Notification or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Following Rules and Best Practices (20110919*)

To learn more about following, please read "[FAQs About Following.](#)"

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that the only automated following behavior that Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. Please review our [Automation Rules and Best Practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and un-follow them. Aggressive follow churn is when an account *repeatedly* follows and then un-follows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 2000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 2000 people total**. Once you've followed 2000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

Tip: If you need to communicate with someone but don't need to see their updates every day, don't follow them. Visit their profile or send them an @reply when you need to; sending @replies doesn't require following, and your reply will appear in the person's @mentions tab so they can reply back.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

The [Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third- party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third- party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Having Trouble?

Check out our [Following Troubleshooting](#) section for solutions to common problems

Automation Rules & Best Practices (2010/01)

Submitted Nov 08, 2009 by [ginger](#)

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted

automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#).

Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

The Twitter Rules (20130813)

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

- **Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others.
- **Trademark:** We reserve the right to reclaim usernames on behalf of businesses or individuals that hold legal claim or trademark on those usernames. Accounts using business names and/or logos to mislead others may be permanently suspended.
- **Private information:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission.
- **Violence and Threats:** You may not publish or post direct, specific threats of violence against others.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Unlawful Use:** You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter Badges:** You may not use badges, such as but not limited to the Promoted or Verified Twitter badge, unless provided by Twitter. Accounts using these badges as part of profile photos, header photos, background images, or in a way that falsely implies affiliation with Twitter may be suspended.

Abuse and Spam

Twitter strives to protect its users from abuse and spam. User abuse and technical abuse are not tolerated on Twitter.com, and may result in permanent suspension. Any accounts engaging in the activities specified below may be subject to permanent suspension.

- **Serial Accounts:** You may not create multiple accounts for disruptive or abusive purposes, or with overlapping use cases. Mass account creation may result in suspension

of all related accounts. Please note that any violation of the Twitter Rules is cause for permanent suspension of all accounts.

- **Targeted Abuse:** You may not engage in targeted abuse or harassment. Some of the factors that we take into account when determining what conduct is considered to be targeted abuse or harassment are:
 - if you are sending messages to a user from multiple accounts;
 - if the sole purpose of your account is to send abusive messages to others;
 - if the reported behavior is one-sided or includes threats
- **Username Squatting:** You may not engage in username squatting. Accounts that are inactive for more than six months may also be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be username squatting are:
 - the number of accounts created
 - creating accounts for the purpose of preventing others from using those account names
 - creating accounts for the purpose of selling those accounts
 - using feeds of third-party content to update and maintain accounts under the names of those third parties
- **Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations.
- **Selling usernames:** You may not buy or sell Twitter usernames.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - If you have followed and/or unfollowed large amounts of users in a short time period, particularly by automated means (aggressive following or follower churn);
 - If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - If your updates consist mainly of links, and not personal updates;
 - If a large number of people are blocking you;
 - If a large number of spam complaints have been filed against you;
 - If you post duplicate content over multiple accounts or multiple duplicate updates on one account;

- If you post multiple unrelated updates to a topic using #, trending or popular topic, or promoted trend;
- If you send large numbers of duplicate @replies or mentions;
- If you send large numbers of unsolicited @replies or mentions in an aggressive attempt to bring attention to a service or link;
- If you add a large number of unrelated users to lists in an attempt to bring attention to an account, service or link;
- If you repeatedly create false or misleading content in an attempt to bring attention to an account, service or link;
- Randomly or aggressively favoriting Tweets through automation in an attempt to bring attention to an account, service or link;
- Randomly or aggressively Retweeting accounts through automation in an attempt to bring attention to an account, service or link;
- If you repeatedly post other users' account information as your own (bio, Tweets, url, etc.);
- If you post misleading links (e.g. affiliate links, links to malware/click jacking pages, etc.);
- Creating multiple misleading accounts in order to gain followers;
- Selling followers;
- Purchasing followers;
- Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising "more followers fast," or any other site that offers to automatically add followers to your account);
- **Pornography:** You may not use obscene or pornographic images in either your profile photo, header photo, or user background.

Your account may be suspended for Terms of Service violations if any of the above is true. Please see our help pages on [Following rules and best practices](#) and [Automation rules and best practices](#) for a more detailed discussion of how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the [Terms of Service](#).

We may revise these Rules from time to time; the most current version will always be at twitter.com/rules.

Have Questions?

Check out our complete list of articles outlining our [policies, guidelines, and best practices](#).

To report an account for violation of the Twitter Rules, [please use our forms](#).

Impersonation policy (201305*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter Impersonation Policy.

An account will not be removed if

- the user shares your name but has no other commonalities, or
- the profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same background or avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to [Twitter's Parody Policy](#) for more information about these types of accounts.

For Frequently Asked Questions about reporting impersonation on Twitter, [click here](#).

Trademark policy (20121003)

Updated October 3, 2012

What is a Trademark Policy Violation on Twitter?

Using a company or business name, logo, or other trademark-protected materials in a manner that may mislead or confuse others with regard to its brand or business affiliation may be considered a trademark policy violation.

How Does Twitter Respond To Reported Trademark Policy Violations?

When we receive reports of trademark policy violations from holders of federal or international trademark registrations, we review the account and may take the following actions:

- When there is a clear intent to mislead others through the unauthorized use of a trademark, Twitter will suspend the account and notify the account holder.
- When we determine that an account appears to be confusing users, but is not purposefully passing itself off as the trademarked good or service, we give the account holder an opportunity to clear up any potential confusion. We may also release a username for the trademark holder's active use.
- We are responsive to reports about confusing or misleading [Promoted Tweet](#) and [Promoted Trend](#) copy, as well as [Promoted Account](#) profile information. For more information, please visit our [Trademark Policy for Promoted Products](#).

What Is not a Trademark Policy Violation?

Using another's trademark in a way that has nothing to do with the product or service for which the trademark was granted is not a violation of Twitter's trademark policy.

- Twitter usernames are provided on a first-come, first-served basis and may not be reserved. For information on why you may not be able to select a certain username, please see our [Why Can't I Register Certain Usernames](#) help page.
- Search and interest keywords advertisers choose are subject to applicable trademark law. Although Twitter may review reports of confusing keywords, we may not disable keywords in every case. For more information about our policies specific to Promoted Products, please see our [Promoted Products Policy](#) and our [Trademark Policy for Promoted Products](#).

Guidelines for News feed, Commentary, and Fan Accounts

Twitter users are allowed to create news feed, commentary, and fan accounts. Twitter provides a platform for its users to share and receive a wide range of ideas and content, and we greatly value and respect our users' expression. Because of these principles, we do not actively monitor users' content and will not edit or remove user content, except in cases of violations of our [Terms of Service](#).

An account's profile information should make it clear that the account is not actually the company or business entity that is the subject of the news feed/commentary/fan account. Here are some suggestions for distinguishing your account:

- **Username:** The username should not be the trademarked name of the subject of the news feed, commentary, or fan account.
- **Name:** The profile name should not be the trademarked name of the company or include the trademarked name in a misleading manner.
- **Bio:** The bio should include a statement to distinguish it from the real company, such as "Unofficial Account," "Fan Account," or "Not affiliated with..."
- **Profile photo, header photo, or background image:** The account should not use another's trademark, logo or other copyright-protected image without express permission.
- **Communication with other users:** The account should not, through private or public communication with other users, try to deceive or mislead others about your identity.

Users may also choose to use different language to indicate that an account is not associated with the actual brand/company/product so long as it is clear and not confusing to others, and does not mislead or deceive.

If an account is reported to be confusing, we may request that the account holder make further changes to bring the account in compliance with these best practices.

How Can I Make My Own Account's Brand or Business Affiliation Clear?

We strongly recommend that you use all of Twitter's account settings (account name, location, web, and bio) to make your account's affiliation clear.

- Please see our [How to Change Your Profile Picture or Information](#) page for instructions on customizing your account. In particular, we recommend clearly stating your location, including your website if you have one, and clearly describing your brand or business in the bio, if applicable.
- You can also consider linking your official website to your Twitter account. [Twitter Buttons](#) are an easy way to add Twitter actions like “Follow” and “Mention” directly to your website.

What Information is Required When Reporting Trademark Policy Violations?

In order to investigate trademark policy violations, please provide all of the following information:

Username of the reported account (e.g., @safety or <http://www.twitter.com/safety>):

Your company name:

Your company Twitter account (if there is one):

Company website:

Your trademarked word, symbol, etc. (e.g. Twitter):

Trademark registration number:

Trademark registration office (e.g., USPTO):

Note: A federal or international trademark registration number is required. If the name you are reporting is not a registered mark (e.g., a government agency or non-profit organization), please let us know:

Your first and last name:

Title:

Address:

Phone:

Fax:

Email (must be from company domain):

Description of confusion (e.g., passing off as your company, including specific descriptions of content or behavior):

Requested Action (e.g., removal of violating account or transfer of trademarked username to an existing company account):

How Do I Report a Trademark Policy Violation?

You do not need a Twitter account to submit a trademark report. Holders of registered trademarks can report possible violations to Twitter's Policy Team through our [support forms](#).

Please submit trademark-related requests from your company email address and follow the format above to help expedite our response. Also, be sure to clearly describe to us why the account or Tweets may cause confusion with your mark.

Private information posted on Twitter (20130523)

Posting another person's private and confidential information is a violation of the [Twitter Rules](#).

Some examples of private and confidential information are:

- credit card information
- social security or other national identity numbers
- addresses or locations that are considered and treated as private
- non-public, personal phone numbers
- non-public, personal email addresses

Keep in mind that although you may consider certain information to be private, not all postings of such information may be a violation of this policy. If information was previously posted or displayed elsewhere on the Internet prior to being put on Twitter, it is not a violation of this policy.

If the private information you're reporting is not your own, please note that you must have documentation that you're authorized to act on behalf of the person whose confidential information is posted.

For frequently asked questions about reporting private information posted on Twitter, [click here](#). To learn more about protecting your private information on Twitter and other websites, [click here](#).

Abusive behavior policy (20130524*)

User disputes and false statements

Twitter is a platform that provides a global communication service which encompasses a variety of users with different voices, ideas and perspectives. As a policy, we do not mediate content or intervene in disputes between users. However, targeted abuse or harassment may constitute a violation of the [Twitter Rules](#) and [Terms of Service](#).

Offensive content

Users are allowed to post content, including potentially inflammatory content, provided they do not violate the [Twitter Terms of Service and Rules](#). Twitter does not screen content and we do not remove potentially offensive content unless such content is a violation of our [Terms of Service](#).

If you believe the content or behavior you are reporting is prohibited in your local jurisdiction, please contact your local authorities so they can accurately assess the content or behavior for possible violations of local law. If Twitter is contacted directly by law enforcement, we can work with them and provide assistance for their investigation as well as guidance around possible options. You can point local law enforcement to our [Law Enforcement Guidelines](#).

For Frequently Asked Questions about reporting abusive behavior on Twitter, [click here](#).

To learn more about what you can do when you encounter abusive behavior on Twitter and other websites, [click here](#).

Copyright and DMCA policy (20130430*)

Topics Covered:

- [What is the DMCA ?](#)
- [What Types of Copyright Complaints Does Twitter Respond To?](#)
- [Am I a Copyright Holder? How Do I Know?](#)
- [What Information Do You Need To Process a Copyright Complaint?](#)
- [How Do I File a Copyright Complaint?](#)
- [How Are Claims Processed?](#)
- [What Happens Next?](#)
- [What Happens If I Receive a DMCA Notification?](#)
- [When Should I File a Counter-Notice?](#)
- [What Information Do You Need to Process a Counter-notice?](#)
- [What Happens After I Submit a Counter-notice?](#)
- [Filing a DMCA Notification or Counter-notice Is Serious Business!](#)

What is the DMCA?

[Section 512](#) of the Digital Millennium Copyright Act (“DMCA”) outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

What Types of Copyright Complaints Does Twitter Respond To?

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile photo, header photo, or background, allegations concerning the unauthorized use of a copyrighted image uploaded through our photo hosting service, or Tweets containing links to allegedly infringing materials.

Tip: If you are concerned about the use of your brand or entity’s name, please review Twitter’s [Trademark Policy](#). If you are concerned about the use of a fictional character, please see our [Parody, Commentary, and Fan Accounts Policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <http://chillingeffects.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);

3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

Tip: If you are reporting the content of a Tweet, please give us a direct link to that Tweet following these instructions: <https://support.twitter.com/articles/80586>. Or please specify if the alleged infringement is in the background, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to Twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate DMCA notifications may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with access to the reporter's complaint along with instructions on how to file a counter-notice, and forward a copy of the complaint to [Chilling Effects](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a copy of the takedown notice, along with instructions for filing a counter-notification.

Tip: If you've not yet received a copy of the DMCA notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a copy of each DMCA notification and counter-notice that we process to [Chilling Effects](#), where they are posted to a public-facing website (with your personal information removed).

Withheld Tweet:

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)

Withheld Media:

**Media not displayed**

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a DMCA Notification?

If you receive a DMCA notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

If you believe that the material reported in the DMCA notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Tip: Re-posting material removed in response to a DMCA notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the DMCA notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of Federal District Court for the judicial district in which the address is

located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

Tip: We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a DMCA Notification or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Targeted abuse: Same as "Violence & Threats/Abusive behavior policy"

Username squatting policy (20130817*)

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Following Rules and Best Practices (20110919*)

To learn more about following, please read "[FAQs About Following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that the only automated following behavior that Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. Please review our [Automation Rules and Best Practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and un-follow them. Aggressive follow churn is when an account *repeatedly* follows and then un-follows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 2000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 2000 people total**. Once you've followed 2000 users, there are limits to the number of additional users

you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

Tip: If you need to communicate with someone but don't need to see their updates every day, don't follow them. Visit their profile or send them an @reply when you need to; sending @replies doesn't require following,

and your reply will appear in the person's @mentions tab so they can reply back.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

The Twitter Rules addresses the use of these programs by disallowing: using or promoting third- party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third- party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Having Trouble?

Check out our [Following Troubleshooting](#) section for solutions to common problems

Automation Rules & Best Practices (2010/01)

Submitted Nov 08, 2009 by [ginger](#)

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account

when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#).

Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

The Twitter Rules

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

- **Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others.
- **Trademark:** We reserve the right to reclaim usernames on behalf of businesses or individuals that hold legal claim or trademark on those usernames. Accounts using business names and/or logos to mislead others may be permanently suspended.
- **Private information:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission.
- **Violence and Threats:** You may not publish or post direct, specific threats of violence against others.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Unlawful Use:** You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter Badges:** You may not use badges, such as but not limited to the Promoted or Verified Twitter badge, unless provided by Twitter. Accounts using these badges as part of profile photos, header photos, background images, or in a way that falsely implies affiliation with Twitter may be suspended.

Abuse and Spam

Twitter strives to protect its users from abuse and spam. User abuse and technical abuse are not tolerated on Twitter.com, and may result in permanent suspension. Any accounts engaging in the activities specified below may be subject to permanent suspension.

- **Serial Accounts:** You may not create multiple accounts for disruptive or abusive purposes, or with overlapping use cases. Mass account creation may result in suspension

of all related accounts. Please note that any violation of the Twitter Rules is cause for permanent suspension of all accounts.

- **Targeted Abuse:** You may not engage in targeted abuse or harassment. Some of the factors that we take into account when determining what conduct is considered to be targeted abuse or harassment are:
 - if you are sending messages to a user from multiple accounts;
 - if the sole purpose of your account is to send abusive messages to others;
 - if the reported behavior is one-sided or includes threats
- **Username Squatting:** You may not engage in username squatting. Accounts that are inactive for more than six months may also be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be username squatting are:
 - the number of accounts created
 - creating accounts for the purpose of preventing others from using those account names
 - creating accounts for the purpose of selling those accounts
 - using feeds of third-party content to update and maintain accounts under the names of those third parties
- **Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations.
- **Selling usernames:** You may not buy or sell Twitter usernames.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - If you have followed and/or unfollowed large amounts of users in a short time period, particularly by automated means (aggressive following or follower churn);
 - If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - If your updates consist mainly of links, and not personal updates;
 - If a large number of people are blocking you;
 - If a large number of spam complaints have been filed against you;
 - If you post duplicate content over multiple accounts or multiple duplicate updates on one account;

- If you post multiple unrelated updates to a topic using #, trending or popular topic, or promoted trend;
- If you send large numbers of duplicate @replies or mentions;
- If you send large numbers of unsolicited @replies or mentions in an aggressive attempt to bring attention to a service or link;
- If you add a large number of unrelated users to lists in an attempt to bring attention to an account, service or link;
- If you repeatedly create false or misleading content in an attempt to bring attention to an account, service or link;
- Randomly or aggressively favoriting Tweets through automation in an attempt to bring attention to an account, service or link;
- Randomly or aggressively Retweeting accounts through automation in an attempt to bring attention to an account, service or link;
- If you repeatedly post other users' account information as your own (bio, Tweets, url, etc.);
- If you post misleading links (e.g. affiliate links, links to malware/click jacking pages, etc.);
- Creating multiple misleading accounts in order to gain followers;
- Selling followers;
- Purchasing followers;
- Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising "more followers fast," or any other site that offers to automatically add followers to your account);
- **Pornography:** You may not use obscene or pornographic images in either your profile photo, header photo, or user background.

Your account may be suspended for Terms of Service violations if any of the above is true. Please see our help pages on [Following rules and best practices](#) and [Automation rules and best practices](#) for a more detailed discussion of how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the [Terms of Service](#).

We may revise these Rules from time to time; the most current version will always be at twitter.com/rules.

Have Questions?

Check out our complete list of articles outlining our [policies, guidelines, and best practices](#).

To report an account for violation of the Twitter Rules, [please use our forms](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter Impersonation Policy.

An account will not be removed if

- the user shares your name but has no other commonalities, or
- the profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same background or avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to [Twitter's Parody Policy](#) for more information about these types of accounts.

For Frequently Asked Questions about reporting impersonation on Twitter, [click here](#).

Trademark policy

Updated October 3, 2012

What is a Trademark Policy Violation on Twitter?

Using a company or business name, logo, or other trademark-protected materials in a manner that may mislead or confuse others with regard to its brand or business affiliation may be considered a trademark policy violation.

How Does Twitter Respond To Reported Trademark Policy Violations?

When we receive reports of trademark policy violations from holders of federal or international trademark registrations, we review the account and may take the following actions:

- When there is a clear intent to mislead others through the unauthorized use of a trademark, Twitter will suspend the account and notify the account holder.
- When we determine that an account appears to be confusing users, but is not purposefully passing itself off as the trademarked good or service, we give the account holder an opportunity to clear up any potential confusion. We may also release a username for the trademark holder's active use.
- We are responsive to reports about confusing or misleading [Promoted Tweet](#) and [Promoted Trend](#) copy, as well as [Promoted Account](#) profile information. For more information, please visit our [Trademark Policy for Promoted Products](#).

What Is not a Trademark Policy Violation?

Using another's trademark in a way that has nothing to do with the product or service for which the trademark was granted is not a violation of Twitter's trademark policy.

- Twitter usernames are provided on a first-come, first-served basis and may not be reserved. For information on why you may not be able to select a certain username, please see our [Why Can't I Register Certain Usernames](#) help page.
- Search and interest keywords advertisers choose are subject to applicable trademark law. Although Twitter may review reports of confusing keywords, we may not disable keywords in every case. For more information about our policies specific to Promoted Products, please see our [Promoted Products Policy](#) and our [Trademark Policy for Promoted Products](#).

Guidelines for News feed, Commentary, and Fan Accounts

Twitter users are allowed to create news feed, commentary, and fan accounts. Twitter provides a platform for its users to share and receive a wide range of ideas and content, and we greatly value and respect our users' expression. Because of these principles, we do not actively monitor users' content and will not edit or remove user content, except in cases of violations of our [Terms of Service](#).

An account's profile information should make it clear that the account is not actually the company or business entity that is the subject of the news feed/commentary/fan account. Here are some suggestions for distinguishing your account:

- **Username:** The username should not be the trademarked name of the subject of the news feed, commentary, or fan account.
- **Name:** The profile name should not be the trademarked name of the company or include the trademarked name in a misleading manner.
- **Bio:** The bio should include a statement to distinguish it from the real company, such as "Unofficial Account," "Fan Account," or "Not affiliated with..."
- **Profile photo, header photo, or background image:** The account should not use another's trademark, logo or other copyright-protected image without express permission.
- **Communication with other users:** The account should not, through private or public communication with other users, try to deceive or mislead others about your identity.

Users may also choose to use different language to indicate that an account is not associated with the actual brand/company/product so long as it is clear and not confusing to others, and does not mislead or deceive.

If an account is reported to be confusing, we may request that the account holder make further changes to bring the account in compliance with these best practices.

How Can I Make My Own Account's Brand or Business Affiliation Clear?

We strongly recommend that you use all of Twitter's account settings (account name, location, web, and bio) to make your account's affiliation clear.

- Please see our [How to Change Your Profile Picture or Information](#) page for instructions on customizing your account. In particular, we recommend clearly stating your location, including your website if you have one, and clearly describing your brand or business in the bio, if applicable.
- You can also consider linking your official website to your Twitter account. [Twitter Buttons](#) are an easy way to add Twitter actions like “Follow” and “Mention” directly to your website.

What Information is Required When Reporting Trademark Policy Violations?

In order to investigate trademark policy violations, please provide all of the following information:

Username of the reported account (e.g., @safety or <http://www.twitter.com/safety>):

Your company name:

Your company Twitter account (if there is one):

Company website:

Your trademarked word, symbol, etc. (e.g. Twitter):

Trademark registration number:

Trademark registration office (e.g., USPTO):

Note: A federal or international trademark registration number is required. If the name you are reporting is not a registered mark (e.g., a government agency or non-profit organization), please let us know:

Your first and last name:

Title:

Address:

Phone:

Fax:

Email (must be from company domain):

Description of confusion (e.g., passing off as your company, including specific descriptions of content or behavior):

Requested Action (e.g., removal of violating account or transfer of trademarked username to an existing company account):

How Do I Report a Trademark Policy Violation?

You do not need a Twitter account to submit a trademark report. Holders of registered trademarks can report possible violations to Twitter's Policy Team through our [support forms](#).

Please submit trademark-related requests from your company email address and follow the format above to help expedite our response. Also, be sure to clearly describe to us why the account or Tweets may cause confusion with your mark.

Private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#).

Some examples of private and confidential information include:

- credit card information
- social security or other national identity numbers
- addresses or locations that are considered and treated as private
- non-public, personal phone numbers
- non-public, personal email addresses
- images or videos that are considered and treated as private under applicable laws

Keep in mind that although you may consider certain information to be private, not all postings of such information may be a violation of this policy. We may consider the context and nature of the information posted, local privacy laws, and other case-specific facts when determining if this policy has been violated. For example, if information was previously posted or displayed elsewhere on the Internet prior to being put on Twitter, it may not be a violation of this policy.

If the private information you're reporting is not your own, please note that you must have documentation that you're authorized to act on behalf of the person whose confidential information is posted.

For frequently asked questions about reporting private information posted on Twitter, [click here](#). To learn more about protecting your private information on Twitter and other websites, [click here](#).

Abusive behavior policy

If you need to report abusive behavior to Twitter, please file a report [here](#).

If you believe you may be in danger, please contact your local law enforcement authority in addition to reporting

the content to Twitter so that the situation can also be addressed offline.

User disputes and false statements

Twitter is a platform that provides a global communication service which encompasses a variety of users with different voices, ideas and perspectives. As a policy, we do not mediate content or intervene in disputes between users. However, targeted abuse or harassment may constitute a violation of the [Twitter Rules](#) and [Terms of Service](#).

Offensive content

Users are allowed to post content, including potentially inflammatory content, provided they do not violate the [Twitter Terms of Service and Rules](#). Twitter does not screen content and we do not remove potentially offensive content unless such content is a violation of our [Terms of Service](#).

If you believe the content or behavior you are reporting is prohibited in your local jurisdiction, please contact your local authorities so they can accurately assess the content or behavior for possible violations of local law. If Twitter is contacted directly by law enforcement, we can work with them and provide assistance for their investigation as well as guidance around possible options. You can point local law enforcement to our [Law Enforcement Guidelines](#).

For Frequently Asked Questions about reporting abusive behavior on Twitter, [click here](#). To learn more about what you can do when you encounter abusive behavior on Twitter and other websites, [click here](#).

Copyright and DMCA policy

Topics Covered:

- [What is the DMCA ?](#)
- [What Types of Copyright Complaints Does Twitter Respond To?](#)
- [Am I a Copyright Holder? How Do I Know?](#)
- [What Information Do You Need To Process a Copyright Complaint?](#)
- [How Do I File a Copyright Complaint?](#)
- [How Are Claims Processed?](#)
- [What Happens Next?](#)
- [What Happens If I Receive a DMCA Notification?](#)
- [When Should I File a Counter-Notice?](#)
- [What Information Do You Need to Process a Counter-notice?](#)
- [What Happens After I Submit a Counter-notice?](#)
- [Filing a DMCA Notification or Counter-notice Is Serious Business!](#)

What is the DMCA?

[Section 512](#) of the Digital Millennium Copyright Act (“DMCA”) outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

What Types of Copyright Complaints Does Twitter Respond To?

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile photo, header photo, or background, allegations concerning the unauthorized use of a copyrighted image uploaded through our photo hosting service, or Tweets containing links to allegedly infringing materials.

Tip: If you are concerned about the use of your brand or entity’s name, please review

Twitter's [Trademark Policy](#). If you are concerned about the use of a fictional character, please see our [Parody, Commentary, and Fan Accounts Policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <http://chillingeffects.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

Tip: If you are reporting the content of a Tweet, please give us a direct link to that Tweet following these instructions: <https://support.twitter.com/articles/80586>. Or please specify if the alleged infringement is in the background, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to Twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate DMCA notifications may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with access to the reporter's complaint

along with instructions on how to file a counter-notice, and forward a copy of the complaint to [Chilling Effects](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a copy of the takedown notice, along with instructions for filing a counter-notification.

Tip: If you've not yet received a copy of the DMCA notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a copy of each DMCA notification and counter-notice that we process to [Chilling Effects](#), where they are posted to a public-facing website (with your personal information removed).

Withheld Tweet:

Tweet withheld 18h
This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)

Withheld Media:

 **Media not displayed**

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a DMCA Notification?

If you receive a DMCA notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

If you believe that the material reported in the DMCA notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Tip: Re-posting material removed in response to a DMCA notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the DMCA notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

Tip: We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a DMCA Notification or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Targeted abuse: Same as “Violence & Threats/Abusive behavior policy”

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Following Rules and Best Practices

To learn more about following, please read "[FAQs About Following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that the only automated following behavior that Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. Please review our [Automation Rules and Best Practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and un-follow them. Aggressive

follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 2000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 2000 people total**. Once you've followed 2000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

Tip: If you need to communicate with someone but don't need to see their updates every day, don't follow them. Visit their profile or send them an @reply when you need to; sending @replies doesn't require following, and your reply will appear in the person's @mentions tab so they can reply back.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

The [Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third- party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third- party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Having Trouble?

Check out our [Following Troubleshooting](#) section for solutions to common problems

Automation Rules & Best Practices

Submitted Nov 08, 2009 by ginger

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

The Twitter Rules

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

- **Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others.
- **Trademark:** We reserve the right to reclaim usernames on behalf of businesses or individuals that hold legal claim or trademark on those usernames. Accounts using business names and/or logos to mislead others may be permanently suspended.
- **Private information:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission.
- **Violence and Threats:** You may not publish or post direct, specific threats of violence against others.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Unlawful Use:** You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter Badges:** You may not use badges, such as but not limited to the Promoted or Verified Twitter badge, unless provided by Twitter. Accounts using these badges as part of profile photos, header photos, background images, or in a way that falsely implies affiliation with Twitter may be suspended.

Abuse and Spam

Twitter strives to protect its users from abuse and spam. User abuse and technical abuse are not tolerated on Twitter.com, and may result in permanent suspension. Any accounts engaging in the activities specified below may be subject to permanent suspension.

- **Serial Accounts:** You may not create multiple accounts for disruptive or abusive purposes, or with overlapping use cases. Mass account creation may result in suspension

of all related accounts. Please note that any violation of the Twitter Rules is cause for permanent suspension of all accounts.

- **Targeted Abuse:** You may not engage in targeted abuse or harassment. Some of the factors that we take into account when determining what conduct is considered to be targeted abuse or harassment are:
 - if you are sending messages to a user from multiple accounts;
 - if the sole purpose of your account is to send abusive messages to others;
 - if the reported behavior is one-sided or includes threats
- **Username Squatting:** You may not engage in username squatting. Accounts that are inactive for more than six months may also be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be username squatting are:
 - the number of accounts created
 - creating accounts for the purpose of preventing others from using those account names
 - creating accounts for the purpose of selling those accounts
 - using feeds of third-party content to update and maintain accounts under the names of those third parties
- **Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations.
- **Selling usernames:** You may not buy or sell Twitter usernames.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - If you have followed and/or unfollowed large amounts of users in a short time period, particularly by automated means (aggressive following or follower churn);
 - If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - If your updates consist mainly of links, and not personal updates;
 - If a large number of people are blocking you;
 - If a large number of spam complaints have been filed against you;
 - If you post duplicate content over multiple accounts or multiple duplicate updates on one account;

- If you post multiple unrelated updates to a topic using #, trending or popular topic, or promoted trend;
- If you send large numbers of duplicate @replies or mentions;
- If you send large numbers of unsolicited @replies or mentions in an aggressive attempt to bring attention to a service or link;
- If you add a large number of unrelated users to lists in an attempt to bring attention to an account, service or link;
- If you repeatedly create false or misleading content in an attempt to bring attention to an account, service or link;
- Randomly or aggressively favoriting Tweets through automation in an attempt to bring attention to an account, service or link;
- Randomly or aggressively Retweeting accounts through automation in an attempt to bring attention to an account, service or link;
- If you repeatedly post other users' account information as your own (bio, Tweets, url, etc.);
- If you post misleading links (e.g. affiliate links, links to malware/click jacking pages, etc.);
- Creating multiple misleading accounts in order to gain followers;
- Selling followers;
- Purchasing followers;
- Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising "more followers fast," or any other site that offers to automatically add followers to your account);
- **Pornography:** You may not use obscene or pornographic images in either your profile photo, header photo, or user background.

Your account may be suspended for Terms of Service violations if any of the above is true. Please see our help pages on [Following rules and best practices](#) and [Automation rules and best practices](#) for a more detailed discussion of how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the [Terms of Service](#).

We may revise these Rules from time to time; the most current version will always be at twitter.com/rules.

Have Questions?

Check out our complete list of articles outlining our [policies, guidelines, and best practices](#).

To report an account for violation of the Twitter Rules, [please use our forms](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter Impersonation Policy.

An account will not be removed if

- the user shares your name but has no other commonalities, or
- the profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same background or avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to [Twitter's Parody Policy](#) for more information about these types of accounts.

For Frequently Asked Questions about reporting impersonation on Twitter, [click here](#).

Trademark policy

Updated October 3, 2012

What is a Trademark Policy Violation on Twitter?

Using a company or business name, logo, or other trademark-protected materials in a manner that may mislead or confuse others with regard to its brand or business affiliation may be considered a trademark policy violation.

How Does Twitter Respond To Reported Trademark Policy Violations?

When we receive reports of trademark policy violations from holders of federal or international trademark registrations, we review the account and may take the following actions:

- When there is a clear intent to mislead others through the unauthorized use of a trademark, Twitter will suspend the account and notify the account holder.
- When we determine that an account appears to be confusing users, but is not purposefully passing itself off as the trademarked good or service, we give the account holder an opportunity to clear up any potential confusion. We may also release a username for the trademark holder's active use.
- We are responsive to reports about confusing or misleading [Promoted Tweet](#) and [Promoted Trend](#) copy, as well as [Promoted Account](#) profile information. For more information, please visit our [Trademark Policy for Promoted Products](#).

What Is not a Trademark Policy Violation?

Using another's trademark in a way that has nothing to do with the product or service for which the trademark was granted is not a violation of Twitter's trademark policy.

- Twitter usernames are provided on a first-come, first-served basis and may not be reserved. For information on why you may not be able to select a certain username, please see our [Why Can't I Register Certain Usernames](#) help page.
- Search and interest keywords advertisers choose are subject to applicable trademark law. Although Twitter may review reports of confusing keywords, we may not disable keywords in every case. For more information about our policies specific to Promoted Products, please see our [Promoted Products Policy](#) and our [Trademark Policy for Promoted Products](#).

Guidelines for News feed, Commentary, and Fan Accounts

Twitter users are allowed to create news feed, commentary, and fan accounts. Twitter provides a platform for its users to share and receive a wide range of ideas and content, and we greatly value and respect our users' expression. Because of these principles, we do not actively monitor users' content and will not edit or remove user content, except in cases of violations of our [Terms of Service](#).

An account's profile information should make it clear that the account is not actually the company or business entity that is the subject of the news feed/commentary/fan account. Here are some suggestions for distinguishing your account:

- **Username:** The username should not be the trademarked name of the subject of the news feed, commentary, or fan account.
- **Name:** The profile name should not be the trademarked name of the company or include the trademarked name in a misleading manner.
- **Bio:** The bio should include a statement to distinguish it from the real company, such as "Unofficial Account," "Fan Account," or "Not affiliated with..."
- **Profile photo, header photo, or background image:** The account should not use another's trademark, logo or other copyright-protected image without express permission.
- **Communication with other users:** The account should not, through private or public communication with other users, try to deceive or mislead others about your identity.

Users may also choose to use different language to indicate that an account is not associated with the actual brand/company/product so long as it is clear and not confusing to others, and does not mislead or deceive.

If an account is reported to be confusing, we may request that the account holder make further changes to bring the account in compliance with these best practices.

How Can I Make My Own Account's Brand or Business Affiliation Clear?

We strongly recommend that you use all of Twitter's account settings (account name, location, web, and bio) to make your account's affiliation clear.

- Please see our [How to Change Your Profile Picture or Information](#) page for instructions on customizing your account. In particular, we recommend clearly stating your location, including your website if you have one, and clearly describing your brand or business in the bio, if applicable.
- You can also consider linking your official website to your Twitter account. [Twitter Buttons](#) are an easy way to add Twitter actions like “Follow” and “Mention” directly to your website.

What Information is Required When Reporting Trademark Policy Violations?

In order to investigate trademark policy violations, please provide all of the following information:

Username of the reported account (e.g., @safety or <http://www.twitter.com/safety>):

Your company name:

Your company Twitter account (if there is one):

Company website:

Your trademarked word, symbol, etc. (e.g. Twitter):

Trademark registration number:

Trademark registration office (e.g., USPTO):

Note: A federal or international trademark registration number is required. If the name you are reporting is not a registered mark (e.g., a government agency or non-profit organization), please let us know:

Your first and last name:

Title:

Address:

Phone:

Fax:

Email (must be from company domain):

Description of confusion (e.g., passing off as your company, including specific descriptions of content or behavior):

Requested Action (e.g., removal of violating account or transfer of trademarked username to an existing company account):

How Do I Report a Trademark Policy Violation?

You do not need a Twitter account to submit a trademark report. Holders of registered trademarks can report possible violations to Twitter's Policy Team through our [support forms](#).

Please submit trademark-related requests from your company email address and follow the format above to help expedite our response. Also, be sure to clearly describe to us why the account or Tweets may cause confusion with your mark.

Private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#).

Some examples of private and confidential information include:

- credit card information
- social security or other national identity numbers
- addresses or locations that are considered and treated as private
- non-public, personal phone numbers
- non-public, personal email addresses
- images or videos that are considered and treated as private under applicable laws

Keep in mind that although you may consider certain information to be private, not all postings of such information may be a violation of this policy. We may consider the context and nature of the information posted, local privacy laws, and other case-specific facts when determining if this policy has been violated. For example, if information was previously posted or displayed elsewhere on the Internet prior to being put on Twitter, it may not be a violation of this policy.

If the private information you're reporting is not your own, please note that you must have documentation that you're authorized to act on behalf of the person whose confidential information is posted.

For frequently asked questions about reporting private information posted on Twitter, [click here](#). To learn more about protecting your private information on Twitter and other websites, [click here](#).

Abusive behavior policy

If you need to report abusive behavior to Twitter, please file a report [here](#).

If you believe you may be in danger, please contact your local law enforcement authority in addition to reporting

the content to Twitter so that the situation can be addressed both online and offline.

User disputes and false statements

Twitter provides a global communication platform which encompasses a variety of users with different voices, ideas and perspectives. As a policy, we do not mediate content or intervene in disputes between users.

Threats and abuse

Users may not make direct, specific threats of violence against others; targeted abuse or harassment is also a violation of the [Twitter Rules](#) and [Terms of Service](#).

For frequently asked questions about reporting abusive behavior on Twitter, [click here](#). To learn more about what you can do when you encounter abusive behavior on Twitter and other websites, [click here](#).

Offensive content

Users are allowed to post content, including potentially inflammatory content, provided they do not violate the [Twitter Rules](#) and [Terms of Service](#). Twitter does not screen content and does not remove potentially offensive content unless such content is in violation of the [Twitter Rules](#) and [Terms of Service](#).

If you believe the content or behavior you are reporting is prohibited in your local jurisdiction, please contact your local authorities so they can accurately assess the content or behavior for possible violations of local law. If Twitter is contacted directly by law enforcement, we can work with them and provide assistance for their investigation as well as guidance around possible options. You can point local law enforcement to our [Law Enforcement Guidelines](#).

Copyright and DMCA policy

Topics Covered:

- [What is the DMCA ?](#)
- [What Types of Copyright Complaints Does Twitter Respond To?](#)
- [Am I a Copyright Holder? How Do I Know?](#)
- [What Information Do You Need To Process a Copyright Complaint?](#)
- [How Do I File a Copyright Complaint?](#)
- [How Are Claims Processed?](#)
- [What Happens Next?](#)
- [What Happens If I Receive a DMCA Notification?](#)
- [When Should I File a Counter-Notice?](#)
- [What Information Do You Need to Process a Counter-notice?](#)
- [What Happens After I Submit a Counter-notice?](#)
- [Filing a DMCA Notification or Counter-notice Is Serious Business!](#)

What is the DMCA?

[Section 512](#) of the Digital Millennium Copyright Act (“DMCA”) outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

What Types of Copyright Complaints Does Twitter Respond To?

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile photo, header photo, or background, allegations concerning the unauthorized use of a copyrighted image uploaded through our photo hosting service, or Tweets containing links to allegedly infringing materials.

Tip: If you are concerned about the use of your brand or entity’s name, please review

Twitter's [Trademark Policy](#). If you are concerned about the use of a fictional character, please see our [Parody, Commentary, and Fan Accounts Policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <http://chillingeffects.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

Tip: If you are reporting the content of a Tweet, please give us a direct link to that Tweet following these instructions: <https://support.twitter.com/articles/80586>. Or please specify if the alleged infringement is in the background, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to Twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate DMCA notifications may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with access to the reporter's complaint

along with instructions on how to file a counter-notice, and forward a copy of the complaint to [Chilling Effects](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a copy of the takedown notice, along with instructions for filing a counter-notification.

Tip: If you've not yet received a copy of the DMCA notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a copy of each DMCA notification and counter-notice that we process to [Chilling Effects](#), where they are posted to a public-facing website (with your personal information removed).

Withheld Tweet:

Tweet withheld 18h
This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)

Withheld Media:

 **Media not displayed**

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a DMCA Notification?

If you receive a DMCA notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

If you believe that the material reported in the DMCA notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Tip: Re-posting material removed in response to a DMCA notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the DMCA notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

Tip: We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a DMCA Notification or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Targeted abuse: Same as “Violence & Threats/Abusive behavior policy”

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Following Rules and Best Practices

To learn more about following, please read "[FAQs About Following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that the only automated following behavior that Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. Please review our [Automation Rules and Best Practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and un-follow them. Aggressive follow churn is when an account *repeatedly* follows and then un-follows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 2000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 2000 people total**. Once you've followed 2000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

Tip: If you need to communicate with someone but don't need to see their updates every day, don't follow them. Visit their profile or send them an @reply when you need to; sending @replies doesn't require following, and your reply will appear in the person's @mentions tab so they can reply back.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

The [Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third- party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third- party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Having Trouble?

Check out our [Following Troubleshooting](#) section for solutions to common problems

Automation Rules & Best Practices

Submitted Nov 08, 2009 by ginger

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the time-stamp, like this:

The Twitter Rules (20130813)

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

- **Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others.
- **Trademark:** We reserve the right to reclaim usernames on behalf of businesses or individuals that hold legal claim or trademark on those usernames. Accounts using business names and/or logos to mislead others may be permanently suspended.
- **Private information:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission.
- **Violence and Threats:** You may not publish or post direct, specific threats of violence against others.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Unlawful Use:** You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter Badges:** You may not use badges, such as but not limited to the Promoted or Verified Twitter badge, unless provided by Twitter. Accounts using these badges as part of profile photos, header photos, background images, or in a way that falsely implies affiliation with Twitter may be suspended.

Abuse and Spam

Twitter strives to protect its users from abuse and spam. User abuse and technical abuse are not tolerated on Twitter.com, and may result in permanent suspension. Any accounts engaging in the activities specified below may be subject to permanent suspension.

- **Serial Accounts:** You may not create multiple accounts for disruptive or abusive purposes, or with overlapping use cases. Mass account creation may result in suspension

of all related accounts. Please note that any violation of the Twitter Rules is cause for permanent suspension of all accounts.

- **Targeted Abuse:** You may not engage in targeted abuse or harassment. Some of the factors that we take into account when determining what conduct is considered to be targeted abuse or harassment are:
 - if you are sending messages to a user from multiple accounts;
 - if the sole purpose of your account is to send abusive messages to others;
 - if the reported behavior is one-sided or includes threats
- **Username Squatting:** You may not engage in username squatting. Accounts that are inactive for more than six months may also be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be username squatting are:
 - the number of accounts created
 - creating accounts for the purpose of preventing others from using those account names
 - creating accounts for the purpose of selling those accounts
 - using feeds of third-party content to update and maintain accounts under the names of those third parties
- **Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations.
- **Selling usernames:** You may not buy or sell Twitter usernames.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - If you have followed and/or unfollowed large amounts of users in a short time period, particularly by automated means (aggressive following or follower churn);
 - If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - If your updates consist mainly of links, and not personal updates;
 - If a large number of people are blocking you;
 - If a large number of spam complaints have been filed against you;
 - If you post duplicate content over multiple accounts or multiple duplicate updates on one account;

- If you post multiple unrelated updates to a topic using #, trending or popular topic, or promoted trend;
- If you send large numbers of duplicate @replies or mentions;
- If you send large numbers of unsolicited @replies or mentions in an aggressive attempt to bring attention to a service or link;
- If you add a large number of unrelated users to lists in an attempt to bring attention to an account, service or link;
- If you repeatedly create false or misleading content in an attempt to bring attention to an account, service or link;
- Randomly or aggressively favoriting Tweets through automation in an attempt to bring attention to an account, service or link;
- Randomly or aggressively Retweeting accounts through automation in an attempt to bring attention to an account, service or link;
- If you repeatedly post other users' account information as your own (bio, Tweets, url, etc.);
- If you post misleading links (e.g. affiliate links, links to malware/click jacking pages, etc.);
- Creating multiple misleading accounts in order to gain followers;
- Selling followers;
- Purchasing followers;
- Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising "more followers fast," or any other site that offers to automatically add followers to your account);
- **Pornography:** You may not use obscene or pornographic images in either your profile photo, header photo, or user background.

Your account may be suspended for Terms of Service violations if any of the above is true. Please see our help pages on [Following rules and best practices](#) and [Automation rules and best practices](#) for a more detailed discussion of how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the [Terms of Service](#).

We may revise these Rules from time to time; the most current version will always be at twitter.com/rules.

Have Questions?

Check out our complete list of articles outlining our [policies, guidelines, and best practices](#).

To report an account for violation of the Twitter Rules, [please use our forms](#).

Impersonation policy (201305*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter Impersonation Policy.

An account will not be removed if

- the user shares your name but has no other commonalities, or
- the profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same background or avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to [Twitter's Parody Policy](#) for more information about these types of accounts.

For Frequently Asked Questions about reporting impersonation on Twitter, [click here](#).

Trademark policy (20121003)

Updated October 3, 2012

What is a Trademark Policy Violation on Twitter?

Using a company or business name, logo, or other trademark-protected materials in a manner that may mislead or confuse others with regard to its brand or business affiliation may be considered a trademark policy violation.

How Does Twitter Respond To Reported Trademark Policy Violations?

When we receive reports of trademark policy violations from holders of federal or international trademark registrations, we review the account and may take the following actions:

- When there is a clear intent to mislead others through the unauthorized use of a trademark, Twitter will suspend the account and notify the account holder.
- When we determine that an account appears to be confusing users, but is not purposefully passing itself off as the trademarked good or service, we give the account holder an opportunity to clear up any potential confusion. We may also release a username for the trademark holder's active use.
- We are responsive to reports about confusing or misleading [Promoted Tweet](#) and [Promoted Trend](#) copy, as well as [Promoted Account](#) profile information. For more information, please visit our [Trademark Policy for Promoted Products](#).

What Is not a Trademark Policy Violation?

Using another's trademark in a way that has nothing to do with the product or service for which the trademark was granted is not a violation of Twitter's trademark policy.

- Twitter usernames are provided on a first-come, first-served basis and may not be reserved. For information on why you may not be able to select a certain username, please see our [Why Can't I Register Certain Usernames](#) help page.
- Search and interest keywords advertisers choose are subject to applicable trademark law. Although Twitter may review reports of confusing keywords, we may not disable keywords in every case. For more information about our policies specific to Promoted Products, please see our [Promoted Products Policy](#) and our [Trademark Policy for Promoted Products](#).

Guidelines for News feed, Commentary, and Fan Accounts

Twitter users are allowed to create news feed, commentary, and fan accounts. Twitter provides a platform for its users to share and receive a wide range of ideas and content, and we greatly value and respect our users' expression. Because of these principles, we do not actively monitor users' content and will not edit or remove user content, except in cases of violations of our [Terms of Service](#).

An account's profile information should make it clear that the account is not actually the company or business entity that is the subject of the news feed/commentary/fan account. Here are some suggestions for distinguishing your account:

- **Username:** The username should not be the trademarked name of the subject of the news feed, commentary, or fan account.
- **Name:** The profile name should not be the trademarked name of the company or include the trademarked name in a misleading manner.
- **Bio:** The bio should include a statement to distinguish it from the real company, such as "Unofficial Account," "Fan Account," or "Not affiliated with..."
- **Profile photo, header photo, or background image:** The account should not use another's trademark, logo or other copyright-protected image without express permission.
- **Communication with other users:** The account should not, through private or public communication with other users, try to deceive or mislead others about your identity.

Users may also choose to use different language to indicate that an account is not associated with the actual brand/company/product so long as it is clear and not confusing to others, and does not mislead or deceive.

If an account is reported to be confusing, we may request that the account holder make further changes to bring the account in compliance with these best practices.

How Can I Make My Own Account's Brand or Business Affiliation Clear?

We strongly recommend that you use all of Twitter's account settings (account name, location, web, and bio) to make your account's affiliation clear.

- Please see our [How to Change Your Profile Picture or Information](#) page for instructions on customizing your account. In particular, we recommend clearly stating your location, including your website if you have one, and clearly describing your brand or business in the bio, if applicable.
- You can also consider linking your official website to your Twitter account. [Twitter Buttons](#) are an easy way to add Twitter actions like “Follow” and “Mention” directly to your website.

What Information is Required When Reporting Trademark Policy Violations?

In order to investigate trademark policy violations, please provide all of the following information:

Username of the reported account (e.g., @safety or <http://www.twitter.com/safety>):

Your company name:

Your company Twitter account (if there is one):

Company website:

Your trademarked word, symbol, etc. (e.g. Twitter):

Trademark registration number:

Trademark registration office (e.g., USPTO):

Note: A federal or international trademark registration number is required. If the name you are reporting is not a registered mark (e.g., a government agency or non-profit organization), please let us know:

Your first and last name:

Title:

Address:

Phone:

Fax:

Email (must be from company domain):

Description of confusion (e.g., passing off as your company, including specific descriptions of content or behavior):

Requested Action (e.g., removal of violating account or transfer of trademarked username to an existing company account):

How Do I Report a Trademark Policy Violation?

You do not need a Twitter account to submit a trademark report. Holders of registered trademarks can report possible violations to Twitter's Policy Team through our [support forms](#).

Please submit trademark-related requests from your company email address and follow the format above to help expedite our response. Also, be sure to clearly describe to us why the account or Tweets may cause confusion with your mark.

Private information posted on Twitter (20130821*)

Posting another person's private and confidential information is a violation of the [Twitter Rules](#).

Some examples of private and confidential information include:

- credit card information
- social security or other national identity numbers
- addresses or locations that are considered and treated as private
- non-public, personal phone numbers
- non-public, personal email addresses
- images or videos that are considered and treated as private under applicable laws

Keep in mind that although you may consider certain information to be private, not all postings of such information may be a violation of this policy. We may consider the context and nature of the information posted, local privacy laws, and other case-specific facts when determining if this policy has been violated. For example, if information was previously posted or displayed elsewhere on the Internet prior to being put on Twitter, it may not be a violation of this policy.

If the private information you're reporting is not your own, please note that you must have documentation that you're authorized to act on behalf of the person whose confidential information is posted.

For frequently asked questions about reporting private information posted on Twitter, [click here](#). To learn more about protecting your private information on Twitter and other websites, [click here](#).

Abusive behavior policy (20130929)

If you need to report abusive behavior to Twitter, please file a report [here](#).

If you believe you may be in danger, please contact your local law enforcement authority in addition to reporting

the content to Twitter so that the situation can be addressed both online and offline.

User disputes and false statements

Twitter provides a global communication platform which encompasses a variety of users with different voices, ideas and perspectives. As a policy, we do not mediate content or intervene in disputes between users.

Threats and abuse

Users may not make direct, specific threats of violence against others; targeted abuse or harassment is also a violation of the [Twitter Rules](#) and [Terms of Service](#).

For frequently asked questions about reporting abusive behavior on Twitter, [click here](#). To learn more about what you can do when you encounter abusive behavior on Twitter and other websites, [click here](#).

Offensive content

Users are allowed to post content, including potentially inflammatory content, provided they do not violate the [Twitter Rules](#) and [Terms of Service](#). Twitter does not screen content and does not remove potentially offensive content unless such content is in violation of the [Twitter Rules](#) and [Terms of Service](#).

If you believe the content or behavior you are reporting is prohibited in your local jurisdiction, please contact your local authorities so they can accurately assess the content or behavior for possible violations of local law. If Twitter is contacted directly by law enforcement, we can work with them and provide assistance for their investigation as well as guidance around possible options. You can point local law enforcement to our [Law Enforcement Guidelines](#).

Copyright and DMCA policy (20130430*)

Topics Covered:

- [What is the DMCA ?](#)
- [What Types of Copyright Complaints Does Twitter Respond To?](#)
- [Am I a Copyright Holder? How Do I Know?](#)
- [What Information Do You Need To Process a Copyright Complaint?](#)
- [How Do I File a Copyright Complaint?](#)
- [How Are Claims Processed?](#)
- [What Happens Next?](#)
- [What Happens If I Receive a DMCA Notification?](#)
- [When Should I File a Counter-Notice?](#)
- [What Information Do You Need to Process a Counter-notice?](#)
- [What Happens After I Submit a Counter-notice?](#)
- [Filing a DMCA Notification or Counter-notice Is Serious Business!](#)

What is the DMCA?

[Section 512](#) of the Digital Millennium Copyright Act (“DMCA”) outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

What Types of Copyright Complaints Does Twitter Respond To?

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile photo, header photo, or background, allegations concerning the unauthorized use of a copyrighted image uploaded through our photo hosting service, or Tweets containing links to allegedly infringing materials.

Tip: If you are concerned about the use of your brand or entity's name, please review Twitter's [Trademark Policy](#). If you are concerned about the use of a fictional character, please see our [Parody, Commentary, and Fan Accounts Policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <http://chillingeffects.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

Tip: If you are reporting the content of a Tweet, please give us a direct link to that Tweet following these instructions: <https://support.twitter.com/articles/80586>. Or please specify if the alleged infringement is in the background, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to Twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate DMCA notifications may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with access to the reporter's complaint

along with instructions on how to file a counter-notice, and forward a copy of the complaint to [Chilling Effects](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a copy of the takedown notice, along with instructions for filing a counter-notification.

Tip: If you've not yet received a copy of the DMCA notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a copy of each DMCA notification and counter-notice that we process to [Chilling Effects](#), where they are posted to a public-facing website (with your personal information removed).

Withheld Tweet:

Tweet withheld 18h
This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)

Withheld Media:

 **Media not displayed**

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a DMCA Notification?

If you receive a DMCA notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

If you believe that the material reported in the DMCA notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Tip: Re-posting material removed in response to a DMCA notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the DMCA notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

Tip: We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a DMCA Notification or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Targeted abuse: Same as “Violence & Threats/Abusive behavior policy”

Username squatting policy (20130817*)

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Following rules and best practices (20131209*)

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or un-following behavior. Please review our [Automation Rules and Best Practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and un-follow them. Aggressive follow churn is when an account *repeatedly* follows and then un-follows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 2000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 2000 people total**. Once you've followed 2000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

Tip: If you need to communicate with someone but don't need to see their updates every day, don't follow them. Visit their profile or send them an @reply when you need to; sending @replies doesn't require following, and your reply will appear in the person's @mentions tab so they can reply back.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

The Twitter Rules addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn’t a race to get the most followers. If you follow users that you’re interested in and post meaningful content, it’s more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Having trouble?

Check out our [Following Troubleshooting](#) section for solutions to common problems.

Automation Rules & Best Practices (2010/01)

Submitted Nov 08, 2009 by ginger

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.

- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

The Twitter Rules

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

- **Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others.
- **Trademark:** We reserve the right to reclaim usernames on behalf of businesses or individuals that hold legal claim or trademark on those usernames. Accounts using business names and/or logos to mislead others may be permanently suspended.
- **Private information:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission.
- **Violence and Threats:** You may not publish or post direct, specific threats of violence against others.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Unlawful Use:** You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter Badges:** You may not use badges, such as but not limited to the Promoted or Verified Twitter badge, unless provided by Twitter. Accounts using these badges as part of profile photos, header photos, background images, or in a way that falsely implies affiliation with Twitter may be suspended.

Abuse and Spam

Twitter strives to protect its users from abuse and spam. User abuse and technical abuse are not tolerated on Twitter.com, and may result in permanent suspension. Any accounts engaging in the activities specified below may be subject to permanent suspension.

- **Serial Accounts:** You may not create multiple accounts for disruptive or abusive purposes, or with overlapping use cases. Mass account creation may result in suspension

of all related accounts. Please note that any violation of the Twitter Rules is cause for permanent suspension of all accounts.

- **Targeted Abuse:** You may not engage in targeted abuse or harassment. Some of the factors that we take into account when determining what conduct is considered to be targeted abuse or harassment are:
 - if you are sending messages to a user from multiple accounts;
 - if the sole purpose of your account is to send abusive messages to others;
 - if the reported behavior is one-sided or includes threats
- **Username Squatting:** You may not engage in username squatting. Accounts that are inactive for more than six months may also be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be username squatting are:
 - the number of accounts created
 - creating accounts for the purpose of preventing others from using those account names
 - creating accounts for the purpose of selling those accounts
 - using feeds of third-party content to update and maintain accounts under the names of those third parties
- **Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations.
- **Selling usernames:** You may not buy or sell Twitter usernames.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - If you have followed and/or unfollowed large amounts of users in a short time period, particularly by automated means (aggressive following or follower churn);
 - If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - If your updates consist mainly of links, and not personal updates;
 - If a large number of people are blocking you;
 - If a large number of spam complaints have been filed against you;
 - If you post duplicate content over multiple accounts or multiple duplicate updates on one account;

- If you post multiple unrelated updates to a topic using #, trending or popular topic, or promoted trend;
- If you send large numbers of duplicate @replies or mentions;
- If you send large numbers of unsolicited @replies or mentions in an aggressive attempt to bring attention to a service or link;
- If you add a large number of unrelated users to lists in an attempt to bring attention to an account, service or link;
- If you repeatedly create false or misleading content in an attempt to bring attention to an account, service or link;
- Randomly or aggressively favoriting Tweets through automation in an attempt to bring attention to an account, service or link;
- Randomly or aggressively Retweeting accounts through automation in an attempt to bring attention to an account, service or link;
- If you repeatedly post other users' account information as your own (bio, Tweets, url, etc.);
- If you post misleading links (e.g. affiliate links, links to malware/click jacking pages, etc.);
- Creating multiple misleading accounts in order to gain followers;
- Selling followers;
- Purchasing followers;
- Using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising "more followers fast," or any other site that offers to automatically add followers to your account);
- **Pornography:** You may not use obscene or pornographic images in either your profile photo, header photo, or user background.

Your account may be suspended for Terms of Service violations if any of the above is true. Please see our help pages on [Following rules and best practices](#) and [Automation rules and best practices](#) for a more detailed discussion of how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the [Terms of Service](#).

We may revise these Rules from time to time; the most current version will always be at twitter.com/rules.

Have Questions?

Check out our complete list of articles outlining our [policies, guidelines, and best practices](#).

To report an account for violation of the Twitter Rules, [please use our forms](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter Impersonation Policy.

An account will not be removed if

- the user shares your name but has no other commonalities, or
- the profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same background or avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to [Twitter's Parody Policy](#) for more information about these types of accounts.

For Frequently Asked Questions about reporting impersonation on Twitter, [click here](#).

Trademark policy

Updated October 3, 2012

What is a Trademark Policy Violation on Twitter?

Using a company or business name, logo, or other trademark-protected materials in a manner that may mislead or confuse others with regard to its brand or business affiliation may be considered a trademark policy violation.

How Does Twitter Respond To Reported Trademark Policy Violations?

When we receive reports of trademark policy violations from holders of federal or international trademark registrations, we review the account and may take the following actions:

- When there is a clear intent to mislead others through the unauthorized use of a trademark, Twitter will suspend the account and notify the account holder.
- When we determine that an account appears to be confusing users, but is not purposefully passing itself off as the trademarked good or service, we give the account holder an opportunity to clear up any potential confusion. We may also release a username for the trademark holder's active use.
- We are responsive to reports about confusing or misleading [Promoted Tweet](#) and [Promoted Trend](#) copy, as well as [Promoted Account](#) profile information. For more information, please visit our [Trademark Policy for Promoted Products](#).

What Is not a Trademark Policy Violation?

Using another's trademark in a way that has nothing to do with the product or service for which the trademark was granted is not a violation of Twitter's trademark policy.

- Twitter usernames are provided on a first-come, first-served basis and may not be reserved. For information on why you may not be able to select a certain username, please see our [Why Can't I Register Certain Usernames](#) help page.
- Search and interest keywords advertisers choose are subject to applicable trademark law. Although Twitter may review reports of confusing keywords, we may not disable keywords in every case. For more information about our policies specific to Promoted Products, please see our [Promoted Products Policy](#) and our [Trademark Policy for Promoted Products](#).

Guidelines for News feed, Commentary, and Fan Accounts

Twitter users are allowed to create news feed, commentary, and fan accounts. Twitter provides a platform for its users to share and receive a wide range of ideas and content, and we greatly value and respect our users' expression. Because of these principles, we do not actively monitor users' content and will not edit or remove user content, except in cases of violations of our [Terms of Service](#).

An account's profile information should make it clear that the account is not actually the company or business entity that is the subject of the news feed/commentary/fan account. Here are some suggestions for distinguishing your account:

- **Username:** The username should not be the trademarked name of the subject of the news feed, commentary, or fan account.
- **Name:** The profile name should not be the trademarked name of the company or include the trademarked name in a misleading manner.
- **Bio:** The bio should include a statement to distinguish it from the real company, such as "Unofficial Account," "Fan Account," or "Not affiliated with..."
- **Profile photo, header photo, or background image:** The account should not use another's trademark, logo or other copyright-protected image without express permission.
- **Communication with other users:** The account should not, through private or public communication with other users, try to deceive or mislead others about your identity.

Users may also choose to use different language to indicate that an account is not associated with the actual brand/company/product so long as it is clear and not confusing to others, and does not mislead or deceive.

If an account is reported to be confusing, we may request that the account holder make further changes to bring the account in compliance with these best practices.

How Can I Make My Own Account's Brand or Business Affiliation Clear?

We strongly recommend that you use all of Twitter's account settings (account name, location, web, and bio) to make your account's affiliation clear.

- Please see our [How to Change Your Profile Picture or Information](#) page for instructions on customizing your account. In particular, we recommend clearly stating your location, including your website if you have one, and clearly describing your brand or business in the bio, if applicable.
- You can also consider linking your official website to your Twitter account. [Twitter Buttons](#) are an easy way to add Twitter actions like “Follow” and “Mention” directly to your website.

What Information is Required When Reporting Trademark Policy Violations?

In order to investigate trademark policy violations, please provide all of the following information:

Username of the reported account (e.g., @safety or <http://www.twitter.com/safety>):

Your company name:

Your company Twitter account (if there is one):

Company website:

Your trademarked word, symbol, etc. (e.g. Twitter):

Trademark registration number:

Trademark registration office (e.g., USPTO):

Note: A federal or international trademark registration number is required. If the name you are reporting is not a registered mark (e.g., a government agency or non-profit organization), please let us know:

Your first and last name:

Title:

Address:

Phone:

Fax:

Email (must be from company domain):

Description of confusion (e.g., passing off as your company, including specific descriptions of content or behavior):

Requested Action (e.g., removal of violating account or transfer of trademarked username to an existing company account):

How Do I Report a Trademark Policy Violation?

You do not need a Twitter account to submit a trademark report. Holders of registered trademarks can report possible violations to Twitter's Policy Team through our [support forms](#).

Please submit trademark-related requests from your company email address and follow the format above to help expedite our response. Also, be sure to clearly describe to us why the account or Tweets may cause confusion with your mark.

Private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#).

Some examples of private and confidential information include:

- credit card information
- social security or other national identity numbers
- addresses or locations that are considered and treated as private
- non-public, personal phone numbers
- non-public, personal email addresses
- images or videos that are considered and treated as private under applicable laws

Keep in mind that although you may consider certain information to be private, not all postings of such information may be a violation of this policy. We may consider the context and nature of the information posted, local privacy laws, and other case-specific facts when determining if this policy has been violated. For example, if information was previously posted or displayed elsewhere on the Internet prior to being put on Twitter, it may not be a violation of this policy.

If the private information you're reporting is not your own, please note that you must have documentation that you're authorized to act on behalf of the person whose confidential information is posted.

For frequently asked questions about reporting private information posted on Twitter, [click here](#). To learn more about protecting your private information on Twitter and other websites, [click here](#).

Abusive behavior policy

If you need to report abusive behavior to Twitter, please file a report [here](#).

If you believe you may be in danger, please contact your local law enforcement authority in addition to reporting

the content to Twitter so that the situation can be addressed both online and offline.

User disputes and false statements

Twitter provides a global communication platform which encompasses a variety of users with different voices, ideas and perspectives. As a policy, we do not mediate content or intervene in disputes between users.

Threats and abuse

Users may not make direct, specific threats of violence against others, including threats against a person or group on the basis of race, ethnicity, national origin, religion, sexual orientation, gender, gender identity, age, or disability. Targeted abuse or harassment is also a violation of the [Twitter Rules](#) and [Terms of Service](#).

For frequently asked questions about reporting abusive behavior on Twitter, [click here](#). To learn more about what you can do when you encounter abusive behavior on Twitter and other websites, [click here](#).

Offensive content

Users are allowed to post content, including potentially inflammatory content, provided they do not violate the [Twitter Rules](#) and [Terms of Service](#). Twitter does not screen content and does not remove potentially offensive content unless such content is in violation of the [Twitter Rules](#) and [Terms of Service](#).

If you believe the content or behavior you are reporting is prohibited in your local jurisdiction, please contact your local authorities so they can accurately assess the content or behavior for possible violations of local law. If Twitter is contacted directly by law enforcement, we can work with them and provide assistance for their investigation as well as guidance around possible options. You can point local law enforcement to our [Law Enforcement Guidelines](#).

Copyright and DMCA policy

Topics Covered:

- [What is the DMCA ?](#)
- [What Types of Copyright Complaints Does Twitter Respond To?](#)
- [Am I a Copyright Holder? How Do I Know?](#)
- [What Information Do You Need To Process a Copyright Complaint?](#)
- [How Do I File a Copyright Complaint?](#)
- [How Are Claims Processed?](#)
- [What Happens Next?](#)
- [What Happens If I Receive a DMCA Notification?](#)
- [When Should I File a Counter-Notice?](#)
- [What Information Do You Need to Process a Counter-notice?](#)
- [What Happens After I Submit a Counter-notice?](#)
- [Filing a DMCA Notification or Counter-notice Is Serious Business!](#)

What is the DMCA?

[Section 512](#) of the Digital Millennium Copyright Act (“DMCA”) outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

[What Types of Copyright Complaints Does Twitter Respond To?](#)

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile photo, header photo, or background, allegations concerning the unauthorized use of a copyrighted image uploaded through our photo hosting service, or Tweets containing links to allegedly infringing materials.

Tip: If you are concerned about the use of your brand or entity's name, please review Twitter's [Trademark Policy](#). If you are concerned about the use of a fictional character, please see our [Parody, Commentary, and Fan Accounts Policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <http://chillingeffects.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

Tip: If you are reporting the content of a Tweet, please give us a direct link to that Tweet following these instructions: <https://support.twitter.com/articles/80586>. Or please specify if the alleged infringement is in the background, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to Twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate DMCA notifications may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with access to the reporter's complaint

along with instructions on how to file a counter-notice, and forward a copy of the complaint to [Chilling Effects](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a copy of the takedown notice, along with instructions for filing a counter-notification.

Tip: If you've not yet received a copy of the DMCA notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a copy of each DMCA notification and counter-notice that we process to [Chilling Effects](#), where they are posted to a public-facing website (with your personal information removed).

Withheld Tweet:

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)

Withheld Media:



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a DMCA Notification?

If you receive a DMCA notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

If you believe that the material reported in the DMCA notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Tip: Re-posting material removed in response to a DMCA notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the DMCA notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

Tip: We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a DMCA Notification or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Targeted abuse: Same as “Violence & Threats/Abusive behavior policy”

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or un-following behavior. Please review our [Automation Rules and Best Practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and un-follow them. Aggressive follow churn is when an account *repeatedly* follows and then un-follows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 2000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 2000 people total**. Once you've followed 2000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this

a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

Tip: If you need to communicate with someone but don't need to see their updates every day, don't follow them. Visit their profile or send them an @reply when you need to; sending @replies doesn't require following, and your reply will appear in the person's @mentions tab so they can reply back.

What are third party applications to “Get More Followers

Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

The [Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Having trouble?

Check out our [Following Troubleshooting section](#) for solutions to common problems.

Automation Rules & Best Practices

Submitted Nov 08, 2009 by **ginger**

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.

3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of

users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

The Twitter Rules

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

- **Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others.
- **Trademark:** We reserve the right to reclaim usernames on behalf of businesses or individuals that hold legal claim or trademark on those usernames. Accounts using business names and/or logos to mislead others may be permanently suspended.
- **Private information:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission.
- **Violence and Threats:** You may not publish or post direct, specific threats of violence against others.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Unlawful Use:** You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter Badges:** You may not use badges, such as but not limited to the Promoted or Verified Twitter badge, unless provided by Twitter. Accounts using these badges as part of profile photos, header photos, background images, or in a way that falsely implies affiliation with Twitter may be suspended.

Abuse and Spam

Twitter strives to protect its users from abuse and spam. User abuse and technical abuse are not tolerated on Twitter.com, and may result in permanent suspension. Any accounts engaging in the activities specified below may be subject to permanent suspension.

- **Serial Accounts:** You may not create multiple accounts for disruptive or abusive purposes, or with overlapping use cases. Mass account creation may result in suspension of all related accounts. Please note that any violation of the Twitter Rules is cause for permanent suspension of all accounts.
- **Targeted Abuse:** You may not engage in targeted abuse or harassment. Some of the factors that we take into account when determining what conduct is considered to be targeted abuse or harassment are:
 - if you are sending messages to a user from multiple accounts;
 - if the sole purpose of your account is to send abusive messages to others;
 - if the reported behavior is one-sided or includes threats
- **Username Squatting:** You may not engage in username squatting. Accounts that are inactive for more than six months may also be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be username squatting are:
 - the number of accounts created
 - creating accounts for the purpose of preventing others from using those account names
 - creating accounts for the purpose of selling those accounts
 - using feeds of third-party content to update and maintain accounts under the names of those third parties
- **Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations.
- **Selling usernames:** You may not buy or sell Twitter usernames.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - If you have followed and/or unfollowed large amounts of users in a short time period, particularly by automated means (aggressive following or follower churn);
 - If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - If your updates consist mainly of links, and not personal updates;
 - If a large number of people are blocking you;
 - If a large number of spam complaints have been filed against you;

- If you post duplicate content over multiple accounts or multiple duplicate updates on one account;
- If you post multiple unrelated updates to a topic using #, trending or popular topic, or promoted trend;
- If you send large numbers of duplicate @replies or mentions;
- If you send large numbers of unsolicited @replies or mentions;
- If you add a large number of unrelated users to lists;
- If you repeatedly create false or misleading content;
- Randomly or aggressively following, favoriting or Retweeting Tweets;
- If you repeatedly post other users' account information as your own (bio, Tweets, url, etc.);
- If you post misleading links (e.g. affiliate links, links to malware/click jacking pages, etc.);
- Creating misleading accounts or account interactions;
- Selling or purchasing account interactions (such as selling or purchasing followers, Retweets, favorites, etc.);
- Using or promoting third-party services or apps that claim to get you more followers (such as follower trains, sites promising "more followers fast" or any other site that offers to automatically add followers to your account);
- **Pornography:** You may not use obscene or pornographic images in either your profile photo, header photo, or user background.

Your account may be suspended for Terms of Service violations if any of the above is true. Please see our help pages on [Following rules and best practices](#) and [Automation rules and best practices](#) for a more detailed discussion of how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the [Terms of Service](#).

We may revise these Rules from time to time; the most current version will always be at twitter.com/rules.

Have Questions?

Check out our complete list of articles outlining our [policies, guidelines, and best practices](#).

To report an account for violation of the Twitter Rules, [please use our forms](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter Impersonation Policy.

An account will not be removed if

- the user shares your name but has no other commonalities, or
- the profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same background or avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to [Twitter's Parody Policy](#) for more information about these types of accounts.

For Frequently Asked Questions about reporting impersonation on Twitter, [click here](#).

Trademark policy

Updated October 3, 2012

What is a Trademark Policy Violation on Twitter?

Using a company or business name, logo, or other trademark-protected materials in a manner that may mislead or confuse others with regard to its brand or business affiliation may be considered a trademark policy violation.

How Does Twitter Respond To Reported Trademark Policy Violations?

When we receive reports of trademark policy violations from holders of federal or international trademark registrations, we review the account and may take the following actions:

- When there is a clear intent to mislead others through the unauthorized use of a trademark, Twitter will suspend the account and notify the account holder.
- When we determine that an account appears to be confusing users, but is not purposefully passing itself off as the trademarked good or service, we give the account holder an opportunity to clear up any potential confusion. We may also release a username for the trademark holder's active use.
- We are responsive to reports about confusing or misleading [Promoted Tweet](#) and [Promoted Trend](#) copy, as well as [Promoted Account](#) profile information. For more information, please visit our [Trademark Policy for Promoted Products](#).

What Is not a Trademark Policy Violation?

Using another's trademark in a way that has nothing to do with the product or service for which the trademark was granted is not a violation of Twitter's trademark policy.

- Twitter usernames are provided on a first-come, first-served basis and may not be reserved. For information on why you may not be able to select a certain username, please see our [Why Can't I Register Certain Usernames](#) help page.
- Search and interest keywords advertisers choose are subject to applicable trademark law. Although Twitter may review reports of confusing keywords, we may not disable keywords in every case. For more information about our policies specific to Promoted Products, please see our [Promoted Products Policy](#) and our [Trademark Policy for Promoted Products](#).

Guidelines for News feed, Commentary, and Fan Accounts

Twitter users are allowed to create news feed, commentary, and fan accounts. Twitter provides a platform for its users to share and receive a wide range of ideas and content, and we greatly value and respect our users' expression. Because of these principles, we do not actively monitor users' content and will not edit or remove user content, except in cases of violations of our [Terms of Service](#). An account's profile information should make it clear that the account is not actually the company or business entity that is the subject of the news feed/commentary/fan account. Here are some suggestions for distinguishing your account:

- **Username:** The username should not be the trademarked name of the subject of the news feed, commentary, or fan account.
- **Name:** The profile name should not be the trademarked name of the company or include the trademarked name in a misleading manner.
- **Bio:** The bio should include a statement to distinguish it from the real company, such as "Unofficial Account," "Fan Account," or "Not affiliated with..."
- **Profile photo, header photo, or background image:** The account should not use another's trademark, logo or other copyright-protected image without express permission.
- **Communication with other users:** The account should not, through private or public communication with other users, try to deceive or mislead others about your identity.

Users may also choose to use different language to indicate that an account is not associated with the actual brand/company/product so long as it is clear and not confusing to others, and does not mislead or deceive.

If an account is reported to be confusing, we may request that the account holder make further changes to bring the account in compliance with these best practices.

How Can I Make My Own Account's Brand or Business Affiliation Clear?

We strongly recommend that you use all of Twitter's account settings (account name, location, web, and bio) to make your account's affiliation clear.

- Please see our [How to Change Your Profile Picture or Information](#) page for instructions on customizing your account. In particular, we recommend clearly stating your location,

- including your website if you have one, and clearly describing your brand or business in the bio, if applicable.
- You can also consider linking your official website to your Twitter account. [Twitter Buttons](#) are an easy way to add Twitter actions like “Follow” and “Mention” directly to your website.

What Information is Required When Reporting Trademark Policy Violations?

In order to investigate trademark policy violations, please provide all of the following information:

Username of the reported account (e.g., @safety or <http://www.twitter.com/safety>):

Your company name:

Your company Twitter account (if there is one):

Company website:

Your trademarked word, symbol, etc. (e.g. Twitter):

Trademark registration number:

Trademark registration office (e.g., USPTO):

Note: A federal or international trademark registration number is required. If the name you are reporting is not a registered mark (e.g., a government agency or non-profit organization), please let us know:

Your first and last name:

Title:

Address:

Phone:

Fax:

Email (must be from company domain):

Description of confusion (e.g., passing off as your company, including specific descriptions of content or behavior):

Requested Action (e.g., removal of violating account or transfer of trademarked username to an existing company account):

How Do I Report a Trademark Policy Violation?

You do not need a Twitter account to submit a trademark report. Holders of registered trademarks can report possible violations to Twitter's Policy Team through our [support forms](#).

Please submit trademark-related requests from your company email address and follow the format above to help expedite our response. Also, be sure to clearly describe to us why the account or Tweets may cause confusion with your mark.

Private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#).

Some examples of private and confidential information include:

- credit card information
- social security or other national identity numbers
- addresses or locations that are considered and treated as private
- non-public, personal phone numbers
- non-public, personal email addresses
- images or videos that are considered and treated as private under applicable laws

Keep in mind that although you may consider certain information to be private, not all postings of such information may be a violation of this policy. We may consider the context and nature of the information posted, local privacy laws, and other case-specific facts when determining if this policy has been violated. For example, if information was previously posted or displayed elsewhere on the Internet prior to being put on Twitter, it may not be a violation of this policy.

If the private information you're reporting is not your own, please note that you must have documentation that you're authorized to act on behalf of the person whose confidential information is posted.

For frequently asked questions about reporting private information posted on Twitter, [click here](#). To learn more about protecting your private information on Twitter and other websites, [click here](#).

Abusive behavior policy

If you need to report abusive behavior to Twitter, please file a report [here](#).

If you believe you may be in danger, please contact your local law enforcement authority in addition to reporting

the content to Twitter so that the situation can be addressed both online and offline.

User disputes and false statements

Twitter provides a global communication platform which encompasses a variety of users with different voices, ideas and perspectives. As a policy, we do not mediate content or intervene in disputes between users.

Threats and abuse

Users may not make direct, specific threats of violence against others, including threats against a person or group on the basis of race, ethnicity, national origin, religion, sexual orientation, gender, gender identity, age, or disability. Targeted abuse or harassment is also a violation of the [Twitter Rules](#) and [Terms of Service](#).

For frequently asked questions about reporting abusive behavior on Twitter, [click here](#). To learn more about what you can do when you encounter abusive behavior on Twitter and other websites, [click here](#).

Offensive content

Users are allowed to post content, including potentially inflammatory content, provided they do not violate the [Twitter Rules](#) and [Terms of Service](#). Twitter does not screen content and does not remove potentially offensive content unless such content is in violation of the [Twitter Rules](#) and [Terms of Service](#).

If you believe the content or behavior you are reporting is prohibited in your local jurisdiction, please contact your local authorities so they can accurately assess the content or behavior for possible violations of local law. If Twitter is contacted directly by law enforcement, we can work with them and provide assistance for their investigation as well as guidance around possible options. You can point local law enforcement to our [Law Enforcement Guidelines](#).

Copyright and DMCA policy

Topics Covered:

- [What is the DMCA ?](#)
- [What Types of Copyright Complaints Does Twitter Respond To?](#)
- [Am I a Copyright Holder? How Do I Know?](#)
- [What Information Do You Need To Process a Copyright Complaint?](#)
- [How Do I File a Copyright Complaint?](#)
- [How Are Claims Processed?](#)
- [What Happens Next?](#)
- [What Happens If I Receive a DMCA Notification?](#)
- [When Should I File a Counter-Notice?](#)
- [What Information Do You Need to Process a Counter-notice?](#)
- [What Happens After I Submit a Counter-notice?](#)
- [Filing a DMCA Notification or Counter-notice Is Serious Business!](#)

What is the DMCA?

[Section 512](#) of the Digital Millennium Copyright Act (“DMCA”) outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

[What Types of Copyright Complaints Does Twitter Respond To?](#)

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile photo, header photo, or background, allegations concerning the unauthorized use of a copyrighted image uploaded through our photo hosting service, or Tweets containing links to allegedly infringing materials.

Tip: If you are concerned about the use of your brand or entity's name, please review Twitter's [Trademark Policy](#). If you are concerned about the use of a fictional character, please see our [Parody, Commentary, and Fan Accounts Policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <http://chillingeffects.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

Tip: If you are reporting the content of a Tweet, please give us a direct link to that Tweet following these instructions: <https://support.twitter.com/articles/80586>. Or please specify if the alleged infringement is in the background, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to Twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate DMCA notifications may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with access to the reporter's complaint

along with instructions on how to file a counter-notice, and forward a copy of the complaint to [Chilling Effects](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a copy of the takedown notice, along with instructions for filing a counter-notification.

Tip: If you've not yet received a copy of the DMCA notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a copy of each DMCA notification and counter-notice that we process to [Chilling Effects](#), where they are posted to a public-facing website (with your personal information removed).

Withheld Tweet:

Tweet withheld 18h
This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)

Withheld Media:

 **Media not displayed**

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a DMCA Notification?

If you receive a DMCA notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

If you believe that the material reported in the DMCA notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Tip: Re-posting material removed in response to a DMCA notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the DMCA notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

Tip: We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a DMCA Notification or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Targeted abuse: Same as “Violence & Threats/Abusive behavior policy”

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or un-following behavior. Please review our [Automation Rules and Best Practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and un-follow them. Aggressive follow churn is when an account *repeatedly* follows and then un-follows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 2000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 2000 people total**. Once you've followed 2000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this

a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

Tip: If you need to communicate with someone but don't need to see their updates every day, don't follow them. Visit their profile or send them an @reply when you need to; sending @replies doesn't require following, and your reply will appear in the person's @mentions tab so they can reply back.

What are third party applications to “Get More Followers

Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

The [Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Having trouble?

Check out our [Following Troubleshooting section](#) for solutions to common problems.

Automation Rules & Best Practices

Submitted Nov 08, 2009 by **ginger**

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.

3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of

users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

The Twitter Rules (20150314*)

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

- **Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others.
- **Trademark:** We reserve the right to reclaim usernames on behalf of businesses or individuals that hold legal claim or trademark on those usernames. Accounts using business names and/or logos to mislead others may be permanently suspended.
- **Private information:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission. You may not post intimate photos or videos that were taken or distributed without the subject's consent.
- **Violence and Threats:** You may not publish or post direct, specific threats of violence against others.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Unlawful Use:** You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter Badges:** You may not use badges, such as but not limited to the Promoted or Verified Twitter badge, unless provided by Twitter. Accounts using these badges as part of profile photos, header photos, background images, or in a way that falsely implies affiliation with Twitter may be suspended.

Abuse and Spam

Twitter strives to protect its users from abuse and spam. User abuse and technical abuse are not tolerated on Twitter.com, and may result in permanent suspension. Any accounts engaging in the activities specified below may be subject to permanent suspension.

- **Serial Accounts:** You may not create multiple accounts for disruptive or abusive purposes, or with overlapping use cases. Mass account creation may result in suspension of all related accounts. Please note that any violation of the Twitter Rules is cause for permanent suspension of all accounts.
- **Targeted Abuse:** You may not engage in targeted abuse or harassment. Some of the factors that we take into account when determining what conduct is considered to be targeted abuse or harassment are:
 - if you are sending messages to a user from multiple accounts;
 - if the sole purpose of your account is to send abusive messages to others;
 - if the reported behavior is one-sided or includes threats
- **Username Squatting:** You may not engage in username squatting. Accounts that are inactive for more than six months may also be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be username squatting are:
 - the number of accounts created
 - creating accounts for the purpose of preventing others from using those account names
 - creating accounts for the purpose of selling those accounts
 - using feeds of third-party content to update and maintain accounts under the names of those third parties
- **Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations.
- **Selling usernames:** You may not buy or sell Twitter usernames.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - If you have followed and/or unfollowed large amounts of users in a short time period, particularly by automated means (aggressive following or follower churn);
 - If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - If your updates consist mainly of links, and not personal updates;
 - If a large number of people are blocking you;
 - If a large number of spam complaints have been filed against you;

- If you post duplicate content over multiple accounts or multiple duplicate updates on one account;
- If you post multiple unrelated updates to a topic using #, trending or popular topic, or promoted trend;
- If you send large numbers of duplicate @replies or mentions;
- If you send large numbers of unsolicited @replies or mentions;
- If you add a large number of unrelated users to lists;
- If you repeatedly create false or misleading content;
- Randomly or aggressively following, favoriting or Retweeting Tweets;
- If you repeatedly post other users' account information as your own (bio, Tweets, url, etc.);
- If you post misleading links (e.g. affiliate links, links to malware/click jacking pages, etc.);
- Creating misleading accounts or account interactions;
- Selling or purchasing account interactions (such as selling or purchasing followers, Retweets, favorites, etc.);
- Using or promoting third-party services or apps that claim to get you more followers (such as follower trains, sites promising "more followers fast" or any other site that offers to automatically add followers to your account);
- **Graphic Content:** You may not use pornographic or excessively violent media in your profile image, header image, or background image.

Your account may be suspended for Terms of Service violations if any of the above is true. Please see our help pages on [Following rules and best practices](#) and [Automation rules and best practices](#) for a more detailed discussion of how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the [Terms of Service](#).

We may revise these Rules from time to time; the most current version will always be at twitter.com/rules.

Have Questions?

Check out our complete list of articles outlining our [policies, guidelines, and best practices](#).

To report an account for violation of the Twitter Rules, [please use our forms](#).

Impersonation policy (201305*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter Impersonation Policy.

An account will not be removed if

- the user shares your name but has no other commonalities, or
- the profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same background or avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to [Twitter's Parody Policy](#) for more information about these types of accounts.

For Frequently Asked Questions about reporting impersonation on Twitter, [click here](#).

Trademark policy (20121003)

Updated October 3, 2012

What is a Trademark Policy Violation on Twitter?

Using a company or business name, logo, or other trademark-protected materials in a manner that may mislead or confuse others with regard to its brand or business affiliation may be considered a trademark policy violation.

How Does Twitter Respond To Reported Trademark Policy Violations?

When we receive reports of trademark policy violations from holders of federal or international trademark registrations, we review the account and may take the following actions:

- When there is a clear intent to mislead others through the unauthorized use of a trademark, Twitter will suspend the account and notify the account holder.
- When we determine that an account appears to be confusing users, but is not purposefully passing itself off as the trademarked good or service, we give the account holder an opportunity to clear up any potential confusion. We may also release a username for the trademark holder's active use.
- We are responsive to reports about confusing or misleading [Promoted Tweet](#) and [Promoted Trend](#) copy, as well as [Promoted Account](#) profile information. For more information, please visit our [Trademark Policy for Promoted Products](#).

What Is not a Trademark Policy Violation?

Using another's trademark in a way that has nothing to do with the product or service for which the trademark was granted is not a violation of Twitter's trademark policy.

- Twitter usernames are provided on a first-come, first-served basis and may not be reserved. For information on why you may not be able to select a certain username, please see our [Why Can't I Register Certain Usernames](#) help page.
- Search and interest keywords advertisers choose are subject to applicable trademark law. Although Twitter may review reports of confusing keywords, we may not disable keywords in every case. For more information about our policies specific to Promoted Products, please see our [Promoted Products Policy](#) and our [Trademark Policy for Promoted Products](#).

Guidelines for News feed, Commentary, and Fan Accounts

Twitter users are allowed to create news feed, commentary, and fan accounts. Twitter provides a platform for its users to share and receive a wide range of ideas and content, and we greatly value and respect our users' expression. Because of these principles, we do not actively monitor users' content and will not edit or remove user content, except in cases of violations of our [Terms of Service](#). An account's profile information should make it clear that the account is not actually the company or business entity that is the subject of the news feed/commentary/fan account. Here are some suggestions for distinguishing your account:

- **Username:** The username should not be the trademarked name of the subject of the news feed, commentary, or fan account.
- **Name:** The profile name should not be the trademarked name of the company or include the trademarked name in a misleading manner.
- **Bio:** The bio should include a statement to distinguish it from the real company, such as "Unofficial Account," "Fan Account," or "Not affiliated with..."
- **Profile photo, header photo, or background image:** The account should not use another's trademark, logo or other copyright-protected image without express permission.
- **Communication with other users:** The account should not, through private or public communication with other users, try to deceive or mislead others about your identity.

Users may also choose to use different language to indicate that an account is not associated with the actual brand/company/product so long as it is clear and not confusing to others, and does not mislead or deceive.

If an account is reported to be confusing, we may request that the account holder make further changes to bring the account in compliance with these best practices.

How Can I Make My Own Account's Brand or Business Affiliation Clear?

We strongly recommend that you use all of Twitter's account settings (account name, location, web, and bio) to make your account's affiliation clear.

- Please see our [How to Change Your Profile Picture or Information](#) page for instructions on customizing your account. In particular, we recommend clearly stating your location,

- including your website if you have one, and clearly describing your brand or business in the bio, if applicable.
- You can also consider linking your official website to your Twitter account. [Twitter Buttons](#) are an easy way to add Twitter actions like “Follow” and “Mention” directly to your website.

What Information is Required When Reporting Trademark Policy Violations?

In order to investigate trademark policy violations, please provide all of the following information:

Username of the reported account (e.g., @safety or <http://www.twitter.com/safety>):

Your company name:

Your company Twitter account (if there is one):

Company website:

Your trademarked word, symbol, etc. (e.g. Twitter):

Trademark registration number:

Trademark registration office (e.g., USPTO):

Note: A federal or international trademark registration number is required. If the name you are reporting is not a registered mark (e.g., a government agency or non-profit organization), please let us know:

Your first and last name:

Title:

Address:

Phone:

Fax:

Email (must be from company domain):

Description of confusion (e.g., passing off as your company, including specific descriptions of content or behavior):

Requested Action (e.g., removal of violating account or transfer of trademarked username to an existing company account):

How Do I Report a Trademark Policy Violation?

You do not need a Twitter account to submit a trademark report. Holders of registered trademarks can report possible violations to Twitter's Policy Team through our [support forms](#).

Please submit trademark-related requests from your company email address and follow the format above to help expedite our response. Also, be sure to clearly describe to us why the account or Tweets may cause confusion with your mark.

Private information posted on Twitter (20130821*)

Posting another person's private and confidential information is a violation of the [Twitter Rules](#).

Some examples of private and confidential information include:

- credit card information
- social security or other national identity numbers
- addresses or locations that are considered and treated as private
- non-public, personal phone numbers
- non-public, personal email addresses
- images or videos that are considered and treated as private under applicable laws

Keep in mind that although you may consider certain information to be private, not all postings of such information may be a violation of this policy. We may consider the context and nature of the information posted, local privacy laws, and other case-specific facts when determining if this policy has been violated. For example, if information was previously posted or displayed elsewhere on the Internet prior to being put on Twitter, it may not be a violation of this policy.

If the private information you're reporting is not your own, please note that you must have documentation that you're authorized to act on behalf of the person whose confidential information is posted.

For frequently asked questions about reporting private information posted on Twitter, [click here](#). To learn more about protecting your private information on Twitter and other websites, [click here](#).

Abusive behavior policy (20140729*)

If you need to report abusive behavior to Twitter, please file a report [here](#).

If you believe you may be in danger, please contact your local law enforcement authority in addition to reporting

the content to Twitter so that the situation can be addressed both online and offline.

User disputes and false statements

Twitter provides a global communication platform which encompasses a variety of users with different voices, ideas and perspectives. As a policy, we do not mediate content or intervene in disputes between users.

Threats and abuse

Users may not make direct, specific threats of violence against others, including threats against a person or group on the basis of race, ethnicity, national origin, religion, sexual orientation, gender, gender identity, age, or disability. Targeted abuse or harassment is also a violation of the [Twitter Rules](#) and [Terms of Service](#).

For frequently asked questions about reporting abusive behavior on Twitter, [click here](#). To learn more about what you can do when you encounter abusive behavior on Twitter and other websites, [click here](#).

Offensive content

Users are allowed to post content, including potentially inflammatory content, provided they do not violate the [Twitter Rules](#) and [Terms of Service](#). Twitter does not screen content and does not remove potentially offensive content unless such content is in violation of the [Twitter Rules](#) and [Terms of Service](#).

If you believe the content or behavior you are reporting is prohibited in your local jurisdiction, please contact your local authorities so they can accurately assess the content or behavior for possible violations of local law. If Twitter is contacted directly by law enforcement, we can work with them and provide assistance for their investigation as well as guidance around possible options. You can point local law enforcement to our [Law Enforcement Guidelines](#).

Copyright and DMCA policy (20130430*)

Topics Covered:

- [What is the DMCA ?](#)
- [What Types of Copyright Complaints Does Twitter Respond To?](#)
- [Am I a Copyright Holder? How Do I Know?](#)
- [What Information Do You Need To Process a Copyright Complaint?](#)
- [How Do I File a Copyright Complaint?](#)
- [How Are Claims Processed?](#)
- [What Happens Next?](#)
- [What Happens If I Receive a DMCA Notification?](#)
- [When Should I File a Counter-Notice?](#)
- [What Information Do You Need to Process a Counter-notice?](#)
- [What Happens After I Submit a Counter-notice?](#)
- [Filing a DMCA Notification or Counter-notice Is Serious Business!](#)

What is the DMCA?

[Section 512](#) of the Digital Millennium Copyright Act (“DMCA”) outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

[What Types of Copyright Complaints Does Twitter Respond To?](#)

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile photo, header photo, or background, allegations concerning the unauthorized use of a copyrighted image uploaded through our photo hosting service,

or Tweets containing links to allegedly infringing materials.

Tip: If you are concerned about the use of your brand or entity's name, please review Twitter's [Trademark Policy](#). If you are concerned about the use of a fictional character, please see our [Parody, Commentary, and Fan Accounts Policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <http://chillingeffects.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

Tip: If you are reporting the content of a Tweet, please give us a direct link to that Tweet following these instructions: <https://support.twitter.com/articles/80586>. Or please specify if the alleged infringement is in the background, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to Twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate DMCA notifications may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with access to the reporter's complaint along with instructions on how to file a counter-notice, and forward a copy of the complaint to [Chilling Effects](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a copy of the takedown notice, along with instructions for filing a counter-notification.

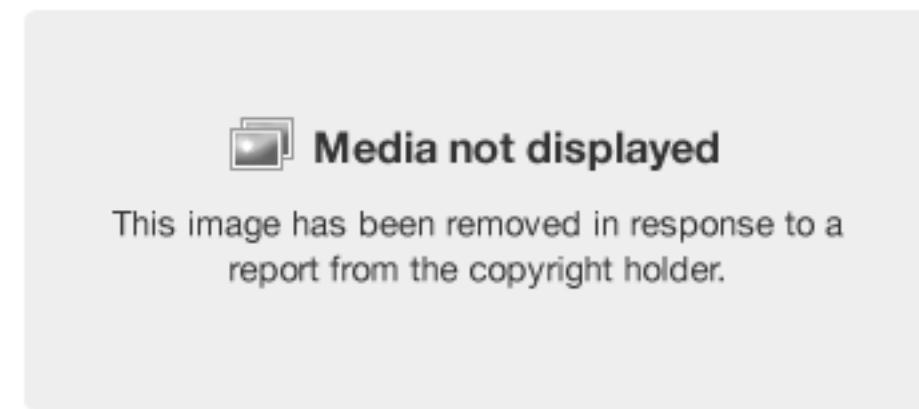
Tip: If you've not yet received a copy of the DMCA notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a copy of each DMCA notification and counter-notice that we process to [Chilling Effects](#), where they are posted to a public-facing website (with your personal information removed).

Withheld Tweet:



Withheld Media:



Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a DMCA Notification?

If you receive a DMCA notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to

read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

If you believe that the material reported in the DMCA notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Tip: Re-posting material removed in response to a DMCA notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the DMCA notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

Tip: We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a DMCA Notification or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Targeted abuse: Same as "Violence & Threats/Abusive behavior policy"

Username squatting policy (20130817*)

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Following rules and best practices (20131209*)

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or un-following behavior. Please review our [Automation Rules and Best Practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and un-follow them. Aggressive follow churn is when an account *repeatedly* follows and then un-follows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 2000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 2000 people total**. Once you've followed 2000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

Tip: If you need to communicate with someone but don't need to see their updates every day, don't follow them. Visit their profile or send them an @reply when you need to; sending @replies doesn't require following, and your reply will appear in the person's @mentions tab so they can reply back.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

The Twitter Rules addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn’t a race to get the most followers. If you follow users that you’re interested in and post meaningful content, it’s more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Having trouble?

Check out our [Following Troubleshooting](#) section for solutions to common problems.

Automation Rules & Best Practices (2010/01)

Submitted Nov 08, 2009 by ginger

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.

- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

The Twitter Rules (20150419*)

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

- **Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others.
- **Trademark:** We reserve the right to reclaim usernames on behalf of businesses or individuals that hold legal claim or trademark on those usernames. Accounts using business names and/or logos to mislead others may be permanently suspended.
- **Private information:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission. You may not post intimate photos or videos that were taken or distributed without the subject's consent.
- **Violence and Threats:** You may not publish or post threats of violence against others or promote violence against others.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Unlawful Use:** You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter Badges:** You may not use badges, such as but not limited to the Promoted or Verified Twitter badge, unless provided by Twitter. Accounts using these badges as part of profile photos, header photos, background images, or in a way that falsely implies affiliation with Twitter may be suspended.

Abuse and Spam

Twitter strives to protect its users from abuse and spam. User abuse and technical abuse are not tolerated on Twitter.com, and may result in permanent suspension. Any accounts engaging in the activities specified below may be subject to permanent suspension.

- **Serial Accounts:** You may not create multiple accounts for disruptive or abusive purposes, or with overlapping use cases. Mass account creation may result in suspension of all related accounts. Please note that any violation of the Twitter Rules is cause for permanent suspension of all accounts.
- **Targeted Abuse:** You may not engage in targeted abuse or harassment. Some of the factors that we take into account when determining what conduct is considered to be targeted abuse or harassment are:
 - if you are sending messages to a user from multiple accounts;
 - if the sole purpose of your account is to send abusive messages to others;
 - if the reported behavior is one-sided or includes threats
- **Username Squatting:** You may not engage in username squatting. Accounts that are inactive for more than six months may also be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be username squatting are:
 - the number of accounts created
 - creating accounts for the purpose of preventing others from using those account names
 - creating accounts for the purpose of selling those accounts
 - using feeds of third-party content to update and maintain accounts under the names of those third parties
- **Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations.
- **Selling usernames:** You may not buy or sell Twitter usernames.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - If you have followed and/or unfollowed large amounts of users in a short time period, particularly by automated means (aggressive following or follower churn);
 - If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - If your updates consist mainly of links, and not personal updates;
 - If a large number of people are blocking you;
 - If a large number of spam complaints have been filed against you;

- If you post duplicate content over multiple accounts or multiple duplicate updates on one account;
- If you post multiple unrelated updates to a topic using #, trending or popular topic, or promoted trend;
- If you send large numbers of duplicate @replies or mentions;
- If you send large numbers of unsolicited @replies or mentions;
- If you add a large number of unrelated users to lists;
- If you repeatedly create false or misleading content;
- Randomly or aggressively following, favoriting or Retweeting Tweets;
- If you repeatedly post other users' account information as your own (bio, Tweets, url, etc.);
- If you post misleading links (e.g. affiliate links, links to malware/click jacking pages, etc.);
- Creating misleading accounts or account interactions;
- Selling or purchasing account interactions (such as selling or purchasing followers, Retweets, favorites, etc.);
- Using or promoting third-party services or apps that claim to get you more followers (such as follower trains, sites promising "more followers fast" or any other site that offers to automatically add followers to your account);
- **Graphic Content:** You may not use pornographic or excessively violent media in your profile image, header image, or background image.

Your account may be suspended for Terms of Service violations if any of the above is true. Please see our help pages on [Following rules and best practices](#) and [Automation rules and best practices](#) for a more detailed discussion of how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the [Terms of Service](#).

We may revise these Rules from time to time; the most current version will always be at twitter.com/rules.

Have Questions?

Check out our complete list of articles outlining our [policies, guidelines, and best practices](#).

To report an account for violation of the Twitter Rules, [please use our forms](#).

Impersonation policy (201305*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter Impersonation Policy.

An account will not be removed if

- the user shares your name but has no other commonalities, or
- the profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same background or avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to [Twitter's Parody Policy](#) for more information about these types of accounts.

For Frequently Asked Questions about reporting impersonation on Twitter, [click here](#).

Trademark policy (20121003)

Updated October 3, 2012

What is a Trademark Policy Violation on Twitter?

Using a company or business name, logo, or other trademark-protected materials in a manner that may mislead or confuse others with regard to its brand or business affiliation may be considered a trademark policy violation.

How Does Twitter Respond To Reported Trademark Policy Violations?

When we receive reports of trademark policy violations from holders of federal or international trademark registrations, we review the account and may take the following actions:

- When there is a clear intent to mislead others through the unauthorized use of a trademark, Twitter will suspend the account and notify the account holder.
- When we determine that an account appears to be confusing users, but is not purposefully passing itself off as the trademarked good or service, we give the account holder an opportunity to clear up any potential confusion. We may also release a username for the trademark holder's active use.
- We are responsive to reports about confusing or misleading [Promoted Tweet](#) and [Promoted Trend](#) copy, as well as [Promoted Account](#) profile information. For more information, please visit our [Trademark Policy for Promoted Products](#).

What Is not a Trademark Policy Violation?

Using another's trademark in a way that has nothing to do with the product or service for which the trademark was granted is not a violation of Twitter's trademark policy.

- Twitter usernames are provided on a first-come, first-served basis and may not be reserved. For information on why you may not be able to select a certain username, please see our [Why Can't I Register Certain Usernames](#) help page.
- Search and interest keywords advertisers choose are subject to applicable trademark law. Although Twitter may review reports of confusing keywords, we may not disable keywords in every case. For more information about our policies specific to Promoted Products, please see our [Promoted Products Policy](#) and our [Trademark Policy for Promoted Products](#).

Guidelines for News feed, Commentary, and Fan Accounts

Twitter users are allowed to create news feed, commentary, and fan accounts. Twitter provides a platform for its users to share and receive a wide range of ideas and content, and we greatly value and respect our users' expression. Because of these principles, we do not actively monitor users' content and will not edit or remove user content, except in cases of violations of our [Terms of Service](#). An account's profile information should make it clear that the account is not actually the company or business entity that is the subject of the news feed/commentary/fan account. Here are some suggestions for distinguishing your account:

- **Username:** The username should not be the trademarked name of the subject of the news feed, commentary, or fan account.
- **Name:** The profile name should not be the trademarked name of the company or include the trademarked name in a misleading manner.
- **Bio:** The bio should include a statement to distinguish it from the real company, such as "Unofficial Account," "Fan Account," or "Not affiliated with..."
- **Profile photo, header photo, or background image:** The account should not use another's trademark, logo or other copyright-protected image without express permission.
- **Communication with other users:** The account should not, through private or public communication with other users, try to deceive or mislead others about your identity.

Users may also choose to use different language to indicate that an account is not associated with the actual brand/company/product so long as it is clear and not confusing to others, and does not mislead or deceive.

If an account is reported to be confusing, we may request that the account holder make further changes to bring the account in compliance with these best practices.

How Can I Make My Own Account's Brand or Business Affiliation Clear?

We strongly recommend that you use all of Twitter's account settings (account name, location, web, and bio) to make your account's affiliation clear.

- Please see our [How to Change Your Profile Picture or Information](#) page for instructions on customizing your account. In particular, we recommend clearly stating your location,

- including your website if you have one, and clearly describing your brand or business in the bio, if applicable.
- You can also consider linking your official website to your Twitter account. [Twitter Buttons](#) are an easy way to add Twitter actions like “Follow” and “Mention” directly to your website.

What Information is Required When Reporting Trademark Policy Violations?

In order to investigate trademark policy violations, please provide all of the following information:

Username of the reported account (e.g., @safety or <http://www.twitter.com/safety>):

Your company name:

Your company Twitter account (if there is one):

Company website:

Your trademarked word, symbol, etc. (e.g. Twitter):

Trademark registration number:

Trademark registration office (e.g., USPTO):

Note: A federal or international trademark registration number is required. If the name you are reporting is not a registered mark (e.g., a government agency or non-profit organization), please let us know:

Your first and last name:

Title:

Address:

Phone:

Fax:

Email (must be from company domain):

Description of confusion (e.g., passing off as your company, including specific descriptions of content or behavior):

Requested Action (e.g., removal of violating account or transfer of trademarked username to an existing company account):

How Do I Report a Trademark Policy Violation?

You do not need a Twitter account to submit a trademark report. Holders of registered trademarks can report possible violations to Twitter's Policy Team through our [support forms](#).

Please submit trademark-related requests from your company email address and follow the format above to help expedite our response. Also, be sure to clearly describe to us why the account or Tweets may cause confusion with your mark.

Private information posted on Twitter (20130821*)

Posting another person's private and confidential information is a violation of the [Twitter Rules](#).

Some examples of private and confidential information include:

- credit card information
- social security or other national identity numbers
- addresses or locations that are considered and treated as private
- non-public, personal phone numbers
- non-public, personal email addresses
- images or videos that are considered and treated as private under applicable laws

Keep in mind that although you may consider certain information to be private, not all postings of such information may be a violation of this policy. We may consider the context and nature of the information posted, local privacy laws, and other case-specific facts when determining if this policy has been violated. For example, if information was previously posted or displayed elsewhere on the Internet prior to being put on Twitter, it may not be a violation of this policy.

If the private information you're reporting is not your own, please note that you must have documentation that you're authorized to act on behalf of the person whose confidential information is posted.

For frequently asked questions about reporting private information posted on Twitter, [click here](#). To learn more about protecting your private information on Twitter and other websites, [click here](#).

Abusive behavior policy (20140729*)

If you need to report abusive behavior to Twitter, please file a report [here](#).

If you believe you may be in danger, please contact your local law enforcement authority in addition to reporting

the content to Twitter so that the situation can be addressed both online and offline.

User disputes and false statements

Twitter provides a global communication platform which encompasses a variety of users with different voices, ideas and perspectives. As a policy, we do not mediate content or intervene in disputes between users.

Threats and abuse

Users may not make direct, specific threats of violence against others, including threats against a person or group on the basis of race, ethnicity, national origin, religion, sexual orientation, gender, gender identity, age, or disability. Targeted abuse or harassment is also a violation of the [Twitter Rules](#) and [Terms of Service](#).

For frequently asked questions about reporting abusive behavior on Twitter, [click here](#). To learn more about what you can do when you encounter abusive behavior on Twitter and other websites, [click here](#).

Offensive content

Users are allowed to post content, including potentially inflammatory content, provided they do not violate the [Twitter Rules](#) and [Terms of Service](#). Twitter does not screen content and does not remove potentially offensive content unless such content is in violation of the [Twitter Rules](#) and [Terms of Service](#).

If you believe the content or behavior you are reporting is prohibited in your local jurisdiction, please contact your local authorities so they can accurately assess the content or behavior for possible violations of local law. If Twitter is contacted directly by law enforcement, we can work with them and provide assistance for their investigation as well as guidance around possible options. You can point local law enforcement to our [Law Enforcement Guidelines](#).

Copyright and DMCA policy (20130430*)

Topics Covered:

- [What is the DMCA ?](#)
- [What Types of Copyright Complaints Does Twitter Respond To?](#)
- [Am I a Copyright Holder? How Do I Know?](#)
- [What Information Do You Need To Process a Copyright Complaint?](#)
- [How Do I File a Copyright Complaint?](#)
- [How Are Claims Processed?](#)
- [What Happens Next?](#)
- [What Happens If I Receive a DMCA Notification?](#)
- [When Should I File a Counter-Notice?](#)
- [What Information Do You Need to Process a Counter-notice?](#)
- [What Happens After I Submit a Counter-notice?](#)
- [Filing a DMCA Notification or Counter-notice Is Serious Business!](#)

What is the DMCA?

[Section 512](#) of the Digital Millennium Copyright Act (“DMCA”) outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

[What Types of Copyright Complaints Does Twitter Respond To?](#)

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile photo, header photo, or background, allegations concerning the unauthorized use of a copyrighted image uploaded through our photo hosting service,

or Tweets containing links to allegedly infringing materials.

Tip: If you are concerned about the use of your brand or entity's name, please review Twitter's [Trademark Policy](#). If you are concerned about the use of a fictional character, please see our [Parody, Commentary, and Fan Accounts Policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <http://chillingeffects.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

Tip: If you are reporting the content of a Tweet, please give us a direct link to that Tweet following these instructions: <https://support.twitter.com/articles/80586>. Or please specify if the alleged infringement is in the background, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to Twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate DMCA notifications may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with access to the reporter's complaint along with instructions on how to file a counter-notice, and forward a copy of the complaint to [Chilling Effects](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a copy of the takedown notice, along with instructions for filing a counter-notification.

Tip: If you've not yet received a copy of the DMCA notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a copy of each DMCA notification and counter-notice that we process to [Chilling Effects](#), where they are posted to a public-facing website (with your personal information removed).

Withheld Tweet:



Withheld Media:



Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a DMCA Notification?

If you receive a DMCA notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to

read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

If you believe that the material reported in the DMCA notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Tip: Re-posting material removed in response to a DMCA notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the DMCA notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

Tip: We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a DMCA Notification or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Targeted abuse: Same as "Violence & Threats/Abusive behavior policy"

Username squatting policy (20130817*)

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Following rules and best practices (20131209*)

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or un-following behavior. Please review our [Automation Rules and Best Practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and un-follow them. Aggressive follow churn is when an account *repeatedly* follows and then un-follows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 2000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 2000 people total**. Once you've followed 2000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

Tip: If you need to communicate with someone but don't need to see their updates every day, don't follow them. Visit their profile or send them an @reply when you need to; sending @replies doesn't require following, and your reply will appear in the person's @mentions tab so they can reply back.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

The Twitter Rules addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn’t a race to get the most followers. If you follow users that you’re interested in and post meaningful content, it’s more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Having trouble?

Check out our [Following Troubleshooting](#) section for solutions to common problems.

Automation Rules & Best Practices (2010/01)

Submitted Nov 08, 2009 by ginger

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.

- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

The Twitter Rules

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

- **Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others.
- **Trademark:** We reserve the right to reclaim usernames on behalf of businesses or individuals that hold legal claim or trademark on those usernames. Accounts using business names and/or logos to mislead others may be permanently suspended.
- **Private information:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission. You may not post intimate photos or videos that were taken or distributed without the subject's consent.
- **Violence and Threats:** You may not publish or post threats of violence against others or promote violence against others.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Unlawful Use:** You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter Badges:** You may not use badges, such as but not limited to the Promoted or Verified Twitter badge, unless provided by Twitter. Accounts using these badges as part of profile photos, header photos, background images, or in a way that falsely implies affiliation with Twitter may be suspended.

Abuse and Spam

Twitter strives to protect its users from abuse and spam. User abuse and technical abuse are not tolerated on Twitter.com, and may result in permanent suspension. Any accounts engaging in the activities specified below may be subject to permanent suspension.

- **Serial Accounts:** You may not create multiple accounts for disruptive or abusive purposes, or with overlapping use cases. Mass account creation may result in suspension of all related accounts. Please note that any violation of the Twitter Rules is cause for permanent suspension of all accounts.
- **Targeted Abuse:** You may not engage in targeted abuse or harassment. Some of the factors that we take into account when determining what conduct is considered to be targeted abuse or harassment are:
 - if you are sending messages to a user from multiple accounts;
 - if the sole purpose of your account is to send abusive messages to others;
 - if the reported behavior is one-sided or includes threats
- **Username Squatting:** You may not engage in username squatting. Accounts that are inactive for more than six months may also be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be username squatting are:
 - the number of accounts created
 - creating accounts for the purpose of preventing others from using those account names
 - creating accounts for the purpose of selling those accounts
 - using feeds of third-party content to update and maintain accounts under the names of those third parties
- **Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations.
- **Selling usernames:** You may not buy or sell Twitter usernames.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - If you have followed and/or unfollowed large amounts of users in a short time period, particularly by automated means (aggressive following or follower churn);
 - If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - If your updates consist mainly of links, and not personal updates;
 - If a large number of people are blocking you;
 - If a large number of spam complaints have been filed against you;

- If you post duplicate content over multiple accounts or multiple duplicate updates on one account;
- If you post multiple unrelated updates to a topic using #, trending or popular topic, or promoted trend;
- If you send large numbers of duplicate @replies or mentions;
- If you send large numbers of unsolicited @replies or mentions;
- If you add a large number of unrelated users to lists;
- If you repeatedly create false or misleading content;
- Randomly or aggressively following, favoriting or Retweeting Tweets;
- If you repeatedly post other users' account information as your own (bio, Tweets, url, etc.);
- If you post misleading links (e.g. affiliate links, links to malware/click jacking pages, etc.);
- Creating misleading accounts or account interactions;
- Selling or purchasing account interactions (such as selling or purchasing followers, Retweets, favorites, etc.);
- Using or promoting third-party services or apps that claim to get you more followers (such as follower trains, sites promising "more followers fast" or any other site that offers to automatically add followers to your account);
- **Graphic Content:** You may not use pornographic or excessively violent media in your profile image, header image, or background image.

Your account may be suspended for Terms of Service violations if any of the above is true. Please see our help pages on [Following rules and best practices](#) and [Automation rules and best practices](#) for a more detailed discussion of how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the [Terms of Service](#).

We may revise these Rules from time to time; the most current version will always be at twitter.com/rules.

Have Questions?

Check out our complete list of articles outlining our [policies, guidelines, and best practices](#).

To report an account for violation of the Twitter Rules, [please use our forms](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter Impersonation Policy.

An account will not be removed if

- the user shares your name but has no other commonalities, or
- the profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same background or avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to [Twitter's Parody Policy](#) for more information about these types of accounts.

For Frequently Asked Questions about reporting impersonation on Twitter, [click here](#).

Trademark policy

Updated October 3, 2012

What is a Trademark Policy Violation on Twitter?

Using a company or business name, logo, or other trademark-protected materials in a manner that may mislead or confuse others with regard to its brand or business affiliation may be considered a trademark policy violation.

How Does Twitter Respond To Reported Trademark Policy Violations?

When we receive reports of trademark policy violations from holders of federal or international trademark registrations, we review the account and may take the following actions:

- When there is a clear intent to mislead others through the unauthorized use of a trademark, Twitter will suspend the account and notify the account holder.
- When we determine that an account appears to be confusing users, but is not purposefully passing itself off as the trademarked good or service, we give the account holder an opportunity to clear up any potential confusion. We may also release a username for the trademark holder's active use.
- We are responsive to reports about confusing or misleading [Promoted Tweet](#) and [Promoted Trend](#) copy, as well as [Promoted Account](#) profile information. For more information, please visit our [Trademark Policy for Promoted Products](#).

What Is not a Trademark Policy Violation?

Using another's trademark in a way that has nothing to do with the product or service for which the trademark was granted is not a violation of Twitter's trademark policy.

- Twitter usernames are provided on a first-come, first-served basis and may not be reserved. For information on why you may not be able to select a certain username, please see our [Why Can't I Register Certain Usernames](#) help page.
- Search and interest keywords advertisers choose are subject to applicable trademark law. Although Twitter may review reports of confusing keywords, we may not disable keywords in every case. For more information about our policies specific to Promoted Products, please see our [Promoted Products Policy](#) and our [Trademark Policy for Promoted Products](#).

Guidelines for News feed, Commentary, and Fan Accounts

Twitter users are allowed to create news feed, commentary, and fan accounts. Twitter provides a platform for its users to share and receive a wide range of ideas and content, and we greatly value and respect our users' expression. Because of these principles, we do not actively monitor users' content and will not edit or remove user content, except in cases of violations of our [Terms of Service](#). An account's profile information should make it clear that the account is not actually the company or business entity that is the subject of the news feed/commentary/fan account. Here are some suggestions for distinguishing your account:

- **Username:** The username should not be the trademarked name of the subject of the news feed, commentary, or fan account.
- **Name:** The profile name should not be the trademarked name of the company or include the trademarked name in a misleading manner.
- **Bio:** The bio should include a statement to distinguish it from the real company, such as "Unofficial Account," "Fan Account," or "Not affiliated with..."
- **Profile photo, header photo, or background image:** The account should not use another's trademark, logo or other copyright-protected image without express permission.
- **Communication with other users:** The account should not, through private or public communication with other users, try to deceive or mislead others about your identity.

Users may also choose to use different language to indicate that an account is not associated with the actual brand/company/product so long as it is clear and not confusing to others, and does not mislead or deceive.

If an account is reported to be confusing, we may request that the account holder make further changes to bring the account in compliance with these best practices.

How Can I Make My Own Account's Brand or Business Affiliation Clear?

We strongly recommend that you use all of Twitter's account settings (account name, location, web, and bio) to make your account's affiliation clear.

- Please see our [How to Change Your Profile Picture or Information](#) page for instructions on customizing your account. In particular, we recommend clearly stating your location,

- including your website if you have one, and clearly describing your brand or business in the bio, if applicable.
- You can also consider linking your official website to your Twitter account. [Twitter Buttons](#) are an easy way to add Twitter actions like “Follow” and “Mention” directly to your website.

What Information is Required When Reporting Trademark Policy Violations?

In order to investigate trademark policy violations, please provide all of the following information:

Username of the reported account (e.g., @safety or <http://www.twitter.com/safety>):

Your company name:

Your company Twitter account (if there is one):

Company website:

Your trademarked word, symbol, etc. (e.g. Twitter):

Trademark registration number:

Trademark registration office (e.g., USPTO):

Note: A federal or international trademark registration number is required. If the name you are reporting is not a registered mark (e.g., a government agency or non-profit organization), please let us know:

Your first and last name:

Title:

Address:

Phone:

Fax:

Email (must be from company domain):

Description of confusion (e.g., passing off as your company, including specific descriptions of content or behavior):

Requested Action (e.g., removal of violating account or transfer of trademarked username to an existing company account):

How Do I Report a Trademark Policy Violation?

You do not need a Twitter account to submit a trademark report. Holders of registered trademarks can report possible violations to Twitter's Policy Team through our [support forms](#).

Please submit trademark-related requests from your company email address and follow the format above to help expedite our response. Also, be sure to clearly describe to us why the account or Tweets may cause confusion with your mark.

Private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#).

Some examples of private and confidential information include:

- credit card information
- social security or other national identity numbers
- addresses or locations that are considered and treated as private
- non-public, personal phone numbers
- non-public, personal email addresses
- images or videos that are considered and treated as private under applicable laws

Keep in mind that although you may consider certain information to be private, not all postings of such information may be a violation of this policy. We may consider the context and nature of the information posted, local privacy laws, and other case-specific facts when determining if this policy has been violated. For example, if information was previously posted or displayed elsewhere on the Internet prior to being put on Twitter, it may not be a violation of this policy.

If the private information you're reporting is not your own, please note that you must have documentation that you're authorized to act on behalf of the person whose confidential information is posted.

For frequently asked questions about reporting private information posted on Twitter, [click here](#). To learn more about protecting your private information on Twitter and other websites, [click here](#).

Abusive behavior policy

If you need to report abusive behavior to Twitter, please file a report [here](#).

If you believe you may be in danger, please contact your local law enforcement authority in addition to reporting

the content to Twitter so that the situation can be addressed both online and offline.

User disputes and false statements

Twitter provides a global communication platform which encompasses a variety of users with different voices, ideas and perspectives. As a policy, we do not mediate content or intervene in disputes between users.

Threats and abuse

Users may not make threats of violence or promote violence, including threatening or promoting terrorism. Users also may not make threats or promote violence against a person or group on the basis of race, ethnicity, national origin, religion, sexual orientation, gender, gender identity, age, or disability. In addition, users may not post intimate photos or videos that were taken or distributed without the subject's consent. Targeted abuse or harassment is also a violation of the [Twitter Rules](#) and [Terms of Service](#).

For frequently asked questions about reporting abusive behavior on Twitter, [click here](#). To learn more about what you can do when you encounter abusive behavior on Twitter and other websites, [click here](#).

Offensive content

Users are allowed to post content, including potentially inflammatory content, provided they do not violate the [Twitter Rules](#) and [Terms of Service](#). Twitter does not screen content and does not remove potentially offensive content unless such content is in violation of the [Twitter Rules](#) and [Terms of Service](#).

If you believe the content or behavior you are reporting is prohibited in your local jurisdiction, please contact your local authorities so they can accurately assess the content or behavior for possible violations of local law. If Twitter is contacted directly by law enforcement, we can work with them and provide assistance for their investigation as well as guidance around possible options. You can point local law enforcement to our [Law Enforcement Guidelines](#).

Copyright and DMCA policy

Topics Covered:

- [What is the DMCA ?](#)
- [What Types of Copyright Complaints Does Twitter Respond To?](#)
- [Am I a Copyright Holder? How Do I Know?](#)
- [What Information Do You Need To Process a Copyright Complaint?](#)
- [How Do I File a Copyright Complaint?](#)
- [How Are Claims Processed?](#)
- [What Happens Next?](#)
- [What Happens If I Receive a DMCA Notification?](#)
- [When Should I File a Counter-Notice?](#)
- [What Information Do You Need to Process a Counter-notice?](#)
- [What Happens After I Submit a Counter-notice?](#)
- [Filing a DMCA Notification or Counter-notice Is Serious Business!](#)

What is the DMCA?

[Section 512](#) of the Digital Millennium Copyright Act (“DMCA”) outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

What Types of Copyright Complaints Does Twitter Respond To?

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as an profile photo, header photo, or background, allegations concerning the unauthorized use of a copyrighted image uploaded through our photo hosting service,

or Tweets containing links to allegedly infringing materials.

Tip: If you are concerned about the use of your brand or entity's name, please review Twitter's [Trademark Policy](#). If you are concerned about the use of a fictional character, please see our [Parody, Commentary, and Fan Accounts Policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <http://chillingeffects.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

Tip: If you are reporting the content of a Tweet, please give us a direct link to that Tweet following these instructions: <https://support.twitter.com/articles/80586>. Or please specify if the alleged infringement is in the background, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to Twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate DMCA notifications may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with access to the reporter's complaint along with instructions on how to file a counter-notice, and forward a copy of the complaint to [Chilling Effects](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a copy of the takedown notice, along with instructions for filing a counter-notification.

Tip: If you've not yet received a copy of the DMCA notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a copy of each DMCA notification and counter-notice that we process to [Chilling Effects](#), where they are posted to a public-facing website (with your personal information removed).

Withheld Tweet:



Withheld Media:



Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a DMCA Notification?

If you receive a DMCA notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to

read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

If you believe that the material reported in the DMCA notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Tip: Re-posting material removed in response to a DMCA notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the DMCA notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

Tip: We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a DMCA Notification or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Targeted abuse: Same as "Violence & Threats/Abusive behavior policy"

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or un-following behavior. Please review our [Automation Rules and Best Practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and un-follow them. Aggressive follow churn is when an account *repeatedly* follows and then un-follows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 2000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 2000 people total**. Once you've followed 2000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

Tip: If you need to communicate with someone but don't need to see their updates every day, don't follow them. Visit their profile or send them an @reply when you need to; sending @replies doesn't require following, and your reply will appear in the person's @mentions tab so they can reply back.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

The Twitter Rules addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn’t a race to get the most followers. If you follow users that you’re interested in and post meaningful content, it’s more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Having trouble?

Check out our [Following Troubleshooting](#) section for solutions to common problems.

Automation Rules & Best Practices

Submitted Nov 08, 2009 by ginger

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.

- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

The Twitter Rules

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

- **Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others.
- **Trademark:** We reserve the right to reclaim usernames on behalf of businesses or individuals that hold legal claim or trademark on those usernames. Accounts using business names and/or logos to mislead others may be permanently suspended.
- **Private information:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission. You may not post intimate photos or videos that were taken or distributed without the subject's consent.
- **Violence and Threats:** You may not publish or post threats of violence against others or promote violence against others.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Unlawful Use:** You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter Badges:** You may not use badges, such as but not limited to the Promoted or Verified Twitter badge, unless provided by Twitter. Accounts using these badges as part of profile photos, header photos, background images, or in a way that falsely implies affiliation with Twitter may be suspended.

Abuse and Spam

Twitter strives to protect its users from abuse and spam. User abuse and technical abuse are not tolerated on Twitter.com, and may result in permanent suspension. Any accounts engaging in the activities specified below may be subject to permanent suspension.

- **Serial Accounts:** You may not create multiple accounts for disruptive or abusive purposes, or with overlapping use cases. Mass account creation may result in suspension of all related accounts. Please note that any violation of the Twitter Rules is cause for permanent suspension of all accounts.
- **Targeted Abuse:** You may not engage in targeted abuse or harassment. Some of the factors that we take into account when determining what conduct is considered to be targeted abuse or harassment are:
 - if you are sending messages to a user from multiple accounts;
 - if the sole purpose of your account is to send abusive messages to others;
 - if the reported behavior is one-sided or includes threats
- **Username Squatting:** You may not engage in username squatting. Accounts that are inactive for more than six months may also be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be username squatting are:
 - the number of accounts created
 - creating accounts for the purpose of preventing others from using those account names
 - creating accounts for the purpose of selling those accounts
 - using feeds of third-party content to update and maintain accounts under the names of those third parties
- **Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations.
- **Selling usernames:** You may not buy or sell Twitter usernames.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - If you have followed and/or unfollowed large amounts of users in a short time period, particularly by automated means (aggressive following or follower churn);
 - If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - If your updates consist mainly of links, and not personal updates;
 - If a large number of people are blocking you;
 - If a large number of spam complaints have been filed against you;

- If you post duplicate content over multiple accounts or multiple duplicate updates on one account;
- If you post multiple unrelated updates to a topic using #, trending or popular topic, or promoted trend;
- If you send large numbers of duplicate @replies or mentions;
- If you send large numbers of unsolicited @replies or mentions;
- If you add a large number of unrelated users to lists;
- If you repeatedly create false or misleading content;
- Randomly or aggressively following, favoriting or Retweeting Tweets;
- If you repeatedly post other users' account information as your own (bio, Tweets, url, etc.);
- If you post misleading links (e.g. affiliate links, links to malware/click jacking pages, etc.);
- Creating misleading accounts or account interactions;
- Selling or purchasing account interactions (such as selling or purchasing followers, Retweets, favorites, etc.);
- Using or promoting third-party services or apps that claim to get you more followers (such as follower trains, sites promising "more followers fast" or any other site that offers to automatically add followers to your account);
- **Graphic Content:** You may not use pornographic or excessively violent media in your profile image, header image, or background image.

Your account may be suspended for Terms of Service violations if any of the above is true. Please see our help pages on [Following rules and best practices](#) and [Automation rules and best practices](#) for a more detailed discussion of how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the [Terms of Service](#).

We may revise these Rules from time to time; the most current version will always be at twitter.com/rules.

Have Questions?

Check out our complete list of articles outlining our [policies, guidelines, and best practices](#).

To report an account for violation of the Twitter Rules, [please use our forms](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter Impersonation Policy.

An account will not be removed if

- the user shares your name but has no other commonalities, or
- the profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same background or avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to [Twitter's Parody Policy](#) for more information about these types of accounts.

For Frequently Asked Questions about reporting impersonation on Twitter, [click here](#).

Trademark policy

Updated October 3, 2012

What is a Trademark Policy Violation on Twitter?

Using a company or business name, logo, or other trademark-protected materials in a manner that may mislead or confuse others with regard to its brand or business affiliation may be considered a trademark policy violation.

How Does Twitter Respond To Reported Trademark Policy Violations?

When we receive reports of trademark policy violations from holders of federal or international trademark registrations, we review the account and may take the following actions:

- When there is a clear intent to mislead others through the unauthorized use of a trademark, Twitter will suspend the account and notify the account holder.
- When we determine that an account appears to be confusing users, but is not purposefully passing itself off as the trademarked good or service, we give the account holder an opportunity to clear up any potential confusion. We may also release a username for the trademark holder's active use.
- We are responsive to reports about confusing or misleading [Promoted Tweet](#) and [Promoted Trend](#) copy, as well as [Promoted Account](#) profile information. For more information, please visit our [Trademark Policy for Promoted Products](#).

What Is not a Trademark Policy Violation?

Using another's trademark in a way that has nothing to do with the product or service for which the trademark was granted is not a violation of Twitter's trademark policy.

- Twitter usernames are provided on a first-come, first-served basis and may not be reserved. For information on why you may not be able to select a certain username, please see our [Why Can't I Register Certain Usernames](#) help page.
- Search and interest keywords advertisers choose are subject to applicable trademark law. Although Twitter may review reports of confusing keywords, we may not disable keywords in every case. For more information about our policies specific to Promoted Products, please see our [Promoted Products Policy](#) and our [Trademark Policy for Promoted Products](#).

Guidelines for News feed, Commentary, and Fan Accounts

Twitter users are allowed to create news feed, commentary, and fan accounts. Twitter provides a platform for its users to share and receive a wide range of ideas and content, and we greatly value and respect our users' expression. Because of these principles, we do not actively monitor users' content and will not edit or remove user content, except in cases of violations of our [Terms of Service](#). An account's profile information should make it clear that the account is not actually the company or business entity that is the subject of the news feed/commentary/fan account. Here are some suggestions for distinguishing your account:

- **Username:** The username should not be the trademarked name of the subject of the news feed, commentary, or fan account.
- **Name:** The profile name should not be the trademarked name of the company or include the trademarked name in a misleading manner.
- **Bio:** The bio should include a statement to distinguish it from the real company, such as "Unofficial Account," "Fan Account," or "Not affiliated with..."
- **Profile photo, header photo, or background image:** The account should not use another's trademark, logo or other copyright-protected image without express permission.
- **Communication with other users:** The account should not, through private or public communication with other users, try to deceive or mislead others about your identity.

Users may also choose to use different language to indicate that an account is not associated with the actual brand/company/product so long as it is clear and not confusing to others, and does not mislead or deceive.

If an account is reported to be confusing, we may request that the account holder make further changes to bring the account in compliance with these best practices.

How Can I Make My Own Account's Brand or Business Affiliation Clear?

We strongly recommend that you use all of Twitter's account settings (account name, location, web, and bio) to make your account's affiliation clear.

- Please see our [How to Change Your Profile Picture or Information](#) page for instructions on customizing your account. In particular, we recommend clearly stating your location,

- including your website if you have one, and clearly describing your brand or business in the bio, if applicable.
- You can also consider linking your official website to your Twitter account. [Twitter Buttons](#) are an easy way to add Twitter actions like “Follow” and “Mention” directly to your website.

What Information is Required When Reporting Trademark Policy Violations?

In order to investigate trademark policy violations, please provide all of the following information:

Username of the reported account (e.g., @safety or <http://www.twitter.com/safety>):

Your company name:

Your company Twitter account (if there is one):

Company website:

Your trademarked word, symbol, etc. (e.g. Twitter):

Trademark registration number:

Trademark registration office (e.g., USPTO):

Note: A federal or international trademark registration number is required. If the name you are reporting is not a registered mark (e.g., a government agency or non-profit organization), please let us know:

Your first and last name:

Title:

Address:

Phone:

Fax:

Email (must be from company domain):

Description of confusion (e.g., passing off as your company, including specific descriptions of content or behavior):

Requested Action (e.g., removal of violating account or transfer of trademarked username to an existing company account):

How Do I Report a Trademark Policy Violation?

You do not need a Twitter account to submit a trademark report. Holders of registered trademarks can report possible violations to Twitter's Policy Team through our [support forms](#).

Please submit trademark-related requests from your company email address and follow the format above to help expedite our response. Also, be sure to clearly describe to us why the account or Tweets may cause confusion with your mark.

Private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Some examples of private and confidential information include:

- credit card information
- social security or other national identity numbers
- addresses or locations that are considered and treated as private
- non-public, personal phone numbers
- non-public, personal email addresses
- images or videos that are considered and treated as private under applicable laws intimate photos or videos that were taken or distributed without the subject's consent

Keep in mind that although you may consider certain information to be private, not all postings of such information may be a violation of this policy. We may consider the context and nature of the information posted, local privacy laws, and other case-specific facts when determining if this policy has been violated. For example, if information was previously posted or displayed elsewhere on the Internet prior to being put on Twitter, it may not be a violation of this policy.

For frequently asked questions about reporting private information posted on Twitter, [click here](#). To learn more about protecting your private information on Twitter and other websites, [click here](#).

Abusive behavior policy

If you need to report abusive behavior to Twitter, please file a report [here](#).

If you believe you may be in danger, please contact your local law enforcement authority in addition to reporting

the content to Twitter so that the situation can be addressed both online and offline.

User disputes and false statements

Twitter provides a global communication platform which encompasses a variety of users with different voices, ideas and perspectives. As a policy, we do not mediate content or intervene in disputes between users.

Threats and abuse

Users may not make threats of violence or promote violence, including threatening or promoting terrorism. Users also may not make threats or promote violence against a person or group on the basis of race, ethnicity, national origin, religion, sexual orientation, gender, gender identity, age, or disability. In addition, users may not post intimate photos or videos

that were taken or distributed without the subject's consent. Targeted abuse or harassment is also a violation of the [Twitter Rules](#) and [Terms of Service](#).

For frequently asked questions about reporting abusive behavior on Twitter, [click here](#). To learn more about what you can do when you encounter abusive behavior on Twitter and other websites, [click here](#).

Offensive content

Users are allowed to post content, including potentially inflammatory content, provided they do not violate the [Twitter Rules](#) and [Terms of Service](#). Twitter does not screen content and does not remove potentially offensive content unless such content is in violation of the [Twitter Rules](#) and [Terms of Service](#).

If you believe the content or behavior you are reporting is prohibited in your local jurisdiction, please contact your local authorities so they can accurately assess the content or behavior for possible violations of local law. If Twitter is contacted directly by law enforcement, we can work with them and provide assistance for their investigation as well as guidance around possible options. You can point local law enforcement to our [Law Enforcement Guidelines](#).

Copyright and DMCA policy

Topics Covered:

- [What is the DMCA ?](#)
- [What Types of Copyright Complaints Does Twitter Respond To?](#)
- [Am I a Copyright Holder? How Do I Know?](#)
- [What Information Do You Need To Process a Copyright Complaint?](#)
- [How Do I File a Copyright Complaint?](#)
- [How Are Claims Processed?](#)
- [What Happens Next?](#)
- [What Happens If I Receive a DMCA Notification?](#)
- [When Should I File a Counter-Notice?](#)
- [What Information Do You Need to Process a Counter-notice?](#)
- [What Happens After I Submit a Counter-notice?](#)
- [Filing a DMCA Notification or Counter-notice Is Serious Business!](#)

What is the DMCA?

[Section 512](#) of the Digital Millennium Copyright Act (“DMCA”) outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

What Types of Copyright Complaints Does Twitter Respond To?

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile photo, header photo, or background, allegations concerning the unauthorized use of a copyrighted image uploaded through our photo hosting service, or Tweets containing links to allegedly infringing materials.

Tip: If you are concerned about the use of your brand or entity's name, please review

Twitter's [Trademark Policy](#). If you are concerned about the use of a fictional character, please see our [Parody, Commentary, and Fan Accounts Policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <http://chillingeffects.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

Tip: If you are reporting the content of a Tweet, please give us a direct link to that Tweet following these instructions: <https://support.twitter.com/articles/80586>. Or please specify if the alleged infringement is in the background, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to Twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate DMCA notifications may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with access to the reporter's complaint

along with instructions on how to file a counter-notice, and forward a copy of the complaint to [Chilling Effects](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a copy of the takedown notice, along with instructions for filing a counter-notification.

Tip: If you've not yet received a copy of the DMCA notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a copy of each DMCA notification and counter-notice that we process to [Chilling Effects](#), where they are posted to a public-facing website (with your personal information removed).

Withheld Tweet:

Tweet withheld 18h
This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)

Withheld Media:

 **Media not displayed**

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a DMCA Notification?

If you receive a DMCA notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

If you believe that the material reported in the DMCA notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Tip: Re-posting material removed in response to a DMCA notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the DMCA notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

Tip: We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a DMCA Notification or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Targeted abuse: Same as “Violence & Threats/Abusive behavior policy”

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or un-following behavior. Please review our [Automation Rules and Best Practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and un-follow them. Aggressive follow churn is when an account *repeatedly* follows and then un-follows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 2000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 2000 people total**. Once you've followed 2000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this

a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

Tip: If you need to communicate with someone but don't need to see their updates every day, don't follow them. Visit their profile or send them an @reply when you need to; sending @replies doesn't require following, and your reply will appear in the person's @mentions tab so they can reply back.

What are third party applications to “Get More Followers

Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

The [Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Having trouble?

Check out our [Following Troubleshooting section](#) for solutions to common problems.

Automation Rules & Best Practices

Submitted Nov 08, 2009 by **ginger**

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.

3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of

users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

The Twitter Rules

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to share your content with others. We respect the ownership of the content that users share and each user is responsible for the content he or she provides. Because of these principles, we do not actively monitor and will not censor user content, except in limited circumstances described below.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter. These limitations comply with legal requirements and make Twitter a better experience for all. We may need to change these rules from time to time and reserve the right to do so. Please check back here to see the latest.

- **Impersonation:** You may not impersonate others through the Twitter service in a manner that does or is intended to mislead, confuse, or deceive others.
- **Trademark:** We reserve the right to reclaim usernames on behalf of businesses or individuals that hold legal claim or trademark on those usernames. Accounts using business names and/or logos to mislead others may be permanently suspended.
- **Private information:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address or Social Security/National Identity numbers, without their express authorization and permission. You may not post intimate photos or videos that were taken or distributed without the subject's consent.
- **Violence and Threats:** You may not publish or post threats of violence against others or promote violence against others.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Unlawful Use:** You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter Badges:** You may not use badges, such as but not limited to the Promoted or Verified Twitter badge, unless provided by Twitter. Accounts using these badges as part of profile photos, header photos, background images, or in a way that falsely implies affiliation with Twitter may be suspended.

Abuse and Spam

Twitter strives to protect its users from abuse and spam. User abuse and technical abuse are not tolerated on Twitter.com, and may result in permanent suspension. Any accounts engaging in the activities specified below may be subject to permanent suspension.

- **Serial Accounts:** You may not create multiple accounts for disruptive or abusive purposes, or with overlapping use cases. Mass account creation may result in suspension of all related accounts. Please note that any violation of the Twitter Rules is cause for permanent suspension of all accounts.
- **Targeted Abuse:** You may not engage in targeted abuse or harassment. Some of the factors that we take into account when determining what conduct is considered to be targeted abuse or harassment are:
 - if you are sending messages to a user from multiple accounts;
 - if the sole purpose of your account is to send abusive messages to others;
 - if the reported behavior is one-sided or includes threats
- **Username Squatting:** You may not engage in username squatting. Accounts that are inactive for more than six months may also be removed without further notice. Some of the factors that we take into account when determining what conduct is considered to be username squatting are:
 - the number of accounts created
 - creating accounts for the purpose of preventing others from using those account names
 - creating accounts for the purpose of selling those accounts
 - using feeds of third-party content to update and maintain accounts under the names of those third parties
- **Invitation spam:** You may not use Twitter.com's address book contact import to send repeat, mass invitations.
- **Selling usernames:** You may not buy or sell Twitter usernames.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another user's browser or computer or to compromise a user's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - If you have followed and/or unfollowed large amounts of users in a short time period, particularly by automated means (aggressive following or follower churn);
 - If you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - If your updates consist mainly of links, and not personal updates;
 - If a large number of people are blocking you;
 - If a large number of spam complaints have been filed against you;

- If you post duplicate content over multiple accounts or multiple duplicate updates on one account;
- If you post multiple unrelated updates to a topic using #, trending or popular topic, or promoted trend;
- If you send large numbers of duplicate @replies or mentions;
- If you send large numbers of unsolicited @replies or mentions;
- If you add a large number of unrelated users to lists;
- If you repeatedly create false or misleading content;
- Randomly or aggressively following, favoriting or Retweeting Tweets;
- If you repeatedly post other users' account information as your own (bio, Tweets, url, etc.);
- If you post misleading links (e.g. affiliate links, links to malware/click jacking pages, etc.);
- Creating misleading accounts or account interactions;
- Selling or purchasing account interactions (such as selling or purchasing followers, Retweets, favorites, etc.);
- Using or promoting third-party services or apps that claim to get you more followers (such as follower trains, sites promising "more followers fast" or any other site that offers to automatically add followers to your account);
- **Graphic Content:** You may not use pornographic or excessively violent media in your profile image, header image, or background image.

Your account may be suspended for Terms of Service violations if any of the above is true. Please see our help pages on [Following rules and best practices](#) and [Automation rules and best practices](#) for a more detailed discussion of how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts engaging in any of these behaviors may be investigated for abuse. Accounts under investigation may be removed from Search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the [Terms of Service](#).

We may revise these Rules from time to time; the most current version will always be at twitter.com/rules.

Have Questions?

Check out our complete list of articles outlining our [policies, guidelines, and best practices](#).

To report an account for violation of the Twitter Rules, [please use our forms](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter Impersonation Policy.

An account will not be removed if

- the user shares your name but has no other commonalities, or
- the profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same background or avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to [Twitter's Parody Policy](#) for more information about these types of accounts.

For Frequently Asked Questions about reporting impersonation on Twitter, [click here](#).

Trademark policy

Updated October 3, 2012

What is a Trademark Policy Violation on Twitter?

Using a company or business name, logo, or other trademark-protected materials in a manner that may mislead or confuse others with regard to its brand or business affiliation may be considered a trademark policy violation.

How Does Twitter Respond To Reported Trademark Policy Violations?

When we receive reports of trademark policy violations from holders of federal or international trademark registrations, we review the account and may take the following actions:

- When there is a clear intent to mislead others through the unauthorized use of a trademark, Twitter will suspend the account and notify the account holder.
- When we determine that an account appears to be confusing users, but is not purposefully passing itself off as the trademarked good or service, we give the account holder an opportunity to clear up any potential confusion. We may also release a username for the trademark holder's active use.
- We are responsive to reports about confusing or misleading [Promoted Tweet](#) and [Promoted Trend](#) copy, as well as [Promoted Account](#) profile information. For more information, please visit our [Trademark Policy for Promoted Products](#).

What Is not a Trademark Policy Violation?

Using another's trademark in a way that has nothing to do with the product or service for which the trademark was granted is not a violation of Twitter's trademark policy.

- Twitter usernames are provided on a first-come, first-served basis and may not be reserved. For information on why you may not be able to select a certain username, please see our [Why Can't I Register Certain Usernames](#) help page.
- Search and interest keywords advertisers choose are subject to applicable trademark law. Although Twitter may review reports of confusing keywords, we may not disable keywords in every case. For more information about our policies specific to Promoted Products, please see our [Promoted Products Policy](#) and our [Trademark Policy for Promoted Products](#).

Guidelines for News feed, Commentary, and Fan Accounts

Twitter users are allowed to create news feed, commentary, and fan accounts. Twitter provides a platform for its users to share and receive a wide range of ideas and content, and we greatly value and respect our users' expression. Because of these principles, we do not actively monitor users' content and will not edit or remove user content, except in cases of violations of our [Terms of Service](#). An account's profile information should make it clear that the account is not actually the company or business entity that is the subject of the news feed/commentary/fan account. Here are some suggestions for distinguishing your account:

- **Username:** The username should not be the trademarked name of the subject of the news feed, commentary, or fan account.
- **Name:** The profile name should not be the trademarked name of the company or include the trademarked name in a misleading manner.
- **Bio:** The bio should include a statement to distinguish it from the real company, such as "Unofficial Account," "Fan Account," or "Not affiliated with..."
- **Profile photo, header photo, or background image:** The account should not use another's trademark, logo or other copyright-protected image without express permission.
- **Communication with other users:** The account should not, through private or public communication with other users, try to deceive or mislead others about your identity.

Users may also choose to use different language to indicate that an account is not associated with the actual brand/company/product so long as it is clear and not confusing to others, and does not mislead or deceive.

If an account is reported to be confusing, we may request that the account holder make further changes to bring the account in compliance with these best practices.

How Can I Make My Own Account's Brand or Business Affiliation Clear?

We strongly recommend that you use all of Twitter's account settings (account name, location, web, and bio) to make your account's affiliation clear.

- Please see our [How to Change Your Profile Picture or Information](#) page for instructions on customizing your account. In particular, we recommend clearly stating your location,

- including your website if you have one, and clearly describing your brand or business in the bio, if applicable.
- You can also consider linking your official website to your Twitter account. [Twitter Buttons](#) are an easy way to add Twitter actions like “Follow” and “Mention” directly to your website.

What Information is Required When Reporting Trademark Policy Violations?

In order to investigate trademark policy violations, please provide all of the following information:

Username of the reported account (e.g., @safety or <http://www.twitter.com/safety>):

Your company name:

Your company Twitter account (if there is one):

Company website:

Your trademarked word, symbol, etc. (e.g. Twitter):

Trademark registration number:

Trademark registration office (e.g., USPTO):

Note: A federal or international trademark registration number is required. If the name you are reporting is not a registered mark (e.g., a government agency or non-profit organization), please let us know:

Your first and last name:

Title:

Address:

Phone:

Fax:

Email (must be from company domain):

Description of confusion (e.g., passing off as your company, including specific descriptions of content or behavior):

Requested Action (e.g., removal of violating account or transfer of trademarked username to an existing company account):

How Do I Report a Trademark Policy Violation?

You do not need a Twitter account to submit a trademark report. Holders of registered trademarks can report possible violations to Twitter's Policy Team through our [support forms](#).

Please submit trademark-related requests from your company email address and follow the format above to help expedite our response. Also, be sure to clearly describe to us why the account or Tweets may cause confusion with your mark.

Private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Some examples of private and confidential information include:

- credit card information
- social security or other national identity numbers
- addresses or locations that are considered and treated as private
- non-public, personal phone numbers
- non-public, personal email addresses
- images or videos that are considered and treated as private under applicable laws intimate
- photos or videos that were taken or distributed without the subject's consent

Keep in mind that although you may consider certain information to be private, not all postings of such information may be a violation of this policy. We may consider the context and nature of the information posted, local privacy laws, and other case-specific facts when determining if this policy has been violated. For example, if information was previously posted or displayed elsewhere on the Internet prior to being put on Twitter, it may not be a violation of this policy.

For frequently asked questions about reporting private information posted on Twitter, [click here](#). To learn more about protecting your private information on Twitter and other websites, [click here](#).

Abusive behavior policy

Engaging in targeted abuse or harassment on Twitter is a violation of the [Twitter Rules](#) and [Terms of Service](#), so we've established a policy that helps us evaluate abusive behavior in its many forms. Maintaining the diversity and openness of the Twitter community is our responsibility, and so is creating a safe environment for every user.

Abusive behavior

We take action on accounts found to be in violation of the [Twitter Rules](#) and [Terms of Service](#), including (but not limited to) the following types of abusive behavior:

Violent threats (direct or indirect): Users may not make threats of violence or promote violence, including threatening or promoting terrorism. Users also may not make threats or promote violence against a person or group on the basis of race, ethnicity, national origin, religion, sexual orientation, gender, gender identity, age, or disability.

Abuse and harassment: Users may not engage in targeted abuse or harassment. Some of the factors that we take into account when determining what conduct is considered to be targeted abuse or harassment are:

if a primary purpose of the reported account is to send abusive messages to others; if the reported behavior is one-sided or includes threats;
if the reported user is inciting others to harass another user; and
if the reported user is sending harassing messages to a user from multiple accounts.

Self-harm: Users may encounter someone considering suicide or self harm on Twitter. When we receive reports that another user is threatening suicide or self harm, we may take a number of steps to assist the reported user, such as reaching out to that user expressing our concern and the concern of other users on Twitter or providing resources such as contact information for our mental health partners.

Private information: You may not publish or post other people's private and confidential information, such as credit card numbers, street address, or Social Security/National Identity numbers, without their express authorization and permission. In addition, you may not post intimate photos or videos that were taken or distributed without the subject's consent. Read more about our private information policy [here](#).

Impersonation: You may not impersonate others through the Twitter service in a manner that is intended to or does mislead, confuse, or deceive others. Read more about our impersonation policy [here](#).

Please note that when a Tweet or account is reported, we may investigate the full account reported, including profile, media and other Tweets.

Offensive content and mediation

Offensive content is tolerated as long as it does not violate the [Twitter Rules](#) and [Terms of Service](#). Additionally, we do not mediate content or intervene in disputes between users.

Twitter provides a global communication platform which encompasses a variety of users with different voices, ideas and perspectives. Because of this diversity, you may encounter content you consider to be inflammatory or inappropriate that is not considered a violation of our rules. To help you control your experience, we've collected information about the tools you have that you can use to tailor your Twitter experience. Read about them in this [article](#).

Accounts found to be in violation

When an account violates our abusive behavior policy, the actions we take against that account depend on the severity of the violation. We may ask users to verify or provide information, delete specific Tweets, or we may suspend or lock the account temporarily or permanently.

Report abuse to us

If you encounter something abusive on Twitter, report it to us. For instructions about how to report abusive behavior, read this [article](#).

We've outlined information about how to report other types of violations (for example, copyright or trademark violations, spam, etc.) in this [article](#).

Copyright and DMCA policy

Topics Covered:

- [What is the DMCA ?](#)
- [What Types of Copyright Complaints Does Twitter Respond To?](#)
- [Am I a Copyright Holder? How Do I Know?](#)
- [What Information Do You Need To Process a Copyright Complaint?](#)
- [How Do I File a Copyright Complaint?](#)
- [How Are Claims Processed?](#)
- [What Happens Next?](#)
- [What Happens If I Receive a DMCA Notification?](#)
- [When Should I File a Counter-Notice?](#)
- [What Information Do You Need to Process a Counter-notice?](#)
- [What Happens After I Submit a Counter-notice?](#)
- [Filing a DMCA Notification or Counter-notice Is Serious Business!](#)

What is the DMCA?

[Section 512](#) of the Digital Millennium Copyright Act (“DMCA”) outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

What Types of Copyright Complaints Does Twitter Respond To?

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile photo, header photo, or background, allegations concerning the unauthorized use of a copyrighted image uploaded through our photo hosting service, or Tweets containing links to allegedly infringing materials.

Tip: If you are concerned about the use of your brand or entity’s name, please review Twitter’s [Trademark Policy](#). If you are concerned about the use of a fictional character, please see our [Parody, Commentary, and Fan Accounts Policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <http://chillingeffects.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);

3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

Tip: If you are reporting the content of a Tweet, please give us a direct link to that Tweet following these instructions: <https://support.twitter.com/articles/80586>. Or please specify if the alleged infringement is in the background, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to Twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate DMCA notifications may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with access to the reporter's complaint along with instructions on how to file a counter-notice, and forward a copy of the complaint to [Chilling Effects](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a copy of the takedown notice, along with instructions for filing a counter-notification.

Tip: If you've not yet received a copy of the DMCA notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a copy of each DMCA notification and counter-notice that we process to [Chilling Effects](#), where they are posted to a public-facing website (with your personal information removed).

Withheld Tweet:

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)

Withheld Media:

**Media not displayed**

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a DMCA Notification?

If you receive a DMCA notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

If you believe that the material reported in the DMCA notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Tip: Re-posting material removed in response to a DMCA notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the DMCA notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of Federal District Court for the judicial district in which the address is

located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

Tip: We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a DMCA Notification or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Targeted abuse: Same as "Violence & Threats/Abusive behavior policy"

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or un-following behavior. Please review our [Automation Rules and Best Practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and un-follow them. Aggressive follow churn is when an account *repeatedly* follows and then un-follows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 2000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 2000 people total**. Once you've followed 2000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

Tip: If you need to communicate with someone but don't need to see their updates every day, don't follow them. Visit their profile or send them an @reply when you need to; sending @replies doesn't require following, and your reply will appear in the person's @mentions tab so they can reply back.

What are third party applications to “Get More Followers

Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate.

Using these is not allowed according to the Twitter Rules.

The [Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Having trouble?

Check out our [Following Troubleshooting](#) section for solutions to common problems.

Automation Rules & Best Practices

Submitted Nov 08, 2009 by [ginger](#)

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find

out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social

action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for

the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. All users must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in the temporary locking and/or permanent suspension of account(s).

Please note that we may need to change these rules from time to time and reserve the right to do so. The most current version will always be available at twitter.com/rules.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter.

- **Trademark:** We reserve the right to reclaim usernames on behalf of businesses or individuals that hold legal claim or trademark on those usernames. Accounts using business names and/or logos to mislead others may be permanently suspended.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Graphic content:** You may not use pornographic or excessively violent media in your profile image, header image, or background image. Twitter may allow some forms of graphic content in Tweets marked as [sensitive media](#). When content crosses the line into gratuitous images of death, Twitter may ask that you remove the content out of respect for the deceased.
- **Unlawful use:** You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter badges:** You may not use badges, such as but not limited to the “promoted” or “verified” Twitter badge, unless provided by Twitter. Accounts using these badges as part of profile photos, header photos, background images, or in a way that falsely implies affiliation with Twitter may be suspended.

Abusive Behavior

We believe in freedom of expression and in speaking truth to power, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we do not tolerate behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Any accounts and related accounts engaging in the activities specified below may be temporarily locked and/or subject to permanent suspension.

- **Violent threats (direct or indirect):** You may not make threats of violence or promote violence, including threatening or promoting terrorism.

- **Harassment:** You may not incite or engage in the targeted abuse or harassment of others. Some of the factors that we may consider when evaluating abusive behavior include:
 - if a primary purpose of the reported account is to harass or send abusive messages to others;
 - if the reported behavior is one-sided or includes threats;
 - if the reported account is inciting others to harass another account; and
 - if the reported account is sending harassing messages to an account from multiple accounts.
- **Hateful conduct:** You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.
- **Multiple account abuse:** Creating multiple accounts with overlapping uses or in order to evade the temporary or permanent suspension of a separate account is not allowed.
- **Private information:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address, or Social Security/National Identity numbers, without their express authorization and permission. In addition, you may not post intimate photos or videos that were taken or distributed without the subject's consent. Read more about our private information policy [here](#).
- **Impersonation:** You may not impersonate others through the Twitter service in a manner that is intended to or does mislead, confuse, or deceive others. Read more about our impersonation policy [here](#).
- **Self-harm:** You may encounter someone considering suicide or self harm on Twitter. When we receive reports that a person is threatening suicide or self harm, we may take a number of steps to assist them, such as reaching out to that person expressing our concern and the concern of other users on Twitter or providing resources such as contact information for our mental health partners.

Spam

We strive to protect people on Twitter from technical abuse and spam. Any accounts engaging in the activities specified below may be temporarily locked or subject to permanent suspension.

- **Username squatting:** You may not engage in username squatting. Accounts that are inactive for more than six months may also be removed without further notice. Some of the factors we take into consideration when determining what conduct is considered to be username squatting are:
 - the number of accounts created;
 - creating accounts for the purpose of preventing others from using those account names;
 - creating accounts for the purpose of selling those accounts; and

- using feeds of third-party content to update and maintain accounts under the names of those third parties.
- **Invitation spam:** You may not use twitter.com's address book contact import to send repeat, mass invitations.
- **Selling usernames:** You may not buy or sell Twitter usernames.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - if you have followed and/or unfollowed large amounts of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - if your updates consist mainly of links, and not personal updates;
 - if a large number of people are blocking you;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicate content over multiple accounts or multiple duplicate updates on one account;
 - if you post multiple unrelated updates to a topic using #, trending or popular topic, or promoted trend;
 - if you send large numbers of duplicate replies or mentions;
 - if you send large numbers of unsolicited replies or mentions;
 - if you add a large number of unrelated users to lists;
 - if you repeatedly create false or misleading content;
 - if you are randomly or aggressively following, liking, or Retweeting Tweets;
 - if you repeatedly post other people's account information as your own (bio, Tweets, URL, etc.);
 - if you post misleading links (e.g. affiliate links, links to malware/clickjacking pages, etc.);
 - if you are creating misleading accounts or account interactions;
 - if you are selling or purchasing account interactions (such as selling or purchasing followers, Retweets, likes, etc.); and

- o if you are using or promoting third-party services or apps that claim to get you more followers (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers to your account).

See our support articles on [Following rules and best practices](#) and [Automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts under investigation may be removed from search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the [Terms of Service](#).

Trademark policy

Updated October 3, 2012

What is a Trademark Policy Violation on Twitter?

Using a company or business name, logo, or other trademark-protected materials in a manner that may mislead or confuse others with regard to its brand or business affiliation may be considered a trademark policy violation.

How Does Twitter Respond To Reported Trademark Policy Violations?

When we receive reports of trademark policy violations from holders of federal or international trademark registrations, we review the account and may take the following actions:

- When there is a clear intent to mislead others through the unauthorized use of a trademark, Twitter will suspend the account and notify the account holder.
- When we determine that an account appears to be confusing users, but is not purposefully passing itself off as the trademarked good or service, we give the account holder an opportunity to clear up any potential confusion. We may also release a username for the trademark holder's active use.
- We are responsive to reports about confusing or misleading [Promoted Tweet](#) and [Promoted Trend](#) copy, as well as [Promoted Account](#) profile information. For more information, please visit our [Trademark Policy for Promoted Products](#).

What Is not a Trademark Policy Violation?

Using another's trademark in a way that has nothing to do with the product or service for which the trademark was granted is not a violation of Twitter's trademark policy.

- Twitter usernames are provided on a first-come, first-served basis and may not be reserved. For information on why you may not be able to select a certain username, please see our [Why Can't I Register Certain Usernames](#) help page.
- Search and interest keywords advertisers choose are subject to applicable trademark law. Although Twitter may review reports of confusing keywords, we may not disable keywords in every case. For more information about our policies specific to Promoted

Products, please see our [Promoted Products Policy](#) and our [Trademark Policy for Promoted Products](#).

Guidelines for News feed, Commentary, and Fan Accounts

Twitter users are allowed to create news feed, commentary, and fan accounts. Twitter provides a platform for its users to share and receive a wide range of ideas and content, and we greatly value and respect our users' expression. Because of these principles, we do not actively monitor users' content and will not edit or remove user content, except in cases of violations of our [Terms of Service](#).

An account's profile information should make it clear that the account is not actually the company or business entity that is the subject of the news feed/commentary/fan account. Here are some suggestions for distinguishing your account:

- **Username:** The username should not be the trademarked name of the subject of the news feed, commentary, or fan account.
- **Name:** The profile name should not be the trademarked name of the company or include the trademarked name in a misleading manner.
- **Bio:** The bio should include a statement to distinguish it from the real company, such as "Unofficial Account," "Fan Account," or "Not affiliated with..."
- **Profile photo, header photo, or background image:** The account should not use another's trademark, logo or other copyright-protected image without express permission.
- **Communication with other users:** The account should not, through private or public communication with other users, try to deceive or mislead others about your identity.

Users may also choose to use different language to indicate that an account is not associated with the actual brand/company/product so long as it is clear and not confusing to others, and does not mislead or deceive.

If an account is reported to be confusing, we may request that the account holder make further changes to bring the account in compliance with these best practices.

How Can I Make My Own Account's Brand or Business Affiliation Clear?

We strongly recommend that you use all of Twitter's account settings (account name, location, web, and bio) to make your account's affiliation clear.

- Please see our [How to Change Your Profile Picture or Information](#) page for instructions on customizing your account. In particular, we recommend clearly stating your location, including your website if you have one, and clearly describing your brand or business in the bio, if applicable.
- You can also consider linking your official website to your Twitter account. [Twitter Buttons](#) are an easy way to add Twitter actions like "Follow" and "Mention" directly to your website.

What Information is Required When Reporting Trademark Policy Violations?

In order to investigate trademark policy violations, please provide all of the following information:

Username of the reported account (e.g., @safety or <http://www.twitter.com/safety>):

Your company name:

Your company Twitter account (if there is one):

Company website:

Your trademarked word, symbol, etc. (e.g. Twitter):

Trademark registration number:

Trademark registration office (e.g., USPTO):

Note: A federal or international trademark registration number is required. If the name you are reporting is not a registered mark (e.g., a government agency or non-profit organization), please let us know:

Your first and last name:

Title:

Address:

Phone:

Fax:

Email (must be from company domain):

Description of confusion (e.g., passing off as your company, including specific descriptions of content or behavior):

Requested Action (e.g., removal of violating account or transfer of trademarked username to an existing company account):

How Do I Report a Trademark Policy Violation?

You do not need a Twitter account to submit a trademark report. Holders of registered trademarks can report possible violations to Twitter's Policy Team through our [support forms](#).

Please submit trademark-related requests from your company email address and follow the format above to help expedite our response. Also, be sure to clearly describe to us why the account or Tweets may cause confusion with your mark.

Copyright and DMCA policy

Topics Covered:

- [What is the DMCA ?](#)
- [What Types of Copyright Complaints Does Twitter Respond To?](#)
- [Am I a Copyright Holder? How Do I Know?](#)
- [What Information Do You Need To Process a Copyright Complaint?](#)
- [How Do I File a Copyright Complaint?](#)
- [How Are Claims Processed?](#)
- [What Happens Next?](#)
- [What Happens If I Receive a DMCA Notification?](#)
- [When Should I File a Counter-Notice?](#)

- [What Information Do You Need to Process a Counter-notice?](#)
- [What Happens After I Submit a Counter-notice?](#)
- [Filing a DMCA Notification or Counter-notice Is Serious Business!](#)

What is the DMCA?

[Section 512](#) of the Digital Millennium Copyright Act (“DMCA”) outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

What Types of Copyright Complaints Does Twitter Respond To?

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile photo, header photo, or background, allegations concerning the unauthorized use of a copyrighted image uploaded through our photo hosting service, or Tweets containing links to allegedly infringing materials.

Tip: If you are concerned about the use of your brand or entity’s name, please review Twitter’s [Trademark Policy](#). If you are concerned about the use of a fictional character, please see our [Parody, Commentary, and Fan Accounts Policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <http://chillingeffects.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

Tip: If you are reporting the content of a Tweet, please give us a direct link to that Tweet following these instructions: <https://support.twitter.com/articles/80586>. Or please specify if the alleged infringement is in the background, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to Twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate DMCA notifications may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with access to the reporter's complaint along with instructions on how to file a counter-notice, and forward a copy of the complaint to [Chilling Effects](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a copy of the takedown notice, along with instructions for filing a counter-notification.

Tip: If you've not yet received a copy of the DMCA notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a copy of each DMCA notification and counter-notice that we process to [Chilling Effects](#), where they are posted to a public-facing website (with your personal information removed).

Withheld Tweet:

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)

Withheld Media:



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a DMCA Notification?

If you receive a DMCA notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

If you believe that the material reported in the DMCA notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Tip: Re-posting material removed in response to a DMCA notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the DMCA notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

Tip: We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a DMCA Notification or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Reporting sensitive media

If you encounter media in Tweets that you believe should be treated as sensitive under [Twitter's media policy](#), please report it using the process described below.

How do I report media in Tweets?

Follow these instructions to report media:

1. From the Tweet you would like to report, click or tap the **More icon** (••• on web or  on iOS;  on Android).
2. Select **Report**.
3. Select **It displays a sensitive [image/video/media]**.
4. Next we'll provide recommendations for additional actions you can take to improve your Twitter experience.

Please remember that if you are seeing something you don't like, and Twitter hasn't placed a warning label before it, it's possible that it doesn't meet our threshold for a warning on the media.

What happens to sensitive media I report?

The Twitter team reviews reports of media flagged by users to determine if that media requires a warning message in order to comply with Twitter's media policies. Please remember that by flagging media you are bringing it to the attention of the Twitter team. The flagged content will not automatically receive a warning message or be removed from the site.

For information about what you can do if your content is reported as sensitive, visit this [article](#).

How to control whether you see sensitive media in Tweets:

You can choose to view media in Tweets that may contain sensitive content without a warning. If you are concerned about viewing sensitive content, you should keep your default settings. Your default settings will help by providing warnings before seeing content that has been marked as possibly sensitive.

1. Log in to Twitter and go to your [Account settings](#).
2. Under the **Content** section, look for **Tweet media** and check the box next to **Do not inform me before showing media that may be sensitive**.
3. Save your settings at the bottom of the page.

Note: Remember you can also [unfollow](#) or [block](#) users that are disrupting your Twitter experience.

How to report content that is illegal or otherwise violates the Twitter Rules

For information outlining other types of violations and how you can report them to us, please see this [article](#).

Private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Some examples of private and confidential information include:

credit card information
social security or other national identity numbers
addresses or locations that are considered and treated as private
non-public, personal phone numbers
non-public, personal email addresses
images or videos that are considered and treated as private under applicable laws intimate photos or videos that were taken or distributed without the subject's consent

Keep in mind that although you may consider certain information to be private, not all postings of such information may be a violation of this policy. We may consider the context and nature of the information posted, local privacy laws, and other case-specific facts when determining if this policy has been violated. For example, if information was previously posted or displayed elsewhere on the Internet prior to being put on Twitter, it may not be a violation of this policy.

For frequently asked questions about reporting private information posted on Twitter, [click here](#). To learn more about protecting your private information on Twitter and other websites, [click here](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter Impersonation Policy.

An account will not be removed if

- the user shares your name but has no other commonalities, or
- the profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same background or avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to [Twitter's Parody Policy](#) for more information about these types of accounts.

For Frequently Asked Questions about reporting impersonation on Twitter, [click here](#).

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and un-following large numbers of other users). You can read more about these below, but if you don't follow or un-follow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or un-following behavior. Please review our [Automation Rules and Best Practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and un-follow them. Aggressive follow churn is when an account *repeatedly* follows and then un-follows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 2000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 2000 people total**. Once you've followed 2000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

Tip: If you need to communicate with someone but don't need to see their updates every day, don't follow them. Visit their profile or send them an @reply when you need to; sending @replies doesn't require following,

and your reply will appear in the person's @mentions tab so they can reply back.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

The [Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an

application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and un-following accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Having trouble?

Check out our [Following Troubleshooting](#) section for solutions to common problems.

Automation Rules & Best Practices

Submitted Nov 08, 2009 by [ginger](#)

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating

the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate

offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the time-stamp, like this:

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. All users must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in the temporary locking and/or permanent suspension of account(s).

Please note that we may need to change these rules from time to time and reserve the right to do so. The most current version will always be available at twitter.com/rules.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter.

- **Trademark:** We reserve the right to reclaim usernames on behalf of businesses or individuals that hold legal claim or trademark on those usernames. Accounts using business names and/or logos to mislead others may be permanently suspended.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Graphic content:** You may not use pornographic or excessively violent media in your profile image, header image, or background image. Twitter may allow some forms of graphic content in Tweets marked as [sensitive media](#). When content crosses the line into gratuitous images of death, Twitter may ask that you remove the content out of respect for the deceased.
- **Unlawful use:** You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter badges:** You may not use badges, such as but not limited to the “promoted” or “verified” Twitter badge, unless provided by Twitter. Accounts using these badges as part of profile photos, header photos, background images, or in a way that falsely implies affiliation with Twitter may be suspended.

Abusive Behavior

We believe in freedom of expression and in speaking truth to power, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we do not tolerate behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Any accounts and related accounts engaging in the activities specified below may be temporarily locked and/or subject to permanent suspension.

- **Violent threats (direct or indirect):** You may not make threats of violence or promote violence, including threatening or promoting terrorism.

- **Harassment:** You may not incite or engage in the targeted abuse or harassment of others. Some of the factors that we may consider when evaluating abusive behavior include:
 - if a primary purpose of the reported account is to harass or send abusive messages to others;
 - if the reported behavior is one-sided or includes threats;
 - if the reported account is inciting others to harass another account; and
 - if the reported account is sending harassing messages to an account from multiple accounts.
- **Hateful conduct:** You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.
- **Multiple account abuse:** Creating multiple accounts with overlapping uses or in order to evade the temporary or permanent suspension of a separate account is not allowed.
- **Private information:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address, or Social Security/National Identity numbers, without their express authorization and permission. In addition, you may not post intimate photos or videos that were taken or distributed without the subject's consent. Read more about our private information policy [here](#).
- **Impersonation:** You may not impersonate others through the Twitter service in a manner that is intended to or does mislead, confuse, or deceive others. Read more about our impersonation policy [here](#).
- **Self-harm:** You may encounter someone considering suicide or self harm on Twitter. When we receive reports that a person is threatening suicide or self harm, we may take a number of steps to assist them, such as reaching out to that person expressing our concern and the concern of other users on Twitter or providing resources such as contact information for our mental health partners.

Spam

We strive to protect people on Twitter from technical abuse and spam. Any accounts engaging in the activities specified below may be temporarily locked or subject to permanent suspension.

- **Username squatting:** You may not engage in username squatting. Accounts that are inactive for more than six months may also be removed without further notice. Some of the factors we take into consideration when determining what conduct is considered to be username squatting are:
 - the number of accounts created;
 - creating accounts for the purpose of preventing others from using those account names;
 - creating accounts for the purpose of selling those accounts; and

- using feeds of third-party content to update and maintain accounts under the names of those third parties.
- **Invitation spam:** You may not use twitter.com's address book contact import to send repeat, mass invitations.
- **Selling usernames:** You may not buy or sell Twitter usernames.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - if you have followed and/or unfollowed large amounts of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - if your updates consist mainly of links, and not personal updates;
 - if a large number of people are blocking you;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicate content over multiple accounts or multiple duplicate updates on one account;
 - if you post multiple unrelated updates to a topic using #, trending or popular topic, or promoted trend;
 - if you send large numbers of duplicate replies or mentions;
 - if you send large numbers of unsolicited replies or mentions;
 - if you add a large number of unrelated users to lists;
 - if you repeatedly create false or misleading content;
 - if you are randomly or aggressively following, liking, or Retweeting Tweets;
 - if you repeatedly post other people's account information as your own (bio, Tweets, URL, etc.);
 - if you post misleading links (e.g. affiliate links, links to malware/clickjacking pages, etc.);
 - if you are creating misleading accounts or account interactions;
 - if you are selling or purchasing account interactions (such as selling or purchasing followers, Retweets, likes, etc.); and

- o if you are using or promoting third-party services or apps that claim to get you more followers (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers to your account).

See our support articles on [Following rules and best practices](#) and [Automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts under investigation may be removed from search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the [Terms of Service](#).

Trademark policy

Updated October 3, 2012

What is a Trademark Policy Violation on Twitter?

Using a company or business name, logo, or other trademark-protected materials in a manner that may mislead or confuse others with regard to its brand or business affiliation may be considered a trademark policy violation.

How Does Twitter Respond To Reported Trademark Policy Violations?

When we receive reports of trademark policy violations from holders of federal or international trademark registrations, we review the account and may take the following actions:

- When there is a clear intent to mislead others through the unauthorized use of a trademark, Twitter will suspend the account and notify the account holder.
- When we determine that an account appears to be confusing users, but is not purposefully passing itself off as the trademarked good or service, we give the account holder an opportunity to clear up any potential confusion. We may also release a username for the trademark holder's active use.
- We are responsive to reports about confusing or misleading [Promoted Tweet](#) and [Promoted Trend](#) copy, as well as [Promoted Account](#) profile information. For more information, please visit our [Trademark Policy for Promoted Products](#).

What Is not a Trademark Policy Violation?

Using another's trademark in a way that has nothing to do with the product or service for which the trademark was granted is not a violation of Twitter's trademark policy.

- Twitter usernames are provided on a first-come, first-served basis and may not be reserved. For information on why you may not be able to select a certain username, please see our [Why Can't I Register Certain Usernames](#) help page.
- Search and interest keywords advertisers choose are subject to applicable trademark law. Although Twitter may review reports of confusing keywords, we may not disable keywords in every case. For more information about our policies specific to Promoted

Products, please see our [Promoted Products Policy](#) and our [Trademark Policy for Promoted Products](#).

Guidelines for News feed, Commentary, and Fan Accounts

Twitter users are allowed to create news feed, commentary, and fan accounts. Twitter provides a platform for its users to share and receive a wide range of ideas and content, and we greatly value and respect our users' expression. Because of these principles, we do not actively monitor users' content and will not edit or remove user content, except in cases of violations of our [Terms of Service](#).

An account's profile information should make it clear that the account is not actually the company or business entity that is the subject of the news feed/commentary/fan account. Here are some suggestions for distinguishing your account:

- **Username:** The username should not be the trademarked name of the subject of the news feed, commentary, or fan account.
- **Name:** The profile name should not be the trademarked name of the company or include the trademarked name in a misleading manner.
- **Bio:** The bio should include a statement to distinguish it from the real company, such as "Unofficial Account," "Fan Account," or "Not affiliated with..."
- **Profile photo, header photo, or background image:** The account should not use another's trademark, logo or other copyright-protected image without express permission.
- **Communication with other users:** The account should not, through private or public communication with other users, try to deceive or mislead others about your identity.

Users may also choose to use different language to indicate that an account is not associated with the actual brand/company/product so long as it is clear and not confusing to others, and does not mislead or deceive.

If an account is reported to be confusing, we may request that the account holder make further changes to bring the account in compliance with these best practices.

How Can I Make My Own Account's Brand or Business Affiliation Clear?

We strongly recommend that you use all of Twitter's account settings (account name, location, web, and bio) to make your account's affiliation clear.

- Please see our [How to Change Your Profile Picture or Information](#) page for instructions on customizing your account. In particular, we recommend clearly stating your location, including your website if you have one, and clearly describing your brand or business in the bio, if applicable.
- You can also consider linking your official website to your Twitter account. [Twitter Buttons](#) are an easy way to add Twitter actions like "Follow" and "Mention" directly to your website.

What Information is Required When Reporting Trademark Policy Violations?

In order to investigate trademark policy violations, please provide all of the following information:

Username of the reported account (e.g., @safety or <http://www.twitter.com/safety>):

Your company name:

Your company Twitter account (if there is one):

Company website:

Your trademarked word, symbol, etc. (e.g. Twitter):

Trademark registration number:

Trademark registration office (e.g., USPTO):

Note: A federal or international trademark registration number is required. If the name you are reporting is not a registered mark (e.g., a government agency or non-profit organization), please let us know:

Your first and last name:

Title:

Address:

Phone:

Fax:

Email (must be from company domain):

Description of confusion (e.g., passing off as your company, including specific descriptions of content or behavior):

Requested Action (e.g., removal of violating account or transfer of trademarked username to an existing company account):

How Do I Report a Trademark Policy Violation?

You do not need a Twitter account to submit a trademark report. Holders of registered trademarks can report possible violations to Twitter's Policy Team through our [support forms](#).

Please submit trademark-related requests from your company email address and follow the format above to help expedite our response. Also, be sure to clearly describe to us why the account or Tweets may cause confusion with your mark.

[Copyright and DMCA policy](#)

Topics Covered:

- [What is the DMCA ?](#)
- [What Types of Copyright Complaints Does Twitter Respond To?](#)
- [Am I a Copyright Holder? How Do I Know?](#)
- [What Information Do You Need To Process a Copyright Complaint?](#)
- [How Do I File a Copyright Complaint?](#)
- [How Are Claims Processed?](#)
- [What Happens Next?](#)
- [What Happens If I Receive a DMCA Notification?](#)
- [When Should I File a Counter-Notice?](#)

- [What Information Do You Need to Process a Counter-notice?](#)
- [What Happens After I Submit a Counter-notice?](#)
- [Filing a DMCA Notification or Counter-notice Is Serious Business!](#)

What is the DMCA?

[Section 512](#) of the Digital Millennium Copyright Act (“DMCA”) outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

What Types of Copyright Complaints Does Twitter Respond To?

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile photo, header photo, or background, allegations concerning the unauthorized use of a copyrighted image uploaded through our photo hosting service, or Tweets containing links to allegedly infringing materials.

Tip: If you are concerned about the use of your brand or entity’s name, please review Twitter’s [Trademark Policy](#). If you are concerned about the use of a fictional character, please see our [Parody, Commentary, and Fan Accounts Policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <http://chillingeffects.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

Tip: If you are reporting the content of a Tweet, please give us a direct link to that Tweet following these instructions: <https://support.twitter.com/articles/80586>. Or please specify if the alleged infringement is in the background, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to Twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate DMCA notifications may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with access to the reporter's complaint along with instructions on how to file a counter-notice, and forward a copy of the complaint to [Chilling Effects](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a copy of the takedown notice, along with instructions for filing a counter-notification.

Tip: If you've not yet received a copy of the DMCA notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a copy of each DMCA notification and counter-notice that we process to [Chilling Effects](#), where they are posted to a public-facing website (with your personal information removed).

Withheld Tweet:

Tweet withheld 18h
This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)

Withheld Media:



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a DMCA Notification?

If you receive a DMCA notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

If you believe that the material reported in the DMCA notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Tip: Re-posting material removed in response to a DMCA notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the DMCA notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

Tip: We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a DMCA Notification or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Reporting sensitive media

If you encounter media in Tweets that you believe should be treated as sensitive under [Twitter's media policy](#), please report it using the process described below.

How do I report media in Tweets?

Follow these instructions to report media:

1. From the Tweet you would like to report, click or tap the **More icon** (••• on web or  on iOS;  on Android).
2. Select **Report**.
3. Select **It displays a sensitive [image/video/media]**.
4. Next we'll provide recommendations for additional actions you can take to improve your Twitter experience.

Please remember that if you are seeing something you don't like, and Twitter hasn't placed a warning label before it, it's possible that it doesn't meet our threshold for a warning on the media.

What happens to sensitive media I report?

The Twitter team reviews reports of media flagged by users to determine if that media requires a warning message in order to comply with Twitter's media policies. Please remember that by flagging media you are bringing it to the attention of the Twitter team. The flagged content will not automatically receive a warning message or be removed from the site.

For information about what you can do if your content is reported as sensitive, visit this [article](#).

How to control whether you see sensitive media in Tweets:

You can choose to view media in Tweets that may contain sensitive content without a warning. If you are concerned about viewing sensitive content, you should keep your default settings. Your default settings will help by providing warnings before seeing content that has been marked as possibly sensitive.

1. Log in to Twitter and go to your [Account settings](#).
2. Under the **Content** section, look for **Tweet media** and check the box next to **Do not inform me before showing media that may be sensitive**.
3. Save your settings at the bottom of the page.

Note: Remember you can also [unfollow](#) or [block](#) users that are disrupting your Twitter experience.

How to report content that is illegal or otherwise violates the Twitter Rules

For information outlining other types of violations and how you can report them to us, please see this [article](#).

Private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Some examples of private and confidential information include:

credit card information
social security or other national identity numbers
addresses or locations that are considered and treated as private
non-public, personal phone numbers
non-public, personal email addresses
images or videos that are considered and treated as private under applicable laws intimate photos or videos that were taken or distributed without the subject's consent

Keep in mind that although you may consider certain information to be private, not all postings of such information may be a violation of this policy. We may consider the context and nature of the information posted, local privacy laws, and other case-specific facts when determining if this policy has been violated. For example, if information was previously posted or displayed elsewhere on the Internet prior to being put on Twitter, it may not be a violation of this policy.

For frequently asked questions about reporting private information posted on Twitter, [click here](#). To learn more about protecting your private information on Twitter and other websites, [click here](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter Impersonation Policy.

An account will not be removed if

- the user shares your name but has no other commonalities, or
- the profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same background or avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to [Twitter's Parody Policy](#) for more information about these types of accounts.

For Frequently Asked Questions about reporting impersonation on Twitter, [click here](#).

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow

or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation Rules & Best Practices

Submitted Nov 08, 2009 by [ginger](#)

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.

- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. All users must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in the temporary locking and/or permanent suspension of account(s).

Please note that we may need to change these rules from time to time and reserve the right to do so. The most current version will always be available at twitter.com/rules.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter.

- **Trademark:** We reserve the right to reclaim usernames on behalf of businesses or individuals that hold legal claim or trademark on those usernames. Accounts using business names and/or logos to mislead others may be permanently suspended.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Graphic content:** You may not use pornographic or excessively violent media in your profile image, header image, or background image. Twitter may allow some forms of graphic content in Tweets marked as [sensitive media](#). When content crosses the line into gratuitous images of death, Twitter may ask that you remove the content out of respect for the deceased.
- **Unlawful use:** You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter badges:** You may not use badges, such as but not limited to the “promoted” or “verified” Twitter badge, unless provided by Twitter. Accounts using these badges as part of profile photos, header photos, background images, or in a way that falsely implies affiliation with Twitter may be suspended.

Abusive Behavior

We believe in freedom of expression and in speaking truth to power, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we do not tolerate behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Any accounts and related accounts engaging in the activities specified below may be temporarily locked and/or subject to permanent suspension.

- **Violent threats (direct or indirect):** You may not make threats of violence or promote violence, including threatening or promoting terrorism.

- **Harassment:** You may not incite or engage in the targeted abuse or harassment of others. Some of the factors that we may consider when evaluating abusive behavior include:
 - if a primary purpose of the reported account is to harass or send abusive messages to others;
 - if the reported behavior is one-sided or includes threats;
 - if the reported account is inciting others to harass another account; and
 - if the reported account is sending harassing messages to an account from multiple accounts.
- **Hateful conduct:** You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.
- **Multiple account abuse:** Creating multiple accounts with overlapping uses or in order to evade the temporary or permanent suspension of a separate account is not allowed.
- **Private information:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address, or Social Security/National Identity numbers, without their express authorization and permission. In addition, you may not post intimate photos or videos that were taken or distributed without the subject's consent. Read more about our private information policy [here](#).
- **Impersonation:** You may not impersonate others through the Twitter service in a manner that is intended to or does mislead, confuse, or deceive others. Read more about our impersonation policy [here](#).
- **Self-harm:** You may encounter someone considering suicide or self harm on Twitter. When we receive reports that a person is threatening suicide or self harm, we may take a number of steps to assist them, such as reaching out to that person expressing our concern and the concern of other users on Twitter or providing resources such as contact information for our mental health partners.

Spam

We strive to protect people on Twitter from technical abuse and spam. Any accounts engaging in the activities specified below may be temporarily locked or subject to permanent suspension.

- **Username squatting:** You may not engage in username squatting. Accounts that are inactive for more than six months may also be removed without further notice. Some of the factors we take into consideration when determining what conduct is considered to be username squatting are:
 - the number of accounts created;
 - creating accounts for the purpose of preventing others from using those account names;
 - creating accounts for the purpose of selling those accounts; and

- using feeds of third-party content to update and maintain accounts under the names of those third parties.
- **Invitation spam:** You may not use twitter.com's address book contact import to send repeat, mass invitations.
- **Selling usernames:** You may not buy or sell Twitter usernames.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - if you have followed and/or unfollowed large amounts of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - if your updates consist mainly of links, and not personal updates;
 - if a large number of people are blocking you;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicate content over multiple accounts or multiple duplicate updates on one account;
 - if you post multiple unrelated updates to a topic using #, trending or popular topic, or promoted trend;
 - if you send large numbers of duplicate replies or mentions;
 - if you send large numbers of unsolicited replies or mentions;
 - if you add a large number of unrelated users to lists;
 - if you repeatedly create false or misleading content;
 - if you are randomly or aggressively following, liking, or Retweeting Tweets;
 - if you repeatedly post other people's account information as your own (bio, Tweets, URL, etc.);
 - if you post misleading links (e.g. affiliate links, links to malware/clickjacking pages, etc.);
 - if you are creating misleading accounts or account interactions;
 - if you are selling or purchasing account interactions (such as selling or purchasing followers, Retweets, likes, etc.); and

- o if you are using or promoting third-party services or apps that claim to get you more followers (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers to your account).

See our support articles on [Following rules and best practices](#) and [Automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts under investigation may be removed from search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the [Terms of Service](#).

Trademark policy

Updated October 3, 2012

What is a Trademark Policy Violation on Twitter?

Using a company or business name, logo, or other trademark-protected materials in a manner that may mislead or confuse others with regard to its brand or business affiliation may be considered a trademark policy violation.

How Does Twitter Respond To Reported Trademark Policy Violations?

When we receive reports of trademark policy violations from holders of federal or international trademark registrations, we review the account and may take the following actions:

- When there is a clear intent to mislead others through the unauthorized use of a trademark, Twitter will suspend the account and notify the account holder.
- When we determine that an account appears to be confusing users, but is not purposefully passing itself off as the trademarked good or service, we give the account holder an opportunity to clear up any potential confusion. We may also release a username for the trademark holder's active use.
- We are responsive to reports about confusing or misleading [Promoted Tweet](#) and [Promoted Trend](#) copy, as well as [Promoted Account](#) profile information. For more information, please visit our [Trademark Policy for Promoted Products](#).

What Is not a Trademark Policy Violation?

Using another's trademark in a way that has nothing to do with the product or service for which the trademark was granted is not a violation of Twitter's trademark policy.

- Twitter usernames are provided on a first-come, first-served basis and may not be reserved. For information on why you may not be able to select a certain username, please see our [Why Can't I Register Certain Usernames](#) help page.
- Search and interest keywords advertisers choose are subject to applicable trademark law. Although Twitter may review reports of confusing keywords, we may not disable keywords in every case. For more information about our policies specific to Promoted

Products, please see our [Promoted Products Policy](#) and our [Trademark Policy for Promoted Products](#).

Guidelines for News feed, Commentary, and Fan Accounts

Twitter users are allowed to create news feed, commentary, and fan accounts. Twitter provides a platform for its users to share and receive a wide range of ideas and content, and we greatly value and respect our users' expression. Because of these principles, we do not actively monitor users' content and will not edit or remove user content, except in cases of violations of our [Terms of Service](#).

An account's profile information should make it clear that the account is not actually the company or business entity that is the subject of the news feed/commentary/fan account. Here are some suggestions for distinguishing your account:

- **Username:** The username should not be the trademarked name of the subject of the news feed, commentary, or fan account.
- **Name:** The profile name should not be the trademarked name of the company or include the trademarked name in a misleading manner.
- **Bio:** The bio should include a statement to distinguish it from the real company, such as "Unofficial Account," "Fan Account," or "Not affiliated with..."
- **Profile photo, header photo, or background image:** The account should not use another's trademark, logo or other copyright-protected image without express permission.
- **Communication with other users:** The account should not, through private or public communication with other users, try to deceive or mislead others about your identity.

Users may also choose to use different language to indicate that an account is not associated with the actual brand/company/product so long as it is clear and not confusing to others, and does not mislead or deceive.

If an account is reported to be confusing, we may request that the account holder make further changes to bring the account in compliance with these best practices.

How Can I Make My Own Account's Brand or Business Affiliation Clear?

We strongly recommend that you use all of Twitter's account settings (account name, location, web, and bio) to make your account's affiliation clear.

- Please see our [How to Change Your Profile Picture or Information](#) page for instructions on customizing your account. In particular, we recommend clearly stating your location, including your website if you have one, and clearly describing your brand or business in the bio, if applicable.
- You can also consider linking your official website to your Twitter account. [Twitter Buttons](#) are an easy way to add Twitter actions like "Follow" and "Mention" directly to your website.

What Information is Required When Reporting Trademark Policy Violations?

In order to investigate trademark policy violations, please provide all of the following information:

Username of the reported account (e.g., @safety or <http://www.twitter.com/safety>):

Your company name:

Your company Twitter account (if there is one):

Company website:

Your trademarked word, symbol, etc. (e.g. Twitter):

Trademark registration number:

Trademark registration office (e.g., USPTO):

Note: A federal or international trademark registration number is required. If the name you are reporting is not a registered mark (e.g., a government agency or non-profit organization), please let us know:

Your first and last name:

Title:

Address:

Phone:

Fax:

Email (must be from company domain):

Description of confusion (e.g., passing off as your company, including specific descriptions of content or behavior):

Requested Action (e.g., removal of violating account or transfer of trademarked username to an existing company account):

How Do I Report a Trademark Policy Violation?

You do not need a Twitter account to submit a trademark report. Holders of registered trademarks can report possible violations to Twitter's Policy Team through our [support forms](#).

Please submit trademark-related requests from your company email address and follow the format above to help expedite our response. Also, be sure to clearly describe to us why the account or Tweets may cause confusion with your mark.

Copyright policy

Topics Covered:

- [What Types of Copyright Complaints Does Twitter Respond To?](#)
- [Am I a Copyright Holder? How Do I Know?](#)
- [What Information Do You Need To Process a Copyright Complaint?](#)
- [How Do I File a Copyright Complaint?](#)

- [How Are Claims Processed?](#)
- [What Happens Next?](#)
- [What Happens If I Receive a Copyright Notification?](#)
- [When Should I File a Counter-notice?](#)
- [What Information Do You Need to Process a Counter-notice?](#)
- [What Happens After I Submit a Counter-notice?](#)
- [Filing a Copyright Notice or Counter-notice Is Serious Business!](#)

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile photo, header photo, or background, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [Fair Use](#) page for more information).

Tip: If you are concerned about the use of your brand or entity’s name, please review Twitter’s [Trademark Policy](#). If you are concerned about the use of a fictional character, please see our [Parody, Commentary, and Fan Accounts Policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

Tip: If you are reporting the content of a Tweet, please give us a direct link to that Tweet following these instructions: <https://support.twitter.com/articles/80586>. Or please specify if the alleged infringement is in the background, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to Twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with access to the reporter's complaint along with instructions on how to file a counter-notice, and forward a copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a copy of the takedown notice, along with instructions for filing a counter-notification.

Tip: If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a copy of each copyright complaint and counter-notice that we process to [Lumen](#), where they are posted to a public-facing website (with your personal information removed).

Withheld Tweet:

Tweet withheld 18h
This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)

Withheld Media:



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Tip: Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

Tip: We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Reporting sensitive media

If you encounter media in Tweets that you believe should be treated as sensitive under [Twitter's media policy](#), please report it using the process described below.

How do I report media in Tweets?

Follow these instructions to report media:

1. From the Tweet you would like to report, click or tap the **More icon** (••• on web or  on iOS;  on Android).
2. Select **Report**.
3. Select **It displays a sensitive [image/video/media]**.
4. Next we'll provide recommendations for additional actions you can take to improve your Twitter experience.

Please remember that if you are seeing something you don't like, and Twitter hasn't placed a warning label before it, it's possible that it doesn't meet our threshold for a warning on the media.

What happens to sensitive media I report?

The Twitter team reviews reports of media flagged by users to determine if that media requires a warning message in order to comply with Twitter's media policies. Please remember that by flagging media you are bringing it to the attention of the Twitter team. The flagged content will not automatically receive a warning message or be removed from the site.

For information about what you can do if your content is reported as sensitive, visit this [article](#).

How to control whether you see sensitive media in Tweets:

You can choose to view media in Tweets that may contain sensitive content without a warning. If you are concerned about viewing sensitive content, you should keep your default settings. Your default settings will help by providing warnings before seeing content that has been marked as possibly sensitive.

1. Log in to Twitter and go to your [Account settings](#).
2. Under the **Content** section, look for **Tweet media** and check the box next to **Do not inform me before showing media that may be sensitive**.
3. Save your settings at the bottom of the page.

Note: Remember you can also [unfollow](#) or [block](#) users that are disrupting your Twitter experience.

How to report content that is illegal or otherwise violates the Twitter Rules

For information outlining other types of violations and how you can report them to us, please see this [article](#).

Private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Some examples of private and confidential information include:

credit card information
social security or other national identity numbers
addresses or locations that are considered and treated as private
non-public, personal phone numbers
non-public, personal email addresses
images or videos that are considered and treated as private under applicable laws intimate
photos or videos that were taken or distributed without the subject's consent

Keep in mind that although you may consider certain information to be private, not all postings of such information may be a violation of this policy. We may consider the context and nature of the information posted, local privacy laws, and other case-specific facts when determining if this policy has been violated. For example, if information was previously posted or displayed elsewhere on the Internet prior to being put on Twitter, it may not be a violation of this policy.

For frequently asked questions about reporting private information posted on Twitter, [click here](#). To learn more about protecting your private information on Twitter and other websites, [click here](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter Impersonation Policy.

An account will not be removed if

- the user shares your name but has no other commonalities, or
- the profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same background or avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to [Twitter's Parody Policy](#) for more information about these types of accounts.

For Frequently Asked Questions about reporting impersonation on Twitter, [click here](#).

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or

squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

[Automation Rules & Best Practices](#)

Submitted Nov 08, 2009 by [ginger](#)

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the time-stamp, like this:

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. All users must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in the temporary locking and/or permanent suspension of account(s).

Please note that we may need to change these rules from time to time and reserve the right to do so. The most current version will always be available at twitter.com/rules.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter.

- **Trademark:** We reserve the right to reclaim usernames on behalf of businesses or individuals that hold legal claim or trademark on those usernames. Accounts using business names and/or logos to mislead others may be permanently suspended.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Graphic content:** You may not use pornographic or excessively violent media in your profile image or header image. Twitter may allow some forms of graphic content in Tweets marked as [sensitive media](#). When content crosses the line into gratuitous images of death, Twitter may ask that you remove the content out of respect for the deceased.
- **Unlawful use:** You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter badges:** You may not use badges, such as but not limited to the “promoted” or “verified” Twitter badge, unless provided by Twitter. Accounts using these badges as part of profile photos, header photos, or in a way that falsely implies affiliation with Twitter, may be suspended.

Abusive Behavior

We believe in freedom of expression and in speaking truth to power, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we do not tolerate behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Any accounts and related accounts engaging in the activities specified below may be temporarily locked and/or subject to permanent suspension.

- **Violent threats (direct or indirect):** You may not make threats of violence or promote violence, including threatening or promoting terrorism.

- **Harassment:** You may not incite or engage in the targeted abuse or harassment of others. Some of the factors that we may consider when evaluating abusive behavior include:
 - if a primary purpose of the reported account is to harass or send abusive messages to others;
 - if the reported behavior is one-sided or includes threats;
 - if the reported account is inciting others to harass another account; and
 - if the reported account is sending harassing messages to an account from multiple accounts.
- **Hateful conduct:** You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.
- **Multiple account abuse:** Creating multiple accounts with overlapping uses or in order to evade the temporary or permanent suspension of a separate account is not allowed.
- **Private information:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address, or Social Security/National Identity numbers, without their express authorization and permission. In addition, you may not post intimate photos or videos that were taken or distributed without the subject's consent. Read more about our private information policy [here](#).
- **Impersonation:** You may not impersonate others through the Twitter service in a manner that is intended to or does mislead, confuse, or deceive others. Read more about our impersonation policy [here](#).
- **Self-harm:** You may encounter someone considering suicide or self harm on Twitter. When we receive reports that a person is threatening suicide or self harm, we may take a number of steps to assist them, such as reaching out to that person expressing our concern and the concern of other users on Twitter or providing resources such as contact information for our mental health partners.

Spam

We strive to protect people on Twitter from technical abuse and spam. Any accounts engaging in the activities specified below may be temporarily locked or subject to permanent suspension.

- **Username squatting:** You may not engage in username squatting. Accounts that are inactive for more than six months may also be removed without further notice. Some of the factors we take into consideration when determining what conduct is considered to be username squatting are:
 - the number of accounts created;
 - creating accounts for the purpose of preventing others from using those account names;
 - creating accounts for the purpose of selling those accounts; and

- using feeds of third-party content to update and maintain accounts under the names of those third parties.
- **Invitation spam:** You may not use twitter.com's address book contact import to send repeat, mass invitations.
- **Selling usernames:** You may not buy or sell Twitter usernames.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - if you have followed and/or unfollowed large amounts of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - if your updates consist mainly of links, and not personal updates;
 - if a large number of people are blocking you;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicate content over multiple accounts or multiple duplicate updates on one account;
 - if you post multiple unrelated updates to a topic using #, trending or popular topic, or promoted trend;
 - if you send large numbers of duplicate replies or mentions;
 - if you send large numbers of unsolicited replies or mentions;
 - if you add a large number of unrelated users to lists;
 - if you repeatedly create false or misleading content;
 - if you are randomly or aggressively following, liking, or Retweeting Tweets;
 - if you repeatedly post other people's account information as your own (bio, Tweets, URL, etc.);
 - if you post misleading links (e.g. affiliate links, links to malware/clickjacking pages, etc.);
 - if you are creating misleading accounts or account interactions;
 - if you are selling or purchasing account interactions (such as selling or purchasing followers, Retweets, likes, etc.); and

- o if you are using or promoting third-party services or apps that claim to get you more followers (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers to your account).

See our support articles on [Following rules and best practices](#) and [Automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts under investigation may be removed from search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the [Terms of Service](#).

Trademark policy

Updated October 3, 2012

What is a Trademark Policy Violation on Twitter?

Using a company or business name, logo, or other trademark-protected materials in a manner that may mislead or confuse others with regard to its brand or business affiliation may be considered a trademark policy violation.

How Does Twitter Respond To Reported Trademark Policy Violations?

When we receive reports of trademark policy violations from holders of federal or international trademark registrations, we review the account and may take the following actions:

- When there is a clear intent to mislead others through the unauthorized use of a trademark, Twitter will suspend the account and notify the account holder.
- When we determine that an account appears to be confusing users, but is not purposefully passing itself off as the trademarked good or service, we give the account holder an opportunity to clear up any potential confusion. We may also release a username for the trademark holder's active use.
- We are responsive to reports about confusing or misleading [Promoted Tweet](#) and [Promoted Trend](#) copy, as well as [Promoted Account](#) profile information. For more information, please visit our [Trademark Policy for Promoted Products](#).

What Is not a Trademark Policy Violation?

Using another's trademark in a way that has nothing to do with the product or service for which the trademark was granted is not a violation of Twitter's trademark policy.

- Twitter usernames are provided on a first-come, first-served basis and may not be reserved. For information on why you may not be able to select a certain username, please see our [Why Can't I Register Certain Usernames](#) help page.
- Search and interest keywords advertisers choose are subject to applicable trademark law. Although Twitter may review reports of confusing keywords, we may not disable keywords in every case. For more information about our policies specific to Promoted

Products, please see our [Promoted Products Policy](#) and our [Trademark Policy for Promoted Products](#).

Guidelines for News feed, Commentary, and Fan Accounts

Twitter users are allowed to create news feed, commentary, and fan accounts. Twitter provides a platform for its users to share and receive a wide range of ideas and content, and we greatly value and respect our users' expression. Because of these principles, we do not actively monitor users' content and will not edit or remove user content, except in cases of violations of our [Terms of Service](#).

An account's profile information should make it clear that the account is not actually the company or business entity that is the subject of the news feed/commentary/fan account. Here are some suggestions for distinguishing your account:

- **Username:** The username should not be the trademarked name of the subject of the news feed, commentary, or fan account.
- **Name:** The profile name should not be the trademarked name of the company or include the trademarked name in a misleading manner.
- **Bio:** The bio should include a statement to distinguish it from the real company, such as "Unofficial Account," "Fan Account," or "Not affiliated with..."
- **Profile photo, header photo, or background image:** The account should not use another's trademark, logo or other copyright-protected image without express permission.
- **Communication with other users:** The account should not, through private or public communication with other users, try to deceive or mislead others about your identity.

Users may also choose to use different language to indicate that an account is not associated with the actual brand/company/product so long as it is clear and not confusing to others, and does not mislead or deceive.

If an account is reported to be confusing, we may request that the account holder make further changes to bring the account in compliance with these best practices.

How Can I Make My Own Account's Brand or Business Affiliation Clear?

We strongly recommend that you use all of Twitter's account settings (account name, location, web, and bio) to make your account's affiliation clear.

- Please see our [How to Change Your Profile Picture or Information](#) page for instructions on customizing your account. In particular, we recommend clearly stating your location, including your website if you have one, and clearly describing your brand or business in the bio, if applicable.
- You can also consider linking your official website to your Twitter account. [Twitter Buttons](#) are an easy way to add Twitter actions like "Follow" and "Mention" directly to your website.

What Information is Required When Reporting Trademark Policy Violations?

In order to investigate trademark policy violations, please provide all of the following information:

Username of the reported account (e.g., @safety or <http://www.twitter.com/safety>):

Your company name:

Your company Twitter account (if there is one):

Company website:

Your trademarked word, symbol, etc. (e.g. Twitter):

Trademark registration number:

Trademark registration office (e.g., USPTO):

Note: A federal or international trademark registration number is required. If the name you are reporting is not a registered mark (e.g., a government agency or non-profit organization), please let us know:

Your first and last name:

Title:

Address:

Phone:

Fax:

Email (must be from company domain):

Description of confusion (e.g., passing off as your company, including specific descriptions of content or behavior):

Requested Action (e.g., removal of violating account or transfer of trademarked username to an existing company account):

How Do I Report a Trademark Policy Violation?

You do not need a Twitter account to submit a trademark report. Holders of registered trademarks can report possible violations to Twitter's Policy Team through our [support forms](#).

Please submit trademark-related requests from your company email address and follow the format above to help expedite our response. Also, be sure to clearly describe to us why the account or Tweets may cause confusion with your mark.

Copyright policy

Topics Covered:

- [What Types of Copyright Complaints Does Twitter Respond To?](#)
- [Am I a Copyright Holder? How Do I Know?](#)
- [What Information Do You Need To Process a Copyright Complaint?](#)
- [How Do I File a Copyright Complaint?](#)

- [How Are Claims Processed?](#)
- [What Happens Next?](#)
- [What Happens If I Receive a Copyright Notification?](#)
- [When Should I File a Counter-notice?](#)
- [What Information Do You Need to Process a Counter-notice?](#)
- [What Happens After I Submit a Counter-notice?](#)
- [Filing a Copyright Notice or Counter-notice Is Serious Business!](#)

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile photo, header photo, or background, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [Fair Use](#) page for more information).

Tip: If you are concerned about the use of your brand or entity’s name, please review Twitter’s [Trademark Policy](#). If you are concerned about the use of a fictional character, please see our [Parody, Commentary, and Fan Accounts Policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

Tip: If you are reporting the content of a Tweet, please give us a direct link to that Tweet following these instructions: <https://support.twitter.com/articles/80586>. Or please specify if the alleged infringement is in the background, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to Twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with access to the reporter's complaint along with instructions on how to file a counter-notice, and forward a copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a copy of the takedown notice, along with instructions for filing a counter-notification.

Tip: If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a copy of each copyright complaint and counter-notice that we process to [Lumen](#), where they are posted to a public-facing website (with your personal information removed).

Withheld Tweet:

Tweet withheld 18h
This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)

Withheld Media:



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Tip: Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

Tip: We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Reporting sensitive media

If you encounter media in Tweets that you believe should be treated as sensitive under [Twitter's media policy](#), please report it using the process described below.

How do I report media in Tweets?

Follow these instructions to report media:

1. From the Tweet you would like to report, click or tap the **More icon** (••• on web or  on iOS;  on Android).
2. Select **Report**.
3. Select **It displays a sensitive [image/video/media]**.
4. Next we'll provide recommendations for additional actions you can take to improve your Twitter experience.

Please remember that if you are seeing something you don't like, and Twitter hasn't placed a warning label before it, it's possible that it doesn't meet our threshold for a warning on the media.

What happens to sensitive media I report?

The Twitter team reviews reports of media flagged by users to determine if that media requires a warning message in order to comply with Twitter's media policies. Please remember that by flagging media you are bringing it to the attention of the Twitter team. The flagged content will not automatically receive a warning message or be removed from the site.

For information about what you can do if your content is reported as sensitive, visit this [article](#).

How to control whether you see sensitive media in Tweets:

You can choose to view media in Tweets that may contain sensitive content without a warning. If you are concerned about viewing sensitive content, you should keep your default settings. Your default settings will help by providing warnings before seeing content that has been marked as possibly sensitive.

1. Log in to Twitter and go to your [Account settings](#).
2. Under the **Content** section, look for **Tweet media** and check the box next to **Do not inform me before showing media that may be sensitive**.
3. Save your settings at the bottom of the page.

Note: Remember you can also [unfollow](#) or [block](#) users that are disrupting your Twitter experience.

How to report content that is illegal or otherwise violates the Twitter Rules

For information outlining other types of violations and how you can report them to us, please see this [article](#).

Private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Some examples of private and confidential information include:

credit card information
social security or other national identity numbers
addresses or locations that are considered and treated as private
non-public, personal phone numbers
non-public, personal email addresses
images or videos that are considered and treated as private under applicable laws intimate
photos or videos that were taken or distributed without the subject's consent

Keep in mind that although you may consider certain information to be private, not all postings of such information may be a violation of this policy. We may consider the context and nature of the information posted, local privacy laws, and other case-specific facts when determining if this policy has been violated. For example, if information was previously posted or displayed elsewhere on the Internet prior to being put on Twitter, it may not be a violation of this policy.

For frequently asked questions about reporting private information posted on Twitter, [click here](#). To learn more about protecting your private information on Twitter and other websites, [click here](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter Impersonation Policy.

An account will not be removed if

- the user shares your name but has no other commonalities, or
- the profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same background or avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to [Twitter's Parody Policy](#) for more information about these types of accounts.

For Frequently Asked Questions about reporting impersonation on Twitter, [click here](#).

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or

squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

[Automation Rules & Best Practices](#)

Submitted Nov 08, 2009 by [ginger](#)

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. All users must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in the temporary locking and/or permanent suspension of account(s).

Please note that we may need to change these rules from time to time and reserve the right to do so. The most current version will always be available at twitter.com/rules.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter.

- **Trademark:** We reserve the right to reclaim usernames on behalf of businesses or individuals that hold legal claim or trademark on those usernames. Accounts using business names and/or logos to mislead others may be permanently suspended.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Graphic content:** You may not use pornographic or excessively violent media in your profile image or header image. Twitter may allow some forms of graphic content in Tweets marked as [sensitive media](#). When content crosses the line into gratuitous images of death, Twitter may ask that you remove the content out of respect for the deceased.
- **Unlawful use:** You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter badges:** You may not use badges, such as but not limited to the “promoted” or “verified” Twitter badge, unless provided by Twitter. Accounts using these badges as part of profile photos, header photos, or in a way that falsely implies affiliation with Twitter, may be suspended.

Abusive Behavior

We believe in freedom of expression and in speaking truth to power, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we do not tolerate behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Any accounts and related accounts engaging in the activities specified below may be temporarily locked and/or subject to permanent suspension.

- **Violent threats (direct or indirect):** You may not make threats of violence or promote violence, including threatening or promoting terrorism.

- **Harassment:** You may not incite or engage in the targeted abuse or harassment of others. Some of the factors that we may consider when evaluating abusive behavior include:
 - if a primary purpose of the reported account is to harass or send abusive messages to others;
 - if the reported behavior is one-sided or includes threats;
 - if the reported account is inciting others to harass another account; and
 - if the reported account is sending harassing messages to an account from multiple accounts.
- **Hateful conduct:** You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.
- **Multiple account abuse:** Creating multiple accounts with overlapping uses or in order to evade the temporary or permanent suspension of a separate account is not allowed.
- **Private information:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address, or Social Security/National Identity numbers, without their express authorization and permission. In addition, you may not post intimate photos or videos that were taken or distributed without the subject's consent. Read more about our private information policy [here](#).
- **Impersonation:** You may not impersonate others through the Twitter service in a manner that is intended to or does mislead, confuse, or deceive others. Read more about our impersonation policy [here](#).
- **Self-harm:** You may encounter someone considering suicide or self harm on Twitter. When we receive reports that a person is threatening suicide or self harm, we may take a number of steps to assist them, such as reaching out to that person expressing our concern and the concern of other users on Twitter or providing resources such as contact information for our mental health partners.

Spam

We strive to protect people on Twitter from technical abuse and spam. Any accounts engaging in the activities specified below may be temporarily locked or subject to permanent suspension.

- **Username squatting:** You may not engage in username squatting. Accounts that are inactive for more than six months may also be removed without further notice. Some of the factors we take into consideration when determining what conduct is considered to be username squatting are:
 - the number of accounts created;
 - creating accounts for the purpose of preventing others from using those account names;
 - creating accounts for the purpose of selling those accounts; and

- using feeds of third-party content to update and maintain accounts under the names of those third parties.
- **Invitation spam:** You may not use twitter.com's address book contact import to send repeat, mass invitations.
- **Selling usernames:** You may not buy or sell Twitter usernames.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - if you have followed and/or unfollowed large amounts of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - if your updates consist mainly of links, and not personal updates;
 - if a large number of people are blocking you;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicate content over multiple accounts or multiple duplicate updates on one account;
 - if you post multiple unrelated updates to a topic using #, trending or popular topic, or promoted trend;
 - if you send large numbers of duplicate replies or mentions;
 - if you send large numbers of unsolicited replies or mentions;
 - if you add a large number of unrelated users to lists;
 - if you repeatedly create false or misleading content;
 - if you are randomly or aggressively following, liking, or Retweeting Tweets;
 - if you repeatedly post other people's account information as your own (bio, Tweets, URL, etc.);
 - if you post misleading links (e.g. affiliate links, links to malware/clickjacking pages, etc.);
 - if you are creating misleading accounts or account interactions;
 - if you are selling or purchasing account interactions (such as selling or purchasing followers, Retweets, likes, etc.); and

- if you are using or promoting third-party services or apps that claim to get you more followers (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers to your account).

See our support articles on [Following rules and best practices](#) and [Automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts under investigation may be removed from search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the [Terms of Service](#).

Trademark policy

Updated June 30, 2016

What is a Trademark Policy Violation on Twitter?

Using another's trademark in a manner that may mislead or confuse others about your brand affiliation may be a violation of our trademark policy.

How Does Twitter Respond to Trademark Reports?

We review each report and do the following:

- If we determine that there is a violation of our trademark policy, we may suspend the account. In some instances, we may give the account holder an opportunity to comply with our policies before suspending the account. For more information, see our [Parody, commentary, and fan account policy](#).
- We may also provide the account holder with your name and information included in the copy of the report.
- If you would like to report a trademark violation in a [Twitter Ad](#), please visit our [Trademark Policy for Promoted Ads](#) to learn more.

Please note that Twitter only investigates requests that are submitted by the trademark holder or by an authorized representative of the trademark holder.

What is not a Trademark Policy Violation?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples include:

- Using a trademark in a way that is outside the scope of the trademark registration (e.g. territory, or goods and services identified in the registration).
- Nominative and other fair uses of trademarks are protected uses under our trademark policy, so long as the account is clearly distinguished from the trademark owner. This includes use by resellers in certain regions and accounts engaging in parody, commentary, or news. For more information, see our [Parody, commentary, and fan account policy](#).

How do I Report a Trademark Policy Violation?

Submit trademark reports through our [Trademark Issue support form](#). Using the appropriate form ensures we have all the information necessary to process your report quickly. Please provide all information requested in the form. If you submit a report with incomplete information, we will follow up with you to request that information. Please note that this will result in a delay in processing your report.

If you have trouble using the form, try updating your browser or switching to another browser.

Copyright policy

Topics Covered:

- [What Types of Copyright Complaints Does Twitter Respond To?](#)
- [Am I a Copyright Holder? How Do I Know?](#)
- [What Information Do You Need To Process a Copyright Complaint?](#)
- [How Do I File a Copyright Complaint?](#)
- [How Are Claims Processed?](#)
- [What Happens Next?](#)
- [What Happens If I Receive a Copyright Notification?](#)
- [When Should I File a Counter-notice?](#)
- [What Information Do You Need to Process a Counter-notice?](#)
- [What Happens After I Submit a Counter-notice?](#)
- [Filing a Copyright Notice or Counter-notice Is Serious Business!](#)

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile photo, header photo, or background, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [Fair Use](#) page for more information).

Tip: If you are concerned about the use of your brand or entity's name, please review Twitter's [Trademark Policy](#). If you are concerned about the use of a fictional character, please see our [Parody, Commentary, and Fan Accounts Policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

Tip: If you are reporting the content of a Tweet, please give us a direct link to that Tweet following these instructions: <https://support.twitter.com/articles/80586>. Or please specify if the alleged infringement is in the background, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to Twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether

the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with access to the reporter's complaint along with instructions on how to file a counter-notice, and forward a copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a copy of the takedown notice, along with instructions for filing a counter-notification.

Tip: If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a copy of each copyright complaint and counter-notice that we process to [Lumen](#), where they are posted to a public-facing website (with your personal information removed).

Withheld Tweet:

Tweet withheld 18h
This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)

Withheld Media:



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Tip: Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

Tip: We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Reporting sensitive media

If you encounter media in Tweets that you believe should be treated as sensitive under [Twitter's media policy](#), please report it using the process described below.

How do I report media in Tweets?

Follow these instructions to report media:

1. From the Tweet you would like to report, click or tap the **More icon** (••• on web or iOS;  on Android).
2. Select **Report**.
3. Select **It displays a sensitive [image/video/media]**.
4. Next we'll provide recommendations for additional actions you can take to improve your Twitter experience.

Please remember that if you are seeing something you don't like, and Twitter hasn't placed a warning label before it, it's possible that it doesn't meet our threshold for a warning on the media.

What happens to sensitive media I report?

The Twitter team reviews reports of media flagged by users to determine if that media requires a warning message in order to comply with Twitter's media policies. Please remember that by flagging media you

are bringing it to the attention of the Twitter team. The flagged content will not automatically receive a warning message or be removed from the site.

For information about what you can do if your content is reported as sensitive, visit this [article](#).

How to control whether you see sensitive media in Tweets:

You can choose to view media in Tweets that may contain sensitive content without a warning. If you are concerned about viewing sensitive content, you should keep your default settings. Your default settings will help by providing warnings before seeing content that has been marked as possibly sensitive.

1. Log in to Twitter and go to your [Account settings](#).
2. Under the **Content** section, look for **Tweet media** and check the box next to **Do not inform me before showing media that may be sensitive**.
3. Save your settings at the bottom of the page.

Note: Remember you can also [unfollow](#) or [block](#) users that are disrupting your Twitter experience.

How to report content that is illegal or otherwise violates the Twitter Rules

For information outlining other types of violations and how you can report them to us, please see this [article](#).

Private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Some examples of private and confidential information include:

credit card information
social security or other national identity numbers
addresses or locations that are considered and treated as private
non-public, personal phone numbers
non-public, personal email addresses
images or videos that are considered and treated as private under applicable laws
intimate photos or videos that were taken or distributed without the subject's consent

Keep in mind that although you may consider certain information to be private, not all postings of such information may be a violation of this policy. We may consider the context and nature of the information posted, local privacy laws, and other case-specific facts when determining if this policy has been violated. For example, if information was previously posted or displayed elsewhere on the Internet prior to being put on Twitter, it may not be a violation of this policy.

For frequently asked questions about reporting private information posted on Twitter, [click here](#). To learn more about protecting your private information on Twitter and other websites, [click here](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter Impersonation Policy.

An account will not be removed if

- the user shares your name but has no other commonalities, or
- the profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same background or avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to [Twitter's Parody Policy](#) for more information about these types of accounts.

For Frequently Asked Questions about reporting impersonation on Twitter, [click here](#).

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation Rules & Best Practices

Submitted Nov 08, 2009 by [ginger](#)

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password.

Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

The Twitter Rules (20160504*)

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. All users must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in the temporary locking and/or permanent suspension of account(s).

Please note that we may need to change these rules from time to time and reserve the right to do so. The most current version will always be available at twitter.com/rules.

Content Boundaries and Use of Twitter

In order to provide the Twitter service and the ability to communicate and stay connected with others, there are some limitations on the type of content that can be published with Twitter.

- **Trademark:** We reserve the right to reclaim usernames on behalf of businesses or individuals that hold legal claim or trademark on those usernames. Accounts using business names and/or logos to mislead others may be permanently suspended.
- **Copyright:** We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in the Terms of Service.
- **Graphic content:** You may not use pornographic or excessively violent media in your profile image or header image. Twitter may allow some forms of graphic content in Tweets marked as [sensitive media](#). When content crosses the line into gratuitous images of death, Twitter may ask that you remove the content out of respect for the deceased.
- **Unlawful use:** You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.
- **Misuse of Twitter badges:** You may not use badges, such as but not limited to the “promoted” or “verified” Twitter badge, unless provided by Twitter. Accounts using these badges as part of profile photos, header photos, or in a way that falsely implies affiliation with Twitter, may be suspended.

Abusive Behavior

We believe in freedom of expression and in speaking truth to power, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we do not tolerate behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Any accounts and related accounts engaging in the activities specified below may be temporarily locked and/or subject to permanent suspension.

- **Violent threats (direct or indirect):** You may not make threats of violence or promote violence, including threatening or promoting terrorism.

- **Harassment:** You may not incite or engage in the targeted abuse or harassment of others. Some of the factors that we may consider when evaluating abusive behavior include:
 - if a primary purpose of the reported account is to harass or send abusive messages to others;
 - if the reported behavior is one-sided or includes threats;
 - if the reported account is inciting others to harass another account; and
 - if the reported account is sending harassing messages to an account from multiple accounts.
- **Hateful conduct:** You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.
- **Multiple account abuse:** Creating multiple accounts with overlapping uses or in order to evade the temporary or permanent suspension of a separate account is not allowed.
- **Private information:** You may not publish or post other people's private and confidential information, such as credit card numbers, street address, or Social Security/National Identity numbers, without their express authorization and permission. In addition, you may not post intimate photos or videos that were taken or distributed without the subject's consent. Read more about our private information policy [here](#).
- **Impersonation:** You may not impersonate others through the Twitter service in a manner that is intended to or does mislead, confuse, or deceive others. Read more about our impersonation policy [here](#).
- **Self-harm:** You may encounter someone considering suicide or self harm on Twitter. When we receive reports that a person is threatening suicide or self harm, we may take a number of steps to assist them, such as reaching out to that person expressing our concern and the concern of other users on Twitter or providing resources such as contact information for our mental health partners.

Spam

We strive to protect people on Twitter from technical abuse and spam. Any accounts engaging in the activities specified below may be temporarily locked or subject to permanent suspension.

- **Username squatting:** You may not engage in username squatting. Accounts that are inactive for more than six months may also be removed without further notice. Some of the factors we take into consideration when determining what conduct is considered to be username squatting are:
 - the number of accounts created;
 - creating accounts for the purpose of preventing others from using those account names;
 - creating accounts for the purpose of selling those accounts; and

- using feeds of third-party content to update and maintain accounts under the names of those third parties.
- **Invitation spam:** You may not use twitter.com's address book contact import to send repeat, mass invitations.
- **Selling usernames:** You may not buy or sell Twitter usernames.
- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use the Twitter service for the purpose of spamming anyone. What constitutes "spamming" will evolve as we respond to new tricks and tactics by spammers. Some of the factors that we take into account when determining what conduct is considered to be spamming are:
 - if you have followed and/or unfollowed large amounts of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if you repeatedly follow and unfollow people, whether to build followers or to garner more attention for your profile;
 - if your updates consist mainly of links, and not personal updates;
 - if a large number of people are blocking you;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicate content over multiple accounts or multiple duplicate updates on one account;
 - if you post multiple unrelated updates to a topic using #, trending or popular topic, or promoted trend;
 - if you send large numbers of duplicate replies or mentions;
 - if you send large numbers of unsolicited replies or mentions;
 - if you add a large number of unrelated users to lists;
 - if you repeatedly create false or misleading content;
 - if you are randomly or aggressively following, liking, or Retweeting Tweets;
 - if you repeatedly post other people's account information as your own (bio, Tweets, URL, etc.);
 - if you post misleading links (e.g. affiliate links, links to malware/clickjacking pages, etc.);
 - if you are creating misleading accounts or account interactions;
 - if you are selling or purchasing account interactions (such as selling or purchasing followers, Retweets, likes, etc.); and

- if you are using or promoting third-party services or apps that claim to get you more followers (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers to your account).

See our support articles on [Following rules and best practices](#) and [Automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts will be permanently suspended.

Accounts under investigation may be removed from search for quality. Twitter reserves the right to immediately terminate your account without further notice in the event that, in its judgment, you violate these Rules or the [Terms of Service](#).

Trademark policy (20160630)

Updated June 30, 2016

What is a Trademark Policy Violation on Twitter?

Using another's trademark in a manner that may mislead or confuse others about your brand affiliation may be a violation of our trademark policy.

How Does Twitter Respond to Trademark Reports?

We review each report and do the following:

- If we determine that there is a violation of our trademark policy, we may suspend the account. In some instances, we may give the account holder an opportunity to comply with our policies before suspending the account. For more information, see our [Parody, commentary, and fan account policy](#).
- We may also provide the account holder with your name and information included in the copy of the report.
- If you would like to report a trademark violation in a [Twitter Ad](#), please visit our [Trademark Policy for Promoted Ads](#) to learn more.

Please note that Twitter only investigates requests that are submitted by the trademark holder or by an authorized representative of the trademark holder.

What is not a Trademark Policy Violation?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples include:

- Using a trademark in a way that is outside the scope of the trademark registration (e.g. territory, or goods and services identified in the registration).
- Nominative and other fair uses of trademarks are protected uses under our trademark policy, so long as the account is clearly distinguished from the trademark owner. This includes use by resellers in certain regions and accounts engaging in parody, commentary, or news. For more information, see our [Parody, commentary, and fan account policy](#).

How do I Report a Trademark Policy Violation?

Submit trademark reports through our [Trademark Issue support form](#). Using the appropriate form ensures we have all the information necessary to process your report quickly. Please provide all information requested in the form. If you submit a report with incomplete information, we will follow up with you to request that information. Please note that this will result in a delay in processing your report.

If you have trouble using the form, try updating your browser or switching to another browser.

Copyright policy (20160206*)

Topics Covered:

- [What Types of Copyright Complaints Does Twitter Respond To?](#)
- [Am I a Copyright Holder? How Do I Know?](#)
- [What Information Do You Need To Process a Copyright Complaint?](#)
- [How Do I File a Copyright Complaint?](#)
- [How Are Claims Processed?](#)
- [What Happens Next?](#)
- [What Happens If I Receive a Copyright Notification?](#)
- [When Should I File a Counter-notice?](#)
- [What Information Do You Need to Process a Counter-notice?](#)
- [What Happens After I Submit a Counter-notice?](#)
- [Filing a Copyright Notice or Counter-notice Is Serious Business!](#)

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile photo, header photo, or background, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [Fair Use](#) page for more information).

Tip: If you are concerned about the use of your brand or entity's name, please review Twitter's [Trademark Policy](#). If you are concerned about the use of a fictional character, please see our [Parody, Commentary, and Fan Accounts Policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

Tip: If you are reporting the content of a Tweet, please give us a direct link to that Tweet following these instructions: <https://support.twitter.com/articles/80586>. Or please specify if the alleged infringement is in the background, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to Twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether

the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with access to the reporter's complaint along with instructions on how to file a counter-notice, and forward a copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a copy of the takedown notice, along with instructions for filing a counter-notification.

Tip: If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a copy of each copyright complaint and counter-notice that we process to [Lumen](#), where they are posted to a public-facing website (with your personal information removed).

Withheld Tweet:

Tweet withheld 18h
This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)

Withheld Media:



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Tip: Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

Tip: We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Reporting sensitive media (20151229*)

If you encounter media in Tweets that you believe should be treated as sensitive under [Twitter's media policy](#), please report it using the process described below.

How do I report media in Tweets?

Follow these instructions to report media:

1. From the Tweet you would like to report, click or tap the **More icon** (••• on web or  on iOS;  on Android).
2. Select **Report**.
3. Select **It displays a sensitive [image/video/media]**.
4. Next we'll provide recommendations for additional actions you can take to improve your Twitter experience.

Please remember that if you are seeing something you don't like, and Twitter hasn't placed a warning label before it, it's possible that it doesn't meet our threshold for a warning on the media.

What happens to sensitive media I report?

The Twitter team reviews reports of media flagged by users to determine if that media requires a warning message in order to comply with Twitter's media policies. Please remember that by flagging media you are bringing it to the attention of the Twitter team. The flagged content will not automatically receive a warning message or be removed from the site.

For information about what you can do if your content is reported as sensitive, visit this [article](#).

How to control whether you see sensitive media in Tweets:

You can choose to view media in Tweets that may contain sensitive content without a warning. If you are concerned about viewing sensitive content, you should keep your default settings. Your default settings will help by providing warnings before seeing content that has been marked as possibly sensitive.

1. Log in to Twitter and go to your [Account settings](#).
2. Under the **Content** section, look for **Tweet media** and check the box next to **Do not inform me before showing media that may be sensitive**.
3. Save your settings at the bottom of the page.

Note: Remember you can also [unfollow](#) or [block](#) users that are disrupting your Twitter experience.

How to report content that is illegal or otherwise violates the Twitter Rules

For information outlining other types of violations and how you can report them to us, please see this [article](#).

Private information posted on Twitter (20170714*)

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Some examples of private and confidential information include, but may not be limited to:

Intimate photos or videos that were taken or distributed without the subject's consent.
Images or videos that are considered and treated as private under applicable laws. Private contact or financial information, such as:

credit card information
social security or other national identity numbers
addresses or locations that are considered and treated as private non-public, personal phone numbers
non-public, personal email addresses

How our enforcement works

Keep in mind that although you may consider certain information to be private, not all postings of such information may be a violation of this policy.

Context matters.

We may consider the context and nature of the information posted, local privacy laws, and other case-specific facts when determining if this policy has been violated. For example, if information was previously posted or displayed elsewhere on the Internet prior to being put on Twitter, it may not be a violation of this policy.

We focus on behavior.

We enforce policies when someone reports content where the behavior may be abusive.

We have a range of enforcement options.

The consequences for violating our rules vary depending on the severity of the violation and the person's previous record of violations. For example, we may ask someone to remove the offending Tweet before they can Tweet again. For other cases, we may suspend an account.

For frequently asked questions about reporting private information posted on Twitter, [click here](#). To learn more about protecting your private information on Twitter and other websites, [click here](#).

Impersonation policy (201305*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter Impersonation Policy.

An account will not be removed if

- the user shares your name but has no other commonalities, or
- the profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same background or avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to [Twitter's Parody Policy](#) for more information about these types of accounts.

For Frequently Asked Questions about reporting impersonation on Twitter, [click here](#).

Username squatting policy (20130817*)

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Following rules and best practices (20160117*)

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, [we do monitor how aggressively users follow other users](#). We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then

unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation Rules & Best Practices (2010/01)

Submitted Nov 08, 2009 by [ginger](#)

We're constantly amazed by the applications and services that grow up around the Twitter platform. External developers have created tools and applications enabling an incredible array of uses; we've seen Twitter accounts that help people stay updated on their [city's emergency services](#), connect to [national health alerts](#), and catch up on the happenings of their [local bridge](#).

Most users will not need to review this document; however, some people are keen on automation. Unfortunately, spammers have also taken advantage of automation, to the frustration of genuine users. This page has

information on what account automations are allowed and which ones could get your account filtered from search or even suspended or terminated.

How Automation Works

In order for an outside party to post to your account or perform other account actions (like following), you have to grant them access by either establishing a Twitter Connection or by giving them your username and password. Please see our help page on [Connecting to Third-Party Applications](#) for information on the difference, and how to revoke access for each method.

Ultimately, you are responsible for the updates to your timeline and the actions of your account, which means you shouldn't give control of your account to a third-party unless you've thoroughly investigated the application. If you decide to automate any account actions, you will need to manage your automations carefully to avoid violating the [Twitter Rules](#). Violating these rules may result in account suspension or termination, or your account being filtered from search, regardless of whether you performed the violation manually, or allowed an automation to do so!

Automating your Tweets

Automating Your Tweets Based on an Action You've Completed:

Maybe you want to post to your Twitter timeline whenever you upload a video to YouTube, create a new photo album, or beat a high score in your favorite game? Sounds great! It's generally fine to update your Twitter account when you complete a particular action. Just be careful—if you post similar updates too frequently, your followers might get a little bored!

Automating Your Tweets by Streaming External Information:

If you want to automate your account with a feed from your personal blog or updates from your RSS feed, please consider the following:

- **Feeds of your own blog:** It's generally fine to automate your account with updates of your own content; just be sure to review the [Twitter Rules](#) before setting up your automation.
- **Feeds that go through redirects or ad pages:** If you're posting automated links that redirect through landing or ad pages before the final content, your account may be suspended for posting misleading links, a violation of the [Twitter Rules](#).
- **Feeds for community benefit:** We welcome feeds that are used for community benefit or provide non-commercial information to a niche group of users, such as local weather feeds or transit information.

Automatically Tweeting to Trending Topics:

Trending topics are the most tweeted-about topics on Twitter. Posting to current trending topics in an automated fashion can degrade the experience for other users, and may result in your account being automatically [filtered from search](#). If this happens, your followers will still see your tweets but you won't show in search results.

Automating your tweets to include references to **unrelated** trending topics will result in account suspension.

Mass-creating Automated Accounts:

There are legitimate reasons why users may have multiple accounts. Serial account creation without prior approval, however, is not allowed. Please [file a support ticket](#) if you would like permission for creating a series of accounts. Include the list of accounts and your planned use of each account; accounts with overlapping use cases and names are generally not approved.

We have allowed multiple, automated accounts for community benefit. For example, Twitter users can now find out when the Hubble passes overhead, whether they're in [San Francisco](#) or [Hong Kong](#). If the content is available elsewhere on Twitter, however, and does not clearly benefit a community or provide unique information to a niche group of users, these accounts may not be allowed.

Automated Advertising Updates:

We are continuing to consider the issue of advertising on the Twitter platform. We're committed to encouraging and supporting individuals and businesses who want to discuss and promote the products they care about. At the same time, we've also seen an increase in the number of fully-automated, spam accounts promoting affiliate offers. These accounts use automated following, automated @replying, and serial account creation to spam other users with mass-created advertising. These behaviors are all prohibited by the [Twitter Rules](#).

The guidelines in this section are meant to both reduce automated spam on Twitter, and provide guidance for users interested in including advertising for others in their timeline.

Automated or mass-created affiliate advertising is not permitted on Twitter. Updates that are posted automatically to your account through a tool or third-party application, for which you are compensated (whether on a "per-click," "per-sale," "per-tweet," or other basis), are prohibited.

We're still thinking about the full spectrum of advertising on Twitter; in the meantime, other forms of advertising and promotion are allowed, but subject to the rules in this document. The following forms of advertising are not prohibited:

- Promoting your own business or website.
- Tweeting unpaid consumer recommendations.
- Tweeting sponsored or compensated links and updates that you have manually posted or approved.

Your tweets may be sponsored by a third-party if you manually post or approve each sponsored tweet before it is posted. These updates may not be automated or scheduled in advance. The following rules also apply:

1. All sponsored or paid updates must be manually approved. You should either manually post these to your account, or individually approve the tweets to be posted to your account.
2. You shouldn't repeatedly post the same ad updates.
3. If you're using a third-party tool or application to approve and post your tweets, it must be an application **with Twitter OAuth approval**. Please see the FAQ below for more information.
4. You should disclose when you post a compensated or sponsored update.

Users combining advertising with other automations should be particularly careful to review the [Twitter Rules](#) and [Follow Rules and Best Practices](#). Account suspensions may be appealed by [filing a support ticket](#). Finally, please do not attempt to circumvent our guidelines on automated advertising by creating lots of bot accounts that promote similar or duplicate websites (or websites with scraped content); these accounts may be permanently suspended or terminated.

Automating other Account Actions

When you're automating account behaviors like @replies, DMs, and retweets, we recommend you follow a general guideline: be nice to other Twitter users. In general, we take a skeptical view of disingenuous social action. If your actions are annoying a bunch of Twitter users, you probably will want to rethink or adjust your automation. Here are some more specific guidelines:

Automated @Replies:

The @reply function is intended to make communication between users easier, and automating this process to put unsolicited messages into lots of users' reply tabs is considered an abuse of feature. If you are automatically sending @reply messages to a bunch of users, the recipients must request or approve this action in advance. For example, sending automated @replies based on keyword searches is not permitted.

Users should also have an easy way to opt-out of your service (in addition to the requirement that all users must opt-in before receiving the messages). We review blocks and reports of spam, so you'll need to provide a clear way for users to stop your messages.

Automated Retweeting of Other Users:

Retweeting someone means you think your followers should see their tweet. When you're retweeting in an automated fashion, and especially based on an automated search or keyword, this becomes less true.

We discourage the automatic retweeting of other users based on a particular keyword and may suspend or terminate accounts that engage in this behavior, particularly if they are being frequently blocked and reported as spam by the users they are retweeting. If you are retweeting users just so they will see your update in their replies folder and look at your account, this is akin to @reply spam and your account may be suspended or terminated.

Automated Following and Un-Following:

Twitter users often review the profiles and recent tweets of their new followers. Being followed in a bulk, automated fashion can make this process bothersome and makes Twitter a less fun place to hang out.

The only auto-following behavior Twitter allows is auto-follow-back (following a user after they have followed you). Automated un-following is also not permitted. If you find yourself frequently needing to un-follow large numbers of users, you might consider reviewing your criteria for following them. Please review our [Follow Rules and Best Practices](#) for a detailed discussion of following recommendations.

Also note that in the general [Twitter Rules](#) that "get followers fast" applications and services are not allowed. Do not surrender your username and password to them.

Automating Your DMs (Direct Messages):

Including an automated "thanks for following" message to your new followers might be annoying to some users. We do not recommend, but generally do not regulate, this behavior; if you receive a DM you don't like, you can un-follow that user and they will no longer be able to send you messages.

Automation that Causes Your Account to Violate the Twitter Rules:

If your account automation is causing your account to violate the [Twitter Rules](#) (by retweeting spam updates, repeatedly posting duplicate links, etc.), your account may be suspended or terminated. You are responsible for the updates on your account, so please do not add automations unless you are confident you will be able to manage your account within the rules.

Other Questions (FAQ)

What is an application with Twitter OAuth approval?

As discussed above, you may post manually-approved, compensated tweets through a third-party application with Twitter OAuth approval. If an application has OAuth approval, the following will be true

1. You allow the application access by approving their connection, not by giving out your username and password.
2. Once approved, you'll be able to see and revoke the connection from your [connections tab](#).

When tweets come from a third-party application, the name of the source is included to the right of the timestamp, like this:

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. These limitations are set forth in the Twitter Rules below.

The Twitter Rules (along with all incorporated policies), [Privacy Policy](#), and [Terms of Service](#) collectively make up the "Twitter User Agreement" that governs a user's access to and use of Twitter's services.

All individuals accessing or using Twitter's services must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in Twitter taking one or more of the following enforcement actions:

- requiring you to delete prohibited content before you can again create new posts and interact with other Twitter users;
- temporarily limiting your ability to create posts or interact with other Twitter users;
- asking you to verify account ownership with a phone number or email address; or
- permanently suspending your account(s).

If you attempt to evade a permanent suspension by creating new accounts, we will suspend your new accounts.

Please note that we may need to change these Rules from time to time and reserve the right to do so. The most current version is always available at: <https://twitter.com/rules>.

The policies set forth in these Twitter Rules govern organic content on our platform. To learn more about the rules which govern ads and promoted content, please review our [Ads policies](#).

Content Boundaries and Use of Twitter

Intellectual property

Trademark: We reserve the right to suspend accounts or take other appropriate action when someone's brand or trademark, including business name and/or logo, is used in a manner that may mislead or confuse others about your brand affiliation. Read more about our [trademark policy and how to report a violation](#).

Copyright: We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in our [Terms of Service](#). Read more about our [copyright policy](#).

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Learn more about our [media policy](#).

Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images. Additionally, Twitter may sometimes require you to remove excessively graphic violence out of respect for the

deceased and their families if we receive a request from their family or an authorized representative. Learn more about [how to make such a request](#), and [how to mark your media as sensitive](#).

Unlawful use

You may not use our service for any unlawful purposes or in furtherance of illegal activities. By using Twitter, you agree to comply with all applicable laws governing your online conduct and content.

Trends

At times, we may prevent certain content from trending. This includes content that violates the Twitter Rules, as well as content that may attempt to manipulate trends. Read more about [what we allow and do not allow to trend](#).

Misuse of Twitter badges

You may not use badges, including but not limited to the “promoted” or “verified” Twitter badges, unless provided by Twitter. Accounts using unauthorized badges as part of their profile photos, header photos, display names, or in any way that falsely implies affiliation with Twitter or authorization from Twitter to display these badges, may be suspended.

Misuse of usernames

Selling usernames: You may not buy or sell Twitter usernames.

Username squatting: You may not engage in username squatting. Some of the factors we take into consideration when determining whether conduct is username squatting include:

- the number of accounts created;
- the creation of accounts for the purpose of preventing others from using those account names;
- the creation of accounts for the purpose of selling those accounts; and
- the use of third-party content feeds to update and maintain accounts under the names of those third parties.

Please note that Twitter may also remove accounts that are inactive for more than six months. Learn more about [username squatting](#).

Abusive Behavior

We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we prohibit behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Context matters when evaluating for abusive behavior and determining appropriate enforcement actions. Factors we may take into consideration include, but are not limited to whether:

- the behavior is targeted at an individual or group of people;

- the report has been filed by the target of the abuse or a bystander;
- the behavior is newsworthy and in the legitimate public interest.

Violence and physical harm

Violence: You may not make specific threats of violence or wish for the serious physical harm, death, or disease of an individual or group of people. This includes, but is not limited to, threatening or promoting terrorism.

Suicide or self-harm: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Child sexual exploitation: You may not promote child sexual exploitation. Learn more about our zero-tolerance [child sexual exploitation policy](#).

Abuse and hateful conduct

Abuse: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Unwanted sexual advances: You may not direct abuse at someone by sending unwanted sexual content, objectifying them in a sexually explicit manner, or otherwise engaging in sexual misconduct.

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. Read more about our [hateful conduct policy](#).

Private information and intimate media

Private information: You may not publish or post other people's private information without their express authorization and permission. Definitions of private information may vary depending on local laws. Read more about our [private information policy](#).

Intimate media: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. Read more about intimate media on Twitter.

Threats to expose / hack: You may not threaten to expose someone's private information or intimate media. You also may not threaten to hack or break into someone's digital information.

Impersonation

You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. While you may maintain parody, fan, commentary, or newsfeed accounts, you may not do so if the intent of the account is to engage in spamming or abusive behavior. Read more about our [impersonation policy](#).

Spam and Security

We strive to protect people on Twitter from technical abuse and spam.

To promote a stable and secure environment on Twitter, you may not do, or attempt to do, any of the following while accessing or using Twitter:

- Access, tamper with, or use non-public areas of Twitter, Twitter's computer systems, or the technical delivery systems of Twitter's providers (except as expressly permitted by the Twitter Bug Bounty program).
- Probe, scan, or test the vulnerability of any system or network, or breach or circumvent any security or authentication measures (except as expressly permitted by the Twitter Bug Bounty program).
- Access or search, or attempt to access or search, Twitter by any means (automated or otherwise) other than through our currently available, published interfaces that are provided by Twitter (and only pursuant to the applicable terms and conditions), unless you have been specifically allowed to do so in a separate agreement with Twitter. Note that crawling Twitter is permissible if done in accordance with the provisions of the robots.txt file; however, scraping Twitter without our prior consent is expressly prohibited.
- Forge any TCP/IP packet header or any part of the header information in any email or posting, or in any way use Twitter to send altered, deceptive, or false source-identifying information.
- Interfere with or disrupt the access of any user, host or network, including, without limitation, sending a virus, overloading, flooding, spamming, mail-bombing Twitter's services, or by scripting the creation of content in such a manner as to interfere with or create an undue burden on Twitter.

Any accounts engaging in the following activities may be temporarily locked or subject to permanent suspension:

- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use Twitter's services for the purpose of spamming anyone. Spam is generally defined on Twitter as bulk or aggressive activity that attempts to manipulate or disrupt Twitter or the experience of users on Twitter to drive traffic or attention to unrelated accounts, products, services, or initiatives. Some of the factors that we take into account when determining what conduct is considered to be spamming include:
 - if you have followed and/or unfollowed a large number of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if your Tweets or Direct Messages consist mainly of links shared without commentary;
 - if a large number of people have blocked you in response to high volumes of untargeted, unsolicited, or duplicative content or engagements from your account;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicative or substantially similar content, replies, or mentions over multiple accounts or multiple duplicate updates on one account, or create duplicate or substantially similar accounts;

- if you post multiple updates to a trending or popular topic with an intent to subvert or manipulate the topic to drive traffic or attention to unrelated accounts, products, services, or initiatives;
- if you send large numbers of unsolicited replies or mentions;
- if you add users to lists in a bulk or aggressive manner;
- if you are randomly or aggressively engaging with Tweets (e.g., likes, Retweets, etc.) or users (e.g., following, adding to lists or Moments, etc.) to drive traffic or attention to unrelated accounts, products, services, or initiatives;
- if you repeatedly post other people's account information as your own (e.g., bio, Tweets, profile URL, etc.);
- if you post misleading, deceptive, or malicious links (e.g., affiliate links, links to malware/clickjacking pages, etc.);
- if you create fake accounts, account interactions, or impressions;
- if you sell, purchase, or attempt to artificially inflate account interactions (such as followers, Retweets, likes, etc.); and
- if you use or promote third-party services or apps that claim to get you more followers, Retweets, or likes (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers or engagements to your account or Tweets).

Please see our support articles on [following rules and best practices](#) and [automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts may be permanently suspended.

Content Visibility

Accounts under investigation or which have been detected as sharing content in violation of these Rules may have their account or Tweet visibility limited in various parts of Twitter, including search. To learn more about situations in which content may be restricted on Twitter, please see our support article on [search rules and restrictions](#).

Trademark policy

Updated June 30, 2016

What is a Trademark Policy Violation on Twitter?

Using another's trademark in a manner that may mislead or confuse others about your brand affiliation may be a violation of our trademark policy.

How Does Twitter Respond to Trademark Reports?

We review each report and do the following:

- If we determine that there is a violation of our trademark policy, we may suspend the account. In some instances, we may give the account holder an opportunity to comply with our policies before suspending the account. For more information, see our [Parody, commentary, and fan account policy](#).
- We may also provide the account holder with your name and information included in the copy of the report.
- If you would like to report a trademark violation in a [Twitter Ad](#), please visit our [Trademark Policy for Promoted Ads](#) to learn more.

Please note that Twitter only investigates requests that are submitted by the trademark holder or by an authorized representative of the trademark holder.

What is not a Trademark Policy Violation?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples include:

- Using a trademark in a way that is outside the scope of the trademark registration (e.g. territory, or goods and services identified in the registration).
- Nominative and other fair uses of trademarks are protected uses under our trademark policy, so long as the account is clearly distinguished from the trademark owner. This includes use by resellers in certain regions and accounts engaging in parody, commentary, or news. For more information, see our [Parody, commentary, and fan account policy](#).

How do I Report a Trademark Policy Violation?

Submit trademark reports through our [Trademark Issue support form](#). Using the appropriate form ensures we have all the information necessary to process your report quickly. Please provide all information requested in the form. If you submit a report with incomplete information, we will follow up with you to request that information. Please note that this will result in a delay in processing your report.

If you have trouble using the form, try updating your browser or switching to another browser.

Copyright policy

Topics Covered:

- [What Types of Copyright Complaints Does Twitter Respond To?](#)
- [Am I a Copyright Holder? How Do I Know?](#)
- [What Information Do You Need To Process a Copyright Complaint?](#)
- [How Do I File a Copyright Complaint?](#)
- [How Are Claims Processed?](#)
- [What Happens Next?](#)

- [What Happens If I Receive a Copyright Notification?](#)
- [When Should I File a Counter-notice?](#)
- [What Information Do You Need to Process a Counter-notice?](#)
- [What Happens After I Submit a Counter-notice?](#)
- [Filing a Copyright Notice or Counter-notice Is Serious Business!](#)

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile photo, header photo, or background, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [Fair Use](#) page for more information).

Tip: If you are concerned about the use of your brand or entity’s name, please review Twitter’s [Trademark Policy](#). If you are concerned about the use of a fictional character, please see our [Parody, Commentary, and Fan Accounts Policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);

3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

Tip: If you are reporting the content of a Tweet, please give us a direct link to that Tweet following these instructions: <https://support.twitter.com/articles/80586>. Or please specify if the alleged infringement is in the background, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to Twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with access to the reporter's complaint along with instructions on how to file a counter-notice, and forward a copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a copy of the takedown notice, along with instructions for filing a counter-notification.

Tip: If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been

withheld (examples below). We also send a copy of each copyright complaint and counter-notice that we process to [Lumen](#), where they are posted to a public-facing website (with your personal information removed).

Withheld Tweet:

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)

Withheld Media:



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Tip: Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);

2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

Tip: We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Twitter media policy

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to easily share your content with others. Below are our guidelines for making sure your content is labeled appropriately.

Media on Twitter

You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.

In addition, you may not feature graphic content (such as media containing pornography or excessive violence) in live video, or in your profile image or header image.

If you see media you think may be in violation of our policy, see [this article](#) for information about how to report it.

Your media settings

All Twitter accounts have media settings to help Twitter understand the type of media you are posting and to help us assist others to find the type of content that they're interested in.

By appropriately marking your media settings, Twitter can identify potentially sensitive content that others users may not wish to see, such as violence or nudity. If you intend to regularly post such content, we request that you adjust your media setting:

1. Log in to your account on the web at twitter.com.
2. Click your **profile icon** in the top right navigation bar.
3. Select **Settings and privacy** from the drop-down menu.
4. Go to your [Privacy and safety](#) settings.
5. Look for the **Tweet media** section and check the box next to **Mark media you Tweet as containing material that may be sensitive**. **Note:** Enabling this setting will result in others seeing a message when they view your profile, alerting them of potentially sensitive media and asking them to confirm they still want to view it.

You can control your preferences for how and when you may see sensitive content in your [media settings](#).

What if I don't change my media settings and post sensitive media?

If one of your Tweets containing media is reported as sensitive, it will be sent to Twitter for review. If we find that the media wasn't marked as being potentially sensitive at the time of the upload:

- We label the media as being potentially sensitive, or in the case of live video, remove the content entirely.
- We may also change your account setting to **Mark media you Tweet as containing material that may be sensitive** so that future posts are marked accordingly.

You will receive an email letting you know that we've labeled your media and if we changed your media setting. You will be able to change your media settings on your account settings page, though this image and any other images that have been labeled will remain marked as potentially sensitive.

If you repeatedly upload media that are mislabeled, the Twitter team may:

- Label your media as having potentially sensitive content.
- Permanently adjust your account setting to **Mark media you Tweet as containing material that may be sensitive** so that your future posts have a warning message users must click through before viewing your media.

Twitter may also use automated techniques to detect and label potentially sensitive media.

Additionally, people who visit your profile may see a message letting them know the account may include potentially sensitive content and asking them to confirm they still want to view it. People who have opted in to see possibly sensitive content will still see your account without the message.

What happens if my profile image is in violation of the Twitter Rules?

If your profile or header image contains graphic content in violation of the Twitter Rules, we may suspend your account and remove the offending content. Repeat violations will result in a permanent suspension.

Additionally, your account may be locked if media in your profile is in violation of Twitter's media policy. In order to unlock your account, you will need to follow our instructions and delete the media in violation. If your account has been locked, see this [article](#) for more information.

Can I appeal a decision made by the Twitter team?

If you are suspended for having a pornographic profile, header, or background image, you can appeal by logging in to your account and submitting an appeal [here](#).

There is currently no way to appeal a decision by the Twitter team that permanently changes your account setting to **Mark media you Tweet as containing material that may be sensitive** in response to repeated mislabeling of sensitive content.

Contacting Twitter about media concerning a deceased family member

In order to respect the wishes of loved ones, Twitter will remove imagery of deceased individuals in certain circumstances. Immediate family members and other authorized individuals may request the removal of images or video of deceased individuals, from when critical injury occurs to the moments before or after death, by sending a request to Twitter Inc. via our [privacy form](#). When reviewing such media removal requests, Twitter considers public interest factors such as the newsworthiness of the content and may not be able to honor every request.

Note: For information about contacting Twitter to deactivate the account of a deceased or incapacitated user, please read [this article](#).

How to mark media as sensitive links to Media policy

FAQs about trends on Twitter

Where can I find trends?

Trends are available on the Twitter for iOS app, the Twitter for Android app, and twitter.com.

On Twitter's mobile apps, you can find trends listed under the trends section



of the **Explore** tab When signed in to twitter.com on a desktop or laptop computer, trends are listed in many places, including the Home timeline, Notifications, search results, and profile pages.

How are trends determined?

Trends are determined by an algorithm and, by default, are tailored for you based on who you follow, your interests, and your location. This algorithm identifies topics that are popular now, rather than topics that have been popular for a while or on a daily basis, to help you discover the hottest emerging topics of discussion on Twitter.

Note: The number of Tweets that are related to the trends is just one of the factors the algorithm looks at when ranking and determining trends. Algorithmically, trends and hashtags are grouped together if they are related to the same topic. For instance, #MondayMotivation and #MotivationMonday may both be represented by #MondayMotivation.

You can choose to see trends that are not tailored for you by selecting a [specific trends location](#) on twitter.com, iOS, or Android. Location trends identify popular topics among people in a specific geographic location.

What does the # sign mean?

You may notice that some trends have # sign before the word or phrase. This is called a [hashtag](#) and is included specifically in Tweets to mark them as relating to a topic, so that people can follow the conversation in search.

Will I see world and local events in my trends?

Yes, when signed into your account on the twitter.com, iOS, or Android, trends will be tailored for you based on your location and who you follow. There will be many world and local news events and conversations that will appear in your trends regardless of your personalization.

Can I see trends for a specific location?

Yes, you can view trends for a specific location on twitter.com, iOS, or Android.

View trends via twitter.com:

To view trends for a specific location, click **Change** in your trends box. You can choose to **Keep tailored trends** based on your location and who you follow, or you can select **Change** to choose a nearby location or search locations.

Adjust your trends settings via Twitter for iOS:

1. In the top menu, tap your **profile** icon.
2. Tap **Settings and privacy**, then tap **Content preferences**.
3. Under **Explore**, tap **Trends**.
4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.
5. When the feature is disabled, you can change your location by tapping **Change location**.
 1. In the search box, type your desired trends location or choose from the list.

Adjust your trends settings via Twitter for Android:



1. In the top menu, you will either see a **navigation menu** icon or your **profile** icon. Tap whichever icon you have, and select **Settings and privacy**.
2. Tap **Content preferences**.
3. Under **Explore**, tap **Trends**.
4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.
5. When the feature is disabled, you can change your location by tapping **Change location**.
 1. In the search box, type your desired trends location or choose from the list.

Note: If you don't find your preferred country or city, it means we're not receiving enough Tweets from that geographical area to create a list. You can look up local Tweets on any topic by using [advanced Twitter Search operators](#).

What happens when I click or tap on a trend?

Clicking or tapping any of the trends takes you to the Twitter search results for that trend. You'll see all Tweets including that phrase or hashtag. To see what people are saying about a previous trend, perform a [search for that keyword](#).

How can I participate in a trend?

Simply post a Tweet including the exact word or phrase as it appears in the trends list (with the hashtag, if you see one). Due to the large number of people Tweeting about these specific trends, you may not always be able to find your particular Tweet in search, but your followers will always see your Tweets.

Note: Twitter also filters searches for quality. To understand what types of behavior could cause your Tweets to be filtered out, read about our [search rules and restrictions](#).

Are there rules for trends?

Yes. We want trends to promote healthy discussions on Twitter. This means that at times, we may prevent certain content from trending. These include trends that:

- Contain profanity or adult/graphic references.

- Incite hate on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease.
- Violate the [Twitter Rules](#).

In some cases, we may also consider the newsworthiness of the content, or if it is in the public interest when evaluating potential violations. In these cases, the content might continue to trend on our platform.

Note that even if we prevent the hashtag or the content from appearing on the list of trending topics, you may still be able to access conversations around that content on Twitter.

Trends descriptions

Below some trends, you may see the approximate number of Tweets associated with the trend, or personalized context such as who in your network is Tweeting about it.

Additionally, when a Moment is found to be relevant to a trend, we may use the title of the Moment as the trend's description.

Having trouble?

Check out our [troubleshooting section](#) for solutions to common problems.

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Child sexual exploitation policy

Twitter removes content promoting child sexual exploitation

We do not tolerate child sexual exploitation on Twitter. When we are made aware of links to images of or content promoting child sexual exploitation they will be removed from the site without further notice and reported to The National Center for Missing & Exploited Children ("NCMEC"); we permanently suspend accounts promoting or containing updates with links to child sexual exploitation.

Reporting content promoting child sexual exploitation

If you think you've found a Twitter account distributing or promoting child sexual exploitation while using Twitter, please click [here](#) to report it. Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. To find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with the National Center for Missing and Exploited Children [here](#), or your local law enforcement. If you believe that the content was hosted on a website outside of the United States, you can report it [here](#) on the International Association of Internet Hotlines website.

NOTE: Please do NOT tweet, retweet or repost child sexual exploitation for any reason. Click [here](#) to report it to us immediately and we will take steps to remove it.

Law Enforcement Agencies

If you are a member of a law enforcement agency seeking information regarding Twitter accounts promoting or distributing child sexual exploitation, please see our [Guidelines for Law Enforcement](#) for information on our process; we also suggest you contact the National Center for Missing and Exploited Children (NCMEC) [here](#), regardless of jurisdiction, to determine whether or not information on those accounts is already on file.

For general law enforcement inquiries, please see our [Guidelines for Law Enforcement](#) for more information.

Hateful conduct policy

Freedom of expression means little if voices are silenced because people are afraid to speak up. We do not tolerate behavior that harasses, intimidates, or uses fear to silence another person's voice. If you see something on Twitter that violates these rules, please report it to us.

How our policy works

As explained in the Twitter Rules,

- **Hateful conduct:** You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Examples of what we do not tolerate includes, but is not limited to behavior that harasses individuals or groups of people with:

- violent threats;
- wishes for the physical harm, death, or disease of individuals or groups;

- references to mass murder, violent events, or specific means of violence in which/with which such groups have been the primary targets or victims;
- behavior that incites fear about a protected group;
- repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone.

How our enforcement works

- **Context matters.**
 - Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. While we accept reports of violations from anyone, sometimes we also need to hear directly from the target to ensure that we have proper context.
 - The number of reports we receive does not impact whether or not something will be removed. However, it may help us prioritize the order in which it gets reviewed.
 - **We focus on behavior.**
 - We enforce policies when someone reports behavior that is abusive and targets an entire protected group and/or individuals who may be members.
 - This targeting can happen in any manner (for example, @mentions, tagging a photo, and more).
 - **We have a range of enforcement options.**
 - The consequences for violating our rules vary depending on the severity of the violation and the person's previous record of violations. For example, we may ask someone to remove the offending Tweet before they can Tweet again. For other cases, we may suspend an account.
-

Private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Some examples of private and confidential information include, but may not be limited to:

Intimate photos or videos that were taken or distributed without the subject's consent.
Images or videos that are considered and treated as private under applicable laws. Private contact or financial information, such as:

credit card information
social security or other national identity numbers
addresses or locations that are considered and treated as private non-public, personal phone numbers
non-public, personal email addresses

How our enforcement works

Keep in mind that although you may consider certain information to be private, not all postings of such information may be a violation of this policy.

Context matters.

We may consider the context and nature of the information posted, local privacy laws, and other case-specific facts when determining if this policy has been violated. For example, if information was previously posted or displayed elsewhere on the Internet prior to being put on Twitter, it may not be a violation of this policy.

We focus on behavior.

We enforce policies when someone reports content where the behavior may be abusive.

We have a range of enforcement options.

The consequences for violating our rules vary depending on the severity of the violation and the person's previous record of violations. For example, we may ask someone to remove the offending Tweet before they can Tweet again. For other cases, we may suspend an account.

For frequently asked questions about reporting private information posted on Twitter, [click here](#). To learn more about protecting your private information on Twitter and other websites, [click here](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter Impersonation Policy.

An account will not be removed if

- the user shares your name but has no other commonalities, or
- the profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same background or avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to [Twitter's Parody Policy](#) for more information about these types of accounts.

For Frequently Asked Questions about reporting impersonation on Twitter, [click here](#).

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000

people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation rules

Updated November 3, 2017

This page is primarily intended for **developers**.

For Twitter users: You are ultimately responsible for the actions taken with your account, or by applications associated with your account. Before authorizing a third-party application to access or use your account, make sure you've thoroughly investigated the application and understand what it will do. If automated activity on your account violates the [Twitter Rules](#) or these automation rules, Twitter may take action on your account, including [filtering your Tweets from search results](#) or suspending your account.

For more information on third-party applications, please see our article on [connecting and revoking third-party applications](#).

If you are a developer and have questions about these Automation Rules, please reach out via <https://twittercommunity.com/c/rules-and-policies>.

If your application has been suspended or restricted from performing write actions, you may submit an appeal at </forms/platform>.

I. Ground Rules Do!

Build solutions that automatically broadcast helpful information in Tweets.

Run creative campaigns that auto-reply to users who engage with your content.

Build solutions that automatically respond to users in Direct Messages.

Try new things that help people (and comply with our rules).

Make sure your application provides a good user experience and performs well — and confirm that remains the case over time.

Don't!

Violate these or other policies. Be extra mindful of our rules about abuse and user privacy.

Abuse the Twitter API or attempt to circumvent rate limits.

Use non-API-based forms of automation, such as scripting the Twitter website. The use of these techniques may result in the permanent suspension of your account.

Spam or bother users, or otherwise send them unsolicited messages.

A. The Twitter Rules and the Developer Agreement and Policy

As with all activity on Twitter, automated activity is subject to the [Twitter Rules](#) and, if you're a developer using the Twitter API, the [Developer Agreement and Policy](#).

You should carefully review these policies to ensure that your automated activity is compliant. Automated applications or activities that violate these policies, or that facilitate or induce users to violate them, may be subject to enforcement action, potentially including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these policies.

Although all aspects of the Twitter Rules and the Developer Agreement and Policy apply to automated activity, you should keep the following rules top of mind:

Spamming: You may not send automated Tweets or Direct Messages that are spam, or otherwise engage in spamming activity. Some examples of spammy behavior to avoid with automation include:

Trending topics: You may not automatically post about trending topics on Twitter, or use automation to attempt to influence or manipulate trending topics.

Multiple posts/accounts: You may not post duplicative or substantially similar Tweets on one account or over multiple accounts you operate.

Duplicate accounts: You may not create and/or automate multiple accounts for duplicative or substantially similar use cases.

However, automating multiple accounts for related but non-duplicative use cases is permitted. For example, you may automate separate accounts to Tweet when the Hubble Space Telescope passes over different cities, such as [San Francisco](#) or [Hong Kong](#).

Misleading links: You may not send automated Tweets or Direct Messages containing links that are misleading, including links that maliciously or deceptively redirect through landing pages or ad pages before displaying the final content.

Sensitive media: Automated Tweets and Direct Messages must comply with the [Twitter media policy](#), and you should mark your account as potentially sensitive if you intend to post graphic, pornographic, or potentially sensitive media.

Abusive behavior: You may not engage in any automated activity that encourages, promotes, or incites abuse, violence, hateful conduct, or harassment, on or off Twitter.

Private information: You may not post private or confidential information about a person without their prior express authorization.

B. Other Ground Rules for Automated Activity

In addition to the policies above, the following ground rules apply to all automated activity on Twitter:

Don't surprise or mislead users: Automated activity should honor users' expectations. Ask for the user's permission before taking an action if you aren't sure.

Mature content or profanity: Don't Direct Message, mention, or reply to users with potentially sensitive content (including profanity), unless they've clearly indicated an intent to receive it in advance.

Be thoughtful about the information you request or exchange on Twitter

Tweets: Don't ask users to send you personal or private information via a public Tweet. If you need additional personal or private information from a user to provide them with customer service (or other similar use cases), you should ask the user to share such information by Direct Message or another private channel. You might even consider adding a [Direct Message deep link](#) to your Tweet.

Direct Messages: You should only ask users for the minimum amount of information you need to provide them with service. If you need to request or exchange particularly sensitive information (such as credit card information), you should consider directing users to your website or other appropriate channel to do so.

II. Activity-Specific Rules

The activity-specific rules in this section apply to taking specific automated actions on Twitter. Please read these rules carefully, as they outline both permitted and prohibited use cases of automation.

Automated applications or activities that violate these rules, or that facilitate or induce users to violate them, may be subject to enforcement action, including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access the Twitter API based on violations of these rules. As a reminder, you should also carefully review the spam guidelines in the [Twitter Rules](#) to avoid having activities performed by you, your app, or other users through your app or service flagged as spam.

A. Automated Actions Through Another User's Account

Twitter users may authorize your app or service to [access their Twitter account through OAuth](#). A user authorizing your app or service to access their Twitter account through OAuth does not

by itself constitute sufficient consent to take automated actions through that user's account.

You may only take automated actions through another Twitter user's account if you:

clearly describe to the user the types of automated actions that will occur; receive express consent from the user to take those automated actions; and immediately honor a user's request to opt-out of further automated actions.

If you substantially change the purpose or functionality of your app or service, you must re-obtain express consent from the user to take automated action through their account before doing so.

These requirements apply to any automated action taken through another Twitter user's account, including posting Tweets, sending Direct Messages, deleting Tweets or Direct Messages, or following/unfollowing other accounts. For applications that offer users the ability to delete Tweets in a bulk or automated manner, you must also clearly state that Tweets are not recoverable once deleted.

B. Automated Tweets

1. Posting automated Tweets

Automated Tweets that cross-post outside information:

You may post automated Tweets based on sources of outside information — such as an RSS feed, weather data, etc. — as long as you are sufficiently authorized to publish such information.

Other automated Tweets (excluding mentions or replies):

Provided you comply with all other rules, you may post automated Tweets for entertainment, informational, or novelty purposes. As a reminder, accounts posting duplicative, spammy, or otherwise prohibited content may be subject to suspension.

2. Posting automated mentions and replies

The reply and mention functions are intended to make communication between Twitter users easier. Automating these actions to reach many users on an unsolicited basis is an abuse of the feature, and is not permitted. For example, sending automated replies to Tweets based on keyword searches alone is not permitted. Spammy or duplicative use of mentions and replies may result in enforcement action, such as the removal of your Tweets from Search or the suspension of your app or account.

However, you may send automated replies or mentions to Twitter users so long as:

in advance of sending the automated reply, the recipient or mentioned user(s) have requested or have clearly indicated an intent on Twitter to be contacted by you (i.e. opted in), for example by replying to a Tweet from your account, or by sending you a Direct Message;

you provide a clear and easy way for such users to opt-out of receiving automated replies and mentions, and promptly honor all such opt-out requests;

you only send one automated reply or mention per user interaction; and

the automated reply or mention is a reply to the user's original Tweet (if your campaign is based on users posting a reply to your Tweet).

Opt-in techniques and indications of user intent take many different forms, depending on the specifics of your use case and implementation. Some examples include:

A Tweet from your account that clearly indicates that a user taking a specific action on that Tweet (such as Retweeting it) will opt the user into receiving an automated response.

A mention of your account by the user in a manner suggesting the user clearly wishes or intends to receive a response. If you want to run an auto-reply campaign with a campaign- or use-case-specific hashtag, users should also mention you in their Tweets.

Note that a user following your account is not on its own a sufficient indication of user intent to receive an automated response.

: Advertisers, publishers, and brands using auto-response campaigns must request approval from Twitter

and may be subject to additional rules. Please reach out to your account or partner manager for assistance. If you do not have a partner manager but are interested in managed client services, you can learn more [here](#).

Additionally, we recommend that any accounts that will communicate with

users via automated mentions or replies:

Note

1. Appropriately filter responses based on potentially sensitive language in user handles,

display names, and Tweet text, as well as potentially sensitive media;

2. Check that the Tweet you are mentioning or replying to still exists (for example, using the statuses/lookup endpoint on the Twitter API).

C. Automated Direct Messages

Sending automated Direct Messages to users

You may not send unsolicited Direct Messages in a bulk or automated manner, and should be thoughtful about the frequency with which you contact users via Direct Message.

You may send automated Direct Messages to users so long as:

in advance of sending the Direct Message, the recipient(s) have requested or have clearly indicated an intent on Twitter to be contacted by you via Direct Message, for example by sending you a Direct Message; **and**

you provide a clear and easy way for such users to opt-out of receiving automated Direct Messages, and promptly honor all such opt-out requests.

The fact that a user is technically able to receive a Direct Message from you (e.g. because the user follows you, has enabled the ability to receive Direct Messages from any account, or because the user is in a pre-existing Direct Message conversation with you) does not necessarily mean they have requested or expect to receive automated Direct Messages from you.

Interacting with users via Direct Message

Per the Ground Rules, remember to be thoughtful about the amount and type of information you request or exchange with users via Direct Messages. If you will be asking a user to provide personal or private information via an automated Direct Message, you must clearly explain how you will use the information you're collecting. Consider including a link to your privacy policy in your Direct Message to the user, as well as in your Twitter profile bio.

Don't publicly share information received in a Direct Message conversation with a user without first obtaining explicit consent from the user. For example, if a user asks you via Direct Message about a purchase they made from you, you may not mention the user in a public Tweet that includes information about their purchase unless you have the user's explicit consent to do so.

After a user-initiated interaction ends, don't send additional follow-up Direct Messages or mention users in a Tweet unless you get permission from the user.

D. Automated actions you take on Tweets or accounts

Automated likes: You may not like Tweets in an automated manner.

Automated Retweets: Provided you comply with all other rules, you may Retweet or Quote Tweet in an automated manner for entertainment, informational, or novelty purposes. Automated Retweets often lead to negative user experiences, and bulk, aggressive, or spammy Retweeting is a violation of the [Twitter Rules](#).

Automated following/unfollowing: You may not follow or unfollow Twitter accounts in a bulk, aggressive, or indiscriminate manner. Aggressive following is a violation of the [Twitter Rules](#). Please also review our [following rules and best practices](#) to ensure you are in compliance. Note that applications that claim to get users more followers are also prohibited under the [Twitter Rules](#).

Automated adding to lists or collections: You may not add Twitter users to lists or add Tweets to collections in a bulk or indiscriminate manner. Adding a large number of unrelated users to lists is a violation of the [Twitter Rules](#).

Twitter search rules and restrictions

This article details the rules for using Twitter search, and explains reasons why content may not show up in Twitter search.

Why doesn't all content show up in Twitter search?

In order to keep your search results relevant, Twitter filters search results for quality Tweets and accounts. Material that jeopardizes search quality or creates a bad search experience for other people may be automatically removed from Twitter search. Read more about this [here](#).

General rules for using Twitter search

The [Twitter Rules](#) explain what behaviors are permitted on Twitter. In addition to these rules, we've included some tips below to keep your content relevant (and your Tweets in Twitter search).

Please do not:

- Repeatedly post duplicate or near-duplicate content (links or Tweets).
- Abuse trending topics or hashtags (topic words with a # sign).
- Send automated Tweets or replies.
- Use bots or applications to post similar messages based on keywords.
- Post similar messages over multiple accounts.
- Aggressively follow and unfollow people.

Twitter may automatically remove accounts engaging in these behaviors from search (or even suspend in some cases) in order to ensure the best experience for everyone.

Search results filtering

To learn how to filter your search results, including how to disable or enable safe search, read our [Using Twitter search](#) article.

Search results filtering in third-party applications

Accounts may be filtered from search for posting similar messages over several accounts. Your account may be filtered from search if you are logged into a third-party application that is updating many accounts with similar or duplicate content.

Rules for participating in contests on Twitter

Businesses sometimes host contests and offer prizes to Retweet an update, following a particular account, or posting Tweets with a specific hashtag topic or reply. If you're running a contest, check out this [article](#) for guidelines for contests on Twitter.

If you repeatedly post near-duplicate updates or duplicate links in order to enter contests, you may be filtered from search. Please keep in mind that you are responsible for the updates that you post to your Twitter account.

What to do if you are missing from search

Please see this [article](#) for more information related to this issue.

If you are being filtered from search, rest assured that your followers will still receive your updates. If you post a reply or mention to a specific person, it will be delivered to that person.

Having trouble?

You can learn more about searching by keyword, location, and trending topics in our [Using Twitter search](#) help article. Our [search troubleshooting](#) section includes solutions to common issues related to search.

The Twitter Rules (20171103)

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. These limitations are set forth in the Twitter Rules below.

The Twitter Rules (along with all incorporated policies), [Privacy Policy](#), and [Terms of Service](#) collectively make up the "Twitter User Agreement" that governs a user's access to and use of Twitter's services.

All individuals accessing or using Twitter's services must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in Twitter taking one or more of the following enforcement actions:

- requiring you to delete prohibited content before you can again create new posts and interact with other Twitter users;
- temporarily limiting your ability to create posts or interact with other Twitter users;
- asking you to verify account ownership with a phone number or email address; or
- permanently suspending your account(s).

If you attempt to evade a permanent suspension by creating new accounts, we will suspend your new accounts.

Please note that we may need to change these Rules from time to time and reserve the right to do so. The most current version is always available at: <https://twitter.com/rules>.

The policies set forth in these Twitter Rules govern organic content on our platform. To learn more about the rules which govern ads and promoted content, please review our [Ads policies](#).

Content Boundaries and Use of Twitter

Intellectual property

Trademark: We reserve the right to suspend accounts or take other appropriate action when someone's brand or trademark, including business name and/or logo, is used in a manner that may mislead or confuse others about your brand affiliation. Read more about our [trademark policy and how to report a violation](#).

Copyright: We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in our [Terms of Service](#). Read more about our [copyright policy](#).

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Learn more about our [media policy](#).

Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images. Additionally, Twitter may sometimes require you to remove excessively graphic violence out of respect for the

deceased and their families if we receive a request from their family or an authorized representative. Learn more about [how to make such a request](#), and [how to mark your media as sensitive](#).

Unlawful use

You may not use our service for any unlawful purposes or in furtherance of illegal activities. By using Twitter, you agree to comply with all applicable laws governing your online conduct and content.

Trends

At times, we may prevent certain content from trending. This includes content that violates the Twitter Rules, as well as content that may attempt to manipulate trends. Read more about [what we allow and do not allow to trend](#).

Misuse of Twitter badges

You may not use badges, including but not limited to the “promoted” or “verified” Twitter badges, unless provided by Twitter. Accounts using unauthorized badges as part of their profile photos, header photos, display names, or in any way that falsely implies affiliation with Twitter or authorization from Twitter to display these badges, may be suspended.

Misuse of usernames

Selling usernames: You may not buy or sell Twitter usernames.

Username squatting: You may not engage in username squatting. Some of the factors we take into consideration when determining whether conduct is username squatting include:

- the number of accounts created;
- the creation of accounts for the purpose of preventing others from using those account names;
- the creation of accounts for the purpose of selling those accounts; and
- the use of third-party content feeds to update and maintain accounts under the names of those third parties.

Please note that Twitter may also remove accounts that are inactive for more than six months. Learn more about [username squatting](#).

Abusive Behavior

We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we prohibit behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Context matters when evaluating for abusive behavior and determining appropriate enforcement actions. Factors we may take into consideration include, but are not limited to whether:

- the behavior is targeted at an individual or group of people;

- the report has been filed by the target of the abuse or a bystander;
- the behavior is newsworthy and in the legitimate public interest.

Violence and physical harm

Violence: You may not make specific threats of violence or wish for the serious physical harm, death, or disease of an individual or group of people. This includes, but is not limited to, threatening or promoting terrorism.

Suicide or self-harm: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Child sexual exploitation: You may not promote child sexual exploitation. Learn more about our zero-tolerance [child sexual exploitation policy](#).

Abuse and hateful conduct

Abuse: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Unwanted sexual advances: You may not direct abuse at someone by sending unwanted sexual content, objectifying them in a sexually explicit manner, or otherwise engaging in sexual misconduct.

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. Read more about our [hateful conduct policy](#).

Private information and intimate media

Private information: You may not publish or post other people's private information without their express authorization and permission. Definitions of private information may vary depending on local laws. Read more about our [private information policy](#).

Intimate media: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. Read more about intimate media on Twitter.

Threats to expose / hack: You may not threaten to expose someone's private information or intimate media. You also may not threaten to hack or break into someone's digital information.

Impersonation

You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. While you may maintain parody, fan, commentary, or newsfeed accounts, you may not do so if the intent of the account is to engage in spamming or abusive behavior. Read more about our [impersonation policy](#).

Spam and Security

We strive to protect people on Twitter from technical abuse and spam.

To promote a stable and secure environment on Twitter, you may not do, or attempt to do, any of the following while accessing or using Twitter:

- Access, tamper with, or use non-public areas of Twitter, Twitter's computer systems, or the technical delivery systems of Twitter's providers (except as expressly permitted by the Twitter Bug Bounty program).
- Probe, scan, or test the vulnerability of any system or network, or breach or circumvent any security or authentication measures (except as expressly permitted by the Twitter Bug Bounty program).
- Access or search, or attempt to access or search, Twitter by any means (automated or otherwise) other than through our currently available, published interfaces that are provided by Twitter (and only pursuant to the applicable terms and conditions), unless you have been specifically allowed to do so in a separate agreement with Twitter. Note that crawling Twitter is permissible if done in accordance with the provisions of the robots.txt file; however, scraping Twitter without our prior consent is expressly prohibited.
- Forge any TCP/IP packet header or any part of the header information in any email or posting, or in any way use Twitter to send altered, deceptive, or false source-identifying information.
- Interfere with or disrupt the access of any user, host or network, including, without limitation, sending a virus, overloading, flooding, spamming, mail-bombing Twitter's services, or by scripting the creation of content in such a manner as to interfere with or create an undue burden on Twitter.

Any accounts engaging in the following activities may be temporarily locked or subject to permanent suspension:

- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use Twitter's services for the purpose of spamming anyone. Spam is generally defined on Twitter as bulk or aggressive activity that attempts to manipulate or disrupt Twitter or the experience of users on Twitter to drive traffic or attention to unrelated accounts, products, services, or initiatives. Some of the factors that we take into account when determining what conduct is considered to be spamming include:
 - if you have followed and/or unfollowed a large number of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if your Tweets or Direct Messages consist mainly of links shared without commentary;
 - if a large number of people have blocked you in response to high volumes of untargeted, unsolicited, or duplicative content or engagements from your account;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicative or substantially similar content, replies, or mentions over multiple accounts or multiple duplicate updates on one account, or create duplicate or substantially similar accounts;

- if you post multiple updates to a trending or popular topic with an intent to subvert or manipulate the topic to drive traffic or attention to unrelated accounts, products, services, or initiatives;
- if you send large numbers of unsolicited replies or mentions;
- if you add users to lists in a bulk or aggressive manner;
- if you are randomly or aggressively engaging with Tweets (e.g., likes, Retweets, etc.) or users (e.g., following, adding to lists or Moments, etc.) to drive traffic or attention to unrelated accounts, products, services, or initiatives;
- if you repeatedly post other people's account information as your own (e.g., bio, Tweets, profile URL, etc.);
- if you post misleading, deceptive, or malicious links (e.g., affiliate links, links to malware/clickjacking pages, etc.);
- if you create fake accounts, account interactions, or impressions;
- if you sell, purchase, or attempt to artificially inflate account interactions (such as followers, Retweets, likes, etc.); and
- if you use or promote third-party services or apps that claim to get you more followers, Retweets, or likes (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers or engagements to your account or Tweets).

Please see our support articles on [following rules and best practices](#) and [automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts may be permanently suspended.

Content Visibility

Accounts under investigation or which have been detected as sharing content in violation of these Rules may have their account or Tweet visibility limited in various parts of Twitter, including search. To learn more about situations in which content may be restricted on Twitter, please see our support article on [search rules and restrictions](#).

Trademark policy (20160630)

Updated June 30, 2016

What is a Trademark Policy Violation on Twitter?

Using another's trademark in a manner that may mislead or confuse others about your brand affiliation may be a violation of our trademark policy.

How Does Twitter Respond to Trademark Reports?

We review each report and do the following:

- If we determine that there is a violation of our trademark policy, we may suspend the account. In some instances, we may give the account holder an opportunity to comply with our policies before suspending the account. For more information, see our [Parody, commentary, and fan account policy](#).
- We may also provide the account holder with your name and information included in the copy of the report.
- If you would like to report a trademark violation in a [Twitter Ad](#), please visit our [Trademark Policy for Promoted Ads](#) to learn more.

Please note that Twitter only investigates requests that are submitted by the trademark holder or by an authorized representative of the trademark holder.

What is not a Trademark Policy Violation?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples include:

- Using a trademark in a way that is outside the scope of the trademark registration (e.g. territory, or goods and services identified in the registration).
- Nominative and other fair uses of trademarks are protected uses under our trademark policy, so long as the account is clearly distinguished from the trademark owner. This includes use by resellers in certain regions and accounts engaging in parody, commentary, or news. For more information, see our [Parody, commentary, and fan account policy](#).

How do I Report a Trademark Policy Violation?

Submit trademark reports through our [Trademark Issue support form](#). Using the appropriate form ensures we have all the information necessary to process your report quickly. Please provide all information requested in the form. If you submit a report with incomplete information, we will follow up with you to request that information. Please note that this will result in a delay in processing your report.

If you have trouble using the form, try updating your browser or switching to another browser.

Copyright policy (20171222)

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from **@Username** has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notification.

When Should I File a Counter-notice?

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);

2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Twitter media policy (20171102*)

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to easily share your content with others. Below are our guidelines for making sure your content is labeled appropriately.

Media on Twitter

You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.

In addition, you may not feature graphic content (such as media containing pornography or excessive violence) in live video, or in your profile image or header image.

If you see media you think may be in violation of our policy, see [this article](#) for information about how to report it.

Your media settings

All Twitter accounts have media settings to help Twitter understand the type of media you are posting and to help us assist others to find the type of content that they're interested in.

By appropriately marking your media settings, Twitter can identify potentially sensitive content that others users may not wish to see, such as violence or nudity. If you intend to regularly post such content, we request that you adjust your media setting:

1. Log in to your account on the web at twitter.com.
2. Click your **profile icon** in the top right navigation bar.
3. Select **Settings and privacy** from the drop-down menu.
4. Go to your [Privacy and safety](#) settings.
5. Look for the **Tweet media** section and check the box next to **Mark media you Tweet as containing material that may be sensitive**. **Note:** Enabling this setting will result in others seeing a message when they view your profile, alerting them of potentially sensitive media and asking them to confirm they still want to view it.

You can control your preferences for how and when you may see sensitive content in your [media settings](#).

What if I don't change my media settings and post sensitive media?

If one of your Tweets containing media is reported as sensitive, it will be sent to Twitter for review. If we find that the media wasn't marked as being potentially sensitive at the time of the upload:

- We label the media as being potentially sensitive, or in the case of live video, remove the content entirely.
- We may also change your account setting to **Mark media you Tweet as containing material that may be sensitive** so that future posts are marked accordingly.

You will receive an email letting you know that we've labeled your media and if we changed your media setting. You will be able to change your media settings on your account settings page, though this image and any other images that have been labeled will remain marked as potentially sensitive.

If you repeatedly upload media that are mislabeled, the Twitter team may:

- Label your media as having potentially sensitive content.
- Permanently adjust your account setting to **Mark media you Tweet as containing material that may be sensitive** so that your future posts have a warning message users must click through before viewing your media.

Twitter may also use automated techniques to detect and label potentially sensitive media.

Additionally, people who visit your profile may see a message letting them know the account may include potentially sensitive content and asking them to confirm they still want to view it. People who have opted in to see possibly sensitive content will still see your account without the message.

What happens if my profile image is in violation of the Twitter Rules?

If your profile or header image contains graphic content in violation of the Twitter Rules, we may suspend your account and remove the offending content. Repeat violations will result in a permanent suspension.

Additionally, your account may be locked if media in your profile is in violation of Twitter's media policy. In order to unlock your account, you will need to follow our instructions and delete the media in violation. If your account has been locked, see this [article](#) for more information.

Can I appeal a decision made by the Twitter team?

If you are suspended for having a pornographic profile, header, or background image, you can appeal by logging in to your account and submitting an appeal [here](#).

There is currently no way to appeal a decision by the Twitter team that permanently changes your account setting to **Mark media you Tweet as containing material that may be sensitive** in response to repeated mislabeling of sensitive content.

Contacting Twitter about media concerning a deceased family member (20171004*)

In order to respect the wishes of loved ones, Twitter will remove imagery of deceased individuals in certain circumstances. Immediate family members and other authorized individuals may request the removal of images or video of deceased individuals, from when critical injury occurs to the moments before or after death, by sending a request to Twitter Inc. via our [privacy form](#). When reviewing such media removal requests, Twitter considers public interest factors such as the newsworthiness of the content and may not be able to honor every request.

Note: For information about contacting Twitter to deactivate the account of a deceased or incapacitated user, please read [this article](#).

How to mark media as sensitive links to Media policy

FAQs about trends on Twitter (20171027)

Where can I find trends?

Trends are available on the Twitter for iOS app, the Twitter for Android app, and twitter.com.

On Twitter's mobile apps, you can find trends listed under the trends section



of the **Explore** tab When signed in to twitter.com on a desktop or laptop computer, trends are listed in many places, including the Home timeline, Notifications, search results, and profile pages.

How are trends determined?

Trends are determined by an algorithm and, by default, are tailored for you based on who you follow, your interests, and your location. This algorithm identifies topics that are popular now, rather than topics that have been popular for a while or on a daily basis, to help you discover the hottest emerging topics of discussion on Twitter.

Note: The number of Tweets that are related to the trends is just one of the factors the algorithm looks at when ranking and determining trends. Algorithmically, trends and hashtags are grouped together if they are related to the same topic. For instance, #MondayMotivation and #MotivationMonday may both be represented by #MondayMotivation.

You can choose to see trends that are not tailored for you by selecting a [specific trends location](#) on twitter.com, iOS, or Android. Location trends identify popular topics among people in a specific geographic location.

What does the # sign mean?

You may notice that some trends have # sign before the word or phrase. This is called a [hashtag](#) and is included specifically in Tweets to mark them as relating to a topic, so that people can follow the conversation in search.

Will I see world and local events in my trends?

Yes, when signed into your account on the twitter.com, iOS, or Android, trends will be tailored for you based on your location and who you follow. There will be many world and local news events and conversations that will appear in your trends regardless of your personalization.

Can I see trends for a specific location?

Yes, you can view trends for a specific location on twitter.com, iOS, or Android.

View trends via twitter.com:

To view trends for a specific location, click **Change** in your trends box. You can choose to **Keep tailored trends** based on your location and who you follow, or you can select **Change** to choose a nearby location or search locations.

Adjust your trends settings via Twitter for iOS:

1. In the top menu, tap your **profile** icon.
2. Tap **Settings and privacy**, then tap **Content preferences**.
3. Under **Explore**, tap **Trends**.

4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.
5. When the feature is disabled, you can change your location by tapping **Change location**.
 1. In the search box, type your desired trends location or choose from the list.

Adjust your trends settings via Twitter for Android:

1. In the top menu, you will either see a **navigation menu icon**  or your **profile icon**. Tap whichever icon you have, and select **Settings and privacy**.
2. Tap **Content preferences**.
3. Under **Explore**, tap **Trends**.
4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.
5. When the feature is disabled, you can change your location by tapping **Change location**.
 1. In the search box, type your desired trends location or choose from the list.

Note: If you don't find your preferred country or city, it means we're not receiving enough Tweets from that geographical area to create a list. You can look up local Tweets on any topic by using [advanced Twitter Search operators](#).

What happens when I click or tap on a trend?

Clicking or tapping any of the trends takes you to the Twitter search results for that trend. You'll see all Tweets including that phrase or hashtag. To see what people are saying about a previous trend, perform a [search for that keyword](#).

How can I participate in a trend?

Simply post a Tweet including the exact word or phrase as it appears in the trends list (with the hashtag, if you see one). Due to the large number of people Tweeting about these specific trends, you may not always be able to find your particular Tweet in search, but your followers will always see your Tweets.

Note: Twitter also filters searches for quality. To understand what types of behavior could cause your Tweets to be filtered out, read about our [search rules and restrictions](#).

Are there rules for trends?

Yes. We want trends to promote healthy discussions on Twitter. This means that at times, we may prevent certain content from trending. These include trends that:

- Contain profanity or adult/graphic references.
- Incite hate on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease.
- Violate the [Twitter Rules](#).

In some cases, we may also consider the newsworthiness of the content, or if it is in the public interest when evaluating potential violations. In these cases, the content might continue to trend on our platform.

Note that even if we prevent the hashtag or the content from appearing on the list of trending topics, you may still be able to access conversations around that content on Twitter.

Trends descriptions

Below some trends, you may see the approximate number of Tweets associated with the trend, or personalized context such as who in your network is Tweeting about it.

Additionally, when a Moment is found to be relevant to a trend, we may use the title of the Moment as the trend's description.

Having trouble?

Check out our [troubleshooting section](#) for solutions to common problems.

Username squatting policy (20130817*)

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Child sexual exploitation policy (20171025*)

Twitter removes content promoting child sexual exploitation

We do not tolerate child sexual exploitation on Twitter. When we are made aware of links to images of or content promoting child sexual exploitation they will be removed from the site without further notice and reported to The National Center for Missing & Exploited Children ("NCMEC"); we permanently suspend accounts promoting or containing updates with links to child sexual exploitation.

Reporting content promoting child sexual exploitation

If you think you've found a Twitter account distributing or promoting child sexual exploitation while using Twitter, please click [here](#) to report it. Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. To find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with the National Center for Missing and Exploited Children [here](#), or your local law enforcement. If you believe that the content was hosted on a website outside of the United States, you can report it [here](#) on the International Association of Internet Hotlines website.

NOTE: Please do NOT tweet, retweet or repost child sexual exploitation for any reason. Click [here](#) to report it to us immediately and we will take steps to remove it.

Law Enforcement Agencies

If you are a member of a law enforcement agency seeking information regarding Twitter accounts promoting or distributing child sexual exploitation, please see our [Guidelines for Law Enforcement](#) for information on our process; we also suggest you contact the National Center for Missing and Exploited Children (NCMEC) [here](#), regardless of jurisdiction, to determine whether or not information on those accounts is already on file.

For general law enforcement inquiries, please see our [Guidelines for Law Enforcement](#) for more information.

Hateful conduct policy (20171101*)

Freedom of expression means little if voices are silenced because people are afraid to speak up. We do not tolerate behavior that harasses, intimidates, or uses fear to silence another person's voice. If you see something on Twitter that violates these rules, please report it to us.

How our policy works

As explained in the Twitter Rules,

- **Hateful conduct:** You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Examples of what we do not tolerate includes, but is not limited to behavior that harasses individuals or groups of people with:

- violent threats;
- wishes for the physical harm, death, or disease of individuals or groups;

- references to mass murder, violent events, or specific means of violence in which/with which such groups have been the primary targets or victims;
- behavior that incites fear about a protected group;
- repeated and/or or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone.

How our enforcement works

- **Context matters.**
 - Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. While we accept reports of violations from anyone, sometimes we also need to hear directly from the target to ensure that we have proper context.
 - The number of reports we receive does not impact whether or not something will be removed. However, it may help us prioritize the order in which it gets reviewed.
 - **We focus on behavior.**
 - We enforce policies when someone reports behavior that is abusive and targets an entire protected group and/or individuals who may be members.
 - This targeting can happen in any manner (for example, @mentions, tagging a photo, and more).
 - **We have a range of enforcement options.**
 - The consequences for violating our rules vary depending on the severity of the violation and the person's previous record of violations. For example, we may ask someone to remove the offending Tweet before they can Tweet again. For other cases, we may suspend an account.
-

Private information posted on Twitter (20170714*)

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Some examples of private and confidential information include, but may not be limited to:

Intimate photos or videos that were taken or distributed without the subject's consent.
Images or videos that are considered and treated as private under applicable laws. Private contact or financial information, such as:

credit card information
social security or other national identity numbers
addresses or locations that are considered and treated as private non-public, personal phone numbers
non-public, personal email addresses

How our enforcement works

Keep in mind that although you may consider certain information to be private, not all postings of such information may be a violation of this policy.

Context matters.

We may consider the context and nature of the information posted, local privacy laws, and other case-specific facts when determining if this policy has been violated. For example, if information was previously posted or displayed elsewhere on the Internet prior to being put on Twitter, it may not be a violation of this policy.

We focus on behavior.

We enforce policies when someone reports content where the behavior may be abusive.

We have a range of enforcement options.

The consequences for violating our rules vary depending on the severity of the violation and the person's previous record of violations. For example, we may ask someone to remove the offending Tweet before they can Tweet again. For other cases, we may suspend an account.

For frequently asked questions about reporting private information posted on Twitter, [click here](#). To learn more about protecting your private information on Twitter and other websites, [click here](#).

Impersonation policy (201305*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter Impersonation Policy.

An account will not be removed if

- the user shares your name but has no other commonalities, or
- the profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same background or avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to [Twitter's Parody Policy](#) for more information about these types of accounts.

For Frequently Asked Questions about reporting impersonation on Twitter, [click here](#).

Following rules and best practices (20160117*)

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation rules (20171103)

Updated November 3, 2017

This page is primarily intended for **developers**.

For Twitter users: You are ultimately responsible for the actions taken with your account, or by applications associated with your account. Before authorizing a third-party application to access or use your account, make sure you've thoroughly investigated the application and understand what it will do. If automated activity on your account violates the [Twitter Rules](#) or these automation rules, Twitter may take action on your

account, including [filtering your Tweets from search results](#) or suspending your account.

For more information on third-party applications, please see our article on [connecting and revoking third-party applications](#).

If you are a developer and have questions about these Automation Rules, please reach out via <https://twittercommunity.com/c/rules-and-policies>.

If your application has been suspended or restricted from performing write actions, you may submit an appeal at </forms/platform>.

I. Ground Rules Do!

Build solutions that automatically broadcast helpful information in Tweets.

Run creative campaigns that auto-reply to users who engage with your content.

Build solutions that automatically respond to users in Direct Messages.

Try new things that help people (and comply with our rules).

Make sure your application provides a good user experience and performs well — and confirm that remains the case over time.

Don't!

Violate these or other policies. Be extra mindful of our rules about abuse and user privacy.

Abuse the Twitter API or attempt to circumvent rate limits.

Use non-API-based forms of automation, such as scripting the Twitter website. The use of these techniques may result in the permanent suspension of your account.

Spam or bother users, or otherwise send them unsolicited messages.

A. The Twitter Rules and the Developer Agreement and Policy

As with all activity on Twitter, automated activity is subject to the [Twitter Rules](#) and, if you're a developer using the Twitter API, the [Developer Agreement and Policy](#).

You should carefully review these policies to ensure that your automated activity is compliant. Automated applications or activities that violate these policies, or that facilitate or induce users to violate them, may be subject to enforcement action,

potentially including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these policies.

Although all aspects of the Twitter Rules and the Developer Agreement and Policy apply to automated activity, you should keep the following rules top of mind:

Spamming: You may not send automated Tweets or Direct Messages that are spam, or otherwise engage in spamming activity. Some examples of spammy behavior to avoid with automation include:

Trending topics: You may not automatically post about trending topics on Twitter, or use automation to attempt to influence or manipulate trending topics.

Multiple posts/accounts: You may not post duplicative or substantially similar Tweets on one account or over multiple accounts you operate.

Duplicate accounts: You may not create and/or automate multiple accounts for duplicative or substantially similar use cases.

However, automating multiple accounts for related but non-duplicative use cases is permitted. For example, you may automate separate accounts to Tweet when the Hubble Space Telescope passes over different cities, such as [San Francisco](#) or [Hong Kong](#).

Misleading links: You may not send automated Tweets or Direct Messages containing links that are misleading, including links that maliciously or deceptively redirect through landing pages or ad pages before displaying the final content.

Sensitive media: Automated Tweets and Direct Messages must comply with the [Twitter media policy](#), and you should mark your account as potentially sensitive if you intend to post graphic, pornographic, or potentially sensitive media.

Abusive behavior: You may not engage in any automated activity that encourages, promotes, or incites abuse, violence, hateful conduct, or harassment, on or off Twitter.

Private information: You may not post private or confidential information about a person without their prior express authorization.

B. Other Ground Rules for Automated Activity

In addition to the policies above, the following ground rules apply to all automated activity on Twitter:

Don't surprise or mislead users: Automated activity should honor users' expectations. Ask for the user's permission before taking an action if you aren't sure.

Mature content or profanity: Don't Direct Message, mention, or reply to users with potentially sensitive content (including profanity), unless they've clearly indicated an intent to receive it in advance.

Be thoughtful about the information you request or exchange on Twitter

Tweets: Don't ask users to send you personal or private information via a public Tweet. If you need additional personal or private information from a user to provide them with customer service (or other similar use cases), you should ask the user to share such information by Direct Message or another private channel. You might even consider adding a [Direct Message deep link](#) to your Tweet.

Direct Messages: You should only ask users for the minimum amount of information you need to provide them with service. If you need to request or exchange particularly sensitive information (such as credit card information), you should consider directing users to your website or other appropriate channel to do so.

II. Activity-Specific Rules

The activity-specific rules in this section apply to taking specific automated actions on Twitter. Please read these rules carefully, as they outline both permitted and prohibited use cases of automation.

Automated applications or activities that violate these rules, or that facilitate or induce users to violate them, may be subject to enforcement action, including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these rules. As a reminder, you should also carefully review the spam guidelines in the [Twitter Rules](#) to avoid having activities performed by you, your app, or other users through your app or service flagged as spam.

A. Automated Actions Through Another User's Account

Twitter users may authorize your app or service to [access their Twitter account through OAuth](#). A user authorizing your app or service to access their Twitter account through OAuth does not by itself constitute sufficient consent to take automated actions through that user's account.

You may only take automated actions through another Twitter user's account if you:

clearly describe to the user the types of automated actions that will occur; receive express consent from the user to take those automated actions; and immediately honor a user's request to opt-out of further automated actions.

If you substantially change the purpose or functionality of your app or service, you must re-obtain express consent from the user to take automated action through their account before doing so.

These requirements apply to any automated action taken through another Twitter user's account, including posting Tweets, sending Direct Messages, deleting Tweets or Direct Messages, or following/unfollowing other accounts. For applications that offer users the ability to delete Tweets in a bulk or automated manner, you must also clearly state that Tweets are not recoverable once deleted.

B. Automated Tweets

1. Posting automated Tweets

Automated Tweets that cross-post outside information:

You may post automated Tweets based on sources of outside information — such as an RSS feed, weather data, etc. — as long as you are sufficiently authorized to publish such information.

Other automated Tweets (excluding mentions or replies):

Provided you comply with all other rules, you may post automated Tweets for entertainment, informational, or novelty

purposes. As a reminder, accounts posting duplicative, spammy, or otherwise prohibited content may be subject to suspension.

2. Posting automated mentions and replies

The reply and mention functions are intended to make communication between Twitter users easier. Automating these actions to reach many users on an unsolicited basis is an abuse of the feature, and is not permitted. For example, sending automated replies to Tweets based on keyword searches alone is not permitted. Spamy or duplicative use of mentions and replies may result in enforcement action, such as the removal of your Tweets from Search or the suspension of your app or account.

However, you may send automated replies or mentions to Twitter users so long as:

in advance of sending the automated reply, the recipient or mentioned user(s) have requested or have clearly indicated an intent on Twitter to be contacted by you (i.e. opted in), for example by replying to a Tweet from your account, or by sending you a Direct Message;

you provide a clear and easy way for such users to opt-out of receiving automated replies and mentions, and promptly honor all such opt-out requests;

you only send one automated reply or mention per user interaction; and

the automated reply or mention is a reply to the user's original Tweet (if your campaign is based on users posting a reply to your Tweet).

Opt-in techniques and indications of user intent take many different forms, depending on the specifics of your use case and implementation. Some examples include:

A Tweet from your account that clearly indicates that a user taking a specific action on that Tweet (such as Retweeting it) will opt the user into receiving an automated response.

A mention of your account by the user in a manner suggesting the user clearly wishes or intends to receive a response. If you want to run an auto-reply campaign with a campaign- or use-case-specific hashtag, users should also mention you in their Tweets.

Note that a user following your account is not on its own a sufficient indication of user intent to receive an automated response.

: Advertisers, publishers, and brands using auto-response campaigns must request approval from Twitter and may be subject to additional rules. Please reach out to your account or partner manager for assistance. If you do not have a partner manager but are interested in managed client services, you can learn more [here](#).

Additionally, we recommend that any accounts that will communicate with

users via automated mentions or replies:

Note

1. Appropriately filter responses based on potentially sensitive language in user handles,

display names, and Tweet text, as well as potentially sensitive media;

2. Check that the Tweet you are mentioning or replying to still exists (for example, using the statuses/lookup endpoint on the Twitter API).

C. Automated Direct Messages

Sending automated Direct Messages to users

You may not send unsolicited Direct Messages in a bulk or automated manner, and should be thoughtful about the frequency with which you contact users via Direct Message.

You may send automated Direct Messages to users so long as:

in advance of sending the Direct Message, the recipient(s) have requested or have clearly indicated an intent on Twitter to be contacted by you via Direct Message, for example by sending you a Direct Message; **and**

you provide a clear and easy way for such users to opt-out of receiving automated Direct Messages, and promptly honor all such opt-out requests.

The fact that a user is technically able to receive a Direct Message from you (e.g. because the user follows you, has enabled the ability to receive Direct Messages from any

account, or because the user is in a pre-existing Direct Message conversation with you) does not necessarily mean they have requested or expect to receive automated Direct Messages from you.

Interacting with users via Direct Message

Per the Ground Rules, remember to be thoughtful about the amount and type of information you request or exchange with users via Direct Messages. If you will be asking a user to provide personal or private information via an automated Direct Message, you must clearly explain how you will use the information you're collecting. Consider including a link to your privacy policy in your Direct Message to the user, as well as in your Twitter profile bio.

Don't publicly share information received in a Direct Message conversation with a user without first obtaining explicit consent from the user. For example, if a user asks you via Direct Message about a purchase they made from you, you may not mention the user in a public Tweet that includes information about their purchase unless you have the user's explicit consent to do so.

After a user-initiated interaction ends, don't send additional follow-up Direct Messages or mention users in a Tweet unless you get permission from the user.

D. Automated actions you take on Tweets or accounts

Automated likes: You may not like Tweets in an automated manner.

Automated Retweets: Provided you comply with all other rules, you may Retweet or Quote Tweet in an automated manner for entertainment, informational, or novelty purposes. Automated Retweets often lead to negative user experiences, and bulk, aggressive, or spammy Retweeting is a violation of the [Twitter Rules](#).

Automated following/unfollowing: You may not follow or unfollow Twitter accounts in a bulk, aggressive, or indiscriminate manner. Aggressive following is a violation of the [Twitter Rules](#). Please also review our [following rules and best practices](#) to ensure you are in compliance. Note that applications that claim to get users more followers are also prohibited under the [Twitter Rules](#).

Automated adding to lists or collections: You may not add Twitter users to lists or add Tweets to collections in a bulk or indiscriminate manner. Adding a large number of unrelated users to lists is a violation of the [Twitter Rules](#).

[Twitter search rules and restrictions](#) (20171103)

This article details the rules for using Twitter search, and explains reasons why content may not show up in Twitter search.

Why doesn't all content show up in Twitter search?

In order to keep your search results relevant, Twitter filters search results for quality Tweets and accounts. Material that jeopardizes search quality or creates a bad search experience for other people may be automatically removed from Twitter search. Read more about this [here](#).

General rules for using Twitter search

The [Twitter Rules](#) explain what behaviors are permitted on Twitter. In addition to these rules, we've included some tips below to keep your content relevant (and your Tweets in Twitter search).

Please do not:

- Repeatedly post duplicate or near-duplicate content (links or Tweets).
- Abuse trending topics or hashtags (topic words with a # sign).
- Send automated Tweets or replies.
- Use bots or applications to post similar messages based on keywords.
- Post similar messages over multiple accounts.
- Aggressively follow and unfollow people.

Twitter may automatically remove accounts engaging in these behaviors from search (or even suspend in some cases) in order to ensure the best experience for everyone.

Search results filtering

To learn how to filter your search results, including how to disable or enable safe search, read our [Using Twitter search](#) article.

Search results filtering in third-party applications

Accounts may be filtered from search for posting similar messages over several accounts. Your account may be filtered from search if you are logged into a third-party application that is updating many accounts with similar or duplicate content.

Rules for participating in contests on Twitter

Businesses sometimes host contests and offer prizes to Retweet an update, following a particular account, or posting Tweets with a specific hashtag topic or reply. If you're running a contest, check out this [article](#) for guidelines for contests on Twitter.

If you repeatedly post near-duplicate updates or duplicate links in order to enter contests, you may be filtered from search. Please keep in mind that you are responsible for the updates that you post to your Twitter account.

What to do if you are missing from search

Please see this [article](#) for more information related to this issue.

If you are being filtered from search, rest assured that your followers will still receive your updates. If you post a reply or mention to a specific person, it will be delivered to that person.

Having trouble?

You can learn more about searching by keyword, location, and trending topics in our [Using Twitter search](#) help article. Our [search troubleshooting](#) section includes solutions to common issues related to search.

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. These limitations are set forth in the Twitter Rules below.

The Twitter Rules (along with all incorporated policies), [Privacy Policy](#), and [Terms of Service](#) collectively make up the "Twitter User Agreement" that governs a user's access to and use of Twitter's services.

All individuals accessing or using Twitter's services must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in Twitter taking one or more of the following enforcement actions:

- requiring you to delete prohibited content before you can again create new posts and interact with other Twitter users;
- temporarily limiting your ability to create posts or interact with other Twitter users;
- asking you to verify account ownership with a phone number or email address; or
- permanently suspending your account(s).

If you attempt to evade a permanent suspension by creating new accounts, we will suspend your new accounts.

Please note that we may need to change these Rules from time to time and reserve the right to do so. The most current version is always available at: <https://twitter.com/rules>.

The policies set forth in these Twitter Rules govern organic content on our platform. To learn more about the rules which govern ads and promoted content, please review our [Ads policies](#).

Content Boundaries and Use of Twitter

Intellectual property

Trademark: We reserve the right to suspend accounts or take other appropriate action when someone's brand or trademark, including business name and/or logo, is used in a manner that may mislead or confuse others about your brand affiliation. Read more about our [trademark policy and how to report a violation](#).

Copyright: We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in our [Terms of Service](#). Read more about our [copyright policy](#).

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Learn more about our [media policy](#).

Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images. Additionally, Twitter may sometimes require you to remove excessively graphic violence out of respect for the

deceased and their families if we receive a request from their family or an authorized representative. Learn more about [how to make such a request](#), and [how to mark your media as sensitive](#).

Unlawful use

You may not use our service for any unlawful purposes or in furtherance of illegal activities. By using Twitter, you agree to comply with all applicable laws governing your online conduct and content.

Trends

At times, we may prevent certain content from trending. This includes content that violates the Twitter Rules, as well as content that may attempt to manipulate trends. Read more about [what we allow and do not allow to trend](#).

Misuse of Twitter badges

You may not use badges, including but not limited to the “promoted” or “verified” Twitter badges, unless provided by Twitter. Accounts using unauthorized badges as part of their profile photos, header photos, display names, or in any way that falsely implies affiliation with Twitter or authorization from Twitter to display these badges, may be suspended.

Misuse of usernames

Selling usernames: You may not buy or sell Twitter usernames.

Username squatting: You may not engage in username squatting. Some of the factors we take into consideration when determining whether conduct is username squatting include:

- the number of accounts created;
- the creation of accounts for the purpose of preventing others from using those account names;
- the creation of accounts for the purpose of selling those accounts; and
- the use of third-party content feeds to update and maintain accounts under the names of those third parties.

Please note that Twitter may also remove accounts that are inactive for more than six months. Learn more about [username squatting](#).

Abusive Behavior

We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we prohibit behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Context matters when evaluating for abusive behavior and determining appropriate enforcement actions. Factors we may take into consideration include, but are not limited to whether:

- the behavior is targeted at an individual or group of people;

- the report has been filed by the target of the abuse or a bystander;
- the behavior is newsworthy and in the legitimate public interest.

Violence and physical harm

Violence: You may not make specific threats of violence or wish for the serious physical harm, death, or disease of an individual or group of people. This includes, but is not limited to, threatening or promoting terrorism.

Suicide or self-harm: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Child sexual exploitation: You may not promote child sexual exploitation. Learn more about our zero-tolerance [child sexual exploitation policy](#).

Abuse and hateful conduct

Abuse: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Unwanted sexual advances: You may not direct abuse at someone by sending unwanted sexual content, objectifying them in a sexually explicit manner, or otherwise engaging in sexual misconduct.

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. Read more about our [hateful conduct policy](#).

Private information and intimate media

Private information: You may not publish or post other people's private information without their express authorization and permission. Definitions of private information may vary depending on local laws. Read more about our [private information policy](#).

Intimate media: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. Read more about intimate media on Twitter.

Threats to expose / hack: You may not threaten to expose someone's private information or intimate media. You also may not threaten to hack or break into someone's digital information.

Impersonation

You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. While you may maintain parody, fan, commentary, or newsfeed accounts, you may not do so if the intent of the account is to engage in spamming or abusive behavior. Read more about our [impersonation policy](#).

Spam and Security

We strive to protect people on Twitter from technical abuse and spam.

To promote a stable and secure environment on Twitter, you may not do, or attempt to do, any of the following while accessing or using Twitter:

- Access, tamper with, or use non-public areas of Twitter, Twitter's computer systems, or the technical delivery systems of Twitter's providers (except as expressly permitted by the Twitter Bug Bounty program).
- Probe, scan, or test the vulnerability of any system or network, or breach or circumvent any security or authentication measures (except as expressly permitted by the Twitter Bug Bounty program).
- Access or search, or attempt to access or search, Twitter by any means (automated or otherwise) other than through our currently available, published interfaces that are provided by Twitter (and only pursuant to the applicable terms and conditions), unless you have been specifically allowed to do so in a separate agreement with Twitter. Note that crawling Twitter is permissible if done in accordance with the provisions of the robots.txt file; however, scraping Twitter without our prior consent is expressly prohibited.
- Forge any TCP/IP packet header or any part of the header information in any email or posting, or in any way use Twitter to send altered, deceptive, or false source-identifying information.
- Interfere with or disrupt the access of any user, host or network, including, without limitation, sending a virus, overloading, flooding, spamming, mail-bombing Twitter's services, or by scripting the creation of content in such a manner as to interfere with or create an undue burden on Twitter.

Any accounts engaging in the following activities may be temporarily locked or subject to permanent suspension:

- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use Twitter's services for the purpose of spamming anyone. Spam is generally defined on Twitter as bulk or aggressive activity that attempts to manipulate or disrupt Twitter or the experience of users on Twitter to drive traffic or attention to unrelated accounts, products, services, or initiatives. Some of the factors that we take into account when determining what conduct is considered to be spamming include:
 - if you have followed and/or unfollowed a large number of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if your Tweets or Direct Messages consist mainly of links shared without commentary;
 - if a large number of people have blocked you in response to high volumes of untargeted, unsolicited, or duplicative content or engagements from your account;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicative or substantially similar content, replies, or mentions over multiple accounts or multiple duplicate updates on one account, or create duplicate or substantially similar accounts;

- if you post multiple updates to a trending or popular topic with an intent to subvert or manipulate the topic to drive traffic or attention to unrelated accounts, products, services, or initiatives;
- if you send large numbers of unsolicited replies or mentions;
- if you add users to lists in a bulk or aggressive manner;
- if you are randomly or aggressively engaging with Tweets (e.g., likes, Retweets, etc.) or users (e.g., following, adding to lists or Moments, etc.) to drive traffic or attention to unrelated accounts, products, services, or initiatives;
- if you repeatedly post other people's account information as your own (e.g., bio, Tweets, profile URL, etc.);
- if you post misleading, deceptive, or malicious links (e.g., affiliate links, links to malware/clickjacking pages, etc.);
- if you create fake accounts, account interactions, or impressions;
- if you sell, purchase, or attempt to artificially inflate account interactions (such as followers, Retweets, likes, etc.); and
- if you use or promote third-party services or apps that claim to get you more followers, Retweets, or likes (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers or engagements to your account or Tweets).

Please see our support articles on [following rules and best practices](#) and [automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts may be permanently suspended.

Content Visibility

Accounts under investigation or which have been detected as sharing content in violation of these Rules may have their account or Tweet visibility limited in various parts of Twitter, including search. To learn more about situations in which content may be restricted on Twitter, please see our support article on [search rules and restrictions](#).

Trademark policy

Updated June 30, 2016

What is a Trademark Policy Violation on Twitter?

Using another's trademark in a manner that may mislead or confuse others about your brand affiliation may be a violation of our trademark policy.

How Does Twitter Respond to Trademark Reports?

We review each report and do the following:

- If we determine that there is a violation of our trademark policy, we may suspend the account. In some instances, we may give the account holder an opportunity to comply with our policies before suspending the account. For more information, see our [Parody, commentary, and fan account policy](#).
- We may also provide the account holder with your name and information included in the copy of the report.
- If you would like to report a trademark violation in a [Twitter Ad](#), please visit our [Trademark Policy for Promoted Ads](#) to learn more.

Please note that Twitter only investigates requests that are submitted by the trademark holder or by an authorized representative of the trademark holder.

What is not a Trademark Policy Violation?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples include:

- Using a trademark in a way that is outside the scope of the trademark registration (e.g. territory, or goods and services identified in the registration).
- Nominative and other fair uses of trademarks are protected uses under our trademark policy, so long as the account is clearly distinguished from the trademark owner. This includes use by resellers in certain regions and accounts engaging in parody, commentary, or news. For more information, see our [Parody, commentary, and fan account policy](#).

How do I Report a Trademark Policy Violation?

Submit trademark reports through our [Trademark Issue support form](#). Using the appropriate form ensures we have all the information necessary to process your report quickly. Please provide all information requested in the form. If you submit a report with incomplete information, we will follow up with you to request that information. Please note that this will result in a delay in processing your report.

If you have trouble using the form, try updating your browser or switching to another browser.

[Copyright policy](#)

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from **@Username** has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notification.

When Should I File a Counter-notice?

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);

2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Twitter media policy

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to easily share your content with others. Below are our guidelines for making sure your content is labeled appropriately.

Media on Twitter

You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.

In addition, you may not feature graphic content (such as media containing pornography or excessive violence) in live video, or in your profile image or header image.

If you see media you think may be in violation of our policy, see [this article](#) for information about how to report it.

Your media settings

All Twitter accounts have media settings to help Twitter understand the type of media you are posting and to help us assist others to find the type of content that they're interested in.

By appropriately marking your media settings, Twitter can identify potentially sensitive content that others users may not wish to see, such as violence or nudity. If you intend to regularly post such content, we request that you adjust your media setting:

1. Log in to your account on the web at twitter.com.
2. Click your **profile icon** in the top right navigation bar.
3. Select **Settings and privacy** from the drop-down menu.
4. Go to your [Privacy and safety](#) settings.
5. Look for the **Tweet media** section and check the box next to **Mark media you Tweet as containing material that may be sensitive**. **Note:** Enabling this setting will result in others seeing a message when they view your profile, alerting them of potentially sensitive media and asking them to confirm they still want to view it.

You can control your preferences for how and when you may see sensitive content in your [media settings](#).

What if I don't change my media settings and post sensitive media?

If one of your Tweets containing media is reported as sensitive, it will be sent to Twitter for review. If we find that the media wasn't marked as being potentially sensitive at the time of the upload:

- We label the media as being potentially sensitive, or in the case of live video, remove the content entirely.
- We may also change your account setting to **Mark media you Tweet as containing material that may be sensitive** so that future posts are marked accordingly.

You will receive an email letting you know that we've labeled your media and if we changed your media setting. You will be able to change your media settings on your account settings page, though this image and any other images that have been labeled will remain marked as potentially sensitive.

If you repeatedly upload media that are mislabeled, the Twitter team may:

- Label your media as having potentially sensitive content.
- Permanently adjust your account setting to **Mark media you Tweet as containing material that may be sensitive** so that your future posts have a warning message users must click through before viewing your media.

Twitter may also use automated techniques to detect and label potentially sensitive media.

Additionally, people who visit your profile may see a message letting them know the account may include potentially sensitive content and asking them to confirm they still want to view it. People who have opted in to see possibly sensitive content will still see your account without the message.

What happens if my profile image is in violation of the Twitter Rules?

If your profile or header image contains graphic content in violation of the Twitter Rules, we may suspend your account and remove the offending content. Repeat violations will result in a permanent suspension.

Additionally, your account may be locked if media in your profile is in violation of Twitter's media policy. In order to unlock your account, you will need to follow our instructions and delete the media in violation. If your account has been locked, see this [article](#) for more information.

Can I appeal a decision made by the Twitter team?

If you are suspended for having a pornographic profile, header, or background image, you can appeal by logging in to your account and submitting an appeal [here](#).

There is currently no way to appeal a decision by the Twitter team that permanently changes your account setting to **Mark media you Tweet as containing material that may be sensitive** in response to repeated mislabeling of sensitive content.

Contacting Twitter about media concerning a deceased family member

In order to respect the wishes of loved ones, Twitter will remove imagery of deceased individuals in certain circumstances. Immediate family members and other authorized individuals may request the removal of images or video of deceased individuals, from when critical injury occurs to the moments before or after death, by sending a request to Twitter Inc. via our [privacy form](#). When reviewing such media removal requests, Twitter considers public interest factors such as the newsworthiness of the content and may not be able to honor every request.

Note: For information about contacting Twitter to deactivate the account of a deceased or incapacitated user, please read [this article](#).

[How to mark media as sensitive links to Media policy](#)

FAQs about trends on Twitter

Where can I find trends?

Trends are available on the Twitter for iOS app, the Twitter for Android app, and [twitter.com](#).

On Twitter's mobile apps, you can find trends listed under the trends section



of the **Explore** tab When signed in to twitter.com on a desktop or laptop computer, trends are listed in many places, including the Home timeline, Notifications, search results, and profile pages.

How are trends determined?

Trends are determined by an algorithm and, by default, are tailored for you based on who you follow, your interests, and your location. This algorithm identifies topics that are popular now, rather than topics that have been popular for a while or on a daily basis, to help you discover the hottest emerging topics of discussion on Twitter.

Note: The number of Tweets that are related to the trends is just one of the factors the algorithm looks at when ranking and determining trends. Algorithmically, trends and hashtags are grouped together if they are related to the same topic. For instance, #MondayMotivation and #MotivationMonday may both be represented by #MondayMotivation.

You can choose to see trends that are not tailored for you by selecting a [specific trends location](#) on twitter.com, iOS, or Android. Location trends identify popular topics among people in a specific geographic location.

What does the # sign mean?

You may notice that some trends have # sign before the word or phrase. This is called a [hashtag](#) and is included specifically in Tweets to mark them as relating to a topic, so that people can follow the conversation in search.

Will I see world and local events in my trends?

Yes, when signed into your account on the twitter.com, iOS, or Android, trends will be tailored for you based on your location and who you follow. There will be many world and local news events and conversations that will appear in your trends regardless of your personalization.

Can I see trends for a specific location?

Yes, you can view trends for a specific location on twitter.com, iOS, or Android.

View trends via twitter.com:

To view trends for a specific location, click **Change** in your trends box. You can choose to **Keep tailored trends** based on your location and who you follow, or you can select **Change** to choose a nearby location or search locations.

Adjust your trends settings via Twitter for iOS:

1. In the top menu, tap your **profile** icon.
2. Tap **Settings and privacy**, then tap **Content preferences**.
3. Under **Explore**, tap **Trends**.
4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.

5. When the feature is disabled, you can change your location by tapping **Change location**.
 1. In the search box, type your desired trends location or choose from the list.

Adjust your trends settings via Twitter for Android:



1. In the top menu, you will either see a **navigation menu** icon or your **profile** icon. Tap whichever icon you have, and select **Settings and privacy**.
2. Tap **Content preferences**.
3. Under **Explore**, tap **Trends**.
4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.
5. When the feature is disabled, you can change your location by tapping **Change location**.
 1. In the search box, type your desired trends location or choose from the list.

Note: If you don't find your preferred country or city, it means we're not receiving enough Tweets from that geographical area to create a list. You can look up local Tweets on any topic by using [advanced Twitter Search operators](#).

What happens when I click or tap on a trend?

Clicking or tapping any of the trends takes you to the Twitter search results for that trend. You'll see all Tweets including that phrase or hashtag. To see what people are saying about a previous trend, perform a [search for that keyword](#).

How can I participate in a trend?

Simply post a Tweet including the exact word or phrase as it appears in the trends list (with the hashtag, if you see one). Due to the large number of people Tweeting about these specific trends, you may not always be able to find your particular Tweet in search, but your followers will always see your Tweets.

Note: Twitter also filters searches for quality. To understand what types of behavior could cause your Tweets to be filtered out, read about our [search rules and restrictions](#).

Are there rules for trends?

Yes. We want trends to promote healthy discussions on Twitter. This means that at times, we may prevent certain content from trending. These include trends that:

- Contain profanity or adult/graphic references.
- Incite hate on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease.
- Violate the [Twitter Rules](#).

In some cases, we may also consider the newsworthiness of the content, or if it is in the public interest when evaluating potential violations. In these cases, the content might continue to trend on our platform.

Note that even if we prevent the hashtag or the content from appearing on the list of trending topics, you may still be able to access conversations around that content on Twitter.

Trends descriptions

Below some trends, you may see the approximate number of Tweets associated with the trend, or personalized context such as who in your network is Tweeting about it.

Additionally, when a Moment is found to be relevant to a trend, we may use the title of the Moment as the trend's description.

Having trouble?

Check out our [troubleshooting section](#) for solutions to common problems.

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Child sexual exploitation policy

Twitter removes content promoting child sexual exploitation

We do not tolerate child sexual exploitation on Twitter. When we are made aware of links to images of or content promoting child sexual exploitation they will be removed from the site without further notice and reported to The National Center for Missing & Exploited Children ("NCMEC"); we permanently suspend accounts promoting or containing updates with links to child sexual exploitation.

Reporting content promoting child sexual exploitation

If you think you've found a Twitter account distributing or promoting child sexual exploitation while using Twitter, please click [here](#) to report it. Please provide the username and links to the relevant tweets

that lead you to believe the account should be investigated. To find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with the National Center for Missing and Exploited Children [here](#), or your local law enforcement. If you believe that the content was hosted on a website outside of the United States, you can report it [here](#) on the International Association of Internet Hotlines website.

NOTE: Please do NOT tweet, retweet or repost child sexual exploitation for any reason. Click [here](#) to report it to us immediately and we will take steps to remove it.

Law Enforcement Agencies

If you are a member of a law enforcement agency seeking information regarding Twitter accounts promoting or distributing child sexual exploitation, please see our [Guidelines for Law Enforcement](#) for information on our process; we also suggest you contact the National Center for Missing and Exploited Children (NCMEC) [here](#), regardless of jurisdiction, to determine whether or not information on those accounts is already on file.

For general law enforcement inquiries, please see our [Guidelines for Law Enforcement](#) for more information.

Hateful conduct policy

Freedom of expression means little if voices are silenced because people are afraid to speak up. We do not tolerate behavior that harasses, intimidates, or uses fear to silence another person's voice. If you see something on Twitter that violates these rules, please report it to us.

How our policy works

As explained in the Twitter Rules,

- **Hateful conduct:** You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Examples of what we do not tolerate includes, but is not limited to behavior that harasses individuals or groups of people with:

- violent threats;
- wishes for the physical harm, death, or disease of individuals or groups;
- references to mass murder, violent events, or specific means of violence in which/with which such groups have been the primary targets or victims;
- behavior that incites fear about a protected group;

- repeated and/or or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone.

How our enforcement works

- **Context matters.**
 - Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. While we accept reports of violations from anyone, sometimes we also need to hear directly from the target to ensure that we have proper context.
 - The number of reports we receive does not impact whether or not something will be removed. However, it may help us prioritize the order in which it gets reviewed.
- **We focus on behavior.**
 - We enforce policies when someone reports behavior that is abusive and targets an entire protected group and/or individuals who may be members.
 - This targeting can happen in any manner (for example, @mentions, tagging a photo, and more).
- **We have a range of enforcement options.**
 - The consequences for violating our rules vary depending on the severity of the violation and the person's previous record of violations. For example, we may ask someone to remove the offending Tweet before they can Tweet again. For other cases, we may suspend an account.

Private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Some examples of private and confidential information include, but may not be limited to:

Intimate photos or videos that were taken or distributed without the subject's consent.
Images or videos that are considered and treated as private under applicable laws. Private contact or financial information, such as:

credit card information
social security or other national identity numbers
addresses or locations that are considered and treated as private non-public, personal phone

numbers
non-public, personal email addresses

How our enforcement works

Keep in mind that although you may consider certain information to be private, not all postings of such information may be a violation of this policy.

Context matters.

We may consider the context and nature of the information posted, local privacy laws, and other case-specific facts when determining if this policy has been violated. For example, if information was previously posted or displayed elsewhere on the Internet prior to being put on Twitter, it may not be a violation of this policy.

We focus on behavior.

We enforce policies when someone reports content where the behavior may be abusive.

We have a range of enforcement options.

The consequences for violating our rules vary depending on the severity of the violation and the person's previous record of violations. For example, we may ask someone to remove the offending Tweet before they can Tweet again. For other cases, we may suspend an account.

For frequently asked questions about reporting private information posted on Twitter, [click here](#). To learn more about protecting your private information on Twitter and other websites, [click here](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter impersonation policy.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to Twitter's [parody, commentary, and fan account policy](#) for more information about these types of accounts.

Find more information about [reporting impersonation on Twitter](#).

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation rules

Updated November 3, 2017

This page is primarily intended for **developers**.

For Twitter users: You are ultimately responsible for the actions taken with your account, or by applications associated with your account. Before authorizing a third-party application to access or use your account, make sure you've thoroughly investigated the application and understand what it will do. If automated activity on your account violates the [Twitter Rules](#) or these automation rules, Twitter may take action on your account, including [filtering your Tweets from search results](#) or suspending your account.

For more information on third-party applications, please see our article on [connecting and revoking third-party applications](#).

If you are a developer and have questions about these Automation Rules, please reach out via <https://twittercommunity.com/c/rules-and-policies>.

If your application has been suspended or restricted from performing write actions, you may submit an appeal at </forms/platform>.

I. Ground Rules Do!

Build solutions that automatically broadcast helpful information in Tweets.

Run creative campaigns that auto-reply to users who engage with your content.

Build solutions that automatically respond to users in Direct Messages.

Try new things that help people (and comply with our rules).

Make sure your application provides a good user experience and performs well — and confirm that remains the case over time.

Don't!

Violate these or other policies. Be extra mindful of our rules about abuse and user privacy.

Abuse the Twitter API or attempt to circumvent rate limits.

Use non-API-based forms of automation, such as scripting the Twitter website. The use of these techniques may result in the permanent suspension of your account.

Spam or bother users, or otherwise send them unsolicited messages.

A. The Twitter Rules and the Developer Agreement and Policy

As with all activity on Twitter, automated activity is subject to the [Twitter Rules](#) and, if you're a developer using the Twitter API, the [Developer Agreement and Policy](#).

You should carefully review these policies to ensure that your automated activity is compliant. Automated applications or activities that violate these policies, or that facilitate or induce users to violate them, may be subject to enforcement action, potentially including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these policies.

Although all aspects of the Twitter Rules and the Developer Agreement and Policy apply to automated activity, you should keep the following rules top of mind:

Spamming: You may not send automated Tweets or Direct Messages that are spam, or otherwise engage in spamming activity. Some examples of spammy behavior to avoid with automation include:

Trending topics: You may not automatically post about trending topics on Twitter, or use automation to attempt to influence or manipulate trending topics.

Multiple posts/accounts: You may not post duplicative or substantially similar Tweets on one account or over multiple accounts you operate.

Duplicate accounts: You may not create and/or automate multiple accounts for duplicative or substantially similar use cases.

However, automating multiple accounts for related but non-duplicative use cases is permitted. For example, you may automate separate accounts to Tweet when the Hubble Space Telescope passes over different cities, such as [San Francisco](#) or [Hong Kong](#).

Misleading links: You may not send automated Tweets or Direct Messages containing links that are misleading, including links that maliciously or deceptively redirect through landing pages or ad pages before displaying the final content.

Sensitive media: Automated Tweets and Direct Messages must comply with the [Twitter media policy](#), and you should mark your account as potentially sensitive if you intend to post graphic, pornographic, or potentially sensitive media.

Abusive behavior: You may not engage in any automated activity that encourages, promotes, or incites abuse, violence, hateful conduct, or harassment, on or off Twitter.

Private information: You may not post private or confidential information about a person without their prior express authorization.

B. Other Ground Rules for Automated Activity

In addition to the policies above, the following ground rules apply to all automated activity on Twitter:

Don't surprise or mislead users: Automated activity should honor users' expectations. Ask for the user's permission before taking an action if you aren't sure.

Mature content or profanity: Don't Direct Message, mention, or reply to users with potentially sensitive content (including profanity), unless they've clearly indicated an intent to receive it in advance.

Be thoughtful about the information you request or exchange on Twitter

Tweets: Don't ask users to send you personal or private information via a public Tweet. If you need additional personal or private information from a user to provide them with customer service (or other similar use cases), you should ask the user to share such information by Direct Message or another private channel. You might even consider adding a [Direct Message deep link](#) to your Tweet.

Direct Messages: You should only ask users for the minimum amount of information you need to provide them with service. If you need to request or exchange particularly sensitive information (such as credit card information), you should consider directing users to your website or other appropriate channel to do so.

II. Activity-Specific Rules

The activity-specific rules in this section apply to taking specific automated actions on Twitter. Please read these rules carefully, as they outline both permitted and prohibited use cases of automation.

Automated applications or activities that violate these rules, or that facilitate or induce users to violate them, may be subject to enforcement action, including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access the Twitter API based on violations of these rules. As a reminder, you should also carefully review the spam guidelines in the [Twitter Rules](#) to avoid having activities performed by you, your app, or other users through your app or service flagged as spam.

A. Automated Actions Through Another User's Account

Twitter users may authorize your app or service to [access their Twitter account through OAuth](#). A user authorizing your app or service to access their Twitter account through OAuth does not by itself constitute sufficient consent to take automated actions through that user's account.

You may only take automated actions through another Twitter user's account if you:

clearly describe to the user the types of automated actions that will occur; receive express consent from the user to take those automated actions; and immediately honor a user's request to opt-out of further automated actions.

If you substantially change the purpose or functionality of your app or service, you must re-obtain express consent from the user to take automated action through their account before doing so.

These requirements apply to any automated action taken through another Twitter user's account, including posting Tweets, sending Direct Messages, deleting Tweets or Direct Messages, or following/unfollowing other accounts. For applications that offer users the ability to delete Tweets in a bulk or automated manner, you must also clearly state that Tweets are not recoverable once deleted.

B. Automated Tweets

1. Posting automated Tweets

Automated Tweets that cross-post outside information: You may post automated Tweets based on sources of outside information — such as an RSS feed, weather data, etc. — as long as you are sufficiently authorized to publish such information.

Other automated Tweets (excluding mentions or replies): Provided you comply with all other rules, you may post automated Tweets for entertainment, informational, or novelty purposes. As a reminder, accounts posting duplicative, spammy, or otherwise prohibited content may be subject to suspension.

2. Posting automated mentions and replies

The reply and mention functions are intended to make communication between Twitter users easier. Automating these actions to reach many users on an unsolicited basis is an abuse of the feature, and is not permitted. For example, sending automated replies to Tweets based on keyword searches alone is not permitted. Spamy or duplicative use of mentions and replies may result in enforcement action, such as the removal of

your Tweets from Search or the suspension of your app or account.

However, you may send automated replies or mentions to Twitter users so long as:

in advance of sending the automated reply, the recipient or mentioned user(s) have requested or have clearly indicated an intent on Twitter to be contacted by you (i.e. opted in), for example by replying to a Tweet from your account, or by sending you a Direct Message;

you provide a clear and easy way for such users to opt-out of receiving automated replies and mentions, and promptly honor all such opt-out requests;

you only send one automated reply or mention per user interaction; and

the automated reply or mention is a reply to the user's original Tweet (if your campaign is based on users posting a reply to your Tweet).

Opt-in techniques and indications of user intent take many different forms, depending on the specifics of your use case and implementation. Some examples include:

A Tweet from your account that clearly indicates that a user taking a specific action on that Tweet (such as Retweeting it) will opt the user into receiving an automated response.

A mention of your account by the user in a manner suggesting the user clearly wishes or intends to receive a response. If you want to run an auto-reply campaign with a campaign- or use-case-specific hashtag, users should also mention you in their Tweets.

Note that a user following your account is not on its own a sufficient indication of user intent to receive an automated response.

: Advertisers, publishers, and brands using auto-response campaigns must request approval from Twitter and may be subject to additional rules. Please reach out to your account or partner manager for assistance. If you do not have a partner manager but are interested in managed client services, you can learn more [here](#).

Additionally, we recommend that any accounts that will communicate with

users via automated mentions or replies:

Note

1. Appropriately filter responses based on potentially sensitive language in user handles,

display names, and Tweet text, as well as potentially sensitive media;

2. Check that the Tweet you are mentioning or replying to still exists (for example, using the statuses/lookup endpoint on the Twitter API).

C. Automated Direct Messages

Sending automated Direct Messages to users

You may not send unsolicited Direct Messages in a bulk or automated manner, and should be thoughtful about the frequency with which you contact users via Direct Message.

You may send automated Direct Messages to users so long as:

in advance of sending the Direct Message, the recipient(s) have requested or have clearly indicated an intent on Twitter to be contacted by you via Direct Message, for example by sending you a Direct Message; **and**

you provide a clear and easy way for such users to opt-out of receiving automated Direct Messages, and promptly honor all such opt-out requests.

The fact that a user is technically able to receive a Direct Message from you (e.g. because the user follows you, has enabled the ability to receive Direct Messages from any account, or because the user is in a pre-existing Direct Message conversation with you) does not necessarily mean they have requested or expect to receive automated Direct Messages from you.

Interacting with users via Direct Message

Per the Ground Rules, remember to be thoughtful about the amount and type of information you request or exchange with users via Direct Messages. If you will be asking a user to provide personal or private information via an automated Direct Message, you must clearly explain how you will use the information you're collecting. Consider including a link to your

privacy policy in your Direct Message to the user, as well as in your Twitter profile bio.

Don't publicly share information received in a Direct Message conversation with a user without first obtaining explicit consent from the user. For example, if a user asks you via Direct Message about a purchase they made from you, you may not mention the user in a public Tweet that includes information about their purchase unless you have the user's explicit consent to do so.

After a user-initiated interaction ends, don't send additional follow-up Direct Messages or mention users in a Tweet unless you get permission from the user.

D. Automated actions you take on Tweets or accounts

Automated likes: You may not like Tweets in an automated manner.

Automated Retweets: Provided you comply with all other rules, you may Retweet or Quote Tweet in an automated manner for entertainment, informational, or novelty purposes. Automated Retweets often lead to negative user experiences, and bulk, aggressive, or spammy Retweeting is a violation of the [Twitter Rules](#).

Automated following/unfollowing: You may not follow or unfollow Twitter accounts in a bulk, aggressive, or indiscriminate manner. Aggressive following is a violation of the [Twitter Rules](#). Please also review our [following rules and best practices](#) to ensure you are in compliance. Note that applications that claim to get users more followers are also prohibited under the [Twitter Rules](#).

Automated adding to lists or collections: You may not add Twitter users to lists or add Tweets to collections in a bulk or indiscriminate manner. Adding a large number of unrelated users to lists is a violation of the [Twitter Rules](#).

Twitter search rules and restrictions

This article details the rules for using Twitter search, and explains reasons why content may not show up in Twitter search.

Why doesn't all content show up in Twitter search?

In order to keep your search results relevant, Twitter filters search results for quality Tweets and accounts. Material that jeopardizes search quality or creates a bad search experience for other people may be automatically removed from Twitter search. Read more about this [here](#).

General rules for using Twitter search

The [Twitter Rules](#) explain what behaviors are permitted on Twitter. In addition to these rules, we've included some tips below to keep your content relevant (and your Tweets in Twitter search).

Please do not:

- Repeatedly post duplicate or near-duplicate content (links or Tweets).
- Abuse trending topics or hashtags (topic words with a # sign).
- Send automated Tweets or replies.
- Use bots or applications to post similar messages based on keywords.
- Post similar messages over multiple accounts.
- Aggressively follow and unfollow people.

Twitter may automatically remove accounts engaging in these behaviors from search (or even suspend in some cases) in order to ensure the best experience for everyone.

Search results filtering

To learn how to filter your search results, including how to disable or enable safe search, read our [Using Twitter search](#) article.

Search results filtering in third-party applications

Accounts may be filtered from search for posting similar messages over several accounts. Your account may be filtered from search if you are logged into a third-party application that is updating many accounts with similar or duplicate content.

Rules for participating in contests on Twitter

Businesses sometimes host contests and offer prizes to Retweet an update, following a particular account, or posting Tweets with a specific hashtag topic or reply. If you're running a contest, check out this [article](#) for guidelines for contests on Twitter.

If you repeatedly post near-duplicate updates or duplicate links in order to enter contests, you may be filtered from search. Please keep in mind that you are responsible for the updates that you post to your Twitter account.

What to do if you are missing from search

Please see this [article](#) for more information related to this issue.

If you are being filtered from search, rest assured that your followers will still receive your updates. If you post a reply or mention to a specific person, it will be delivered to that person.

Having trouble?

You can learn more about searching by keyword, location, and trending topics in our [Using Twitter search](#) help article. Our [search troubleshooting](#) section includes solutions to common issues related to search.

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. These limitations are set forth in the Twitter Rules below.

The Twitter Rules (along with all incorporated policies), [Privacy Policy](#), and [Terms of Service](#) collectively make up the "Twitter User Agreement" that governs a user's access to and use of Twitter's services.

All individuals accessing or using Twitter's services must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in Twitter taking one or more of the following enforcement actions:

- requiring you to delete prohibited content before you can again create new posts and interact with other Twitter users;
- temporarily limiting your ability to create posts or interact with other Twitter users;
- asking you to verify account ownership with a phone number or email address; or
- permanently suspending your account(s).

If you attempt to evade a permanent suspension by creating new accounts, we will suspend your new accounts.

Please note that we may need to change these Rules from time to time and reserve the right to do so. The most current version is always available at: <https://twitter.com/rules>.

The policies set forth in these Twitter Rules govern organic content on our platform. To learn more about the rules which govern ads and promoted content, please review our [Ads policies](#).

Content Boundaries and Use of Twitter

Intellectual property

Trademark: We reserve the right to suspend accounts or take other appropriate action when someone's brand or trademark, including business name and/or logo, is used in a manner that may mislead or confuse others about your brand affiliation. Read more about our [trademark policy](#) and [how to report a violation](#).

Copyright: We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in our [Terms of Service](#). Read more about our [copyright policy](#).

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Learn more about our [media policy](#).

Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images. Additionally, Twitter may sometimes require you to remove excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about [how to make such a request](#), and [how to mark your media as sensitive](#).

Unlawful use

You may not use our service for any unlawful purposes or in furtherance of illegal activities. By using Twitter, you agree to comply with all applicable laws governing your online conduct and content.

Trends

At times, we may prevent certain content from trending. This includes content that violates the Twitter Rules, as well as content that may attempt to manipulate trends. Read more about [what we allow and do not allow to trend](#).

Misuse of Twitter badges

You may not use badges, including but not limited to the “promoted” or “verified” Twitter badges, unless provided by Twitter. Accounts using unauthorized badges as part of their profile photos, header photos, display names, or in any way that falsely implies affiliation with Twitter or authorization from Twitter to display these badges, may be suspended.

Misuse of usernames

Selling usernames: You may not buy or sell Twitter usernames.

Username squatting: You may not engage in username squatting. Some of the factors we take into consideration when determining whether conduct is username squatting include:

- the number of accounts created;
- the creation of accounts for the purpose of preventing others from using those account names;
- the creation of accounts for the purpose of selling those accounts; and
- the use of third-party content feeds to update and maintain accounts under the names of those third parties.

Please note that Twitter may also remove accounts that are inactive for more than six months. Learn more about [username squatting](#).

Abusive Behavior

We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we prohibit behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Context matters when evaluating for abusive behavior and determining appropriate enforcement actions. Factors we may take into consideration include, but are not limited to whether:

- the behavior is targeted at an individual or group of people;
- the report has been filed by the target of the abuse or a bystander;
- the behavior is newsworthy and in the legitimate public interest.

Violence and physical harm

Violence: You may not make specific threats of violence or wish for the serious physical harm, death, or disease of an individual or group of people. This includes, but is not limited to, threatening or promoting terrorism. You also may not affiliate with organizations that — whether by their own statements or activity both on and off the platform — use or promote violence against civilians to further their causes. We will begin enforcing this rule around affiliation with such organizations on December 18, 2017.

Suicide or self-harm: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Child sexual exploitation: You may not promote child sexual exploitation. Learn more about our zero-tolerance [child sexual exploitation policy](#).

Abuse and hateful conduct

Abuse: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Unwanted sexual advances: You may not direct abuse at someone by sending unwanted sexual content, objectifying them in a sexually explicit manner, or otherwise engaging in sexual misconduct.

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. Read more about our [hateful conduct policy](#).

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category. We will begin enforcing this rule on December 18, 2017.

Private information and intimate media

Private information: You may not publish or post other people's private information without their express authorization and permission. Definitions of private information may vary depending on local laws. Read more about our [private information policy](#).

Intimate media: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. Read more about intimate media on Twitter.

Threats to expose / hack: You may not threaten to expose someone's private information or intimate media. You also may not threaten to hack or break into someone's digital information.

Impersonation

You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. While you may maintain parody, fan, commentary, or newsfeed accounts, you may not do so if the intent of the account is to engage in spamming or abusive behavior. Read more about our [impersonation policy](#).

Spam and Security

We strive to protect people on Twitter from technical abuse and spam.

To promote a stable and secure environment on Twitter, you may not do, or attempt to do, any of the following while accessing or using Twitter:

- Access, tamper with, or use non-public areas of Twitter, Twitter's computer systems, or the technical delivery systems of Twitter's providers (except as expressly permitted by the Twitter Bug Bounty program).
- Probe, scan, or test the vulnerability of any system or network, or breach or circumvent any security or authentication measures (except as expressly permitted by the Twitter Bug Bounty program).
- Access or search, or attempt to access or search, Twitter by any means (automated or otherwise) other than through our currently available, published interfaces that are provided by Twitter (and only pursuant to the applicable terms and conditions), unless you have been specifically allowed to do so in a separate agreement with Twitter. Note that crawling Twitter is permissible if done in accordance with the provisions of the robots.txt file; however, scraping Twitter without our prior consent is expressly prohibited.
- Forge any TCP/IP packet header or any part of the header information in any email or posting, or in any way use Twitter to send altered, deceptive, or false source-identifying information.
- Interfere with or disrupt the access of any user, host or network, including, without limitation, sending a virus, overloading, flooding, spamming, mail-bombing Twitter's services, or by scripting the creation of content in such a manner as to interfere with or create an undue burden on Twitter.

Any accounts engaging in the following activities may be temporarily locked or subject to permanent suspension:

- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use Twitter's services for the purpose of spamming anyone. Spam is generally defined on Twitter as bulk or aggressive activity that attempts to manipulate or disrupt Twitter or the experience of users on Twitter to drive traffic or attention to unrelated accounts, products, services, or initiatives. Some of the factors that we take into account when determining what conduct is considered to be spamming include:
 - if you have followed and/or unfollowed a large number of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if your Tweets or Direct Messages consist mainly of links shared without commentary;
 - if a large number of people have blocked you in response to high volumes of untargeted, unsolicited, or duplicative content or engagements from your account;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicative or substantially similar content, replies, or mentions over multiple accounts or multiple duplicate updates on one account, or create duplicate or substantially similar accounts;
 - if you post multiple updates to a trending or popular topic with an intent to subvert or manipulate the topic to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you send large numbers of unsolicited replies or mentions;
 - if you add users to lists in a bulk or aggressive manner;
 - if you are randomly or aggressively engaging with Tweets (e.g., likes, Retweets, etc.) or users (e.g., following, adding to lists or Moments, etc.) to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you repeatedly post other people's account information as your own (e.g., bio, Tweets, profile URL, etc.);
 - if you post misleading, deceptive, or malicious links (e.g., affiliate links, links to malware/clickjacking pages, etc.);
 - if you create fake accounts, account interactions, or impressions;
 - if you sell, purchase, or attempt to artificially inflate account interactions (such as followers, Retweets, likes, etc.); and
 - if you use or promote third-party services or apps that claim to get you more followers, Retweets, or likes (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers or engagements to your account or Tweets).

Please see our support articles on [following rules and best practices](#) and [automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts may be permanently suspended.

Content Visibility

Accounts under investigation or which have been detected as sharing content in violation of these Rules may have their account or Tweet visibility limited in various parts of Twitter, including search. To learn more about situations in which content may be restricted on Twitter, please see our support article on [search rules and restrictions](#).

Trademark policy

Updated June 30, 2016

What is a Trademark Policy Violation on Twitter?

Using another's trademark in a manner that may mislead or confuse others about your brand affiliation may be a violation of our trademark policy.

How Does Twitter Respond to Trademark Reports?

We review each report and do the following:

- If we determine that there is a violation of our trademark policy, we may suspend the account. In some instances, we may give the account holder an opportunity to comply with our policies before suspending the account. For more information, see our [Parody, commentary, and fan account policy](#).
- We may also provide the account holder with your name and information included in the copy of the report.
- If you would like to report a trademark violation in a [Twitter Ad](#), please visit our [Trademark Policy for Promoted Ads](#) to learn more.

Please note that Twitter only investigates requests that are submitted by the trademark holder or by an authorized representative of the trademark holder.

What is not a Trademark Policy Violation?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples include:

- Using a trademark in a way that is outside the scope of the trademark registration (e.g. territory, or goods and services identified in the registration).
- Nominative and other fair uses of trademarks are protected uses under our trademark policy, so long as the account is clearly distinguished from the trademark owner. This includes use by resellers in certain regions and accounts engaging in parody, commentary, or news. For more information, see our [Parody, commentary, and fan account policy](#).

How do I Report a Trademark Policy Violation?

Submit trademark reports through our [Trademark Issue support form](#). Using the appropriate form ensures we have all the information necessary to process your report quickly. Please provide all information requested in the form. If you submit a report with incomplete information, we will follow up with you to request that information. Please note that this will result in a delay in processing your report.

If you have trouble using the form, try updating your browser or switching to another browser.

Copyright policy

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged

infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notification.

When Should I File a Counter-notice?

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);

3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Twitter media policy

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to easily share your content with others. Below are our guidelines for making sure your content is labeled appropriately.

Media on Twitter

You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.

In addition, you may not feature graphic content (such as media containing pornography or excessive violence) in live video, or in your profile image or header image.

If you see media you think may be in violation of our policy, see [this article](#) for information about how to report it.

Your media settings

All Twitter accounts have media settings to help Twitter understand the type of media you are posting and to help us assist others to find the type of content that they're interested in.

By appropriately marking your media settings, Twitter can identify potentially sensitive content that other users may not wish to see, such as violence or nudity. If you intend to regularly post such content, we request that you adjust your media setting:

1. Log in to your account on the web at twitter.com.
2. Click your **profile icon** in the top right navigation bar.
3. Select **Settings and privacy** from the drop-down menu.
4. Go to your [Privacy and safety](#) settings.
5. Look for the **Tweet media** section and check the box next to **Mark media you Tweet as containing material that may be sensitive**. **Note:** Enabling this setting will result in others seeing a message when they view your profile, alerting them of potentially sensitive media and asking them to confirm they still want to view it.

You can control your preferences for how and when you may see sensitive content in your [media settings](#).

What if I don't change my media settings and post sensitive media?

If one of your Tweets containing media is reported as sensitive, it will be sent to Twitter for review. If we find that the media wasn't marked as being potentially sensitive at the time of the upload:

- We label the media as being potentially sensitive, or in the case of live video, remove the content entirely.
- We may also change your account setting to **Mark media you Tweet as containing material that may be sensitive** so that future posts are marked accordingly.

You will receive an email letting you know that we've labeled your media and if we changed your media setting. You will be able to change your media settings on your account settings page, though this image and any other images that have been labeled will remain marked as potentially sensitive.

If you repeatedly upload media that are mislabeled, the Twitter team may:

- Label your media as having potentially sensitive content.
- Permanently adjust your account setting to **Mark media you Tweet as containing material that may be sensitive** so that your future posts have a warning message users must click through before viewing your media.

Twitter may also use automated techniques to detect and label potentially sensitive media.

Additionally, people who visit your profile may see a message letting them know the account may include potentially sensitive content and asking them to confirm they still want to view it. People who have opted in to see possibly sensitive content will still see your account without the message.

What happens if my profile image is in violation of the Twitter Rules?

If your profile or header image contains graphic content in violation of the Twitter Rules, we may suspend your account and remove the offending content. Repeat violations will result in a permanent suspension.

Additionally, your account may be locked if media in your profile is in violation of Twitter's media policy. In order to unlock your account, you will need to follow our instructions and delete the media in violation. If your account has been locked, see this [article](#) for more information.

Can I appeal a decision made by the Twitter team?

If you are suspended for having a pornographic profile, header, or background image, you can appeal by logging in to your account and submitting an appeal [here](#).

There is currently no way to appeal a decision by the Twitter team that permanently changes your account setting to **Mark media you Tweet as containing material that may be sensitive** in response to repeated mislabeling of sensitive content.

How to contact Twitter about media concerning a deceased family member

In order to respect the wishes of loved ones, Twitter will remove imagery of deceased individuals in certain circumstances. Immediate family members and other authorized individuals may request the removal of images or video of deceased individuals, from when critical injury occurs to the moments before or after death, by sending a request to Twitter Inc. via our [privacy form](#). When reviewing such media removal requests, Twitter considers public interest factors such as the newsworthiness of the content and may not be able to honor every request. Note: Find information about contacting Twitter to [deactivate the account of a deceased or incapacitated user](#).

How to mark media as sensitive links to Media policy

FAQs about trends on Twitter

Where can I find trends?

Trends are available on the Twitter for iOS app, the Twitter for Android app, and [twitter.com](#).

On Twitter's mobile apps, you can find trends listed under the trends section



of the **Explore** tab. When signed in to [twitter.com](#) on a desktop or laptop computer, trends are listed in many places, including the Home timeline, Notifications, search results, and profile pages.

How are trends determined?

Trends are determined by an algorithm and, by default, are tailored for you based on who you follow, your interests, and your location. This algorithm identifies topics that are popular now, rather

than topics that have been popular for a while or on a daily basis, to help you discover the hottest emerging topics of discussion on Twitter.

Note: The number of Tweets that are related to the trends is just one of the factors the algorithm looks at when ranking and determining trends. Algorithmically, trends and hashtags are grouped together if they are related to the same topic. For instance, #MondayMotivation and #MotivationMonday may both be represented by #MondayMotivation.

You can choose to see trends that are not tailored for you by selecting a [specific trends location](#) on twitter.com, iOS, or Android. Location trends identify popular topics among people in a specific geographic location.

What does the # sign mean?

You may notice that some trends have # sign before the word or phrase. This is called a [hashtag](#) and is included specifically in Tweets to mark them as relating to a topic, so that people can follow the conversation in search.

Will I see world and local events in my trends?

Yes, when signed into your account on the twitter.com, iOS, or Android, trends will be tailored for you based on your location and who you follow. There will be many world and local news events and conversations that will appear in your trends regardless of your personalization.

Can I see trends for a specific location?

Yes, you can view trends for a specific location on twitter.com, iOS, or Android.

View trends via twitter.com:

To view trends for a specific location, click **Change** in your trends box. You can choose to **Keep tailored trends** based on your location and who you follow, or you can select **Change** to choose a nearby location or search locations.

Adjust your trends settings via Twitter for iOS:

1. In the top menu, tap your **profile** icon.
2. Tap **Settings and privacy**, then tap **Content preferences**.
3. Under **Explore**, tap **Trends**.
4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.
5. When the feature is disabled, you can change your location by tapping **Change location**.
 1. In the search box, type your desired trends location or choose from the list.

Adjust your trends settings via Twitter for Android:

1. In the top menu, you will either see a **navigation menu** icon  or your **profile** icon. Tap whichever icon you have, and select **Settings and privacy**.

2. Tap **Content preferences**.
3. Under **Explore**, tap **Trends**.
4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.
5. When the feature is disabled, you can change your location by tapping **Change location**.
 1. In the search box, type your desired trends location or choose from the list.

Note: If you don't find your preferred country or city, it means we're not receiving enough Tweets from that geographical area to create a list. You can look up local Tweets on any topic by using [advanced Twitter Search operators](#).

What happens when I click or tap on a trend?

Clicking or tapping any of the trends takes you to the Twitter search results for that trend. You'll see all Tweets including that phrase or hashtag. To see what people are saying about a previous trend, perform a [search for that keyword](#).

How can I participate in a trend?

Simply post a Tweet including the exact word or phrase as it appears in the trends list (with the hashtag, if you see one). Due to the large number of people Tweeting about these specific trends, you may not always be able to find your particular Tweet in search, but your followers will always see your Tweets.

Note: Twitter also filters searches for quality. To understand what types of behavior could cause your Tweets to be filtered out, read about our [search rules and restrictions](#).

Are there rules for trends?

Yes. We want trends to promote healthy discussions on Twitter. This means that at times, we may prevent certain content from trending. These include trends that:

- Contain profanity or adult/graphic references.
- Incite hate on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease.
- Violate the [Twitter Rules](#).

In some cases, we may also consider the newsworthiness of the content, or if it is in the public interest when evaluating potential violations. In these cases, the content might continue to trend on our platform.

Note that even if we prevent the hashtag or the content from appearing on the list of trending topics, you may still be able to access conversations around that content on Twitter.

Trends descriptions

Below some trends, you may see the approximate number of Tweets associated with the trend, or personalized context such as who in your network is Tweeting about it.

Additionally, when a Moment is found to be relevant to a trend, we may use the title of the Moment as the trend's description.

Having trouble?

Check out our [troubleshooting section](#) for solutions to common problems.

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Child sexual exploitation policy

Twitter removes content promoting child sexual exploitation

We do not tolerate child sexual exploitation on Twitter. When we are made aware of links to images of or content promoting child sexual exploitation they will be removed from the site without further notice and reported to The National Center for Missing & Exploited Children ("NCMEC"); we permanently suspend accounts promoting or containing updates with links to child sexual exploitation.

Reporting content promoting child sexual exploitation

If you think you've found a Twitter account distributing or promoting child sexual exploitation while using Twitter, please click [here](#) to report it. Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. To find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with the National Center for Missing and Exploited Children [here](#), or your local law enforcement. If you believe that the content was hosted on a website outside of the United States, you can report it [here](#) on the International Association of Internet Hotlines website.

NOTE: Please do NOT tweet, retweet or repost child sexual exploitation for any reason. Click [here](#) to report it to us immediately and we will take steps to remove it.

Law Enforcement Agencies

If you are a member of a law enforcement agency seeking information regarding Twitter accounts promoting or distributing child sexual exploitation, please see our [Guidelines for Law Enforcement](#) for information on our process; we also suggest you contact the National Center for Missing and Exploited Children (NCMEC) [here](#), regardless of jurisdiction, to determine whether or not information on those accounts is already on file.

For general law enforcement inquiries, please see our [Guidelines for Law Enforcement](#) for more information.

Hateful conduct policy

Freedom of expression means little if voices are silenced because people are afraid to speak up. We do not tolerate behavior that harasses, intimidates, or uses fear to silence another person's voice. If you see something on Twitter that violates these rules, please report it to us.

How our policy works

As explained in the Twitter Rules,

- **Hateful conduct:** You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Examples of what we do not tolerate includes, but is not limited to behavior that harasses individuals or groups of people with:

- violent threats;
- wishes for the physical harm, death, or disease of individuals or groups;
- references to mass murder, violent events, or specific means of violence in which/with which such groups have been the primary targets or victims;
- behavior that incites fear about a protected group;
- repeated and/or or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone.

How our enforcement works

- **Context matters.**
 - Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. While we accept reports of violations

from anyone, sometimes we also need to hear directly from the target to ensure that we have proper context.

- The number of reports we receive does not impact whether or not something will be removed. However, it may help us prioritize the order in which it gets reviewed.

- **We focus on behavior.**

- We enforce policies when someone reports behavior that is abusive and targets an entire protected group and/or individuals who may be members.
- This targeting can happen in any manner (for example, @mentions, tagging a photo, and more).

- **We have a range of enforcement options.**

- The consequences for violating our rules vary depending on the severity of the violation and the person's previous record of violations. For example, we may ask someone to remove the offending Tweet before they can Tweet again. For other cases, we may suspend an account.
-

Report private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Read about our [policy regarding private information posted on Twitter](#).

Do I need to have a Twitter account to make a report?

No, you do not need to have a Twitter account to file a report that your private and confidential information has been posted on Twitter.

Can I report unauthorized posting of Copyrighted material?

If your copyrighted material has been Tweeted without your permission, please review our [Copyright and DMCA policy](#).

What can I do if my information is posted on other websites?

We cannot take down any information that is posted on a website other than Twitter. If your information is posted on another website, please contact the help center or support staff for the website where your information is posted.

How do I report that my private information is posted on Twitter?

You can report that your private information has been posted on Twitter directly from the offending Tweet(s). Learn [how to report individual Tweets](#).

File a [report that someone has posted your private information](#).

What happens when Twitter receives a valid report?

When we receive a complete and valid report that private information has been posted on Twitter, we'll investigate the account and Tweets reported. We will review where, if anywhere, the information has been made publicly available

before taking action on the account or Tweets. If the information you reported was previously posted elsewhere on the Internet, it is not a violation of our policy and we will not take action.

Learn more about [protecting your private information on Twitter and other websites](#).

Read our [policy regarding private information posted on Twitter](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter impersonation policy.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to Twitter's [parody, commentary, and fan account policy](#) for more information about these types of accounts.

Find more information about [reporting impersonation on Twitter](#).

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation rules

Updated November 3, 2017

This page is primarily intended for **developers**.

For Twitter users: You are ultimately responsible for the actions taken with your account, or by applications associated with your account. Before authorizing a third-party application to access or use your account, make sure you've thoroughly investigated the application and understand what it will do. If automated activity on your account violates the [Twitter Rules](#) or these automation rules, Twitter may take action on your account, including [filtering your Tweets from search results](#) or suspending your account.

For more information on third-party applications, please see our article on [connecting and revoking third-party applications](#).

If you are a developer and have questions about these Automation Rules, please reach out via <https://twittercommunity.com/c/rules-and-policies>.

If your application has been suspended or restricted from performing write actions, you may submit an appeal at </forms/platform>.

I. Ground Rules Do!

Build solutions that automatically broadcast helpful information in Tweets.

Run creative campaigns that auto-reply to users who engage with your content.

Build solutions that automatically respond to users in Direct Messages.

Try new things that help people (and comply with our rules).

Make sure your application provides a good user experience and performs well — and confirm that remains the case over time.

Don't!

Violate these or other policies. Be extra mindful of our rules about abuse and user privacy.

Abuse the Twitter API or attempt to circumvent rate limits.

Use non-API-based forms of automation, such as scripting the Twitter website. The use of these techniques may result in the permanent suspension of your account.

Spam or bother users, or otherwise send them unsolicited messages.

A. The Twitter Rules and the Developer Agreement and Policy

As with all activity on Twitter, automated activity is subject to the [Twitter Rules](#) and, if you're a developer using the Twitter API, the [Developer Agreement and Policy](#).

You should carefully review these policies to ensure that your automated activity is compliant. Automated applications or activities that violate these policies, or that facilitate or induce users to violate them, may be subject to enforcement action, potentially including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these policies.

Although all aspects of the Twitter Rules and the Developer Agreement and Policy apply to automated activity, you should keep the following rules top of mind:

Spamming: You may not send automated Tweets or Direct Messages that are spam, or otherwise engage in spamming activity. Some examples of spammy behavior to avoid with automation include:

Trending topics: You may not automatically post about trending topics on Twitter, or use automation to attempt to influence or manipulate trending topics.

Multiple posts/accounts: You may not post duplicative or substantially similar Tweets on one account or over multiple accounts you operate.

Duplicate accounts: You may not create and/or automate multiple accounts for duplicative or substantially similar use cases.

However, automating multiple accounts for related but non-duplicative use cases is permitted. For example, you may automate separate accounts to Tweet when the Hubble Space Telescope passes over different cities, such as [San Francisco](#) or [Hong Kong](#).

Misleading links: You may not send automated Tweets or Direct Messages containing links that are misleading, including links that maliciously or deceptively redirect through landing pages or ad pages before displaying the final content.

Sensitive media: Automated Tweets and Direct Messages must comply with the [Twitter media policy](#), and you should mark your account as potentially sensitive if you intend to post graphic, pornographic, or potentially sensitive media.

Abusive behavior: You may not engage in any automated activity that encourages, promotes, or incites abuse, violence, hateful conduct, or harassment, on or off Twitter.

Private information: You may not post private or confidential information about a person without their prior express authorization.

B. Other Ground Rules for Automated Activity

In addition to the policies above, the following ground rules apply to all automated activity on Twitter:

Don't surprise or mislead users: Automated activity should honor users' expectations. Ask for the user's permission before taking an action if you aren't sure.

Mature content or profanity: Don't Direct Message, mention, or reply to users with potentially sensitive content (including profanity), unless they've clearly indicated an intent to receive it in advance.

Be thoughtful about the information you request or exchange on Twitter

Tweets: Don't ask users to send you personal or private information via a public Tweet. If you need additional personal or private information from a user to provide them with customer service (or other similar use cases), you should ask the user to share such information by Direct Message or another private channel. You might even consider adding a [Direct Message deep link](#) to your Tweet.

Direct Messages: You should only ask users for the minimum amount of information you need to provide them with service. If you need to request or exchange particularly sensitive information (such as credit card information), you should consider directing users to your website or other appropriate channel to do so.

II. Activity-Specific Rules

The activity-specific rules in this section apply to taking specific automated actions on Twitter. Please read these rules carefully, as they outline both permitted and prohibited use cases of automation.

Automated applications or activities that violate these rules, or that facilitate or induce users to violate them, may be subject to enforcement action, including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access the Twitter API based on violations of these rules. As a reminder, you should also carefully review the spam guidelines in the [Twitter Rules](#) to avoid having activities performed by you, your app, or other users through your app or service flagged as spam.

A. Automated Actions Through Another User's Account

Twitter users may authorize your app or service to [access their Twitter account through OAuth](#). A user authorizing your app or service to access their Twitter account through OAuth does not by itself constitute sufficient consent to take automated actions through that user's account.

You may only take automated actions through another Twitter user's account if you:

clearly describe to the user the types of automated actions that will occur; receive express consent from the user to take those automated actions; and immediately honor a user's request to opt-out of further automated actions.

If you substantially change the purpose or functionality of your app or service, you must re-obtain express consent from the user to take automated action through their account before doing so.

These requirements apply to any automated action taken through another Twitter user's account, including posting Tweets, sending Direct Messages, deleting Tweets or Direct Messages, or following/unfollowing other accounts. For applications that offer users the ability to delete Tweets in a

bulk or automated manner, you must also clearly state that Tweets are not recoverable once deleted.

B. Automated Tweets

1. Posting automated Tweets

Automated Tweets that cross-post outside information:

You may post automated Tweets based on sources of outside information — such as an RSS feed, weather data, etc. — as long as you are sufficiently authorized to publish such information.

Other automated Tweets (excluding mentions or replies):

Provided you comply with all other rules, you may post automated Tweets for entertainment, informational, or novelty purposes. As a reminder, accounts posting duplicative, spammy, or otherwise prohibited content may be subject to suspension.

2. Posting automated mentions and replies

The reply and mention functions are intended to make communication between Twitter users easier. Automating these actions to reach many users on an unsolicited basis is an abuse of the feature, and is not permitted. For example, sending automated replies to Tweets based on keyword searches alone is not permitted. Spamy or duplicative use of mentions and replies may result in enforcement action, such as the removal of your Tweets from Search or the suspension of your app or account.

However, you may send automated replies or mentions to Twitter users so long as:

in advance of sending the automated reply, the recipient or mentioned user(s) have requested or have clearly indicated an intent on Twitter to be contacted by you (i.e. opted in), for example by replying to a Tweet from your account, or by sending you a Direct Message;

you provide a clear and easy way for such users to opt-out of receiving automated replies and mentions, and promptly honor all such opt-out requests;

you only send one automated reply or mention per user interaction; and

the automated reply or mention is a reply to the user's original Tweet (if your campaign is based on users posting a reply to your Tweet).

Opt-in techniques and indications of user intent take many different forms, depending on the specifics of your use case and implementation. Some examples include:

A Tweet from your account that clearly indicates that a user taking a specific action on that Tweet (such as Retweeting it) will opt the user into receiving an automated response.

A mention of your account by the user in a manner suggesting the user clearly wishes or intends to receive a response. If you want to run an auto-reply campaign with a campaign- or use-case-specific hashtag, users should also mention you in their Tweets.

Note that a user following your account is not on its own a sufficient indication of user intent to receive an automated response.

: Advertisers, publishers, and brands using auto-response campaigns must request approval from Twitter and may be subject to additional rules. Please reach out to your account or partner manager for assistance. If you do not have a partner manager but are interested in managed client services, you can learn more [here](#).

Additionally, we recommend that any accounts that will communicate with

users via automated mentions or replies:

Note

1. Appropriately filter responses based on potentially sensitive language in user handles,

display names, and Tweet text, as well as potentially sensitive media;

2. Check that the Tweet you are mentioning or replying to still exists (for example, using the statuses/lookup endpoint on the Twitter API).

C. Automated Direct Messages

Sending automated Direct Messages to users

You may not send unsolicited Direct Messages in a bulk or automated manner, and should be thoughtful about the frequency with which you contact users via Direct Message.

You may send automated Direct Messages to users so long as:

in advance of sending the Direct Message, the recipient(s) have requested or have clearly indicated an intent on Twitter to be contacted by you via Direct Message, for example by sending you a Direct Message; **and**

you provide a clear and easy way for such users to opt-out of receiving automated Direct Messages, and promptly honor all such opt-out requests.

The fact that a user is technically able to receive a Direct Message from you (e.g. because the user follows you, has enabled the ability to receive Direct Messages from any account, or because the user is in a pre-existing Direct Message conversation with you) does not necessarily mean they have requested or expect to receive automated Direct Messages from you.

Interacting with users via Direct Message

Per the Ground Rules, remember to be thoughtful about the amount and type of information you request or exchange with users via Direct Messages. If you will be asking a user to provide personal or private information via an automated Direct Message, you must clearly explain how you will use the information you're collecting. Consider including a link to your privacy policy in your Direct Message to the user, as well as in your Twitter profile bio.

Don't publicly share information received in a Direct Message conversation with a user without first obtaining explicit consent from the user. For example, if a user asks you via Direct Message about a purchase they made from you, you may not mention the user in a public Tweet that includes information about their purchase unless you have the user's explicit consent to do so.

After a user-initiated interaction ends, don't send additional follow-up Direct Messages or mention users in a Tweet unless you get permission from the user.

D. Automated actions you take on Tweets or accounts

Automated likes: You may not like Tweets in an automated manner.

Automated Retweets: Provided you comply with all other rules, you may Retweet or Quote Tweet in an automated manner for entertainment, informational, or novelty purposes. Automated Retweets often lead to negative user experiences, and bulk, aggressive, or spammy Retweeting is a violation of the [Twitter Rules](#).

Automated following/unfollowing: You may not follow or unfollow Twitter accounts in a bulk, aggressive, or indiscriminate manner. Aggressive following is a violation of the [Twitter Rules](#). Please also review our [following rules and best practices](#) to ensure you are in compliance. Note that applications that claim to get users more followers are also prohibited under the [Twitter Rules](#).

Automated adding to lists or collections: You may not add Twitter users to lists or add Tweets to collections in a bulk or indiscriminate manner. Adding a large number of unrelated users to lists is a violation of the [Twitter Rules](#).

Twitter search rules and restrictions

This article details the rules for using Twitter search, and explains reasons why content may not show up in Twitter search.

Why doesn't all content show up in Twitter search?

In order to keep your search results relevant, Twitter filters search results for quality Tweets and accounts. Material that jeopardizes search quality or creates a bad search experience for other people may be automatically removed from Twitter search. Read more about this [here](#).

General rules for using Twitter search

The [Twitter Rules](#) explain what behaviors are permitted on Twitter. In addition to these rules, we've included some tips below to keep your content relevant (and your Tweets in Twitter search).

Please do not:

- Repeatedly post duplicate or near-duplicate content (links or Tweets).
- Abuse trending topics or hashtags (topic words with a # sign).
- Send automated Tweets or replies.
- Use bots or applications to post similar messages based on keywords.
- Post similar messages over multiple accounts.
- Aggressively follow and unfollow people.

Twitter may automatically remove accounts engaging in these behaviors from search (or even suspend in some cases) in order to ensure the best experience for everyone.

Search results filtering

To learn how to filter your search results, including how to disable or enable safe search, read our [Using Twitter search](#) article.

Search results filtering in third-party applications

Accounts may be filtered from search for posting similar messages over several accounts. Your account may be filtered from search if you are logged into a third-party application that is updating many accounts with similar or duplicate content.

Rules for participating in contests on Twitter

Businesses sometimes host contests and offer prizes to Retweet an update, following a particular account, or posting Tweets with a specific hashtag topic or reply. If you're running a contest, check out this [article](#) for guidelines for contests on Twitter.

If you repeatedly post near-duplicate updates or duplicate links in order to enter contests, you may be filtered from search. Please keep in mind that you are responsible for the updates that you post to your Twitter account.

What to do if you are missing from search

Please see this [article](#) for more information related to this issue.

If you are being filtered from search, rest assured that your followers will still receive your updates. If you post a reply or mention to a specific person, it will be delivered to that person.

Having trouble?

You can learn more about searching by keyword, location, and trending topics in our [Using Twitter search](#) help article. Our [search troubleshooting](#) section includes solutions to common issues related to search.

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. These limitations are set forth in the Twitter Rules below.

The Twitter Rules (along with all incorporated policies), [Privacy Policy](#), and [Terms of Service](#) collectively make up the "Twitter User Agreement" that governs a user's access to and use of Twitter's services.

All individuals accessing or using Twitter's services must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in Twitter taking one or more of the following enforcement actions:

- requiring you to delete prohibited content before you can again create new posts and interact with other Twitter users;
- temporarily limiting your ability to create posts or interact with other Twitter users;
- asking you to verify account ownership with a phone number or email address; or
- permanently suspending your account(s).

If you attempt to evade a permanent suspension by creating new accounts, we will suspend your new accounts.

Please note that we may need to change these Rules from time to time and reserve the right to do so. The most current version is always available at: <https://twitter.com/rules>.

The policies set forth in these Twitter Rules govern organic content on our platform. To learn more about the rules which govern ads and promoted content, please review our [Ads policies](#).

Content Boundaries and Use of Twitter

Intellectual property

Trademark: We reserve the right to suspend accounts or take other appropriate action when someone's brand or trademark, including business name and/or logo, is used in a manner that may mislead or confuse others about your brand affiliation. Read more about our [trademark policy](#) and [how to report a violation](#).

Copyright: We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in our [Terms of Service](#). Read more about our [copyright policy](#).

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Learn more about our [media policy](#).

Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images. Additionally, Twitter may sometimes require you to remove excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about [how to make such a request](#), and [how to mark your media as sensitive](#).

Unlawful use

You may not use our service for any unlawful purposes or in furtherance of illegal activities. By using Twitter, you agree to comply with all applicable laws governing your online conduct and content.

Trends

At times, we may prevent certain content from trending. This includes content that violates the Twitter Rules, as well as content that may attempt to manipulate trends. Read more about [what we allow and do not allow to trend](#).

Misuse of Twitter badges

You may not use badges, including but not limited to the “promoted” or “verified” Twitter badges, unless provided by Twitter. Accounts using unauthorized badges as part of their profile photos, header photos, display names, or in any way that falsely implies affiliation with Twitter or authorization from Twitter to display these badges, may be suspended.

Misuse of usernames

Selling usernames: You may not buy or sell Twitter usernames.

Username squatting: You may not engage in username squatting. Some of the factors we take into consideration when determining whether conduct is username squatting include:

- the number of accounts created;
- the creation of accounts for the purpose of preventing others from using those account names;
- the creation of accounts for the purpose of selling those accounts; and
- the use of third-party content feeds to update and maintain accounts under the names of those third parties.

Please note that Twitter may also remove accounts that are inactive for more than six months. Learn more about [username squatting](#).

Abusive Behavior

We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we prohibit behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Context matters when evaluating for abusive behavior and determining appropriate enforcement actions. Factors we may take into consideration include, but are not limited to whether:

- the behavior is targeted at an individual or group of people;
- the report has been filed by the target of the abuse or a bystander;
- the behavior is newsworthy and in the legitimate public interest.

Violence and physical harm

Violence: You may not make specific threats of violence or wish for the serious physical harm, death, or disease of an individual or group of people. This includes, but is not limited to, threatening or promoting terrorism. You also may not affiliate with organizations that — whether by their own statements or activity both on and off the platform — use or promote violence against civilians to further their causes. We will begin enforcing this rule around affiliation with such organizations on December 18, 2017.

Suicide or self-harm: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Child sexual exploitation: You may not promote child sexual exploitation. Learn more about our zero-tolerance [child sexual exploitation policy](#).

Abuse and hateful conduct

Abuse: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Unwanted sexual advances: You may not direct abuse at someone by sending unwanted sexual content, objectifying them in a sexually explicit manner, or otherwise engaging in sexual misconduct.

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. Read more about our [hateful conduct policy](#).

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category. We will begin enforcing this rule on December 18, 2017.

Private information and intimate media

Private information: You may not publish or post other people's private information without their express authorization and permission. Definitions of private information may vary depending on local laws. Read more about our [private information policy](#).

Intimate media: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. Read more about intimate media on Twitter.

Threats to expose / hack: You may not threaten to expose someone's private information or intimate media. You also may not threaten to hack or break into someone's digital information.

Impersonation

You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. While you may maintain parody, fan, commentary, or newsfeed accounts, you may not do so if the intent of the account is to engage in spamming or abusive behavior. Read more about our [impersonation policy](#).

Spam and Security

We strive to protect people on Twitter from technical abuse and spam.

To promote a stable and secure environment on Twitter, you may not do, or attempt to do, any of the following while accessing or using Twitter:

- Access, tamper with, or use non-public areas of Twitter, Twitter's computer systems, or the technical delivery systems of Twitter's providers (except as expressly permitted by the Twitter Bug Bounty program).
- Probe, scan, or test the vulnerability of any system or network, or breach or circumvent any security or authentication measures (except as expressly permitted by the Twitter Bug Bounty program).
- Access or search, or attempt to access or search, Twitter by any means (automated or otherwise) other than through our currently available, published interfaces that are provided by Twitter (and only pursuant to the applicable terms and conditions), unless you have been specifically allowed to do so in a separate agreement with Twitter. Note that crawling Twitter is permissible if done in accordance with the provisions of the robots.txt file; however, scraping Twitter without our prior consent is expressly prohibited.
- Forge any TCP/IP packet header or any part of the header information in any email or posting, or in any way use Twitter to send altered, deceptive, or false source-identifying information.
- Interfere with or disrupt the access of any user, host or network, including, without limitation, sending a virus, overloading, flooding, spamming, mail-bombing Twitter's services, or by scripting the creation of content in such a manner as to interfere with or create an undue burden on Twitter.

Any accounts engaging in the following activities may be temporarily locked or subject to permanent suspension:

- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use Twitter's services for the purpose of spamming anyone. Spam is generally defined on Twitter as bulk or aggressive activity that attempts to manipulate or disrupt Twitter or the experience of users on Twitter to drive traffic or attention to unrelated accounts, products, services, or initiatives. Some of the factors that we take into account when determining what conduct is considered to be spamming include:
 - if you have followed and/or unfollowed a large number of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if your Tweets or Direct Messages consist mainly of links shared without commentary;
 - if a large number of people have blocked you in response to high volumes of untargeted, unsolicited, or duplicative content or engagements from your account;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicative or substantially similar content, replies, or mentions over multiple accounts or multiple duplicate updates on one account, or create duplicate or substantially similar accounts;
 - if you post multiple updates to a trending or popular topic with an intent to subvert or manipulate the topic to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you send large numbers of unsolicited replies or mentions;
 - if you add users to lists in a bulk or aggressive manner;
 - if you are randomly or aggressively engaging with Tweets (e.g., likes, Retweets, etc.) or users (e.g., following, adding to lists or Moments, etc.) to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you repeatedly post other people's account information as your own (e.g., bio, Tweets, profile URL, etc.);
 - if you post misleading, deceptive, or malicious links (e.g., affiliate links, links to malware/clickjacking pages, etc.);
 - if you create fake accounts, account interactions, or impressions;
 - if you sell, purchase, or attempt to artificially inflate account interactions (such as followers, Retweets, likes, etc.); and
 - if you use or promote third-party services or apps that claim to get you more followers, Retweets, or likes (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers or engagements to your account or Tweets).

Please see our support articles on [following rules and best practices](#) and [automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts may be permanently suspended.

Content Visibility

Accounts under investigation or which have been detected as sharing content in violation of these Rules may have their account or Tweet visibility limited in various parts of Twitter, including search. To learn more about situations in which content may be restricted on Twitter, please see our support article on [search rules and restrictions](#).

Trademark policy

Updated June 30, 2016

What is a Trademark Policy Violation on Twitter?

Using another's trademark in a manner that may mislead or confuse others about your brand affiliation may be a violation of our trademark policy.

How Does Twitter Respond to Trademark Reports?

We review each report and do the following:

- If we determine that there is a violation of our trademark policy, we may suspend the account. In some instances, we may give the account holder an opportunity to comply with our policies before suspending the account. For more information, see our [Parody, commentary, and fan account policy](#).
- We may also provide the account holder with your name and information included in the copy of the report.
- If you would like to report a trademark violation in a [Twitter Ad](#), please visit our [Trademark Policy for Promoted Ads](#) to learn more.

Please note that Twitter only investigates requests that are submitted by the trademark holder or by an authorized representative of the trademark holder.

What is not a Trademark Policy Violation?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples include:

- Using a trademark in a way that is outside the scope of the trademark registration (e.g. territory, or goods and services identified in the registration).
- Nominative and other fair uses of trademarks are protected uses under our trademark policy, so long as the account is clearly distinguished from the trademark owner. This includes use by resellers in certain regions and accounts engaging in parody, commentary, or news. For more information, see our [Parody, commentary, and fan account policy](#).

How do I Report a Trademark Policy Violation?

Submit trademark reports through our [Trademark Issue support form](#). Using the appropriate form ensures we have all the information necessary to process your report quickly. Please provide all information requested in the form. If you submit a report with incomplete information, we will follow up with you to request that information. Please note that this will result in a delay in processing your report.

If you have trouble using the form, try updating your browser or switching to another browser.

Copyright policy

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged

infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notification.

When Should I File a Counter-notice?

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);

3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Twitter media policy

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to easily share your content with others. Below are our guidelines for making sure your content is labeled appropriately.

Media on Twitter

You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.

In addition, you may not feature graphic content (such as media containing pornography or excessive violence) in live video, or in your profile image or header image.

If you see media you think may be in violation of our policy, see [this article](#) for information about how to report it.

Your media settings

All Twitter accounts have media settings to help Twitter understand the type of media you are posting and to help us assist others to find the type of content that they're interested in.

By appropriately marking your media settings, Twitter can identify potentially sensitive content that other users may not wish to see, such as violence or nudity. If you intend to regularly post such content, we request that you adjust your media setting:

1. Log in to your account on the web at twitter.com.
2. Click your **profile icon** in the top right navigation bar.
3. Select **Settings and privacy** from the drop-down menu.
4. Go to your [Privacy and safety](#) settings.
5. Look for the **Tweet media** section and check the box next to **Mark media you Tweet as containing material that may be sensitive**. **Note:** Enabling this setting will result in others seeing a message when they view your profile, alerting them of potentially sensitive media and asking them to confirm they still want to view it.

You can control your preferences for how and when you may see sensitive content in your [media settings](#).

What if I don't change my media settings and post sensitive media?

If one of your Tweets containing media is reported as sensitive, it will be sent to Twitter for review. If we find that the media wasn't marked as being potentially sensitive at the time of the upload:

- We label the media as being potentially sensitive, or in the case of live video, remove the content entirely.
- We may also change your account setting to **Mark media you Tweet as containing material that may be sensitive** so that future posts are marked accordingly.

You will receive an email letting you know that we've labeled your media and if we changed your media setting. You will be able to change your media settings on your account settings page, though this image and any other images that have been labeled will remain marked as potentially sensitive.

If you repeatedly upload media that are mislabeled, the Twitter team may:

- Label your media as having potentially sensitive content.
- Permanently adjust your account setting to **Mark media you Tweet as containing material that may be sensitive** so that your future posts have a warning message users must click through before viewing your media.

Twitter may also use automated techniques to detect and label potentially sensitive media.

Additionally, people who visit your profile may see a message letting them know the account may include potentially sensitive content and asking them to confirm they still want to view it. People who have opted in to see possibly sensitive content will still see your account without the message.

What happens if my profile image is in violation of the Twitter Rules?

If your profile or header image contains graphic content in violation of the Twitter Rules, we may suspend your account and remove the offending content. Repeat violations will result in a permanent suspension.

Additionally, your account may be locked if media in your profile is in violation of Twitter's media policy. In order to unlock your account, you will need to follow our instructions and delete the media in violation. If your account has been locked, see this [article](#) for more information.

Can I appeal a decision made by the Twitter team?

If you are suspended for having a pornographic profile, header, or background image, you can appeal by logging in to your account and submitting an appeal [here](#).

There is currently no way to appeal a decision by the Twitter team that permanently changes your account setting to **Mark media you Tweet as containing material that may be sensitive** in response to repeated mislabeling of sensitive content.

How to contact Twitter about media concerning a deceased family member

In order to respect the wishes of loved ones, Twitter will remove imagery of deceased individuals in certain circumstances. Immediate family members and other authorized individuals may request the removal of images or video of deceased individuals, from when critical injury occurs to the moments before or after death, by sending a request to Twitter Inc. via our [privacy form](#). When reviewing such media removal requests, Twitter considers public interest factors such as the newsworthiness of the content and may not be able to honor every request. Note: Find information about contacting Twitter to [deactivate the account of a deceased or incapacitated user](#).

How to mark media as sensitive links to Media policy

FAQs about trends on Twitter

Where can I find trends?

Trends are available on the Twitter for iOS app, the Twitter for Android app, and [twitter.com](#).

On Twitter's mobile apps, you can find trends listed under the trends section



of the **Explore** tab. When signed in to [twitter.com](#) on a desktop or laptop computer, trends are listed in many places, including the Home timeline, Notifications, search results, and profile pages.

How are trends determined?

Trends are determined by an algorithm and, by default, are tailored for you based on who you follow, your interests, and your location. This algorithm identifies topics that are popular now, rather

than topics that have been popular for a while or on a daily basis, to help you discover the hottest emerging topics of discussion on Twitter.

Note: The number of Tweets that are related to the trends is just one of the factors the algorithm looks at when ranking and determining trends. Algorithmically, trends and hashtags are grouped together if they are related to the same topic. For instance, #MondayMotivation and #MotivationMonday may both be represented by #MondayMotivation.

You can choose to see trends that are not tailored for you by selecting a [specific trends location](#) on twitter.com, iOS, or Android. Location trends identify popular topics among people in a specific geographic location.

What does the # sign mean?

You may notice that some trends have # sign before the word or phrase. This is called a [hashtag](#) and is included specifically in Tweets to mark them as relating to a topic, so that people can follow the conversation in search.

Will I see world and local events in my trends?

Yes, when signed into your account on the twitter.com, iOS, or Android, trends will be tailored for you based on your location and who you follow. There will be many world and local news events and conversations that will appear in your trends regardless of your personalization.

Can I see trends for a specific location?

Yes, you can view trends for a specific location on twitter.com, iOS, or Android.

View trends via twitter.com:

To view trends for a specific location, click **Change** in your trends box. You can choose to **Keep tailored trends** based on your location and who you follow, or you can select **Change** to choose a nearby location or search locations.

Adjust your trends settings via Twitter for iOS:

1. In the top menu, tap your **profile** icon.
2. Tap **Settings and privacy**, then tap **Content preferences**.
3. Under **Explore**, tap **Trends**.
4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.
5. When the feature is disabled, you can change your location by tapping **Change location**.
 1. In the search box, type your desired trends location or choose from the list.

Adjust your trends settings via Twitter for Android:

1. In the top menu, you will either see a **navigation menu** icon  or your **profile** icon. Tap whichever icon you have, and select **Settings and privacy**.

2. Tap **Content preferences**.
3. Under **Explore**, tap **Trends**.
4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.
5. When the feature is disabled, you can change your location by tapping **Change location**.
 1. In the search box, type your desired trends location or choose from the list.

Note: If you don't find your preferred country or city, it means we're not receiving enough Tweets from that geographical area to create a list. You can look up local Tweets on any topic by using [advanced Twitter Search operators](#).

What happens when I click or tap on a trend?

Clicking or tapping any of the trends takes you to the Twitter search results for that trend. You'll see all Tweets including that phrase or hashtag. To see what people are saying about a previous trend, perform a [search for that keyword](#).

How can I participate in a trend?

Simply post a Tweet including the exact word or phrase as it appears in the trends list (with the hashtag, if you see one). Due to the large number of people Tweeting about these specific trends, you may not always be able to find your particular Tweet in search, but your followers will always see your Tweets.

Note: Twitter also filters searches for quality. To understand what types of behavior could cause your Tweets to be filtered out, read about our [search rules and restrictions](#).

Are there rules for trends?

Yes. We want trends to promote healthy discussions on Twitter. This means that at times, we may prevent certain content from trending. These include trends that:

- Contain profanity or adult/graphic references.
- Incite hate on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease.
- Violate the [Twitter Rules](#).

In some cases, we may also consider the newsworthiness of the content, or if it is in the public interest when evaluating potential violations. In these cases, the content might continue to trend on our platform.

Note that even if we prevent the hashtag or the content from appearing on the list of trending topics, you may still be able to access conversations around that content on Twitter.

Trends descriptions

Below some trends, you may see the approximate number of Tweets associated with the trend, or personalized context such as who in your network is Tweeting about it.

Additionally, when a Moment is found to be relevant to a trend, we may use the title of the Moment as the trend's description.

Having trouble?

Check out our [troubleshooting section](#) for solutions to common problems.

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Child sexual exploitation policy

Twitter removes content promoting child sexual exploitation

We do not tolerate child sexual exploitation on Twitter. When we are made aware of links to images of or content promoting child sexual exploitation they will be removed from the site without further notice and reported to The National Center for Missing & Exploited Children ("NCMEC"); we permanently suspend accounts promoting or containing updates with links to child sexual exploitation.

Reporting content promoting child sexual exploitation

If you think you've found a Twitter account distributing or promoting child sexual exploitation while using Twitter, please click [here](#) to report it. Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. To find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with the National Center for Missing and Exploited Children [here](#), or your local law enforcement. If you believe that the content was hosted on a website outside of the United States, you can report it [here](#) on the International Association of Internet Hotlines website.

NOTE: Please do NOT tweet, retweet or repost child sexual exploitation for any reason. Click [here](#) to report it to us immediately and we will take steps to remove it.

Law Enforcement Agencies

If you are a member of a law enforcement agency seeking information regarding Twitter accounts promoting or distributing child sexual exploitation, please see our [Guidelines for Law Enforcement](#) for information on our process; we also suggest you contact the National Center for Missing and Exploited Children (NCMEC) [here](#), regardless of jurisdiction, to determine whether or not information on those accounts is already on file.

For general law enforcement inquiries, please see our [Guidelines for Law Enforcement](#) for more information.

Hateful conduct policy

Freedom of expression means little if voices are silenced because people are afraid to speak up. We do not tolerate behavior that harasses, intimidates, or uses fear to silence another person's voice. If you see something on Twitter that violates these rules, please report it to us.

How our policy works

As explained in the Twitter Rules,

- **Hateful conduct:** You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Examples of what we do not tolerate includes, but is not limited to behavior that harasses individuals or groups of people with:

- violent threats;
- wishes for the physical harm, death, or disease of individuals or groups;
- references to mass murder, violent events, or specific means of violence in which/with which such groups have been the primary targets or victims;
- behavior that incites fear about a protected group;
- repeated and/or or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone.

How our enforcement works

- **Context matters.**
 - Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. While we accept reports of violations

from anyone, sometimes we also need to hear directly from the target to ensure that we have proper context.

- The number of reports we receive does not impact whether or not something will be removed. However, it may help us prioritize the order in which it gets reviewed.

- **We focus on behavior.**

- We enforce policies when someone reports behavior that is abusive and targets an entire protected group and/or individuals who may be members.
- This targeting can happen in any manner (for example, @mentions, tagging a photo, and more).

- **We have a range of enforcement options.**

- The consequences for violating our rules vary depending on the severity of the violation and the person's previous record of violations. For example, we may ask someone to remove the offending Tweet before they can Tweet again. For other cases, we may suspend an account.
-

Report private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Read about our [policy regarding private information posted on Twitter](#).

Do I need to have a Twitter account to make a report?

No, you do not need to have a Twitter account to file a report that your private and confidential information has been posted on Twitter.

Can I report unauthorized posting of Copyrighted material?

If your copyrighted material has been Tweeted without your permission, please review our [Copyright and DMCA policy](#).

What can I do if my information is posted on other websites?

We cannot take down any information that is posted on a website other than Twitter. If your information is posted on another website, please contact the help center or support staff for the website where your information is posted.

How do I report that my private information is posted on Twitter?

You can report that your private information has been posted on Twitter directly from the offending Tweet(s). Learn [how to report individual Tweets](#).

File a [report that someone has posted your private information](#).

What happens when Twitter receives a valid report?

When we receive a complete and valid report that private information has been posted on Twitter, we'll investigate the account and Tweets reported. We will review where, if anywhere, the information has been made publicly available

before taking action on the account or Tweets. If the information you reported was previously posted elsewhere on the Internet, it is not a violation of our policy and we will not take action.

Learn more about [protecting your private information on Twitter and other websites](#).

Read our [policy regarding private information posted on Twitter](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter impersonation policy.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to Twitter's [parody, commentary, and fan account policy](#) for more information about these types of accounts.

Find more information about [reporting impersonation on Twitter](#).

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation rules

Updated November 3, 2017

This page is primarily intended for **developers**.

For Twitter users: You are ultimately responsible for the actions taken with your account, or by applications associated with your account. Before authorizing a third-party application to access or use your account, make sure you've thoroughly investigated the application and understand what it will do. If automated activity on your account violates the [Twitter Rules](#) or these automation rules, Twitter may take action on your account, including [filtering your Tweets from search results](#) or suspending your account.

For more information on third-party applications, please see our article on [connecting and revoking third-party applications](#).

If you are a developer and have questions about these Automation Rules, please reach out via <https://twittercommunity.com/c/rules-and-policies>.

If your application has been suspended or restricted from performing write actions, you may submit an appeal at </forms/platform>.

I. Ground Rules Do!

Build solutions that automatically broadcast helpful information in Tweets.

Run creative campaigns that auto-reply to users who engage with your content.

Build solutions that automatically respond to users in Direct Messages.

Try new things that help people (and comply with our rules).

Make sure your application provides a good user experience and performs well — and confirm that remains the case over time.

Don't!

Violate these or other policies. Be extra mindful of our rules about abuse and user privacy.

Abuse the Twitter API or attempt to circumvent rate limits.

Use non-API-based forms of automation, such as scripting the Twitter website. The use of these techniques may result in the permanent suspension of your account.

Spam or bother users, or otherwise send them unsolicited messages.

A. The Twitter Rules and the Developer Agreement and Policy

As with all activity on Twitter, automated activity is subject to the [Twitter Rules](#) and, if you're a developer using the Twitter API, the [Developer Agreement and Policy](#).

You should carefully review these policies to ensure that your automated activity is compliant. Automated applications or activities that violate these policies, or that facilitate or induce users to violate them, may be subject to enforcement action, potentially including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these policies.

Although all aspects of the Twitter Rules and the Developer Agreement and Policy apply to automated activity, you should keep the following rules top of mind:

Spamming: You may not send automated Tweets or Direct Messages that are spam, or otherwise engage in spamming activity. Some examples of spammy behavior to avoid with automation include:

Trending topics: You may not automatically post about trending topics on Twitter, or use automation to attempt to influence or manipulate trending topics.

Multiple posts/accounts: You may not post duplicative or substantially similar Tweets on one account or over multiple accounts you operate.

Duplicate accounts: You may not create and/or automate multiple accounts for duplicative or substantially similar use cases.

However, automating multiple accounts for related but non-duplicative use cases is permitted. For example, you may automate separate accounts to Tweet when the Hubble Space Telescope passes over different cities, such as [San Francisco](#) or [Hong Kong](#).

Misleading links: You may not send automated Tweets or Direct Messages containing links that are misleading, including links that maliciously or deceptively redirect through landing pages or ad pages before displaying the final content.

Sensitive media: Automated Tweets and Direct Messages must comply with the [Twitter media policy](#), and you should mark your account as potentially sensitive if you intend to post graphic, pornographic, or potentially sensitive media.

Abusive behavior: You may not engage in any automated activity that encourages, promotes, or incites abuse, violence, hateful conduct, or harassment, on or off Twitter.

Private information: You may not post private or confidential information about a person without their prior express authorization.

B. Other Ground Rules for Automated Activity

In addition to the policies above, the following ground rules apply to all automated activity on Twitter:

Don't surprise or mislead users: Automated activity should honor users' expectations. Ask for the user's permission before taking an action if you aren't sure.

Mature content or profanity: Don't Direct Message, mention, or reply to users with potentially sensitive content (including profanity), unless they've clearly indicated an intent to receive it in advance.

Be thoughtful about the information you request or exchange on Twitter

Tweets: Don't ask users to send you personal or private information via a public Tweet. If you need additional personal or private information from a user to provide them with customer service (or other similar use cases), you should ask the user to share such information by Direct Message or another private channel. You might even consider adding a [Direct Message deep link](#) to your Tweet.

Direct Messages: You should only ask users for the minimum amount of information you need to provide them with service. If you need to request or exchange particularly sensitive information (such as credit card information), you should consider directing users to your website or other appropriate channel to do so.

II. Activity-Specific Rules

The activity-specific rules in this section apply to taking specific automated actions on Twitter. Please read these rules carefully, as they outline both permitted and prohibited use cases of automation.

Automated applications or activities that violate these rules, or that facilitate or induce users to violate them, may be subject to enforcement action, including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access the Twitter API based on violations of these rules. As a reminder, you should also carefully review the spam guidelines in the [Twitter Rules](#) to avoid having activities performed by you, your app, or other users through your app or service flagged as spam.

A. Automated Actions Through Another User's Account

Twitter users may authorize your app or service to [access their Twitter account through OAuth](#). A user authorizing your app or service to access their Twitter account through OAuth does not by itself constitute sufficient consent to take automated actions through that user's account.

You may only take automated actions through another Twitter user's account if you:

clearly describe to the user the types of automated actions that will occur; receive express consent from the user to take those automated actions; and immediately honor a user's request to opt-out of further automated actions.

If you substantially change the purpose or functionality of your app or service, you must re-obtain express consent from the user to take automated action through their account before doing so.

These requirements apply to any automated action taken through another Twitter user's account, including posting Tweets, sending Direct Messages, deleting Tweets or Direct Messages, or following/unfollowing other accounts. For applications that offer users the ability to delete Tweets in a

bulk or automated manner, you must also clearly state that Tweets are not recoverable once deleted.

B. Automated Tweets

1. Posting automated Tweets

Automated Tweets that cross-post outside information:

You may post automated Tweets based on sources of outside information — such as an RSS feed, weather data, etc. — as long as you are sufficiently authorized to publish such information.

Other automated Tweets (excluding mentions or replies):

Provided you comply with all other rules, you may post automated Tweets for entertainment, informational, or novelty purposes. As a reminder, accounts posting duplicative, spammy, or otherwise prohibited content may be subject to suspension.

2. Posting automated mentions and replies

The reply and mention functions are intended to make communication between Twitter users easier. Automating these actions to reach many users on an unsolicited basis is an abuse of the feature, and is not permitted. For example, sending automated replies to Tweets based on keyword searches alone is not permitted. Spamy or duplicative use of mentions and replies may result in enforcement action, such as the removal of your Tweets from Search or the suspension of your app or account.

However, you may send automated replies or mentions to Twitter users so long as:

in advance of sending the automated reply, the recipient or mentioned user(s) have requested or have clearly indicated an intent on Twitter to be contacted by you (i.e. opted in), for example by replying to a Tweet from your account, or by sending you a Direct Message;

you provide a clear and easy way for such users to opt-out of receiving automated replies and mentions, and promptly honor all such opt-out requests;

you only send one automated reply or mention per user interaction; and

the automated reply or mention is a reply to the user's original Tweet (if your campaign is based on users posting a reply to your Tweet).

Opt-in techniques and indications of user intent take many different forms, depending on the specifics of your use case and implementation. Some examples include:

A Tweet from your account that clearly indicates that a user taking a specific action on that Tweet (such as Retweeting it) will opt the user into receiving an automated response.

A mention of your account by the user in a manner suggesting the user clearly wishes or intends to receive a response. If you want to run an auto-reply campaign with a campaign- or use-case-specific hashtag, users should also mention you in their Tweets.

Note that a user following your account is not on its own a sufficient indication of user intent to receive an automated response.

: Advertisers, publishers, and brands using auto-response campaigns must request approval from Twitter and may be subject to additional rules. Please reach out to your account or partner manager for assistance. If you do not have a partner manager but are interested in managed client services, you can learn more [here](#).

Additionally, we recommend that any accounts that will communicate with

users via automated mentions or replies:

Note

1. Appropriately filter responses based on potentially sensitive language in user handles,

display names, and Tweet text, as well as potentially sensitive media;

2. Check that the Tweet you are mentioning or replying to still exists (for example, using the statuses/lookup endpoint on the Twitter API).

C. Automated Direct Messages

Sending automated Direct Messages to users

You may not send unsolicited Direct Messages in a bulk or automated manner, and should be thoughtful about the frequency with which you contact users via Direct Message.

You may send automated Direct Messages to users so long as:

in advance of sending the Direct Message, the recipient(s) have requested or have clearly indicated an intent on Twitter to be contacted by you via Direct Message, for example by sending you a Direct Message; **and**

you provide a clear and easy way for such users to opt-out of receiving automated Direct Messages, and promptly honor all such opt-out requests.

The fact that a user is technically able to receive a Direct Message from you (e.g. because the user follows you, has enabled the ability to receive Direct Messages from any account, or because the user is in a pre-existing Direct Message conversation with you) does not necessarily mean they have requested or expect to receive automated Direct Messages from you.

Interacting with users via Direct Message

Per the Ground Rules, remember to be thoughtful about the amount and type of information you request or exchange with users via Direct Messages. If you will be asking a user to provide personal or private information via an automated Direct Message, you must clearly explain how you will use the information you're collecting. Consider including a link to your privacy policy in your Direct Message to the user, as well as in your Twitter profile bio.

Don't publicly share information received in a Direct Message conversation with a user without first obtaining explicit consent from the user. For example, if a user asks you via Direct Message about a purchase they made from you, you may not mention the user in a public Tweet that includes information about their purchase unless you have the user's explicit consent to do so.

After a user-initiated interaction ends, don't send additional follow-up Direct Messages or mention users in a Tweet unless you get permission from the user.

D. Automated actions you take on Tweets or accounts

Automated likes: You may not like Tweets in an automated manner.

Automated Retweets: Provided you comply with all other rules, you may Retweet or Quote Tweet in an automated manner for entertainment, informational, or novelty purposes. Automated Retweets often lead to negative user experiences, and bulk, aggressive, or spammy Retweeting is a violation of the [Twitter Rules](#).

Automated following/unfollowing: You may not follow or unfollow Twitter accounts in a bulk, aggressive, or indiscriminate manner. Aggressive following is a violation of the [Twitter Rules](#). Please also review our [following rules and best practices](#) to ensure you are in compliance. Note that applications that claim to get users more followers are also prohibited under the [Twitter Rules](#).

Automated adding to lists or collections: You may not add Twitter users to lists or add Tweets to collections in a bulk or indiscriminate manner. Adding a large number of unrelated users to lists is a violation of the [Twitter Rules](#).

About search rules and restrictions

This article details the rules for using Twitter search, and explains reasons why content may not show up in Twitter search.

Why doesn't all content show up in Twitter search?

In order to keep your search results relevant, Twitter filters search results for quality Tweets and accounts. Material that

jeopardizes search quality or creates a bad search experience for other people may be automatically removed from Twitter search. Read more about this [here](#).

General rules for using Twitter search

The [Twitter Rules](#) explain what behaviors are permitted on Twitter. In addition to these rules, we've included some tips below to keep your content relevant (and your Tweets in Twitter search).

Please do not:

Repeatedly post duplicate or near-duplicate content (links or Tweets). Abuse trending topics or hashtags (topic words with a # sign).

Send automated Tweets or replies.

Use bots or applications to post similar messages based on keywords. Post similar messages over multiple accounts.

Aggressively follow and unfollow people.

Twitter may automatically remove accounts engaging in these behaviors from search (or even suspend in some cases) in order to ensure the best experience for everyone.

Search results filtering

To learn how to filter your search results, including how to disable or enable safe search, read our [Using Twitter search](#) article.

Search results filtering in third-party applications

Accounts may be filtered from search for posting similar messages over several accounts. Your account may be filtered from search if you are logged into a third-party application that is updating many accounts with similar or duplicate content.

Rules for participating in contests on Twitter

Businesses sometimes host contests and offer prizes to Retweet an update, following a particular account, or posting Tweets with a specific hashtag topic or reply. If you're running a

contest, check out this [article](#) for guidelines for contests on Twitter.

If you repeatedly post near-duplicate updates or duplicate links in order to enter contests, you may be filtered from search. Please keep in mind that you are responsible for the updates that you post to your Twitter account.

What to do if you are missing from search

Please see this [article](#) for more information related to this issue.

If you are being filtered from search, rest assured that your followers will still receive your updates. If you post a reply or mention to a specific person, it will be delivered to that person.

Having trouble?

You can learn more about searching by keyword, location, and trending topics in our [Using Twitter search](#) help article.

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. These limitations are set forth in the Twitter Rules below.

The Twitter Rules (along with all incorporated policies), [Privacy Policy](#), and [Terms of Service](#) collectively make up the "Twitter User Agreement" that governs a user's access to and use of Twitter's services.

All individuals accessing or using Twitter's services must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in Twitter taking one or more of the following enforcement actions:

- requiring you to delete prohibited content before you can again create new posts and interact with other Twitter users;
- temporarily limiting your ability to create posts or interact with other Twitter users;
- asking you to verify account ownership with a phone number or email address; or
- permanently suspending your account(s).

If you attempt to evade a permanent suspension by creating new accounts, we will suspend your new accounts.

Please note that we may need to change these Rules from time to time and reserve the right to do so. The most current version is always available at: <https://twitter.com/rules>.

The policies set forth in these Twitter Rules govern organic content on our platform. To learn more about the rules which govern ads and promoted content, please review our [Ads policies](#).

Content Boundaries and Use of Twitter

Intellectual property

Trademark: We reserve the right to suspend accounts or take other appropriate action when someone's brand or trademark, including business name and/or logo, is used in a manner that may mislead or confuse others about your brand affiliation. Read more about our [trademark policy](#) and [how to report a violation](#).

Copyright: We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in our [Terms of Service](#). Read more about our [copyright policy](#).

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Learn more about our [media policy](#).

Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images. Additionally, Twitter may sometimes require you to remove excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about [how to make such a request](#), and [how to mark your media as sensitive](#).

Unlawful use

You may not use our service for any unlawful purposes or in furtherance of illegal activities. By using Twitter, you agree to comply with all applicable laws governing your online conduct and content.

Trends

At times, we may prevent certain content from trending. This includes content that violates the Twitter Rules, as well as content that may attempt to manipulate trends. Read more about [what we allow and do not allow to trend](#).

Misuse of Twitter badges

You may not use badges, including but not limited to the “promoted” or “verified” Twitter badges, unless provided by Twitter. Accounts using unauthorized badges as part of their profile photos, header photos, display names, or in any way that falsely implies affiliation with Twitter or authorization from Twitter to display these badges, may be suspended.

Misuse of usernames

Selling usernames: You may not buy or sell Twitter usernames.

Username squatting: You may not engage in username squatting. Some of the factors we take into consideration when determining whether conduct is username squatting include:

- the number of accounts created;
- the creation of accounts for the purpose of preventing others from using those account names;
- the creation of accounts for the purpose of selling those accounts; and
- the use of third-party content feeds to update and maintain accounts under the names of those third parties.

Please note that Twitter may also remove accounts that are inactive for more than six months. Learn more about [username squatting](#).

Abusive Behavior

We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we prohibit behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Context matters when evaluating for abusive behavior and determining appropriate enforcement actions. Factors we may take into consideration include, but are not limited to whether:

- the behavior is targeted at an individual or group of people;
- the report has been filed by the target of the abuse or a bystander;
- the behavior is newsworthy and in the legitimate public interest.

Violence and physical harm

Violence: You may not make specific threats of violence or wish for the serious physical harm, death, or disease of an individual or group of people. This includes, but is not limited to, threatening or promoting terrorism. You also may not affiliate with organizations that — whether by their own statements or activity both on and off the platform — use or promote violence against civilians to further their causes. We will begin enforcing this rule around affiliation with such organizations on December 18, 2017.

Suicide or self-harm: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Child sexual exploitation: You may not promote child sexual exploitation. Learn more about our zero-tolerance [child sexual exploitation policy](#).

Abuse and hateful conduct

Abuse: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Unwanted sexual advances: You may not direct abuse at someone by sending unwanted sexual content, objectifying them in a sexually explicit manner, or otherwise engaging in sexual misconduct.

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. Read more about our [hateful conduct policy](#).

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category. We will begin enforcing this rule on December 18, 2017.

Private information and intimate media

Private information: You may not publish or post other people's private information without their express authorization and permission. Definitions of private information may vary depending on local laws. Read more about our [private information policy](#).

Intimate media: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. Read more about intimate media on Twitter.

Threats to expose / hack: You may not threaten to expose someone's private information or intimate media. You also may not threaten to hack or break into someone's digital information.

Impersonation

You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. While you may maintain parody, fan, commentary, or newsfeed accounts, you may not do so if the intent of the account is to engage in spamming or abusive behavior. Read more about our [impersonation policy](#).

Spam and Security

We strive to protect people on Twitter from technical abuse and spam.

To promote a stable and secure environment on Twitter, you may not do, or attempt to do, any of the following while accessing or using Twitter:

- Access, tamper with, or use non-public areas of Twitter, Twitter's computer systems, or the technical delivery systems of Twitter's providers (except as expressly permitted by the Twitter Bug Bounty program).
- Probe, scan, or test the vulnerability of any system or network, or breach or circumvent any security or authentication measures (except as expressly permitted by the Twitter Bug Bounty program).
- Access or search, or attempt to access or search, Twitter by any means (automated or otherwise) other than through our currently available, published interfaces that are provided by Twitter (and only pursuant to the applicable terms and conditions), unless you have been specifically allowed to do so in a separate agreement with Twitter. Note that crawling Twitter is permissible if done in accordance with the provisions of the robots.txt file; however, scraping Twitter without our prior consent is expressly prohibited.
- Forge any TCP/IP packet header or any part of the header information in any email or posting, or in any way use Twitter to send altered, deceptive, or false source-identifying information.
- Interfere with or disrupt the access of any user, host or network, including, without limitation, sending a virus, overloading, flooding, spamming, mail-bombing Twitter's services, or by scripting the creation of content in such a manner as to interfere with or create an undue burden on Twitter.

Any accounts engaging in the following activities may be temporarily locked or subject to permanent suspension:

- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use Twitter's services for the purpose of spamming anyone. Spam is generally defined on Twitter as bulk or aggressive activity that attempts to manipulate or disrupt Twitter or the experience of users on Twitter to drive traffic or attention to unrelated accounts, products, services, or initiatives. Some of the factors that we take into account when determining what conduct is considered to be spamming include:
 - if you have followed and/or unfollowed a large number of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if your Tweets or Direct Messages consist mainly of links shared without commentary;
 - if a large number of people have blocked you in response to high volumes of untargeted, unsolicited, or duplicative content or engagements from your account;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicative or substantially similar content, replies, or mentions over multiple accounts or multiple duplicate updates on one account, or create duplicate or substantially similar accounts;
 - if you post multiple updates to a trending or popular topic with an intent to subvert or manipulate the topic to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you send large numbers of unsolicited replies or mentions;
 - if you add users to lists in a bulk or aggressive manner;
 - if you are randomly or aggressively engaging with Tweets (e.g., likes, Retweets, etc.) or users (e.g., following, adding to lists or Moments, etc.) to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you repeatedly post other people's account information as your own (e.g., bio, Tweets, profile URL, etc.);
 - if you post misleading, deceptive, or malicious links (e.g., affiliate links, links to malware/clickjacking pages, etc.);
 - if you create fake accounts, account interactions, or impressions;
 - if you sell, purchase, or attempt to artificially inflate account interactions (such as followers, Retweets, likes, etc.); and
 - if you use or promote third-party services or apps that claim to get you more followers, Retweets, or likes (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers or engagements to your account or Tweets).

Please see our support articles on [following rules and best practices](#) and [automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts may be permanently suspended.

Content Visibility

Accounts under investigation or which have been detected as sharing content in violation of these Rules may have their account or Tweet visibility limited in various parts of Twitter, including search. To learn more about situations in which content may be restricted on Twitter, please see our support article on [search rules and restrictions](#).

Trademark policy

Updated June 30, 2016

What is a Trademark Policy Violation on Twitter?

Using another's trademark in a manner that may mislead or confuse others about your brand affiliation may be a violation of our trademark policy.

How Does Twitter Respond to Trademark Reports?

We review each report and do the following:

- If we determine that there is a violation of our trademark policy, we may suspend the account. In some instances, we may give the account holder an opportunity to comply with our policies before suspending the account. For more information, see our [Parody, commentary, and fan account policy](#).
- We may also provide the account holder with your name and information included in the copy of the report.
- If you would like to report a trademark violation in a [Twitter Ad](#), please visit our [Trademark Policy for Promoted Ads](#) to learn more.

Please note that Twitter only investigates requests that are submitted by the trademark holder or by an authorized representative of the trademark holder.

What is not a Trademark Policy Violation?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples include:

- Using a trademark in a way that is outside the scope of the trademark registration (e.g. territory, or goods and services identified in the registration).
- Nominative and other fair uses of trademarks are protected uses under our trademark policy, so long as the account is clearly distinguished from the trademark owner. This includes use by resellers in certain regions and accounts engaging in parody, commentary, or news. For more information, see our [Parody, commentary, and fan account policy](#).

How do I Report a Trademark Policy Violation?

Submit trademark reports through our [Trademark Issue support form](#). Using the appropriate form ensures we have all the information necessary to process your report quickly. Please provide all information requested in the form. If you submit a report with incomplete information, we will follow up with you to request that information. Please note that this will result in a delay in processing your report.

If you have trouble using the form, try updating your browser or switching to another browser.

Copyright policy

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged

infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notification.

When Should I File a Counter-notice?

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);

3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Twitter media policy

Our goal is to provide a service that allows you to discover and receive content from sources that interest you as well as to easily share your content with others. Below are our guidelines for making sure your content is labeled appropriately.

Media on Twitter

You may not use our service for any unlawful purposes or in furtherance of illegal activities. International users agree to comply with all local laws regarding online conduct and acceptable content.

In addition, you may not feature graphic content (such as media containing pornography or excessive violence) in live video, or in your profile image or header image.

If you see media you think may be in violation of our policy, see [this article](#) for information about how to report it.

Your media settings

All Twitter accounts have media settings to help Twitter understand the type of media you are posting and to help us assist others to find the type of content that they're interested in.

By appropriately marking your media settings, Twitter can identify potentially sensitive content that other users may not wish to see, such as violence or nudity. If you intend to regularly post such content, we request that you adjust your media setting:

1. Log in to your account on the web at twitter.com.
2. Click your **profile icon** in the top right navigation bar.
3. Select **Settings and privacy** from the drop-down menu.
4. Go to your [Privacy and safety](#) settings.
5. Look for the **Tweet media** section and check the box next to **Mark media you Tweet as containing material that may be sensitive**. **Note:** Enabling this setting will result in others seeing a message when they view your profile, alerting them of potentially sensitive media and asking them to confirm they still want to view it.

You can control your preferences for how and when you may see sensitive content in your [media settings](#).

What if I don't change my media settings and post sensitive media?

If one of your Tweets containing media is reported as sensitive, it will be sent to Twitter for review. If we find that the media wasn't marked as being potentially sensitive at the time of the upload:

- We label the media as being potentially sensitive, or in the case of live video, remove the content entirely.
- We may also change your account setting to **Mark media you Tweet as containing material that may be sensitive** so that future posts are marked accordingly.

You will receive an email letting you know that we've labeled your media and if we changed your media setting. You will be able to change your media settings on your account settings page, though this image and any other images that have been labeled will remain marked as potentially sensitive.

If you repeatedly upload media that are mislabeled, the Twitter team may:

- Label your media as having potentially sensitive content.
- Permanently adjust your account setting to **Mark media you Tweet as containing material that may be sensitive** so that your future posts have a warning message users must click through before viewing your media.

Twitter may also use automated techniques to detect and label potentially sensitive media.

Additionally, people who visit your profile may see a message letting them know the account may include potentially sensitive content and asking them to confirm they still want to view it. People who have opted in to see possibly sensitive content will still see your account without the message.

What happens if my profile image is in violation of the Twitter Rules?

If your profile or header image contains graphic content in violation of the Twitter Rules, we may suspend your account and remove the offending content. Repeat violations will result in a permanent suspension.

Additionally, your account may be locked if media in your profile is in violation of Twitter's media policy. In order to unlock your account, you will need to follow our instructions and delete the media in violation. If your account has been locked, see this [article](#) for more information.

Can I appeal a decision made by the Twitter team?

If you are suspended for having a pornographic profile, header, or background image, you can appeal by logging in to your account and submitting an appeal [here](#).

There is currently no way to appeal a decision by the Twitter team that permanently changes your account setting to **Mark media you Tweet as containing material that may be sensitive** in response to repeated mislabeling of sensitive content.

How to contact Twitter about media concerning a deceased family member

In order to respect the wishes of loved ones, Twitter will remove imagery of deceased individuals in certain circumstances. Immediate family members and other authorized individuals may request the removal of images or video of deceased individuals, from when critical injury occurs to the moments before or after death, by sending a request to Twitter Inc. via our [privacy form](#). When reviewing such media removal requests, Twitter considers public interest factors such as the newsworthiness of the content and may not be able to honor every request. Note: Find information about contacting Twitter to [deactivate the account of a deceased or incapacitated user](#).

How to mark media as sensitive links to Media policy

FAQs about trends on Twitter

Where can I find trends?

Trends are available on the Twitter for iOS app, the Twitter for Android app, and [twitter.com](#).

On Twitter's mobile apps, you can find trends listed under the trends section



of the **Explore** tab. When signed in to [twitter.com](#) on a desktop or laptop computer, trends are listed in many places, including the Home timeline, Notifications, search results, and profile pages.

How are trends determined?

Trends are determined by an algorithm and, by default, are tailored for you based on who you follow, your interests, and your location. This algorithm identifies topics that are popular now, rather

than topics that have been popular for a while or on a daily basis, to help you discover the hottest emerging topics of discussion on Twitter.

Note: The number of Tweets that are related to the trends is just one of the factors the algorithm looks at when ranking and determining trends. Algorithmically, trends and hashtags are grouped together if they are related to the same topic. For instance, #MondayMotivation and #MotivationMonday may both be represented by #MondayMotivation.

You can choose to see trends that are not tailored for you by selecting a [specific trends location](#) on twitter.com, iOS, or Android. Location trends identify popular topics among people in a specific geographic location.

What does the # sign mean?

You may notice that some trends have # sign before the word or phrase. This is called a [hashtag](#) and is included specifically in Tweets to mark them as relating to a topic, so that people can follow the conversation in search.

Will I see world and local events in my trends?

Yes, when signed into your account on the twitter.com, iOS, or Android, trends will be tailored for you based on your location and who you follow. There will be many world and local news events and conversations that will appear in your trends regardless of your personalization.

Can I see trends for a specific location?

Yes, you can view trends for a specific location on twitter.com, iOS, or Android.

View trends via twitter.com:

To view trends for a specific location, click **Change** in your trends box. You can choose to **Keep tailored trends** based on your location and who you follow, or you can select **Change** to choose a nearby location or search locations.

Adjust your trends settings via Twitter for iOS:

1. In the top menu, tap your **profile** icon.
2. Tap **Settings and privacy**, then tap **Content preferences**.
3. Under **Explore**, tap **Trends**.
4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.
5. When the feature is disabled, you can change your location by tapping **Change location**.
 1. In the search box, type your desired trends location or choose from the list.

Adjust your trends settings via Twitter for Android:

1. In the top menu, you will either see a **navigation menu** icon  or your **profile** icon. Tap whichever icon you have, and select **Settings and privacy**.

2. Tap **Content preferences**.
3. Under **Explore**, tap **Trends**.
4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.
5. When the feature is disabled, you can change your location by tapping **Change location**.
 1. In the search box, type your desired trends location or choose from the list.

Note: If you don't find your preferred country or city, it means we're not receiving enough Tweets from that geographical area to create a list. You can look up local Tweets on any topic by using [advanced Twitter Search operators](#).

What happens when I click or tap on a trend?

Clicking or tapping any of the trends takes you to the Twitter search results for that trend. You'll see all Tweets including that phrase or hashtag. To see what people are saying about a previous trend, perform a [search for that keyword](#).

How can I participate in a trend?

Simply post a Tweet including the exact word or phrase as it appears in the trends list (with the hashtag, if you see one). Due to the large number of people Tweeting about these specific trends, you may not always be able to find your particular Tweet in search, but your followers will always see your Tweets.

Note: Twitter also filters searches for quality. To understand what types of behavior could cause your Tweets to be filtered out, read about our [search rules and restrictions](#).

Are there rules for trends?

Yes. We want trends to promote healthy discussions on Twitter. This means that at times, we may prevent certain content from trending. These include trends that:

- Contain profanity or adult/graphic references.
- Incite hate on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease.
- Violate the [Twitter Rules](#).

In some cases, we may also consider the newsworthiness of the content, or if it is in the public interest when evaluating potential violations. In these cases, the content might continue to trend on our platform.

Note that even if we prevent the hashtag or the content from appearing on the list of trending topics, you may still be able to access conversations around that content on Twitter.

Trends descriptions

Below some trends, you may see the approximate number of Tweets associated with the trend, or personalized context such as who in your network is Tweeting about it.

Additionally, when a Moment is found to be relevant to a trend, we may use the title of the Moment as the trend's description.

Having trouble?

Check out our [troubleshooting section](#) for solutions to common problems.

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Child sexual exploitation policy

Twitter removes content promoting child sexual exploitation

We do not tolerate child sexual exploitation on Twitter. When we are made aware of links to images of or content promoting child sexual exploitation they will be removed from the site without further notice and reported to The National Center for Missing & Exploited Children ("NCMEC"); we permanently suspend accounts promoting or containing updates with links to child sexual exploitation.

Reporting content promoting child sexual exploitation

If you think you've found a Twitter account distributing or promoting child sexual exploitation while using Twitter, please click [here](#) to report it. Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. To find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with the National Center for Missing and Exploited Children [here](#), or your local law enforcement. If you believe that the content was hosted on a website outside of the United States, you can report it [here](#) on the International Association of Internet Hotlines website.

NOTE: Please do NOT tweet, retweet or repost child sexual exploitation for any reason. Click [here](#) to report it to us immediately and we will take steps to remove it.

Law Enforcement Agencies

If you are a member of a law enforcement agency seeking information regarding Twitter accounts promoting or distributing child sexual exploitation, please see our [Guidelines for Law Enforcement](#) for information on our process; we also suggest you contact the National Center for Missing and Exploited Children (NCMEC) [here](#), regardless of jurisdiction, to determine whether or not information on those accounts is already on file.

For general law enforcement inquiries, please see our [Guidelines for Law Enforcement](#) for more information.

Hateful conduct policy

Freedom of expression means little if voices are silenced because people are afraid to speak up. We do not tolerate behavior that harasses, intimidates, or uses fear to silence another person's voice. If you see something on Twitter that violates these rules, please report it to us.

How our policy works

As explained in the Twitter Rules,

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Examples of what we do not tolerate includes, but is not limited to behavior that harasses individuals or groups of people with:

violent threats;
wishes for the physical harm, death, or disease of individuals or groups;

references to mass murder, violent events, or specific means of violence in which/with which such groups have been the primary targets or victims;

behavior that incites fear about a protected group;

repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone.

How our enforcement works Context matters.

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. While we accept reports of violations from anyone, sometimes we also need to hear directly from the target to ensure that we have proper context.

The number of reports we receive does not impact whether or not something will be removed. However, it may help us prioritize the order in which it gets reviewed.

We focus on behavior.

We enforce policies when someone reports behavior that is abusive and targets an entire protected group and/or individuals who may be members.

This targeting can happen in any manner (for example, @mentions, tagging a photo, and more).

We have a range of enforcement options.

The consequences for violating our rules vary depending on the severity of the violation and the person's previous record of violations. For example, we may ask someone to remove the offending Tweet before they can Tweet again. For other cases, we may suspend an account.

Report private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Read about our [policy regarding private information posted on Twitter](#).

Do I need to have a Twitter account to make a report?

No, you do not need to have a Twitter account to file a report that your private and confidential information has been posted on Twitter.

Can I report unauthorized posting of Copyrighted material?

If your copyrighted material has been Tweeted without your permission, please review our [Copyright and DMCA policy](#).

What can I do if my information is posted on other websites?

We cannot take down any information that is posted on a website other than Twitter. If your information is posted on another website, please contact the help center or support staff for the website where your information is posted.

How do I report that my private information is posted on Twitter?

You can report that your private information has been posted on Twitter directly from the offending Tweet(s). Learn [how to report individual Tweets](#).

File a [report that someone has posted your private information](#).

What happens when Twitter receives a valid report?

When we receive a complete and valid report that private information has been posted on Twitter, we'll investigate the account and Tweets reported. We will review where, if anywhere, the information has been made publicly available before taking action on the account or Tweets. If the information you reported was previously posted elsewhere on the Internet, it is not a violation of our policy and we will not take action.

Learn more about [protecting your private information on Twitter and other websites](#).

Read our [policy regarding private information posted on Twitter](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter impersonation policy.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to Twitter's [parody, commentary, and fan account policy](#) for more information about these types of accounts.

Find more information about [reporting impersonation on Twitter](#).

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow

or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation rules

Updated November 3, 2017

This page is primarily intended for **developers**.

For Twitter users: You are ultimately responsible for the actions taken with your account, or by applications associated with your account. Before authorizing a third-party application to access or use your account, make sure you've thoroughly investigated the application and understand what it will do. If automated activity on your account violates the [Twitter Rules](#) or these automation rules, Twitter may take action on your account, including [filtering your Tweets from search results](#) or suspending your account.

For more information on third-party applications, please see our article on [connecting and revoking third-party applications](#).

If you are a developer and have questions about these Automation Rules, please reach out via <https://twittercommunity.com/c/rules-and-policies>.

If your application has been suspended or restricted from performing write actions, you may submit an appeal at </forms/platform>.

I. Ground Rules Do!

Build solutions that automatically broadcast helpful information in Tweets.

Run creative campaigns that auto-reply to users who engage with your content.

Build solutions that automatically respond to users in Direct Messages.

Try new things that help people (and comply with our rules).

Make sure your application provides a good user experience and performs well — and confirm that remains the case over time.

Don't!

Violate these or other policies. Be extra mindful of our rules about abuse and user privacy.

Abuse the Twitter API or attempt to circumvent rate limits.

Use non-API-based forms of automation, such as scripting the Twitter website. The use of these techniques may result in the permanent suspension of your account.

Spam or bother users, or otherwise send them unsolicited messages.

A. The Twitter Rules and the Developer Agreement and Policy

As with all activity on Twitter, automated activity is subject to the [Twitter Rules](#) and, if you're a developer using the Twitter API, the [Developer Agreement and Policy](#).

You should carefully review these policies to ensure that your automated activity is compliant. Automated applications or activities that violate these policies, or that facilitate or induce users to violate them, may be subject to enforcement action, potentially including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these policies.

Although all aspects of the Twitter Rules and the Developer Agreement and Policy apply to automated activity, you should keep the following rules top of mind:

Spamming: You may not send automated Tweets or Direct Messages that are spam, or otherwise engage in spamming activity. Some examples of spammy behavior to avoid with automation include:

Trending topics: You may not automatically post about trending topics on Twitter, or use automation to attempt to influence or manipulate trending topics.

Multiple posts/accounts: You may not post duplicative or substantially similar Tweets on one account or over multiple accounts you operate.

Duplicate accounts: You may not create and/or automate multiple accounts for duplicative or substantially similar use cases.

However, automating multiple accounts for related but non-duplicative use cases is permitted. For example, you may automate separate accounts to Tweet when the Hubble Space Telescope passes over different cities, such as [San Francisco](#) or [Hong Kong](#).

Misleading links: You may not send automated Tweets or Direct Messages containing links that are misleading, including links that maliciously or deceptively redirect through landing pages or ad pages before displaying the final content.

Sensitive media: Automated Tweets and Direct Messages must comply with the [Twitter media policy](#), and you should mark your account as potentially sensitive if you intend to post graphic, pornographic, or potentially sensitive media.

Abusive behavior: You may not engage in any automated activity that encourages, promotes, or incites abuse, violence, hateful conduct, or harassment, on or off Twitter.

Private information: You may not post private or confidential information about a person without their prior express authorization.

B. Other Ground Rules for Automated Activity

In addition to the policies above, the following ground rules apply to all automated activity on Twitter:

Don't surprise or mislead users: Automated activity should honor users' expectations. Ask for the user's permission before taking an action if you aren't sure.

Mature content or profanity: Don't Direct Message, mention, or reply to users with potentially sensitive content (including profanity), unless they've clearly indicated an intent to receive it in advance.

Be thoughtful about the information you request or exchange on Twitter

Tweets: Don't ask users to send you personal or private information via a public Tweet. If you need additional personal or private information from a user to provide them with customer service (or other similar use cases), you should ask the user to share such information by Direct Message or another private channel. You might even consider adding a [Direct Message deep link](#) to your Tweet.

Direct Messages: You should only ask users for the minimum amount of information you need to provide them with service. If you need to request or exchange

particularly sensitive information (such as credit card information), you should consider directing users to your website or other appropriate channel to do so.

II. Activity-Specific Rules

The activity-specific rules in this section apply to taking specific automated actions on Twitter. Please read these rules carefully, as they outline both permitted and prohibited use cases of automation.

Automated applications or activities that violate these rules, or that facilitate or induce users to violate them, may be subject to enforcement action, including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access the Twitter API based on violations of these rules. As a reminder, you should also carefully review the spam guidelines in the [Twitter Rules](#) to avoid having activities performed by you, your app, or other users through your app or service flagged as spam.

A. Automated Actions Through Another User's Account

Twitter users may authorize your app or service to [access their Twitter account through OAuth](#). A user authorizing your app or service to access their Twitter account through OAuth does not by itself constitute sufficient consent to take automated actions through that user's account.

You may only take automated actions through another Twitter user's account if you:

clearly describe to the user the types of automated actions that will occur; receive express consent from the user to take those automated actions; and immediately honor a user's request to opt-out of further automated actions.

If you substantially change the purpose or functionality of your app or service, you must re-obtain express consent from the user to take automated action through their account before doing so.

These requirements apply to any automated action taken through another Twitter user's account, including posting

Tweets, sending Direct Messages, deleting Tweets or Direct Messages, or following/unfollowing other accounts. For applications that offer users the ability to delete Tweets in a bulk or automated manner, you must also clearly state that Tweets are not recoverable once deleted.

B. Automated Tweets

1. Posting automated Tweets

Automated Tweets that cross-post outside information:

You may post automated Tweets based on sources of outside information — such as an RSS feed, weather data, etc. — as long as you are sufficiently authorized to publish such information.

Other automated Tweets (excluding mentions or replies):

Provided you comply with all other rules, you may post automated Tweets for entertainment, informational, or novelty purposes. As a reminder, accounts posting duplicative, spammy, or otherwise prohibited content may be subject to suspension.

2. Posting automated mentions and replies

The reply and mention functions are intended to make communication between Twitter users easier. Automating these actions to reach many users on an unsolicited basis is an abuse of the feature, and is not permitted. For example, sending automated replies to Tweets based on keyword searches alone is not permitted. Spamy or duplicative use of mentions and replies may result in enforcement action, such as the removal of your Tweets from Search or the suspension of your app or account.

However, you may send automated replies or mentions to Twitter users so long as:

in advance of sending the automated reply, the recipient or mentioned user(s) have requested or have clearly indicated an intent on Twitter to be contacted by you (i.e.

opted in), for example by replying to a Tweet from your account, or by sending you a Direct Message;

you provide a clear and easy way for such users to opt-out of receiving automated replies and mentions, and promptly honor all such opt-out requests;

you only send one automated reply or mention per user interaction; and

the automated reply or mention is a reply to the user's original Tweet (if your campaign is based on users posting a reply to your Tweet).

Opt-in techniques and indications of user intent take many different forms, depending on the specifics of your use case and implementation. Some examples include:

A Tweet from your account that clearly indicates that a user taking a specific action on that Tweet (such as Retweeting it) will opt the user into receiving an automated response.

A mention of your account by the user in a manner suggesting the user clearly wishes or intends to receive a response. If you want to run an auto-reply campaign with a campaign- or use-case-specific hashtag, users should also mention you in their Tweets.

Note that a user following your account is not on its own a sufficient indication of user intent to receive an automated response.

: Advertisers, publishers, and brands using auto-response campaigns must request approval from Twitter and may be subject to additional rules. Please reach out to your account or partner manager for assistance. If you do not have a partner manager but are interested in managed client services, you can learn more [here](#).

Additionally, we recommend that any accounts that will communicate with

users via automated mentions or replies:

Note

1. Appropriately filter responses based on potentially sensitive language in user handles,

display names, and Tweet text, as well as potentially sensitive media;

2. Check that the Tweet you are mentioning or replying to still exists (for example, using the statuses/lookup endpoint on the Twitter API).

C. Automated Direct Messages

Sending automated Direct Messages to users

You may not send unsolicited Direct Messages in a bulk or automated manner, and should be thoughtful about the frequency with which you contact users via Direct Message.

You may send automated Direct Messages to users so long as:

in advance of sending the Direct Message, the recipient(s) have requested or have clearly indicated an intent on Twitter to be contacted by you via Direct Message, for example by sending you a Direct Message; **and**

you provide a clear and easy way for such users to opt-out of receiving automated Direct Messages, and promptly honor all such opt-out requests.

The fact that a user is technically able to receive a Direct Message from you (e.g. because the user follows you, has enabled the ability to receive Direct Messages from any account, or because the user is in a pre-existing Direct Message conversation with you) does not necessarily mean they have requested or expect to receive automated Direct Messages from you.

Interacting with users via Direct Message

Per the Ground Rules, remember to be thoughtful about the amount and type of information you request or exchange with users via Direct Messages. If you will be asking a user to provide personal or private information via an automated Direct Message, you must clearly explain how you will use the information you're collecting. Consider including a link to your privacy policy in your Direct Message to the user, as well as in your Twitter profile bio.

Don't publicly share information received in a Direct Message conversation with a user without first obtaining explicit consent from the user. For example, if a user asks you via Direct Message about a purchase they made from you, you may not

mention the user in a public Tweet that includes information about their purchase unless you have the user's explicit consent to do so.

After a user-initiated interaction ends, don't send additional follow-up Direct Messages or mention users in a Tweet unless you get permission from the user.

D. Automated actions you take on Tweets or accounts

Automated likes: You may not like Tweets in an automated manner.

Automated Retweets: Provided you comply with all other rules, you may Retweet or Quote Tweet in an automated manner for entertainment, informational, or novelty purposes. Automated Retweets often lead to negative user experiences, and bulk, aggressive, or spammy Retweeting is a violation of the [Twitter Rules](#).

Automated following/unfollowing: You may not follow or unfollow Twitter accounts in a bulk, aggressive, or indiscriminate manner. Aggressive following is a violation of the [Twitter Rules](#). Please also review our [following rules and best practices](#) to ensure you are in compliance. Note that applications that claim to get users more followers are also prohibited under the [Twitter Rules](#).

Automated adding to lists or collections: You may not add Twitter users to lists or add Tweets to collections in a bulk or indiscriminate manner. Adding a large number of unrelated users to lists is a violation of the [Twitter Rules](#).

About search rules and restrictions

This article details the rules for using Twitter search, and explains reasons why content may not show up in Twitter search.

Why doesn't all content show up in Twitter search?

In order to keep your search results relevant, Twitter filters search results for quality Tweets and accounts. Material that jeopardizes search quality or creates a bad search experience for other people may be automatically removed from Twitter search. Read more about this [here](#).

General rules for using Twitter search

The [Twitter Rules](#) explain what behaviors are permitted on Twitter. In addition to these rules, we've included some tips below to keep your content relevant (and your Tweets in Twitter search).

Please do not:

Repeatedly post duplicate or near-duplicate content (links or Tweets). Abuse trending topics or hashtags (topic words with a # sign).

Send automated Tweets or replies.

Use bots or applications to post similar messages based on keywords. Post similar messages over multiple accounts.

Aggressively follow and unfollow people.

Twitter may automatically remove accounts engaging in these behaviors from search (or even suspend in some cases) in order to ensure the best experience for everyone.

Search results filtering

To learn how to filter your search results, including how to disable or enable safe search, read our [Using Twitter search](#) article.

Search results filtering in third-party applications

Accounts may be filtered from search for posting similar messages over several accounts. Your account may be filtered from search if you are logged into a third-party application that is updating many accounts with similar or duplicate content.

Rules for participating in contests on Twitter

Businesses sometimes host contests and offer prizes to Retweet an update, following a particular account, or posting Tweets with a specific hashtag topic or reply. If you're running a contest, check out this [article](#) for guidelines for contests on Twitter.

If you repeatedly post near-duplicate updates or duplicate links in order to enter contests, you may be filtered from search. Please keep in mind that you are responsible for the updates that you post to your Twitter account.

What to do if you are missing from search

Please see this [article](#) for more information related to this issue.

If you are being filtered from search, rest assured that your followers will still receive your updates. If you post a reply or mention to a specific person, it will be delivered to that person.

Having trouble?

You can learn more about searching by keyword, location, and trending topics in our [Using Twitter search](#) help article.

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. These limitations are set forth in the Twitter Rules below.

The Twitter Rules (along with all incorporated policies), [Privacy Policy](#), and [Terms of Service](#) collectively make up the "Twitter User Agreement" that governs a user's access to and use of Twitter's services.

All individuals accessing or using Twitter's services must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in Twitter taking one or more of the following enforcement actions:

- requiring you to delete prohibited content before you can again create new posts and interact with other Twitter users;
- temporarily limiting your ability to create posts or interact with other Twitter users;
- asking you to verify account ownership with a phone number or email address; or
- permanently suspending your account(s).

If you attempt to evade a permanent suspension by creating new accounts, we will suspend your new accounts.

Please note that we may need to change these Rules from time to time and reserve the right to do so. The most current version is always available at: <https://twitter.com/rules>.

The policies set forth in these Twitter Rules govern organic content on our platform. To learn more about the rules which govern ads and promoted content, please review our [Ads policies](#).

Content Boundaries and Use of Twitter

Intellectual property

Trademark: We reserve the right to suspend accounts or take other appropriate action when someone's brand or trademark, including business name and/or logo, is used in a manner that may mislead or confuse others about your brand affiliation. Read more about our [trademark policy](#) and [how to report a violation](#).

Copyright: We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in our [Terms of Service](#). Read more about our [copyright policy](#).

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Learn more about our [media policy](#).

Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images. Additionally, Twitter may sometimes require you to remove excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about [how to make such a request](#), and [how to mark your media as sensitive](#).

Unlawful use

You may not use our service for any unlawful purposes or in furtherance of illegal activities. By using Twitter, you agree to comply with all applicable laws governing your online conduct and content.

Trends

At times, we may prevent certain content from trending. This includes content that violates the Twitter Rules, as well as content that may attempt to manipulate trends. Read more about [what we allow and do not allow to trend](#).

Misuse of Twitter badges

You may not use badges, including but not limited to the “promoted” or “verified” Twitter badges, unless provided by Twitter. Accounts using unauthorized badges as part of their profile photos, header photos, display names, or in any way that falsely implies affiliation with Twitter or authorization from Twitter to display these badges, may be suspended.

Misuse of usernames

Selling usernames: You may not buy or sell Twitter usernames.

Username squatting: You may not engage in username squatting. Some of the factors we take into consideration when determining whether conduct is username squatting include:

- the number of accounts created;
- the creation of accounts for the purpose of preventing others from using those account names;
- the creation of accounts for the purpose of selling those accounts; and
- the use of third-party content feeds to update and maintain accounts under the names of those third parties.

Please note that Twitter may also remove accounts that are inactive for more than six months. Learn more about [username squatting](#).

Abusive Behavior

We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we prohibit behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Context matters when evaluating for abusive behavior and determining appropriate enforcement actions. Factors we may take into consideration include, but are not limited to whether:

- the behavior is targeted at an individual or group of people;
- the report has been filed by the target of the abuse or a bystander;
- the behavior is newsworthy and in the legitimate public interest.

Violence and physical harm

Violence: You may not make specific threats of violence or wish for the serious physical harm, death, or disease of an individual or group of people. This includes, but is not limited to, threatening or promoting terrorism. You also may not affiliate with organizations that — whether by their own statements or activity both on and off the platform — use or promote violence against civilians to further their causes. We will begin enforcing this rule around affiliation with such organizations on December 18, 2017.

Suicide or self-harm: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Child sexual exploitation: You may not promote child sexual exploitation. Learn more about our zero-tolerance [child sexual exploitation policy](#).

Abuse and hateful conduct

Abuse: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Unwanted sexual advances: You may not direct abuse at someone by sending unwanted sexual content, objectifying them in a sexually explicit manner, or otherwise engaging in sexual misconduct.

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. Read more about our [hateful conduct policy](#).

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category. We will begin enforcing this rule on December 18, 2017.

Private information and intimate media

Private information: You may not publish or post other people's private information without their express authorization and permission. Definitions of private information may vary depending on local laws. Read more about our [private information policy](#).

Intimate media: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. Read more about intimate media on Twitter.

Threats to expose / hack: You may not threaten to expose someone's private information or intimate media. You also may not threaten to hack or break into someone's digital information.

Impersonation

You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. While you may maintain parody, fan, commentary, or newsfeed accounts, you may not do so if the intent of the account is to engage in spamming or abusive behavior. Read more about our [impersonation policy](#).

Spam and Security

We strive to protect people on Twitter from technical abuse and spam.

To promote a stable and secure environment on Twitter, you may not do, or attempt to do, any of the following while accessing or using Twitter:

- Access, tamper with, or use non-public areas of Twitter, Twitter's computer systems, or the technical delivery systems of Twitter's providers (except as expressly permitted by the Twitter Bug Bounty program).
- Probe, scan, or test the vulnerability of any system or network, or breach or circumvent any security or authentication measures (except as expressly permitted by the Twitter Bug Bounty program).
- Access or search, or attempt to access or search, Twitter by any means (automated or otherwise) other than through our currently available, published interfaces that are provided by Twitter (and only pursuant to the applicable terms and conditions), unless you have been specifically allowed to do so in a separate agreement with Twitter. Note that crawling Twitter is permissible if done in accordance with the provisions of the robots.txt file; however, scraping Twitter without our prior consent is expressly prohibited.
- Forge any TCP/IP packet header or any part of the header information in any email or posting, or in any way use Twitter to send altered, deceptive, or false source-identifying information.
- Interfere with or disrupt the access of any user, host or network, including, without limitation, sending a virus, overloading, flooding, spamming, mail-bombing Twitter's services, or by scripting the creation of content in such a manner as to interfere with or create an undue burden on Twitter.

Any accounts engaging in the following activities may be temporarily locked or subject to permanent suspension:

- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use Twitter's services for the purpose of spamming anyone. Spam is generally defined on Twitter as bulk or aggressive activity that attempts to manipulate or disrupt Twitter or the experience of users on Twitter to drive traffic or attention to unrelated accounts, products, services, or initiatives. Some of the factors that we take into account when determining what conduct is considered to be spamming include:
 - if you have followed and/or unfollowed a large number of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if your Tweets or Direct Messages consist mainly of links shared without commentary;
 - if a large number of people have blocked you in response to high volumes of untargeted, unsolicited, or duplicative content or engagements from your account;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicative or substantially similar content, replies, or mentions over multiple accounts or multiple duplicate updates on one account, or create duplicate or substantially similar accounts;
 - if you post multiple updates to a trending or popular topic with an intent to subvert or manipulate the topic to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you send large numbers of unsolicited replies or mentions;
 - if you add users to lists in a bulk or aggressive manner;
 - if you are randomly or aggressively engaging with Tweets (e.g., likes, Retweets, etc.) or users (e.g., following, adding to lists or Moments, etc.) to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you repeatedly post other people's account information as your own (e.g., bio, Tweets, profile URL, etc.);
 - if you post misleading, deceptive, or malicious links (e.g., affiliate links, links to malware/clickjacking pages, etc.);
 - if you create fake accounts, account interactions, or impressions;
 - if you sell, purchase, or attempt to artificially inflate account interactions (such as followers, Retweets, likes, etc.); and
 - if you use or promote third-party services or apps that claim to get you more followers, Retweets, or likes (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers or engagements to your account or Tweets).

Please see our support articles on [following rules and best practices](#) and [automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts may be permanently suspended.

Content Visibility

Accounts under investigation or which have been detected as sharing content in violation of these Rules may have their account or Tweet visibility limited in various parts of Twitter, including search. To learn more about situations in which content may be restricted on Twitter, please see our support article on [search rules and restrictions](#).

Trademark policy

Updated June 30, 2016

What is a Trademark Policy Violation on Twitter?

Using another's trademark in a manner that may mislead or confuse others about your brand affiliation may be a violation of our trademark policy.

How Does Twitter Respond to Trademark Reports?

We review each report and do the following:

- If we determine that there is a violation of our trademark policy, we may suspend the account. In some instances, we may give the account holder an opportunity to comply with our policies before suspending the account. For more information, see our [Parody, commentary, and fan account policy](#).
- We may also provide the account holder with your name and information included in the copy of the report.
- If you would like to report a trademark violation in a [Twitter Ad](#), please visit our [Trademark Policy for Promoted Ads](#) to learn more.

Please note that Twitter only investigates requests that are submitted by the trademark holder or by an authorized representative of the trademark holder.

What is not a Trademark Policy Violation?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples include:

- Using a trademark in a way that is outside the scope of the trademark registration (e.g. territory, or goods and services identified in the registration).
- Nominative and other fair uses of trademarks are protected uses under our trademark policy, so long as the account is clearly distinguished from the trademark owner. This includes use by resellers in certain regions and accounts engaging in parody, commentary, or news. For more information, see our [Parody, commentary, and fan account policy](#).

How do I Report a Trademark Policy Violation?

Submit trademark reports through our [Trademark Issue support form](#). Using the appropriate form ensures we have all the information necessary to process your report quickly. Please provide all information requested in the form. If you submit a report with incomplete information, we will follow up with you to request that information. Please note that this will result in a delay in processing your report.

If you have trouble using the form, try updating your browser or switching to another browser.

Copyright policy

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged

infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notification.

When Should I File a Counter-notice?

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);

3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Twitter media policy

Twitter Rules: Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images.

Twitter Rules: You may not use hateful images or symbols in your profile image or profile header.

Rationale

Some forms of graphic violence, adult content, or hateful imagery may be permitted in Tweets when they are marked as [sensitive media](#). However, you may not include this type of content in live video, or in profile or header images.

While we want people to feel free to share media that reflects their creativity or individuality, or to show what's happening in the world, we will take action when it crosses the line into abuse towards a person, group, or protected category.

Finally, we may sometimes require you to remove media containing excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about how to make such a request [here](#).

When this applies

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. Some examples include, but are not limited to, depictions of:

- the moment at which someone dies
- gruesome crime or accident scenes
- bodily harm, torture, dismemberment, or mutilation

We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity (including close-ups of genitals, buttocks, or breasts)
 - Please note that exceptions may be made for artistic, medical, health, or educational content. Breastfeeding content does not need to be marked as sensitive.
- simulating a sexual act
- intercourse or any sexual act (may involve humans, humanoid animals, cartoons, or anime)

If you see media you think may be in violation of our policy, report it. You can read about how to report sensitive media [here](#).

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, or ethnicity/national origin.

Some examples of hateful imagery include, but are not limited to:

- Symbols historically associated with hate groups (for example, the Nazi swastika)
- Images depicting others as less than human or photoshopped to include hateful symbols
- Photoshopped references to a mass murder that targeted a protected category

Do I need to be the target of this content in order for it to be a violation of the Twitter Rules?
No, we review both first-person and bystander reports of such content.

Consequences

When adult content, graphic violence, or hateful imagery appears in Tweets, we may place this content behind an interstitial advising viewers to be aware that they will see sensitive media if they click through. This allows us to identify potentially sensitive content that some people may not wish to see. Learn more about [how to control whether you see sensitive media](#).

However, if such content appears in live video, header, or profile images, it will violate our media policy. The consequences for violating this policy depend on the severity of the violation and the person's previous record of violations.

The first time someone violates this policy they will be required to remove the imagery and may be asked to verify their account ownership. Subsequent violations could result in permanent suspension. Learn more about [our range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

How to contact Twitter about media concerning a deceased family member

In order to respect the wishes of loved ones, Twitter will remove imagery of deceased individuals in certain circumstances. Immediate family members and other authorized individuals may request the removal of images

or video of deceased individuals, from when critical injury occurs to the moments before or after death, by sending a request to Twitter Inc. via our [privacy form](#). When reviewing such media removal requests, Twitter considers public interest factors such as the newsworthiness of the content and may not be able to honor every request. Note: Find information about contacting Twitter to [deactivate the account of a deceased or incapacitated user](#).

How to mark media as sensitive links to Media policy

FAQs about trends on Twitter

Where can I find trends?

Trends are available on the Twitter for iOS app, the Twitter for Android app, and [twitter.com](#).

On Twitter's mobile apps, you can find trends listed under the trends section



of the **Explore** tab When signed in to [twitter.com](#) on a desktop or laptop computer, trends are listed in many places, including the Home timeline, Notifications, search results, and profile pages.

How are trends determined?

Trends are determined by an algorithm and, by default, are tailored for you based on who you follow, your interests, and your location. This algorithm identifies topics that are popular now, rather than topics that have been popular for a while or on a daily basis, to help you discover the hottest emerging topics of discussion on Twitter.

Note: The number of Tweets that are related to the trends is just one of the factors the algorithm looks at when ranking and determining trends. Algorithmically, trends and hashtags are grouped together if they are related to the same topic. For instance, #MondayMotivation and #MotivationMonday may both be represented by #MondayMotivation.

You can choose to see trends that are not tailored for you by selecting a [specific trends location](#) on [twitter.com](#), iOS, or Android. Location trends identify popular topics among people in a specific geographic location.

What does the # sign mean?

You may notice that some trends have # sign before the word or phrase. This is called a [hashtag](#) and is included specifically in Tweets to mark them as relating to a topic, so that people can follow the conversation in search.

Will I see world and local events in my trends?

Yes, when signed into your account on the [twitter.com](#), iOS, or Android, trends will be tailored for you based on your location and who you follow. There will be many world and local news events and conversations that will appear in your trends regardless of your personalization.

Can I see trends for a specific location?

Yes, you can view trends for a specific location on twitter.com, iOS, or Android.

View trends via twitter.com:

To view trends for a specific location, click **Change** in your trends box. You can choose to **Keep tailored trends** based on your location and who you follow, or you can select **Change** to choose a nearby location or search locations.

Adjust your trends settings via Twitter for iOS:

1. In the top menu, tap your **profile** icon.
2. Tap **Settings and privacy**, then tap **Content preferences**.
3. Under **Explore**, tap **Trends**.
4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.
5. When the feature is disabled, you can change your location by tapping **Change location**.
 1. In the search box, type your desired trends location or choose from the list.

Adjust your trends settings via Twitter for Android:

1. In the top menu, you will either see a **navigation menu** icon  or your **profile** icon. Tap whichever icon you have, and select **Settings and privacy**.
2. Tap **Content preferences**.
3. Under **Explore**, tap **Trends**.
4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.
5. When the feature is disabled, you can change your location by tapping **Change location**.
 1. In the search box, type your desired trends location or choose from the list.

Note: If you don't find your preferred country or city, it means we're not receiving enough Tweets from that geographical area to create a list. You can look up local Tweets on any topic by using [advanced Twitter Search operators](#).

What happens when I click or tap on a trend?

Clicking or tapping any of the trends takes you to the Twitter search results for that trend. You'll see all Tweets including that phrase or hashtag. To see what people are saying about a previous trend, perform a [search for that keyword](#).

How can I participate in a trend?

Simply post a Tweet including the exact word or phrase as it appears in the trends list (with the hashtag, if you see one). Due to the large number of people Tweeting about these specific trends, you may not always be able to find your particular Tweet in search, but your followers will always see your Tweets.

Note: Twitter also filters searches for quality. To understand what types of behavior could cause your Tweets to be filtered out, read about our [search rules and restrictions](#).

Are there rules for trends?

Yes. We want trends to promote healthy discussions on Twitter. This means that at times, we may prevent certain content from trending. These include trends that:

- Contain profanity or adult/graphic references.
- Incite hate on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease.
- Violate the [Twitter Rules](#).

In some cases, we may also consider the newsworthiness of the content, or if it is in the public interest when evaluating potential violations. In these cases, the content might continue to trend on our platform.

Note that even if we prevent the hashtag or the content from appearing on the list of trending topics, you may still be able to access conversations around that content on Twitter.

Trends descriptions

Below some trends, you may see the approximate number of Tweets associated with the trend, or personalized context such as who in your network is Tweeting about it.

Additionally, when a Moment is found to be relevant to a trend, we may use the title of the Moment as the trend's description.

Having trouble?

Check out our [troubleshooting section](#) for solutions to common problems.

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Child sexual exploitation policy

Twitter removes content promoting child sexual exploitation

We do not tolerate child sexual exploitation on Twitter. When we are made aware of links to images of or content promoting child sexual exploitation they will be removed from the site without further notice and reported to The National Center for Missing & Exploited Children ("NCMEC"); we permanently suspend accounts promoting or containing updates with links to child sexual exploitation.

Reporting content promoting child sexual exploitation

If you think you've found a Twitter account distributing or promoting child sexual exploitation while using Twitter, please click [here](#) to report it. Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. To find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with the National Center for Missing and Exploited Children [here](#), or your local law enforcement. If you believe that the content was hosted on a website outside of the United States, you can report it [here](#) on the International Association of Internet Hotlines website.

NOTE: Please do NOT tweet, retweet or repost child sexual exploitation for any reason. Click [here](#) to report it to us immediately and we will take steps to remove it.

Law Enforcement Agencies

If you are a member of a law enforcement agency seeking information regarding Twitter accounts promoting or distributing child sexual exploitation, please see our [Guidelines for Law Enforcement](#) for information on our process; we also suggest you contact the National Center for Missing and Exploited Children (NCMEC) [here](#), regardless of jurisdiction, to determine whether or not information on those accounts is already on file.

For general law enforcement inquiries, please see our [Guidelines for Law Enforcement](#) for more information.

Hateful conduct policy

Freedom of expression means little if voices are silenced because people are afraid to speak up. We do not tolerate behavior that harasses, intimidates, or uses fear to silence another person's voice. If you see something on Twitter that violates these rules, please report it to us.

How our policy works

As explained in the Twitter Rules,

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Examples of what we do not tolerate includes, but is not limited to behavior that harasses individuals or groups of people with:

violent threats;

wishes for the physical harm, death, or disease of individuals or groups;

references to mass murder, violent events, or specific means of violence in which/with which such groups have been the primary targets or victims;

behavior that incites fear about a protected group;

repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone.

How our enforcement works Context matters.

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. While we accept reports of violations from anyone, sometimes we also need to hear directly from the target to ensure that we have proper context.

The number of reports we receive does not impact whether or not something will be removed. However, it may help us prioritize the order in which it gets reviewed.

We focus on behavior.

We enforce policies when someone reports behavior that is abusive and targets an entire protected group and/or individuals who may be members.

This targeting can happen in any manner (for example, @mentions, tagging a photo, and more).

We have a range of enforcement options.

The consequences for violating our rules vary depending on the severity of the violation and the person's previous record of violations. For example, we may ask someone to remove the offending Tweet before they can Tweet again. For other cases, we may suspend an account.

Report private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Read about our [policy regarding private information posted on Twitter](#).

Do I need to have a Twitter account to make a report?

No, you do not need to have a Twitter account to file a report that your private and confidential information has been posted on Twitter.

Can I report unauthorized posting of Copyrighted material?

If your copyrighted material has been Tweeted without your permission, please review our [Copyright and DMCA policy](#).

What can I do if my information is posted on other websites?

We cannot take down any information that is posted on a website other than Twitter. If your information is posted on another website, please contact the help center or support staff for the website where your information is posted.

How do I report that my private information is posted on Twitter?

You can report that your private information has been posted on Twitter directly from the offending Tweet(s). Learn [how to report individual Tweets](#).

File a [report that someone has posted your private information](#).

What happens when Twitter receives a valid report?

When we receive a complete and valid report that private information has been posted on Twitter, we'll investigate the account and Tweets reported. We will review where, if anywhere, the information has been made publicly available before taking action on the account or Tweets. If the information you reported was previously posted elsewhere on the Internet, it is not a violation of our policy and we will not take action.

Learn more about [protecting your private information on Twitter and other websites](#).

Read our [policy regarding private information posted on Twitter](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter impersonation policy.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to Twitter's [parody, commentary, and fan account policy](#) for more information about these types of accounts.

Find more information about [reporting impersonation on Twitter](#).

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation rules

Updated November 3, 2017

This page is primarily intended for **developers**.

For Twitter users: You are ultimately responsible for the actions taken with your account, or by applications associated with your account. Before authorizing a third-party application to access or use your account, make sure you've thoroughly investigated the application and understand what it will do. If automated activity on your account violates the [Twitter Rules](#) or these automation rules, Twitter may take action on your account, including [filtering your Tweets from search results](#) or suspending your account.

For more information on third-party applications, please see our article on [connecting and revoking third-party applications](#).

If you are a developer and have questions about these Automation Rules, please reach out via <https://twittercommunity.com/c/rules-and-policies>.

If your application has been suspended or restricted from performing write actions, you may submit an appeal at </forms/platform>.

I. Ground Rules Do!

Build solutions that automatically broadcast helpful information in Tweets.

Run creative campaigns that auto-reply to users who engage with your content.

Build solutions that automatically respond to users in Direct Messages.

Try new things that help people (and comply with our rules).

Make sure your application provides a good user experience and performs well — and confirm that remains the case over time.

Don't!

Violate these or other policies. Be extra mindful of our rules about abuse and user privacy.

Abuse the Twitter API or attempt to circumvent rate limits.

Use non-API-based forms of automation, such as scripting the Twitter website. The use of these techniques may result in the permanent suspension of your account.

Spam or bother users, or otherwise send them unsolicited messages.

A. The Twitter Rules and the Developer Agreement and Policy

As with all activity on Twitter, automated activity is subject to the [Twitter Rules](#) and, if you're a developer using the Twitter API, the [Developer Agreement and Policy](#).

You should carefully review these policies to ensure that your automated activity is compliant. Automated applications or activities that violate these policies, or that facilitate or induce users to violate them, may be subject to enforcement action, potentially including suspension of associated Twitter accounts.

We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these policies.

Although all aspects of the Twitter Rules and the Developer Agreement and Policy apply to automated activity, you should keep the following rules top of mind:

Spamming: You may not send automated Tweets or Direct Messages that are spam, or otherwise engage in spamming activity. Some examples of spammy behavior to avoid with automation include:

Trending topics: You may not automatically post about trending topics on Twitter, or use automation to attempt to influence or manipulate trending topics.

Multiple posts/accounts: You may not post duplicative or substantially similar Tweets on one account or over multiple accounts you operate.

Duplicate accounts: You may not create and/or automate multiple accounts for duplicative or substantially similar use cases.

However, automating multiple accounts for related but non-duplicative use cases is permitted. For example, you may automate separate accounts to Tweet when the Hubble Space Telescope passes over different cities, such as [San Francisco](#) or [Hong Kong](#).

Misleading links: You may not send automated Tweets or Direct Messages containing links that are misleading, including links that maliciously or deceptively redirect through landing pages or ad pages before displaying the final content.

Sensitive media: Automated Tweets and Direct Messages must comply with the [Twitter media policy](#), and you should mark your account as potentially sensitive if you intend to post graphic, pornographic, or potentially sensitive media.

Abusive behavior: You may not engage in any automated activity that encourages, promotes, or incites abuse, violence, hateful conduct, or harassment, on or off Twitter.

Private information: You may not post private or confidential information about a person without their prior express authorization.

B. Other Ground Rules for Automated Activity

In addition to the policies above, the following ground rules apply to all automated activity on Twitter:

Don't surprise or mislead users: Automated activity should honor users' expectations. Ask for the user's permission before taking an action if you aren't sure.

Mature content or profanity: Don't Direct Message, mention, or reply to users with potentially sensitive content (including profanity), unless they've clearly indicated an intent to receive it in advance.

Be thoughtful about the information you request or exchange on Twitter

Tweets: Don't ask users to send you personal or private information via a public Tweet. If you need additional personal or private information from a user to provide them with customer service (or other similar use cases), you should ask the user to share such information by Direct Message or another private channel. You might even consider adding a [Direct Message deep link](#) to your Tweet.

Direct Messages: You should only ask users for the minimum amount of information you need to provide them with service. If you need to request or exchange particularly sensitive information (such as credit card information), you should consider directing users to your website or other appropriate channel to do so.

II. Activity-Specific Rules

The activity-specific rules in this section apply to taking specific automated actions on Twitter. Please read these rules carefully, as they outline both permitted and prohibited use cases of automation.

Automated applications or activities that violate these rules, or that facilitate or induce users to violate them, may be subject to enforcement action, including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these rules. As a reminder, you should also carefully review the spam guidelines in the [Twitter Rules](#) to avoid having activities performed by you, your app, or other users through your app or service flagged as spam.

A. Automated Actions Through Another User's Account

Twitter users may authorize your app or service to [access their Twitter account through OAuth](#). A user authorizing your app or service to access their Twitter account through OAuth does not by itself constitute sufficient consent to take automated actions through that user's account.

You may only take automated actions through another Twitter user's account if you:

clearly describe to the user the types of automated actions that will occur; receive express consent from the user to take those automated actions; and immediately honor a user's request to opt-out of further automated actions.

If you substantially change the purpose or functionality of your app or service, you must re-obtain express consent from the user to take automated action through their account before doing so.

These requirements apply to any automated action taken through another Twitter user's account, including posting Tweets, sending Direct Messages, deleting Tweets or Direct Messages, or following/unfollowing other accounts. For applications that offer users the ability to delete Tweets in a bulk or automated manner, you must also clearly state that Tweets are not recoverable once deleted.

B. Automated Tweets

1. Posting automated Tweets

Automated Tweets that cross-post outside information:

You may post automated Tweets based on sources of outside information — such as an RSS feed, weather data, etc. — as long as you are sufficiently authorized to publish such information.

Other automated Tweets (excluding mentions or replies):

Provided you comply with all other rules, you may post automated Tweets for entertainment, informational, or novelty

purposes. As a reminder, accounts posting duplicative, spammy, or otherwise prohibited content may be subject to suspension.

2. Posting automated mentions and replies

The reply and mention functions are intended to make communication between Twitter users easier. Automating these actions to reach many users on an unsolicited basis is an abuse of the feature, and is not permitted. For example, sending automated replies to Tweets based on keyword searches alone is not permitted. Spamy or duplicative use of mentions and replies may result in enforcement action, such as the removal of your Tweets from Search or the suspension of your app or account.

However, you may send automated replies or mentions to Twitter users so long as:

in advance of sending the automated reply, the recipient or mentioned user(s) have requested or have clearly indicated an intent on Twitter to be contacted by you (i.e. opted in), for example by replying to a Tweet from your account, or by sending you a Direct Message;

you provide a clear and easy way for such users to opt-out of receiving automated replies and mentions, and promptly honor all such opt-out requests;

you only send one automated reply or mention per user interaction; and

the automated reply or mention is a reply to the user's original Tweet (if your campaign is based on users posting a reply to your Tweet).

Opt-in techniques and indications of user intent take many different forms, depending on the specifics of your use case and implementation. Some examples include:

A Tweet from your account that clearly indicates that a user taking a specific action on that Tweet (such as Retweeting it) will opt the user into receiving an automated response.

A mention of your account by the user in a manner suggesting the user clearly wishes or intends to receive a response. If you want to run an auto-reply campaign with a campaign- or use-case-specific hashtag, users should also mention you in their Tweets.

Note that a user following your account is not on its own a sufficient indication of user intent to receive an automated response.

: Advertisers, publishers, and brands using auto-response campaigns must request approval from Twitter and may be subject to additional rules. Please reach out to your account or partner manager for assistance. If you do not have a partner manager but are interested in managed client services, you can learn more [here](#).

Additionally, we recommend that any accounts that will communicate with

users via automated mentions or replies:

Note

1. Appropriately filter responses based on potentially sensitive language in user handles,

display names, and Tweet text, as well as potentially sensitive media;

2. Check that the Tweet you are mentioning or replying to still exists (for example, using the statuses/lookup endpoint on the Twitter API).

C. Automated Direct Messages

Sending automated Direct Messages to users

You may not send unsolicited Direct Messages in a bulk or automated manner, and should be thoughtful about the frequency with which you contact users via Direct Message.

You may send automated Direct Messages to users so long as:

in advance of sending the Direct Message, the recipient(s) have requested or have clearly indicated an intent on Twitter to be contacted by you via Direct Message, for example by sending you a Direct Message; **and**

you provide a clear and easy way for such users to opt-out of receiving automated Direct Messages, and promptly honor all such opt-out requests.

The fact that a user is technically able to receive a Direct Message from you (e.g. because the user follows you, has enabled the ability to receive Direct Messages from any

account, or because the user is in a pre-existing Direct Message conversation with you) does not necessarily mean they have requested or expect to receive automated Direct Messages from you.

Interacting with users via Direct Message

Per the Ground Rules, remember to be thoughtful about the amount and type of information you request or exchange with users via Direct Messages. If you will be asking a user to provide personal or private information via an automated Direct Message, you must clearly explain how you will use the information you're collecting. Consider including a link to your privacy policy in your Direct Message to the user, as well as in your Twitter profile bio.

Don't publicly share information received in a Direct Message conversation with a user without first obtaining explicit consent from the user. For example, if a user asks you via Direct Message about a purchase they made from you, you may not mention the user in a public Tweet that includes information about their purchase unless you have the user's explicit consent to do so.

After a user-initiated interaction ends, don't send additional follow-up Direct Messages or mention users in a Tweet unless you get permission from the user.

D. Automated actions you take on Tweets or accounts

Automated likes: You may not like Tweets in an automated manner.

Automated Retweets: Provided you comply with all other rules, you may Retweet or Quote Tweet in an automated manner for entertainment, informational, or novelty purposes. Automated Retweets often lead to negative user experiences, and bulk, aggressive, or spammy Retweeting is a violation of the [Twitter Rules](#).

Automated following/unfollowing: You may not follow or unfollow Twitter accounts in a bulk, aggressive, or indiscriminate manner. Aggressive following is a violation of the [Twitter Rules](#). Please also review our [following rules and best practices](#) to ensure you are in compliance. Note that applications that claim to get users more followers are also prohibited under the [Twitter Rules](#).

Automated adding to lists or collections: You may not add Twitter users to lists or add Tweets to collections in a bulk or indiscriminate manner. Adding a large number of unrelated users to lists is a violation of the [Twitter Rules](#).

About search rules and restrictions

This article details the rules for using Twitter search, and explains reasons why content may not show up in Twitter search.

Why doesn't all content show up in Twitter search?

In order to keep your search results relevant, Twitter filters search results for quality Tweets and accounts. Material that jeopardizes search quality or creates a bad search experience for other people may be automatically removed from Twitter search. Read more about this [here](#).

General rules for using Twitter search

The [Twitter Rules](#) explain what behaviors are permitted on Twitter. In addition to these rules, we've included some tips below to keep your content relevant (and your Tweets in Twitter search).

Please do not:

Repeatedly post duplicate or near-duplicate content (links or Tweets). Abuse trending topics or hashtags (topic words with a # sign).
Send automated Tweets or replies.

Use bots or applications to post similar messages based on keywords. Post similar messages over multiple accounts.

Aggressively follow and unfollow people.

Twitter may automatically remove accounts engaging in these behaviors from search (or even suspend in some cases) in order to ensure the best experience for everyone.

Search results filtering

To learn how to filter your search results, including how to disable or enable safe search, read our [Using Twitter search](#) article.

Search results filtering in third-party applications

Accounts may be filtered from search for posting similar messages over several accounts. Your account may be filtered from search if you are logged into a third-party application that is updating many accounts with similar or duplicate content.

Rules for participating in contests on Twitter

Businesses sometimes host contests and offer prizes to Retweet an update, following a particular account, or posting Tweets with a specific hashtag topic or reply. If you're running a contest, check out this [article](#) for guidelines for contests on Twitter.

If you repeatedly post near-duplicate updates or duplicate links in order to enter contests, you may be filtered from search. Please keep in mind that you are responsible for the updates that you post to your Twitter account.

What to do if you are missing from search

Please see this [article](#) for more information related to this issue.

If you are being filtered from search, rest assured that your followers will still receive your updates. If you post a reply or mention to a specific person, it will be delivered to that person.

Having trouble?

You can learn more about searching by keyword, location, and trending topics in our [Using Twitter search](#) help article.

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. These limitations are set forth in the Twitter Rules below.

The Twitter Rules (along with all incorporated policies), [Privacy Policy](#), and [Terms of Service](#) collectively make up the "Twitter User Agreement" that governs a user's access to and use of Twitter's services.

All individuals accessing or using Twitter's services must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in Twitter taking one or more of the following enforcement actions:

- requiring you to delete prohibited content before you can again create new posts and interact with other Twitter users;
- temporarily limiting your ability to create posts or interact with other Twitter users;
- asking you to verify account ownership with a phone number or email address; or
- permanently suspending your account(s).

If you attempt to evade a permanent suspension by creating new accounts, we will suspend your new accounts.

Please note that we may need to change these Rules from time to time and reserve the right to do so. The most current version is always available at: <https://twitter.com/rules>.

The policies set forth in these Twitter Rules govern organic content on our platform. To learn more about the rules which govern ads and promoted content, please review our [Ads policies](#).

Content Boundaries and Use of Twitter

Intellectual property

Trademark: We reserve the right to suspend accounts or take other appropriate action when someone's brand or trademark, including business name and/or logo, is used in a manner that may mislead or confuse others about your brand affiliation. Read more about our [trademark policy](#) and [how to report a violation](#).

Copyright: We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in our [Terms of Service](#). Read more about our [copyright policy](#).

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Learn more about our [media policy](#).

Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images. Additionally, Twitter may sometimes require you to remove excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about [how to make such a request](#), and [how to mark your media as sensitive](#).

Unlawful use

You may not use our service for any unlawful purposes or in furtherance of illegal activities. By using Twitter, you agree to comply with all applicable laws governing your online conduct and content.

Trends

At times, we may prevent certain content from trending. This includes content that violates the Twitter Rules, as well as content that may attempt to manipulate trends. Read more about [what we allow and do not allow to trend](#).

Misuse of Twitter badges

You may not use badges, including but not limited to the “promoted” or “verified” Twitter badges, unless provided by Twitter. Accounts using unauthorized badges as part of their profile photos, header photos, display names, or in any way that falsely implies affiliation with Twitter or authorization from Twitter to display these badges, may be suspended.

Misuse of usernames

Selling usernames: You may not buy or sell Twitter usernames.

Username squatting: You may not engage in username squatting. Some of the factors we take into consideration when determining whether conduct is username squatting include:

- the number of accounts created;
- the creation of accounts for the purpose of preventing others from using those account names;
- the creation of accounts for the purpose of selling those accounts; and
- the use of third-party content feeds to update and maintain accounts under the names of those third parties.

Please note that Twitter may also remove accounts that are inactive for more than six months. Learn more about [username squatting](#).

Abusive Behavior

We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we prohibit behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Context matters when evaluating for abusive behavior and determining appropriate enforcement actions. Factors we may take into consideration include, but are not limited to whether:

- the behavior is targeted at an individual or group of people;
- the report has been filed by the target of the abuse or a bystander;
- the behavior is newsworthy and in the legitimate public interest.

Violence and physical harm

Violence: You may not make specific threats of violence or wish for the serious physical harm, death, or disease of an individual or group of people. This includes, but is not limited to, threatening or promoting terrorism. You also may not affiliate with organizations that — whether by their own statements or activity both on and off the platform — use or promote violence against civilians to further their causes. We will begin enforcing this rule around affiliation with such organizations on December 18, 2017.

Suicide or self-harm: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Child sexual exploitation: You may not promote child sexual exploitation. Learn more about our zero-tolerance [child sexual exploitation policy](#).

Abuse and hateful conduct

Abuse: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Unwanted sexual advances: You may not direct abuse at someone by sending unwanted sexual content, objectifying them in a sexually explicit manner, or otherwise engaging in sexual misconduct.

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. Read more about our [hateful conduct policy](#).

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category. We will begin enforcing this rule on December 18, 2017.

Private information and intimate media

Private information: You may not publish or post other people's private information without their express authorization and permission. Definitions of private information may vary depending on local laws. Read more about our [private information policy](#).

Intimate media: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. Read more about intimate media on Twitter.

Threats to expose / hack: You may not threaten to expose someone's private information or intimate media. You also may not threaten to hack or break into someone's digital information.

Impersonation

You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. While you may maintain parody, fan, commentary, or newsfeed accounts, you may not do so if the intent of the account is to engage in spamming or abusive behavior. Read more about our [impersonation policy](#).

Spam and Security

We strive to protect people on Twitter from technical abuse and spam.

To promote a stable and secure environment on Twitter, you may not do, or attempt to do, any of the following while accessing or using Twitter:

- Access, tamper with, or use non-public areas of Twitter, Twitter's computer systems, or the technical delivery systems of Twitter's providers (except as expressly permitted by the Twitter Bug Bounty program).
- Probe, scan, or test the vulnerability of any system or network, or breach or circumvent any security or authentication measures (except as expressly permitted by the Twitter Bug Bounty program).
- Access or search, or attempt to access or search, Twitter by any means (automated or otherwise) other than through our currently available, published interfaces that are provided by Twitter (and only pursuant to the applicable terms and conditions), unless you have been specifically allowed to do so in a separate agreement with Twitter. Note that crawling Twitter is permissible if done in accordance with the provisions of the robots.txt file; however, scraping Twitter without our prior consent is expressly prohibited.
- Forge any TCP/IP packet header or any part of the header information in any email or posting, or in any way use Twitter to send altered, deceptive, or false source-identifying information.
- Interfere with or disrupt the access of any user, host or network, including, without limitation, sending a virus, overloading, flooding, spamming, mail-bombing Twitter's services, or by scripting the creation of content in such a manner as to interfere with or create an undue burden on Twitter.

Any accounts engaging in the following activities may be temporarily locked or subject to permanent suspension:

- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use Twitter's services for the purpose of spamming anyone. Spam is generally defined on Twitter as bulk or aggressive activity that attempts to manipulate or disrupt Twitter or the experience of users on Twitter to drive traffic or attention to unrelated accounts, products, services, or initiatives. Some of the factors that we take into account when determining what conduct is considered to be spamming include:
 - if you have followed and/or unfollowed a large number of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if your Tweets or Direct Messages consist mainly of links shared without commentary;
 - if a large number of people have blocked you in response to high volumes of untargeted, unsolicited, or duplicative content or engagements from your account;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicative or substantially similar content, replies, or mentions over multiple accounts or multiple duplicate updates on one account, or create duplicate or substantially similar accounts;
 - if you post multiple updates to a trending or popular topic with an intent to subvert or manipulate the topic to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you send large numbers of unsolicited replies or mentions;
 - if you add users to lists in a bulk or aggressive manner;
 - if you are randomly or aggressively engaging with Tweets (e.g., likes, Retweets, etc.) or users (e.g., following, adding to lists or Moments, etc.) to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you repeatedly post other people's account information as your own (e.g., bio, Tweets, profile URL, etc.);
 - if you post misleading, deceptive, or malicious links (e.g., affiliate links, links to malware/clickjacking pages, etc.);
 - if you create fake accounts, account interactions, or impressions;
 - if you sell, purchase, or attempt to artificially inflate account interactions (such as followers, Retweets, likes, etc.); and
 - if you use or promote third-party services or apps that claim to get you more followers, Retweets, or likes (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers or engagements to your account or Tweets).

Please see our support articles on [following rules and best practices](#) and [automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts may be permanently suspended.

Content Visibility

Accounts under investigation or which have been detected as sharing content in violation of these Rules may have their account or Tweet visibility limited in various parts of Twitter, including search. To learn more about situations in which content may be restricted on Twitter, please see our support article on [search rules and restrictions](#).

Trademark policy

Updated June 30, 2016

What is a Trademark Policy Violation on Twitter?

Using another's trademark in a manner that may mislead or confuse others about your brand affiliation may be a violation of our trademark policy.

How Does Twitter Respond to Trademark Reports?

We review each report and do the following:

- If we determine that there is a violation of our trademark policy, we may suspend the account. In some instances, we may give the account holder an opportunity to comply with our policies before suspending the account. For more information, see our [Parody, commentary, and fan account policy](#).
- We may also provide the account holder with your name and information included in the copy of the report.
- If you would like to report a trademark violation in a [Twitter Ad](#), please visit our [Trademark Policy for Promoted Ads](#) to learn more.

Please note that Twitter only investigates requests that are submitted by the trademark holder or by an authorized representative of the trademark holder.

What is not a Trademark Policy Violation?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples include:

- Using a trademark in a way that is outside the scope of the trademark registration (e.g. territory, or goods and services identified in the registration).
- Nominative and other fair uses of trademarks are protected uses under our trademark policy, so long as the account is clearly distinguished from the trademark owner. This includes use by resellers in certain regions and accounts engaging in parody, commentary, or news. For more information, see our [Parody, commentary, and fan account policy](#).

How do I Report a Trademark Policy Violation?

Submit trademark reports through our [Trademark Issue support form](#). Using the appropriate form ensures we have all the information necessary to process your report quickly. Please provide all information requested in the form. If you submit a report with incomplete information, we will follow up with you to request that information. Please note that this will result in a delay in processing your report.

If you have trouble using the form, try updating your browser or switching to another browser.

Copyright policy

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged

infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notification.

When Should I File a Counter-notice?

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);

3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Twitter media policy

Twitter Rules: Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images.

Twitter Rules: You may not use hateful images or symbols in your profile image or profile header.

Rationale

Some forms of graphic violence, adult content, or hateful imagery may be permitted in Tweets when they are marked as [sensitive media](#). However, you may not include this type of content in live video, or in profile or header images.

While we want people to feel free to share media that reflects their creativity or individuality, or to show what's happening in the world, we will take action when it crosses the line into abuse towards a person, group, or protected category.

Finally, we may sometimes require you to remove media containing excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about how to make such a request [here](#).

When this applies

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. Some examples include, but are not limited to, depictions of:

- the moment at which someone dies
- gruesome crime or accident scenes
- bodily harm, torture, dismemberment, or mutilation

We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity (including close-ups of genitals, buttocks, or breasts)
 - Please note that exceptions may be made for artistic, medical, health, or educational content. Breastfeeding content does not need to be marked as sensitive.
- simulating a sexual act
- intercourse or any sexual act (may involve humans, humanoid animals, cartoons, or anime)

If you see media you think may be in violation of our policy, report it. You can read about how to report sensitive media [here](#).

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, or ethnicity/national origin.

Some examples of hateful imagery include, but are not limited to:

- Symbols historically associated with hate groups (for example, the Nazi swastika)
- Images depicting others as less than human or photoshopped to include hateful symbols
- Photoshopped references to a mass murder that targeted a protected category

Do I need to be the target of this content in order for it to be a violation of the Twitter Rules?
No, we review both first-person and bystander reports of such content.

Consequences

When adult content, graphic violence, or hateful imagery appears in Tweets, we may place this content behind an interstitial advising viewers to be aware that they will see sensitive media if they click through. This allows us to identify potentially sensitive content that some people may not wish to see. Learn more about [how to control whether you see sensitive media](#).

However, if such content appears in live video, header, or profile images, it will violate our media policy. The consequences for violating this policy depend on the severity of the violation and the person's previous record of violations.

The first time someone violates this policy they will be required to remove the imagery and may be asked to verify their account ownership. Subsequent violations could result in permanent suspension. Learn more about [our range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

How to contact Twitter about media concerning a deceased family member

In order to respect the wishes of loved ones, Twitter will remove imagery of deceased individuals in certain circumstances. Immediate family members and other authorized individuals may request the removal of images

or video of deceased individuals, from when critical injury occurs to the moments before or after death, by sending a request to Twitter Inc. via our [privacy form](#). When reviewing such media removal requests, Twitter considers public interest factors such as the newsworthiness of the content and may not be able to honor every request. Note: Find information about contacting Twitter to [deactivate the account of a deceased or incapacitated user](#).

How to mark media as sensitive links to Media policy

FAQs about trends on Twitter

Where can I find trends?

Trends are available on the Twitter for iOS app, the Twitter for Android app, and [twitter.com](#).

On Twitter's mobile apps, you can find trends listed under the trends section



of the **Explore** tab. When signed in to [twitter.com](#) on a desktop or laptop computer, trends are listed in many places, including the Home timeline, Notifications, search results, and profile pages.

How are trends determined?

Trends are determined by an algorithm and, by default, are tailored for you based on who you follow, your interests, and your location. This algorithm identifies topics that are popular now, rather than topics that have been popular for a while or on a daily basis, to help you discover the hottest emerging topics of discussion on Twitter.

Note: The number of Tweets that are related to the trends is just one of the factors the algorithm looks at when ranking and determining trends. Algorithmically, trends and hashtags are grouped together if they are related to the same topic. For instance, #MondayMotivation and #MotivationMonday may both be represented by #MondayMotivation.

You can choose to see trends that are not tailored for you by selecting a [specific trends location](#) on [twitter.com](#), iOS, or Android. Location trends identify popular topics among people in a specific geographic location.

What does the # sign mean?

You may notice that some trends have # sign before the word or phrase. This is called a [hashtag](#) and is included specifically in Tweets to mark them as relating to a topic, so that people can follow the conversation in search.

Will I see world and local events in my trends?

Yes, when signed into your account on the [twitter.com](#), iOS, or Android, trends will be tailored for you based on your location and who you follow. There will be many world and local news events and conversations that will appear in your trends regardless of your personalization.

Can I see trends for a specific location?

Yes, you can view trends for a specific location on twitter.com, iOS, or Android.

View trends via twitter.com:

To view trends for a specific location, click **Change** in your trends box. You can choose to **Keep tailored trends** based on your location and who you follow, or you can select **Change** to choose a nearby location or search locations.

Adjust your trends settings via Twitter for iOS:

1. In the top menu, tap your **profile** icon.
2. Tap **Settings and privacy**, then tap **Content preferences**.
3. Under **Explore**, tap **Trends**.
4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.
5. When the feature is disabled, you can change your location by tapping **Change location**.
 1. In the search box, type your desired trends location or choose from the list.

Adjust your trends settings via Twitter for Android:

1. In the top menu, you will either see a **navigation menu** icon  or your **profile** icon. Tap whichever icon you have, and select **Settings and privacy**.
2. Tap **Content preferences**.
3. Under **Explore**, tap **Trends**.
4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.
5. When the feature is disabled, you can change your location by tapping **Change location**.
 1. In the search box, type your desired trends location or choose from the list.

Note: If you don't find your preferred country or city, it means we're not receiving enough Tweets from that geographical area to create a list. You can look up local Tweets on any topic by using [advanced Twitter Search operators](#).

What happens when I click or tap on a trend?

Clicking or tapping any of the trends takes you to the Twitter search results for that trend. You'll see all Tweets including that phrase or hashtag. To see what people are saying about a previous trend, perform a [search for that keyword](#).

How can I participate in a trend?

Simply post a Tweet including the exact word or phrase as it appears in the trends list (with the hashtag, if you see one). Due to the large number of people Tweeting about these specific trends, you may not always be able to find your particular Tweet in search, but your followers will always see your Tweets.

Note: Twitter also filters searches for quality. To understand what types of behavior could cause your Tweets to be filtered out, read about our [search rules and restrictions](#).

Are there rules for trends?

Yes. We want trends to promote healthy discussions on Twitter. This means that at times, we may prevent certain content from trending. These include trends that:

- Contain profanity or adult/graphic references.
- Incite hate on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease.
- Violate the [Twitter Rules](#).

In some cases, we may also consider the newsworthiness of the content, or if it is in the public interest when evaluating potential violations. In these cases, the content might continue to trend on our platform.

Note that even if we prevent the hashtag or the content from appearing on the list of trending topics, you may still be able to access conversations around that content on Twitter.

Trends descriptions

Below some trends, you may see the approximate number of Tweets associated with the trend, or personalized context such as who in your network is Tweeting about it.

Additionally, when a Moment is found to be relevant to a trend, we may use the title of the Moment as the trend's description.

Having trouble?

Check out our [troubleshooting section](#) for solutions to common problems.

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Child sexual exploitation policy

Twitter removes content promoting child sexual exploitation

We do not tolerate child sexual exploitation on Twitter. When we are made aware of links to images of or content promoting child sexual exploitation they will be removed from the site without further notice and reported to The National Center for Missing & Exploited Children ("NCMEC"); we permanently suspend accounts promoting or containing updates with links to child sexual exploitation.

Reporting content promoting child sexual exploitation

If you think you've found a Twitter account distributing or promoting child sexual exploitation while using Twitter, please click [here](#) to report it. Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. To find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with the National Center for Missing and Exploited Children [here](#), or your local law enforcement. If you believe that the content was hosted on a website outside of the United States, you can report it [here](#) on the International Association of Internet Hotlines website.

NOTE: Please do NOT tweet, retweet or repost child sexual exploitation for any reason. Click [here](#) to report it to us immediately and we will take steps to remove it.

Law Enforcement Agencies

If you are a member of a law enforcement agency seeking information regarding Twitter accounts promoting or distributing child sexual exploitation, please see our [Guidelines for Law Enforcement](#) for information on our process; we also suggest you contact the National Center for Missing and Exploited Children (NCMEC) [here](#), regardless of jurisdiction, to determine whether or not information on those accounts is already on file.

For general law enforcement inquiries, please see our [Guidelines for Law Enforcement](#) for more information.

Hateful conduct policy

Freedom of expression means little if voices are silenced because people are afraid to speak up. We do not tolerate behavior that harasses, intimidates, or uses fear to silence another person's voice. If you see something on Twitter that violates these rules, please report it to us.

How our policy works

As explained in the Twitter Rules,

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Examples of what we do not tolerate includes, but is not limited to behavior that harasses individuals or groups of people with:

violent threats;

wishes for the physical harm, death, or disease of individuals or groups;

references to mass murder, violent events, or specific means of violence in which/with which such groups have been the primary targets or victims;

behavior that incites fear about a protected group;

repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone.

How our enforcement works Context matters.

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. While we accept reports of violations from anyone, sometimes we also need to hear directly from the target to ensure that we have proper context.

The number of reports we receive does not impact whether or not something will be removed. However, it may help us prioritize the order in which it gets reviewed.

We focus on behavior.

We enforce policies when someone reports behavior that is abusive and targets an entire protected group and/or individuals who may be members.

This targeting can happen in any manner (for example, @mentions, tagging a photo, and more).

We have a range of enforcement options.

The consequences for violating our rules vary depending on the severity of the violation and the person's previous record of violations. For example, we may ask someone to remove the offending Tweet before they can Tweet again. For other cases, we may suspend an account.

Report private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Read about our [policy regarding private information posted on Twitter](#).

Do I need to have a Twitter account to make a report?

No, you do not need to have a Twitter account to file a report that your private and confidential information has been posted on Twitter.

Can I report unauthorized posting of Copyrighted material?

If your copyrighted material has been Tweeted without your permission, please review our [Copyright and DMCA policy](#).

What can I do if my information is posted on other websites?

We cannot take down any information that is posted on a website other than Twitter. If your information is posted on another website, please contact the help center or support staff for the website where your information is posted.

How do I report that my private information is posted on Twitter?

You can report that your private information has been posted on Twitter directly from the offending Tweet(s). Learn [how to report individual Tweets](#).

File a [report that someone has posted your private information](#).

What happens when Twitter receives a valid report?

When we receive a complete and valid report that private information has been posted on Twitter, we'll investigate the account and Tweets reported. We will review where, if anywhere, the information has been made publicly available before taking action on the account or Tweets. If the information you reported was previously posted elsewhere on the Internet, it is not a violation of our policy and we will not take action.

Learn more about [protecting your private information on Twitter and other websites](#).

Read our [policy regarding private information posted on Twitter](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter impersonation policy.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to Twitter's [parody, commentary, and fan account policy](#) for more information about these types of accounts.

Find more information about [reporting impersonation on Twitter](#).

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation rules

Updated November 3, 2017

This page is primarily intended for **developers**.

For Twitter users: You are ultimately responsible for the actions taken with your account, or by applications associated with your account. Before authorizing a third-party application to access or use your account, make sure you've thoroughly investigated the application and understand what it will do. If automated activity on your account violates the [Twitter Rules](#) or these automation rules, Twitter may take action on your account, including [filtering your Tweets from search results](#) or suspending your account.

For more information on third-party applications, please see our article on [connecting and revoking third-party applications](#).

If you are a developer and have questions about these Automation Rules, please reach out via <https://twittercommunity.com/c/rules-and-policies>.

If your application has been suspended or restricted from performing write actions, you may submit an appeal at </forms/platform>.

I. Ground Rules Do!

Build solutions that automatically broadcast helpful information in Tweets.

Run creative campaigns that auto-reply to users who engage with your content.

Build solutions that automatically respond to users in Direct Messages.

Try new things that help people (and comply with our rules).

Make sure your application provides a good user experience and performs well — and confirm that remains the case over time.

Don't!

Violate these or other policies. Be extra mindful of our rules about abuse and user privacy.

Abuse the Twitter API or attempt to circumvent rate limits.

Use non-API-based forms of automation, such as scripting the Twitter website. The use of these techniques may result in the permanent suspension of your account.

Spam or bother users, or otherwise send them unsolicited messages.

A. The Twitter Rules and the Developer Agreement and Policy

As with all activity on Twitter, automated activity is subject to the [Twitter Rules](#) and, if you're a developer using the Twitter API, the [Developer Agreement and Policy](#).

You should carefully review these policies to ensure that your automated activity is compliant. Automated applications or activities that violate these policies, or that facilitate or induce users to violate them, may be subject to enforcement action, potentially including suspension of associated Twitter accounts.

We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these policies.

Although all aspects of the Twitter Rules and the Developer Agreement and Policy apply to automated activity, you should keep the following rules top of mind:

Spamming: You may not send automated Tweets or Direct Messages that are spam, or otherwise engage in spamming activity. Some examples of spammy behavior to avoid with automation include:

Trending topics: You may not automatically post about trending topics on Twitter, or use automation to attempt to influence or manipulate trending topics.

Multiple posts/accounts: You may not post duplicative or substantially similar Tweets on one account or over multiple accounts you operate.

Duplicate accounts: You may not create and/or automate multiple accounts for duplicative or substantially similar use cases.

However, automating multiple accounts for related but non-duplicative use cases is permitted. For example, you may automate separate accounts to Tweet when the Hubble Space Telescope passes over different cities, such as [San Francisco](#) or [Hong Kong](#).

Misleading links: You may not send automated Tweets or Direct Messages containing links that are misleading, including links that maliciously or deceptively redirect through landing pages or ad pages before displaying the final content.

Sensitive media: Automated Tweets and Direct Messages must comply with the [Twitter media policy](#), and you should mark your account as potentially sensitive if you intend to post graphic, pornographic, or potentially sensitive media.

Abusive behavior: You may not engage in any automated activity that encourages, promotes, or incites abuse, violence, hateful conduct, or harassment, on or off Twitter.

Private information: You may not post private or confidential information about a person without their prior express authorization.

B. Other Ground Rules for Automated Activity

In addition to the policies above, the following ground rules apply to all automated activity on Twitter:

Don't surprise or mislead users: Automated activity should honor users' expectations. Ask for the user's permission before taking an action if you aren't sure.

Mature content or profanity: Don't Direct Message, mention, or reply to users with potentially sensitive content (including profanity), unless they've clearly indicated an intent to receive it in advance.

Be thoughtful about the information you request or exchange on Twitter

Tweets: Don't ask users to send you personal or private information via a public Tweet. If you need additional personal or private information from a user to provide them with customer service (or other similar use cases), you should ask the user to share such information by Direct Message or another private channel. You might even consider adding a [Direct Message deep link](#) to your Tweet.

Direct Messages: You should only ask users for the minimum amount of information you need to provide them with service. If you need to request or exchange particularly sensitive information (such as credit card information), you should consider directing users to your website or other appropriate channel to do so.

II. Activity-Specific Rules

The activity-specific rules in this section apply to taking specific automated actions on Twitter. Please read these rules carefully, as they outline both permitted and prohibited use cases of automation.

Automated applications or activities that violate these rules, or that facilitate or induce users to violate them, may be subject to enforcement action, including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these rules. As a reminder, you should also carefully review the spam guidelines in the [Twitter Rules](#) to avoid having activities performed by you, your app, or other users through your app or service flagged as spam.

A. Automated Actions Through Another User's Account

Twitter users may authorize your app or service to [access their Twitter account through OAuth](#). A user authorizing your app or service to access their Twitter account through OAuth does not by itself constitute sufficient consent to take automated actions through that user's account.

You may only take automated actions through another Twitter user's account if you:

clearly describe to the user the types of automated actions that will occur; receive express consent from the user to take those automated actions; and immediately honor a user's request to opt-out of further automated actions.

If you substantially change the purpose or functionality of your app or service, you must re-obtain express consent from the user to take automated action through their account before doing so.

These requirements apply to any automated action taken through another Twitter user's account, including posting Tweets, sending Direct Messages, deleting Tweets or Direct Messages, or following/unfollowing other accounts. For applications that offer users the ability to delete Tweets in a bulk or automated manner, you must also clearly state that Tweets are not recoverable once deleted.

B. Automated Tweets

1. Posting automated Tweets

Automated Tweets that cross-post outside information:

You may post automated Tweets based on sources of outside information — such as an RSS feed, weather data, etc. — as long as you are sufficiently authorized to publish such information.

Other automated Tweets (excluding mentions or replies):

Provided you comply with all other rules, you may post automated Tweets for entertainment, informational, or novelty

purposes. As a reminder, accounts posting duplicative, spammy, or otherwise prohibited content may be subject to suspension.

2. Posting automated mentions and replies

The reply and mention functions are intended to make communication between Twitter users easier. Automating these actions to reach many users on an unsolicited basis is an abuse of the feature, and is not permitted. For example, sending automated replies to Tweets based on keyword searches alone is not permitted. Spamy or duplicative use of mentions and replies may result in enforcement action, such as the removal of your Tweets from Search or the suspension of your app or account.

However, you may send automated replies or mentions to Twitter users so long as:

in advance of sending the automated reply, the recipient or mentioned user(s) have requested or have clearly indicated an intent on Twitter to be contacted by you (i.e. opted in), for example by replying to a Tweet from your account, or by sending you a Direct Message;

you provide a clear and easy way for such users to opt-out of receiving automated replies and mentions, and promptly honor all such opt-out requests;

you only send one automated reply or mention per user interaction; and

the automated reply or mention is a reply to the user's original Tweet (if your campaign is based on users posting a reply to your Tweet).

Opt-in techniques and indications of user intent take many different forms, depending on the specifics of your use case and implementation. Some examples include:

A Tweet from your account that clearly indicates that a user taking a specific action on that Tweet (such as Retweeting it) will opt the user into receiving an automated response.

A mention of your account by the user in a manner suggesting the user clearly wishes or intends to receive a response. If you want to run an auto-reply campaign with a campaign- or use-case-specific hashtag, users should also mention you in their Tweets.

Note that a user following your account is not on its own a sufficient indication of user intent to receive an automated response.

: Advertisers, publishers, and brands using auto-response campaigns must request approval from Twitter and may be subject to additional rules. Please reach out to your account or partner manager for assistance. If you do not have a partner manager but are interested in managed client services, you can learn more [here](#).

Additionally, we recommend that any accounts that will communicate with

users via automated mentions or replies:

Note

1. Appropriately filter responses based on potentially sensitive language in user handles,

display names, and Tweet text, as well as potentially sensitive media;

2. Check that the Tweet you are mentioning or replying to still exists (for example, using the statuses/lookup endpoint on the Twitter API).

C. Automated Direct Messages

Sending automated Direct Messages to users

You may not send unsolicited Direct Messages in a bulk or automated manner, and should be thoughtful about the frequency with which you contact users via Direct Message.

You may send automated Direct Messages to users so long as:

in advance of sending the Direct Message, the recipient(s) have requested or have clearly indicated an intent on Twitter to be contacted by you via Direct Message, for example by sending you a Direct Message; **and**

you provide a clear and easy way for such users to opt-out of receiving automated Direct Messages, and promptly honor all such opt-out requests.

The fact that a user is technically able to receive a Direct Message from you (e.g. because the user follows you, has enabled the ability to receive Direct Messages from any

account, or because the user is in a pre-existing Direct Message conversation with you) does not necessarily mean they have requested or expect to receive automated Direct Messages from you.

Interacting with users via Direct Message

Per the Ground Rules, remember to be thoughtful about the amount and type of information you request or exchange with users via Direct Messages. If you will be asking a user to provide personal or private information via an automated Direct Message, you must clearly explain how you will use the information you're collecting. Consider including a link to your privacy policy in your Direct Message to the user, as well as in your Twitter profile bio.

Don't publicly share information received in a Direct Message conversation with a user without first obtaining explicit consent from the user. For example, if a user asks you via Direct Message about a purchase they made from you, you may not mention the user in a public Tweet that includes information about their purchase unless you have the user's explicit consent to do so.

After a user-initiated interaction ends, don't send additional follow-up Direct Messages or mention users in a Tweet unless you get permission from the user.

D. Automated actions you take on Tweets or accounts

Automated likes: You may not like Tweets in an automated manner.

Automated Retweets: Provided you comply with all other rules, you may Retweet or Quote Tweet in an automated manner for entertainment, informational, or novelty purposes. Automated Retweets often lead to negative user experiences, and bulk, aggressive, or spammy Retweeting is a violation of the [Twitter Rules](#).

Automated following/unfollowing: You may not follow or unfollow Twitter accounts in a bulk, aggressive, or indiscriminate manner. Aggressive following is a violation of the [Twitter Rules](#). Please also review our [following rules and best practices](#) to ensure you are in compliance. Note that applications that claim to get users more followers are also prohibited under the [Twitter Rules](#).

Automated adding to lists or collections: You may not add Twitter users to lists or add Tweets to collections in a bulk or indiscriminate manner. Adding a large number of unrelated users to lists is a violation of the [Twitter Rules](#).

About search rules and restrictions

This article details the rules for using Twitter search, and explains reasons why content may not show up in Twitter search.

Why doesn't all content show up in Twitter search?

In order to keep your search results relevant, Twitter filters search results for quality Tweets and accounts. Material that jeopardizes search quality or creates a bad search experience for other people may be automatically removed from Twitter search. Read more about this [here](#).

General rules for using Twitter search

The [Twitter Rules](#) explain what behaviors are permitted on Twitter. In addition to these rules, we've included some tips below to keep your content relevant (and your Tweets in Twitter search).

Please do not:

Repeatedly post duplicate or near-duplicate content (links or Tweets). Abuse trending topics or hashtags (topic words with a # sign).
Send automated Tweets or replies.

Use bots or applications to post similar messages based on keywords. Post similar messages over multiple accounts.

Aggressively follow and unfollow people.

Twitter may automatically remove accounts engaging in these behaviors from search (or even suspend in some cases) in order to ensure the best experience for everyone.

Search results filtering

To learn how to filter your search results, including how to disable or enable safe search, read our [Using Twitter search](#) article.

Search results filtering in third-party applications

Accounts may be filtered from search for posting similar messages over several accounts. Your account may be filtered from search if you are logged into a third-party application that is updating many accounts with similar or duplicate content.

Rules for participating in contests on Twitter

Businesses sometimes host contests and offer prizes to Retweet an update, following a particular account, or posting Tweets with a specific hashtag topic or reply. If you're running a contest, check out this [article](#) for guidelines for contests on Twitter.

If you repeatedly post near-duplicate updates or duplicate links in order to enter contests, you may be filtered from search. Please keep in mind that you are responsible for the updates that you post to your Twitter account.

What to do if you are missing from search

Please see this [article](#) for more information related to this issue.

If you are being filtered from search, rest assured that your followers will still receive your updates. If you post a reply or mention to a specific person, it will be delivered to that person.

Having trouble?

You can learn more about searching by keyword, location, and trending topics in our [Using Twitter search](#) help article.

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. These limitations are set forth in the Twitter Rules below.

The Twitter Rules (along with all incorporated policies), [Privacy Policy](#), and [Terms of Service](#) collectively make up the "Twitter User Agreement" that governs a user's access to and use of Twitter's services.

All individuals accessing or using Twitter's services must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in Twitter taking one or more of the following enforcement actions:

- requiring you to delete prohibited content before you can again create new posts and interact with other Twitter users;
- temporarily limiting your ability to create posts or interact with other Twitter users;
- asking you to verify account ownership with a phone number or email address; or
- permanently suspending your account(s).

If you attempt to evade a permanent suspension by creating new accounts, we will suspend your new accounts.

Please note that we may need to change these Rules from time to time and reserve the right to do so. The most current version is always available at: <https://twitter.com/rules>.

The policies set forth in these Twitter Rules govern organic content on our platform. To learn more about the rules which govern ads and promoted content, please review our [Ads policies](#).

Content Boundaries and Use of Twitter

Intellectual property

Trademark: We reserve the right to suspend accounts or take other appropriate action when someone's brand or trademark, including business name and/or logo, is used in a manner that may mislead or confuse others about your brand affiliation. Read more about our [trademark policy](#) and [how to report a violation](#).

Copyright: We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in our [Terms of Service](#). Read more about our [copyright policy](#).

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Learn more about our [media policy](#).

Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images. Additionally, Twitter may sometimes require you to remove excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about [how to make such a request](#), and [how to mark your media as sensitive](#).

Unlawful use

You may not use our service for any unlawful purposes or in furtherance of illegal activities. By using Twitter, you agree to comply with all applicable laws governing your online conduct and content.

Trends

At times, we may prevent certain content from trending. This includes content that violates the Twitter Rules, as well as content that may attempt to manipulate trends. Read more about [what we allow and do not allow to trend](#).

Misuse of Twitter badges

You may not use badges, including but not limited to the “promoted” or “verified” Twitter badges, unless provided by Twitter. Accounts using unauthorized badges as part of their profile photos, header photos, display names, or in any way that falsely implies affiliation with Twitter or authorization from Twitter to display these badges, may be suspended.

Misuse of usernames

Selling usernames: You may not buy or sell Twitter usernames.

Username squatting: You may not engage in username squatting. Some of the factors we take into consideration when determining whether conduct is username squatting include:

- the number of accounts created;
- the creation of accounts for the purpose of preventing others from using those account names;
- the creation of accounts for the purpose of selling those accounts; and
- the use of third-party content feeds to update and maintain accounts under the names of those third parties.

Please note that Twitter may also remove accounts that are inactive for more than six months. Learn more about [username squatting](#).

Abusive Behavior

We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we prohibit behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Context matters when evaluating for abusive behavior and determining appropriate enforcement actions. Factors we may take into consideration include, but are not limited to whether:

- the behavior is targeted at an individual or group of people;
- the report has been filed by the target of the abuse or a bystander;
- the behavior is newsworthy and in the legitimate public interest.

Violence and physical harm

Violence: You may not make specific threats of violence or wish for the serious physical harm, death, or disease of an individual or group of people. This includes, but is not limited to, threatening or promoting terrorism. You also may not affiliate with organizations that — whether by their own statements or activity both on and off the platform — use or promote violence against civilians to further their causes. We will begin enforcing this rule around affiliation with such organizations on December 18, 2017.

Suicide or self-harm: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Child sexual exploitation: You may not promote child sexual exploitation. Learn more about our zero-tolerance [child sexual exploitation policy](#).

Abuse and hateful conduct

Abuse: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Unwanted sexual advances: You may not direct abuse at someone by sending unwanted sexual content, objectifying them in a sexually explicit manner, or otherwise engaging in sexual misconduct.

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. Read more about our [hateful conduct policy](#).

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category. We will begin enforcing this rule on December 18, 2017.

Private information and intimate media

Private information: You may not publish or post other people's private information without their express authorization and permission. Definitions of private information may vary depending on local laws. Read more about our [private information policy](#).

Intimate media: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. Read more about intimate media on Twitter.

Threats to expose / hack: You may not threaten to expose someone's private information or intimate media. You also may not threaten to hack or break into someone's digital information.

Impersonation

You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. While you may maintain parody, fan, commentary, or newsfeed accounts, you may not do so if the intent of the account is to engage in spamming or abusive behavior. Read more about our [impersonation policy](#).

Spam and Security

We strive to protect people on Twitter from technical abuse and spam.

To promote a stable and secure environment on Twitter, you may not do, or attempt to do, any of the following while accessing or using Twitter:

- Access, tamper with, or use non-public areas of Twitter, Twitter's computer systems, or the technical delivery systems of Twitter's providers (except as expressly permitted by the Twitter Bug Bounty program).
- Probe, scan, or test the vulnerability of any system or network, or breach or circumvent any security or authentication measures (except as expressly permitted by the Twitter Bug Bounty program).
- Access or search, or attempt to access or search, Twitter by any means (automated or otherwise) other than through our currently available, published interfaces that are provided by Twitter (and only pursuant to the applicable terms and conditions), unless you have been specifically allowed to do so in a separate agreement with Twitter. Note that crawling Twitter is permissible if done in accordance with the provisions of the robots.txt file; however, scraping Twitter without our prior consent is expressly prohibited.
- Forge any TCP/IP packet header or any part of the header information in any email or posting, or in any way use Twitter to send altered, deceptive, or false source-identifying information.
- Interfere with or disrupt the access of any user, host or network, including, without limitation, sending a virus, overloading, flooding, spamming, mail-bombing Twitter's services, or by scripting the creation of content in such a manner as to interfere with or create an undue burden on Twitter.

Any accounts engaging in the following activities may be temporarily locked or subject to permanent suspension:

- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use Twitter's services for the purpose of spamming anyone. Spam is generally defined on Twitter as bulk or aggressive activity that attempts to manipulate or disrupt Twitter or the experience of users on Twitter to drive traffic or attention to unrelated accounts, products, services, or initiatives. Some of the factors that we take into account when determining what conduct is considered to be spamming include:
 - if you have followed and/or unfollowed a large number of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if your Tweets or Direct Messages consist mainly of links shared without commentary;
 - if a large number of people have blocked you in response to high volumes of untargeted, unsolicited, or duplicative content or engagements from your account;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicative or substantially similar content, replies, or mentions over multiple accounts or multiple duplicate updates on one account, or create duplicate or substantially similar accounts;
 - if you post multiple updates to a trending or popular topic with an intent to subvert or manipulate the topic to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you send large numbers of unsolicited replies or mentions;
 - if you add users to lists in a bulk or aggressive manner;
 - if you are randomly or aggressively engaging with Tweets (e.g., likes, Retweets, etc.) or users (e.g., following, adding to lists or Moments, etc.) to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you repeatedly post other people's account information as your own (e.g., bio, Tweets, profile URL, etc.);
 - if you post misleading, deceptive, or malicious links (e.g., affiliate links, links to malware/clickjacking pages, etc.);
 - if you create fake accounts, account interactions, or impressions;
 - if you sell, purchase, or attempt to artificially inflate account interactions (such as followers, Retweets, likes, etc.); and
 - if you use or promote third-party services or apps that claim to get you more followers, Retweets, or likes (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers or engagements to your account or Tweets).

Please see our support articles on [following rules and best practices](#) and [automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts may be permanently suspended.

Content Visibility

Accounts under investigation or which have been detected as sharing content in violation of these Rules may have their account or Tweet visibility limited in various parts of Twitter, including search. To learn more about situations in which content may be restricted on Twitter, please see our support article on [search rules and restrictions](#).

Trademark policy

Updated June 30, 2016

What is a Trademark Policy Violation on Twitter?

Using another's trademark in a manner that may mislead or confuse others about your brand affiliation may be a violation of our trademark policy.

How Does Twitter Respond to Trademark Reports?

We review each report and do the following:

- If we determine that there is a violation of our trademark policy, we may suspend the account. In some instances, we may give the account holder an opportunity to comply with our policies before suspending the account. For more information, see our [Parody, commentary, and fan account policy](#).
- We may also provide the account holder with your name and information included in the copy of the report.
- If you would like to report a trademark violation in a [Twitter Ad](#), please visit our [Trademark Policy for Promoted Ads](#) to learn more.

Please note that Twitter only investigates requests that are submitted by the trademark holder or by an authorized representative of the trademark holder.

What is not a Trademark Policy Violation?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples include:

- Using a trademark in a way that is outside the scope of the trademark registration (e.g. territory, or goods and services identified in the registration).
- Nominative and other fair uses of trademarks are protected uses under our trademark policy, so long as the account is clearly distinguished from the trademark owner. This includes use by resellers in certain regions and accounts engaging in parody, commentary, or news. For more information, see our [Parody, commentary, and fan account policy](#).

How do I Report a Trademark Policy Violation?

Submit trademark reports through our [Trademark Issue support form](#). Using the appropriate form ensures we have all the information necessary to process your report quickly. Please provide all information requested in the form. If you submit a report with incomplete information, we will follow up with you to request that information. Please note that this will result in a delay in processing your report.

If you have trouble using the form, try updating your browser or switching to another browser.

Copyright policy

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged

infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notification.

When Should I File a Counter-notice?

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);

3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Twitter media policy

Twitter Rules: Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images.

Twitter Rules: You may not use hateful images or symbols in your profile image or profile header.

Rationale

Some forms of graphic violence, adult content, or hateful imagery may be permitted in Tweets when they are marked as [sensitive media](#). However, you may not include this type of content in live video, or in profile or header images.

While we want people to feel free to share media that reflects their creativity or individuality, or to show what's happening in the world, we will take action when it crosses the line into abuse towards a person, group, or protected category.

Finally, we may sometimes require you to remove media containing excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about how to make such a request [here](#).

When this applies

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. Some examples include, but are not limited to, depictions of:

- the moment at which someone dies
- gruesome crime or accident scenes
- bodily harm, torture, dismemberment, or mutilation

We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity (including close-ups of genitals, buttocks, or breasts)
 - Please note that exceptions may be made for artistic, medical, health, or educational content. Breastfeeding content does not need to be marked as sensitive.
- simulating a sexual act
- intercourse or any sexual act (may involve humans, humanoid animals, cartoons, or anime)

If you see media you think may be in violation of our policy, report it. You can read about how to report sensitive media [here](#).

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, or ethnicity/national origin.

Some examples of hateful imagery include, but are not limited to:

- Symbols historically associated with hate groups (for example, the Nazi swastika)
- Images depicting others as less than human or altered to include hateful symbols
- Altered image references to a mass murder that targeted a protected category

Do I need to be the target of this content in order for it to be a violation of the Twitter Rules?
No, we review both first-person and bystander reports of such content.

Consequences

When adult content, graphic violence, or hateful imagery appears in Tweets, we may place this content behind an interstitial advising viewers to be aware that they will see sensitive media if they click through. This allows us to identify potentially sensitive content that some people may not wish to see. Learn more about [how to control whether you see sensitive media](#).

However, if such content appears in live video, header, or profile images, it will violate our media policy. The consequences for violating this policy depend on the severity of the violation and the person's previous record of violations.

The first time someone violates this policy they will be required to remove the imagery and may be asked to verify their account ownership. Subsequent violations could result in permanent suspension. Learn more about [our range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

How to contact Twitter about media concerning a deceased family member

In order to respect the wishes of loved ones, Twitter will remove imagery of deceased individuals in certain circumstances. Immediate family members and other authorized individuals may request the removal of images

or video of deceased individuals, from when critical injury occurs to the moments before or after death, by sending a request to Twitter Inc. via our [privacy form](#). When reviewing such media removal requests, Twitter considers public interest factors such as the newsworthiness of the content and may not be able to honor every request. Note: Find information about contacting Twitter to [deactivate the account of a deceased or incapacitated user](#).

How to mark media as sensitive links to Media policy

FAQs about trends on Twitter

Where can I find trends?

Trends are available on the Twitter for iOS app, the Twitter for Android app, and [twitter.com](#).

On Twitter's mobile apps, you can find trends listed under the trends section



of the **Explore** tab. When signed in to [twitter.com](#) on a desktop or laptop computer, trends are listed in many places, including the Home timeline, Notifications, search results, and profile pages.

How are trends determined?

Trends are determined by an algorithm and, by default, are tailored for you based on who you follow, your interests, and your location. This algorithm identifies topics that are popular now, rather than topics that have been popular for a while or on a daily basis, to help you discover the hottest emerging topics of discussion on Twitter.

Note: The number of Tweets that are related to the trends is just one of the factors the algorithm looks at when ranking and determining trends. Algorithmically, trends and hashtags are grouped together if they are related to the same topic. For instance, #MondayMotivation and #MotivationMonday may both be represented by #MondayMotivation.

You can choose to see trends that are not tailored for you by selecting a [specific trends location](#) on [twitter.com](#), iOS, or Android. Location trends identify popular topics among people in a specific geographic location.

What does the # sign mean?

You may notice that some trends have # sign before the word or phrase. This is called a [hashtag](#) and is included specifically in Tweets to mark them as relating to a topic, so that people can follow the conversation in search.

Will I see world and local events in my trends?

Yes, when signed into your account on the [twitter.com](#), iOS, or Android, trends will be tailored for you based on your location and who you follow. There will be many world and local news events and conversations that will appear in your trends regardless of your personalization.

Can I see trends for a specific location?

Yes, you can view trends for a specific location on twitter.com, iOS, or Android.

View trends via twitter.com:

To view trends for a specific location, click **Change** in your trends box. You can choose to **Keep tailored trends** based on your location and who you follow, or you can select **Change** to choose a nearby location or search locations.

Adjust your trends settings via Twitter for iOS:

1. In the top menu, tap your **profile** icon.
2. Tap **Settings and privacy**, then tap **Content preferences**.
3. Under **Explore**, tap **Trends**.
4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.
5. When the feature is disabled, you can change your location by tapping **Change location**.
 1. In the search box, type your desired trends location or choose from the list.

Adjust your trends settings via Twitter for Android:

1. In the top menu, you will either see a **navigation menu** icon  or your **profile** icon. Tap whichever icon you have, and select **Settings and privacy**.
2. Tap **Content preferences**.
3. Under **Explore**, tap **Trends**.
4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.
5. When the feature is disabled, you can change your location by tapping **Change location**.
 1. In the search box, type your desired trends location or choose from the list.

Note: If you don't find your preferred country or city, it means we're not receiving enough Tweets from that geographical area to create a list. You can look up local Tweets on any topic by using [advanced Twitter Search operators](#).

What happens when I click or tap on a trend?

Clicking or tapping any of the trends takes you to the Twitter search results for that trend. You'll see all Tweets including that phrase or hashtag. To see what people are saying about a previous trend, perform a [search for that keyword](#).

How can I participate in a trend?

Simply post a Tweet including the exact word or phrase as it appears in the trends list (with the hashtag, if you see one). Due to the large number of people Tweeting about these specific trends, you may not always be able to find your particular Tweet in search, but your followers will always see your Tweets.

Note: Twitter also filters searches for quality. To understand what types of behavior could cause your Tweets to be filtered out, read about our [search rules and restrictions](#).

Are there rules for trends?

Yes. We want trends to promote healthy discussions on Twitter. This means that at times, we may prevent certain content from trending. These include trends that:

- Contain profanity or adult/graphic references.
- Incite hate on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease.
- Violate the [Twitter Rules](#).

In some cases, we may also consider the newsworthiness of the content, or if it is in the public interest when evaluating potential violations. In these cases, the content might continue to trend on our platform.

Note that even if we prevent the hashtag or the content from appearing on the list of trending topics, you may still be able to access conversations around that content on Twitter.

Trends descriptions

Below some trends, you may see the approximate number of Tweets associated with the trend, or personalized context such as who in your network is Tweeting about it.

Additionally, when a Moment is found to be relevant to a trend, we may use the title of the Moment as the trend's description.

Having trouble?

Check out our [troubleshooting section](#) for solutions to common problems.

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Child sexual exploitation policy

Twitter removes content promoting child sexual exploitation

We do not tolerate child sexual exploitation on Twitter. When we are made aware of links to images of or content promoting child sexual exploitation they will be removed from the site without further notice and reported to The National Center for Missing & Exploited Children ("NCMEC"); we permanently suspend accounts promoting or containing updates with links to child sexual exploitation.

Reporting content promoting child sexual exploitation

If you think you've found a Twitter account distributing or promoting child sexual exploitation while using Twitter, please click [here](#) to report it. Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. To find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with the National Center for Missing and Exploited Children [here](#), or your local law enforcement. If you believe that the content was hosted on a website outside of the United States, you can report it [here](#) on the International Association of Internet Hotlines website.

NOTE: Please do NOT tweet, retweet or repost child sexual exploitation for any reason. Click [here](#) to report it to us immediately and we will take steps to remove it.

Law Enforcement Agencies

If you are a member of a law enforcement agency seeking information regarding Twitter accounts promoting or distributing child sexual exploitation, please see our [Guidelines for Law Enforcement](#) for information on our process; we also suggest you contact the National Center for Missing and Exploited Children (NCMEC) [here](#), regardless of jurisdiction, to determine whether or not information on those accounts is already on file.

For general law enforcement inquiries, please see our [Guidelines for Law Enforcement](#) for more information.

Hateful conduct policy

Freedom of expression means little if voices are silenced because people are afraid to speak up. We do not tolerate behavior that harasses, intimidates, or uses fear to silence another person's voice. If you see something on Twitter that violates these rules, please report it to us.

How our policy works

As explained in the Twitter Rules,

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Examples of what we do not tolerate includes, but is not limited to behavior that harasses individuals or groups of people with:

violent threats;

wishes for the physical harm, death, or disease of individuals or groups;

references to mass murder, violent events, or specific means of violence in which/with which such groups have been the primary targets or victims;

behavior that incites fear about a protected group;

repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone.

How our enforcement works Context matters.

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. While we accept reports of violations from anyone, sometimes we also need to hear directly from the target to ensure that we have proper context.

The number of reports we receive does not impact whether or not something will be removed. However, it may help us prioritize the order in which it gets reviewed.

We focus on behavior.

We enforce policies when someone reports behavior that is abusive and targets an entire protected group and/or individuals who may be members.

This targeting can happen in any manner (for example, @mentions, tagging a photo, and more).

We have a range of enforcement options.

The consequences for violating our rules vary depending on the severity of the violation and the person's previous record of violations. For example, we may ask someone to remove the offending Tweet before they can Tweet again. For other cases, we may suspend an account.

Report private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Read about our [policy regarding private information posted on Twitter](#).

Do I need to have a Twitter account to make a report?

No, you do not need to have a Twitter account to file a report that your private and confidential information has been posted on Twitter.

Can I report unauthorized posting of Copyrighted material?

If your copyrighted material has been Tweeted without your permission, please review our [Copyright and DMCA policy](#).

What can I do if my information is posted on other websites?

We cannot take down any information that is posted on a website other than Twitter. If your information is posted on another website, please contact the help center or support staff for the website where your information is posted.

How do I report that my private information is posted on Twitter?

You can report that your private information has been posted on Twitter directly from the offending Tweet(s). Learn [how to report individual Tweets](#).

File a [report that someone has posted your private information](#).

What happens when Twitter receives a valid report?

When we receive a complete and valid report that private information has been posted on Twitter, we'll investigate the account and Tweets reported. We will review where, if anywhere, the information has been made publicly available before taking action on the account or Tweets. If the information you reported was previously posted elsewhere on the Internet, it is not a violation of our policy and we will not take action.

Learn more about [protecting your private information on Twitter and other websites](#).

Read our [policy regarding private information posted on Twitter](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter impersonation policy.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to Twitter's [parody, commentary, and fan account policy](#) for more information about these types of accounts.

Find more information about [reporting impersonation on Twitter](#).

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation rules

Updated November 3, 2017

This page is primarily intended for **developers**.

For Twitter users: You are ultimately responsible for the actions taken with your account, or by applications associated with your account. Before authorizing a third-party application to access or use your account, make sure you've thoroughly investigated the application and understand what it will do. If automated activity on your account violates the [Twitter Rules](#) or these automation rules, Twitter may take action on your account, including [filtering your Tweets from search results](#) or suspending your account.

For more information on third-party applications, please see our article on [connecting and revoking third-party applications](#).

If you are a developer and have questions about these Automation Rules, please reach out via <https://twittercommunity.com/c/rules-and-policies>.

If your application has been suspended or restricted from performing write actions, you may submit an appeal at </forms/platform>.

I. Ground Rules Do!

Build solutions that automatically broadcast helpful information in Tweets.

Run creative campaigns that auto-reply to users who engage with your content.

Build solutions that automatically respond to users in Direct Messages.

Try new things that help people (and comply with our rules).

Make sure your application provides a good user experience and performs well — and confirm that remains the case over time.

Don't!

Violate these or other policies. Be extra mindful of our rules about abuse and user privacy.

Abuse the Twitter API or attempt to circumvent rate limits.

Use non-API-based forms of automation, such as scripting the Twitter website. The use of these techniques may result in the permanent suspension of your account.

Spam or bother users, or otherwise send them unsolicited messages.

A. The Twitter Rules and the Developer Agreement and Policy

As with all activity on Twitter, automated activity is subject to the [Twitter Rules](#) and, if you're a developer using the Twitter API, the [Developer Agreement and Policy](#).

You should carefully review these policies to ensure that your automated activity is compliant. Automated applications or activities that violate these policies, or that facilitate or induce users to violate them, may be subject to enforcement action, potentially including suspension of associated Twitter accounts.

We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these policies.

Although all aspects of the Twitter Rules and the Developer Agreement and Policy apply to automated activity, you should keep the following rules top of mind:

Spamming: You may not send automated Tweets or Direct Messages that are spam, or otherwise engage in spamming activity. Some examples of spammy behavior to avoid with automation include:

Trending topics: You may not automatically post about trending topics on Twitter, or use automation to attempt to influence or manipulate trending topics.

Multiple posts/accounts: You may not post duplicative or substantially similar Tweets on one account or over multiple accounts you operate.

Duplicate accounts: You may not create and/or automate multiple accounts for duplicative or substantially similar use cases.

However, automating multiple accounts for related but non-duplicative use cases is permitted. For example, you may automate separate accounts to Tweet when the Hubble Space Telescope passes over different cities, such as [San Francisco](#) or [Hong Kong](#).

Misleading links: You may not send automated Tweets or Direct Messages containing links that are misleading, including links that maliciously or deceptively redirect through landing pages or ad pages before displaying the final content.

Sensitive media: Automated Tweets and Direct Messages must comply with the [Twitter media policy](#), and you should mark your account as potentially sensitive if you intend to post graphic, pornographic, or potentially sensitive media.

Abusive behavior: You may not engage in any automated activity that encourages, promotes, or incites abuse, violence, hateful conduct, or harassment, on or off Twitter.

Private information: You may not post private or confidential information about a person without their prior express authorization.

B. Other Ground Rules for Automated Activity

In addition to the policies above, the following ground rules apply to all automated activity on Twitter:

Don't surprise or mislead users: Automated activity should honor users' expectations. Ask for the user's permission before taking an action if you aren't sure.

Mature content or profanity: Don't Direct Message, mention, or reply to users with potentially sensitive content (including profanity), unless they've clearly indicated an intent to receive it in advance.

Be thoughtful about the information you request or exchange on Twitter

Tweets: Don't ask users to send you personal or private information via a public Tweet. If you need additional personal or private information from a user to provide them with customer service (or other similar use cases), you should ask the user to share such information by Direct Message or another private channel. You might even consider adding a [Direct Message deep link](#) to your Tweet.

Direct Messages: You should only ask users for the minimum amount of information you need to provide them with service. If you need to request or exchange particularly sensitive information (such as credit card information), you should consider directing users to your website or other appropriate channel to do so.

II. Activity-Specific Rules

The activity-specific rules in this section apply to taking specific automated actions on Twitter. Please read these rules carefully, as they outline both permitted and prohibited use cases of automation.

Automated applications or activities that violate these rules, or that facilitate or induce users to violate them, may be subject to enforcement action, including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these rules. As a reminder, you should also carefully review the spam guidelines in the [Twitter Rules](#) to avoid having activities performed by you, your app, or other users through your app or service flagged as spam.

A. Automated Actions Through Another User's Account

Twitter users may authorize your app or service to [access their Twitter account through OAuth](#). A user authorizing your app or service to access their Twitter account through OAuth does not by itself constitute sufficient consent to take automated actions through that user's account.

You may only take automated actions through another Twitter user's account if you:

clearly describe to the user the types of automated actions that will occur; receive express consent from the user to take those automated actions; and immediately honor a user's request to opt-out of further automated actions.

If you substantially change the purpose or functionality of your app or service, you must re-obtain express consent from the user to take automated action through their account before doing so.

These requirements apply to any automated action taken through another Twitter user's account, including posting Tweets, sending Direct Messages, deleting Tweets or Direct Messages, or following/unfollowing other accounts. For applications that offer users the ability to delete Tweets in a bulk or automated manner, you must also clearly state that Tweets are not recoverable once deleted.

B. Automated Tweets

1. Posting automated Tweets

Automated Tweets that cross-post outside information:

You may post automated Tweets based on sources of outside information — such as an RSS feed, weather data, etc. — as long as you are sufficiently authorized to publish such information.

Other automated Tweets (excluding mentions or replies):

Provided you comply with all other rules, you may post automated Tweets for entertainment, informational, or novelty

purposes. As a reminder, accounts posting duplicative, spammy, or otherwise prohibited content may be subject to suspension.

2. Posting automated mentions and replies

The reply and mention functions are intended to make communication between Twitter users easier. Automating these actions to reach many users on an unsolicited basis is an abuse of the feature, and is not permitted. For example, sending automated replies to Tweets based on keyword searches alone is not permitted. Spamy or duplicative use of mentions and replies may result in enforcement action, such as the removal of your Tweets from Search or the suspension of your app or account.

However, you may send automated replies or mentions to Twitter users so long as:

in advance of sending the automated reply, the recipient or mentioned user(s) have requested or have clearly indicated an intent on Twitter to be contacted by you (i.e. opted in), for example by replying to a Tweet from your account, or by sending you a Direct Message;

you provide a clear and easy way for such users to opt-out of receiving automated replies and mentions, and promptly honor all such opt-out requests;

you only send one automated reply or mention per user interaction; and

the automated reply or mention is a reply to the user's original Tweet (if your campaign is based on users posting a reply to your Tweet).

Opt-in techniques and indications of user intent take many different forms, depending on the specifics of your use case and implementation. Some examples include:

A Tweet from your account that clearly indicates that a user taking a specific action on that Tweet (such as Retweeting it) will opt the user into receiving an automated response.

A mention of your account by the user in a manner suggesting the user clearly wishes or intends to receive a response. If you want to run an auto-reply campaign with a campaign- or use-case-specific hashtag, users should also mention you in their Tweets.

Note that a user following your account is not on its own a sufficient indication of user intent to receive an automated response.

: Advertisers, publishers, and brands using auto-response campaigns must request approval from Twitter and may be subject to additional rules. Please reach out to your account or partner manager for assistance. If you do not have a partner manager but are interested in managed client services, you can learn more [here](#).

Additionally, we recommend that any accounts that will communicate with

users via automated mentions or replies:

Note

1. Appropriately filter responses based on potentially sensitive language in user handles,

display names, and Tweet text, as well as potentially sensitive media;

2. Check that the Tweet you are mentioning or replying to still exists (for example, using the statuses/lookup endpoint on the Twitter API).

C. Automated Direct Messages

Sending automated Direct Messages to users

You may not send unsolicited Direct Messages in a bulk or automated manner, and should be thoughtful about the frequency with which you contact users via Direct Message.

You may send automated Direct Messages to users so long as:

in advance of sending the Direct Message, the recipient(s) have requested or have clearly indicated an intent on Twitter to be contacted by you via Direct Message, for example by sending you a Direct Message; **and**

you provide a clear and easy way for such users to opt-out of receiving automated Direct Messages, and promptly honor all such opt-out requests.

The fact that a user is technically able to receive a Direct Message from you (e.g. because the user follows you, has enabled the ability to receive Direct Messages from any

account, or because the user is in a pre-existing Direct Message conversation with you) does not necessarily mean they have requested or expect to receive automated Direct Messages from you.

Interacting with users via Direct Message

Per the Ground Rules, remember to be thoughtful about the amount and type of information you request or exchange with users via Direct Messages. If you will be asking a user to provide personal or private information via an automated Direct Message, you must clearly explain how you will use the information you're collecting. Consider including a link to your privacy policy in your Direct Message to the user, as well as in your Twitter profile bio.

Don't publicly share information received in a Direct Message conversation with a user without first obtaining explicit consent from the user. For example, if a user asks you via Direct Message about a purchase they made from you, you may not mention the user in a public Tweet that includes information about their purchase unless you have the user's explicit consent to do so.

After a user-initiated interaction ends, don't send additional follow-up Direct Messages or mention users in a Tweet unless you get permission from the user.

D. Automated actions you take on Tweets or accounts

Automated likes: You may not like Tweets in an automated manner.

Automated Retweets: Provided you comply with all other rules, you may Retweet or Quote Tweet in an automated manner for entertainment, informational, or novelty purposes. Automated Retweets often lead to negative user experiences, and bulk, aggressive, or spammy Retweeting is a violation of the [Twitter Rules](#).

Automated following/unfollowing: You may not follow or unfollow Twitter accounts in a bulk, aggressive, or indiscriminate manner. Aggressive following is a violation of the [Twitter Rules](#). Please also review our [following rules and best practices](#) to ensure you are in compliance. Note that applications that claim to get users more followers are also prohibited under the [Twitter Rules](#).

Automated adding to lists or collections: You may not add Twitter users to lists or add Tweets to collections in a bulk or indiscriminate manner. Adding a large number of unrelated users to lists is a violation of the [Twitter Rules](#).

About search rules and restrictions

This article details the rules for using Twitter search, and explains reasons why content may not show up in Twitter search.

Why doesn't all content show up in Twitter search?

In order to keep your search results relevant, Twitter filters search results for quality Tweets and accounts. Material that jeopardizes search quality or creates a bad search experience for other people may be automatically removed from Twitter search. Read more about this [here](#).

General rules for using Twitter search

The [Twitter Rules](#) explain what behaviors are permitted on Twitter. In addition to these rules, we've included some tips below to keep your content relevant (and your Tweets in Twitter search).

Please do not:

Repeatedly post duplicate or near-duplicate content (links or Tweets). Abuse trending topics or hashtags (topic words with a # sign).
Send automated Tweets or replies.

Use bots or applications to post similar messages based on keywords. Post similar messages over multiple accounts.

Aggressively follow and unfollow people.

Twitter may automatically remove accounts engaging in these behaviors from search (or even suspend in some cases) in order to ensure the best experience for everyone.

Search results filtering

To learn how to filter your search results, including how to disable or enable safe search, read our [Using Twitter search](#) article.

Search results filtering in third-party applications

Accounts may be filtered from search for posting similar messages over several accounts. Your account may be filtered from search if you are logged into a third-party application that is updating many accounts with similar or duplicate content.

Rules for participating in contests on Twitter

Businesses sometimes host contests and offer prizes to Retweet an update, following a particular account, or posting Tweets with a specific hashtag topic or reply. If you're running a contest, check out this [article](#) for guidelines for contests on Twitter.

If you repeatedly post near-duplicate updates or duplicate links in order to enter contests, you may be filtered from search. Please keep in mind that you are responsible for the updates that you post to your Twitter account.

What to do if you are missing from search

Please see this [article](#) for more information related to this issue.

If you are being filtered from search, rest assured that your followers will still receive your updates. If you post a reply or mention to a specific person, it will be delivered to that person.

Having trouble?

You can learn more about searching by keyword, location, and trending topics in our [Using Twitter search](#) help article.

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. These limitations are set forth in the Twitter Rules below.

The Twitter Rules (along with all incorporated policies), [Privacy Policy](#), and [Terms of Service](#) collectively make up the "Twitter User Agreement" that governs a user's access to and use of Twitter's services.

All individuals accessing or using Twitter's services must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in Twitter taking one or more of the following enforcement actions:

- requiring you to delete prohibited content before you can again create new posts and interact with other Twitter users;
- temporarily limiting your ability to create posts or interact with other Twitter users;
- asking you to verify account ownership with a phone number or email address; or
- permanently suspending your account(s).

If you attempt to evade a permanent suspension by creating new accounts, we will suspend your new accounts.

Please note that we may need to change these Rules from time to time and reserve the right to do so. The most current version is always available at: <https://twitter.com/rules>.

The policies set forth in these Twitter Rules govern organic content on our platform. To learn more about the rules which govern ads and promoted content, please review our [Ads policies](#).

Content Boundaries and Use of Twitter

Intellectual property

Trademark: We reserve the right to suspend accounts or take other appropriate action when someone's brand or trademark, including business name and/or logo, is used in a manner that may mislead or confuse others about your brand affiliation. Read more about our [trademark policy](#) and [how to report a violation](#).

Copyright: We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in our [Terms of Service](#). Read more about our [copyright policy](#).

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Learn more about our [media policy](#).

Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images. Additionally, Twitter may sometimes require you to remove excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about [how to make such a request](#), and [how to mark your media as sensitive](#).

Unlawful use

You may not use our service for any unlawful purposes or in furtherance of illegal activities. By using Twitter, you agree to comply with all applicable laws governing your online conduct and content.

Trends

At times, we may prevent certain content from trending. This includes content that violates the Twitter Rules, as well as content that may attempt to manipulate trends. Read more about [what we allow and do not allow to trend](#).

Misuse of Twitter badges

You may not use badges, including but not limited to the “promoted” or “verified” Twitter badges, unless provided by Twitter. Accounts using unauthorized badges as part of their profile photos, header photos, display names, or in any way that falsely implies affiliation with Twitter or authorization from Twitter to display these badges, may be suspended.

Misuse of usernames

Selling usernames: You may not buy or sell Twitter usernames.

Username squatting: You may not engage in username squatting. Some of the factors we take into consideration when determining whether conduct is username squatting include:

- the number of accounts created;
- the creation of accounts for the purpose of preventing others from using those account names;
- the creation of accounts for the purpose of selling those accounts; and
- the use of third-party content feeds to update and maintain accounts under the names of those third parties.

Please note that Twitter may also remove accounts that are inactive for more than six months. Learn more about [username squatting](#).

Abusive Behavior

We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we prohibit behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Context matters when evaluating for abusive behavior and determining appropriate enforcement actions. Factors we may take into consideration include, but are not limited to whether:

- the behavior is targeted at an individual or group of people;
- the report has been filed by the target of the abuse or a bystander;
- the behavior is newsworthy and in the legitimate public interest.

Violence and physical harm

Violence: You may not make specific threats of violence or wish for the serious physical harm, death, or disease of an individual or group of people. This includes, but is not limited to, threatening or promoting terrorism. You also may not affiliate with organizations that — whether by their own statements or activity both on and off the platform — use or promote violence against civilians to further their causes. We will begin enforcing this rule around affiliation with such organizations on December 18, 2017.

Suicide or self-harm: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Child sexual exploitation: You may not promote child sexual exploitation. Learn more about our zero-tolerance [child sexual exploitation policy](#).

Abuse and hateful conduct

Abuse: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Unwanted sexual advances: You may not direct abuse at someone by sending unwanted sexual content, objectifying them in a sexually explicit manner, or otherwise engaging in sexual misconduct.

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. Read more about our [hateful conduct policy](#).

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category. We will begin enforcing this rule on December 18, 2017.

Private information and intimate media

Private information: You may not publish or post other people's private information without their express authorization and permission. Definitions of private information may vary depending on local laws. Read more about our [private information policy](#).

Intimate media: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. Read more about intimate media on Twitter.

Threats to expose / hack: You may not threaten to expose someone's private information or intimate media. You also may not threaten to hack or break into someone's digital information.

Impersonation

You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. While you may maintain parody, fan, commentary, or newsfeed accounts, you may not do so if the intent of the account is to engage in spamming or abusive behavior. Read more about our [impersonation policy](#).

Spam and Security

We strive to protect people on Twitter from technical abuse and spam.

To promote a stable and secure environment on Twitter, you may not do, or attempt to do, any of the following while accessing or using Twitter:

- Access, tamper with, or use non-public areas of Twitter, Twitter's computer systems, or the technical delivery systems of Twitter's providers (except as expressly permitted by the Twitter Bug Bounty program).
- Probe, scan, or test the vulnerability of any system or network, or breach or circumvent any security or authentication measures (except as expressly permitted by the Twitter Bug Bounty program).
- Access or search, or attempt to access or search, Twitter by any means (automated or otherwise) other than through our currently available, published interfaces that are provided by Twitter (and only pursuant to the applicable terms and conditions), unless you have been specifically allowed to do so in a separate agreement with Twitter. Note that crawling Twitter is permissible if done in accordance with the provisions of the robots.txt file; however, scraping Twitter without our prior consent is expressly prohibited.
- Forge any TCP/IP packet header or any part of the header information in any email or posting, or in any way use Twitter to send altered, deceptive, or false source-identifying information.
- Interfere with or disrupt the access of any user, host or network, including, without limitation, sending a virus, overloading, flooding, spamming, mail-bombing Twitter's services, or by scripting the creation of content in such a manner as to interfere with or create an undue burden on Twitter.

Any accounts engaging in the following activities may be temporarily locked or subject to permanent suspension:

- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use Twitter's services for the purpose of spamming anyone. Spam is generally defined on Twitter as bulk or aggressive activity that attempts to manipulate or disrupt Twitter or the experience of users on Twitter to drive traffic or attention to unrelated accounts, products, services, or initiatives. Some of the factors that we take into account when determining what conduct is considered to be spamming include:
 - if you have followed and/or unfollowed a large number of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if your Tweets or Direct Messages consist mainly of links shared without commentary;
 - if a large number of people have blocked you in response to high volumes of untargeted, unsolicited, or duplicative content or engagements from your account;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicative or substantially similar content, replies, or mentions over multiple accounts or multiple duplicate updates on one account, or create duplicate or substantially similar accounts;
 - if you post multiple updates to a trending or popular topic with an intent to subvert or manipulate the topic to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you send large numbers of unsolicited replies or mentions;
 - if you add users to lists in a bulk or aggressive manner;
 - if you are randomly or aggressively engaging with Tweets (e.g., likes, Retweets, etc.) or users (e.g., following, adding to lists or Moments, etc.) to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you repeatedly post other people's account information as your own (e.g., bio, Tweets, profile URL, etc.);
 - if you post misleading, deceptive, or malicious links (e.g., affiliate links, links to malware/clickjacking pages, etc.);
 - if you create fake accounts, account interactions, or impressions;
 - if you sell, purchase, or attempt to artificially inflate account interactions (such as followers, Retweets, likes, etc.); and
 - if you use or promote third-party services or apps that claim to get you more followers, Retweets, or likes (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers or engagements to your account or Tweets).

Please see our support articles on [following rules and best practices](#) and [automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts may be permanently suspended.

Content Visibility

Accounts under investigation or which have been detected as sharing content in violation of these Rules may have their account or Tweet visibility limited in various parts of Twitter, including search. To learn more about situations in which content may be restricted on Twitter, please see our support article on [search rules and restrictions](#).

Trademark policy

Updated June 30, 2016

What is a Trademark Policy Violation on Twitter?

Using another's trademark in a manner that may mislead or confuse others about your brand affiliation may be a violation of our trademark policy.

How Does Twitter Respond to Trademark Reports?

We review each report and do the following:

- If we determine that there is a violation of our trademark policy, we may suspend the account. In some instances, we may give the account holder an opportunity to comply with our policies before suspending the account. For more information, see our [Parody, commentary, and fan account policy](#).
- We may also provide the account holder with your name and information included in the copy of the report.
- If you would like to report a trademark violation in a [Twitter Ad](#), please visit our [Trademark Policy for Promoted Ads](#) to learn more.

Please note that Twitter only investigates requests that are submitted by the trademark holder or by an authorized representative of the trademark holder.

What is not a Trademark Policy Violation?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples include:

- Using a trademark in a way that is outside the scope of the trademark registration (e.g. territory, or goods and services identified in the registration).
- Nominative and other fair uses of trademarks are protected uses under our trademark policy, so long as the account is clearly distinguished from the trademark owner. This includes use by resellers in certain regions and accounts engaging in parody, commentary, or news. For more information, see our [Parody, commentary, and fan account policy](#).

How do I Report a Trademark Policy Violation?

Submit trademark reports through our [Trademark Issue support form](#). Using the appropriate form ensures we have all the information necessary to process your report quickly. Please provide all information requested in the form. If you submit a report with incomplete information, we will follow up with you to request that information. Please note that this will result in a delay in processing your report.

If you have trouble using the form, try updating your browser or switching to another browser.

Copyright policy

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged

infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notification.

When Should I File a Counter-notice?

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);

3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Twitter media policy

Twitter Rules: Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images.

Twitter Rules: You may not use hateful images or symbols in your profile image or profile header.

Rationale

Some forms of graphic violence, adult content, or hateful imagery may be permitted in Tweets when they are marked as [sensitive media](#). However, you may not include this type of content in live video, or in profile or header images.

While we want people to feel free to share media that reflects their creativity or individuality, or to show what's happening in the world, we will take action when it crosses the line into abuse towards a person, group, or protected category.

Finally, we may sometimes require you to remove media containing excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about how to make such a request [here](#).

When this applies

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. Some examples include, but are not limited to, depictions of:

- the moment at which someone dies
- gruesome crime or accident scenes
- bodily harm, torture, dismemberment, or mutilation

We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity (including close-ups of genitals, buttocks, or breasts)
 - Please note that exceptions may be made for artistic, medical, health, or educational content. Breastfeeding content does not need to be marked as sensitive.
- simulating a sexual act
- intercourse or any sexual act (may involve humans, humanoid animals, cartoons, or anime)

If you see media you think may be in violation of our policy, report it. You can read about how to report sensitive media [here](#).

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, or ethnicity/national origin.

Some examples of hateful imagery include, but are not limited to:

- Symbols historically associated with hate groups (for example, the Nazi swastika)
- Images depicting others as less than human or altered to include hateful symbols
- Altered image references to a mass murder that targeted a protected category

Do I need to be the target of this content in order for it to be a violation of the Twitter Rules?
No, we review both first-person and bystander reports of such content.

Consequences

When adult content, graphic violence, or hateful imagery appears in Tweets, we may place this content behind an interstitial advising viewers to be aware that they will see sensitive media if they click through. This allows us to identify potentially sensitive content that some people may not wish to see. Learn more about [how to control whether you see sensitive media](#).

However, if such content appears in live video, header, or profile images, it will violate our media policy. The consequences for violating this policy depend on the severity of the violation and the person's previous record of violations.

The first time someone violates this policy they will be required to remove the imagery and may be asked to verify their account ownership. Subsequent violations could result in permanent suspension. Learn more about [our range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

How to contact Twitter about media concerning a deceased family member

In order to respect the wishes of loved ones, Twitter will remove imagery of deceased individuals in certain circumstances. Immediate family members and other authorized individuals may request the removal of images

or video of deceased individuals, from when critical injury occurs to the moments before or after death, by sending a request to Twitter Inc. via our [privacy form](#). When reviewing such media removal requests, Twitter considers public interest factors such as the newsworthiness of the content and may not be able to honor every request. Note: Find information about contacting Twitter to [deactivate the account of a deceased or incapacitated user](#).

How to mark media as sensitive links to Media policy

Twitter trends FAQs

Where can I find trends?

Trends are available on the Twitter for iOS app, the Twitter for Android app, and [twitter.com](#).

On Twitter's mobile apps, you can find trends listed under the trends section of the **Explore** tab. When signed in to [twitter.com](#) on a desktop or laptop computer, trends are listed in many places, including the Home timeline, Notifications, search results, and profile pages.

How are trends determined?

Trends are determined by an algorithm and, by default, are tailored for you based on who you follow, your interests, and your location. This algorithm identifies topics that are popular now, rather than topics that have been popular for a while or on a daily basis, to help you discover the hottest emerging topics of discussion on Twitter.

You can choose to see trends that are not tailored for you by selecting a specific trends location on [twitter.com](#), iOS, or Android (instructions below). Location trends identify popular topics among people in a specific geographic location.

Note: The number of Tweets that are related to the trends is just one of the factors the algorithm looks at when ranking and determining trends. Algorithmically, trends and hashtags are grouped together if they are related to the same topic. For

instance, #MondayMotivation and #MotivationMonday may both be represented by #MondayMotivation.

What does the # sign mean?

You may notice that some trends have # sign before the word or phrase. This is called a **hashtag** and is included specifically in Tweets to mark them as relating to a topic, so that people can follow the conversation in search.

Will I see world and local events in my trends?

Yes, when signed into your account on the twitter.com, iOS, or Android, trends will be tailored for you based on your location and who you follow. There will be many world and local news events and conversations that will appear in your trends regardless of your personalization.

Can I see trends for a specific location?

Yes, you can view trends for a specific location on twitter.com, iOS, or Android.

View trends via twitter.com:

To view trends for a specific location, click **Change** in your trends box. You can choose to **Keep tailored trends** based on your location and who you follow, or you can select **Change** to choose a nearby location or search locations.

Adjust your trends settings via Twitter for iOS:

1. In the top menu, tap your **profile** icon.
2. Tap **Settings and privacy**, then tap **Content preferences**.
3. Under **Explore**, tap **Trends**.
4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.
5. When the feature is disabled, you can change your location by tapping **Change**

location. In the search box, type your desired trends location or choose from the list.

Adjust your trends settings via Twitter for Android:

1. In the top menu, you will either see a **navigation menu** icon or your **profile** icon.

Tap whichever icon you have, and select **Settings and privacy**.

2. Tap **Content preferences**.
3. Under **Explore**, tap **Trends**.
4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.
5. When the feature is disabled, you can change your location by tapping **Change location**. In the search box, type your desired trends location or choose from the list.

Note: If you don't find your preferred country or city, it means we're not receiving enough Tweets from that geographical area to create a list. You can look up local Tweets on any topic by using [advanced Twitter search operators](#).

What happens when I click or tap on a trend?

Clicking or tapping any of the trends takes you to the Twitter search results for that trend. You'll see all Tweets including that phrase or hashtag. To see what people are saying about a previous trend, perform a [search for that keyword](#).

How can I participate in a trend?

Simply post a Tweet including the exact word or phrase as it appears in the trends list (with the hashtag, if you see one). Due to the large number of people Tweeting about these specific trends, you may not always be able to find your particular Tweet in search, but your followers will always see your Tweets.

Note: Twitter also filters searches for quality. To understand what types of behavior could cause your Tweets to be filtered out, read about our [search rules and restrictions](#).

Are there rules for trends?

Yes. We want trends to promote healthy discussions on Twitter. This means that at times, we may prevent certain content from trending. These include trends that:

Contain profanity or adult/graphic references.

Incite hate on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease.

Violate the [Twitter Rules](#).

In some cases, we may also consider the newsworthiness of the content, or if it is in the public interest when evaluating potential violations. In these cases, the content might continue to trend on our platform.

Note that even if we prevent the hashtag or the content from appearing on the list of trending topics, you may still be able to access conversations around that content on Twitter.

Do trends have descriptions?

Below some trends, you may see the approximate number of Tweets associated with the trend, or personalized context such as who in your network is Tweeting about it.

Additionally, when a Moment is found to be relevant to a trend, we may use the title of the Moment as the trend's description.

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Child sexual exploitation policy

Twitter removes content promoting child sexual exploitation

We do not tolerate child sexual exploitation on Twitter. When we are made aware of links to images of or content promoting child sexual exploitation they will be removed from the site without further notice and reported to The National Center for Missing & Exploited Children ("NCMEC"); we permanently suspend accounts promoting or containing updates with links to child sexual exploitation.

Reporting content promoting child sexual exploitation

If you think you've found a Twitter account distributing or promoting child sexual exploitation while using Twitter, please click [here](#) to report it. Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. To find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with the National Center for Missing and Exploited Children [here](#), or your local law enforcement. If you believe that the content was hosted on a website outside of the United States, you can report it [here](#) on the International Association of Internet Hotlines website.

NOTE: Please do NOT tweet, retweet or repost child sexual exploitation for any reason. Click [here](#) to report it to us immediately and we will take steps to remove it.

Law Enforcement Agencies

If you are a member of a law enforcement agency seeking information regarding Twitter accounts promoting or distributing child sexual exploitation, please see our [Guidelines for Law Enforcement](#) for information on our process; we also suggest you contact the National Center for Missing and Exploited Children (NCMEC) [here](#), regardless of jurisdiction, to determine whether or not information on those accounts is already on file.

For general law enforcement inquiries, please see our [Guidelines for Law Enforcement](#) for more information.

Hateful conduct policy

Freedom of expression means little if voices are silenced because people are afraid to speak up. We do not tolerate behavior that harasses, intimidates, or uses fear to silence another person's voice. If you see something on Twitter that violates these rules, please report it to us.

How our policy works

As explained in the Twitter Rules,

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Examples of what we do not tolerate includes, but is not limited to behavior that harasses individuals or groups of people with:

violent threats;

wishes for the physical harm, death, or disease of individuals or groups;

references to mass murder, violent events, or specific means of violence in which/with which such groups have been the primary targets or victims;

behavior that incites fear about a protected group;

repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone.

How our enforcement works Context matters.

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. While we accept reports of violations from anyone, sometimes we also need to hear directly from the target to ensure that we have proper context.

The number of reports we receive does not impact whether or not something will be removed. However, it may help us prioritize the order in which it gets reviewed.

We focus on behavior.

We enforce policies when someone reports behavior that is abusive and targets an entire protected group and/or individuals who may be members.

This targeting can happen in any manner (for example, @mentions, tagging a photo, and more).

We have a range of enforcement options.

The consequences for violating our rules vary depending on the severity of the violation and the person's previous record of violations. For example, we may ask someone to remove the offending Tweet before they can Tweet again. For other cases, we may suspend an account.

Report private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Read about our [policy regarding private information posted on Twitter](#).

Do I need to have a Twitter account to make a report?

No, you do not need to have a Twitter account to file a report that your private and confidential information has been posted on Twitter.

Can I report unauthorized posting of Copyrighted material?

If your copyrighted material has been Tweeted without your permission, please review our [Copyright and DMCA policy](#).

What can I do if my information is posted on other websites?

We cannot take down any information that is posted on a website other than Twitter. If your information is posted on another website, please contact the help center or support staff for the website where your information is posted.

How do I report that my private information is posted on Twitter?

You can report that your private information has been posted on Twitter directly from the offending Tweet(s). Learn [how to report individual Tweets](#).

File a [report that someone has posted your private information](#).

What happens when Twitter receives a valid report?

When we receive a complete and valid report that private information has been posted on Twitter, we'll investigate the account and Tweets reported. We will review where, if anywhere, the information has been made publicly available before taking action on the account or Tweets. If the information you reported was previously posted elsewhere on the Internet, it is not a violation of our policy and we will not take action.

Learn more about [protecting your private information on Twitter and other websites](#).

Read our [policy regarding private information posted on Twitter](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter impersonation policy.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to Twitter's [parody, commentary, and fan account policy](#) for more information about these types of accounts.

Find more information about [reporting impersonation on Twitter](#)

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are

prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation rules

Updated November 3, 2017

This page is primarily intended for **developers**.

For Twitter users: You are ultimately responsible for the actions taken with your account, or by applications associated with your account. Before authorizing a third-party application to access or use your account, make sure you've thoroughly investigated the application and understand what it will do. If automated activity on your account violates the [Twitter Rules](#) or these automation rules, Twitter may take action on your account, including [filtering your Tweets from search results](#) or suspending your account.

For more information on third-party applications, please see our article on [connecting and revoking third-party applications](#).

If you are a developer and have questions about these Automation Rules, please reach out via <https://twittercommunity.com/c/rules-and-policies>.

If your application has been suspended or restricted from performing write actions, you may submit an appeal at </forms/platform>.

I. Ground Rules Do!

Build solutions that automatically broadcast helpful information in Tweets.

Run creative campaigns that auto-reply to users who engage with your content.

Build solutions that automatically respond to users in Direct Messages.

Try new things that help people (and comply with our rules).

Make sure your application provides a good user experience and performs well — and confirm that remains the case over time.

Don't!

Violate these or other policies. Be extra mindful of our rules about abuse and user privacy.

Abuse the Twitter API or attempt to circumvent rate limits.

Use non-API-based forms of automation, such as scripting the Twitter website. The use of these techniques may result in the permanent suspension of your account.

Spam or bother users, or otherwise send them unsolicited messages.

A. The Twitter Rules and the Developer Agreement and Policy

As with all activity on Twitter, automated activity is subject to the [Twitter Rules](#) and, if you're a developer using the Twitter API, the [Developer Agreement and Policy](#).

You should carefully review these policies to ensure that your automated activity is compliant. Automated applications or activities that violate these policies, or that facilitate or induce users to violate them, may be subject to enforcement action, potentially including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these policies.

Although all aspects of the Twitter Rules and the Developer Agreement and Policy apply to automated activity, you should keep the following rules top of mind:

Spamming: You may not send automated Tweets or Direct Messages that are spam, or otherwise engage in spamming activity. Some examples of spammy behavior to avoid with automation include:

Trending topics: You may not automatically post about trending topics on Twitter, or use automation to attempt to influence or manipulate trending topics.

Multiple posts/accounts: You may not post duplicative or substantially similar Tweets on one account or over multiple accounts you operate.

Duplicate accounts: You may not create and/or automate multiple accounts for duplicative or substantially similar use cases.

However, automating multiple accounts for related but non-duplicative use cases is permitted. For example, you may automate separate accounts to Tweet when the Hubble Space Telescope passes over different cities, such as [San Francisco](#) or [Hong Kong](#).

Misleading links: You may not send automated Tweets or Direct Messages containing links that are misleading, including links that maliciously or deceptively redirect through landing pages or ad pages before displaying the final content.

Sensitive media: Automated Tweets and Direct Messages must comply with the [Twitter media policy](#), and you should mark your account as potentially sensitive if you intend to post graphic, pornographic, or potentially sensitive media.

Abusive behavior: You may not engage in any automated activity that encourages, promotes, or incites abuse, violence, hateful conduct, or harassment, on or off Twitter.

Private information: You may not post private or confidential information about a person without their prior express authorization.

B. Other Ground Rules for Automated Activity

In addition to the policies above, the following ground rules apply to all automated activity on Twitter:

Don't surprise or mislead users: Automated activity should honor users' expectations. Ask for the user's permission before taking an action if you aren't sure.

Mature content or profanity: Don't Direct Message, mention, or reply to users with potentially sensitive content (including profanity), unless they've clearly indicated an intent to receive it in advance.

Be thoughtful about the information you request or exchange on Twitter

Tweets: Don't ask users to send you personal or private information via a public Tweet. If you need additional personal or private information from a user to provide them with customer service (or other similar use cases), you should ask the user to share such information by Direct Message or another private channel. You might even consider adding a [Direct Message deep link](#) to your Tweet.

Direct Messages: You should only ask users for the minimum amount of information you need to provide them with service. If you need to request or exchange particularly sensitive information (such as credit card information), you should consider directing users to your website or other appropriate channel to do so.

II. Activity-Specific Rules

The activity-specific rules in this section apply to taking specific automated actions on Twitter. Please read these rules carefully, as they outline both permitted and prohibited use cases of automation.

Automated applications or activities that violate these rules, or that facilitate or induce users to violate them, may be subject to enforcement action, including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these rules. As a reminder, you should also carefully review the spam guidelines in the [Twitter Rules](#) to avoid having activities performed by you, your app, or other users through your app or service flagged as spam.

A. Automated Actions Through Another User's Account

Twitter users may authorize your app or service to [access their Twitter account through OAuth](#). A user authorizing your app or service to access their Twitter account through OAuth does not by itself constitute sufficient consent to take automated actions through that user's account.

You may only take automated actions through another Twitter user's account if you:

clearly describe to the user the types of automated actions that will occur; receive express consent from the user to take those automated actions; and immediately honor a user's request to opt-out of further automated actions.

If you substantially change the purpose or functionality of your app or service, you must re-obtain express consent from the user to take automated action through their account before doing so.

These requirements apply to any automated action taken through another Twitter user's account, including posting Tweets, sending Direct Messages, deleting Tweets or Direct Messages, or following/unfollowing other accounts. For applications that offer users the ability to delete Tweets in a bulk or automated manner, you must also clearly state that Tweets are not recoverable once deleted.

B. Automated Tweets

1. Posting automated Tweets

Automated Tweets that cross-post outside information:

You may post automated Tweets based on sources of outside information — such as an RSS feed, weather data, etc. — as long as you are sufficiently authorized to publish such information.

Other automated Tweets (excluding mentions or replies):

Provided you comply with all other rules, you may post automated Tweets for entertainment, informational, or novelty purposes. As a reminder, accounts posting duplicative, spammy, or otherwise prohibited content may be subject to suspension.

2. Posting automated mentions and replies

The reply and mention functions are intended to make communication between Twitter users easier. Automating these

actions to reach many users on an unsolicited basis is an abuse of the feature, and is not permitted. For example, sending automated replies to Tweets based on keyword searches alone is not permitted. Spammy or duplicative use of mentions and replies may result in enforcement action, such as the removal of your Tweets from Search or the suspension of your app or account.

However, you may send automated replies or mentions to Twitter users so long as:

in advance of sending the automated reply, the recipient or mentioned user(s) have requested or have clearly indicated an intent on Twitter to be contacted by you (i.e. opted in), for example by replying to a Tweet from your account, or by sending you a Direct Message;

you provide a clear and easy way for such users to opt-out of receiving automated replies and mentions, and promptly honor all such opt-out requests;

you only send one automated reply or mention per user interaction; and

the automated reply or mention is a reply to the user's original Tweet (if your campaign is based on users posting a reply to your Tweet).

Opt-in techniques and indications of user intent take many different forms, depending on the specifics of your use case and implementation. Some examples include:

A Tweet from your account that clearly indicates that a user taking a specific action on that Tweet (such as Retweeting it) will opt the user into receiving an automated response.

A mention of your account by the user in a manner suggesting the user clearly wishes or intends to receive a response. If you want to run an auto-reply campaign with a campaign- or use-case-specific hashtag, users should also mention you in their Tweets.

Note that a user following your account is not on its own a sufficient indication of user intent to receive an automated response.

: Advertisers, publishers, and brands using auto-response campaigns must request approval from Twitter and may be subject to additional rules. Please reach out to your account or partner manager for assistance. If you do not

have a partner manager but are interested in managed client services, you can learn more [here](#).

Additionally, we recommend that any accounts that will communicate with

users via automated mentions or replies:

Note

1. Appropriately filter responses based on potentially sensitive language in user handles,

display names, and Tweet text, as well as potentially sensitive media;

2. Check that the Tweet you are mentioning or replying to still exists (for example, using the statuses/lookup endpoint on the Twitter API).

C. Automated Direct Messages

Sending automated Direct Messages to users

You may not send unsolicited Direct Messages in a bulk or automated manner, and should be thoughtful about the frequency with which you contact users via Direct Message.

You may send automated Direct Messages to users so long as:

in advance of sending the Direct Message, the recipient(s) have requested or have clearly indicated an intent on Twitter to be contacted by you via Direct Message, for example by sending you a Direct Message; **and**

you provide a clear and easy way for such users to opt-out of receiving automated Direct Messages, and promptly honor all such opt-out requests.

The fact that a user is technically able to receive a Direct Message from you (e.g. because the user follows you, has enabled the ability to receive Direct Messages from any account, or because the user is in a pre-existing Direct Message conversation with you) does not necessarily mean they have requested or expect to receive automated Direct Messages from you.

Interacting with users via Direct Message

Per the Ground Rules, remember to be thoughtful about the amount and type of information you request or exchange with users via Direct Messages. If you will be asking a user to provide personal or private information via an automated Direct Message, you must clearly explain how you will use the information you're collecting. Consider including a link to your privacy policy in your Direct Message to the user, as well as in your Twitter profile bio.

Don't publicly share information received in a Direct Message conversation with a user without first obtaining explicit consent from the user. For example, if a user asks you via Direct Message about a purchase they made from you, you may not mention the user in a public Tweet that includes information about their purchase unless you have the user's explicit consent to do so.

After a user-initiated interaction ends, don't send additional follow-up Direct Messages or mention users in a Tweet unless you get permission from the user.

D. Automated actions you take on Tweets or accounts

Automated likes: You may not like Tweets in an automated manner.

Automated Retweets: Provided you comply with all other rules, you may Retweet or Quote Tweet in an automated manner for entertainment, informational, or novelty purposes. Automated Retweets often lead to negative user experiences, and bulk, aggressive, or spammy Retweeting is a violation of the [Twitter Rules](#).

Automated following/unfollowing: You may not follow or unfollow Twitter accounts in a bulk, aggressive, or indiscriminate manner. Aggressive following is a violation of the [Twitter Rules](#). Please also review our [following rules and best practices](#) to ensure you are in compliance. Note that applications that claim to get users more followers are also prohibited under the [Twitter Rules](#).

Automated adding to lists or collections: You may not add Twitter users to lists or add Tweets to collections in a bulk or indiscriminate manner. Adding a large number of unrelated users to lists is a violation of the [Twitter Rules](#).

About search rules and restrictions

This article details the rules for using Twitter search, and explains reasons why content may not show up in Twitter search.

Why doesn't all content show up in Twitter search?

In order to keep your search results relevant, Twitter filters search results for quality Tweets and accounts. Material that jeopardizes search quality or creates a bad search experience for other people may be automatically removed from Twitter search. Read more about this [here](#).

General rules for using Twitter search

The [Twitter Rules](#) explain what behaviors are permitted on Twitter. In addition to these rules, we've included some tips below to keep your content relevant (and your Tweets in Twitter search).

Please do not:

Repeatedly post duplicate or near-duplicate content (links or Tweets). Abuse trending topics or hashtags (topic words with a # sign).

Send automated Tweets or replies.

Use bots or applications to post similar messages based on keywords. Post similar messages over multiple accounts.

Aggressively follow and unfollow people.

Twitter may automatically remove accounts engaging in these behaviors from search (or even suspend in some cases) in order to ensure the best experience for everyone.

Search results filtering

To learn how to filter your search results, including how to disable or enable safe search, read our [Using Twitter search](#) article.

Search results filtering in third-party applications

Accounts may be filtered from search for posting similar messages over several accounts. Your account may be filtered from search if you are logged into a third-party application that is updating many accounts with similar or duplicate content.

Rules for participating in contests on Twitter

Businesses sometimes host contests and offer prizes to Retweet an update, following a particular account, or posting Tweets with a specific hashtag topic or reply. If you're running a contest, check out this [article](#) for guidelines for contests on Twitter.

If you repeatedly post near-duplicate updates or duplicate links in order to enter contests, you may be filtered from search. Please keep in mind that you are responsible for the updates that you post to your Twitter account.

What to do if you are missing from search

Please see this [article](#) for more information related to this issue.

If you are being filtered from search, rest assured that your followers will still receive your updates. If you post a reply or mention to a specific person, it will be delivered to that person.

Having trouble?

You can learn more about searching by keyword, location, and trending topics in our [Using Twitter search](#) help article.

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. These limitations are set forth in the Twitter Rules below.

The Twitter Rules (along with all incorporated policies), [Privacy Policy](#), and [Terms of Service](#) collectively make up the "Twitter User Agreement" that governs a user's access to and use of Twitter's services.

All individuals accessing or using Twitter's services must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in Twitter taking one or more of the following enforcement actions:

- requiring you to delete prohibited content before you can again create new posts and interact with other Twitter users;
- temporarily limiting your ability to create posts or interact with other Twitter users;
- asking you to verify account ownership with a phone number or email address; or
- permanently suspending your account(s).

If you attempt to evade a permanent suspension by creating new accounts, we will suspend your new accounts.

Please note that we may need to change these Rules from time to time and reserve the right to do so. The most current version is always available at: <https://twitter.com/rules>.

The policies set forth in these Twitter Rules govern organic content on our platform. To learn more about the rules which govern ads and promoted content, please review our [Ads policies](#).

Content Boundaries and Use of Twitter

Intellectual property

Trademark: We reserve the right to suspend accounts or take other appropriate action when someone's brand or trademark, including business name and/or logo, is used in a manner that may mislead or confuse others about your brand affiliation. Read more about our [trademark policy](#) and [how to report a violation](#).

Copyright: We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in our [Terms of Service](#). Read more about our [copyright policy](#).

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Learn more about our [media policy](#).

Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images. Additionally, Twitter may sometimes require you to remove excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about [how to make such a request](#), and [how to mark your media as sensitive](#).

Unlawful use

You may not use our service for any unlawful purposes or in furtherance of illegal activities. By using Twitter, you agree to comply with all applicable laws governing your online conduct and content.

Trends

At times, we may prevent certain content from trending. This includes content that violates the Twitter Rules, as well as content that may attempt to manipulate trends. Read more about [what we allow and do not allow to trend](#).

Misuse of Twitter badges

You may not use badges, including but not limited to the “promoted” or “verified” Twitter badges, unless provided by Twitter. Accounts using unauthorized badges as part of their profile photos, header photos, display names, or in any way that falsely implies affiliation with Twitter or authorization from Twitter to display these badges, may be suspended.

Misuse of usernames

Selling usernames: You may not buy or sell Twitter usernames.

Username squatting: You may not engage in username squatting. Some of the factors we take into consideration when determining whether conduct is username squatting include:

- the number of accounts created;
- the creation of accounts for the purpose of preventing others from using those account names;
- the creation of accounts for the purpose of selling those accounts; and
- the use of third-party content feeds to update and maintain accounts under the names of those third parties.

Please note that Twitter may also remove accounts that are inactive for more than six months. Learn more about [username squatting](#).

Abusive Behavior

We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we prohibit behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Context matters when evaluating for abusive behavior and determining appropriate enforcement actions. Factors we may take into consideration include, but are not limited to whether:

- the behavior is targeted at an individual or group of people;
- the report has been filed by the target of the abuse or a bystander;
- the behavior is newsworthy and in the legitimate public interest.

Violence and physical harm

Violence: You may not make specific threats of violence or wish for the serious physical harm, death, or disease of an individual or group of people. This includes, but is not limited to, threatening or promoting terrorism. You also may not affiliate with organizations that — whether by their own statements or activity both on and off the platform — use or promote violence against civilians to further their causes. We will begin enforcing this rule around affiliation with such organizations on December 18, 2017.

Suicide or self-harm: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Child sexual exploitation: You may not promote child sexual exploitation. Learn more about our zero-tolerance [child sexual exploitation policy](#).

Abuse and hateful conduct

Abuse: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Unwanted sexual advances: You may not direct abuse at someone by sending unwanted sexual content, objectifying them in a sexually explicit manner, or otherwise engaging in sexual misconduct.

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. Read more about our [hateful conduct policy](#).

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category. We will begin enforcing this rule on December 18, 2017.

Private information and intimate media

Private information: You may not publish or post other people's private information without their express authorization and permission. Definitions of private information may vary depending on local laws. Read more about our [private information policy](#).

Intimate media: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. Read more about intimate media on Twitter.

Threats to expose / hack: You may not threaten to expose someone's private information or intimate media. You also may not threaten to hack or break into someone's digital information.

Impersonation

You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. While you may maintain parody, fan, commentary, or newsfeed accounts, you may not do so if the intent of the account is to engage in spamming or abusive behavior. Read more about our [impersonation policy](#).

Spam and Security

We strive to protect people on Twitter from technical abuse and spam.

To promote a stable and secure environment on Twitter, you may not do, or attempt to do, any of the following while accessing or using Twitter:

- Access, tamper with, or use non-public areas of Twitter, Twitter's computer systems, or the technical delivery systems of Twitter's providers (except as expressly permitted by the Twitter Bug Bounty program).
- Probe, scan, or test the vulnerability of any system or network, or breach or circumvent any security or authentication measures (except as expressly permitted by the Twitter Bug Bounty program).
- Access or search, or attempt to access or search, Twitter by any means (automated or otherwise) other than through our currently available, published interfaces that are provided by Twitter (and only pursuant to the applicable terms and conditions), unless you have been specifically allowed to do so in a separate agreement with Twitter. Note that crawling Twitter is permissible if done in accordance with the provisions of the robots.txt file; however, scraping Twitter without our prior consent is expressly prohibited.
- Forge any TCP/IP packet header or any part of the header information in any email or posting, or in any way use Twitter to send altered, deceptive, or false source-identifying information.
- Interfere with or disrupt the access of any user, host or network, including, without limitation, sending a virus, overloading, flooding, spamming, mail-bombing Twitter's services, or by scripting the creation of content in such a manner as to interfere with or create an undue burden on Twitter.

Any accounts engaging in the following activities may be temporarily locked or subject to permanent suspension:

- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use Twitter's services for the purpose of spamming anyone. Spam is generally defined on Twitter as bulk or aggressive activity that attempts to manipulate or disrupt Twitter or the experience of users on Twitter to drive traffic or attention to unrelated accounts, products, services, or initiatives. Some of the factors that we take into account when determining what conduct is considered to be spamming include:
 - if you have followed and/or unfollowed a large number of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if your Tweets or Direct Messages consist mainly of links shared without commentary;
 - if a large number of people have blocked you in response to high volumes of untargeted, unsolicited, or duplicative content or engagements from your account;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicative or substantially similar content, replies, or mentions over multiple accounts or multiple duplicate updates on one account, or create duplicate or substantially similar accounts;
 - if you post multiple updates to a trending or popular topic with an intent to subvert or manipulate the topic to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you send large numbers of unsolicited replies or mentions;
 - if you add users to lists in a bulk or aggressive manner;
 - if you are randomly or aggressively engaging with Tweets (e.g., likes, Retweets, etc.) or users (e.g., following, adding to lists or Moments, etc.) to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you repeatedly post other people's account information as your own (e.g., bio, Tweets, profile URL, etc.);
 - if you post misleading, deceptive, or malicious links (e.g., affiliate links, links to malware/clickjacking pages, etc.);
 - if you create fake accounts, account interactions, or impressions;
 - if you sell, purchase, or attempt to artificially inflate account interactions (such as followers, Retweets, likes, etc.); and
 - if you use or promote third-party services or apps that claim to get you more followers, Retweets, or likes (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers or engagements to your account or Tweets).

Please see our support articles on [following rules and best practices](#) and [automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts may be permanently suspended.

Content Visibility

Accounts under investigation or which have been detected as sharing content in violation of these Rules may have their account or Tweet visibility limited in various parts of Twitter, including search. To learn more about situations in which content may be restricted on Twitter, please see our support article on [search rules and restrictions](#).

Trademark policy

Updated June 30, 2016

What is a Trademark Policy Violation on Twitter?

Using another's trademark in a manner that may mislead or confuse others about your brand affiliation may be a violation of our trademark policy.

How Does Twitter Respond to Trademark Reports?

We review each report and do the following:

- If we determine that there is a violation of our trademark policy, we may suspend the account. In some instances, we may give the account holder an opportunity to comply with our policies before suspending the account. For more information, see our [Parody, commentary, and fan account policy](#).
- We may also provide the account holder with your name and information included in the copy of the report.
- If you would like to report a trademark violation in a [Twitter Ad](#), please visit our [Trademark Policy for Promoted Ads](#) to learn more.

Please note that Twitter only investigates requests that are submitted by the trademark holder or by an authorized representative of the trademark holder.

What is not a Trademark Policy Violation?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples include:

- Using a trademark in a way that is outside the scope of the trademark registration (e.g. territory, or goods and services identified in the registration).
- Nominative and other fair uses of trademarks are protected uses under our trademark policy, so long as the account is clearly distinguished from the trademark owner. This includes use by resellers in certain regions and accounts engaging in parody, commentary, or news. For more information, see our [Parody, commentary, and fan account policy](#).

How do I Report a Trademark Policy Violation?

Submit trademark reports through our [Trademark Issue support form](#). Using the appropriate form ensures we have all the information necessary to process your report quickly. Please provide all information requested in the form. If you submit a report with incomplete information, we will follow up with you to request that information. Please note that this will result in a delay in processing your report.

If you have trouble using the form, try updating your browser or switching to another browser.

Copyright policy

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged

infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notification.

When Should I File a Counter-notice?

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);

3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Twitter media policy

Twitter Rules: Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images.

Twitter Rules: You may not use hateful images or symbols in your profile image or profile header.

Rationale

Some forms of graphic violence, adult content, or hateful imagery may be permitted in Tweets when they are marked as [sensitive media](#). However, you may not include this type of content in live video, or in profile or header images.

While we want people to feel free to share media that reflects their creativity or individuality, or to show what's happening in the world, we will take action when it crosses the line into abuse towards a person, group, or protected category.

Finally, we may sometimes require you to remove media containing excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about how to make such a request [here](#).

When this applies

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. Some examples include, but are not limited to, depictions of:

- the moment at which someone dies
- gruesome crime or accident scenes
- bodily harm, torture, dismemberment, or mutilation

We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity (including close-ups of genitals, buttocks, or breasts)
 - Please note that exceptions may be made for artistic, medical, health, or educational content. Breastfeeding content does not need to be marked as sensitive.
- simulating a sexual act
- intercourse or any sexual act (may involve humans, humanoid animals, cartoons, or anime)

If you see media you think may be in violation of our policy, report it. You can read about how to report sensitive media [here](#).

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, or ethnicity/national origin.

Some examples of hateful imagery include, but are not limited to:

- Symbols historically associated with hate groups (for example, the Nazi swastika)
- Images depicting others as less than human or altered to include hateful symbols
- Altered image references to a mass murder that targeted a protected category

Do I need to be the target of this content in order for it to be a violation of the Twitter Rules?
No, we review both first-person and bystander reports of such content.

Consequences

When adult content, graphic violence, or hateful imagery appears in Tweets, we may place this content behind an interstitial advising viewers to be aware that they will see sensitive media if they click through. This allows us to identify potentially sensitive content that some people may not wish to see. Learn more about [how to control whether you see sensitive media](#).

However, if such content appears in live video, header, or profile images, it will violate our media policy. The consequences for violating this policy depend on the severity of the violation and the person's previous record of violations.

The first time someone violates this policy they will be required to remove the imagery and may be asked to verify their account ownership. Subsequent violations could result in permanent suspension. Learn more about [our range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

How to contact Twitter about media concerning a deceased family member

In order to respect the wishes of loved ones, Twitter will remove imagery of deceased individuals in certain circumstances. Immediate family members and other authorized individuals may request the removal of images

or video of deceased individuals, from when critical injury occurs to the moments before or after death, by sending a request to Twitter Inc. via our [privacy form](#). When reviewing such media removal requests, Twitter considers public interest factors such as the newsworthiness of the content and may not be able to honor every request. Note: Find information about contacting Twitter to [deactivate the account of a deceased or incapacitated user](#).

How to mark media as sensitive links to Media policy

Twitter trends FAQs

Where can I find trends?

Trends are available on the Twitter for iOS app, the Twitter for Android app, and [twitter.com](#).

On Twitter's mobile apps, you can find trends listed under the trends section of the **Explore** tab. When signed in to [twitter.com](#) on a desktop or laptop computer, trends are listed in many places, including the Home timeline, Notifications, search results, and profile pages.

How are trends determined?

Trends are determined by an algorithm and, by default, are tailored for you based on who you follow, your interests, and your location. This algorithm identifies topics that are popular now, rather than topics that have been popular for a while or on a daily basis, to help you discover the hottest emerging topics of discussion on Twitter.

You can choose to see trends that are not tailored for you by selecting a specific trends location on [twitter.com](#), iOS, or Android (instructions below). Location trends identify popular topics among people in a specific geographic location.

Note: The number of Tweets that are related to the trends is just one of the factors the algorithm looks at when ranking and determining trends. Algorithmically, trends and hashtags are grouped together if they are related to the same topic. For

instance, #MondayMotivation and #MotivationMonday may both be represented by #MondayMotivation.

What does the # sign mean?

You may notice that some trends have # sign before the word or phrase. This is called a **hashtag** and is included specifically in Tweets to mark them as relating to a topic, so that people can follow the conversation in search.

Will I see world and local events in my trends?

Yes, when signed into your account on the twitter.com, iOS, or Android, trends will be tailored for you based on your location and who you follow. There will be many world and local news events and conversations that will appear in your trends regardless of your personalization.

Can I see trends for a specific location?

Yes, you can view trends for a specific location on twitter.com, iOS, or Android.

View trends via twitter.com:

To view trends for a specific location, click **Change** in your trends box. You can choose to **Keep tailored trends** based on your location and who you follow, or you can select **Change** to choose a nearby location or search locations.

Adjust your trends settings via Twitter for iOS:

1. In the top menu, tap your **profile** icon.
2. Tap **Settings and privacy**, then tap **Content preferences**.
3. Under **Explore**, tap **Trends**.
4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.
5. When the feature is disabled, you can change your location by tapping **Change**

location. In the search box, type your desired trends location or choose from the list.

Adjust your trends settings via Twitter for Android:

1. In the top menu, you will either see a **navigation menu** icon or your **profile** icon.

Tap whichever icon you have, and select **Settings and privacy**.

2. Tap **Content preferences**.
3. Under **Explore**, tap **Trends**.
4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.
5. When the feature is disabled, you can change your location by tapping **Change location**. In the search box, type your desired trends location or choose from the list.

Note: If you don't find your preferred country or city, it means we're not receiving enough Tweets from that geographical area to create a list. You can look up local Tweets on any topic by using [advanced Twitter search operators](#).

What happens when I click or tap on a trend?

Clicking or tapping any of the trends takes you to the Twitter search results for that trend. You'll see all Tweets including that phrase or hashtag. To see what people are saying about a previous trend, perform a [search for that keyword](#).

How can I participate in a trend?

Simply post a Tweet including the exact word or phrase as it appears in the trends list (with the hashtag, if you see one). Due to the large number of people Tweeting about these specific trends, you may not always be able to find your particular Tweet in search, but your followers will always see your Tweets.

Note: Twitter also filters searches for quality. To understand what types of behavior could cause your Tweets to be filtered out, read about our [search rules and restrictions](#).

Are there rules for trends?

Yes. We want trends to promote healthy discussions on Twitter. This means that at times, we may prevent certain content from trending. These include trends that:

Contain profanity or adult/graphic references.

Incite hate on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease.

Violate the [Twitter Rules](#).

In some cases, we may also consider the newsworthiness of the content, or if it is in the public interest when evaluating potential violations. In these cases, the content might continue to trend on our platform.

Note that even if we prevent the hashtag or the content from appearing on the list of trending topics, you may still be able to access conversations around that content on Twitter.

Do trends have descriptions?

Below some trends, you may see the approximate number of Tweets associated with the trend, or personalized context such as who in your network is Tweeting about it.

Additionally, when a Moment is found to be relevant to a trend, we may use the title of the Moment as the trend's description.

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Child sexual exploitation policy

Twitter removes content promoting child sexual exploitation

We do not tolerate child sexual exploitation on Twitter. When we are made aware of links to images of or content promoting child sexual exploitation they will be removed from the site without further notice and reported to The National Center for Missing & Exploited Children ("NCMEC"); we permanently suspend accounts promoting or containing updates with links to child sexual exploitation.

Reporting content promoting child sexual exploitation

If you think you've found a Twitter account distributing or promoting child sexual exploitation while using Twitter, please click [here](#) to report it. Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. To find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with the National Center for Missing and Exploited Children [here](#), or your local law enforcement. If you believe that the content was hosted on a website outside of the United States, you can report it [here](#) on the International Association of Internet Hotlines website.

NOTE: Please do NOT tweet, retweet or repost child sexual exploitation for any reason. Click [here](#) to report it to us immediately and we will take steps to remove it.

Law Enforcement Agencies

If you are a member of a law enforcement agency seeking information regarding Twitter accounts promoting or distributing child sexual exploitation, please see our [Guidelines for Law Enforcement](#) for information on our process; we also suggest you contact the National Center for Missing and Exploited Children (NCMEC) [here](#), regardless of jurisdiction, to determine whether or not information on those accounts is already on file.

For general law enforcement inquiries, please see our [Guidelines for Law Enforcement](#) for more information.

Hateful conduct policy

Freedom of expression means little if voices are silenced because people are afraid to speak up. We do not tolerate behavior that harasses, intimidates, or uses fear to silence another person's voice. If you see something on Twitter that violates these rules, please report it to us.

How our policy works

As explained in the Twitter Rules,

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual

orientation, gender, gender identity, religious affiliation, age, disability, or disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Examples of what we do not tolerate includes, but is not limited to behavior that harasses individuals or groups of people with:

violent threats;

wishes for the physical harm, death, or disease of individuals or groups;

references to mass murder, violent events, or specific means of violence in which/with which such groups have been the primary targets or victims;

behavior that incites fear about a protected group;

repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone.

How our enforcement works Context matters.

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. While we accept reports of violations from anyone, sometimes we also need to hear directly from the target to ensure that we have proper context.

The number of reports we receive does not impact whether or not something will be removed. However, it may help us prioritize the order in which it gets reviewed.

We focus on behavior.

We enforce policies when someone reports behavior that is abusive and targets an entire protected group and/or individuals who may be members.

This targeting can happen in any manner (for example, @mentions, tagging a photo, and more).

We have a range of enforcement options.

The consequences for violating our rules vary depending on the severity of the violation and the person's previous record of violations. For example, we may ask someone to remove the offending Tweet before they can Tweet again. For other cases, we may suspend an account.

Report private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Read about our [policy regarding private information posted on Twitter](#).

Do I need to have a Twitter account to make a report?

No, you do not need to have a Twitter account to file a report that your private and confidential information has been posted on Twitter.

Can I report unauthorized posting of Copyrighted material?

If your copyrighted material has been Tweeted without your permission, please review our [Copyright and DMCA policy](#).

What can I do if my information is posted on other websites?

We cannot take down any information that is posted on a website other than Twitter. If your information is posted on another website, please contact the help center or support staff for the website where your information is posted.

How do I report that my private information is posted on Twitter?

You can report that your private information has been posted on Twitter directly from the offending Tweet(s).

Learn [how to report individual Tweets](#).

File a [report that someone has posted your private information](#).

What happens when Twitter receives a valid report?

When we receive a complete and valid report that private information has been posted on Twitter, we'll investigate the account and Tweets reported. We will review where, if anywhere, the information has been made publicly available before taking action on the account or Tweets. If the information you reported was previously posted elsewhere on the Internet, it is not a violation of our policy and we will not take action.

Learn more about [protecting your private information on Twitter and other websites](#).

Read our [policy regarding private information posted on Twitter](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter impersonation policy.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to Twitter's [parody, commentary, and fan account policy](#) for more information about these types of accounts.

Find more information about [reporting impersonation on Twitter](#).

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000

people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation rules

Updated November 3, 2017

This page is primarily intended for **developers**.

For Twitter users: You are ultimately responsible for the actions taken with your account, or by applications associated with your account. Before authorizing a third-party application to access or use your account, make sure you've thoroughly investigated the application and understand what it will do. If automated activity on your account violates the [Twitter Rules](#) or these automation rules, Twitter may take action on your account, including [filtering your Tweets from search results](#) or suspending your account.

For more information on third-party applications, please see our article on [connecting and revoking third-party applications](#).

If you are a developer and have questions about these Automation Rules, please reach out via <https://twittercommunity.com/c/rules-and-policies>.

If your application has been suspended or restricted from performing write actions, you may submit an appeal at </forms/platform>.

I. Ground Rules Do!

Build solutions that automatically broadcast helpful information in Tweets.

Run creative campaigns that auto-reply to users who engage with your content.

Build solutions that automatically respond to users in Direct Messages.

Try new things that help people (and comply with our rules).

Make sure your application provides a good user experience and performs well — and confirm that remains the case over time.

Don't!

Violate these or other policies. Be extra mindful of our rules about abuse and user privacy.

Abuse the Twitter API or attempt to circumvent rate limits.

Use non-API-based forms of automation, such as scripting the Twitter website. The use of these techniques may result in the permanent suspension of your account.

Spam or bother users, or otherwise send them unsolicited messages.

A. The Twitter Rules and the Developer Agreement and Policy

As with all activity on Twitter, automated activity is subject to the [Twitter Rules](#) and, if you're a developer using the Twitter API, the [Developer Agreement and Policy](#).

You should carefully review these policies to ensure that your automated activity is compliant. Automated applications or activities that violate these policies, or that facilitate or induce users to violate them, may be subject to enforcement action, potentially including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these policies.

Although all aspects of the Twitter Rules and the Developer Agreement and Policy apply to automated activity, you should keep the following rules top of mind:

Spamming: You may not send automated Tweets or Direct Messages that are spam, or otherwise engage in spamming activity. Some examples of spammy behavior to avoid with automation include:

Trending topics: You may not automatically post about trending topics on Twitter, or use automation to attempt to influence or manipulate trending topics.

Multiple posts/accounts: You may not post duplicative or substantially similar Tweets on one account or over multiple accounts you operate.

Duplicate accounts: You may not create and/or automate multiple accounts for duplicative or substantially similar use cases.

However, automating multiple accounts for related but non-duplicative use cases is permitted. For example, you may automate separate accounts to Tweet when the Hubble Space Telescope passes over different cities, such as [San Francisco](#) or [Hong Kong](#).

Misleading links: You may not send automated Tweets or Direct Messages containing links that are misleading, including links that maliciously or deceptively redirect through landing pages or ad pages before displaying the final content.

Sensitive media: Automated Tweets and Direct Messages must comply with the [Twitter media policy](#), and you should mark your account as potentially sensitive if you intend to post graphic, pornographic, or potentially sensitive media.

Abusive behavior: You may not engage in any automated activity that encourages, promotes, or incites abuse, violence, hateful conduct, or harassment, on or off Twitter.

Private information: You may not post private or confidential information about a person without their prior express authorization.

B. Other Ground Rules for Automated Activity

In addition to the policies above, the following ground rules apply to all automated activity on Twitter:

Don't surprise or mislead users: Automated activity should honor users' expectations. Ask for the user's permission before taking an action if you aren't sure.

Mature content or profanity: Don't Direct Message, mention, or reply to users with potentially sensitive content (including profanity), unless they've clearly indicated an intent to receive it in advance.

Be thoughtful about the information you request or exchange on Twitter

Tweets: Don't ask users to send you personal or private information via a public Tweet. If you need additional personal or private information from a user to provide them with customer service (or other similar use cases), you should ask the user to share such information by Direct Message or another private channel. You might even consider adding a [Direct Message deep link](#) to your Tweet.

Direct Messages: You should only ask users for the minimum amount of information you need to provide them with service. If you need to request or exchange particularly sensitive information (such as credit card information), you should consider directing users to your website or other appropriate channel to do so.

II. Activity-Specific Rules

The activity-specific rules in this section apply to taking specific automated actions on Twitter. Please read these rules carefully, as they outline both permitted and prohibited use cases of automation.

Automated applications or activities that violate these rules, or that facilitate or induce users to violate them, may be subject to enforcement action, including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access the Twitter API based on violations of these rules. As a reminder, you should also carefully review the spam guidelines in the [Twitter Rules](#) to avoid having activities performed by you, your app, or other users through your app or service flagged as spam.

A. Automated Actions Through Another User's Account

Twitter users may authorize your app or service to [access their Twitter account through OAuth](#). A user authorizing your app or service to access their Twitter account through OAuth does not

by itself constitute sufficient consent to take automated actions through that user's account.

You may only take automated actions through another Twitter user's account if you:

clearly describe to the user the types of automated actions that will occur; receive express consent from the user to take those automated actions; and immediately honor a user's request to opt-out of further automated actions.

If you substantially change the purpose or functionality of your app or service, you must re-obtain express consent from the user to take automated action through their account before doing so.

These requirements apply to any automated action taken through another Twitter user's account, including posting Tweets, sending Direct Messages, deleting Tweets or Direct Messages, or following/unfollowing other accounts. For applications that offer users the ability to delete Tweets in a bulk or automated manner, you must also clearly state that Tweets are not recoverable once deleted.

B. Automated Tweets

1. Posting automated Tweets

Automated Tweets that cross-post outside information:

You may post automated Tweets based on sources of outside information — such as an RSS feed, weather data, etc. — as long as you are sufficiently authorized to publish such information.

Other automated Tweets (excluding mentions or replies):

Provided you comply with all other rules, you may post automated Tweets for entertainment, informational, or novelty purposes. As a reminder, accounts posting duplicative, spammy, or otherwise prohibited content may be subject to suspension.

2. Posting automated mentions and replies

The reply and mention functions are intended to make communication between Twitter users easier. Automating these actions to reach many users on an unsolicited basis is an abuse of the feature, and is not permitted. For example, sending automated replies to Tweets based on keyword searches alone is not permitted. Spammy or duplicative use of mentions and replies may result in enforcement action, such as the removal of your Tweets from Search or the suspension of your app or account.

However, you may send automated replies or mentions to Twitter users so long as:

in advance of sending the automated reply, the recipient or mentioned user(s) have requested or have clearly indicated an intent on Twitter to be contacted by you (i.e. opted in), for example by replying to a Tweet from your account, or by sending you a Direct Message;

you provide a clear and easy way for such users to opt-out of receiving automated replies and mentions, and promptly honor all such opt-out requests;

you only send one automated reply or mention per user interaction; and

the automated reply or mention is a reply to the user's original Tweet (if your campaign is based on users posting a reply to your Tweet).

Opt-in techniques and indications of user intent take many different forms, depending on the specifics of your use case and implementation. Some examples include:

A Tweet from your account that clearly indicates that a user taking a specific action on that Tweet (such as Retweeting it) will opt the user into receiving an automated response.

A mention of your account by the user in a manner suggesting the user clearly wishes or intends to receive a response. If you want to run an auto-reply campaign with a campaign- or use-case-specific hashtag, users should also mention you in their Tweets.

Note that a user following your account is not on its own a sufficient indication of user intent to receive an automated response.

: Advertisers, publishers, and brands using auto-response campaigns must request approval from Twitter

and may be subject to additional rules. Please reach out to your account or partner manager for assistance. If you do not have a partner manager but are interested in managed client services, you can learn more [here](#).

Additionally, we recommend that any accounts that will communicate with

users via automated mentions or replies:

Note

1. Appropriately filter responses based on potentially sensitive language in user handles,

display names, and Tweet text, as well as potentially sensitive media;

2. Check that the Tweet you are mentioning or replying to still exists (for example, using the statuses/lookup endpoint on the Twitter API).

C. Automated Direct Messages

Sending automated Direct Messages to users

You may not send unsolicited Direct Messages in a bulk or automated manner, and should be thoughtful about the frequency with which you contact users via Direct Message.

You may send automated Direct Messages to users so long as:

in advance of sending the Direct Message, the recipient(s) have requested or have clearly indicated an intent on Twitter to be contacted by you via Direct Message, for example by sending you a Direct Message; **and**

you provide a clear and easy way for such users to opt-out of receiving automated Direct Messages, and promptly honor all such opt-out requests.

The fact that a user is technically able to receive a Direct Message from you (e.g. because the user follows you, has enabled the ability to receive Direct Messages from any account, or because the user is in a pre-existing Direct Message conversation with you) does not necessarily mean they have requested or expect to receive automated Direct Messages from you.

Interacting with users via Direct Message

Per the Ground Rules, remember to be thoughtful about the amount and type of information you request or exchange with users via Direct Messages. If you will be asking a user to provide personal or private information via an automated Direct Message, you must clearly explain how you will use the information you're collecting. Consider including a link to your privacy policy in your Direct Message to the user, as well as in your Twitter profile bio.

Don't publicly share information received in a Direct Message conversation with a user without first obtaining explicit consent from the user. For example, if a user asks you via Direct Message about a purchase they made from you, you may not mention the user in a public Tweet that includes information about their purchase unless you have the user's explicit consent to do so.

After a user-initiated interaction ends, don't send additional follow-up Direct Messages or mention users in a Tweet unless you get permission from the user.

D. Automated actions you take on Tweets or accounts

Automated likes: You may not like Tweets in an automated manner.

Automated Retweets: Provided you comply with all other rules, you may Retweet or Quote Tweet in an automated manner for entertainment, informational, or novelty purposes. Automated Retweets often lead to negative user experiences, and bulk, aggressive, or spammy Retweeting is a violation of the [Twitter Rules](#).

Automated following/unfollowing: You may not follow or unfollow Twitter accounts in a bulk, aggressive, or indiscriminate manner. Aggressive following is a violation of the [Twitter Rules](#). Please also review our [following rules and best practices](#) to ensure you are in compliance. Note that applications that claim to get users more followers are also prohibited under the [Twitter Rules](#).

Automated adding to lists or collections: You may not add Twitter users to lists or add Tweets to collections in a bulk or indiscriminate manner. Adding a large number of unrelated users to lists is a violation of the [Twitter Rules](#).

[About search rules and restrictions](#)

This article details the rules for using Twitter search, and explains reasons why content may not show up in Twitter search.

Why doesn't all content show up in Twitter search?

In order to keep your search results relevant, Twitter filters search results for quality Tweets and accounts. Material that jeopardizes search quality or creates a bad search experience for other people may be automatically removed from Twitter search. Read more about this [here](#).

General rules for using Twitter search

The [Twitter Rules](#) explain what behaviors are permitted on Twitter. In addition to these rules, we've included some tips below to keep your content relevant (and your Tweets in Twitter search).

Please do not:

- Repeatedly post duplicate or near-duplicate content (links or Tweets).
- Abuse trending topics or hashtags (topic words with a # sign).
- Send automated Tweets or replies.
- Use bots or applications to post similar messages based on keywords.
- Post similar messages over multiple accounts.
- Aggressively follow and unfollow people.

Twitter may automatically remove accounts engaging in these behaviors from search (or even suspend in some cases) in order to ensure the best experience for everyone.

Search results filtering

To learn how to filter your search results, including how to disable or enable safe search, read our [Using Twitter search](#) article.

Search results filtering in third-party applications

Accounts may be filtered from search for posting similar messages over several accounts. Your account may be filtered from search if you are logged into a third-party application that is updating many accounts with similar or duplicate content.

Rules for participating in contests on Twitter

Businesses sometimes host contests and offer prizes to Retweet an update, following a particular account, or posting Tweets with a specific hashtag topic or reply. If you're running a contest, check out this [article](#) for guidelines for contests on Twitter.

If you repeatedly post near-duplicate updates or duplicate links in order to enter contests, you may be filtered from search. Please keep in mind that you are responsible for the updates that you post to your Twitter account.

What to do if you are missing from search

Please see this [article](#) for more information related to this issue.

If you are being filtered from search, rest assured that your followers will still receive your updates. If you post a reply or mention to a specific person, it will be delivered to that person.

Having trouble?

You can learn more about searching by keyword, location, and trending topics in our [Using Twitter search](#) help article.

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. These limitations are set forth in the Twitter Rules below.

The Twitter Rules (along with all incorporated policies), [Privacy Policy](#), and [Terms of Service](#) collectively make up the "Twitter User Agreement" that governs a user's access to and use of Twitter's services.

All individuals accessing or using Twitter's services must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in Twitter taking one or more of the following enforcement actions:

- requiring you to delete prohibited content before you can again create new posts and interact with other Twitter users;
- temporarily limiting your ability to create posts or interact with other Twitter users;
- asking you to verify account ownership with a phone number or email address; or
- permanently suspending your account(s).

If you attempt to evade a permanent suspension by creating new accounts, we will suspend your new accounts.

Please note that we may need to change these Rules from time to time and reserve the right to do so. The most current version is always available at: <https://twitter.com/rules>.

The policies set forth in these Twitter Rules govern organic content on our platform. To learn more about the rules which govern ads and promoted content, please review our [Ads policies](#).

Content Boundaries and Use of Twitter

Intellectual property

Trademark: We reserve the right to suspend accounts or take other appropriate action when someone's brand or trademark, including business name and/or logo, is used in a manner that may mislead or confuse others about your brand affiliation. Read more about our [trademark policy](#) and [how to report a violation](#).

Copyright: We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in our [Terms of Service](#). Read more about our [copyright policy](#).

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Learn more about our [media policy](#).

Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images. Additionally, Twitter may sometimes require you to remove excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about [how to make such a request](#), and [how to mark your media as sensitive](#).

Unlawful use

You may not use our service for any unlawful purposes or in furtherance of illegal activities. By using Twitter, you agree to comply with all applicable laws governing your online conduct and content.

Trends

At times, we may prevent certain content from trending. This includes content that violates the Twitter Rules, as well as content that may attempt to manipulate trends. Read more about [what we allow and do not allow to trend](#).

Misuse of Twitter badges

You may not use badges, including but not limited to the “promoted” or “verified” Twitter badges, unless provided by Twitter. Accounts using unauthorized badges as part of their profile photos, header photos, display names, or in any way that falsely implies affiliation with Twitter or authorization from Twitter to display these badges, may be suspended.

Misuse of usernames

Selling usernames: You may not buy or sell Twitter usernames.

Username squatting: You may not engage in username squatting. Some of the factors we take into consideration when determining whether conduct is username squatting include:

- the number of accounts created;
- the creation of accounts for the purpose of preventing others from using those account names;
- the creation of accounts for the purpose of selling those accounts; and
- the use of third-party content feeds to update and maintain accounts under the names of those third parties.

Please note that Twitter may also remove accounts that are inactive for more than six months. Learn more about [username squatting](#).

Abusive Behavior

We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we prohibit behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Context matters when evaluating for abusive behavior and determining appropriate enforcement actions. Factors we may take into consideration include, but are not limited to whether:

- the behavior is targeted at an individual or group of people;
- the report has been filed by the target of the abuse or a bystander;
- the behavior is newsworthy and in the legitimate public interest.

Violence and physical harm

Violence: You may not make specific threats of violence or wish for the serious physical harm, death, or disease of an individual or group of people. This includes, but is not limited to, threatening or promoting terrorism. You also may not affiliate with organizations that — whether by their own statements or activity both on and off the platform — use or promote violence against civilians to further their causes. We will begin enforcing this rule around affiliation with such organizations on December 18, 2017.

Suicide or self-harm: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Child sexual exploitation: You may not promote child sexual exploitation. Learn more about our zero-tolerance [child sexual exploitation policy](#).

Abuse and hateful conduct

Abuse: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Unwanted sexual advances: You may not direct abuse at someone by sending unwanted sexual content, objectifying them in a sexually explicit manner, or otherwise engaging in sexual misconduct.

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. Read more about our [hateful conduct policy](#).

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category. We will begin enforcing this rule on December 18, 2017.

Private information and intimate media

Private information: You may not publish or post other people's private information without their express authorization and permission. Definitions of private information may vary depending on local laws. Read more about our [private information policy](#).

Intimate media: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. Read more about intimate media on Twitter.

Threats to expose / hack: You may not threaten to expose someone's private information or intimate media. You also may not threaten to hack or break into someone's digital information.

Impersonation

You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. While you may maintain parody, fan, commentary, or newsfeed accounts, you may not do so if the intent of the account is to engage in spamming or abusive behavior. Read more about our [impersonation policy](#).

Spam and Security

We strive to protect people on Twitter from technical abuse and spam.

To promote a stable and secure environment on Twitter, you may not do, or attempt to do, any of the following while accessing or using Twitter:

- Access, tamper with, or use non-public areas of Twitter, Twitter's computer systems, or the technical delivery systems of Twitter's providers (except as expressly permitted by the Twitter Bug Bounty program).
- Probe, scan, or test the vulnerability of any system or network, or breach or circumvent any security or authentication measures (except as expressly permitted by the Twitter Bug Bounty program).
- Access or search, or attempt to access or search, Twitter by any means (automated or otherwise) other than through our currently available, published interfaces that are provided by Twitter (and only pursuant to the applicable terms and conditions), unless you have been specifically allowed to do so in a separate agreement with Twitter. Note that crawling Twitter is permissible if done in accordance with the provisions of the robots.txt file; however, scraping Twitter without our prior consent is expressly prohibited.
- Forge any TCP/IP packet header or any part of the header information in any email or posting, or in any way use Twitter to send altered, deceptive, or false source-identifying information.
- Interfere with or disrupt the access of any user, host or network, including, without limitation, sending a virus, overloading, flooding, spamming, mail-bombing Twitter's services, or by scripting the creation of content in such a manner as to interfere with or create an undue burden on Twitter.

Any accounts engaging in the following activities may be temporarily locked or subject to permanent suspension:

- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use Twitter's services for the purpose of spamming anyone. Spam is generally defined on Twitter as bulk or aggressive activity that attempts to manipulate or disrupt Twitter or the experience of users on Twitter to drive traffic or attention to unrelated accounts, products, services, or initiatives. Some of the factors that we take into account when determining what conduct is considered to be spamming include:
 - if you have followed and/or unfollowed a large number of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if your Tweets or Direct Messages consist mainly of links shared without commentary;
 - if a large number of people have blocked you in response to high volumes of untargeted, unsolicited, or duplicative content or engagements from your account;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicative or substantially similar content, replies, or mentions over multiple accounts or multiple duplicate updates on one account, or create duplicate or substantially similar accounts;
 - if you post multiple updates to a trending or popular topic with an intent to subvert or manipulate the topic to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you send large numbers of unsolicited replies or mentions;
 - if you add users to lists in a bulk or aggressive manner;
 - if you are randomly or aggressively engaging with Tweets (e.g., likes, Retweets, etc.) or users (e.g., following, adding to lists or Moments, etc.) to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you repeatedly post other people's account information as your own (e.g., bio, Tweets, profile URL, etc.);
 - if you post misleading, deceptive, or malicious links (e.g., affiliate links, links to malware/clickjacking pages, etc.);
 - if you create fake accounts, account interactions, or impressions;
 - if you sell, purchase, or attempt to artificially inflate account interactions (such as followers, Retweets, likes, etc.); and
 - if you use or promote third-party services or apps that claim to get you more followers, Retweets, or likes (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers or engagements to your account or Tweets).

Please see our support articles on [following rules and best practices](#) and [automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts may be permanently suspended.

Content Visibility

Accounts under investigation or which have been detected as sharing content in violation of these Rules may have their account or Tweet visibility limited in various parts of Twitter, including search. To learn more about situations in which content may be restricted on Twitter, please see our support article on [search rules and restrictions](#).

Trademark policy

Updated June 30, 2016

What is a Trademark Policy Violation on Twitter?

Using another's trademark in a manner that may mislead or confuse others about your brand affiliation may be a violation of our trademark policy.

How Does Twitter Respond to Trademark Reports?

We review each report and do the following:

- If we determine that there is a violation of our trademark policy, we may suspend the account. In some instances, we may give the account holder an opportunity to comply with our policies before suspending the account. For more information, see our [Parody, commentary, and fan account policy](#).
- We may also provide the account holder with your name and information included in the copy of the report.
- If you would like to report a trademark violation in a [Twitter Ad](#), please visit our [Trademark Policy for Promoted Ads](#) to learn more.

Please note that Twitter only investigates requests that are submitted by the trademark holder or by an authorized representative of the trademark holder.

What is not a Trademark Policy Violation?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples include:

- Using a trademark in a way that is outside the scope of the trademark registration (e.g. territory, or goods and services identified in the registration).
- Nominative and other fair uses of trademarks are protected uses under our trademark policy, so long as the account is clearly distinguished from the trademark owner. This includes use by resellers in certain regions and accounts engaging in parody, commentary, or news. For more information, see our [Parody, commentary, and fan account policy](#).

How do I Report a Trademark Policy Violation?

Submit trademark reports through our [Trademark Issue support form](#). Using the appropriate form ensures we have all the information necessary to process your report quickly. Please provide all information requested in the form. If you submit a report with incomplete information, we will follow up with you to request that information. Please note that this will result in a delay in processing your report.

If you have trouble using the form, try updating your browser or switching to another browser.

Copyright policy

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged

infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notification.

When Should I File a Counter-notice?

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and

4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Twitter media policy

Twitter Rules: Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images.

Twitter Rules: You may not use hateful images or symbols in your profile image or profile header.

Rationale

Some forms of graphic violence, adult content, or hateful imagery may be permitted in Tweets when they are marked as [sensitive media](#). However, you may not include this type of content in live video, or in profile or header images.

While we want people to feel free to share media that reflects their creativity or individuality, or to show what's happening in the world, we will take action when it crosses the line into abuse towards a person, group, or protected category.

Finally, we may sometimes require you to remove media containing excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about how to make such a request [here](#).

When this applies

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. Some examples include, but are not limited to, depictions of:

- the moment at which someone dies

- gruesome crime or accident scenes
- bodily harm, torture, dismemberment, or mutilation

We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity (including close-ups of genitals, buttocks, or breasts)
 - Please note that exceptions may be made for artistic, medical, health, or educational content. Breastfeeding content does not need to be marked as sensitive.
- simulating a sexual act
- intercourse or any sexual act (may involve humans, humanoid animals, cartoons, or anime)

If you see media you think may be in violation of our policy, report it. You can read about how to report sensitive media [here](#).

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, or ethnicity/national origin.

Some examples of hateful imagery include, but are not limited to:

- Symbols historically associated with hate groups (for example, the Nazi swastika)
- Images depicting others as less than human or altered to include hateful symbols
- Altered image references to a mass murder that targeted a protected category

Do I need to be the target of this content in order for it to be a violation of the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When adult content, graphic violence, or hateful imagery appears in Tweets, we may place this content behind an interstitial advising viewers to be aware that they will see sensitive media if they click through. This allows us to identify potentially sensitive content that some people may not wish to see. Learn more about [how to control whether you see sensitive media](#).

However, if such content appears in live video, header, or profile images, it will violate our media policy. The consequences for violating this policy depend on the severity of the violation and the person's previous record of violations.

The first time someone violates this policy they will be required to remove the imagery and may be asked to verify their account ownership. Subsequent violations could result in permanent suspension. Learn more about [our range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

How to contact Twitter about media concerning a deceased family member

In order to respect the wishes of loved ones, Twitter will remove imagery of deceased individuals in certain circumstances. Immediate family members and other authorized individuals may request the removal of images or video of deceased individuals, from when critical injury occurs to the moments before or after death, by sending

a request to Twitter Inc. via our [privacy form](#). When reviewing such media removal requests, Twitter considers public interest factors such as the newsworthiness of the content and may not be able to honor every request. Note: Find information about contacting Twitter to [deactivate the account of a deceased or incapacitated user](#).

How to mark media as sensitive links to Media policy

Twitter trends FAQs

Where can I find trends?

Trends are available on the Twitter for iOS app, the Twitter for Android app, and [twitter.com](#).

On Twitter's mobile apps, you can find trends listed under the trends section of the **Explore** tab. When signed in to [twitter.com](#) on a desktop or laptop computer, trends are listed in many places, including the Home timeline, Notifications, search results, and profile pages.

How are trends determined?

Trends are determined by an algorithm and, by default, are tailored for you based on who you follow, your interests, and your location. This algorithm identifies topics that are popular now, rather than topics that have been popular for a while or on a daily basis, to help you discover the hottest emerging topics of discussion on Twitter.

You can choose to see trends that are not tailored for you by selecting a specific trends location on [twitter.com](#), iOS, or Android (instructions below). Location trends identify popular topics among people in a specific geographic location.

Note: The number of Tweets that are related to the trends is just one of the factors the algorithm looks at when ranking and determining trends. Algorithmically, trends and hashtags are grouped together if they are related to the same topic. For instance, #MondayMotivation and #MotivationMonday may both be represented by #MondayMotivation.

What does the # sign mean?

You may notice that some trends have # sign before the word or phrase. This is called a [hashtag](#) and is included specifically in Tweets to mark them as relating to a topic, so that people can follow the conversation in search.

Will I see world and local events in my trends?

Yes, when signed into your account on the [twitter.com](#), iOS, or Android, trends will be tailored for you based on your location and who you follow. There will be many world and local news events and conversations that will appear in your trends regardless of your personalization.

Can I see trends for a specific location?

Yes, you can view trends for a specific location on [twitter.com](#), iOS, or Android.

View trends via [twitter.com](#):

To view trends for a specific location, click **Change** in your trends box. You can choose to **Keep tailored trends** based on your location and who you follow, or you can select **Change** to choose a nearby location or search locations.

Adjust your trends settings via Twitter for iOS:

1. In the top menu, tap your **profile** icon.
2. Tap **Settings and privacy**, then tap **Content preferences**.
3. Under **Explore**, tap **Trends**.
4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.
5. When the feature is disabled, you can change your location by tapping **Change**

location. In the search box, type your desired trends location or choose from the list.

Adjust your trends settings via Twitter for Android:

1. In the top menu, you will either see a **navigation menu** icon or your **profile** icon.
Tap whichever icon you have, and select **Settings and privacy**.

2. Tap **Content preferences**.
3. Under **Explore**, tap **Trends**.
4. Drag the slider next to **Trends for you** to turn on and receive personalized trends.
5. When the feature is disabled, you can change your location by tapping **Change location**. In the search box, type your desired trends location or choose from the list.

Note: If you don't find your preferred country or city, it means we're not receiving enough Tweets from that geographical area to create a list. You can look up local Tweets on any topic by using [advanced Twitter search operators](#).

What happens when I click or tap on a trend?

Clicking or tapping any of the trends takes you to the Twitter search results for that trend. You'll see all Tweets including that phrase or hashtag. To see what people are saying about a previous trend, perform a [search for that keyword](#).

How can I participate in a trend?

Simply post a Tweet including the exact word or phrase as it appears in the trends list (with the hashtag, if you see one). Due to the large number of people Tweeting about these specific trends, you may not always be able to find your particular Tweet in search, but your followers will always see your Tweets.

Note: Twitter also filters searches for quality. To understand what types of behavior could cause your Tweets to be filtered out, read about our [search rules and restrictions](#).

Are there rules for trends?

Yes. We want trends to promote healthy discussions on Twitter. This means that at times, we may prevent certain content from trending. These include trends that:

Contain profanity or adult/graphic references.

Incite hate on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease.

Violate the [Twitter Rules](#).

In some cases, we may also consider the newsworthiness of the content, or if it is in the public interest when evaluating potential violations. In these cases, the content might continue to trend on our platform.

Note that even if we prevent the hashtag or the content from appearing on the list of trending topics, you may still be able to access conversations around that content on Twitter.

Do trends have descriptions?

Below some trends, you may see the approximate number of Tweets associated with the trend, or personalized context such as who in your network is Tweeting about it.

Additionally, when a Moment is found to be relevant to a trend, we may use the title of the Moment as the trend's description.

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Child sexual exploitation policy

Twitter removes content promoting child sexual exploitation

We do not tolerate child sexual exploitation on Twitter. When we are made aware of links to images of or content promoting child sexual exploitation they will be removed from the site without further notice

and reported to The National Center for Missing & Exploited Children ("NCMEC"); we permanently suspend accounts promoting or containing updates with links to child sexual exploitation.

Reporting content promoting child sexual exploitation

If you think you've found a Twitter account distributing or promoting child sexual exploitation while using Twitter, please click [here](#) to report it. Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. To find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with the National Center for Missing and Exploited Children [here](#), or your local law enforcement. If you believe that the content was hosted on a website outside of the United States, you can report it [here](#) on the International Association of Internet Hotlines website.

NOTE: Please do NOT tweet, retweet or repost child sexual exploitation for any reason. Click [here](#) to report it to us immediately and we will take steps to remove it.

Law Enforcement Agencies

If you are a member of a law enforcement agency seeking information regarding Twitter accounts promoting or distributing child sexual exploitation, please see our [Guidelines for Law Enforcement](#) for information on our process; we also suggest you contact the National Center for Missing and Exploited Children (NCMEC) [here](#), regardless of jurisdiction, to determine whether or not information on those accounts is already on file.

For general law enforcement inquiries, please see our [Guidelines for Law Enforcement](#) for more information.

Hateful conduct policy

Freedom of expression means little if voices are silenced because people are afraid to speak up. We do not tolerate behavior that harasses, intimidates, or uses fear to silence another person's voice. If you see something on Twitter that violates these rules, please report it to us.

How our policy works

As explained in the Twitter Rules,

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease.

We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Examples of what we do not tolerate includes, but is not limited to behavior that harasses individuals or groups of people with:

violent threats;

wishes for the physical harm, death, or disease of individuals or groups;

references to mass murder, violent events, or specific means of violence in which/with which such groups have been the primary targets or victims;

behavior that incites fear about a protected group;

repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone.

How our enforcement works Context matters.

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. While we accept reports of violations from anyone, sometimes we also need to hear directly from the target to ensure that we have proper context.

The number of reports we receive does not impact whether or not something will be removed. However, it may help us prioritize the order in which it gets reviewed.

We focus on behavior.

We enforce policies when someone reports behavior that is abusive and targets an entire protected group and/or individuals who may be members.

This targeting can happen in any manner (for example, @mentions, tagging a photo, and more).

We have a range of enforcement options.

The consequences for violating our rules vary depending on the severity of the violation and the person's previous record of violations. For example, we may ask someone to remove the offending Tweet before they can Tweet again. For other cases, we may suspend an account.

Report private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Read about our [policy regarding private information posted on Twitter](#).

Do I need to have a Twitter account to make a report?

No, you do not need to have a Twitter account to file a report that your private and confidential information has been posted on Twitter.

Can I report unauthorized posting of Copyrighted material?

If your copyrighted material has been Tweeted without your permission, please review our [Copyright and DMCA policy](#).

What can I do if my information is posted on other websites?

We cannot take down any information that is posted on a website other than Twitter. If your information is posted on another website, please contact the help center or support staff for the website where your information is posted.

How do I report that my private information is posted on Twitter?

You can report that your private information has been posted on Twitter directly from the offending Tweet(s). Learn [how to report individual Tweets](#).

File a [report that someone has posted your private information](#).

What happens when Twitter receives a valid report?

When we receive a complete and valid report that private information has been posted on Twitter, we'll investigate the account and Tweets reported. We will review where, if anywhere, the information has been made publicly available before taking action on the account or Tweets. If the information you reported was previously posted elsewhere on the Internet, it is not a violation of our policy and we will not take action.

Learn more about [protecting your private information on Twitter and other websites](#).

Read our [policy regarding private information posted on Twitter](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter impersonation policy.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to Twitter's [parody, commentary, and fan account policy](#) for more information about these types of accounts.

Find more information about [reporting impersonation on Twitter](#).

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation rules

Updated November 3, 2017

This page is primarily intended for **developers**.

For Twitter users: You are ultimately responsible for the actions taken with your account, or by applications associated with your account. Before authorizing a third-party application to access or use your account, make sure you've thoroughly investigated the application and understand what it will do. If automated activity on your account violates the [Twitter Rules](#) or these automation rules, Twitter may take action on your account, including [filtering your Tweets from search results](#) or suspending your account.

For more information on third-party applications, please see our article on [connecting and revoking third-party applications](#).

If you are a developer and have questions about these Automation Rules, please reach out via <https://twittercommunity.com/c/rules-and-policies>.

If your application has been suspended or restricted from performing write actions, you may submit an appeal at </forms/platform>.

I. Ground Rules Do!

Build solutions that automatically broadcast helpful information in Tweets.

Run creative campaigns that auto-reply to users who engage with your content.

Build solutions that automatically respond to users in Direct Messages.

Try new things that help people (and comply with our rules).

Make sure your application provides a good user experience and performs well — and confirm that remains the case over time.

Don't!

Violate these or other policies. Be extra mindful of our rules about abuse and user privacy.

Abuse the Twitter API or attempt to circumvent rate limits.

Use non-API-based forms of automation, such as scripting the Twitter website. The use of these techniques may result in the permanent suspension of your account.

Spam or bother users, or otherwise send them unsolicited messages.

A. The Twitter Rules and the Developer Agreement and Policy

As with all activity on Twitter, automated activity is subject to the [Twitter Rules](#) and, if you're a developer using the Twitter API, the [Developer Agreement and Policy](#).

You should carefully review these policies to ensure that your automated activity is compliant. Automated applications or activities that violate these policies, or that facilitate or induce users to violate them, may be subject to enforcement action, potentially including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these policies.

Although all aspects of the Twitter Rules and the Developer Agreement and Policy apply to automated activity, you should keep the following rules top of mind:

Spamming: You may not send automated Tweets or Direct Messages that are spam, or otherwise engage in spamming activity. Some examples of spammy behavior to avoid with automation include:

Trending topics: You may not automatically post about trending topics on Twitter, or use automation to attempt to influence or manipulate trending topics.

Multiple posts/accounts: You may not post duplicative or substantially similar Tweets on one account or over multiple accounts you operate.

Duplicate accounts: You may not create and/or automate multiple accounts for duplicative or substantially similar use cases.

However, automating multiple accounts for related but non-duplicative use cases is permitted. For example, you may automate separate accounts to Tweet when the Hubble Space Telescope passes over different cities, such as [San Francisco](#) or [Hong Kong](#).

Misleading links: You may not send automated Tweets or Direct Messages containing links that are misleading, including links that maliciously or deceptively redirect through landing pages or ad pages before displaying the final content.

Sensitive media: Automated Tweets and Direct Messages must comply with the [Twitter media policy](#), and you should mark your account as potentially sensitive if you intend to post graphic, pornographic, or potentially sensitive media.

Abusive behavior: You may not engage in any automated activity that encourages, promotes, or incites abuse, violence, hateful conduct, or harassment, on or off Twitter.

Private information: You may not post private or confidential information about a person without their prior express authorization.

B. Other Ground Rules for Automated Activity

In addition to the policies above, the following ground rules apply to all automated activity on Twitter:

Don't surprise or mislead users: Automated activity should honor users' expectations. Ask for the user's permission before taking an action if you aren't sure.

Mature content or profanity: Don't Direct Message, mention, or reply to users with potentially sensitive content (including profanity), unless they've clearly indicated an intent to receive it in advance.

Be thoughtful about the information you request or exchange on Twitter

Tweets: Don't ask users to send you personal or private information via a public Tweet. If you need additional personal or private information from a user to provide them with customer service (or other similar use cases), you should ask the user to share such information by Direct Message or another private channel. You might even consider adding a [Direct Message deep link](#) to your Tweet.

Direct Messages: You should only ask users for the minimum amount of information you need to provide them with service. If you need to request or exchange particularly sensitive information (such as credit card information), you should consider directing users to your website or other appropriate channel to do so.

II. Activity-Specific Rules

The activity-specific rules in this section apply to taking specific automated actions on Twitter. Please read these rules carefully, as they outline both permitted and prohibited use cases of automation.

Automated applications or activities that violate these rules, or that facilitate or induce users to violate them, may be subject to enforcement action, including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access the Twitter API based on violations of these rules. As a reminder, you should also carefully review the spam guidelines in the [Twitter Rules](#) to avoid having activities performed by you, your app, or other users through your app or service flagged as spam.

A. Automated Actions Through Another User's Account

Twitter users may authorize your app or service to [access their Twitter account through OAuth](#). A user authorizing your app or service to access their Twitter account through OAuth does not by itself constitute sufficient consent to take automated actions through that user's account.

You may only take automated actions through another Twitter user's account if you:

clearly describe to the user the types of automated actions that will occur; receive express consent from the user to take those automated actions; and immediately honor a user's request to opt-out of further automated actions.

If you substantially change the purpose or functionality of your app or service, you must re-obtain express consent from the user to take automated action through their account before doing so.

These requirements apply to any automated action taken through another Twitter user's account, including posting Tweets, sending Direct Messages, deleting Tweets or Direct Messages, or following/unfollowing other accounts. For applications that offer users the ability to delete Tweets in a bulk or automated manner, you must also clearly state that Tweets are not recoverable once deleted.

B. Automated Tweets

1. Posting automated Tweets

Automated Tweets that cross-post outside information: You may post automated Tweets based on sources of outside information — such as an RSS feed, weather data, etc. — as long as you are sufficiently authorized to publish such information.

Other automated Tweets (excluding mentions or replies): Provided you comply with all other rules, you may post automated Tweets for entertainment, informational, or novelty purposes. As a reminder, accounts posting duplicative, spammy, or otherwise prohibited content may be subject to suspension.

2. Posting automated mentions and replies

The reply and mention functions are intended to make communication between Twitter users easier. Automating these actions to reach many users on an unsolicited basis is an abuse of the feature, and is not permitted. For example, sending automated replies to Tweets based on keyword searches alone is not permitted. Spamy or duplicative use of mentions and replies may result in enforcement action, such as the removal of

your Tweets from Search or the suspension of your app or account.

However, you may send automated replies or mentions to Twitter users so long as:

in advance of sending the automated reply, the recipient or mentioned user(s) have requested or have clearly indicated an intent on Twitter to be contacted by you (i.e. opted in), for example by replying to a Tweet from your account, or by sending you a Direct Message;

you provide a clear and easy way for such users to opt-out of receiving automated replies and mentions, and promptly honor all such opt-out requests;

you only send one automated reply or mention per user interaction; and

the automated reply or mention is a reply to the user's original Tweet (if your campaign is based on users posting a reply to your Tweet).

Opt-in techniques and indications of user intent take many different forms, depending on the specifics of your use case and implementation. Some examples include:

A Tweet from your account that clearly indicates that a user taking a specific action on that Tweet (such as Retweeting it) will opt the user into receiving an automated response.

A mention of your account by the user in a manner suggesting the user clearly wishes or intends to receive a response. If you want to run an auto-reply campaign with a campaign- or use-case-specific hashtag, users should also mention you in their Tweets.

Note that a user following your account is not on its own a sufficient indication of user intent to receive an automated response.

: Advertisers, publishers, and brands using auto-response campaigns must request approval from Twitter and may be subject to additional rules. Please reach out to your account or partner manager for assistance. If you do not have a partner manager but are interested in managed client services, you can learn more [here](#).

Additionally, we recommend that any accounts that will communicate with

users via automated mentions or replies:

Note

1. Appropriately filter responses based on potentially sensitive language in user handles,

display names, and Tweet text, as well as potentially sensitive media;

2. Check that the Tweet you are mentioning or replying to still exists (for example, using the statuses/lookup endpoint on the Twitter API).

C. Automated Direct Messages

Sending automated Direct Messages to users

You may not send unsolicited Direct Messages in a bulk or automated manner, and should be thoughtful about the frequency with which you contact users via Direct Message.

You may send automated Direct Messages to users so long as:

in advance of sending the Direct Message, the recipient(s) have requested or have clearly indicated an intent on Twitter to be contacted by you via Direct Message, for example by sending you a Direct Message; **and**

you provide a clear and easy way for such users to opt-out of receiving automated Direct Messages, and promptly honor all such opt-out requests.

The fact that a user is technically able to receive a Direct Message from you (e.g. because the user follows you, has enabled the ability to receive Direct Messages from any account, or because the user is in a pre-existing Direct Message conversation with you) does not necessarily mean they have requested or expect to receive automated Direct Messages from you.

Interacting with users via Direct Message

Per the Ground Rules, remember to be thoughtful about the amount and type of information you request or exchange with users via Direct Messages. If you will be asking a user to provide personal or private information via an automated Direct Message, you must clearly explain how you will use the information you're collecting. Consider including a link to your

privacy policy in your Direct Message to the user, as well as in your Twitter profile bio.

Don't publicly share information received in a Direct Message conversation with a user without first obtaining explicit consent from the user. For example, if a user asks you via Direct Message about a purchase they made from you, you may not mention the user in a public Tweet that includes information about their purchase unless you have the user's explicit consent to do so.

After a user-initiated interaction ends, don't send additional follow-up Direct Messages or mention users in a Tweet unless you get permission from the user.

D. Automated actions you take on Tweets or accounts

Automated likes: You may not like Tweets in an automated manner.

Automated Retweets: Provided you comply with all other rules, you may Retweet or Quote Tweet in an automated manner for entertainment, informational, or novelty purposes. Automated Retweets often lead to negative user experiences, and bulk, aggressive, or spammy Retweeting is a violation of the [Twitter Rules](#).

Automated following/unfollowing: You may not follow or unfollow Twitter accounts in a bulk, aggressive, or indiscriminate manner. Aggressive following is a violation of the [Twitter Rules](#). Please also review our [following rules and best practices](#) to ensure you are in compliance. Note that applications that claim to get users more followers are also prohibited under the [Twitter Rules](#).

Automated adding to lists or collections: You may not add Twitter users to lists or add Tweets to collections in a bulk or indiscriminate manner. Adding a large number of unrelated users to lists is a violation of the [Twitter Rules](#).

About search rules and restrictions

This article details the rules for using Twitter search, and explains reasons why content may not show up in Twitter search.

Why doesn't all content show up in Twitter search?

In order to keep your search results relevant, Twitter filters search results for quality Tweets and accounts. Material that jeopardizes search quality or creates a bad search experience for other people may be automatically removed from Twitter search. Read more about this [here](#).

General rules for using Twitter search

The [Twitter Rules](#) explain what behaviors are permitted on Twitter. In addition to these rules, we've included some tips below to keep your content relevant (and your Tweets in Twitter search).

Please do not:

- Repeatedly post duplicate or near-duplicate content (links or Tweets).
- Abuse trending topics or hashtags (topic words with a # sign).
- Send automated Tweets or replies.
- Use bots or applications to post similar messages based on keywords.
- Post similar messages over multiple accounts.
- Aggressively follow and unfollow people.

Twitter may automatically remove accounts engaging in these behaviors from search (or even suspend in some cases) in order to ensure the best experience for everyone.

Search results filtering

To learn how to filter your search results, including how to disable or enable safe search, read our [Using Twitter search](#) article.

Search results filtering in third-party applications

Accounts may be filtered from search for posting similar messages over several accounts. Your account may be filtered from search if you are logged into a third-party application that is updating many accounts with similar or duplicate content.

Rules for participating in contests on Twitter

Businesses sometimes host contests and offer prizes to Retweet an update, following a particular account, or posting Tweets with a specific hashtag topic or reply. If you're running a contest, check out this [article](#) for guidelines for contests on Twitter.

If you repeatedly post near-duplicate updates or duplicate links in order to enter contests, you may be filtered from search. Please keep in mind that you are responsible for the updates that you post to your Twitter account.

What to do if you are missing from search

Please see this [article](#) for more information related to this issue.

If you are being filtered from search, rest assured that your followers will still receive your updates. If you post a reply or mention to a specific person, it will be delivered to that person.

Having trouble?

You can learn more about searching by keyword, location, and trending topics in our [Using Twitter search](#) help article.

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. These limitations are set forth in the Twitter Rules below.

The Twitter Rules (along with all incorporated policies), [Privacy Policy](#), and [Terms of Service](#) collectively make up the "Twitter User Agreement" that governs a user's access to and use of Twitter's services.

All individuals accessing or using Twitter's services must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in Twitter taking one or more of the following enforcement actions:

- requiring you to delete prohibited content before you can again create new posts and interact with other Twitter users;
- temporarily limiting your ability to create posts or interact with other Twitter users;
- asking you to verify account ownership with a phone number or email address; or
- permanently suspending your account(s).

If you attempt to evade a permanent suspension by creating new accounts, we will suspend your new accounts.

Please note that we may need to change these Rules from time to time and reserve the right to do so. The most current version is always available at: <https://twitter.com/rules>.

The policies set forth in these Twitter Rules govern organic content on our platform. To learn more about the rules which govern ads and promoted content, please review our [Ads policies](#).

Content Boundaries and Use of Twitter

Intellectual property

Trademark: We reserve the right to suspend accounts or take other appropriate action when someone's brand or trademark, including business name and/or logo, is used in a manner that may mislead or confuse others about your brand affiliation. Read more about our [trademark policy](#) and [how to report a violation](#).

Copyright: We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in our [Terms of Service](#). Read more about our [copyright policy](#).

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Learn more about our [media policy](#).

Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images. Additionally, Twitter may sometimes require you to remove excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about [how to make such a request](#), and [how to mark your media as sensitive](#).

Unlawful use

You may not use our service for any unlawful purposes or in furtherance of illegal activities. By using Twitter, you agree to comply with all applicable laws governing your online conduct and content.

Trends

At times, we may prevent certain content from trending. This includes content that violates the Twitter Rules, as well as content that may attempt to manipulate trends. Read more about [what we allow and do not allow to trend](#).

Misuse of Twitter badges

You may not use badges, including but not limited to the “promoted” or “verified” Twitter badges, unless provided by Twitter. Accounts using unauthorized badges as part of their profile photos, header photos, display names, or in any way that falsely implies affiliation with Twitter or authorization from Twitter to display these badges, may be suspended.

Misuse of usernames

Selling usernames: You may not buy or sell Twitter usernames.

Username squatting: You may not engage in username squatting. Some of the factors we take into consideration when determining whether conduct is username squatting include:

- the number of accounts created;
- the creation of accounts for the purpose of preventing others from using those account names;
- the creation of accounts for the purpose of selling those accounts; and
- the use of third-party content feeds to update and maintain accounts under the names of those third parties.

Please note that Twitter may also remove accounts that are inactive for more than six months. Learn more about [username squatting](#).

Abusive Behavior

We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we prohibit behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Context matters when evaluating for abusive behavior and determining appropriate enforcement actions. Factors we may take into consideration include, but are not limited to whether:

- the behavior is targeted at an individual or group of people;
- the report has been filed by the target of the abuse or a bystander;
- the behavior is newsworthy and in the legitimate public interest.

Violence and physical harm

Violence: You may not make specific threats of violence or wish for the serious physical harm, death, or disease of an individual or group of people. This includes, but is not limited to, threatening or promoting terrorism. You also may not affiliate with organizations that — whether by their own statements or activity both on and off the platform — use or promote violence against civilians to further their causes. We will begin enforcing this rule around affiliation with such organizations on December 18, 2017.

Suicide or self-harm: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Child sexual exploitation: You may not promote child sexual exploitation. Learn more about our zero-tolerance [child sexual exploitation policy](#).

Abuse and hateful conduct

Abuse: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Unwanted sexual advances: You may not direct abuse at someone by sending unwanted sexual content, objectifying them in a sexually explicit manner, or otherwise engaging in sexual misconduct.

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. Read more about our [hateful conduct policy](#).

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category. We will begin enforcing this rule on December 18, 2017.

Private information and intimate media

Private information: You may not publish or post other people's private information without their express authorization and permission. Definitions of private information may vary depending on local laws. Read more about our [private information policy](#).

Intimate media: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. Read more about intimate media on Twitter.

Threats to expose / hack: You may not threaten to expose someone's private information or intimate media. You also may not threaten to hack or break into someone's digital information.

Impersonation

You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. While you may maintain parody, fan, commentary, or newsfeed accounts, you may not do so if the intent of the account is to engage in spamming or abusive behavior. Read more about our [impersonation policy](#).

Spam and Security

We strive to protect people on Twitter from technical abuse and spam.

To promote a stable and secure environment on Twitter, you may not do, or attempt to do, any of the following while accessing or using Twitter:

- Access, tamper with, or use non-public areas of Twitter, Twitter's computer systems, or the technical delivery systems of Twitter's providers (except as expressly permitted by the Twitter Bug Bounty program).
- Probe, scan, or test the vulnerability of any system or network, or breach or circumvent any security or authentication measures (except as expressly permitted by the Twitter Bug Bounty program).
- Access or search, or attempt to access or search, Twitter by any means (automated or otherwise) other than through our currently available, published interfaces that are provided by Twitter (and only pursuant to the applicable terms and conditions), unless you have been specifically allowed to do so in a separate agreement with Twitter. Note that crawling Twitter is permissible if done in accordance with the provisions of the robots.txt file; however, scraping Twitter without our prior consent is expressly prohibited.
- Forge any TCP/IP packet header or any part of the header information in any email or posting, or in any way use Twitter to send altered, deceptive, or false source-identifying information.
- Interfere with or disrupt the access of any user, host or network, including, without limitation, sending a virus, overloading, flooding, spamming, mail-bombing Twitter's services, or by scripting the creation of content in such a manner as to interfere with or create an undue burden on Twitter.

Any accounts engaging in the following activities may be temporarily locked or subject to permanent suspension:

- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use Twitter's services for the purpose of spamming anyone. Spam is generally defined on Twitter as bulk or aggressive activity that attempts to manipulate or disrupt Twitter or the experience of users on Twitter to drive traffic or attention to unrelated accounts, products, services, or initiatives. Some of the factors that we take into account when determining what conduct is considered to be spamming include:
 - if you have followed and/or unfollowed a large number of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if your Tweets or Direct Messages consist mainly of links shared without commentary;
 - if a large number of people have blocked you in response to high volumes of untargeted, unsolicited, or duplicative content or engagements from your account;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicative or substantially similar content, replies, or mentions over multiple accounts or multiple duplicate updates on one account, or create duplicate or substantially similar accounts;
 - if you post multiple updates to a trending or popular topic with an intent to subvert or manipulate the topic to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you send large numbers of unsolicited replies or mentions;
 - if you add users to lists in a bulk or aggressive manner;
 - if you are randomly or aggressively engaging with Tweets (e.g., likes, Retweets, etc.) or users (e.g., following, adding to lists or Moments, etc.) to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you repeatedly post other people's account information as your own (e.g., bio, Tweets, profile URL, etc.);
 - if you post misleading, deceptive, or malicious links (e.g., affiliate links, links to malware/clickjacking pages, etc.);
 - if you create fake accounts, account interactions, or impressions;
 - if you sell, purchase, or attempt to artificially inflate account interactions (such as followers, Retweets, likes, etc.); and
 - if you use or promote third-party services or apps that claim to get you more followers, Retweets, or likes (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers or engagements to your account or Tweets).

Please see our support articles on [following rules and best practices](#) and [automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts may be permanently suspended.

Content Visibility

Accounts under investigation or which have been detected as sharing content in violation of these Rules may have their account or Tweet visibility limited in various parts of Twitter, including search. To learn more about situations in which content may be restricted on Twitter, please see our support article on [search rules and restrictions](#).

Trademark policy

Updated June 30, 2016

What is a Trademark Policy Violation on Twitter?

Using another's trademark in a manner that may mislead or confuse others about your brand affiliation may be a violation of our trademark policy.

How Does Twitter Respond to Trademark Reports?

We review each report and do the following:

- If we determine that there is a violation of our trademark policy, we may suspend the account. In some instances, we may give the account holder an opportunity to comply with our policies before suspending the account. For more information, see our [Parody, commentary, and fan account policy](#).
- We may also provide the account holder with your name and information included in the copy of the report.
- If you would like to report a trademark violation in a [Twitter Ad](#), please visit our [Trademark Policy for Promoted Ads](#) to learn more.

Please note that Twitter only investigates requests that are submitted by the trademark holder or by an authorized representative of the trademark holder.

What is not a Trademark Policy Violation?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples include:

- Using a trademark in a way that is outside the scope of the trademark registration (e.g. territory, or goods and services identified in the registration).
- Nominative and other fair uses of trademarks are protected uses under our trademark policy, so long as the account is clearly distinguished from the trademark owner. This includes use by resellers in certain regions and accounts engaging in parody, commentary, or news. For more information, see our [Parody, commentary, and fan account policy](#).

How do I Report a Trademark Policy Violation?

Submit trademark reports through our [Trademark Issue support form](#). Using the appropriate form ensures we have all the information necessary to process your report quickly. Please provide all information requested in the form. If you submit a report with incomplete information, we will follow up with you to request that information. Please note that this will result in a delay in processing your report.

If you have trouble using the form, try updating your browser or switching to another browser.

Copyright policy

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged

infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notification.

When Should I File a Counter-notice?

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and

4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Twitter media policy

Twitter Rules: Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images.

Twitter Rules: You may not use hateful images or symbols in your profile image or profile header.

Rationale

Some forms of graphic violence, adult content, or hateful imagery may be permitted in Tweets when they are marked as [sensitive media](#). However, you may not include this type of content in live video, or in profile or header images.

While we want people to feel free to share media that reflects their creativity or individuality, or to show what's happening in the world, we will take action when it crosses the line into abuse towards a person, group, or protected category.

Finally, we may sometimes require you to remove media containing excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about how to make such a request [here](#).

When this applies

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. Some examples include, but are not limited to, depictions of:

- the moment at which someone dies

- gruesome crime or accident scenes
- bodily harm, torture, dismemberment, or mutilation

We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity (including close-ups of genitals, buttocks, or breasts)
 - Please note that exceptions may be made for artistic, medical, health, or educational content. Breastfeeding content does not need to be marked as sensitive.
- simulating a sexual act
- intercourse or any sexual act (may involve humans, humanoid animals, cartoons, or anime)

If you see media you think may be in violation of our policy, report it. You can read about how to report sensitive media [here](#).

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, or ethnicity/national origin.

Some examples of hateful imagery include, but are not limited to:

- Symbols historically associated with hate groups (for example, the Nazi swastika)
- Images depicting others as less than human or altered to include hateful symbols
- Altered image references to a mass murder that targeted a protected category

Do I need to be the target of this content in order for it to be a violation of the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When adult content, graphic violence, or hateful imagery appears in Tweets, we may place this content behind an interstitial advising viewers to be aware that they will see sensitive media if they click through. This allows us to identify potentially sensitive content that some people may not wish to see. Learn more about [how to control whether you see sensitive media](#).

However, if such content appears in live video, header, or profile images, it will violate our media policy. The consequences for violating this policy depend on the severity of the violation and the person's previous record of violations.

The first time someone violates this policy they will be required to remove the imagery and may be asked to verify their account ownership. Subsequent violations could result in permanent suspension. Learn more about [our range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

How to contact Twitter about media concerning a deceased family member

In order to respect the wishes of loved ones, Twitter will remove imagery of deceased individuals in certain circumstances. Immediate family members and other authorized individuals may request the removal of images or video of deceased individuals, from when critical injury occurs to the moments before or after death, by sending

a request to Twitter Inc. via our [privacy form](#). When reviewing such media removal requests, Twitter considers public interest factors such as the newsworthiness of the content and may not be able to honor every request. Note: Find information about contacting Twitter to [deactivate the account of a deceased or incapacitated user](#).

How to mark media as sensitive links to Media policy

Twitter trends FAQs

Where can I find trends?

Trends are available on the Twitter for iOS app, the Twitter for Android app, and [twitter.com](#).

On Twitter's mobile apps, you can find trends listed under the trends section of the **Explore** tab. When signed in to [twitter.com](#) on a desktop or laptop computer, trends are listed in many places, including the Home timeline, Notifications, search results, and profile pages.

How are trends determined?

Trends are determined by an algorithm and, by default, are tailored for you based on who you follow, your interests, and your location. This algorithm identifies topics that are popular now, rather than topics that have been popular for a while or on a daily basis, to help you discover the hottest emerging topics of discussion on Twitter.

You can choose to see trends that are not tailored for you by selecting a specific trends location on [twitter.com](#), iOS, or Android (instructions below). Location trends identify popular topics among people in a specific geographic location.

Note: The number of Tweets that are related to the trends is just one of the factors the algorithm looks at when ranking and determining trends. Algorithmically, trends and hashtags are grouped together if they are related to the same topic. For instance, #MondayMotivation and #MotivationMonday may both be represented by #MondayMotivation.

What does the # sign mean?

You may notice that some trends have # sign before the word or phrase. This is called a **hashtag** and is included specifically in Tweets to mark them as relating to a topic, so that people can follow the conversation in search.

Will I see world and local events in my trends?

Yes, when signed into your account on the twitter.com, iOS, or Android, trends will be tailored for you based on your location and who you follow. There will be many world and local news events and conversations that will appear in your trends regardless of your personalization.

Can I see trends for a specific location?

Yes, you can view trends for a specific location on twitter.com, iOS, or Android.

View trends via twitter.com:

To view trends for a specific location, click **Change** in your trends box. You can choose to **Keep tailored trends** based on your location and who you follow, or you can select **Change** to choose a nearby location or search locations.

Adjust your trends settings via Twitter for iOS: 0.

In the top menu, tap your profile icon.

17. Tap **Settings and privacy**, then tap **Content preferences**.
18. Under **Explore**, tap **Trends**.

T. Drag the slider next to **Trends for you** to turn on and receive personalized trends.

V. When the feature is disabled, you can change your location by tapping **Change location**. In the search box, type your desired trends location or choose from the list.

Adjust your trends settings via Twitter for Android:

0. In the top menu, you will either see a navigation menu icon or your profile

icon. Tap whichever icon you have, and select **Settings and privacy**.

17. Tap **Content preferences**.
18. Under **Explore**, tap **Trends**.

T. Drag the slider next to **Trends for you** to turn on and receive personalized trends.

V. When the feature is disabled, you can change your location by tapping **Change location**. In the search box, type your desired trends location or choose from the list.

Note: If you don't find your preferred country or city, it means we're not receiving enough Tweets from that geographical area to create a list. You can look up local Tweets on any topic by using [advanced Twitter search operators](#).

What happens when I click or tap on a trend?

Clicking or tapping any of the trends takes you to the Twitter search results for that trend. You'll see all Tweets including that phrase or hashtag. To see what people are saying about a previous trend, perform a [search for that keyword](#).

How can I participate in a trend?

Simply post a Tweet including the exact word or phrase as it appears in the trends list (with the hashtag, if you see one). Due to the large number of people Tweeting about these specific trends, you may not always be able to find your particular Tweet in search, but your followers will always see your Tweets.

Note: Twitter also filters searches for quality. To understand what types of behavior could cause your Tweets to be filtered out, read about our [search rules and restrictions](#).

Are there rules for trends?

Yes. We want trends to promote healthy discussions on Twitter. This means that at times, we may prevent certain content from trending. These include trends that:

Contain profanity or adult/graphic references.

Incite hate on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease.

[Violate the Twitter Rules](#).

In some cases, we may also consider the newsworthiness of the content, or if it is in the public interest when evaluating potential violations. In these cases, the content might continue to trend on our platform.

Note that even if we prevent the hashtag or the content from appearing on the list of trending topics, you may still be able to access conversations around that content on Twitter.

Do trends have descriptions?

Below some trends, you may see the approximate number of Tweets associated with the trend, or personalized context such as who in your network is Tweeting about it.

Additionally, when a Moment is found to be relevant to a trend, we may use the title of the Moment as the trend's description.

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Child sexual exploitation policy

Twitter removes content promoting child sexual exploitation

We do not tolerate child sexual exploitation on Twitter. When we are made aware of links to images of or content promoting child sexual exploitation they will be removed from the site without further notice and reported to The National Center for Missing & Exploited Children ("NCMEC"); we permanently suspend accounts promoting or containing updates with links to child sexual exploitation.

Reporting content promoting child sexual exploitation

If you think you've found a Twitter account distributing or promoting child sexual exploitation while using Twitter, please click [here](#) to report it. Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. To find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with the National Center for Missing and Exploited Children [here](#), or your local law enforcement. If you believe that the content was hosted on a website outside of the United States, you can report it [here](#) on the International Association of Internet Hotlines website.

NOTE: Please do NOT tweet, retweet or repost child sexual exploitation for any reason. Click [here](#) to report it to us immediately and we will take steps to remove it.

Law Enforcement Agencies

If you are a member of a law enforcement agency seeking information regarding Twitter accounts promoting or distributing child sexual exploitation, please see our [Guidelines for Law Enforcement](#) for information on our process; we also suggest you contact the National Center for Missing and Exploited Children (NCMEC) [here](#), regardless of jurisdiction, to determine whether or not information on those accounts is already on file.

For general law enforcement inquiries, please see our [Guidelines for Law Enforcement](#) for more information.

Hateful conduct policy

Freedom of expression means little if voices are silenced because people are afraid to speak up. We do not tolerate behavior that harasses, intimidates, or uses fear to silence another person's voice. If you see something on Twitter that violates these rules, please report it to us.

How our policy works

As explained in the Twitter Rules,

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Examples of what we do not tolerate includes, but is not limited to behavior that harasses individuals or groups of people with:

violent threats;

wishes for the physical harm, death, or disease of individuals or groups;

references to mass murder, violent events, or specific means of violence in which/with which such groups have been the primary targets or victims;

behavior that incites fear about a protected group;

repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone.

How our enforcement works Context matters.

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. While we accept reports of violations from anyone, sometimes we also need to hear directly from the target to ensure that we have proper context.

The number of reports we receive does not impact whether or not something will be removed. However, it may help us prioritize the order in which it gets reviewed.

We focus on behavior.

We enforce policies when someone reports behavior that is abusive and targets an entire protected group and/or individuals who may be members.

This targeting can happen in any manner (for example, @mentions, tagging a photo, and more).

We have a range of enforcement options.

The consequences for violating our rules vary depending on the severity of the violation and the person's previous record of violations. For example, we may ask someone to remove the offending Tweet before they can Tweet again. For other cases, we may suspend an account.

Report private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Read about our [policy regarding private information posted on Twitter](#).

Do I need to have a Twitter account to make a report?

No, you do not need to have a Twitter account to file a report that your private and confidential information has been posted on Twitter.

Can I report unauthorized posting of Copyrighted material?

If your copyrighted material has been Tweeted without your permission, please review our [Copyright and DMCA policy](#).

What can I do if my information is posted on other websites?

We cannot take down any information that is posted on a website other than Twitter. If your information is posted on another website, please contact the help center or support staff for the website where your information is posted.

How do I report that my private information is posted on Twitter?

You can report that your private information has been posted on Twitter directly from the offending Tweet(s). Learn [how to report individual Tweets](#).

File a [report that someone has posted your private information](#).

What happens when Twitter receives a valid report?

When we receive a complete and valid report that private information has been posted on Twitter, we'll investigate the account and Tweets reported. We will review where, if anywhere, the information has been made publicly available before taking action on the account or Tweets. If the information you reported was previously posted elsewhere on the Internet, it is not a violation of our policy and we will not take action.

Learn more about [protecting your private information on Twitter and other websites](#).

Read our [policy regarding private information posted on Twitter](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter impersonation policy.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to Twitter's [parody, commentary, and fan account policy](#) for more information about these types of accounts.

Find more information about [reporting impersonation on Twitter](#).

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation rules

Updated November 3, 2017

This page is primarily intended for **developers**.

For Twitter users: You are ultimately responsible for the actions taken with your account, or by applications associated with your account. Before authorizing a third-party application to access or use your account, make sure you've thoroughly investigated the application and understand what it will do. If automated activity on your account violates the [Twitter Rules](#) or these automation rules, Twitter may take action on your account, including [filtering your Tweets from search results](#) or suspending your account.

For more information on third-party applications, please see our article on [connecting and revoking third-party applications](#).

If you are a developer and have questions about these Automation Rules, please reach out via <https://twittercommunity.com/c/rules-and-policies>.

If your application has been suspended or restricted from performing write actions, you may submit an appeal at </forms/platform>.

I. Ground Rules Do!

Build solutions that automatically broadcast helpful information in Tweets.

Run creative campaigns that auto-reply to users who engage with your content.

Build solutions that automatically respond to users in Direct Messages.

Try new things that help people (and comply with our rules).

Make sure your application provides a good user experience and performs well — and confirm that remains the case over time.

Don't!

Violate these or other policies. Be extra mindful of our rules about abuse and user privacy.

Abuse the Twitter API or attempt to circumvent rate limits.

Use non-API-based forms of automation, such as scripting the Twitter website. The use of these techniques may result in the permanent suspension of your account.

Spam or bother users, or otherwise send them unsolicited messages.

A. The Twitter Rules and the Developer Agreement and Policy

As with all activity on Twitter, automated activity is subject to the [Twitter Rules](#) and, if you're a developer using the Twitter API, the [Developer Agreement and Policy](#).

You should carefully review these policies to ensure that your automated activity is compliant. Automated applications or activities that violate these policies, or that facilitate or induce users to violate them, may be subject to enforcement action, potentially including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these policies.

Although all aspects of the Twitter Rules and the Developer Agreement and Policy apply to automated activity, you should keep the following rules top of mind:

Spamming: You may not send automated Tweets or Direct Messages that are spam, or otherwise engage in spamming activity. Some examples of spammy behavior to avoid with automation include:

Trending topics: You may not automatically post about trending topics on Twitter, or use automation to attempt to influence or manipulate trending topics.

Multiple posts/accounts: You may not post duplicative or substantially similar Tweets on one account or over multiple accounts you operate.

Duplicate accounts: You may not create and/or automate multiple accounts for duplicative or substantially similar use cases.

However, automating multiple accounts for related but non-duplicative use cases is permitted. For example, you may automate separate accounts to Tweet when the Hubble Space Telescope passes over different cities, such as [San Francisco](#) or [Hong Kong](#).

Misleading links: You may not send automated Tweets or Direct Messages containing links that are misleading, including links that maliciously or deceptively redirect through landing pages or ad pages before displaying the final content.

Sensitive media: Automated Tweets and Direct Messages must comply with the [Twitter media policy](#), and you should mark your account as potentially sensitive if you intend to post graphic, pornographic, or potentially sensitive media.

Abusive behavior: You may not engage in any automated activity that encourages, promotes, or incites abuse, violence, hateful conduct, or harassment, on or off Twitter.

Private information: You may not post private or confidential information about a person without their prior express authorization.

B. Other Ground Rules for Automated Activity

In addition to the policies above, the following ground rules apply to all automated activity on Twitter:

Don't surprise or mislead users: Automated activity should honor users' expectations. Ask for the user's permission before taking an action if you aren't sure.

Mature content or profanity: Don't Direct Message, mention, or reply to users with potentially sensitive content (including profanity), unless they've clearly indicated an intent to receive it in advance.

Be thoughtful about the information you request or exchange on Twitter

Tweets: Don't ask users to send you personal or private information via a public Tweet. If you need additional personal or private information from a user to provide them with customer service (or other similar use cases), you should ask the user to share such information by Direct Message or another private channel. You might even consider adding a [Direct Message deep link](#) to your Tweet.

Direct Messages: You should only ask users for the minimum amount of information you need to provide them with service. If you need to request or exchange particularly sensitive information (such as credit card information), you should consider directing users to your website or other appropriate channel to do so.

II. Activity-Specific Rules

The activity-specific rules in this section apply to taking specific automated actions on Twitter. Please read these rules carefully, as they outline both permitted and prohibited use cases of automation.

Automated applications or activities that violate these rules, or that facilitate or induce users to violate them, may be subject to enforcement action, including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access the Twitter API based on violations of these rules. As a reminder, you should also carefully review the spam guidelines in the [Twitter Rules](#) to avoid having activities performed by you, your app, or other users through your app or service flagged as spam.

A. Automated Actions Through Another User's Account

Twitter users may authorize your app or service to [access their Twitter account through OAuth](#). A user authorizing your app or service to access their Twitter account through OAuth does not by itself constitute sufficient consent to take automated actions through that user's account.

You may only take automated actions through another Twitter user's account if you:

clearly describe to the user the types of automated actions that will occur; receive express consent from the user to take those automated actions; and immediately honor a user's request to opt-out of further automated actions.

If you substantially change the purpose or functionality of your app or service, you must re-obtain express consent from the user to take automated action through their account before doing so.

These requirements apply to any automated action taken through another Twitter user's account, including posting Tweets, sending Direct Messages, deleting Tweets or Direct Messages, or following/unfollowing other accounts. For applications that offer users the ability to delete Tweets in a bulk or automated manner, you must also clearly state that Tweets are not recoverable once deleted.

B. Automated Tweets

1. Posting automated Tweets

Automated Tweets that cross-post outside information: You may post automated Tweets based on sources of outside information — such as an RSS feed, weather data, etc. — as long as you are sufficiently authorized to publish such information.

Other automated Tweets (excluding mentions or replies): Provided you comply with all other rules, you may post automated Tweets for entertainment, informational, or novelty purposes. As a reminder, accounts posting duplicative, spammy, or otherwise prohibited content may be subject to suspension.

2. Posting automated mentions and replies

The reply and mention functions are intended to make communication between Twitter users easier. Automating these actions to reach many users on an unsolicited basis is an abuse of the feature, and is not permitted. For example, sending automated replies to Tweets based on keyword searches alone is not permitted. Spamy or duplicative use of mentions and replies may result in enforcement action, such as the removal of

your Tweets from Search or the suspension of your app or account.

However, you may send automated replies or mentions to Twitter users so long as:

in advance of sending the automated reply, the recipient or mentioned user(s) have requested or have clearly indicated an intent on Twitter to be contacted by you (i.e. opted in), for example by replying to a Tweet from your account, or by sending you a Direct Message;

you provide a clear and easy way for such users to opt-out of receiving automated replies and mentions, and promptly honor all such opt-out requests;

you only send one automated reply or mention per user interaction; and

the automated reply or mention is a reply to the user's original Tweet (if your campaign is based on users posting a reply to your Tweet).

Opt-in techniques and indications of user intent take many different forms, depending on the specifics of your use case and implementation. Some examples include:

A Tweet from your account that clearly indicates that a user taking a specific action on that Tweet (such as Retweeting it) will opt the user into receiving an automated response.

A mention of your account by the user in a manner suggesting the user clearly wishes or intends to receive a response. If you want to run an auto-reply campaign with a campaign- or use-case-specific hashtag, users should also mention you in their Tweets.

Note that a user following your account is not on its own a sufficient indication of user intent to receive an automated response.

: Advertisers, publishers, and brands using auto-response campaigns must request approval from Twitter and may be subject to additional rules. Please reach out to your account or partner manager for assistance. If you do not have a partner manager but are interested in managed client services, you can learn more [here](#).

Additionally, we recommend that any accounts that will communicate with

users via automated mentions or replies:

Note

1. Appropriately filter responses based on potentially sensitive language in user handles,

display names, and Tweet text, as well as potentially sensitive media;

2. Check that the Tweet you are mentioning or replying to still exists (for example, using the statuses/lookup endpoint on the Twitter API).

C. Automated Direct Messages

Sending automated Direct Messages to users

You may not send unsolicited Direct Messages in a bulk or automated manner, and should be thoughtful about the frequency with which you contact users via Direct Message.

You may send automated Direct Messages to users so long as:

in advance of sending the Direct Message, the recipient(s) have requested or have clearly indicated an intent on Twitter to be contacted by you via Direct Message, for example by sending you a Direct Message; **and**

you provide a clear and easy way for such users to opt-out of receiving automated Direct Messages, and promptly honor all such opt-out requests.

The fact that a user is technically able to receive a Direct Message from you (e.g. because the user follows you, has enabled the ability to receive Direct Messages from any account, or because the user is in a pre-existing Direct Message conversation with you) does not necessarily mean they have requested or expect to receive automated Direct Messages from you.

Interacting with users via Direct Message

Per the Ground Rules, remember to be thoughtful about the amount and type of information you request or exchange with users via Direct Messages. If you will be asking a user to provide personal or private information via an automated Direct Message, you must clearly explain how you will use the information you're collecting. Consider including a link to your

privacy policy in your Direct Message to the user, as well as in your Twitter profile bio.

Don't publicly share information received in a Direct Message conversation with a user without first obtaining explicit consent from the user. For example, if a user asks you via Direct Message about a purchase they made from you, you may not mention the user in a public Tweet that includes information about their purchase unless you have the user's explicit consent to do so.

After a user-initiated interaction ends, don't send additional follow-up Direct Messages or mention users in a Tweet unless you get permission from the user.

D. Automated actions you take on Tweets or accounts

Automated likes: You may not like Tweets in an automated manner.

Automated Retweets: Provided you comply with all other rules, you may Retweet or Quote Tweet in an automated manner for entertainment, informational, or novelty purposes. Automated Retweets often lead to negative user experiences, and bulk, aggressive, or spammy Retweeting is a violation of the [Twitter Rules](#).

Automated following/unfollowing: You may not follow or unfollow Twitter accounts in a bulk, aggressive, or indiscriminate manner. Aggressive following is a violation of the [Twitter Rules](#). Please also review our [following rules and best practices](#) to ensure you are in compliance. Note that applications that claim to get users more followers are also prohibited under the [Twitter Rules](#).

Automated adding to lists or collections: You may not add Twitter users to lists or add Tweets to collections in a bulk or indiscriminate manner. Adding a large number of unrelated users to lists is a violation of the [Twitter Rules](#).

About search rules and restrictions

This article details the rules for using Twitter search, and explains reasons why content may not show up in Twitter search.

Why doesn't all content show up in Twitter search?

In order to keep your search results relevant, Twitter filters search results for quality Tweets and accounts. Material that jeopardizes search quality or creates a bad search experience for other people may be automatically removed from Twitter search. Read more about this [here](#).

General rules for using Twitter search

The [Twitter Rules](#) explain what behaviors are permitted on Twitter. In addition to these rules, we've included some tips below to keep your content relevant (and your Tweets in Twitter search).

Please do not:

- Repeatedly post duplicate or near-duplicate content (links or Tweets).
- Abuse trending topics or hashtags (topic words with a # sign).
- Send automated Tweets or replies.
- Use bots or applications to post similar messages based on keywords.
- Post similar messages over multiple accounts.
- Aggressively follow and unfollow people.

Twitter may automatically remove accounts engaging in these behaviors from search (or even suspend in some cases) in order to ensure the best experience for everyone.

Search results filtering

To learn how to filter your search results, including how to disable or enable safe search, read our [Using Twitter search](#) article.

Search results filtering in third-party applications

Accounts may be filtered from search for posting similar messages over several accounts. Your account may be filtered from search if you are logged into a third-party application that is updating many accounts with similar or duplicate content.

Rules for participating in contests on Twitter

Businesses sometimes host contests and offer prizes to Retweet an update, following a particular account, or posting Tweets with a specific hashtag topic or reply. If you're running a contest, check out this [article](#) for guidelines for contests on Twitter.

If you repeatedly post near-duplicate updates or duplicate links in order to enter contests, you may be filtered from search. Please keep in mind that you are responsible for the updates that you post to your Twitter account.

What to do if you are missing from search

Please see this [article](#) for more information related to this issue.

If you are being filtered from search, rest assured that your followers will still receive your updates. If you post a reply or mention to a specific person, it will be delivered to that person.

Having trouble?

You can learn more about searching by keyword, location, and trending topics in our [Using Twitter search](#) help article.

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. These limitations are set forth in the Twitter Rules below.

The Twitter Rules (along with all incorporated policies), [Privacy Policy](#), and [Terms of Service](#) collectively make up the "Twitter User Agreement" that governs a user's access to and use of Twitter's services.

All individuals accessing or using Twitter's services must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in Twitter taking one or more of the following enforcement actions:

- requiring you to delete prohibited content before you can again create new posts and interact with other Twitter users;
- temporarily limiting your ability to create posts or interact with other Twitter users;
- asking you to verify account ownership with a phone number or email address; or
- permanently suspending your account(s).

If you attempt to evade a permanent suspension by creating new accounts, we will suspend your new accounts.

Please note that we may need to change these Rules from time to time and reserve the right to do so. The most current version is always available at: <https://twitter.com/rules>.

The policies set forth in these Twitter Rules govern organic content on our platform. To learn more about the rules which govern ads and promoted content, please review our [Ads policies](#).

Content Boundaries and Use of Twitter

Intellectual property

Trademark: We reserve the right to suspend accounts or take other appropriate action when someone's brand or trademark, including business name and/or logo, is used in a manner that may mislead or confuse others about your brand affiliation. Read more about our [trademark policy](#) and [how to report a violation](#).

Copyright: We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in our [Terms of Service](#). Read more about our [copyright policy](#).

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Learn more about our [media policy](#).

Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images. Additionally, Twitter may sometimes require you to remove excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about [how to make such a request](#), and [how to mark your media as sensitive](#).

Unlawful use

You may not use our service for any unlawful purposes or in furtherance of illegal activities. By using Twitter, you agree to comply with all applicable laws governing your online conduct and content.

Trends

At times, we may prevent certain content from trending. This includes content that violates the Twitter Rules, as well as content that may attempt to manipulate trends. Read more about [what we allow and do not allow to trend](#).

Misuse of Twitter badges

You may not use badges, including but not limited to the “promoted” or “verified” Twitter badges, unless provided by Twitter. Accounts using unauthorized badges as part of their profile photos, header photos, display names, or in any way that falsely implies affiliation with Twitter or authorization from Twitter to display these badges, may be suspended.

Misuse of usernames

Selling usernames: You may not buy or sell Twitter usernames.

Username squatting: You may not engage in username squatting. Some of the factors we take into consideration when determining whether conduct is username squatting include:

- the number of accounts created;
- the creation of accounts for the purpose of preventing others from using those account names;
- the creation of accounts for the purpose of selling those accounts; and
- the use of third-party content feeds to update and maintain accounts under the names of those third parties.

Please note that Twitter may also remove accounts that are inactive for more than six months. Learn more about [username squatting](#).

Abusive Behavior

We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we prohibit behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Context matters when evaluating for abusive behavior and determining appropriate enforcement actions. Factors we may take into consideration include, but are not limited to whether:

- the behavior is targeted at an individual or group of people;
- the report has been filed by the target of the abuse or a bystander;
- the behavior is newsworthy and in the legitimate public interest.

Violence and physical harm

Violence: You may not make specific threats of violence or wish for the serious physical harm, death, or disease of an individual or group of people. This includes, but is not limited to, threatening or promoting terrorism. You also may not affiliate with organizations that — whether by their own statements or activity both on and off the platform — use or promote violence against civilians to further their causes. We will begin enforcing this rule around affiliation with such organizations on December 18, 2017.

Suicide or self-harm: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Child sexual exploitation: You may not promote child sexual exploitation. Learn more about our zero-tolerance [child sexual exploitation policy](#).

Abuse and hateful conduct

Abuse: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Unwanted sexual advances: You may not direct abuse at someone by sending unwanted sexual content, objectifying them in a sexually explicit manner, or otherwise engaging in sexual misconduct.

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. Read more about our [hateful conduct policy](#).

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category. We will begin enforcing this rule on December 18, 2017.

Private information and intimate media

Private information: You may not publish or post other people's private information without their express authorization and permission. Definitions of private information may vary depending on local laws. Read more about our [private information policy](#).

Intimate media: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. Read more about intimate media on Twitter.

Threats to expose / hack: You may not threaten to expose someone's private information or intimate media. You also may not threaten to hack or break into someone's digital information.

Impersonation

You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. While you may maintain parody, fan, commentary, or newsfeed accounts, you may not do so if the intent of the account is to engage in spamming or abusive behavior. Read more about our [impersonation policy](#).

Spam and Security

We strive to protect people on Twitter from technical abuse and spam.

To promote a stable and secure environment on Twitter, you may not do, or attempt to do, any of the following while accessing or using Twitter:

- Access, tamper with, or use non-public areas of Twitter, Twitter's computer systems, or the technical delivery systems of Twitter's providers (except as expressly permitted by the Twitter Bug Bounty program).
- Probe, scan, or test the vulnerability of any system or network, or breach or circumvent any security or authentication measures (except as expressly permitted by the Twitter Bug Bounty program).
- Access or search, or attempt to access or search, Twitter by any means (automated or otherwise) other than through our currently available, published interfaces that are provided by Twitter (and only pursuant to the applicable terms and conditions), unless you have been specifically allowed to do so in a separate agreement with Twitter. Note that crawling Twitter is permissible if done in accordance with the provisions of the robots.txt file; however, scraping Twitter without our prior consent is expressly prohibited.
- Forge any TCP/IP packet header or any part of the header information in any email or posting, or in any way use Twitter to send altered, deceptive, or false source-identifying information.
- Interfere with or disrupt the access of any user, host or network, including, without limitation, sending a virus, overloading, flooding, spamming, mail-bombing Twitter's services, or by scripting the creation of content in such a manner as to interfere with or create an undue burden on Twitter.

Any accounts engaging in the following activities may be temporarily locked or subject to permanent suspension:

- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use Twitter's services for the purpose of spamming anyone. Spam is generally defined on Twitter as bulk or aggressive activity that attempts to manipulate or disrupt Twitter or the experience of users on Twitter to drive traffic or attention to unrelated accounts, products, services, or initiatives. Some of the factors that we take into account when determining what conduct is considered to be spamming include:
 - if you have followed and/or unfollowed a large number of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if your Tweets or Direct Messages consist mainly of links shared without commentary;
 - if a large number of people have blocked you in response to high volumes of untargeted, unsolicited, or duplicative content or engagements from your account;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicative or substantially similar content, replies, or mentions over multiple accounts or multiple duplicate updates on one account, or create duplicate or substantially similar accounts;
 - if you post multiple updates to a trending or popular topic with an intent to subvert or manipulate the topic to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you send large numbers of unsolicited replies or mentions;
 - if you add users to lists in a bulk or aggressive manner;
 - if you are randomly or aggressively engaging with Tweets (e.g., likes, Retweets, etc.) or users (e.g., following, adding to lists or Moments, etc.) to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you repeatedly post other people's account information as your own (e.g., bio, Tweets, profile URL, etc.);
 - if you post misleading, deceptive, or malicious links (e.g., affiliate links, links to malware/clickjacking pages, etc.);
 - if you create fake accounts, account interactions, or impressions;
 - if you sell, purchase, or attempt to artificially inflate account interactions (such as followers, Retweets, likes, etc.); and
 - if you use or promote third-party services or apps that claim to get you more followers, Retweets, or likes (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers or engagements to your account or Tweets).

Please see our support articles on [following rules and best practices](#) and [automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts may be permanently suspended.

Content Visibility

Accounts under investigation or which have been detected as sharing content in violation of these Rules may have their account or Tweet visibility limited in various parts of Twitter, including search. To learn more about situations in which content may be restricted on Twitter, please see our support article on [search rules and restrictions](#).

Trademark policy

Updated June 30, 2016

What is a Trademark Policy Violation on Twitter?

Using another's trademark in a manner that may mislead or confuse others about your brand affiliation may be a violation of our trademark policy.

How Does Twitter Respond to Trademark Reports?

We review each report and do the following:

- If we determine that there is a violation of our trademark policy, we may suspend the account. In some instances, we may give the account holder an opportunity to comply with our policies before suspending the account. For more information, see our [Parody, commentary, and fan account policy](#).
- We may also provide the account holder with your name and information included in the copy of the report.
- If you would like to report a trademark violation in a [Twitter Ad](#), please visit our [Trademark Policy for Promoted Ads](#) to learn more.

Please note that Twitter only investigates requests that are submitted by the trademark holder or by an authorized representative of the trademark holder.

What is not a Trademark Policy Violation?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples include:

- Using a trademark in a way that is outside the scope of the trademark registration (e.g. territory, or goods and services identified in the registration).
- Nominative and other fair uses of trademarks are protected uses under our trademark policy, so long as the account is clearly distinguished from the trademark owner. This includes use by resellers in certain regions and accounts engaging in parody, commentary, or news. For more information, see our [Parody, commentary, and fan account policy](#).

How do I Report a Trademark Policy Violation?

Submit trademark reports through our [Trademark Issue support form](#). Using the appropriate form ensures we have all the information necessary to process your report quickly. Please provide all information requested in the form. If you submit a report with incomplete information, we will follow up with you to request that information. Please note that this will result in a delay in processing your report.

If you have trouble using the form, try updating your browser or switching to another browser.

Copyright policy

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged

infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notification.

When Should I File a Counter-notice?

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and

4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Twitter media policy

Twitter Rules: Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images.

Twitter Rules: You may not use hateful images or symbols in your profile image or profile header.

Rationale

Some forms of graphic violence, adult content, or hateful imagery may be permitted in Tweets when they are marked as [sensitive media](#). However, you may not include this type of content in live video, or in profile or header images.

While we want people to feel free to share media that reflects their creativity or individuality, or to show what's happening in the world, we will take action when it crosses the line into abuse towards a person, group, or protected category.

Finally, we may sometimes require you to remove media containing excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about how to make such a request [here](#).

When this applies

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. Some examples include, but are not limited to, depictions of:

- the moment at which someone dies
- gruesome crime or accident scenes
- bodily harm, torture, dismemberment, or mutilation

We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity (including close-ups of genitals, buttocks, or breasts)
 - Please note that exceptions may be made for artistic, medical, health, or educational content. Breastfeeding content does not need to be marked as sensitive.
- simulating a sexual act
- intercourse or any sexual act (may involve humans, humanoid animals, cartoons, or anime)

If you see media you think may be in violation of our policy, report it. You can read about how to report sensitive media [here](#).

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, or ethnicity/national origin.

Some examples of hateful imagery include, but are not limited to:

- Symbols historically associated with hate groups (for example, the Nazi swastika)
- Images depicting others as less than human or altered to include hateful symbols
- Altered image references to a mass murder that targeted a protected category

Do I need to be the target of this content in order for it to be a violation of the Twitter Rules?
No, we review both first-person and bystander reports of such content.

Consequences

When adult content, graphic violence, or hateful imagery appears in Tweets, we may place this content behind an interstitial advising viewers to be aware that they will see sensitive media if they click through. This allows us to identify potentially sensitive content that some people may not wish to see. Learn more about [how to control whether you see sensitive media](#).

However, if such content appears in live video, header, or profile images, it will violate our media policy. The consequences for violating this policy depend on the severity of the violation and the person's previous record of violations.

The first time someone violates this policy they will be required to remove the imagery and may be asked to verify their account ownership. Subsequent violations could result in permanent suspension. Learn more about [our range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

How to contact Twitter about media concerning a deceased family member

In order to respect the wishes of loved ones, Twitter will remove imagery of deceased individuals in certain circumstances. Immediate family members and other authorized individuals may request the removal of images or video of deceased individuals, from when critical injury occurs to the moments before or after death, by sending a request to Twitter Inc. via our [privacy form](#). When reviewing such media removal requests, Twitter considers public interest factors such as the newsworthiness of the content and may not be able to honor every request.

Note: Find information about contacting Twitter to [deactivate the account of a deceased or incapacitated user](#).

How to mark media as sensitive links to Media policy

Twitter trends FAQs

Where can I find trends?

Trends are available on the Twitter for iOS app, the Twitter for Android app, and [twitter.com](#).

On Twitter's mobile apps, you can find trends listed under the trends section of the **Explore** tab. When signed in to [twitter.com](#) on a desktop or laptop computer, trends are listed in many places, including the Home timeline, Notifications, search results, and profile pages.

How are trends determined?

Trends are determined by an algorithm and, by default, are tailored for you based on who you follow, your interests, and your location. This algorithm identifies topics that are popular now, rather than topics that have been popular for a while or on a daily basis, to help you discover the hottest emerging topics of discussion on Twitter.

You can choose to see trends that are not tailored for you by selecting a specific trends location on [twitter.com](#), iOS, or Android (instructions below). Location trends identify popular topics among people in a specific geographic location.

Note: The number of Tweets that are related to the trends is just one of the factors the algorithm looks at when ranking and determining trends. Algorithmically, trends and hashtags are grouped together if they are related to the same topic. For instance, #MondayMotivation and #MotivationMonday may both be represented by #MondayMotivation.

What does the # sign mean?

You may notice that some trends have # sign before the word or phrase. This is called a **hashtag** and is included specifically in Tweets to mark them as relating to a topic, so that people can follow the conversation in search.

Will I see world and local events in my trends?

Yes, when signed into your account on the twitter.com, iOS, or Android, trends will be tailored for you based on your location and who you follow. There will be many world and local news events and conversations that will appear in your trends regardless of your personalization.

Can I see trends for a specific location?

Yes, you can view trends for a specific location on twitter.com, iOS, or Android.

View trends via twitter.com:

To view trends for a specific location, click **Change** in your trends box. You can choose to **Keep tailored trends** based on your location and who you follow, or you can select **Change** to choose a nearby location or search locations.

Adjust your trends settings via Twitter for iOS: 0.

In the top menu, tap your profile icon.

17. Tap **Settings and privacy**, then tap **Content preferences**.
18. Under **Explore**, tap **Trends**.

T. Drag the slider next to **Trends for you** to turn it off and receive personalized trends.

V. When the feature is disabled, you can change your location by tapping **Change location**. In the search box, type your desired trends location or choose from the list.

Adjust your trends settings via Twitter for Android:

O. In the top menu, you will either see a navigation menu icon or your profile

icon. Tap whichever icon you have, and select **Settings and privacy**.

17. Tap **Content preferences**.

18. Under **Explore**, tap **Trends**.

T. Drag the slider next to **Trends for you** to turn on and receive personalized trends.

V. When the feature is disabled, you can change your location by tapping **Change location**. In the search box, type your desired trends location or choose from the list.

Note: If you don't find your preferred country or city, it means we're not receiving enough Tweets from that geographical area to create a list. You can look up local Tweets on any topic by using [advanced Twitter search operators](#).

What happens when I click or tap on a trend?

Clicking or tapping any of the trends takes you to the Twitter search results for that trend. You'll see all Tweets including that phrase or hashtag. To see what people are saying about a previous trend, perform a [search for that keyword](#).

How can I participate in a trend?

Simply post a Tweet including the exact word or phrase as it appears in the trends list (with the hashtag, if you see one). Due to the large number of people Tweeting about these specific trends, you may not always be able to find your particular Tweet in search, but your followers will always see your Tweets.

Note: Twitter also filters searches for quality. To understand what types of behavior could cause your Tweets to be filtered out, read about our [search rules and restrictions](#).

Are there rules for trends?

Yes. We want trends to promote healthy discussions on Twitter. This means that at times, we may prevent certain content from trending. These include trends that:

Contain profanity or adult/graphic references.

Incite hate on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease.

Violate the [Twitter Rules](#).

In some cases, we may also consider the newsworthiness of the content, or if it is in the public interest when evaluating potential violations. In these cases, the content might continue to trend on our platform.

Note that even if we prevent the hashtag or the content from appearing on the list of trending topics, you may still be able to access conversations around that content on Twitter.

Do trends have descriptions?

Below some trends, you may see the approximate number of Tweets associated with the trend, or personalized context such as who in your network is Tweeting about it.

Additionally, when a Moment is found to be relevant to a trend, we may use the title of the Moment as the trend's description.

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Child sexual exploitation policy

Twitter removes content promoting child sexual exploitation

We do not tolerate child sexual exploitation on Twitter. When we are made aware of links to images of or content promoting child sexual exploitation they will be removed from the site without further notice and reported to The National Center for Missing & Exploited Children ("NCMEC"); we permanently suspend accounts promoting or containing updates with links to child sexual exploitation.

Reporting content promoting child sexual exploitation

If you think you've found a Twitter account distributing or promoting child sexual exploitation while using Twitter, please click [here](#) to report it. Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. To find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with the National Center for Missing and Exploited Children [here](#), or your local law enforcement. If you believe that the content was hosted on a website outside of the United States, you can report it [here](#) on the International Association of Internet Hotlines website.

NOTE: Please do NOT tweet, retweet or repost child sexual exploitation for any reason. Click [here](#) to report it to us immediately and we will take steps to remove it.

Law Enforcement Agencies

If you are a member of a law enforcement agency seeking information regarding Twitter accounts promoting or distributing child sexual exploitation, please see our [Guidelines for Law Enforcement](#) for information on our process; we also suggest you contact the National Center for Missing and Exploited Children (NCMEC) [here](#), regardless of jurisdiction, to determine whether or not information on those accounts is already on file.

For general law enforcement inquiries, please see our [Guidelines for Law Enforcement](#) for more information.

Hateful conduct policy

Freedom of expression means little if voices are silenced because people are afraid to speak up. We do not tolerate behavior that harasses, intimidates, or uses fear to silence another person's voice. If you see something on Twitter that violates these rules, please report it to us.

How our policy works

As explained in the Twitter Rules,

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual

orientation, gender, gender identity, religious affiliation, age, disability, or disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Examples of what we do not tolerate includes, but is not limited to behavior that harasses individuals or groups of people with:

violent threats;

wishes for the physical harm, death, or disease of individuals or groups;

references to mass murder, violent events, or specific means of violence in which/with which such groups have been the primary targets or victims;

behavior that incites fear about a protected group;

repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone.

How our enforcement works Context matters.

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. While we accept reports of violations from anyone, sometimes we also need to hear directly from the target to ensure that we have proper context.

The number of reports we receive does not impact whether or not something will be removed. However, it may help us prioritize the order in which it gets reviewed.

We focus on behavior.

We enforce policies when someone reports behavior that is abusive and targets an entire protected group and/or individuals who may be members.

This targeting can happen in any manner (for example, @mentions, tagging a photo, and more).

We have a range of enforcement options.

The consequences for violating our rules vary depending on the severity of the violation and the person's previous record of violations. For example, we may ask someone to remove the offending Tweet before they can Tweet again. For other cases, we may suspend an account.

Report private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Read about our [policy regarding private information posted on Twitter](#).

Do I need to have a Twitter account to make a report?

No, you do not need to have a Twitter account to file a report that your private and confidential information has been posted on Twitter.

Can I report unauthorized posting of Copyrighted material?

If your copyrighted material has been Tweeted without your permission, please review our [Copyright and DMCA policy](#).

What can I do if my information is posted on other websites?

We cannot take down any information that is posted on a website other than Twitter. If your information is posted on another website, please contact the help center or support staff for the website where your information is posted.

How do I report that my private information is posted on Twitter?

You can report that your private information has been posted on Twitter directly from the offending Tweet(s). Learn [how to report individual Tweets](#).

File a [report that someone has posted your private information](#).

What happens when Twitter receives a valid report?

When we receive a complete and valid report that private information has been posted on Twitter, we'll investigate the account and Tweets reported. We will review where, if anywhere, the information has been made publicly available before taking action on the account or Tweets. If the information you reported was previously posted elsewhere on the Internet, it is not a violation of our policy and we will not take action.

Learn more about [protecting your private information on Twitter and other websites](#).

Read our [policy regarding private information posted on Twitter](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter impersonation policy.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to Twitter's [parody, commentary, and fan account policy](#) for more information about these types of accounts.

Find more information about [reporting impersonation on Twitter](#).

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000

people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation rules

Updated November 3, 2017

This page is primarily intended for **developers**.

For Twitter users: You are ultimately responsible for the actions taken with your account, or by applications associated with your account. Before authorizing a third-party application to access or use your account, make sure you've thoroughly investigated the application and understand what it will do. If automated activity on your account violates the [Twitter Rules](#) or these automation rules, Twitter may take action on your account, including [filtering your Tweets from search results](#) or suspending your account.

For more information on third-party applications, please see our article on [connecting and revoking third-party applications](#).

If you are a developer and have questions about these Automation Rules, please reach out via <https://twittercommunity.com/c/rules-and-policies>.

If your application has been suspended or restricted from performing write actions, you may submit an appeal at </forms/platform>.

I. Ground Rules Do!

Build solutions that automatically broadcast helpful information in Tweets.

Run creative campaigns that auto-reply to users who engage with your content.

Build solutions that automatically respond to users in Direct Messages.

Try new things that help people (and comply with our rules).

Make sure your application provides a good user experience and performs well — and confirm that remains the case over time.

Don't!

Violate these or other policies. Be extra mindful of our rules about abuse and user privacy.

Abuse the Twitter API or attempt to circumvent rate limits.

Use non-API-based forms of automation, such as scripting the Twitter website. The use of these techniques may result in the permanent suspension of your account.

Spam or bother users, or otherwise send them unsolicited messages.

A. The Twitter Rules and the Developer Agreement and Policy

As with all activity on Twitter, automated activity is subject to the [Twitter Rules](#) and, if you're a developer using the Twitter API, the [Developer Agreement and Policy](#).

You should carefully review these policies to ensure that your automated activity is compliant. Automated applications or activities that violate these policies, or that facilitate or induce users to violate them, may be subject to enforcement action, potentially including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these policies.

Although all aspects of the Twitter Rules and the Developer Agreement and Policy apply to automated activity, you should keep the following rules top of mind:

Spamming: You may not send automated Tweets or Direct Messages that are spam, or otherwise engage in spamming activity. Some examples of spammy behavior to avoid with automation include:

Trending topics: You may not automatically post about trending topics on Twitter, or use automation to attempt to influence or manipulate trending topics.

Multiple posts/accounts: You may not post duplicative or substantially similar Tweets on one account or over multiple accounts you operate.

Duplicate accounts: You may not create and/or automate multiple accounts for duplicative or substantially similar use cases.

However, automating multiple accounts for related but non-duplicative use cases is permitted. For example, you may automate separate accounts to Tweet when the Hubble Space Telescope passes over different cities, such as [San Francisco](#) or [Hong Kong](#).

Misleading links: You may not send automated Tweets or Direct Messages containing links that are misleading, including links that maliciously or deceptively redirect through landing pages or ad pages before displaying the final content.

Sensitive media: Automated Tweets and Direct Messages must comply with the [Twitter media policy](#), and you should mark your account as potentially sensitive if you intend to post graphic, pornographic, or potentially sensitive media.

Abusive behavior: You may not engage in any automated activity that encourages, promotes, or incites abuse, violence, hateful conduct, or harassment, on or off Twitter.

Private information: You may not post private or confidential information about a person without their prior express authorization.

B. Other Ground Rules for Automated Activity

In addition to the policies above, the following ground rules apply to all automated activity on Twitter:

Don't surprise or mislead users: Automated activity should honor users' expectations. Ask for the user's permission before taking an action if you aren't sure.

Mature content or profanity: Don't Direct Message, mention, or reply to users with potentially sensitive content (including profanity), unless they've clearly indicated an intent to receive it in advance.

Be thoughtful about the information you request or exchange on Twitter

Tweets: Don't ask users to send you personal or private information via a public Tweet. If you need additional personal or private information from a user to provide them with customer service (or other similar use cases), you should ask the user to share such information by Direct Message or another private channel. You might even consider adding a [Direct Message deep link](#) to your Tweet.

Direct Messages: You should only ask users for the minimum amount of information you need to provide them with service. If you need to request or exchange particularly sensitive information (such as credit card information), you should consider directing users to your website or other appropriate channel to do so.

II. Activity-Specific Rules

The activity-specific rules in this section apply to taking specific automated actions on Twitter. Please read these rules carefully, as they outline both permitted and prohibited use cases of automation.

Automated applications or activities that violate these rules, or that facilitate or induce users to violate them, may be subject to enforcement action, including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access the Twitter API based on violations of these rules. As a reminder, you should also carefully review the spam guidelines in the [Twitter Rules](#) to avoid having activities performed by you, your app, or other users through your app or service flagged as spam.

A. Automated Actions Through Another User's Account

Twitter users may authorize your app or service to [access their Twitter account through OAuth](#). A user authorizing your app or service to access their Twitter account through OAuth does not

by itself constitute sufficient consent to take automated actions through that user's account.

You may only take automated actions through another Twitter user's account if you:

clearly describe to the user the types of automated actions that will occur; receive express consent from the user to take those automated actions; and immediately honor a user's request to opt-out of further automated actions.

If you substantially change the purpose or functionality of your app or service, you must re-obtain express consent from the user to take automated action through their account before doing so.

These requirements apply to any automated action taken through another Twitter user's account, including posting Tweets, sending Direct Messages, deleting Tweets or Direct Messages, or following/unfollowing other accounts. For applications that offer users the ability to delete Tweets in a bulk or automated manner, you must also clearly state that Tweets are not recoverable once deleted.

B. Automated Tweets

1. Posting automated Tweets

Automated Tweets that cross-post outside information:

You may post automated Tweets based on sources of outside information — such as an RSS feed, weather data, etc. — as long as you are sufficiently authorized to publish such information.

Other automated Tweets (excluding mentions or replies):

Provided you comply with all other rules, you may post automated Tweets for entertainment, informational, or novelty purposes. As a reminder, accounts posting duplicative, spammy, or otherwise prohibited content may be subject to suspension.

2. Posting automated mentions and replies

The reply and mention functions are intended to make communication between Twitter users easier. Automating these actions to reach many users on an unsolicited basis is an abuse of the feature, and is not permitted. For example, sending automated replies to Tweets based on keyword searches alone is not permitted. Spammy or duplicative use of mentions and replies may result in enforcement action, such as the removal of your Tweets from Search or the suspension of your app or account.

However, you may send automated replies or mentions to Twitter users so long as:

in advance of sending the automated reply, the recipient or mentioned user(s) have requested or have clearly indicated an intent on Twitter to be contacted by you (i.e. opted in), for example by replying to a Tweet from your account, or by sending you a Direct Message;

you provide a clear and easy way for such users to opt-out of receiving automated replies and mentions, and promptly honor all such opt-out requests;

you only send one automated reply or mention per user interaction; and

the automated reply or mention is a reply to the user's original Tweet (if your campaign is based on users posting a reply to your Tweet).

Opt-in techniques and indications of user intent take many different forms, depending on the specifics of your use case and implementation. Some examples include:

A Tweet from your account that clearly indicates that a user taking a specific action on that Tweet (such as Retweeting it) will opt the user into receiving an automated response.

A mention of your account by the user in a manner suggesting the user clearly wishes or intends to receive a response. If you want to run an auto-reply campaign with a campaign- or use-case-specific hashtag, users should also mention you in their Tweets.

Note that a user following your account is not on its own a sufficient indication of user intent to receive an automated response.

: Advertisers, publishers, and brands using auto-response campaigns must request approval from Twitter

and may be subject to additional rules. Please reach out to your account or partner manager for assistance. If you do not have a partner manager but are interested in managed client services, you can learn more [here](#).

Additionally, we recommend that any accounts that will communicate with

users via automated mentions or replies:

Note

1. Appropriately filter responses based on potentially sensitive language in user handles,

display names, and Tweet text, as well as potentially sensitive media;

2. Check that the Tweet you are mentioning or replying to still exists (for example, using the statuses/lookup endpoint on the Twitter API).

C. Automated Direct Messages

Sending automated Direct Messages to users

You may not send unsolicited Direct Messages in a bulk or automated manner, and should be thoughtful about the frequency with which you contact users via Direct Message.

You may send automated Direct Messages to users so long as:

in advance of sending the Direct Message, the recipient(s) have requested or have clearly indicated an intent on Twitter to be contacted by you via Direct Message, for example by sending you a Direct Message; **and**

you provide a clear and easy way for such users to opt-out of receiving automated Direct Messages, and promptly honor all such opt-out requests.

The fact that a user is technically able to receive a Direct Message from you (e.g. because the user follows you, has enabled the ability to receive Direct Messages from any account, or because the user is in a pre-existing Direct Message conversation with you) does not necessarily mean they have requested or expect to receive automated Direct Messages from you.

Interacting with users via Direct Message

Per the Ground Rules, remember to be thoughtful about the amount and type of information you request or exchange with users via Direct Messages. If you will be asking a user to provide personal or private information via an automated Direct Message, you must clearly explain how you will use the information you're collecting. Consider including a link to your privacy policy in your Direct Message to the user, as well as in your Twitter profile bio.

Don't publicly share information received in a Direct Message conversation with a user without first obtaining explicit consent from the user. For example, if a user asks you via Direct Message about a purchase they made from you, you may not mention the user in a public Tweet that includes information about their purchase unless you have the user's explicit consent to do so.

After a user-initiated interaction ends, don't send additional follow-up Direct Messages or mention users in a Tweet unless you get permission from the user.

D. Automated actions you take on Tweets or accounts

Automated likes: You may not like Tweets in an automated manner.

Automated Retweets: Provided you comply with all other rules, you may Retweet or Quote Tweet in an automated manner for entertainment, informational, or novelty purposes. Automated Retweets often lead to negative user experiences, and bulk, aggressive, or spammy Retweeting is a violation of the [Twitter Rules](#).

Automated following/unfollowing: You may not follow or unfollow Twitter accounts in a bulk, aggressive, or indiscriminate manner. Aggressive following is a violation of the [Twitter Rules](#). Please also review our [following rules and best practices](#) to ensure you are in compliance. Note that applications that claim to get users more followers are also prohibited under the [Twitter Rules](#).

Automated adding to lists or collections: You may not add Twitter users to lists or add Tweets to collections in a bulk or indiscriminate manner. Adding a large number of unrelated users to lists is a violation of the [Twitter Rules](#).

About search rules and restrictions

This article details the rules for using Twitter search, and explains reasons why content may not show up in Twitter search.

Why doesn't all content show up in Twitter search?

In order to keep your search results relevant, Twitter filters search results for quality Tweets and accounts. Material that jeopardizes search quality or creates a bad search experience for other people may be automatically removed from Twitter search. Read more about this [here](#).

General rules for using Twitter search

The [Twitter Rules](#) explain what behaviors are permitted on Twitter. In addition to these rules, we've included some tips below to keep your content relevant (and your Tweets in Twitter search).

Please do not:

- Repeatedly post duplicate or near-duplicate content (links or Tweets).
- Abuse trending topics or hashtags (topic words with a # sign).
- Send automated Tweets or replies.
- Use bots or applications to post similar messages based on keywords.
- Post similar messages over multiple accounts.
- Aggressively follow and unfollow people.

Twitter may automatically remove accounts engaging in these behaviors from search (or even suspend in some cases) in order to ensure the best experience for everyone.

Search results filtering

To learn how to filter your search results, including how to disable or enable safe search, read our [Using Twitter search](#) article.

Search results filtering in third-party applications

Accounts may be filtered from search for posting similar messages over several accounts. Your account may be filtered from search if you are logged into a third-party application that is updating many accounts with similar or duplicate content.

Rules for participating in contests on Twitter

Businesses sometimes host contests and offer prizes to Retweet an update, following a particular account, or posting Tweets with a specific hashtag topic or reply. If you're running a contest, check out this [article](#) for guidelines for contests on Twitter.

If you repeatedly post near-duplicate updates or duplicate links in order to enter contests, you may be filtered from search. Please keep in mind that you are responsible for the updates that you post to your Twitter account.

What to do if you are missing from search

Please see this [article](#) for more information related to this issue.

If you are being filtered from search, rest assured that your followers will still receive your updates. If you post a reply or mention to a specific person, it will be delivered to that person.

Having trouble?

You can learn more about searching by keyword, location, and trending topics in our [Using Twitter search](#) help article.

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. These limitations are set forth in the Twitter Rules below.

The Twitter Rules (along with all incorporated policies), [Privacy Policy](#), and [Terms of Service](#) collectively make up the "Twitter User Agreement" that governs a user's access to and use of Twitter's services.

All individuals accessing or using Twitter's services must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in Twitter taking one or more of the following enforcement actions:

- requiring you to delete prohibited content before you can again create new posts and interact with other Twitter users;
- temporarily limiting your ability to create posts or interact with other Twitter users;
- asking you to verify account ownership with a phone number or email address; or
- permanently suspending your account(s).

If you attempt to evade a permanent suspension by creating new accounts, we will suspend your new accounts.

Please note that we may need to change these Rules from time to time and reserve the right to do so. The most current version is always available at: <https://twitter.com/rules>.

The policies set forth in these Twitter Rules govern organic content on our platform. To learn more about the rules which govern ads and promoted content, please review our [Ads policies](#).

Content Boundaries and Use of Twitter

Intellectual property

Trademark: We reserve the right to suspend accounts or take other appropriate action when someone's brand or trademark, including business name and/or logo, is used in a manner that may mislead or confuse others about your brand affiliation. Read more about our [trademark policy](#) and [how to report a violation](#).

Copyright: We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in our [Terms of Service](#). Read more about our [copyright policy](#).

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Learn more about our [media policy](#).

Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images. Additionally, Twitter may sometimes require you to remove excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about [how to make such a request](#), and [how to mark your media as sensitive](#).

Unlawful use

You may not use our service for any unlawful purposes or in furtherance of illegal activities. By using Twitter, you agree to comply with all applicable laws governing your online conduct and content.

Trends

At times, we may prevent certain content from trending. This includes content that violates the Twitter Rules, as well as content that may attempt to manipulate trends. Read more about [what we allow and do not allow to trend](#).

Misuse of Twitter badges

You may not use badges, including but not limited to the “promoted” or “verified” Twitter badges, unless provided by Twitter. Accounts using unauthorized badges as part of their profile photos, header photos, display names, or in any way that falsely implies affiliation with Twitter or authorization from Twitter to display these badges, may be suspended.

Misuse of usernames

Selling usernames: You may not buy or sell Twitter usernames.

Username squatting: You may not engage in username squatting. Some of the factors we take into consideration when determining whether conduct is username squatting include:

- the number of accounts created;
- the creation of accounts for the purpose of preventing others from using those account names;
- the creation of accounts for the purpose of selling those accounts; and
- the use of third-party content feeds to update and maintain accounts under the names of those third parties.

Please note that Twitter may also remove accounts that are inactive for more than six months. Learn more about [username squatting](#).

Abusive Behavior

We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we prohibit behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Context matters when evaluating for abusive behavior and determining appropriate enforcement actions. Factors we may take into consideration include, but are not limited to whether:

- the behavior is targeted at an individual or group of people;
- the report has been filed by the target of the abuse or a bystander;
- the behavior is newsworthy and in the legitimate public interest.

Violence and physical harm

Violence: You may not make specific threats of violence or wish for the serious physical harm, death, or disease of an individual or group of people. This includes, but is not limited to, threatening or promoting terrorism. You also may not affiliate with organizations that — whether by their own statements or activity both on and off the platform — use or promote violence against civilians to further their causes. We will begin enforcing this rule around affiliation with such organizations on December 18, 2017.

Suicide or self-harm: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Child sexual exploitation: You may not promote child sexual exploitation. Learn more about our zero-tolerance [child sexual exploitation policy](#).

Abuse and hateful conduct

Abuse: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Unwanted sexual advances: You may not direct abuse at someone by sending unwanted sexual content, objectifying them in a sexually explicit manner, or otherwise engaging in sexual misconduct.

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. Read more about our [hateful conduct policy](#).

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category. We will begin enforcing this rule on December 18, 2017.

Private information and intimate media

Private information: You may not publish or post other people's private information without their express authorization and permission. Definitions of private information may vary depending on local laws. Read more about our [private information policy](#).

Intimate media: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. Read more about intimate media on Twitter.

Threats to expose / hack: You may not threaten to expose someone's private information or intimate media. You also may not threaten to hack or break into someone's digital information.

Impersonation

You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. While you may maintain parody, fan, commentary, or newsfeed accounts, you may not do so if the intent of the account is to engage in spamming or abusive behavior. Read more about our [impersonation policy](#).

Spam and Security

We strive to protect people on Twitter from technical abuse and spam.

To promote a stable and secure environment on Twitter, you may not do, or attempt to do, any of the following while accessing or using Twitter:

- Access, tamper with, or use non-public areas of Twitter, Twitter's computer systems, or the technical delivery systems of Twitter's providers (except as expressly permitted by the Twitter Bug Bounty program).
- Probe, scan, or test the vulnerability of any system or network, or breach or circumvent any security or authentication measures (except as expressly permitted by the Twitter Bug Bounty program).
- Access or search, or attempt to access or search, Twitter by any means (automated or otherwise) other than through our currently available, published interfaces that are provided by Twitter (and only pursuant to the applicable terms and conditions), unless you have been specifically allowed to do so in a separate agreement with Twitter. Note that crawling Twitter is permissible if done in accordance with the provisions of the robots.txt file; however, scraping Twitter without our prior consent is expressly prohibited.
- Forge any TCP/IP packet header or any part of the header information in any email or posting, or in any way use Twitter to send altered, deceptive, or false source-identifying information.
- Interfere with or disrupt the access of any user, host or network, including, without limitation, sending a virus, overloading, flooding, spamming, mail-bombing Twitter's services, or by scripting the creation of content in such a manner as to interfere with or create an undue burden on Twitter.

Any accounts engaging in the following activities may be temporarily locked or subject to permanent suspension:

- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use Twitter's services for the purpose of spamming anyone. Spam is generally defined on Twitter as bulk or aggressive activity that attempts to manipulate or disrupt Twitter or the experience of users on Twitter to drive traffic or attention to unrelated accounts, products, services, or initiatives. Some of the factors that we take into account when determining what conduct is considered to be spamming include:
 - if you have followed and/or unfollowed a large number of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if your Tweets or Direct Messages consist mainly of links shared without commentary;
 - if a large number of people have blocked you in response to high volumes of untargeted, unsolicited, or duplicative content or engagements from your account;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicative or substantially similar content, replies, or mentions over multiple accounts or multiple duplicate updates on one account, or create duplicate or substantially similar accounts;
 - if you post multiple updates to a trending or popular topic with an intent to subvert or manipulate the topic to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you send large numbers of unsolicited replies or mentions;
 - if you add users to lists in a bulk or aggressive manner;
 - if you are randomly or aggressively engaging with Tweets (e.g., likes, Retweets, etc.) or users (e.g., following, adding to lists or Moments, etc.) to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you repeatedly post other people's account information as your own (e.g., bio, Tweets, profile URL, etc.);
 - if you post misleading, deceptive, or malicious links (e.g., affiliate links, links to malware/clickjacking pages, etc.);
 - if you create fake accounts, account interactions, or impressions;
 - if you sell, purchase, or attempt to artificially inflate account interactions (such as followers, Retweets, likes, etc.); and
 - if you use or promote third-party services or apps that claim to get you more followers, Retweets, or likes (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers or engagements to your account or Tweets).

Please see our support articles on [following rules and best practices](#) and [automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts may be permanently suspended.

Content Visibility

Accounts under investigation or which have been detected as sharing content in violation of these Rules may have their account or Tweet visibility limited in various parts of Twitter, including search. To learn more about situations in which content may be restricted on Twitter, please see our support article on [search rules and restrictions](#).

Trademark policy

Updated June 30, 2016

What is a Trademark Policy Violation on Twitter?

Using another's trademark in a manner that may mislead or confuse others about your brand affiliation may be a violation of our trademark policy.

How Does Twitter Respond to Trademark Reports?

We review each report and do the following:

- If we determine that there is a violation of our trademark policy, we may suspend the account. In some instances, we may give the account holder an opportunity to comply with our policies before suspending the account. For more information, see our [Parody, commentary, and fan account policy](#).
- We may also provide the account holder with your name and information included in the copy of the report.
- If you would like to report a trademark violation in a [Twitter Ad](#), please visit our [Trademark Policy for Promoted Ads](#) to learn more.

Please note that Twitter only investigates requests that are submitted by the trademark holder or by an authorized representative of the trademark holder.

What is not a Trademark Policy Violation?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples include:

- Using a trademark in a way that is outside the scope of the trademark registration (e.g. territory, or goods and services identified in the registration).
- Nominative and other fair uses of trademarks are protected uses under our trademark policy, so long as the account is clearly distinguished from the trademark owner. This includes use by resellers in certain regions and accounts engaging in parody, commentary, or news. For more information, see our [Parody, commentary, and fan account policy](#).

How do I Report a Trademark Policy Violation?

Submit trademark reports through our [Trademark Issue support form](#). Using the appropriate form ensures we have all the information necessary to process your report quickly. Please provide all information requested in the form. If you submit a report with incomplete information, we will follow up with you to request that information. Please note that this will result in a delay in processing your report.

If you have trouble using the form, try updating your browser or switching to another browser.

Copyright policy

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged

infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notification.

When Should I File a Counter-notice?

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and

4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Twitter media policy

Twitter Rules: Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images.

Twitter Rules: You may not use hateful images or symbols in your profile image or profile header.

Rationale

Some forms of graphic violence, adult content, or hateful imagery may be permitted in Tweets when they are marked as [sensitive media](#). However, you may not include this type of content in live video, or in profile or header images.

While we want people to feel free to share media that reflects their creativity or individuality, or to show what's happening in the world, we will take action when it crosses the line into abuse towards a person, group, or protected category.

Finally, we may sometimes require you to remove media containing excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about how to make such a request [here](#).

When this applies

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. Some examples include, but are not limited to, depictions of:

- the moment at which someone dies
- gruesome crime or accident scenes
- bodily harm, torture, dismemberment, or mutilation

We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity (including close-ups of genitals, buttocks, or breasts)
 - Please note that exceptions may be made for artistic, medical, health, or educational content. Breastfeeding content does not need to be marked as sensitive.
- simulating a sexual act
- intercourse or any sexual act (may involve humans, humanoid animals, cartoons, or anime)

If you see media you think may be in violation of our policy, report it. You can read about how to report sensitive media [here](#).

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, or ethnicity/national origin.

Some examples of hateful imagery include, but are not limited to:

- Symbols historically associated with hate groups (for example, the Nazi swastika)
- Images depicting others as less than human or altered to include hateful symbols
- Altered image references to a mass murder that targeted a protected category

Do I need to be the target of this content in order for it to be a violation of the Twitter Rules?
No, we review both first-person and bystander reports of such content.

Consequences

When adult content, graphic violence, or hateful imagery appears in Tweets, we may place this content behind an interstitial advising viewers to be aware that they will see sensitive media if they click through. This allows us to identify potentially sensitive content that some people may not wish to see. Learn more about [how to control whether you see sensitive media](#).

However, if such content appears in live video, header, or profile images, it will violate our media policy. The consequences for violating this policy depend on the severity of the violation and the person's previous record of violations.

The first time someone violates this policy they will be required to remove the imagery and may be asked to verify their account ownership. Subsequent violations could result in permanent suspension. Learn more about [our range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

How to contact Twitter about media concerning a deceased family member

In order to respect the wishes of loved ones, Twitter will remove imagery of deceased individuals in certain circumstances. Immediate family members and other authorized individuals may request the removal of images or video of deceased individuals, from when critical injury occurs to the moments before or after death, by sending a request to Twitter Inc. via our [privacy form](#). When reviewing such media removal requests, Twitter considers public interest factors such as the newsworthiness of the content and may not be able to honor every request.

Note: Find information about contacting Twitter to [deactivate the account of a deceased or incapacitated user](#).

How to mark media as sensitive links to Media policy

Twitter trends FAQs

Where can I find trends?

Trends are available on the Twitter for iOS app, the Twitter for Android app, and [twitter.com](#).

On Twitter's mobile apps, you can find trends listed under the trends section of the **Explore** tab. When signed in to [twitter.com](#) on a desktop or laptop computer, trends are listed in many places, including the Home timeline, Notifications, search results, and profile pages.

How are trends determined?

Trends are determined by an algorithm and, by default, are tailored for you based on who you follow, your interests, and your location. This algorithm identifies topics that are popular now, rather than topics that have been popular for a while or on a daily basis, to help you discover the hottest emerging topics of discussion on Twitter.

You can choose to see trends that are not tailored for you by selecting a specific trends location on [twitter.com](#), iOS, or Android (instructions below). Location trends identify popular topics among people in a specific geographic location.

Note: The number of Tweets that are related to the trends is just one of the factors the algorithm looks at when ranking and determining trends. Algorithmically, trends and hashtags are grouped together if they are related to the same topic. For instance, #MondayMotivation and #MotivationMonday may both be represented by #MondayMotivation.

What does the # sign mean?

You may notice that some trends have # sign before the word or phrase. This is called a **hashtag** and is included specifically in Tweets to mark them as relating to a topic, so that people can follow the conversation in search.

Will I see world and local events in my trends?

Yes, when signed into your account on the twitter.com, iOS, or Android, trends will be tailored for you based on your location and who you follow. There will be many world and local news events and conversations that will appear in your trends regardless of your personalization.

Can I see trends for a specific location?

Yes, you can view trends for a specific location on twitter.com, iOS, or Android.

View trends via twitter.com:

To view trends for a specific location, click **Change** in your trends box. You can choose to **Keep tailored trends** based on your location and who you follow, or you can select **Change** to choose a nearby location or search locations.

Adjust your trends settings via Twitter for iOS: 0.

In the top menu, tap your profile icon.

17. Tap **Settings and privacy**, then tap **Content preferences**.
18. Under **Explore**, tap **Trends**.

T. Drag the slider next to **Trends for you** to turn it off and receive personalized trends.

V. When the feature is disabled, you can change your location by tapping **Change location**. In the search box, type your desired trends location or choose from the list.

Adjust your trends settings via Twitter for Android:

O. In the top menu, you will either see a navigation menu icon or your profile

icon. Tap whichever icon you have, and select **Settings and privacy**.

17. Tap **Content preferences**.

18. Under **Explore**, tap **Trends**.

T. Drag the slider next to **Trends for you** to turn on and receive personalized trends.

V. When the feature is disabled, you can change your location by tapping **Change location**. In the search box, type your desired trends location or choose from the list.

Note: If you don't find your preferred country or city, it means we're not receiving enough Tweets from that geographical area to create a list. You can look up local Tweets on any topic by using [advanced Twitter search operators](#).

What happens when I click or tap on a trend?

Clicking or tapping any of the trends takes you to the Twitter search results for that trend. You'll see all Tweets including that phrase or hashtag. To see what people are saying about a previous trend, perform a [search for that keyword](#).

How can I participate in a trend?

Simply post a Tweet including the exact word or phrase as it appears in the trends list (with the hashtag, if you see one). Due to the large number of people Tweeting about these specific trends, you may not always be able to find your particular Tweet in search, but your followers will always see your Tweets.

Note: Twitter also filters searches for quality. To understand what types of behavior could cause your Tweets to be filtered out, read about our [search rules and restrictions](#).

Are there rules for trends?

Yes. We want trends to promote healthy discussions on Twitter. This means that at times, we may prevent certain content from trending. These include trends that:

Contain profanity or adult/graphic references.

Incite hate on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease.

Violate the [Twitter Rules](#).

In some cases, we may also consider the newsworthiness of the content, or if it is in the public interest when evaluating potential violations. In these cases, the content might continue to trend on our platform.

Note that even if we prevent the hashtag or the content from appearing on the list of trending topics, you may still be able to access conversations around that content on Twitter.

Do trends have descriptions?

Below some trends, you may see the approximate number of Tweets associated with the trend, or personalized context such as who in your network is Tweeting about it.

Additionally, when a Moment is found to be relevant to a trend, we may use the title of the Moment as the trend's description.

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Child sexual exploitation policy

Twitter removes content promoting child sexual exploitation

We do not tolerate child sexual exploitation on Twitter. When we are made aware of links to images of or content promoting child sexual exploitation they will be removed from the site without further notice and reported to The National Center for Missing & Exploited Children ("NCMEC"); we permanently suspend accounts promoting or containing updates with links to child sexual exploitation.

Reporting content promoting child sexual exploitation

If you think you've found a Twitter account distributing or promoting child sexual exploitation while using Twitter, please click [here](#) to report it. Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. To find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with the National Center for Missing and Exploited Children [here](#), or your local law enforcement. If you believe that the content was hosted on a website outside of the United States, you can report it [here](#) on the International Association of Internet Hotlines website.

NOTE: Please do NOT tweet, retweet or repost child sexual exploitation for any reason. Click [here](#) to report it to us immediately and we will take steps to remove it.

Law Enforcement Agencies

If you are a member of a law enforcement agency seeking information regarding Twitter accounts promoting or distributing child sexual exploitation, please see our [Guidelines for Law Enforcement](#) for information on our process; we also suggest you contact the National Center for Missing and Exploited Children (NCMEC) [here](#), regardless of jurisdiction, to determine whether or not information on those accounts is already on file.

For general law enforcement inquiries, please see our [Guidelines for Law Enforcement](#) for more information.

Hateful conduct policy

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You

also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identity with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you".

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., "I hope you get cancer and die."

Wishing for someone to fall victim to a serious accident, e.g., "I wish that you would get run over by a car next time you run your mouth."

Saying that a group of individuals deserve serious physical injury, e.g., "If this group of protesters don't shut up, they deserve to be shot."

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are

more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is

not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to delete the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#). If someone believes their account was suspended in error, they can

[submit an appeal.](#)

Report private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Read about our [policy regarding private information posted on Twitter](#).

Do I need to have a Twitter account to make a report?

No, you do not need to have a Twitter account to file a report that your private and confidential information has been posted on Twitter.

Can I report unauthorized posting of Copyrighted material?

If your copyrighted material has been Tweeted without your permission, please review our [Copyright and DMCA policy](#).

What can I do if my information is posted on other websites?

We cannot take down any information that is posted on a website other than Twitter. If your information is posted on another website, please contact the help center or support staff for the website where your information is posted.

How do I report that my private information is posted on Twitter?

You can report that your private information has been posted on Twitter directly from the offending Tweet(s). Learn [how to report individual Tweets](#).

File a [report that someone has posted your private information](#).

What happens when Twitter receives a valid report?

When we receive a complete and valid report that private information has been posted on Twitter, we'll investigate the account and Tweets reported. We will review where, if anywhere, the information has been made publicly available before taking action on the account or Tweets. If the information you reported was previously posted elsewhere on the Internet, it is not a violation of our policy and we will not take action.

Learn more about [protecting your private information on Twitter and other websites](#).

Read our [policy regarding private information posted on Twitter](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter impersonation policy.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to Twitter's [parody, commentary, and fan account policy](#) for more information about these types of accounts.

Find more information about [reporting impersonation on Twitter](#).

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation rules

Updated November 3, 2017

This page is primarily intended for **developers**.

For Twitter users: You are ultimately responsible for the actions taken with your account, or by applications associated with your account. Before authorizing a third-party application to access or use your account, make sure you've thoroughly investigated the application and understand what it will do. If automated activity on your account violates the [Twitter Rules](#) or these automation rules, Twitter may take action on your

account, including [filtering your Tweets from search results](#) or suspending your account.

For more information on third-party applications, please see our article on [connecting and revoking third-party applications](#).

If you are a developer and have questions about these Automation Rules, please reach out via <https://twittercommunity.com/c/rules-and-policies>.

If your application has been suspended or restricted from performing write actions, you may submit an appeal at </forms/platform>.

I. Ground Rules Do!

Build solutions that automatically broadcast helpful information in Tweets.

Run creative campaigns that auto-reply to users who engage with your content.

Build solutions that automatically respond to users in Direct Messages.

Try new things that help people (and comply with our rules).

Make sure your application provides a good user experience and performs well — and confirm that remains the case over time.

Don't!

Violate these or other policies. Be extra mindful of our rules about abuse and user privacy.

Abuse the Twitter API or attempt to circumvent rate limits.

Use non-API-based forms of automation, such as scripting the Twitter website. The use of these techniques may result in the permanent suspension of your account.

Spam or bother users, or otherwise send them unsolicited messages.

A. The Twitter Rules and the Developer Agreement and Policy

As with all activity on Twitter, automated activity is subject to the [Twitter Rules](#) and, if you're a developer using the Twitter API, the [Developer Agreement and Policy](#).

You should carefully review these policies to ensure that your automated activity is compliant. Automated applications or activities that violate these policies, or that facilitate or induce users to violate them, may be subject to enforcement action,

potentially including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these policies.

Although all aspects of the Twitter Rules and the Developer Agreement and Policy apply to automated activity, you should keep the following rules top of mind:

Spamming: You may not send automated Tweets or Direct Messages that are spam, or otherwise engage in spamming activity. Some examples of spammy behavior to avoid with automation include:

Trending topics: You may not automatically post about trending topics on Twitter, or use automation to attempt to influence or manipulate trending topics.

Multiple posts/accounts: You may not post duplicative or substantially similar Tweets on one account or over multiple accounts you operate.

Duplicate accounts: You may not create and/or automate multiple accounts for duplicative or substantially similar use cases.

However, automating multiple accounts for related but non-duplicative use cases is permitted. For example, you may automate separate accounts to Tweet when the Hubble Space Telescope passes over different cities, such as [San Francisco](#) or [Hong Kong](#).

Misleading links: You may not send automated Tweets or Direct Messages containing links that are misleading, including links that maliciously or deceptively redirect through landing pages or ad pages before displaying the final content.

Sensitive media: Automated Tweets and Direct Messages must comply with the [Twitter media policy](#), and you should mark your account as potentially sensitive if you intend to post graphic, pornographic, or potentially sensitive media.

Abusive behavior: You may not engage in any automated activity that encourages, promotes, or incites abuse, violence, hateful conduct, or harassment, on or off Twitter.

Private information: You may not post private or confidential information about a person without their prior express authorization.

B. Other Ground Rules for Automated Activity

In addition to the policies above, the following ground rules apply to all automated activity on Twitter:

Don't surprise or mislead users: Automated activity should honor users' expectations. Ask for the user's permission before taking an action if you aren't sure.

Mature content or profanity: Don't Direct Message, mention, or reply to users with potentially sensitive content (including profanity), unless they've clearly indicated an intent to receive it in advance.

Be thoughtful about the information you request or exchange on Twitter

Tweets: Don't ask users to send you personal or private information via a public Tweet. If you need additional personal or private information from a user to provide them with customer service (or other similar use cases), you should ask the user to share such information by Direct Message or another private channel. You might even consider adding a [Direct Message deep link](#) to your Tweet.

Direct Messages: You should only ask users for the minimum amount of information you need to provide them with service. If you need to request or exchange particularly sensitive information (such as credit card information), you should consider directing users to your website or other appropriate channel to do so.

II. Activity-Specific Rules

The activity-specific rules in this section apply to taking specific automated actions on Twitter. Please read these rules carefully, as they outline both permitted and prohibited use cases of automation.

Automated applications or activities that violate these rules, or that facilitate or induce users to violate them, may be subject to enforcement action, including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these rules. As a reminder, you should also carefully review the spam guidelines in the [Twitter Rules](#) to avoid having activities performed by you, your app, or other users through your app or service flagged as spam.

A. Automated Actions Through Another User's Account

Twitter users may authorize your app or service to [access their Twitter account through OAuth](#). A user authorizing your app or service to access their Twitter account through OAuth does not by itself constitute sufficient consent to take automated actions through that user's account.

You may only take automated actions through another Twitter user's account if you:

clearly describe to the user the types of automated actions that will occur; receive express consent from the user to take those automated actions; and immediately honor a user's request to opt-out of further automated actions.

If you substantially change the purpose or functionality of your app or service, you must re-obtain express consent from the user to take automated action through their account before doing so.

These requirements apply to any automated action taken through another Twitter user's account, including posting Tweets, sending Direct Messages, deleting Tweets or Direct Messages, or following/unfollowing other accounts. For applications that offer users the ability to delete Tweets in a bulk or automated manner, you must also clearly state that Tweets are not recoverable once deleted.

B. Automated Tweets

1. Posting automated Tweets

Automated Tweets that cross-post outside information:

You may post automated Tweets based on sources of outside information — such as an RSS feed, weather data, etc. — as long as you are sufficiently authorized to publish such information.

Other automated Tweets (excluding mentions or replies):

Provided you comply with all other rules, you may post automated Tweets for entertainment, informational, or novelty

purposes. As a reminder, accounts posting duplicative, spammy, or otherwise prohibited content may be subject to suspension.

2. Posting automated mentions and replies

The reply and mention functions are intended to make communication between Twitter users easier. Automating these actions to reach many users on an unsolicited basis is an abuse of the feature, and is not permitted. For example, sending automated replies to Tweets based on keyword searches alone is not permitted. Spamy or duplicative use of mentions and replies may result in enforcement action, such as the removal of your Tweets from Search or the suspension of your app or account.

However, you may send automated replies or mentions to Twitter users so long as:

in advance of sending the automated reply, the recipient or mentioned user(s) have requested or have clearly indicated an intent on Twitter to be contacted by you (i.e. opted in), for example by replying to a Tweet from your account, or by sending you a Direct Message;

you provide a clear and easy way for such users to opt-out of receiving automated replies and mentions, and promptly honor all such opt-out requests;

you only send one automated reply or mention per user interaction; and

the automated reply or mention is a reply to the user's original Tweet (if your campaign is based on users posting a reply to your Tweet).

Opt-in techniques and indications of user intent take many different forms, depending on the specifics of your use case and implementation. Some examples include:

A Tweet from your account that clearly indicates that a user taking a specific action on that Tweet (such as Retweeting it) will opt the user into receiving an automated response.

A mention of your account by the user in a manner suggesting the user clearly wishes or intends to receive a response. If you want to run an auto-reply campaign with a campaign- or use-case-specific hashtag, users should also mention you in their Tweets.

Note that a user following your account is not on its own a sufficient indication of user intent to receive an automated response.

: Advertisers, publishers, and brands using auto-response campaigns must request approval from Twitter and may be subject to additional rules. Please reach out to your account or partner manager for assistance. If you do not have a partner manager but are interested in managed client services, you can learn more [here](#).

Additionally, we recommend that any accounts that will communicate with

users via automated mentions or replies:

Note

1. Appropriately filter responses based on potentially sensitive language in user handles,

display names, and Tweet text, as well as potentially sensitive media;

2. Check that the Tweet you are mentioning or replying to still exists (for example, using the statuses/lookup endpoint on the Twitter API).

C. Automated Direct Messages

Sending automated Direct Messages to users

You may not send unsolicited Direct Messages in a bulk or automated manner, and should be thoughtful about the frequency with which you contact users via Direct Message.

You may send automated Direct Messages to users so long as:

in advance of sending the Direct Message, the recipient(s) have requested or have clearly indicated an intent on Twitter to be contacted by you via Direct Message, for example by sending you a Direct Message; **and**

you provide a clear and easy way for such users to opt-out of receiving automated Direct Messages, and promptly honor all such opt-out requests.

The fact that a user is technically able to receive a Direct Message from you (e.g. because the user follows you, has enabled the ability to receive Direct Messages from any

account, or because the user is in a pre-existing Direct Message conversation with you) does not necessarily mean they have requested or expect to receive automated Direct Messages from you.

Interacting with users via Direct Message

Per the Ground Rules, remember to be thoughtful about the amount and type of information you request or exchange with users via Direct Messages. If you will be asking a user to provide personal or private information via an automated Direct Message, you must clearly explain how you will use the information you're collecting. Consider including a link to your privacy policy in your Direct Message to the user, as well as in your Twitter profile bio.

Don't publicly share information received in a Direct Message conversation with a user without first obtaining explicit consent from the user. For example, if a user asks you via Direct Message about a purchase they made from you, you may not mention the user in a public Tweet that includes information about their purchase unless you have the user's explicit consent to do so.

After a user-initiated interaction ends, don't send additional follow-up Direct Messages or mention users in a Tweet unless you get permission from the user.

D. Automated actions you take on Tweets or accounts

Automated likes: You may not like Tweets in an automated manner.

Automated Retweets: Provided you comply with all other rules, you may Retweet or Quote Tweet in an automated manner for entertainment, informational, or novelty purposes. Automated Retweets often lead to negative user experiences, and bulk, aggressive, or spammy Retweeting is a violation of the [Twitter Rules](#).

Automated following/unfollowing: You may not follow or unfollow Twitter accounts in a bulk, aggressive, or indiscriminate manner. Aggressive following is a violation of the [Twitter Rules](#). Please also review our [following rules and best practices](#) to ensure you are in compliance. Note that applications that claim to get users more followers are also prohibited under the [Twitter Rules](#).

Automated adding to lists or collections: You may not add Twitter users to lists or add Tweets to collections in a bulk or indiscriminate manner. Adding a large number of unrelated users to lists is a violation of the [Twitter Rules](#).

About search rules and restrictions

This article details the rules for using Twitter search, and explains reasons why content may not show up in Twitter search.

Why doesn't all content show up in Twitter search?

In order to keep your search results relevant, Twitter filters search results for quality Tweets and accounts. Material that jeopardizes search quality or creates a bad search experience for other people may be automatically removed from Twitter search. Read more about this [here](#).

General rules for using Twitter search

The [Twitter Rules](#) explain what behaviors are permitted on Twitter. In addition to these rules, we've included some tips below to keep your content relevant (and your Tweets in Twitter search).

Please do not:

- Repeatedly post duplicate or near-duplicate content (links or Tweets).
- Abuse trending topics or hashtags (topic words with a # sign).
- Send automated Tweets or replies.
- Use bots or applications to post similar messages based on keywords.
- Post similar messages over multiple accounts.
- Aggressively follow and unfollow people.

Twitter may automatically remove accounts engaging in these behaviors from search (or even suspend in some cases) in order to ensure the best experience for everyone.

Search results filtering

To learn how to filter your search results, including how to disable or enable safe search, read our [Using Twitter search](#) article.

Search results filtering in third-party applications

Accounts may be filtered from search for posting similar messages over several accounts. Your account may be filtered from search if you are logged into a third-party application that is updating many accounts with similar or duplicate content.

Rules for participating in contests on Twitter

Businesses sometimes host contests and offer prizes to Retweet an update, following a particular account, or posting Tweets with a specific hashtag topic or reply. If you're running a contest, check out this [article](#) for guidelines for contests on Twitter.

If you repeatedly post near-duplicate updates or duplicate links in order to enter contests, you may be filtered from search. Please keep in mind that you are responsible for the updates that you post to your Twitter account.

What to do if you are missing from search

Please see this [article](#) for more information related to this issue.

If you are being filtered from search, rest assured that your followers will still receive your updates. If you post a reply or mention to a specific person, it will be delivered to that person.

Having trouble?

You can learn more about searching by keyword, location, and trending topics in our [Using Twitter search](#) help article.

The Twitter Rules (20180424*)

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. These limitations are set forth in the Twitter Rules below.

The Twitter Rules (along with all incorporated policies), [Privacy Policy](#), and [Terms of Service](#) collectively make up the "Twitter User Agreement" that governs a user's access to and use of Twitter's services.

All individuals accessing or using Twitter's services must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in Twitter taking one or more of the following enforcement actions:

- requiring you to delete prohibited content before you can again create new posts and interact with other Twitter users;
- temporarily limiting your ability to create posts or interact with other Twitter users;
- asking you to verify account ownership with a phone number or email address; or
- permanently suspending your account(s).

If you attempt to evade a permanent suspension by creating new accounts, we will suspend your new accounts.

Please note that we may need to change these Rules from time to time and reserve the right to do so. The most current version is always available at: <https://twitter.com/rules>.

The policies set forth in these Twitter Rules govern organic content on our platform. To learn more about the rules which govern ads and promoted content, please review our [Ads policies](#).

Content Boundaries and Use of Twitter

Intellectual property

Trademark: We reserve the right to suspend accounts or take other appropriate action when someone's brand or trademark, including business name and/or logo, is used in a manner that may mislead or confuse others about your brand affiliation. Read more about our [trademark policy](#) and [how to report a violation](#).

Copyright: We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in our [Terms of Service](#). Read more about our [copyright policy](#).

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Learn more about our [media policy](#).

Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images. Additionally, Twitter may sometimes require you to remove excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about [how to make such a request](#), and [how to mark your media as sensitive](#).

Unlawful use

You may not use our service for any unlawful purposes or in furtherance of illegal activities. By using Twitter, you agree to comply with all applicable laws governing your online conduct and content.

Trends

At times, we may prevent certain content from trending. This includes content that violates the Twitter Rules, as well as content that may attempt to manipulate trends. Read more about [what we allow and do not allow to trend](#).

Misuse of Twitter badges

You may not use badges, including but not limited to the “promoted” or “verified” Twitter badges, unless provided by Twitter. Accounts using unauthorized badges as part of their profile photos, header photos, display names, or in any way that falsely implies affiliation with Twitter or authorization from Twitter to display these badges, may be suspended.

Misuse of usernames

Selling usernames: You may not buy or sell Twitter usernames.

Username squatting: You may not engage in username squatting. Some of the factors we take into consideration when determining whether conduct is username squatting include:

- the number of accounts created;
- the creation of accounts for the purpose of preventing others from using those account names;
- the creation of accounts for the purpose of selling those accounts; and
- the use of third-party content feeds to update and maintain accounts under the names of those third parties.

Please note that Twitter may also remove accounts that are inactive for more than six months. Learn more about [username squatting](#).

Abusive Behavior

We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we prohibit behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Context matters when evaluating for abusive behavior and determining appropriate enforcement actions. Factors we may take into consideration include, but are not limited to whether:

- the behavior is targeted at an individual or group of people;
- the report has been filed by the target of the abuse or a bystander;
- the behavior is newsworthy and in the legitimate public interest.

Violence and physical harm

Violence: You may not make specific threats of violence or wish for the serious physical harm, death, or disease of an individual or group of people. This includes, but is not limited to, threatening or promoting terrorism. You also may not affiliate with organizations that — whether by their own statements or activity both on and off the platform — use or promote violence against civilians to further their causes. We will begin enforcing this rule around affiliation with such organizations on December 18, 2017.

Suicide or self-harm: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Child sexual exploitation: You may not promote child sexual exploitation. Learn more about our zero-tolerance [child sexual exploitation policy](#).

Abuse and hateful conduct

Abuse: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Unwanted sexual advances: You may not direct abuse at someone by sending unwanted sexual content, objectifying them in a sexually explicit manner, or otherwise engaging in sexual misconduct.

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. Read more about our [hateful conduct policy](#).

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category. We will begin enforcing this rule on December 18, 2017.

Private information and intimate media

Private information: You may not publish or post other people's private information without their express authorization and permission. Definitions of private information may vary depending on local laws. Read more about our [private information policy](#).

Intimate media: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. Read more about intimate media on Twitter.

Threats to expose / hack: You may not threaten to expose someone's private information or intimate media. You also may not threaten to hack or break into someone's digital information.

Impersonation

You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. While you may maintain parody, fan, commentary, or newsfeed accounts, you may not do so if the intent of the account is to engage in spamming or abusive behavior. Read more about our [impersonation policy](#).

Spam and Security

We strive to protect people on Twitter from technical abuse and spam.

To promote a stable and secure environment on Twitter, you may not do, or attempt to do, any of the following while accessing or using Twitter:

- Access, tamper with, or use non-public areas of Twitter, Twitter's computer systems, or the technical delivery systems of Twitter's providers (except as expressly permitted by the Twitter Bug Bounty program).
- Probe, scan, or test the vulnerability of any system or network, or breach or circumvent any security or authentication measures (except as expressly permitted by the Twitter Bug Bounty program).
- Access or search, or attempt to access or search, Twitter by any means (automated or otherwise) other than through our currently available, published interfaces that are provided by Twitter (and only pursuant to the applicable terms and conditions), unless you have been specifically allowed to do so in a separate agreement with Twitter. Note that crawling Twitter is permissible if done in accordance with the provisions of the robots.txt file; however, scraping Twitter without our prior consent is expressly prohibited.
- Forge any TCP/IP packet header or any part of the header information in any email or posting, or in any way use Twitter to send altered, deceptive, or false source-identifying information.
- Interfere with or disrupt the access of any user, host or network, including, without limitation, sending a virus, overloading, flooding, spamming, mail-bombing Twitter's services, or by scripting the creation of content in such a manner as to interfere with or create an undue burden on Twitter.

Any accounts engaging in the following activities may be temporarily locked or subject to permanent suspension:

- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use Twitter's services for the purpose of spamming anyone. Spam is generally defined on Twitter as bulk or aggressive activity that attempts to manipulate or disrupt Twitter or the experience of users on Twitter to drive traffic or attention to unrelated accounts, products, services, or initiatives. Some of the factors that we take into account when determining what conduct is considered to be spamming include:
 - if you have followed and/or unfollowed a large number of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if your Tweets or Direct Messages consist mainly of links shared without commentary;
 - if a large number of people have blocked you in response to high volumes of untargeted, unsolicited, or duplicative content or engagements from your account;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicative or substantially similar content, replies, or mentions over multiple accounts or multiple duplicate updates on one account, or create duplicate or substantially similar accounts;
 - if you post multiple updates to a trending or popular topic with an intent to subvert or manipulate the topic to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you send large numbers of unsolicited replies or mentions;
 - if you add users to lists in a bulk or aggressive manner;
 - if you are randomly or aggressively engaging with Tweets (e.g., likes, Retweets, etc.) or users (e.g., following, adding to lists or Moments, etc.) to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you repeatedly post other people's account information as your own (e.g., bio, Tweets, profile URL, etc.);
 - if you post misleading, deceptive, or malicious links (e.g., affiliate links, links to malware/clickjacking pages, etc.);
 - if you create fake accounts, account interactions, or impressions;
 - if you sell, purchase, or attempt to artificially inflate account interactions (such as followers, Retweets, likes, etc.); and
 - if you use or promote third-party services or apps that claim to get you more followers, Retweets, or likes (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers or engagements to your account or Tweets).

Please see our support articles on [following rules and best practices](#) and [automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts may be permanently suspended.

Content Visibility

Accounts under investigation or which have been detected as sharing content in violation of these Rules may have their account or Tweet visibility limited in various parts of Twitter, including search. To learn more about situations in which content may be restricted on Twitter, please see our support article on [search rules and restrictions](#).

Trademark policy (20160630)

Updated June 30, 2016

What is a Trademark Policy Violation on Twitter?

Using another's trademark in a manner that may mislead or confuse others about your brand affiliation may be a violation of our trademark policy.

How Does Twitter Respond to Trademark Reports?

We review each report and do the following:

- If we determine that there is a violation of our trademark policy, we may suspend the account. In some instances, we may give the account holder an opportunity to comply with our policies before suspending the account. For more information, see our [Parody, commentary, and fan account policy](#).
- We may also provide the account holder with your name and information included in the copy of the report.
- If you would like to report a trademark violation in a [Twitter Ad](#), please visit our [Trademark Policy for Promoted Ads](#) to learn more.

Please note that Twitter only investigates requests that are submitted by the trademark holder or by an authorized representative of the trademark holder.

What is not a Trademark Policy Violation?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples include:

- Using a trademark in a way that is outside the scope of the trademark registration (e.g. territory, or goods and services identified in the registration).
- Nominative and other fair uses of trademarks are protected uses under our trademark policy, so long as the account is clearly distinguished from the trademark owner. This includes use by resellers in certain regions and accounts engaging in parody, commentary, or news. For more information, see our [Parody, commentary, and fan account policy](#).

How do I Report a Trademark Policy Violation?

Submit trademark reports through our [Trademark Issue support form](#). Using the appropriate form ensures we have all the information necessary to process your report quickly. Please provide all information requested in the form. If you submit a report with incomplete information, we will follow up with you to request that information. Please note that this will result in a delay in processing your report.

If you have trouble using the form, try updating your browser or switching to another browser.

Copyright policy (20180601*)

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged

infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notification.

When Should I File a Counter-notice?

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and

4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Twitter media policy (20181024*)

Twitter Rules: Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images.

Twitter Rules: You may not use hateful images or symbols in your profile image or profile header.

Rationale

Some forms of graphic violence, adult content, or hateful imagery may be permitted in Tweets when they are marked as [sensitive media](#). However, you may not include this type of content in live video, or in profile or header images.

While we want people to feel free to share media that reflects their creativity or individuality, or to show what's happening in the world, we will take action when it crosses the line into abuse towards a person, group, or protected category.

Finally, we may sometimes require you to remove media containing excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about how to make such a request [here](#).

When this applies

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. Some examples include, but are not limited to, depictions of:

- the moment at which someone dies
- gruesome crime or accident scenes
- bodily harm, torture, dismemberment, or mutilation

We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity (including close-ups of genitals, buttocks, or breasts)
 - Please note that exceptions may be made for artistic, medical, health, or educational content. Breastfeeding content does not need to be marked as sensitive.
- simulating a sexual act
- intercourse or any sexual act (may involve humans, humanoid animals, cartoons, or anime)

If you see media you think may be in violation of our policy, report it. You can read about how to report sensitive media [here](#).

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, or ethnicity/national origin.

Some examples of hateful imagery include, but are not limited to:

- Symbols historically associated with hate groups (for example, the Nazi swastika)
- Images depicting others as less than human or altered to include hateful symbols
- Altered image references to a mass murder that targeted a protected category

Do I need to be the target of this content in order for it to be a violation of the Twitter Rules?
No, we review both first-person and bystander reports of such content.

Consequences

When adult content, graphic violence, or hateful imagery appears in Tweets, we may place this content behind an interstitial advising viewers to be aware that they will see sensitive media if they click through. This allows us to identify potentially sensitive content that some people may not wish to see. Learn more about [how to control whether you see sensitive media](#).

However, if such content appears in live video, header, or profile images, it will violate our media policy. The consequences for violating this policy depend on the severity of the violation and the person's previous record of violations.

The first time someone violates this policy they will be required to remove the imagery and may be asked to verify their account ownership. Subsequent violations could result in permanent suspension. Learn more about [our range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

How to contact Twitter about media concerning a deceased family member

In order to respect the wishes of loved ones, Twitter will remove imagery of deceased individuals in certain circumstances. Immediate family members and other authorized individuals may request the removal of images or video of deceased individuals, from when critical injury occurs to the moments before or after death, by sending a request to Twitter Inc. via our [privacy form](#). When reviewing such media removal requests, Twitter considers public interest factors such as the newsworthiness of the content and may not be able to honor every request.

Note: Find information about contacting Twitter to [deactivate the account of a deceased or incapacitated user](#).

How to mark media as sensitive links to Media policy

Twitter trends FAQs (20180722*)

Where can I find trends?

Trends are available on the Twitter for iOS app, the Twitter for Android app, and twitter.com.

On Twitter's mobile apps, you can find trends listed under the trends section of the **Explore** tab. When signed in to twitter.com on a desktop or laptop computer, trends are listed in many places, including the Home timeline, Notifications, search results, and profile pages.

How are trends determined?

Trends are determined by an algorithm and, by default, are tailored for you based on who you follow, your interests, and your location. This algorithm identifies topics that are popular now, rather than topics that have been popular for a while or on a daily basis, to help you discover the hottest emerging topics of discussion on Twitter.

You can choose to see trends that are not tailored for you by selecting a specific trends location on twitter.com, iOS, or Android (instructions below). Location trends identify popular topics among people in a specific geographic location.

Note: The number of Tweets that are related to the trends is just one of the factors the algorithm looks at when ranking and determining trends. Algorithmically, trends and hashtags are grouped together if they are related to the same topic. For instance, #MondayMotivation and #MotivationMonday may both be represented by #MondayMotivation.

What does the # sign mean?

You may notice that some trends have # sign before the word or phrase. This is called a **hashtag** and is included specifically in Tweets to mark them as relating to a topic, so that people can follow the conversation in search.

Will I see world and local events in my trends?

Yes, when signed into your account on the twitter.com, iOS, or Android, trends will be tailored for you based on your location and who you follow. There will be many world and local news events and conversations that will appear in your trends regardless of your personalization.

Can I see trends for a specific location?

Yes, you can view trends for a specific location on twitter.com, iOS, or Android.

View trends via twitter.com:

To view trends for a specific location, click **Change** in your trends box. You can choose to **Keep tailored trends** based on your location and who you follow, or you can select **Change** to choose a nearby location or search locations.

Adjust your trends settings via Twitter for iOS: 0.

In the top menu, tap your profile icon.

17. Tap **Settings and privacy**, then tap **Content preferences**.
18. Under **Explore**, tap **Trends**.

T. Drag the slider next to **Trends for you** to turn on and receive personalized trends.

V. When the feature is disabled, you can change your location by tapping **Change location**. In the search box, type your desired trends location or choose from the list.

Adjust your trends settings via Twitter for Android:

O. In the top menu, you will either see a navigation menu icon or your profile

icon. Tap whichever icon you have, and select **Settings and privacy**.

17. Tap **Content preferences**.

18. Under **Explore**, tap **Trends**.

T. Drag the slider next to **Trends for you** to turn on and receive personalized trends.

V. When the feature is disabled, you can change your location by tapping **Change location**. In the search box, type your desired trends location or choose from the list.

Note: If you don't find your preferred country or city, it means we're not receiving enough Tweets from that geographical area to create a list. You can look up local Tweets on any topic by using [advanced Twitter search operators](#).

What happens when I click or tap on a trend?

Clicking or tapping any of the trends takes you to the Twitter search results for that trend. You'll see all Tweets including that phrase or hashtag. To see what people are saying about a previous trend, perform a [search for that keyword](#).

How can I participate in a trend?

Simply post a Tweet including the exact word or phrase as it appears in the trends list (with the hashtag, if you see one). Due to the large number of people Tweeting about these specific trends, you may not always be able to find your particular Tweet in search, but your followers will always see your Tweets.

Note: Twitter also filters searches for quality. To understand what types of behavior could cause your Tweets to be filtered out, read about our [search rules and restrictions](#).

Are there rules for trends?

Yes. We want trends to promote healthy discussions on Twitter. This means that at times, we may prevent certain content from trending. These include trends that:

Contain profanity or adult/graphic references.

Incite hate on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease.

Violate the [Twitter Rules](#).

In some cases, we may also consider the newsworthiness of the content, or if it is in the public interest when evaluating potential violations. In these cases, the content might continue to trend on our platform.

Note that even if we prevent the hashtag or the content from appearing on the list of trending topics, you may still be able to access conversations around that content on Twitter.

Do trends have descriptions?

Below some trends, you may see the approximate number of Tweets associated with the trend, or personalized context such as who in your network is Tweeting about it.

Additionally, when a Moment is found to be relevant to a trend, we may use the title of the Moment as the trend's description.

Username squatting policy (20130817*)

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Child sexual exploitation policy (20171025*)

Twitter removes content promoting child sexual exploitation

We do not tolerate child sexual exploitation on Twitter. When we are made aware of links to images of or content promoting child sexual exploitation they will be removed from the site without further notice and reported to The National Center for Missing & Exploited Children ("NCMEC"); we permanently suspend accounts promoting or containing updates with links to child sexual exploitation.

Reporting content promoting child sexual exploitation

If you think you've found a Twitter account distributing or promoting child sexual exploitation while using Twitter, please click [here](#) to report it. Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. To find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with the National Center for Missing and Exploited Children [here](#), or your local law enforcement. If you believe that the content was hosted on a website outside of the United States, you can report it [here](#) on the International Association of Internet Hotlines website.

NOTE: Please do NOT tweet, retweet or repost child sexual exploitation for any reason. Click [here](#) to report it to us immediately and we will take steps to remove it.

Law Enforcement Agencies

If you are a member of a law enforcement agency seeking information regarding Twitter accounts promoting or distributing child sexual exploitation, please see our [Guidelines for Law Enforcement](#) for information on our process; we also suggest you contact the National Center for Missing and Exploited Children (NCMEC) [here](#), regardless of jurisdiction, to determine whether or not information on those accounts is already on file.

For general law enforcement inquiries, please see our [Guidelines for Law Enforcement](#) for more information.

Hateful conduct policy (20181113*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We

also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., “I will kill you”.

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., “I hope you get cancer and die.”

Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”

Saying that a group of individuals deserve serious physical injury, e.g., “If this group of protesters don’t shut up, they deserve to be shot.”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to delete the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#). If someone believes their account was suspended in error, they can [submit an appeal](#).

Report private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Read about our [policy regarding private information posted on Twitter](#).

Do I need to have a Twitter account to make a report?

No, you do not need to have a Twitter account to file a report that your private and confidential information has been posted on Twitter.

Can I report unauthorized posting of Copyrighted material?

If your copyrighted material has been Tweeted without your permission, please review our [Copyright and DMCA policy](#).

What can I do if my information is posted on other websites?

We cannot take down any information that is posted on a website other than Twitter. If your information is posted on another website, please contact the help center or support staff for the website where your information is posted.

How do I report that my private information is posted on Twitter?

You can report that your private information has been posted on Twitter directly from the offending Tweet(s). Learn [how to report individual Tweets](#).

File a [report that someone has posted your private information](#).

What happens when Twitter receives a valid report?

When we receive a complete and valid report that private information has been posted on Twitter, we'll investigate the account and Tweets reported. We will review where, if anywhere, the information has been made publicly available before taking action on the account or Tweets. If the information you reported was previously posted elsewhere on the Internet, it is not a violation of our policy and we will not take action.

Learn more about [protecting your private information on Twitter and other websites](#).

Read our [policy regarding private information posted on Twitter](#).

Impersonation policy (20171223*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter impersonation policy.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to Twitter's [parody, commentary, and fan account policy](#) for more information about these types of accounts.

Find more information about [reporting impersonation on Twitter](#).

Following rules and best practices (20160117*)

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then

unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation rules (20171103)

Updated November 3, 2017

This page is primarily intended for **developers**.

For Twitter users: You are ultimately responsible for the actions taken with your account, or by applications associated

with your account. Before authorizing a third-party application to access or use your account, make sure you've thoroughly investigated the application and understand what it will do. If automated activity on your account violates the [Twitter Rules](#) or these automation rules, Twitter may take action on your account, including [filtering your Tweets from search results](#) or suspending your account.

For more information on third-party applications, please see our article on [connecting and revoking third-party applications](#).

If you are a developer and have questions about these Automation Rules, please reach out via <https://twittercommunity.com/c/rules-and-policies>.

If your application has been suspended or restricted from performing write actions, you may submit an appeal at </forms/platform>.

I. Ground Rules Do!

Build solutions that automatically broadcast helpful information in Tweets.

Run creative campaigns that auto-reply to users who engage with your content.

Build solutions that automatically respond to users in Direct Messages.

Try new things that help people (and comply with our rules).

Make sure your application provides a good user experience and performs well — and confirm that remains the case over time.

Don't!

Violate these or other policies. Be extra mindful of our rules about abuse and user privacy.

Abuse the Twitter API or attempt to circumvent rate limits.

Use non-API-based forms of automation, such as scripting the Twitter website. The use of these techniques may result in the permanent suspension of your account.

Spam or bother users, or otherwise send them unsolicited messages.

A. The Twitter Rules and the Developer Agreement and Policy

As with all activity on Twitter, automated activity is subject to the [Twitter Rules](#) and, if you're a developer using the Twitter API, the [Developer Agreement and Policy](#).

You should carefully review these policies to ensure that your automated activity is compliant. Automated applications or activities that violate these policies, or that facilitate or induce users to violate them, may be subject to enforcement action, potentially including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these policies.

Although all aspects of the Twitter Rules and the Developer Agreement and Policy apply to automated activity, you should keep the following rules top of mind:

Spamming: You may not send automated Tweets or Direct Messages that are spam, or otherwise engage in spamming activity. Some examples of spammy behavior to avoid with automation include:

Trending topics: You may not automatically post about trending topics on Twitter, or use automation to attempt to influence or manipulate trending topics.

Multiple posts/accounts: You may not post duplicative or substantially similar Tweets on one account or over multiple accounts you operate.

Duplicate accounts: You may not create and/or automate multiple accounts for duplicative or substantially similar use cases.

However, automating multiple accounts for related but non-duplicative use cases is permitted. For example, you may automate separate accounts to Tweet when the Hubble Space Telescope passes over different cities, such as [San Francisco](#) or [Hong Kong](#).

Misleading links: You may not send automated Tweets or Direct Messages containing links that are misleading, including links that maliciously or deceptively redirect through landing pages or ad pages before displaying the final content.

Sensitive media: Automated Tweets and Direct Messages must comply with the [Twitter media policy](#), and you should mark your account as potentially sensitive if you intend to post graphic, pornographic, or potentially sensitive media.

Abusive behavior: You may not engage in any automated activity that encourages, promotes, or incites abuse, violence, hateful conduct, or harassment, on or off Twitter.

Private information: You may not post private or confidential information about a person without their prior express authorization.

B. Other Ground Rules for Automated Activity

In addition to the policies above, the following ground rules apply to all automated activity on Twitter:

Don't surprise or mislead users: Automated activity should honor users' expectations. Ask for the user's permission before taking an action if you aren't sure.

Mature content or profanity: Don't Direct Message, mention, or reply to users with potentially sensitive content (including profanity), unless they've clearly indicated an intent to receive it in advance.

Be thoughtful about the information you request or exchange on Twitter

Tweets: Don't ask users to send you personal or private information via a public Tweet. If you need additional personal or private information from a user to provide them with customer service (or other similar use cases), you should ask the user to share such information by Direct Message or another private channel. You might even consider adding a [Direct Message deep link](#) to your Tweet.

Direct Messages: You should only ask users for the minimum amount of information you need to provide them with service. If you need to request or exchange particularly sensitive information (such as credit card information), you should consider directing users to your website or other appropriate channel to do so.

II. Activity-Specific Rules

The activity-specific rules in this section apply to taking specific automated actions on Twitter. Please read these rules carefully, as they outline both permitted and prohibited use cases of automation.

Automated applications or activities that violate these rules, or that facilitate or induce users to violate them, may be subject to enforcement action, including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access the Twitter API based on violations of these rules. As a reminder, you should also carefully review the spam guidelines in the [Twitter Rules](#) to avoid having activities performed by you, your app, or other users through your app or service flagged as spam.

A. Automated Actions Through Another User's Account

Twitter users may authorize your app or service to [access their Twitter account through OAuth](#). A user authorizing your app or service to access their Twitter account through OAuth does not by itself constitute sufficient consent to take automated actions through that user's account.

You may only take automated actions through another Twitter user's account if you:

clearly describe to the user the types of automated actions that will occur; receive express consent from the user to take those automated actions; and immediately honor a user's request to opt-out of further automated actions.

If you substantially change the purpose or functionality of your app or service, you must re-obtain express consent from the user to take automated action through their account before doing so.

These requirements apply to any automated action taken through another Twitter user's account, including posting Tweets, sending Direct Messages, deleting Tweets or Direct Messages, or following/unfollowing other accounts. For applications that offer users the ability to delete Tweets in a bulk or automated manner, you must also clearly state that Tweets are not recoverable once deleted.

B. Automated Tweets

1. Posting automated Tweets

Automated Tweets that cross-post outside information:

You may post automated Tweets based on sources of outside information — such as an RSS feed, weather data, etc. — as long as you are sufficiently authorized to publish such information.

Other automated Tweets (excluding mentions or replies):

Provided you comply with all other rules, you may post automated Tweets for entertainment, informational, or novelty purposes. As a reminder, accounts posting duplicative, spammy, or otherwise prohibited content may be subject to suspension.

2. Posting automated mentions and replies

The reply and mention functions are intended to make communication between Twitter users easier. Automating these actions to reach many users on an unsolicited basis is an abuse of the feature, and is not permitted. For example, sending automated replies to Tweets based on keyword searches alone is not permitted. Spamy or duplicative use of mentions and replies may result in enforcement action, such as the removal of your Tweets from Search or the suspension of your app or account.

However, you may send automated replies or mentions to Twitter users so long as:

in advance of sending the automated reply, the recipient or mentioned user(s) have requested or have clearly indicated an intent on Twitter to be contacted by you (i.e. opted in), for example by replying to a Tweet from your account, or by sending you a Direct Message;

you provide a clear and easy way for such users to opt-out of receiving automated replies and mentions, and promptly honor all such opt-out requests;

you only send one automated reply or mention per user interaction; and

the automated reply or mention is a reply to the user's original Tweet (if your campaign is based on users posting a reply to your Tweet).

Opt-in techniques and indications of user intent take many different forms, depending on the specifics of your use case and implementation. Some examples include:

A Tweet from your account that clearly indicates that a user taking a specific action on that Tweet (such as Retweeting it) will opt the user into receiving an automated response.

A mention of your account by the user in a manner suggesting the user clearly wishes or intends to receive a response. If you want to run an auto-reply campaign with a campaign- or use-case-specific hashtag, users should also mention you in their Tweets.

Note that a user following your account is not on its own a sufficient indication of user intent to receive an automated response.

: Advertisers, publishers, and brands using auto-response campaigns must request approval from Twitter and may be subject to additional rules. Please reach out to your account or partner manager for assistance. If you do not have a partner manager but are interested in managed client services, you can learn more [here](#).

Additionally, we recommend that any accounts that will communicate with

users via automated mentions or replies:

Note

1. Appropriately filter responses based on potentially sensitive language in user handles,

display names, and Tweet text, as well as potentially sensitive media;

2. Check that the Tweet you are mentioning or replying to still exists (for example, using the statuses/lookup endpoint on the Twitter API).

C. Automated Direct Messages

Sending automated Direct Messages to users

You may not send unsolicited Direct Messages in a bulk or automated manner, and should be thoughtful about the frequency with which you contact users via Direct Message.

You may send automated Direct Messages to users so long as:

in advance of sending the Direct Message, the recipient(s) have requested or have clearly indicated an intent on Twitter to be contacted by you via Direct Message, for example by sending you a Direct Message; **and**

you provide a clear and easy way for such users to opt-out of receiving automated Direct Messages, and promptly honor all such opt-out requests.

The fact that a user is technically able to receive a Direct Message from you (e.g. because the user follows you, has enabled the ability to receive Direct Messages from any account, or because the user is in a pre-existing Direct Message conversation with you) does not necessarily mean they have requested or expect to receive automated Direct Messages from you.

Interacting with users via Direct Message

Per the Ground Rules, remember to be thoughtful about the amount and type of information you request or exchange with users via Direct Messages. If you will be asking a user to provide personal or private information via an automated Direct Message, you must clearly explain how you will use the information you're collecting. Consider including a link to your privacy policy in your Direct Message to the user, as well as in your Twitter profile bio.

Don't publicly share information received in a Direct Message conversation with a user without first obtaining explicit consent from the user. For example, if a user asks you via Direct Message about a purchase they made from you, you may not mention the user in a public Tweet that includes information about their purchase unless you have the user's explicit consent to do so.

After a user-initiated interaction ends, don't send additional follow-up Direct Messages or mention users in a Tweet unless you get permission from the user.

D. Automated actions you take on Tweets or accounts

Automated likes: You may not like Tweets in an automated manner.

Automated Retweets: Provided you comply with all other rules, you may Retweet or Quote Tweet in an automated manner for entertainment, informational, or novelty purposes. Automated Retweets often lead to negative user experiences, and bulk, aggressive, or spammy Retweeting is a violation of the [Twitter Rules](#).

Automated following/unfollowing: You may not follow or unfollow Twitter accounts in a bulk, aggressive, or indiscriminate manner. Aggressive following is a violation of the [Twitter Rules](#). Please also review our [following rules and best practices](#) to ensure you are in compliance. Note that applications that claim to get users more followers are also prohibited under the [Twitter Rules](#).

Automated adding to lists or collections: You may not add Twitter users to lists or add Tweets to collections in a bulk or indiscriminate manner. Adding a large number of unrelated users to lists is a violation of the [Twitter Rules](#).

About search rules and restrictions

(20180513*)

This article details the rules for using Twitter search, and explains reasons why content may not show up in Twitter search.

Why doesn't all content show up in Twitter search?

In order to keep your search results relevant, Twitter filters search results for quality Tweets and accounts. Material that jeopardizes search quality or creates a bad search experience for other people may be automatically removed from Twitter search. Read more about this [here](#).

General rules for using Twitter search

The [Twitter Rules](#) explain what behaviors are permitted on Twitter. In addition to these rules, we've included some tips below to keep your content relevant (and your Tweets in Twitter search).

Please do not:

- Repeatedly post duplicate or near-duplicate content (links or Tweets).
- Abuse trending topics or hashtags (topic words with a # sign).
- Send automated Tweets or replies.
- Use bots or applications to post similar messages based on keywords.
- Post similar messages over multiple accounts.
- Aggressively follow and unfollow people.

Twitter may automatically remove accounts engaging in these behaviors from search (or even suspend in some cases) in order to ensure the best experience for everyone.

Search results filtering

To learn how to filter your search results, including how to disable or enable safe search, read our [Using Twitter search](#) article.

Search results filtering in third-party applications

Accounts may be filtered from search for posting similar messages over several accounts. Your account may be filtered from search if you are logged into a third-party application that is updating many accounts with similar or duplicate content.

Rules for participating in contests on Twitter

Businesses sometimes host contests and offer prizes to Retweet an update, following a particular account, or posting Tweets with a specific hashtag topic or reply. If you're running a contest, check out this [article](#) for guidelines for contests on Twitter.

If you repeatedly post near-duplicate updates or duplicate links in order to enter contests, you may be filtered from search. Please keep in mind that you are responsible for the updates that you post to your Twitter account.

What to do if you are missing from search

Please see this [article](#) for more information related to this issue.

If you are being filtered from search, rest assured that your followers will still receive your updates. If you post a reply or mention to a specific person, it will be delivered to that person.

Having trouble?

You can learn more about searching by keyword, location, and trending topics in our [Using Twitter search](#) help article.

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. These limitations are set forth in the Twitter Rules below.

The Twitter Rules (along with all incorporated policies), [Privacy Policy](#), and [Terms of Service](#) collectively make up the "Twitter User Agreement" that governs a user's access to and use of Twitter's services.

All individuals accessing or using Twitter's services must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in Twitter taking one or more of the following enforcement actions:

- requiring you to delete prohibited content before you can again create new posts and interact with other Twitter users;
- temporarily limiting your ability to create posts or interact with other Twitter users;
- asking you to verify account ownership with a phone number or email address; or
- permanently suspending your account(s).

If you attempt to evade a permanent suspension by creating new accounts, we will suspend your new accounts.

Please note that we may need to change these Rules from time to time and reserve the right to do so. The most current version is always available at: <https://twitter.com/rules>.

The policies set forth in these Twitter Rules govern organic content on our platform. To learn more about the rules which govern ads and promoted content, please review our [Ads policies](#).

Content Boundaries and Use of Twitter

Intellectual property

Trademark: We reserve the right to suspend accounts or take other appropriate action when someone's brand or trademark, including business name and/or logo, is used in a manner that may mislead or confuse others about your brand affiliation. Read more about our [trademark policy](#) and [how to report a violation](#).

Copyright: We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in our [Terms of Service](#). Read more about our [copyright policy](#).

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Learn more about our [media policy](#).

Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images. Additionally, Twitter may sometimes require you to remove excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about [how to make such a request](#), and [how to mark your media as sensitive](#).

Unlawful use

You may not use our service for any unlawful purposes or in furtherance of illegal activities. By using Twitter, you agree to comply with all applicable laws governing your online conduct and content.

Trends

At times, we may prevent certain content from trending. This includes content that violates the Twitter Rules, as well as content that may attempt to manipulate trends. Read more about [what we allow and do not allow to trend](#).

Misuse of Twitter badges

You may not use badges, including but not limited to the “promoted” or “verified” Twitter badges, unless provided by Twitter. Accounts using unauthorized badges as part of their profile photos, header photos, display names, or in any way that falsely implies affiliation with Twitter or authorization from Twitter to display these badges, may be suspended.

Misuse of usernames

Selling usernames: You may not buy or sell Twitter usernames.

Username squatting: You may not engage in username squatting. Some of the factors we take into consideration when determining whether conduct is username squatting include:

- the number of accounts created;
- the creation of accounts for the purpose of preventing others from using those account names;
- the creation of accounts for the purpose of selling those accounts; and
- the use of third-party content feeds to update and maintain accounts under the names of those third parties.

Please note that Twitter may also remove accounts that are inactive for more than six months. Learn more about [username squatting](#).

Abusive Behavior

We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we prohibit behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Context matters when evaluating for abusive behavior and determining appropriate enforcement actions. Factors we may take into consideration include, but are not limited to whether:

- the behavior is targeted at an individual or group of people;
- the report has been filed by the target of the abuse or a bystander;
- the behavior is newsworthy and in the legitimate public interest.

Violence and physical harm

Violence: You may not make specific threats of violence or wish for the serious physical harm, death, or disease of an individual or group of people. This includes, but is not limited to, threatening or promoting terrorism. You also may not affiliate with organizations that — whether by their own statements or activity both on and off the platform — use or promote violence against civilians to further their causes. We will begin enforcing this rule around affiliation with such organizations on December 18, 2017.

Suicide or self-harm: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Child sexual exploitation: You may not promote child sexual exploitation. Learn more about our zero-tolerance [child sexual exploitation policy](#).

Abuse and hateful conduct

Abuse: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Unwanted sexual advances: You may not direct abuse at someone by sending unwanted sexual content, objectifying them in a sexually explicit manner, or otherwise engaging in sexual misconduct.

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. Read more about our [hateful conduct policy](#).

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category. We will begin enforcing this rule on December 18, 2017.

Private information and intimate media

Private information: You may not publish or post other people's private information without their express authorization and permission. Definitions of private information may vary depending on local laws. Read more about our [private information policy](#).

Intimate media: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. Read more about intimate media on Twitter.

Threats to expose / hack: You may not threaten to expose someone's private information or intimate media. You also may not threaten to hack or break into someone's digital information.

Impersonation

You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. While you may maintain parody, fan, commentary, or newsfeed accounts, you may not do so if the intent of the account is to engage in spamming or abusive behavior. Read more about our [impersonation policy](#).

Spam and Security

We strive to protect people on Twitter from technical abuse and spam.

To promote a stable and secure environment on Twitter, you may not do, or attempt to do, any of the following while accessing or using Twitter:

- Access, tamper with, or use non-public areas of Twitter, Twitter's computer systems, or the technical delivery systems of Twitter's providers (except as expressly permitted by the Twitter Bug Bounty program).
- Probe, scan, or test the vulnerability of any system or network, or breach or circumvent any security or authentication measures (except as expressly permitted by the Twitter Bug Bounty program).
- Access or search, or attempt to access or search, Twitter by any means (automated or otherwise) other than through our currently available, published interfaces that are provided by Twitter (and only pursuant to the applicable terms and conditions), unless you have been specifically allowed to do so in a separate agreement with Twitter. Note that crawling Twitter is permissible if done in accordance with the provisions of the robots.txt file; however, scraping Twitter without our prior consent is expressly prohibited.
- Forge any TCP/IP packet header or any part of the header information in any email or posting, or in any way use Twitter to send altered, deceptive, or false source-identifying information.
- Interfere with or disrupt the access of any user, host or network, including, without limitation, sending a virus, overloading, flooding, spamming, mail-bombing Twitter's services, or by scripting the creation of content in such a manner as to interfere with or create an undue burden on Twitter.

Any accounts engaging in the following activities may be temporarily locked or subject to permanent suspension:

- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use Twitter's services for the purpose of spamming anyone. Spam is generally defined on Twitter as bulk or aggressive activity that attempts to manipulate or disrupt Twitter or the experience of users on Twitter to drive traffic or attention to unrelated accounts, products, services, or initiatives. Some of the factors that we take into account when determining what conduct is considered to be spamming include:
 - if you have followed and/or unfollowed a large number of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if your Tweets or Direct Messages consist mainly of links shared without commentary;
 - if a large number of people have blocked you in response to high volumes of untargeted, unsolicited, or duplicative content or engagements from your account;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicative or substantially similar content, replies, or mentions over multiple accounts or multiple duplicate updates on one account, or create duplicate or substantially similar accounts;
 - if you post multiple updates to a trending or popular topic with an intent to subvert or manipulate the topic to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you send large numbers of unsolicited replies or mentions;
 - if you add users to lists in a bulk or aggressive manner;
 - if you are randomly or aggressively engaging with Tweets (e.g., likes, Retweets, etc.) or users (e.g., following, adding to lists or Moments, etc.) to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you repeatedly post other people's account information as your own (e.g., bio, Tweets, profile URL, etc.);
 - if you post misleading, deceptive, or malicious links (e.g., affiliate links, links to malware/clickjacking pages, etc.);
 - if you create fake accounts, account interactions, or impressions;
 - if you sell, purchase, or attempt to artificially inflate account interactions (such as followers, Retweets, likes, etc.); and
 - if you use or promote third-party services or apps that claim to get you more followers, Retweets, or likes (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers or engagements to your account or Tweets).

Please see our support articles on [following rules and best practices](#) and [automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts may be permanently suspended.

Content Visibility

Accounts under investigation or which have been detected as sharing content in violation of these Rules may have their account or Tweet visibility limited in various parts of Twitter, including search. To learn more about situations in which content may be restricted on Twitter, please see our support article on [search rules and restrictions](#).

Trademark policy

Updated June 30, 2016

What is a Trademark Policy Violation on Twitter?

Using another's trademark in a manner that may mislead or confuse others about your brand affiliation may be a violation of our trademark policy.

How Does Twitter Respond to Trademark Reports?

We review each report and do the following:

- If we determine that there is a violation of our trademark policy, we may suspend the account. In some instances, we may give the account holder an opportunity to comply with our policies before suspending the account. For more information, see our [Parody, commentary, and fan account policy](#).
- We may also provide the account holder with your name and information included in the copy of the report.
- If you would like to report a trademark violation in a [Twitter Ad](#), please visit our [Trademark Policy for Promoted Ads](#) to learn more.

Please note that Twitter only investigates requests that are submitted by the trademark holder or by an authorized representative of the trademark holder.

What is not a Trademark Policy Violation?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples include:

- Using a trademark in a way that is outside the scope of the trademark registration (e.g. territory, or goods and services identified in the registration).
- Nominative and other fair uses of trademarks are protected uses under our trademark policy, so long as the account is clearly distinguished from the trademark owner. This includes use by resellers in certain regions and accounts engaging in parody, commentary, or news. For more information, see our [Parody, commentary, and fan account policy](#).

How do I Report a Trademark Policy Violation?

Submit trademark reports through our [Trademark Issue support form](#). Using the appropriate form ensures we have all the information necessary to process your report quickly. Please provide all information requested in the form. If you submit a report with incomplete information, we will follow up with you to request that information. Please note that this will result in a delay in processing your report.

If you have trouble using the form, try updating your browser or switching to another browser.

Copyright policy

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged

infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and

4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Twitter media policy

Twitter Rules: Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images.

Twitter Rules: You may not use hateful images or symbols in your profile image or profile header.

Rationale

Some forms of graphic violence, adult content, or hateful imagery may be permitted in Tweets when they are marked as [sensitive media](#). However, you may not include this type of content in live video, or in profile or header images.

While we want people to feel free to share media that reflects their creativity or individuality, or to show what's happening in the world, we will take action when it crosses the line into abuse towards a person, group, or protected category.

Finally, we may sometimes require you to remove media containing excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about how to make such a request [here](#).

When this applies

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. Some examples include, but are not limited to, depictions of:

- the moment at which someone dies
- gruesome crime or accident scenes
- bodily harm, torture, dismemberment, or mutilation

We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity (including close-ups of genitals, buttocks, or breasts)
 - Please note that exceptions may be made for artistic, medical, health, or educational content. Breastfeeding content does not need to be marked as sensitive.
- simulating a sexual act
- intercourse or any sexual act (may involve humans, humanoid animals, cartoons, or anime)

If you see media you think may be in violation of our policy, report it. You can read about how to report sensitive media [here](#).

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, or ethnicity/national origin.

Some examples of hateful imagery include, but are not limited to:

- Symbols historically associated with hate groups (for example, the Nazi swastika)
- Images depicting others as less than human or altered to include hateful symbols
- Altered image references to a mass murder that targeted a protected category

Do I need to be the target of this content in order for it to be a violation of the Twitter Rules?
No, we review both first-person and bystander reports of such content.

Consequences

When adult content, graphic violence, or hateful imagery appears in Tweets, we may place this content behind an interstitial advising viewers to be aware that they will see sensitive media if they click through. This allows us to identify potentially sensitive content that some people may not wish to see. Learn more about [how to control whether you see sensitive media](#).

However, if such content appears in live video, header, or profile images, it will violate our media policy. The consequences for violating this policy depend on the severity of the violation and the person's previous record of violations.

The first time someone violates this policy they will be required to remove the imagery and may be asked to verify their account ownership. Subsequent violations could result in permanent suspension. Learn more about [our range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

How to contact Twitter about media concerning a deceased family member

In order to respect the wishes of loved ones, Twitter will remove imagery of deceased individuals in certain circumstances. Immediate family members and other authorized individuals may request the removal of images or video of deceased individuals, from when critical injury occurs to the moments before or after death, by sending a request to Twitter Inc. via our [privacy form](#). When reviewing such media removal requests, Twitter considers public interest factors such as the newsworthiness of the content and may not be able to honor every request.

Note: Find information about contacting Twitter to [deactivate the account of a deceased or incapacitated user](#).

How to mark media as sensitive links to Media policy

Twitter trends FAQs

Where can I find trends?

Trends are available on the Twitter for iOS app, the Twitter for Android app, and [twitter.com](#).

On Twitter's mobile apps, you can find trends listed under the trends section of the **Explore** tab. When signed in to [twitter.com](#) on a desktop or laptop computer, trends are listed in many places, including the Home timeline, Notifications, search results, and profile pages.

How are trends determined?

Trends are determined by an algorithm and, by default, are tailored for you based on who you follow, your interests, and your location. This algorithm identifies topics that are popular now, rather than topics that have been popular for a while or on a daily basis, to help you discover the hottest emerging topics of discussion on Twitter.

You can choose to see trends that are not tailored for you by selecting a specific trends location on [twitter.com](#), iOS, or Android (instructions below). Location trends identify popular topics among people in a specific geographic location.

Note: The number of Tweets that are related to the trends is just one of the factors the algorithm looks at when ranking and determining trends. Algorithmically, trends and hashtags are grouped together if they are related to the same topic. For instance, #MondayMotivation and #MotivationMonday may both be represented by #MondayMotivation.

What does the # sign mean?

You may notice that some trends have # sign before the word or phrase. This is called a **hashtag** and is included specifically in Tweets to mark them as relating to a topic, so that people can follow the conversation in search.

Will I see world and local events in my trends?

Yes, when signed into your account on the twitter.com, iOS, or Android, trends will be tailored for you based on your location and who you follow. There will be many world and local news events and conversations that will appear in your trends regardless of your personalization.

Can I see trends for a specific location?

Yes, you can view trends for a specific location on twitter.com, iOS, or Android.

View trends via twitter.com:

To view trends for a specific location, click **Change** in your trends box. You can choose to **Keep tailored trends** based on your location and who you follow, or you can select **Change** to choose a nearby location or search locations.

Adjust your trends settings via Twitter for iOS: 0.

In the top menu, tap your profile icon.

17. Tap **Settings and privacy**, then tap **Content preferences**.
18. Under **Explore**, tap **Trends**.

T. Drag the slider next to **Trends for you** to turn on and receive personalized trends.

V. When the feature is disabled, you can change your location by tapping **Change location**. In the search box, type your desired trends location or choose from the list.

Adjust your trends settings via Twitter for Android:

O. In the top menu, you will either see a navigation menu icon or your profile

icon. Tap whichever icon you have, and select **Settings and privacy**.

17. Tap **Content preferences**.

18. Under **Explore**, tap **Trends**.

T. Drag the slider next to **Trends for you** to turn on and receive personalized trends.

V. When the feature is disabled, you can change your location by tapping **Change location**. In the search box, type your desired trends location or choose from the list.

Note: If you don't find your preferred country or city, it means we're not receiving enough Tweets from that geographical area to create a list. You can look up local Tweets on any topic by using [advanced Twitter search operators](#).

What happens when I click or tap on a trend?

Clicking or tapping any of the trends takes you to the Twitter search results for that trend. You'll see all Tweets including that phrase or hashtag. To see what people are saying about a previous trend, perform a [search for that keyword](#).

How can I participate in a trend?

Simply post a Tweet including the exact word or phrase as it appears in the trends list (with the hashtag, if you see one). Due to the large number of people Tweeting about these specific trends, you may not always be able to find your particular Tweet in search, but your followers will always see your Tweets.

Note: Twitter also filters searches for quality. To understand what types of behavior could cause your Tweets to be filtered out, read about our [search rules and restrictions](#).

Are there rules for trends?

Yes. We want trends to promote healthy discussions on Twitter. This means that at times, we may prevent certain content from trending. These include trends that:

Contain profanity or adult/graphic references.

Incite hate on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease.

Violate the [Twitter Rules](#).

In some cases, we may also consider the newsworthiness of the content, or if it is in the public interest when evaluating potential violations. In these cases, the content might continue to trend on our platform.

Note that even if we prevent the hashtag or the content from appearing on the list of trending topics, you may still be able to access conversations around that content on Twitter.

Do trends have descriptions?

Below some trends, you may see the approximate number of Tweets associated with the trend, or personalized context such as who in your network is Tweeting about it.

Additionally, when a Moment is found to be relevant to a trend, we may use the title of the Moment as the trend's description.

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Child sexual exploitation policy

Twitter rules

Twitter does not tolerate any material that features or promotes child sexual exploitation. This may include media, text, illustrated, or computer generated images.

Rationale

Regardless of the intent, viewing, sharing, or linking to child sexual exploitation media contributes to the re-victimization of the depicted minor and, therefore, is not allowed on Twitter. This also applies to content that may further contribute to victimization of children through promotion or glorification of child sexual exploitation.

When this applies

For the purposes of this policy, a minor means any person under the age of 18 years. Examples of content that depicts or promotes child sexual exploitation include, but are not limited to:

Visual depiction of a minor engaging in sexually explicit or sexually suggestive act.

Illustrated, computer-generated or other forms of realistic depictions of a human minor in a sexually explicit context, or engaged in a sexually explicit act.

Links to third-party sites that host child sexual exploitation material.

The following behaviors are also not allowed on Twitter: Fantasizing about or promoting engagement in child sexual exploitation.

Expressing a desire to obtain materials that feature child sexual exploitation.

Recruiting, advertising or expressing interest in: (1) a commercial sex act involving a child, or (2) in harboring and/or transporting a child for sexual purposes.

Sending sexually explicit images to a child. Engaging a minor in a sexually explicit conversation.

Trying to obtain sexually explicit images from a minor or trying to recruit a minor for sexual services through incentive or blackmail.

Identifying alleged victims of childhood sexual exploitation by name or photo.

Consequences

In the majority of cases, the consequence for violating our child sexual exploitation policy is immediate and permanent suspension. In addition, violators will not be allowed to create new accounts in the future.

In a limited number of situations when no malicious intent has been identified, users may instead be asked to remove the violating Tweet and be temporarily locked out of their account before they can Tweet again.

When we are made aware of links depicting or promoting child sexual exploitation they will be removed from the site without further notice and reported to [The National Center for Missing & Exploited Children](#) (NCMEC).

What you can do

If you think you've found a Twitter account distributing or promoting child sexual exploitation, please report it through our designated [form](#). Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. For help on how to find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with NCMEC [here](#) or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the International Association of Internet Hotlines [website](#).

Our partners

We partner with multiple organisations around the world whose work is dedicated to fighting child sexual exploitation. You can find a full list of our safety partners [here](#).

Hateful conduct policy

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you".

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., "I hope you get cancer and die."

Wishing for someone to fall victim to a serious accident, e.g., "I wish that you would get run over by a car next time you run your mouth."

Saying that a group of individuals deserve serious physical injury, e.g., "If this group of protesters don't shut up, they deserve to be shot."

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references

forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an

individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode

before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#). If someone believes their account was suspended in error, they can [submit an appeal](#).

Report private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Read about our [policy regarding private information posted on Twitter](#).

Do I need to have a Twitter account to make a report?

No, you do not need to have a Twitter account to file a report that your private and confidential information has been posted on Twitter.

Can I report unauthorized posting of Copyrighted material?

If your copyrighted material has been Tweeted without your permission, please review our [Copyright and DMCA policy](#).

What can I do if my information is posted on other websites?

We cannot take down any information that is posted on a website other than Twitter. If your information is posted on another website, please contact the help center or support staff for the website where your information is posted.

How do I report that my private information is posted on Twitter?

You can report that your private information has been posted on Twitter directly from the offending Tweet(s). Learn [how to report individual Tweets](#).

File a [report that someone has posted your private information](#).

What happens when Twitter receives a valid report?

When we receive a complete and valid report that private information has been posted on Twitter, we'll investigate the account and Tweets reported. We will review where, if anywhere, the information has been made publicly available before taking action on the account or Tweets. If the information you reported was previously posted elsewhere on the Internet, it is not a violation of our policy and we will not take action.

Learn more about [protecting your private information on Twitter and other websites](#).

Read our [policy regarding private information posted on Twitter](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter impersonation policy.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to Twitter's [parody, commentary, and fan account policy](#) for more information about these types of accounts.

Find more information about [reporting impersonation on Twitter](#).

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation rules

Updated November 3, 2017

This page is primarily intended for **developers**.

For Twitter users: You are ultimately responsible for the actions taken with your account, or by applications associated with your account. Before authorizing a third-party application to access or use your account, make sure you've thoroughly investigated the application and understand what it will do. If automated activity on your account violates the [Twitter Rules](#) or these automation rules, Twitter may take action on your account, including [filtering your Tweets from search results](#) or suspending your account.

For more information on third-party applications, please see our article on [connecting and revoking third-party applications](#).

If you are a developer and have questions about these Automation Rules, please reach out via <https://twittercommunity.com/c/rules-and-policies>.

If your application has been suspended or restricted from performing write actions, you may submit an appeal at </forms/platform>.

I. Ground Rules Do!

Build solutions that automatically broadcast helpful information in Tweets.

Run creative campaigns that auto-reply to users who engage with your content.

Build solutions that automatically respond to users in Direct Messages.

Try new things that help people (and comply with our rules).

Make sure your application provides a good user experience and performs well — and confirm that remains the case over time.

Don't!

Violate these or other policies. Be extra mindful of our rules about abuse and user privacy.

Abuse the Twitter API or attempt to circumvent rate limits.

Use non-API-based forms of automation, such as scripting the Twitter website. The use of these techniques may result in the permanent suspension of your account.

Spam or bother users, or otherwise send them unsolicited messages.

A. The Twitter Rules and the Developer Agreement and Policy

As with all activity on Twitter, automated activity is subject to the [Twitter Rules](#) and, if you're a developer using the Twitter API, the [Developer Agreement and Policy](#).

You should carefully review these policies to ensure that your automated activity is compliant. Automated applications or activities that violate these policies, or that facilitate or induce users to violate them, may be subject to enforcement action, potentially including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these policies.

Although all aspects of the Twitter Rules and the Developer Agreement and Policy apply to automated activity, you should keep the following rules top of mind:

Spamming: You may not send automated Tweets or Direct Messages that are spam, or otherwise engage in spamming activity. Some examples of spammy behavior to avoid with automation include:

Trending topics: You may not automatically post about trending topics on Twitter, or use automation to attempt to influence or manipulate trending topics.

Multiple posts/accounts: You may not post duplicative or substantially similar Tweets on one account or over multiple accounts you operate.

Duplicate accounts: You may not create and/or automate multiple accounts for duplicative or substantially similar use cases.

However, automating multiple accounts for related but non-duplicative use cases is permitted. For example, you may automate separate accounts to Tweet when the Hubble Space Telescope passes over different cities, such as [San Francisco](#) or [Hong Kong](#).

Misleading links: You may not send automated Tweets or Direct Messages containing links that are misleading, including

links that maliciously or deceptively redirect through landing pages or ad pages before displaying the final content.

Sensitive media: Automated Tweets and Direct Messages must comply with the [Twitter media policy](#), and you should mark your account as potentially sensitive if you intend to post graphic, pornographic, or potentially sensitive media.

Abusive behavior: You may not engage in any automated activity that encourages, promotes, or incites abuse, violence, hateful conduct, or harassment, on or off Twitter.

Private information: You may not post private or confidential information about a person without their prior express authorization.

B. Other Ground Rules for Automated Activity

In addition to the policies above, the following ground rules apply to all automated activity on Twitter:

Don't surprise or mislead users: Automated activity should honor users' expectations. Ask for the user's permission before taking an action if you aren't sure.

Mature content or profanity: Don't Direct Message, mention, or reply to users with potentially sensitive content (including profanity), unless they've clearly indicated an intent to receive it in advance.

Be thoughtful about the information you request or exchange on Twitter

Tweets: Don't ask users to send you personal or private information via a public Tweet. If you need additional personal or private information from a user to provide them with customer service (or other similar use cases), you should ask the user to share such information by Direct Message or another private channel. You might even consider adding a [Direct Message deep link](#) to your Tweet.

Direct Messages: You should only ask users for the minimum amount of information you need to provide them with service. If you need to request or exchange particularly sensitive information (such as credit card information), you should consider directing users to your website or other appropriate channel to do so.

II. Activity-Specific Rules

The activity-specific rules in this section apply to taking specific automated actions on Twitter. Please read these rules carefully,

as they outline both permitted and prohibited use cases of automation.

Automated applications or activities that violate these rules, or that facilitate or induce users to violate them, may be subject to enforcement action, including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access the Twitter API based on violations of these rules. As a reminder, you should also carefully review the spam guidelines in the [Twitter Rules](#) to avoid having activities performed by you, your app, or other users through your app or service flagged as spam.

A. Automated Actions Through Another User's Account

Twitter users may authorize your app or service to [access their Twitter account through OAuth](#). A user authorizing your app or service to access their Twitter account through OAuth does not by itself constitute sufficient consent to take automated actions through that user's account.

You may only take automated actions through another Twitter user's account if you:

clearly describe to the user the types of automated actions that will occur; receive express consent from the user to take those automated actions; and immediately honor a user's request to opt-out of further automated actions.

If you substantially change the purpose or functionality of your app or service, you must re-obtain express consent from the user to take automated action through their account before doing so.

These requirements apply to any automated action taken through another Twitter user's account, including posting Tweets, sending Direct Messages, deleting Tweets or Direct Messages, or following/unfollowing other accounts. For applications that offer users the ability to delete Tweets in a bulk or automated manner, you must also clearly state that Tweets are not recoverable once deleted.

B. Automated Tweets

1. Posting automated Tweets

Automated Tweets that cross-post outside information:

You may post automated Tweets based on sources of outside information — such as an RSS feed, weather data, etc. — as long as you are sufficiently authorized to publish such information.

Other automated Tweets (excluding mentions or replies):

Provided you comply with all other rules, you may post automated Tweets for entertainment, informational, or novelty purposes. As a reminder, accounts posting duplicative, spammy, or otherwise prohibited content may be subject to suspension.

2. Posting automated mentions and replies

The reply and mention functions are intended to make communication between Twitter users easier. Automating these actions to reach many users on an unsolicited basis is an abuse of the feature, and is not permitted. For example, sending automated replies to Tweets based on keyword searches alone is not permitted. Spamy or duplicative use of mentions and replies may result in enforcement action, such as the removal of your Tweets from Search or the suspension of your app or account.

However, you may send automated replies or mentions to Twitter users so long as:

in advance of sending the automated reply, the recipient or mentioned user(s) have requested or have clearly indicated an intent on Twitter to be contacted by you (i.e. opted in), for example by replying to a Tweet from your account, or by sending you a Direct Message;

you provide a clear and easy way for such users to opt-out of receiving automated replies and mentions, and promptly honor all such opt-out requests;

you only send one automated reply or mention per user interaction; and

the automated reply or mention is a reply to the user's original Tweet (if your campaign is based on users posting a reply to your Tweet).

Opt-in techniques and indications of user intent take many different forms, depending on the specifics of your use case and implementation. Some examples include:

A Tweet from your account that clearly indicates that a user taking a specific action on that Tweet (such as Retweeting it) will opt the user into receiving an automated response.

A mention of your account by the user in a manner suggesting the user clearly wishes or intends to receive a response. If you want to run an auto-reply campaign with a campaign- or use-case-specific hashtag, users should also mention you in their Tweets.

Note that a user following your account is not on its own a sufficient indication of user intent to receive an automated response.

: Advertisers, publishers, and brands using auto-response campaigns must request approval from Twitter and may be subject to additional rules. Please reach out to your account or partner manager for assistance. If you do not have a partner manager but are interested in managed client services, you can learn more [here](#).

Additionally, we recommend that any accounts that will communicate with

users via automated mentions or replies:

Note

1. Appropriately filter responses based on potentially sensitive language in user handles,

display names, and Tweet text, as well as potentially sensitive media;

2. Check that the Tweet you are mentioning or replying to still exists (for example, using the statuses/lookup endpoint on the Twitter API).

C. Automated Direct Messages

Sending automated Direct Messages to users

You may not send unsolicited Direct Messages in a bulk or automated manner, and should be thoughtful about the frequency with which you contact users via Direct Message.

You may send automated Direct Messages to users so long as:

in advance of sending the Direct Message, the recipient(s) have requested or have clearly indicated an intent on Twitter to be contacted by you via Direct Message, for example by sending you a Direct Message; **and**

you provide a clear and easy way for such users to opt-out of receiving automated Direct Messages, and promptly honor all such opt-out requests.

The fact that a user is technically able to receive a Direct Message from you (e.g. because the user follows you, has enabled the ability to receive Direct Messages from any account, or because the user is in a pre-existing Direct Message conversation with you) does not necessarily mean they have requested or expect to receive automated Direct Messages from you.

Interacting with users via Direct Message

Per the Ground Rules, remember to be thoughtful about the amount and type of information you request or exchange with users via Direct Messages. If you will be asking a user to provide personal or private information via an automated Direct Message, you must clearly explain how you will use the information you're collecting. Consider including a link to your privacy policy in your Direct Message to the user, as well as in your Twitter profile bio.

Don't publicly share information received in a Direct Message conversation with a user without first obtaining explicit consent from the user. For example, if a user asks you via Direct Message about a purchase they made from you, you may not mention the user in a public Tweet that includes information about their purchase unless you have the user's explicit consent to do so.

After a user-initiated interaction ends, don't send additional follow-up Direct Messages or mention users in a Tweet unless you get permission from the user.

D. Automated actions you take on Tweets or accounts

Automated likes: You may not like Tweets in an automated manner.

Automated Retweets: Provided you comply with all other rules, you may Retweet or Quote Tweet in an automated manner for entertainment, informational, or novelty purposes. Automated Retweets often lead to negative user experiences, and bulk, aggressive, or spammy Retweeting is a violation of the [Twitter Rules](#).

Automated following/unfollowing: You may not follow or unfollow Twitter accounts in a bulk, aggressive, or indiscriminate manner. Aggressive following is a violation of the [Twitter Rules](#). Please also review our [following rules and best practices](#) to ensure you are in compliance. Note that applications that claim to get users more followers are also prohibited under the [Twitter Rules](#).

Automated adding to lists or collections: You may not add Twitter users to lists or add Tweets to collections in a bulk or indiscriminate manner. Adding a large number of unrelated users to lists is a violation of the [Twitter Rules](#).

About search rules and restrictions

This article details the rules for using Twitter search, and explains reasons why content may not show up in Twitter search.

Why doesn't all content show up in Twitter search?

In order to keep your search results relevant, Twitter filters search results for quality Tweets and accounts. Material that jeopardizes search quality or creates a bad search experience for other people may be automatically removed from Twitter search. Read more about this [here](#).

General rules for using Twitter search

The [Twitter Rules](#) explain what behaviors are permitted on Twitter. In addition to these rules, we've included some tips below to keep your content relevant (and your Tweets in Twitter search).

Please do not:

- Repeatedly post duplicate or near-duplicate content (links or Tweets).
- Abuse trending topics or hashtags (topic words with a # sign).
- Send automated Tweets or replies.
- Use bots or applications to post similar messages based on keywords.
- Post similar messages over multiple accounts.
- Aggressively follow and unfollow people.

Twitter may automatically remove accounts engaging in these behaviors from search (or even suspend in some cases) in order to ensure the best experience for everyone.

Search results filtering

To learn how to filter your search results, including how to disable or enable safe search, read our [Using Twitter search](#) article.

Search results filtering in third-party applications

Accounts may be filtered from search for posting similar messages over several accounts. Your account may be filtered from search if you are logged into a third-party application that is updating many accounts with similar or duplicate content.

Rules for participating in contests on Twitter

Businesses sometimes host contests and offer prizes to Retweet an update, following a particular account, or posting Tweets with a specific hashtag topic or reply. If you're running a contest, check out this [article](#) for guidelines for contests on Twitter.

If you repeatedly post near-duplicate updates or duplicate links in order to enter contests, you may be filtered from search. Please keep in mind that you are responsible for the updates that you post to your Twitter account.

What to do if you are missing from search

Please see this [article](#) for more information related to this issue.

If you are being filtered from search, rest assured that your followers will still receive your updates. If you post a reply or mention to a specific person, it will be delivered to that person.

Having trouble?

You can learn more about searching by keyword, location, and trending topics in our [Using Twitter search](#) help article.

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. These limitations are set forth in the Twitter Rules below.

The Twitter Rules (along with all incorporated policies), [Privacy Policy](#), and [Terms of Service](#) collectively make up the "Twitter User Agreement" that governs a user's access to and use of Twitter's services.

All individuals accessing or using Twitter's services must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in Twitter taking one or more of the following enforcement actions:

- requiring you to delete prohibited content before you can again create new posts and interact with other Twitter users;
- temporarily limiting your ability to create posts or interact with other Twitter users;
- asking you to verify account ownership with a phone number or email address; or
- permanently suspending your account(s).

If you attempt to evade a permanent suspension by creating new accounts, we will suspend your new accounts.

Please note that we may need to change these Rules from time to time and reserve the right to do so. The most current version is always available at: <https://twitter.com/rules>.

The policies set forth in these Twitter Rules govern organic content on our platform. To learn more about the rules which govern ads and promoted content, please review our [Ads policies](#).

Content Boundaries and Use of Twitter

Intellectual property

Trademark: We reserve the right to suspend accounts or take other appropriate action when someone's brand or trademark, including business name and/or logo, is used in a manner that may mislead or confuse others about your brand affiliation. Read more about our [trademark policy](#) and [how to report a violation](#).

Copyright: We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in our [Terms of Service](#). Read more about our [copyright policy](#).

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Learn more about our [media policy](#).

Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images. Additionally, Twitter may sometimes require you to remove excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about [how to make such a request](#), and [how to mark your media as sensitive](#).

Unlawful use

You may not use our service for any unlawful purposes or in furtherance of illegal activities. By using Twitter, you agree to comply with all applicable laws governing your online conduct and content.

Trends

At times, we may prevent certain content from trending. This includes content that violates the Twitter Rules, as well as content that may attempt to manipulate trends. Read more about [what we allow and do not allow to trend](#).

Misuse of Twitter badges

You may not use badges, including but not limited to the “promoted” or “verified” Twitter badges, unless provided by Twitter. Accounts using unauthorized badges as part of their profile photos, header photos, display names, or in any way that falsely implies affiliation with Twitter or authorization from Twitter to display these badges, may be suspended.

Misuse of usernames

Selling usernames: You may not buy or sell Twitter usernames.

Username squatting: You may not engage in username squatting. Some of the factors we take into consideration when determining whether conduct is username squatting include:

- the number of accounts created;
- the creation of accounts for the purpose of preventing others from using those account names;
- the creation of accounts for the purpose of selling those accounts; and
- the use of third-party content feeds to update and maintain accounts under the names of those third parties.

Please note that Twitter may also remove accounts that are inactive for more than six months. Learn more about [username squatting](#).

Abusive Behavior

We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we prohibit behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Context matters when evaluating for abusive behavior and determining appropriate enforcement actions. Factors we may take into consideration include, but are not limited to whether:

- the behavior is targeted at an individual or group of people;
- the report has been filed by the target of the abuse or a bystander;
- the behavior is newsworthy and in the legitimate public interest.

Violence and physical harm

Violence: You may not make specific threats of violence or wish for the serious physical harm, death, or disease of an individual or group of people. This includes, but is not limited to, threatening or promoting terrorism. You also may not affiliate with organizations that — whether by their own statements or activity both on and off the platform — use or promote violence against civilians to further their causes. We will begin enforcing this rule around affiliation with such organizations on December 18, 2017.

Suicide or self-harm: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Child sexual exploitation: You may not promote child sexual exploitation. Learn more about our zero-tolerance [child sexual exploitation policy](#).

Abuse and hateful conduct

Abuse: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Unwanted sexual advances: You may not direct abuse at someone by sending unwanted sexual content, objectifying them in a sexually explicit manner, or otherwise engaging in sexual misconduct.

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. Read more about our [hateful conduct policy](#).

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category. We will begin enforcing this rule on December 18, 2017.

Private information and intimate media

Private information: You may not publish or post other people's private information without their express authorization and permission. Definitions of private information may vary depending on local laws. Read more about our [private information policy](#).

Intimate media: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. Read more about intimate media on Twitter.

Threats to expose / hack: You may not threaten to expose someone's private information or intimate media. You also may not threaten to hack or break into someone's digital information.

Impersonation

You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. While you may maintain parody, fan, commentary, or newsfeed accounts, you may not do so if the intent of the account is to engage in spamming or abusive behavior. Read more about our [impersonation policy](#).

Spam and Security

We strive to protect people on Twitter from technical abuse and spam.

To promote a stable and secure environment on Twitter, you may not do, or attempt to do, any of the following while accessing or using Twitter:

- Access, tamper with, or use non-public areas of Twitter, Twitter's computer systems, or the technical delivery systems of Twitter's providers (except as expressly permitted by the Twitter Bug Bounty program).
- Probe, scan, or test the vulnerability of any system or network, or breach or circumvent any security or authentication measures (except as expressly permitted by the Twitter Bug Bounty program).
- Access or search, or attempt to access or search, Twitter by any means (automated or otherwise) other than through our currently available, published interfaces that are provided by Twitter (and only pursuant to the applicable terms and conditions), unless you have been specifically allowed to do so in a separate agreement with Twitter. Note that crawling Twitter is permissible if done in accordance with the provisions of the robots.txt file; however, scraping Twitter without our prior consent is expressly prohibited.
- Forge any TCP/IP packet header or any part of the header information in any email or posting, or in any way use Twitter to send altered, deceptive, or false source-identifying information.
- Interfere with or disrupt the access of any user, host or network, including, without limitation, sending a virus, overloading, flooding, spamming, mail-bombing Twitter's services, or by scripting the creation of content in such a manner as to interfere with or create an undue burden on Twitter.

Any accounts engaging in the following activities may be temporarily locked or subject to permanent suspension:

- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use Twitter's services for the purpose of spamming anyone. Spam is generally defined on Twitter as bulk or aggressive activity that attempts to manipulate or disrupt Twitter or the experience of users on Twitter to drive traffic or attention to unrelated accounts, products, services, or initiatives. Some of the factors that we take into account when determining what conduct is considered to be spamming include:
 - if you have followed and/or unfollowed a large number of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if your Tweets or Direct Messages consist mainly of links shared without commentary;
 - if a large number of people have blocked you in response to high volumes of untargeted, unsolicited, or duplicative content or engagements from your account;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicative or substantially similar content, replies, or mentions over multiple accounts or multiple duplicate updates on one account, or create duplicate or substantially similar accounts;
 - if you post multiple updates to a trending or popular topic with an intent to subvert or manipulate the topic to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you send large numbers of unsolicited replies or mentions;
 - if you add users to lists in a bulk or aggressive manner;
 - if you are randomly or aggressively engaging with Tweets (e.g., likes, Retweets, etc.) or users (e.g., following, adding to lists or Moments, etc.) to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you repeatedly post other people's account information as your own (e.g., bio, Tweets, profile URL, etc.);
 - if you post misleading, deceptive, or malicious links (e.g., affiliate links, links to malware/clickjacking pages, etc.);
 - if you create fake accounts, account interactions, or impressions;
 - if you sell, purchase, or attempt to artificially inflate account interactions (such as followers, Retweets, likes, etc.); and
 - if you use or promote third-party services or apps that claim to get you more followers, Retweets, or likes (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers or engagements to your account or Tweets).

Please see our support articles on [following rules and best practices](#) and [automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts may be permanently suspended.

Content Visibility

Accounts under investigation or which have been detected as sharing content in violation of these Rules may have their account or Tweet visibility limited in various parts of Twitter, including search. To learn more about situations in which content may be restricted on Twitter, please see our support article on [search rules and restrictions](#).

Trademark policy

Updated June 30, 2016

What is a Trademark Policy Violation on Twitter?

Using another's trademark in a manner that may mislead or confuse others about your brand affiliation may be a violation of our trademark policy.

How Does Twitter Respond to Trademark Reports?

We review each report and do the following:

- If we determine that there is a violation of our trademark policy, we may suspend the account. In some instances, we may give the account holder an opportunity to comply with our policies before suspending the account. For more information, see our [Parody, commentary, and fan account policy](#).
- We may also provide the account holder with your name and information included in the copy of the report.
- If you would like to report a trademark violation in a [Twitter Ad](#), please visit our [Trademark Policy for Promoted Ads](#) to learn more.

Please note that Twitter only investigates requests that are submitted by the trademark holder or by an authorized representative of the trademark holder.

What is not a Trademark Policy Violation?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples include:

- Using a trademark in a way that is outside the scope of the trademark registration (e.g. territory, or goods and services identified in the registration).
- Nominative and other fair uses of trademarks are protected uses under our trademark policy, so long as the account is clearly distinguished from the trademark owner. This includes use by resellers in certain regions and accounts engaging in parody, commentary, or news. For more information, see our [Parody, commentary, and fan account policy](#).

How do I Report a Trademark Policy Violation?

Submit trademark reports through our [Trademark Issue support form](#). Using the appropriate form ensures we have all the information necessary to process your report quickly. Please provide all information requested in the form. If you submit a report with incomplete information, we will follow up with you to request that information. Please note that this will result in a delay in processing your report.

If you have trouble using the form, try updating your browser or switching to another browser.

Copyright policy

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information

concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Twitter media policy

Twitter Rules: Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images.

Twitter Rules: You may not use hateful images or symbols in your profile image or profile header.

Rationale

Some forms of graphic violence, adult content, or hateful imagery may be permitted in Tweets when they are marked as [sensitive media](#). However, you may not include this type of content in live video, or in profile or header images.

While we want people to feel free to share media that reflects their creativity or individuality, or to show what's happening in the world, we will take action when it crosses the line into abuse towards a person, group, or protected category.

Finally, we may sometimes require you to remove media containing excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about how to make such a request [here](#).

When this applies

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. Some examples include, but are not limited to, depictions of:

- the moment at which someone dies
- gruesome crime or accident scenes
- bodily harm, torture, dismemberment, or mutilation

We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity (including close-ups of genitals, buttocks, or breasts)
 - Please note that exceptions may be made for artistic, medical, health, or educational content. Breastfeeding content does not need to be marked as sensitive.
- simulating a sexual act
- intercourse or any sexual act (may involve humans, humanoid animals, cartoons, or anime)

If you see media you think may be in violation of our policy, report it. You can read about how to report sensitive media [here](#).

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, or ethnicity/national origin.

Some examples of hateful imagery include, but are not limited to:

- Symbols historically associated with hate groups (for example, the Nazi swastika)
- Images depicting others as less than human or altered to include hateful symbols
- Altered image references to a mass murder that targeted a protected category

Do I need to be the target of this content in order for it to be a violation of the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When adult content, graphic violence, or hateful imagery appears in Tweets, we may place this content behind an interstitial advising viewers to be aware that they will see sensitive media if they click through. This allows us to identify potentially sensitive content that some people may not wish to see. Learn more about [how to control whether you see sensitive media](#).

However, if such content appears in live video, header, or profile images, it will violate our media policy. The consequences for violating this policy depend on the severity of the violation and the person's previous record of violations.

The first time someone violates this policy they will be required to remove the imagery and may be asked to verify their account ownership. Subsequent violations could result in permanent suspension. Learn more about [our range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

How to contact Twitter about media concerning a deceased family member

In order to respect the wishes of loved ones, Twitter will remove imagery of deceased individuals in certain circumstances. Immediate family members and other authorized individuals may request the removal of images or video of deceased individuals, from when critical injury occurs to the moments before or after death, by sending a request to Twitter Inc. via our [privacy form](#). When reviewing such media removal requests, Twitter considers public interest factors such as the newsworthiness of the content and may not be able to honor every request. Note: Find information about contacting Twitter to [deactivate the account of a deceased or incapacitated user](#).

How to mark media as sensitive links to Media policy

Twitter trends FAQs

Where can I find trends?

Trends are available on the Twitter for iOS app, the Twitter for Android app, and [twitter.com](#).

On Twitter's mobile apps, you can find trends listed under the trends section of the **Explore** tab. When signed in to [twitter.com](#) on a desktop or laptop computer, trends are listed in many places, including the Home timeline, Notifications, search results, and profile pages.

How are trends determined?

Trends are determined by an algorithm and, by default, are tailored for you based on who you follow, your interests, and your location. This algorithm identifies topics that are popular now, rather than topics that have been popular for a while or on a daily basis, to help you discover the hottest emerging topics of discussion on Twitter.

You can choose to see trends that are not tailored for you by selecting a specific trends location on [twitter.com](#), iOS, or Android (instructions below). Location trends identify popular topics among people in a specific geographic location.

Note: The number of Tweets that are related to the trends is just one of the factors the algorithm looks at when ranking and determining trends. Algorithmically, trends and hashtags are

grouped together if they are related to the same topic. For instance, #MondayMotivation and #MotivationMonday may both be represented by #MondayMotivation.

What does the # sign mean?

You may notice that some trends have # sign before the word or phrase. This is called a **hashtag** and is included specifically in Tweets to mark them as relating to a topic, so that people can follow the conversation in search.

Will I see world and local events in my trends?

Yes, when signed into your account on the twitter.com, iOS, or Android, trends will be tailored for you based on your location and who you follow. There will be many world and local news events and conversations that will appear in your trends regardless of your personalization.

Can I see trends for a specific location?

Yes, you can view trends for a specific location on twitter.com, iOS, or Android.

View trends via twitter.com:

To view trends for a specific location, click **Change** in your trends box. You can choose to **Keep tailored trends** based on your location and who you follow, or you can select **Change** to choose a nearby location or search locations.

Adjust your trends settings via Twitter for iOS:

0. In the top menu, tap your **profile** icon.

17. Tap **Settings and privacy**, then tap **Content preferences**.
18. Under **Explore**, tap **Trends**.

T. Drag the slider next to **Trends for you** to turn on and receive personalized trends.

V. When the feature is disabled, you can change your location by tapping **Change location**. In the search box, type your desired trends location or choose from the list.

Adjust your trends settings via Twitter for Android:

0. In the top menu, you will either see a **navigation menu** icon or your **profile**

icon. Tap whichever icon you have, and select **Settings and privacy**.

17. Tap **Content preferences**.
18. Under **Explore**, tap **Trends**.

T. Drag the slider next to **Trends for you** to turn on and receive personalized trends.

V. When the feature is disabled, you can change your location by tapping **Change location**. In the search box, type your desired trends location or choose from the list.

Note: If you don't find your preferred country or city, it means we're not receiving enough Tweets from that geographical area to create a list. You can look up local Tweets on any topic by using [advanced Twitter search operators](#).

What happens when I click or tap on a trend?

Clicking or tapping any of the trends takes you to the Twitter search results for that trend. You'll see all Tweets including that phrase or hashtag. To see what people are saying about a previous trend, perform a [search for that keyword](#).

How can I participate in a trend?

Simply post a Tweet including the exact word or phrase as it appears in the trends list (with the hashtag, if you see one). Due to the large number of people Tweeting about these specific trends, you may not always be able to find your particular Tweet in search, but your followers will always see your Tweets.

Note: Twitter also filters searches for quality. To understand what types of behavior could cause your Tweets to be filtered out, read about our [search rules and restrictions](#).

Are there rules for trends?

Yes. We want trends to promote healthy discussions on Twitter. This means that at times, we may prevent certain content from trending. These include trends that:

Contain profanity or adult/graphic references.

Incite hate on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease.

Violate the [Twitter Rules](#).

In some cases, we may also consider the newsworthiness of the content, or if it is in the public interest when evaluating potential violations. In these cases, the content might continue to trend on our platform.

Note that even if we prevent the hashtag or the content from appearing on the list of trending topics, you may still be able to access conversations around that content on Twitter.

Do trends have descriptions?

Below some trends, you may see the approximate number of Tweets associated with the trend, or personalized context such as who in your network is Tweeting about it.

Additionally, when a Moment is found to be relevant to a trend, we may use the title of the Moment as the trend's description.

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Child sexual exploitation policy

Twitter rules

Twitter does not tolerate any material that features or promotes child sexual exploitation. This may include media, text, illustrated, or computer generated images.

Rationale

Regardless of the intent, viewing, sharing, or linking to child sexual exploitation media contributes to the re-victimization of the depicted minor and, therefore, is not allowed on Twitter. This also applies to content that may further contribute to victimization of children through promotion or glorification of child sexual exploitation.

When this applies

For the purposes of this policy, a minor means any person under the age of 18 years. Examples of content that depicts or promotes child sexual exploitation include, but are not limited to:

Visual depiction of a minor engaging in sexually explicit or sexually suggestive act.

Illustrated, computer-generated or other forms of realistic depictions of a human minor in a sexually explicit context, or engaged in a sexually explicit act.

Links to third-party sites that host child sexual exploitation material.

The following behaviors are also not allowed on Twitter: Fantasizing about or promoting engagement in child sexual exploitation.

Expressing a desire to obtain materials that feature child sexual exploitation.

Recruiting, advertising or expressing interest in: (1) a commercial sex act involving a child, or (2) in harboring and/or transporting a child for sexual purposes.

Sending sexually explicit images to a child. Engaging a minor in a sexually explicit conversation.

Trying to obtain sexually explicit images from a minor or trying to recruit a minor for sexual services through incentive or blackmail.

Identifying alleged victims of childhood sexual exploitation by name or photo.

Consequences

In the majority of cases, the consequence for violating our child sexual exploitation policy is immediate and permanent suspension. In addition, violators will not be allowed to create new accounts in the future.

In a limited number of situations when no malicious intent has been identified, users may instead be asked to remove the violating Tweet and be temporarily locked out of their account before they can Tweet again.

When we are made aware of links depicting or promoting child sexual exploitation they will be removed from the site without further notice and reported to [The National Center for Missing & Exploited Children](#) (NCMEC).

What you can do

If you think you've found a Twitter account distributing or promoting child sexual exploitation, please report it through our designated [form](#). Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. For help on how to find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with NCMEC [here](#) or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the International Association of Internet Hotlines [website](#).

Our partners

We partner with multiple organisations around the world whose work is dedicated to fighting child sexual exploitation. You can find a full list of our safety partners [here](#).

Hateful conduct policy

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., “I will kill you”.

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., “I hope you get cancer and die.”

Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”

Saying that a group of individuals deserve serious physical injury, e.g., “If this group of protesters don’t shut up, they deserve to be shot.”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references

forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an

individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode

before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#). If someone believes their account was suspended in error, they can [submit an appeal](#).

Report private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Read about our [policy regarding private information posted on Twitter](#).

Do I need to have a Twitter account to make a report?

No, you do not need to have a Twitter account to file a report that your private and confidential information has been posted on Twitter.

Can I report unauthorized posting of Copyrighted material?

If your copyrighted material has been Tweeted without your permission, please review our [Copyright and DMCA policy](#).

What can I do if my information is posted on other websites?

We cannot take down any information that is posted on a website other than Twitter. If your information is posted on another website, please contact the help center or support staff for the website where your information is posted.

How do I report that my private information is posted on Twitter?

You can report that your private information has been posted on Twitter directly from the offending Tweet(s). Learn [how to report individual Tweets](#).

File a [report that someone has posted your private information](#).

What happens when Twitter receives a valid report?

When we receive a complete and valid report that private information has been posted on Twitter, we'll investigate the account and Tweets reported. We will review where, if anywhere, the information has been made publicly available before taking action on the account or Tweets. If the information you reported was previously posted elsewhere on the Internet, it is not a violation of our policy and we will not take action.

Learn more about [protecting your private information on Twitter and other websites](#).

Read our [policy regarding private information posted on Twitter](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts portraying another person in a confusing or deceptive manner may be permanently suspended under the Twitter impersonation policy.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals.

Accounts with similar usernames or that are similar in appearance (e.g. the same avatar image) are not automatically in violation of the impersonation policy. In order to be impersonation, the account must also portray another person in a misleading or deceptive manner.

Twitter users are allowed to create parody, commentary, or fan accounts. Please refer to Twitter's [parody, commentary, and fan account policy](#) for more information about these types of accounts.

Find more information about [reporting impersonation on Twitter](#).

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation rules

Updated November 3, 2017

This page is primarily intended for **developers**.

For Twitter users: You are ultimately responsible for the actions taken with your account, or by applications associated with your account. Before authorizing a third-party application to access or use your account, make sure you've thoroughly investigated the application and understand what it will do. If automated activity on your account violates the [Twitter Rules](#) or these automation rules, Twitter may take action on your account, including [filtering your Tweets from search results](#) or suspending your account.

For more information on third-party applications, please see our article on [connecting and revoking third-party applications](#).

If you are a developer and have questions about these Automation Rules, please reach out via <https://twittercommunity.com/c/rules-and-policies>.

If your application has been suspended or restricted from performing write actions, you may submit an appeal at </forms/platform>.

I. Ground Rules Do!

Build solutions that automatically broadcast helpful information in Tweets.

Run creative campaigns that auto-reply to users who engage with your content.

Build solutions that automatically respond to users in Direct Messages.

Try new things that help people (and comply with our rules).

Make sure your application provides a good user experience and performs well — and confirm that remains the case over time.

Don't!

Violate these or other policies. Be extra mindful of our rules about abuse and user privacy.

Abuse the Twitter API or attempt to circumvent rate limits.

Use non-API-based forms of automation, such as scripting the Twitter website. The use of these techniques may result in the permanent suspension of your account.

Spam or bother users, or otherwise send them unsolicited messages.

A. The Twitter Rules and the Developer Agreement and Policy

As with all activity on Twitter, automated activity is subject to the [Twitter Rules](#) and, if you're a developer using the Twitter API, the [Developer Agreement and Policy](#).

You should carefully review these policies to ensure that your automated activity is compliant. Automated applications or activities that violate these policies, or that facilitate or induce users to violate them, may be subject to enforcement action, potentially including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these policies.

Although all aspects of the Twitter Rules and the Developer Agreement and Policy apply to automated activity, you should keep the following rules top of mind:

Spamming: You may not send automated Tweets or Direct Messages that are spam, or otherwise engage in spamming activity. Some examples of spammy behavior to avoid with automation include:

Trending topics: You may not automatically post about trending topics on Twitter, or use automation to attempt to influence or manipulate trending topics.

Multiple posts/accounts: You may not post duplicative or substantially similar Tweets on one account or over multiple accounts you operate.

Duplicate accounts: You may not create and/or automate multiple accounts for duplicative or substantially similar use cases.

However, automating multiple accounts for related but non-duplicative use cases is permitted. For example, you may automate separate accounts to Tweet when the Hubble Space Telescope passes over different cities, such as [San Francisco](#) or [Hong Kong](#).

Misleading links: You may not send automated Tweets or Direct Messages containing links that are misleading, including

links that maliciously or deceptively redirect through landing pages or ad pages before displaying the final content.

Sensitive media: Automated Tweets and Direct Messages must comply with the [Twitter media policy](#), and you should mark your account as potentially sensitive if you intend to post graphic, pornographic, or potentially sensitive media.

Abusive behavior: You may not engage in any automated activity that encourages, promotes, or incites abuse, violence, hateful conduct, or harassment, on or off Twitter.

Private information: You may not post private or confidential information about a person without their prior express authorization.

B. Other Ground Rules for Automated Activity

In addition to the policies above, the following ground rules apply to all automated activity on Twitter:

Don't surprise or mislead users: Automated activity should honor users' expectations. Ask for the user's permission before taking an action if you aren't sure.

Mature content or profanity: Don't Direct Message, mention, or reply to users with potentially sensitive content (including profanity), unless they've clearly indicated an intent to receive it in advance.

Be thoughtful about the information you request or exchange on Twitter

Tweets: Don't ask users to send you personal or private information via a public Tweet. If you need additional personal or private information from a user to provide them with customer service (or other similar use cases), you should ask the user to share such information by Direct Message or another private channel. You might even consider adding a [Direct Message deep link](#) to your Tweet.

Direct Messages: You should only ask users for the minimum amount of information you need to provide them with service. If you need to request or exchange particularly sensitive information (such as credit card information), you should consider directing users to your website or other appropriate channel to do so.

II. Activity-Specific Rules

The activity-specific rules in this section apply to taking specific automated actions on Twitter. Please read these rules carefully,

as they outline both permitted and prohibited use cases of automation.

Automated applications or activities that violate these rules, or that facilitate or induce users to violate them, may be subject to enforcement action, including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access the Twitter API based on violations of these rules. As a reminder, you should also carefully review the spam guidelines in the [Twitter Rules](#) to avoid having activities performed by you, your app, or other users through your app or service flagged as spam.

A. Automated Actions Through Another User's Account

Twitter users may authorize your app or service to [access their Twitter account through OAuth](#). A user authorizing your app or service to access their Twitter account through OAuth does not by itself constitute sufficient consent to take automated actions through that user's account.

You may only take automated actions through another Twitter user's account if you:

clearly describe to the user the types of automated actions that will occur; receive express consent from the user to take those automated actions; and immediately honor a user's request to opt-out of further automated actions.

If you substantially change the purpose or functionality of your app or service, you must re-obtain express consent from the user to take automated action through their account before doing so.

These requirements apply to any automated action taken through another Twitter user's account, including posting Tweets, sending Direct Messages, deleting Tweets or Direct Messages, or following/unfollowing other accounts. For applications that offer users the ability to delete Tweets in a bulk or automated manner, you must also clearly state that Tweets are not recoverable once deleted.

B. Automated Tweets

1. Posting automated Tweets

Automated Tweets that cross-post outside information:

You may post automated Tweets based on sources of outside information — such as an RSS feed, weather data, etc. — as long as you are sufficiently authorized to publish such information.

Other automated Tweets (excluding mentions or replies):

Provided you comply with all other rules, you may post automated Tweets for entertainment, informational, or novelty purposes. As a reminder, accounts posting duplicative, spammy, or otherwise prohibited content may be subject to suspension.

2. Posting automated mentions and replies

The reply and mention functions are intended to make communication between Twitter users easier. Automating these actions to reach many users on an unsolicited basis is an abuse of the feature, and is not permitted. For example, sending automated replies to Tweets based on keyword searches alone is not permitted. Spamy or duplicative use of mentions and replies may result in enforcement action, such as the removal of your Tweets from Search or the suspension of your app or account.

However, you may send automated replies or mentions to Twitter users so long as:

in advance of sending the automated reply, the recipient or mentioned user(s) have requested or have clearly indicated an intent on Twitter to be contacted by you (i.e. opted in), for example by replying to a Tweet from your account, or by sending you a Direct Message;

you provide a clear and easy way for such users to opt-out of receiving automated replies and mentions, and promptly honor all such opt-out requests;

you only send one automated reply or mention per user interaction; and

the automated reply or mention is a reply to the user's original Tweet (if your campaign is based on users posting a reply to your Tweet).

Opt-in techniques and indications of user intent take many different forms, depending on the specifics of your use case and implementation. Some examples include:

A Tweet from your account that clearly indicates that a user taking a specific action on that Tweet (such as Retweeting it) will opt the user into receiving an automated response.

A mention of your account by the user in a manner suggesting the user clearly wishes or intends to receive a response. If you want to run an auto-reply campaign with a campaign- or use-case-specific hashtag, users should also mention you in their Tweets.

Note that a user following your account is not on its own a sufficient indication of user intent to receive an automated response.

: Advertisers, publishers, and brands using auto-response campaigns must request approval from Twitter and may be subject to additional rules. Please reach out to your account or partner manager for assistance. If you do not have a partner manager but are interested in managed client services, you can learn more [here](#).

Additionally, we recommend that any accounts that will communicate with

users via automated mentions or replies:

Note

1. Appropriately filter responses based on potentially sensitive language in user handles,

display names, and Tweet text, as well as potentially sensitive media;

2. Check that the Tweet you are mentioning or replying to still exists (for example, using the statuses/lookup endpoint on the Twitter API).

C. Automated Direct Messages

Sending automated Direct Messages to users

You may not send unsolicited Direct Messages in a bulk or automated manner, and should be thoughtful about the frequency with which you contact users via Direct Message.

You may send automated Direct Messages to users so long as:

in advance of sending the Direct Message, the recipient(s) have requested or have clearly indicated an intent on Twitter to be contacted by you via Direct Message, for example by sending you a Direct Message; **and**

you provide a clear and easy way for such users to opt-out of receiving automated Direct Messages, and promptly honor all such opt-out requests.

The fact that a user is technically able to receive a Direct Message from you (e.g. because the user follows you, has enabled the ability to receive Direct Messages from any account, or because the user is in a pre-existing Direct Message conversation with you) does not necessarily mean they have requested or expect to receive automated Direct Messages from you.

Interacting with users via Direct Message

Per the Ground Rules, remember to be thoughtful about the amount and type of information you request or exchange with users via Direct Messages. If you will be asking a user to provide personal or private information via an automated Direct Message, you must clearly explain how you will use the information you're collecting. Consider including a link to your privacy policy in your Direct Message to the user, as well as in your Twitter profile bio.

Don't publicly share information received in a Direct Message conversation with a user without first obtaining explicit consent from the user. For example, if a user asks you via Direct Message about a purchase they made from you, you may not mention the user in a public Tweet that includes information about their purchase unless you have the user's explicit consent to do so.

After a user-initiated interaction ends, don't send additional follow-up Direct Messages or mention users in a Tweet unless you get permission from the user.

D. Automated actions you take on Tweets or accounts

Automated likes: You may not like Tweets in an automated manner.

Automated Retweets: Provided you comply with all other rules, you may Retweet or Quote Tweet in an automated manner for entertainment, informational, or novelty purposes. Automated Retweets often lead to negative user experiences, and bulk, aggressive, or spammy Retweeting is a violation of the [Twitter Rules](#).

Automated following/unfollowing: You may not follow or unfollow Twitter accounts in a bulk, aggressive, or indiscriminate manner. Aggressive following is a violation of the [Twitter Rules](#). Please also review our [following rules and best practices](#) to ensure you are in compliance. Note that applications that claim to get users more followers are also prohibited under the [Twitter Rules](#).

Automated adding to lists or collections: You may not add Twitter users to lists or add Tweets to collections in a bulk or indiscriminate manner. Adding a large number of unrelated users to lists is a violation of the [Twitter Rules](#).

About search rules and restrictions

This article details the rules for using Twitter search, and explains reasons why content may not show up in Twitter search.

Why doesn't all content show up in Twitter search?

In order to keep your search results relevant, Twitter filters search results for quality Tweets and accounts. Material that jeopardizes search quality or creates a bad search experience for other people may be automatically removed from Twitter search. Read more about this [here](#).

General rules for using Twitter search

The [Twitter Rules](#) explain what behaviors are permitted on Twitter. In addition to these rules, we've included some tips below to keep your content relevant (and your Tweets in Twitter search).

Please do not:

- Repeatedly post duplicate or near-duplicate content (links or Tweets).
- Abuse trending topics or hashtags (topic words with a # sign).
- Send automated Tweets or replies.
- Use bots or applications to post similar messages based on keywords.
- Post similar messages over multiple accounts.
- Aggressively follow and unfollow people.

Twitter may automatically remove accounts engaging in these behaviors from search (or even suspend in some cases) in order to ensure the best experience for everyone.

Search results filtering

To learn how to filter your search results, including how to disable or enable safe search, read our [Using Twitter search](#) article.

Search results filtering in third-party applications

Accounts may be filtered from search for posting similar messages over several accounts. Your account may be filtered from search if you are logged into a third-party application that is updating many accounts with similar or duplicate content.

Rules for participating in contests on Twitter

Businesses sometimes host contests and offer prizes to Retweet an update, following a particular account, or posting Tweets with a specific hashtag topic or reply. If you're running a contest, check out this [article](#) for guidelines for contests on Twitter.

If you repeatedly post near-duplicate updates or duplicate links in order to enter contests, you may be filtered from search. Please keep in mind that you are responsible for the updates that you post to your Twitter account.

What to do if you are missing from search

Please see this [article](#) for more information related to this issue.

If you are being filtered from search, rest assured that your followers will still receive your updates. If you post a reply or mention to a specific person, it will be delivered to that person.

Having trouble?

You can learn more about searching by keyword, location, and trending topics in our [Using Twitter search](#) help article.

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. These limitations are set forth in the Twitter Rules below.

The Twitter Rules (along with all incorporated policies), [Privacy Policy](#), and [Terms of Service](#) collectively make up the "Twitter User Agreement" that governs a user's access to and use of Twitter's services.

All individuals accessing or using Twitter's services must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in Twitter taking one or more of the following enforcement actions:

- requiring you to delete prohibited content before you can again create new posts and interact with other Twitter users;
- temporarily limiting your ability to create posts or interact with other Twitter users;
- asking you to verify account ownership with a phone number or email address; or
- permanently suspending your account(s).

If you attempt to evade a permanent suspension by creating new accounts, we will suspend your new accounts.

Please note that we may need to change these Rules from time to time and reserve the right to do so. The most current version is always available at: <https://twitter.com/rules>.

The policies set forth in these Twitter Rules govern organic content on our platform. To learn more about the rules which govern ads and promoted content, please review our [Ads policies](#).

Content Boundaries and Use of Twitter

Intellectual property

Trademark: We reserve the right to suspend accounts or take other appropriate action when someone's brand or trademark, including business name and/or logo, is used in a manner that may mislead or confuse others about your brand affiliation. Read more about our [trademark policy](#) and [how to report a violation](#).

Copyright: We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in our [Terms of Service](#). Read more about our [copyright policy](#).

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Learn more about our [media policy](#).

Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images. Additionally, Twitter may sometimes require you to remove excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about [how to make such a request](#), and [how to mark your media as sensitive](#).

Unlawful use

You may not use our service for any unlawful purposes or in furtherance of illegal activities. By using Twitter, you agree to comply with all applicable laws governing your online conduct and content.

Trends

At times, we may prevent certain content from trending. This includes content that violates the Twitter Rules, as well as content that may attempt to manipulate trends. Read more about [what we allow and do not allow to trend](#).

Misuse of Twitter badges

You may not use badges, including but not limited to the “promoted” or “verified” Twitter badges, unless provided by Twitter. Accounts using unauthorized badges as part of their profile photos, header photos, display names, or in any way that falsely implies affiliation with Twitter or authorization from Twitter to display these badges, may be suspended.

Misuse of usernames

Selling usernames: You may not buy or sell Twitter usernames.

Username squatting: You may not engage in username squatting. Some of the factors we take into consideration when determining whether conduct is username squatting include:

- the number of accounts created;
- the creation of accounts for the purpose of preventing others from using those account names;
- the creation of accounts for the purpose of selling those accounts; and
- the use of third-party content feeds to update and maintain accounts under the names of those third parties.

Please note that Twitter may also remove accounts that are inactive for more than six months. Learn more about [username squatting](#).

Abusive Behavior

We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we prohibit behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user’s voice.

Context matters when evaluating for abusive behavior and determining appropriate enforcement actions. Factors we may take into consideration include, but are not limited to whether:

- the behavior is targeted at an individual or group of people;
- the report has been filed by the target of the abuse or a bystander;
- the behavior is newsworthy and in the legitimate public interest.

Violence and physical harm

Violence: You may not make specific threats of violence or wish for the serious physical harm, death, or disease of an individual or group of people. This includes, but is not limited to, threatening or promoting terrorism. You also may not affiliate with organizations that — whether by their own statements or activity both on and off the platform — use or promote violence against civilians to further their causes. We will begin enforcing this rule around affiliation with such organizations on December 18, 2017.

Suicide or self-harm: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Child sexual exploitation: You may not promote child sexual exploitation. Learn more about our zero-tolerance [child sexual exploitation policy](#).

Abuse and hateful conduct

Abuse: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Unwanted sexual advances: You may not direct abuse at someone by sending unwanted sexual content, objectifying them in a sexually explicit manner, or otherwise engaging in sexual misconduct.

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. Read more about our [hateful conduct policy](#).

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category. We will begin enforcing this rule on December 18, 2017.

Private information and intimate media

Private information: You may not publish or post other people's private information without their express authorization and permission. Definitions of private information may vary depending on local laws. Read more about our [private information policy](#).

Intimate media: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. Read more about intimate media on Twitter.

Threats to expose / hack: You may not threaten to expose someone's private information or intimate media. You also may not threaten to hack or break into someone's digital information.

Impersonation

You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. While you may maintain parody, fan, commentary, or newsfeed accounts, you may not do so if the intent of the account is to engage in spamming or abusive behavior. Read more about our [impersonation policy](#).

Spam and Security

We strive to protect people on Twitter from technical abuse and spam.

To promote a stable and secure environment on Twitter, you may not do, or attempt to do, any of the following while accessing or using Twitter:

- Access, tamper with, or use non-public areas of Twitter, Twitter's computer systems, or the technical delivery systems of Twitter's providers (except as expressly permitted by the Twitter Bug Bounty program).
- Probe, scan, or test the vulnerability of any system or network, or breach or circumvent any security or authentication measures (except as expressly permitted by the Twitter Bug Bounty program).
- Access or search, or attempt to access or search, Twitter by any means (automated or otherwise) other than through our currently available, published interfaces that are provided by Twitter (and only pursuant to the applicable terms and conditions), unless you have been specifically allowed to do so in a separate agreement with Twitter. Note that crawling Twitter is permissible if done in accordance with the provisions of the robots.txt file; however, scraping Twitter without our prior consent is expressly prohibited.
- Forge any TCP/IP packet header or any part of the header information in any email or posting, or in any way use Twitter to send altered, deceptive, or false source-identifying information.
- Interfere with or disrupt the access of any user, host or network, including, without limitation, sending a virus, overloading, flooding, spamming, mail-bombing Twitter's services, or by scripting the creation of content in such a manner as to interfere with or create an undue burden on Twitter.

Any accounts engaging in the following activities may be temporarily locked or subject to permanent suspension:

- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Spam:** You may not use Twitter's services for the purpose of spamming anyone. Spam is generally defined on Twitter as bulk or aggressive activity that attempts to manipulate or disrupt Twitter or the experience of users on Twitter to drive traffic or attention to unrelated accounts, products, services, or initiatives. Some of the factors that we take into account when determining what conduct is considered to be spamming include:
 - if you have followed and/or unfollowed a large number of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
 - if your Tweets or Direct Messages consist mainly of links shared without commentary;
 - if a large number of people have blocked you in response to high volumes of untargeted, unsolicited, or duplicative content or engagements from your account;
 - if a large number of spam complaints have been filed against you;
 - if you post duplicative or substantially similar content, replies, or mentions over multiple accounts or multiple duplicate updates on one account, or create duplicate or substantially similar accounts;
 - if you post multiple updates to a trending or popular topic with an intent to subvert or manipulate the topic to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you send large numbers of unsolicited replies or mentions;
 - if you add users to lists in a bulk or aggressive manner;
 - if you are randomly or aggressively engaging with Tweets (e.g., likes, Retweets, etc.) or users (e.g., following, adding to lists or Moments, etc.) to drive traffic or attention to unrelated accounts, products, services, or initiatives;
 - if you repeatedly post other people's account information as your own (e.g., bio, Tweets, profile URL, etc.);
 - if you post misleading, deceptive, or malicious links (e.g., affiliate links, links to malware/clickjacking pages, etc.);
 - if you create fake accounts, account interactions, or impressions;
 - if you sell, purchase, or attempt to artificially inflate account interactions (such as followers, Retweets, likes, etc.); and
 - if you use or promote third-party services or apps that claim to get you more followers, Retweets, or likes (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers or engagements to your account or Tweets).

Please see our support articles on [following rules and best practices](#) and [automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors. Accounts created to replace suspended accounts may be permanently suspended.

Content Visibility

Accounts under investigation or which have been detected as sharing content in violation of these Rules may have their account or Tweet visibility limited in various parts of Twitter, including search. To learn more about situations in which content may be restricted on Twitter, please see our support article on [search rules and restrictions](#).

Trademark policy

Updated June 30, 2016

What is a Trademark Policy Violation on Twitter?

Using another's trademark in a manner that may mislead or confuse others about your brand affiliation may be a violation of our trademark policy.

How Does Twitter Respond to Trademark Reports?

We review each report and do the following:

- If we determine that there is a violation of our trademark policy, we may suspend the account. In some instances, we may give the account holder an opportunity to comply with our policies before suspending the account. For more information, see our [Parody, commentary, and fan account policy](#).
- We may also provide the account holder with your name and information included in the copy of the report.
- If you would like to report a trademark violation in a [Twitter Ad](#), please visit our [Trademark Policy for Promoted Ads](#) to learn more.

Please note that Twitter only investigates requests that are submitted by the trademark holder or by an authorized representative of the trademark holder.

What is not a Trademark Policy Violation?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples include:

- Using a trademark in a way that is outside the scope of the trademark registration (e.g. territory, or goods and services identified in the registration).
- Nominative and other fair uses of trademarks are protected uses under our trademark policy, so long as the account is clearly distinguished from the trademark owner. This includes use by resellers in certain regions and accounts engaging in parody, commentary, or news. For more information, see our [Parody, commentary, and fan account policy](#).

How do I Report a Trademark Policy Violation?

Submit trademark reports through our [Trademark Issue support form](#). Using the appropriate form ensures we have all the information necessary to process your report quickly. Please provide all information requested in the form. If you submit a report with incomplete information, we will follow up with you to request that information. Please note that this will result in a delay in processing your report.

If you have trouble using the form, try updating your browser or switching to another browser.

Copyright policy

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

<https://web.archive.org/web/20180601194918/https://help.twitter.com/en/rules-and-policies/copyright-policy> 1/3

27.11.2020 Copyright policy

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

<https://web.archive.org/web/20180601194918/https://help.twitter.com/en/rules-and-policies/copyright-policy> 2/3

27.11.2020 Copyright policy

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);

2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Twitter media policy

Twitter Rules: Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images.

Twitter Rules: You may not use hateful images or symbols in your profile image or profile header.

Rationale

Some forms of graphic violence, adult content, or hateful imagery may be permitted in Tweets when they are marked as [sensitive media](#). However, you may not include this type of content in live video, or in profile or header images.

While we want people to feel free to share media that reflects their creativity or individuality, or to show what's happening in the world, we will take action when it crosses the line into abuse towards a person, group, or protected category.

Finally, we may sometimes require you to remove media containing excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about how to make such a request [here](#).

When this applies

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. Some examples include, but are not limited to, depictions of:

- the moment at which someone dies
- gruesome crime or accident scenes
- bodily harm, torture, dismemberment, or mutilation

We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity (including close-ups of genitals, buttocks, or breasts)
 - Please note that exceptions may be made for artistic, medical, health, or educational content. Breastfeeding content does not need to be marked as sensitive.
- simulating a sexual act
- intercourse or any sexual act (may involve humans, humanoid animals, cartoons, or anime)

If you see media you think may be in violation of our policy, report it. You can read about how to report sensitive media [here](#).

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, or ethnicity/national origin.

Some examples of hateful imagery include, but are not limited to:

- Symbols historically associated with hate groups (for example, the Nazi swastika)
- Images depicting others as less than human or altered to include hateful symbols
- Altered image references to a mass murder that targeted a protected category

Do I need to be the target of this content in order for it to be a violation of the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When adult content, graphic violence, or hateful imagery appears in Tweets, we may place this content behind an interstitial advising viewers to be aware that they will see sensitive media if they click through. This allows us to identify potentially sensitive content that some people may not wish to see. Learn more about [how to control whether you see sensitive media](#).

However, if such content appears in live video, header, or profile images, it will violate our media policy. The consequences for violating this policy depend on the severity of the violation and the person's previous record of violations.

The first time someone violates this policy they will be required to remove the imagery and may be asked to verify their account ownership. Subsequent violations could result in permanent suspension. Learn more about [our range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

How to contact Twitter about media concerning a deceased family member

In order to respect the wishes of loved ones, Twitter will remove imagery of deceased individuals in certain circumstances. Immediate family members and other authorized individuals may request the removal of images or video of deceased individuals, from when critical injury occurs to the moments before or after death, by sending a request to Twitter Inc. via our [privacy form](#). When reviewing such media removal requests, Twitter considers public interest factors such as the newsworthiness of the content and may not be able to honor every request.

Note: Find information about contacting Twitter to [deactivate the account of a deceased or incapacitated user](#).

How to mark media as sensitive links to Media policy

Twitter trends FAQs

Where can I find trends?

Trends are available on the Twitter for iOS app, the Twitter for Android app, and [twitter.com](#).

On Twitter's mobile apps, you can find trends listed under the trends section of the **Explore** tab. When signed in to [twitter.com](#) on a desktop or laptop computer, trends are listed in many places, including the Home timeline, Notifications, search results, and profile pages.

How are trends determined?

Trends are determined by an algorithm and, by default, are tailored for you based on who you follow, your interests, and your location. This algorithm identifies topics that are popular now, rather than topics that have been popular for a while or on a daily basis, to help you discover the hottest emerging topics of discussion on Twitter.

You can choose to see trends that are not tailored for you by selecting a specific trends location on [twitter.com](#), iOS, or Android

(instructions below). Location trends identify popular topics among people in a specific geographic location.

Note: The number of Tweets that are related to the trends is just one of the factors the algorithm looks at when ranking and determining trends. Algorithmically, trends and hashtags are grouped together if they are related to the same topic. For instance, #MondayMotivation and #MotivationMonday may both be represented by #MondayMotivation.

What does the # sign mean?

You may notice that some trends have # sign before the word or phrase. This is called a **hashtag** and is included specifically in Tweets to mark them as relating to a topic, so that people can follow the conversation in search.

Will I see world and local events in my trends?

Yes, when signed into your account on the twitter.com, iOS, or Android, trends will be tailored for you based on your location and who you follow. There will be many world and local news events and conversations that will appear in your trends regardless of your personalization.

Can I see trends for a specific location?

Yes, you can view trends for a specific location on twitter.com, iOS, or Android.

View trends via twitter.com:

To view trends for a specific location, click **Change** in your trends box. You can choose to **Keep tailored trends** based on your location and who you follow, or you can select **Change** to choose a nearby location or search locations.

Adjust your trends settings via Twitter for iOS: o.

In the top menu, tap your profile icon.

17. Tap **Settings and privacy**, then tap **Content preferences**.

18. Under **Explore**, tap **Trends**.

T. Drag the slider next to **Trends for you** to turn on and receive personalized trends.

V. When the feature is disabled, you can change your location by tapping **Change location**. In the search box, type your desired trends location or choose from the list.

Adjust your trends settings via Twitter for Android:

O. In the top menu, you will either see a navigation menu icon or your **profile**

icon. Tap whichever icon you have, and select **Settings and privacy**.

17. Tap **Content preferences**.

18. Under **Explore**, tap **Trends**.

T. Drag the slider next to **Trends for you** to turn on and receive personalized trends.

V. When the feature is disabled, you can change your location by tapping **Change location**. In the search box, type your desired trends location or choose from the list.

Note: If you don't find your preferred country or city, it means we're not receiving enough Tweets from that geographical area to create a list. You can look up local Tweets on any topic by using [advanced Twitter search operators](#).

What happens when I click or tap on a trend?

Clicking or tapping any of the trends takes you to the Twitter search results for that trend. You'll see all Tweets including that phrase or hashtag. To see what people are saying about a previous trend, perform a [search for that keyword](#).

How can I participate in a trend?

Simply post a Tweet including the exact word or phrase as it appears in the trends list (with the hashtag, if you see one). Due to the large number of people Tweeting about these specific trends, you may not always be able to find your particular Tweet in search, but your followers will always see your Tweets.

Note: Twitter also filters searches for quality. To understand what types of behavior could cause your Tweets to be filtered out, read about our [search rules and restrictions](#).

Are there rules for trends?

Yes. We want trends to promote healthy discussions on Twitter. This means that at times, we may prevent certain content from trending. These include trends that:

Contain profanity or adult/graphic references.

Incite hate on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease.

Violate the [Twitter Rules](#).

In some cases, we may also consider the newsworthiness of the content, or if it is in the public interest when evaluating potential violations. In these cases, the content might continue to trend on our platform.

Note that even if we prevent the hashtag or the content from appearing on the list of trending topics, you may still be able to access conversations around that content on Twitter.

Do trends have descriptions?

Below some trends, you may see the approximate number of Tweets associated with the trend, or personalized context such as who in your network is Tweeting about it.

Additionally, when a Moment is found to be relevant to a trend, we may use the title of the Moment as the trend's description.

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Child sexual exploitation policy

Twitter rules

Twitter does not tolerate any material that features or promotes child sexual exploitation. This may include media, text, illustrated, or computer generated images.

Rationale

Regardless of the intent, viewing, sharing, or linking to child sexual exploitation media contributes to the re-victimization of the depicted minor and, therefore, is not allowed on Twitter. This also applies to content that may further contribute to victimization of children through promotion or glorification of child sexual exploitation.

When this applies

For the purposes of this policy, a minor means any person under the age of 18 years. Examples of content that depicts or promotes child sexual exploitation include, but are not limited to:

Visual depiction of a minor engaging in sexually explicit or sexually suggestive act.

Illustrated, computer-generated or other forms of realistic depictions of a human minor in a sexually explicit context, or engaged in a sexually explicit act.

Links to third-party sites that host child sexual exploitation material.

The following behaviors are also not allowed on Twitter: Fantasizing about or promoting engagement in child sexual exploitation.

Expressing a desire to obtain materials that feature child sexual exploitation.

Recruiting, advertising or expressing interest in: (1) a commercial sex act involving a child, or (2) in harboring and/or transporting a child for sexual purposes.

Sending sexually explicit images to a child. Engaging a minor in a sexually explicit conversation.

Trying to obtain sexually explicit images from a minor or trying to recruit a minor for sexual services through incentive or blackmail.

Identifying alleged victims of childhood sexual exploitation by name or photo.

Consequences

In the majority of cases, the consequence for violating our child sexual exploitation policy is immediate and permanent suspension. In addition, violators will not be allowed to create new accounts in the future.

In a limited number of situations when no malicious intent has been identified, users may instead be asked to remove the violating Tweet and be temporarily locked out of their account before they can Tweet again.

When we are made aware of links depicting or promoting child sexual exploitation they will be removed from the site without further notice and reported to [The National Center for Missing & Exploited Children](#) (NCMEC).

What you can do

If you think you've found a Twitter account distributing or promoting child sexual exploitation, please report it through our designated [form](#). Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. For help on how to find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with NCMEC [here](#) or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the International Association of Internet Hotlines [website](#).

Our partners

We partner with multiple organisations around the world whose work is dedicated to fighting child sexual exploitation. You can find a full list of our safety partners [here](#).

Hateful conduct policy

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple

underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you".

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., "I hope you get cancer and die."

Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”

Saying that a group of individuals deserve serious physical injury, e.g., “If this group of protesters don’t shut up, they deserve to be shot.”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct,

as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can

[submit an appeal](#).

Report private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Read about our [policy regarding private information posted on Twitter](#).

Do I need to have a Twitter account to make a report?

No, you do not need to have a Twitter account to file a report that your private and confidential information has been posted on Twitter.

Can I report unauthorized posting of Copyrighted material?

If your copyrighted material has been Tweeted without your permission, please review our [Copyright and DMCA policy](#).

What can I do if my information is posted on other websites?

We cannot take down any information that is posted on a website other than Twitter. If your information is posted on another website, please contact the help center or support staff for the website where your information is posted.

How do I report that my private information is posted on Twitter?

You can report that your private information has been posted on Twitter directly from the offending Tweet(s). Learn [how to report individual Tweets](#).

File a [report that someone has posted your private information](#).

What happens when Twitter receives a valid report?

When we receive a complete and valid report that private information has been posted on Twitter, we'll investigate the account and Tweets reported. We will review where, if anywhere, the information has been made publicly available before taking action on the account or Tweets. If the information you reported was previously posted elsewhere on the Internet, it is not a violation of our policy and we will not take action.

Learn more about [protecting your private information on Twitter and other websites](#).

Read our [policy regarding private information posted on Twitter](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

The user shares your name but has no other commonalities, or

The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation rules

Updated November 3, 2017

This page is primarily intended for **developers**.

For Twitter users: You are ultimately responsible for the actions taken with your account, or by applications associated with your account. Before authorizing a third-party application to access or use your account, make sure you've thoroughly investigated the application and understand what it will do. If automated activity on your account violates the [Twitter Rules](#) or these automation rules, Twitter may take action on your

account, including [filtering your Tweets from search results](#) or suspending your account.

For more information on third-party applications, please see our article on [connecting and revoking third-party applications](#).

If you are a developer and have questions about these Automation Rules, please reach out via <https://twittercommunity.com/c/rules-and-policies>.

If your application has been suspended or restricted from performing write actions, you may submit an appeal at </forms/platform>.

I. Ground Rules Do!

Build solutions that automatically broadcast helpful information in Tweets.

Run creative campaigns that auto-reply to users who engage with your content.

Build solutions that automatically respond to users in Direct Messages.

Try new things that help people (and comply with our rules).

Make sure your application provides a good user experience and performs well — and confirm that remains the case over time.

Don't!

Violate these or other policies. Be extra mindful of our rules about abuse and user privacy.

Abuse the Twitter API or attempt to circumvent rate limits.

Use non-API-based forms of automation, such as scripting the Twitter website. The use of these techniques may result in the permanent suspension of your account.

Spam or bother users, or otherwise send them unsolicited messages.

A. The Twitter Rules and the Developer Agreement and Policy

As with all activity on Twitter, automated activity is subject to the [Twitter Rules](#) and, if you're a developer using the Twitter API, the [Developer Agreement and Policy](#).

You should carefully review these policies to ensure that your automated activity is compliant. Automated applications or activities that violate these policies, or that facilitate or induce users to violate them, may be subject to enforcement action,

potentially including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these policies.

Although all aspects of the Twitter Rules and the Developer Agreement and Policy apply to automated activity, you should keep the following rules top of mind:

Spamming: You may not send automated Tweets or Direct Messages that are spam, or otherwise engage in spamming activity. Some examples of spammy behavior to avoid with automation include:

Trending topics: You may not automatically post about trending topics on Twitter, or use automation to attempt to influence or manipulate trending topics.

Multiple posts/accounts: You may not post duplicative or substantially similar Tweets on one account or over multiple accounts you operate.

Duplicate accounts: You may not create and/or automate multiple accounts for duplicative or substantially similar use cases.

However, automating multiple accounts for related but non-duplicative use cases is permitted. For example, you may automate separate accounts to Tweet when the Hubble Space Telescope passes over different cities, such as [San Francisco](#) or [Hong Kong](#).

Misleading links: You may not send automated Tweets or Direct Messages containing links that are misleading, including links that maliciously or deceptively redirect through landing pages or ad pages before displaying the final content.

Sensitive media: Automated Tweets and Direct Messages must comply with the [Twitter media policy](#), and you should mark your account as potentially sensitive if you intend to post graphic, pornographic, or potentially sensitive media.

Abusive behavior: You may not engage in any automated activity that encourages, promotes, or incites abuse, violence, hateful conduct, or harassment, on or off Twitter.

Private information: You may not post private or confidential information about a person without their prior express authorization.

B. Other Ground Rules for Automated Activity

In addition to the policies above, the following ground rules apply to all automated activity on Twitter:

Don't surprise or mislead users: Automated activity should honor users' expectations. Ask for the user's permission before taking an action if you aren't sure.

Mature content or profanity: Don't Direct Message, mention, or reply to users with potentially sensitive content (including profanity), unless they've clearly indicated an intent to receive it in advance.

Be thoughtful about the information you request or exchange on Twitter

Tweets: Don't ask users to send you personal or private information via a public Tweet. If you need additional personal or private information from a user to provide them with customer service (or other similar use cases), you should ask the user to share such information by Direct Message or another private channel. You might even consider adding a [Direct Message deep link](#) to your Tweet.

Direct Messages: You should only ask users for the minimum amount of information you need to provide them with service. If you need to request or exchange particularly sensitive information (such as credit card information), you should consider directing users to your website or other appropriate channel to do so.

II. Activity-Specific Rules

The activity-specific rules in this section apply to taking specific automated actions on Twitter. Please read these rules carefully, as they outline both permitted and prohibited use cases of automation.

Automated applications or activities that violate these rules, or that facilitate or induce users to violate them, may be subject to enforcement action, including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these rules. As a reminder, you should also carefully review the spam guidelines in the [Twitter Rules](#) to avoid having activities performed by you, your app, or other users through your app or service flagged as spam.

A. Automated Actions Through Another User's Account

Twitter users may authorize your app or service to [access their Twitter account through OAuth](#). A user authorizing your app or service to access their Twitter account through OAuth does not by itself constitute sufficient consent to take automated actions through that user's account.

You may only take automated actions through another Twitter user's account if you:

clearly describe to the user the types of automated actions that will occur; receive express consent from the user to take those automated actions; and immediately honor a user's request to opt-out of further automated actions.

If you substantially change the purpose or functionality of your app or service, you must re-obtain express consent from the user to take automated action through their account before doing so.

These requirements apply to any automated action taken through another Twitter user's account, including posting Tweets, sending Direct Messages, deleting Tweets or Direct Messages, or following/unfollowing other accounts. For applications that offer users the ability to delete Tweets in a bulk or automated manner, you must also clearly state that Tweets are not recoverable once deleted.

B. Automated Tweets

1. Posting automated Tweets

Automated Tweets that cross-post outside information:

You may post automated Tweets based on sources of outside information — such as an RSS feed, weather data, etc. — as long as you are sufficiently authorized to publish such information.

Other automated Tweets (excluding mentions or replies):

Provided you comply with all other rules, you may post automated Tweets for entertainment, informational, or novelty

purposes. As a reminder, accounts posting duplicative, spammy, or otherwise prohibited content may be subject to suspension.

2. Posting automated mentions and replies

The reply and mention functions are intended to make communication between Twitter users easier. Automating these actions to reach many users on an unsolicited basis is an abuse of the feature, and is not permitted. For example, sending automated replies to Tweets based on keyword searches alone is not permitted. Spamy or duplicative use of mentions and replies may result in enforcement action, such as the removal of your Tweets from Search or the suspension of your app or account.

However, you may send automated replies or mentions to Twitter users so long as:

in advance of sending the automated reply, the recipient or mentioned user(s) have requested or have clearly indicated an intent on Twitter to be contacted by you (i.e. opted in), for example by replying to a Tweet from your account, or by sending you a Direct Message;

you provide a clear and easy way for such users to opt-out of receiving automated replies and mentions, and promptly honor all such opt-out requests;

you only send one automated reply or mention per user interaction; and

the automated reply or mention is a reply to the user's original Tweet (if your campaign is based on users posting a reply to your Tweet).

Opt-in techniques and indications of user intent take many different forms, depending on the specifics of your use case and implementation. Some examples include:

A Tweet from your account that clearly indicates that a user taking a specific action on that Tweet (such as Retweeting it) will opt the user into receiving an automated response.

A mention of your account by the user in a manner suggesting the user clearly wishes or intends to receive a response. If you want to run an auto-reply campaign with a campaign- or use-case-specific hashtag, users should also mention you in their Tweets.

Note that a user following your account is not on its own a sufficient indication of user intent to receive an automated response.

: Advertisers, publishers, and brands using auto-response campaigns must request approval from Twitter and may be subject to additional rules. Please reach out to your account or partner manager for assistance. If you do not have a partner manager but are interested in managed client services, you can learn more [here](#).

Additionally, we recommend that any accounts that will communicate with

users via automated mentions or replies:

Note

1. Appropriately filter responses based on potentially sensitive language in user handles,

display names, and Tweet text, as well as potentially sensitive media;

2. Check that the Tweet you are mentioning or replying to still exists (for example, using the statuses/lookup endpoint on the Twitter API).

C. Automated Direct Messages

Sending automated Direct Messages to users

You may not send unsolicited Direct Messages in a bulk or automated manner, and should be thoughtful about the frequency with which you contact users via Direct Message.

You may send automated Direct Messages to users so long as:

in advance of sending the Direct Message, the recipient(s) have requested or have clearly indicated an intent on Twitter to be contacted by you via Direct Message, for example by sending you a Direct Message; **and**

you provide a clear and easy way for such users to opt-out of receiving automated Direct Messages, and promptly honor all such opt-out requests.

The fact that a user is technically able to receive a Direct Message from you (e.g. because the user follows you, has enabled the ability to receive Direct Messages from any

account, or because the user is in a pre-existing Direct Message conversation with you) does not necessarily mean they have requested or expect to receive automated Direct Messages from you.

Interacting with users via Direct Message

Per the Ground Rules, remember to be thoughtful about the amount and type of information you request or exchange with users via Direct Messages. If you will be asking a user to provide personal or private information via an automated Direct Message, you must clearly explain how you will use the information you're collecting. Consider including a link to your privacy policy in your Direct Message to the user, as well as in your Twitter profile bio.

Don't publicly share information received in a Direct Message conversation with a user without first obtaining explicit consent from the user. For example, if a user asks you via Direct Message about a purchase they made from you, you may not mention the user in a public Tweet that includes information about their purchase unless you have the user's explicit consent to do so.

After a user-initiated interaction ends, don't send additional follow-up Direct Messages or mention users in a Tweet unless you get permission from the user.

D. Automated actions you take on Tweets or accounts

Automated likes: You may not like Tweets in an automated manner.

Automated Retweets: Provided you comply with all other rules, you may Retweet or Quote Tweet in an automated manner for entertainment, informational, or novelty purposes. Automated Retweets often lead to negative user experiences, and bulk, aggressive, or spammy Retweeting is a violation of the [Twitter Rules](#).

Automated following/unfollowing: You may not follow or unfollow Twitter accounts in a bulk, aggressive, or indiscriminate manner. Aggressive following is a violation of the [Twitter Rules](#). Please also review our [following rules and best practices](#) to ensure you are in compliance. Note that applications that claim to get users more followers are also prohibited under the [Twitter Rules](#).

Automated adding to lists or collections: You may not add Twitter users to lists or add Tweets to collections in a bulk or indiscriminate manner. Adding a large number of unrelated users to lists is a violation of the [Twitter Rules](#).

About search rules and restrictions

This article details the rules for using Twitter search, and explains reasons why content may not show up in Twitter search.

Why doesn't all content show up in Twitter search?

In order to keep your search results relevant, Twitter filters search results for quality Tweets and accounts. Material that jeopardizes search quality or creates a bad search experience for other people may be automatically removed from Twitter search. Read more about this [here](#).

General rules for using Twitter search

The [Twitter Rules](#) explain what behaviors are permitted on Twitter. In addition to these rules, we've included some tips below to keep your content relevant (and your Tweets in Twitter search).

Please do not:

- Repeatedly post duplicate or near-duplicate content (links or Tweets).
- Abuse trending topics or hashtags (topic words with a # sign).
- Send automated Tweets or replies.
- Use bots or applications to post similar messages based on keywords.
- Post similar messages over multiple accounts.
- Aggressively follow and unfollow people.

Twitter may automatically remove accounts engaging in these behaviors from search (or even suspend in some cases) in order to ensure the best experience for everyone.

Search results filtering

To learn how to filter your search results, including how to disable or enable safe search, read our [Using Twitter search](#) article.

Search results filtering in third-party applications

Accounts may be filtered from search for posting similar messages over several accounts. Your account may be filtered from search if you are logged into a third-party application that is updating many accounts with similar or duplicate content.

Rules for participating in contests on Twitter

Businesses sometimes host contests and offer prizes to Retweet an update, following a particular account, or posting Tweets with a specific hashtag topic or reply. If you're running a contest, check out this [article](#) for guidelines for contests on Twitter.

If you repeatedly post near-duplicate updates or duplicate links in order to enter contests, you may be filtered from search. Please keep in mind that you are responsible for the updates that you post to your Twitter account.

What to do if you are missing from search

Please see this [article](#) for more information related to this issue.

If you are being filtered from search, rest assured that your followers will still receive your updates. If you post a reply or mention to a specific person, it will be delivered to that person.

Having trouble?

You can learn more about searching by keyword, location, and trending topics in our [Using Twitter search](#) help article.

The Twitter Rules

We believe that everyone should have the power to create and share ideas and information instantly, without barriers. In order to protect the experience and safety of people who use Twitter, there are some limitations on the type of content and behavior that we allow. These limitations are set forth in the Twitter Rules below.

The Twitter Rules (along with all incorporated policies), [Privacy Policy](#), and [Terms of Service](#) collectively make up the "Twitter User Agreement" that governs a user's access to and use of Twitter's services.

All individuals accessing or using Twitter's services must adhere to the policies set forth in the Twitter Rules. Failure to do so may result in Twitter taking one or more of the following enforcement actions:

- requiring you to remove prohibited content before you can again create new posts and interact with other Twitter users;
- temporarily limiting your ability to create posts or interact with other Twitter users;
- asking you to verify account ownership with a phone number or email address; or
- permanently suspending your account(s).

If you attempt to evade a permanent suspension by creating new accounts, we will suspend your new accounts.

Please note that we may need to change these Rules from time to time and reserve the right to do so. The most current version is always available at: <https://twitter.com/rules>.

The policies set forth in these Twitter Rules govern organic content on our platform. To learn more about the rules which govern ads and promoted content, please review our [Ads policies](#).

Content Boundaries and Use of Twitter

Intellectual property

Trademark: We reserve the right to suspend accounts or take other appropriate action when someone's brand or trademark, including business name and/or logo, is used in a manner that may mislead or confuse others about your brand affiliation. Read more about our [trademark policy and how to report a violation](#).

Copyright: We will respond to clear and complete notices of alleged copyright infringement. Our copyright procedures are set forth in our [Terms of Service](#). Read more about our [copyright policy](#).

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal.

Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in live video, your profile, or header images. Additionally, Twitter may sometimes require you to remove excessively graphic violence. [Read more](#) about how we define graphic violence and adult content and [how to mark your media as sensitive](#).

Media depicting deceased individuals: We may require you to remove media that depicts the death of an identifiable individual if we receive a request from their family or an authorized representative. Learn more about [how to make such a request](#).

Unlawful use

You may not use our service for any unlawful purposes or in furtherance of illegal activities. By using Twitter, you agree to comply with all applicable laws governing your online conduct and content.

Distribution of hacked materials

We do not permit the use of our services to directly distribute content obtained through hacking that contains personally identifiable information, may put people in imminent harm or danger, or contains trade secrets. Direct distribution of hacked materials includes posting hacked content on Twitter (for instance, in the text of a Tweet, or in an image), or directly linking to hacked content hosted on other websites.

We may suspend accounts which directly distribute hacked materials where the account has made a claim of responsibility for a hack, or where Twitter is able to reliably attribute a hack to the account distributing that content.

Trends

At times, we may prevent certain content from trending. This includes content that violates the Twitter Rules, as well as content that may attempt to manipulate trends. Read more about [what we allow and do not allow to trend](#).

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Misuse of Twitter badges

You may not use badges, including but not limited to the “promoted” or “verified” Twitter badges, unless provided by Twitter. Accounts using unauthorized badges as part of their profile photos, header photos, display names, or in any way that falsely implies affiliation with Twitter or authorization from Twitter to display these badges, may be suspended.

Misuse of usernames

Selling usernames: You may not buy or sell Twitter usernames.

Username squatting: You may not engage in username squatting. Some of the factors we take into consideration when determining whether conduct is username squatting include:

- the number of accounts created;
- the creation of accounts for the purpose of preventing others from using those account names;
- the creation of accounts for the purpose of selling those accounts; and
- the use of third-party content feeds to update and maintain accounts under the names of those third parties.

Please note that Twitter may also remove accounts that are inactive for more than six months. Learn more about [username squatting](#).

Abusive Behavior

We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up. In order to ensure that people feel safe expressing diverse opinions and beliefs, we prohibit behavior that crosses the line into abuse, including behavior that harasses, intimidates, or uses fear to silence another user's voice.

Context matters when evaluating for abusive behavior and determining appropriate enforcement actions. Factors we may take into consideration include, but are not limited to whether:

- the behavior is targeted at an individual or group of people;
- the report has been filed by the target of the abuse or a bystander;
- the behavior is newsworthy and in the legitimate public interest.

Violence and physical harm

Violence: You may not make specific threats of violence or wish for the serious physical harm, death, or disease of an individual or group of people. This includes, but is not limited to, threatening or promoting terrorism. You also may not affiliate with organizations that — whether by their own statements or activity both on and off the platform — use or promote violence against civilians to further their causes.

Suicide or self-harm: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Child sexual exploitation: You may not promote child sexual exploitation. Learn more about our zero-tolerance [child sexual exploitation policy](#).

Abuse and hateful conduct

Abuse: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Unwanted sexual advances: You may not direct abuse at someone by sending unwanted sexual content, objectifying them in a sexually explicit manner, or otherwise engaging in sexual misconduct.

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. Read more about our [hateful conduct policy](#).

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Private information and intimate media

Private information: You may not publish or post other people's private information without their express authorization and permission. Definitions of private information may vary depending on local laws. Read more about our [private information policy](#).

Intimate media: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. Media depicting sexual violence and/or assault is also not permitted. Note:

limited exceptions may apply if there is clear context that the interaction is consensual. Read more [about intimate media](#) on Twitter.

Threats to expose / hack: You may not threaten to expose someone's private information or intimate media. You also may not threaten to hack or break into someone's digital information or attempt to incentivize others to do so (e.g., through setting a bounty or reward on such actions).

Impersonation

You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. While you may maintain parody, fan, commentary, or newsfeed accounts, you may not do so if the intent of the account is to engage in spamming or abusive behavior. Read more about our [impersonation policy](#).

Spam and Security

We strive to protect people on Twitter from technical abuse and spam.

To promote a stable and secure environment on Twitter, you may not do, or attempt to do, any of the following while accessing or using Twitter:

- Access, tamper with, or use non-public areas of Twitter, Twitter's computer systems, or the technical delivery systems of Twitter's providers (except as expressly permitted by the Twitter Bug Bounty program).
- Probe, scan, or test the vulnerability of any system or network, or breach or circumvent any security or authentication measures (except as expressly permitted by the Twitter Bug Bounty program).
- Access or search, or attempt to access or search, Twitter by any means (automated or otherwise) other than through our currently available, published interfaces that are provided by Twitter (and only pursuant to the applicable terms and conditions), unless you have been specifically allowed to do so in a separate agreement with Twitter. Note that crawling Twitter is permissible if done in accordance with the provisions of the robots.txt file; however, scraping Twitter without our prior consent is expressly prohibited.
- Forge any TCP/IP packet header or any part of the header information in any email or posting, or in any way use Twitter to send altered, deceptive, or false source-identifying information.
- Interfere with or disrupt the access of any user, host or network, including, without limitation, sending a virus, overloading, flooding, spamming, mail-bombing Twitter's services, or by scripting the creation of content in such a manner as to interfere with or create an undue burden on Twitter.

Any accounts engaging in the following activities may be temporarily locked or subject to permanent suspension:

- **Malware/Phishing:** You may not publish or link to malicious content intended to damage or disrupt another person's browser or computer or to compromise a person's privacy.
- **Fake accounts:** You may not register or create fake and misleading accounts. While you may use Twitter pseudonymously or as a [parody, commentary, or fan account](#), you may not use misleading account information in order to engage in spamming, abusive, or disruptive behavior, including attempts to manipulate the conversations on Twitter. Some of the factors that we take into account when determining whether an account is fake include:
 - Use of stock or stolen avatar photos
 - Use of stolen or copied profile bios
 - Use of intentionally misleading profile information, including profile location
- **Spam:** You may not use Twitter's services for the purpose of spamming anyone. Spam is generally defined on Twitter as bulk or aggressive activity that attempts to manipulate or disrupt Twitter or the experience of users on Twitter to drive traffic or attention to unrelated accounts, products, services, or initiatives. Some of the factors that we take into account when determining what conduct is considered to be spamming include:

- if you have followed and/or unfollowed a large number of accounts in a short time period, particularly by automated means (aggressive following or follower churn);
- if your Tweets or Direct Messages consist mainly of links shared without commentary;
- if a large number of people have blocked you in response to high volumes of untargeted, unsolicited, or duplicative content or engagements from your account;
- if a large number of spam complaints have been filed against you;
- if you post duplicative or substantially similar content, replies, or mentions over multiple accounts or multiple duplicate updates on one account, or create duplicate or substantially similar accounts;
- if you post multiple updates to a trending or popular topic with an intent to subvert or manipulate the topic to drive traffic or attention to unrelated accounts, products, services, or initiatives;
- if you send large numbers of unsolicited replies or mentions;
- if you add users to lists in a bulk or aggressive manner;
- if you are randomly or aggressively engaging with Tweets (e.g., likes, Retweets, etc.) or users (e.g., following, adding to lists or Moments, etc.) to drive traffic or attention to unrelated accounts, products, services, or initiatives;
- if you repeatedly post other people's account information as your own (e.g., bio, Tweets, profile URL, etc.);
- if you post misleading, deceptive, or malicious links (e.g., affiliate links, links to malware/clickjacking pages, etc.);
- if you sell, purchase, or attempt to artificially inflate account interactions (such as followers, Retweets, likes, etc.); and
- if you use or promote third-party services or apps that claim to get you more followers, Retweets, or likes (such as follower trains, sites promising "more followers fast", or any other site that offers to automatically add followers or engagements to your account or Tweets).

Please see our support articles on [following rules and best practices](#) and [automation rules and best practices](#) for more detailed information about how the Rules apply to those particular account behaviors.

Accounts created to replace or mimic suspended accounts may be permanently suspended. We may also remove accounts which Twitter is able to reliably attribute to entities known to violate the Twitter Rules.

Content Visibility

Accounts under investigation or which have been detected as sharing content in violation of these Rules may have their account or Tweet visibility limited in various parts of Twitter, including search. To learn more about situations in which content may be restricted on Twitter, please see our support article on [search rules and restrictions](#).

Trademark policy

Updated June 30, 2016

What is a Trademark Policy Violation on Twitter?

Using another's trademark in a manner that may mislead or confuse others about your brand affiliation may be a violation of our trademark policy.

How Does Twitter Respond to Trademark Reports?

We review each report and do the following:

- If we determine that there is a violation of our trademark policy, we may suspend the account. In some instances, we may give the account holder an opportunity to comply with our policies before suspending the account. For more information, see our [Parody, commentary, and fan account policy](#).
- We may also provide the account holder with your name and information included in the copy of the report.
- If you would like to report a trademark violation in a [Twitter Ad](#), please visit our [Trademark Policy for Promoted Ads](#) to learn more.

Please note that Twitter only investigates requests that are submitted by the trademark holder or by an authorized representative of the trademark holder.

What is not a Trademark Policy Violation?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples include:

- Using a trademark in a way that is outside the scope of the trademark registration (e.g. territory, or goods and services identified in the registration).
- Nominative and other fair uses of trademarks are protected uses under our trademark policy, so long as the account is clearly distinguished from the trademark owner. This includes use by resellers in certain regions and accounts engaging in parody, commentary, or news. For more information, see our [Parody, commentary, and fan account policy](#).

How do I Report a Trademark Policy Violation?

Submit trademark reports through our [Trademark Issue support form](#). Using the appropriate form ensures we have all the information necessary to process your report quickly. Please provide all information requested in the form. If you submit a report with incomplete information, we will follow up with you to request that information. Please note that this will result in a delay in processing your report.

If you have trouble using the form, try updating your browser or switching to another browser.

Copyright policy

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including

<http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

<https://web.archive.org/web/20180601194918/https://help.twitter.com/en/rules-and-policies/copyright-policy> 1/3

27.11.2020 Copyright policy

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

<https://web.archive.org/web/20180601194918/https://help.twitter.com/en/rules-and-policies/copyright-policy> 2/3

27.11.2020 Copyright policy

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);

2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Twitter media policy

Twitter Rules: Twitter allows some forms of graphic violence and/or adult content in Tweets marked as containing sensitive media. However, you may not use such content in your profile or header images.

Twitter Rules: You may not use hateful images or symbols in your profile image or profile header.

Rationale

Some forms of graphic violence, adult content, or hateful imagery may be permitted in Tweets when they are marked as [sensitive media](#). However, you may not include this type of content in live video, or in profile or header images.

While we want people to feel free to share media that reflects their creativity or individuality, or to show what's happening in the world, we will take action when it crosses the line into abuse towards a person, group, or protected category.

Finally, we may sometimes require you to remove media containing excessively graphic violence out of respect for the deceased and their families if we receive a request from their family or an authorized representative. Learn more about how to make such a request [here](#).

When this applies

Graphic violence and adult content

We consider graphic violence to be any form of gory media related to death, serious injury, violence, or surgical procedures. Some examples include, but are not limited to, depictions of:

- the moment at which someone dies
- gruesome crime or accident scenes
- bodily harm, torture, dismemberment, or mutilation

We consider adult content to be any media that is pornographic and/or may be intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity (including close-ups of genitals, buttocks, or breasts)
 - Please note that exceptions may be made for artistic, medical, health, or educational content. Breastfeeding content does not need to be marked as sensitive.
- simulating a sexual act
- intercourse or any sexual act (may involve humans, humanoid animals, cartoons, or anime)

If you see media you think may be in violation of our policy, report it. You can read about how to report sensitive media [here](#).

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, or ethnicity/national origin.

Some examples of hateful imagery include, but are not limited to:

- Symbols historically associated with hate groups (for example, the Nazi swastika)
- Images depicting others as less than human or altered to include hateful symbols
- Altered image references to a mass murder that targeted a protected category

Do I need to be the target of this content in order for it to be a violation of the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When adult content, graphic violence, or hateful imagery appears in Tweets, we may place this content behind an interstitial advising viewers to be aware that they will see sensitive media if they click through. This allows us to identify potentially sensitive content that some people may not wish to see. Learn more about [how to control whether you see sensitive media](#).

However, if such content appears in live video, header, or profile images, it will violate our media policy. The consequences for violating this policy depend on the severity of the violation and the person's previous record of violations.

The first time someone violates this policy they will be required to remove the imagery and may be asked to verify their account ownership. Subsequent violations could result in permanent suspension. Learn more about [our range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

How to contact Twitter about media concerning a deceased family member

In order to respect the wishes of loved ones, Twitter will remove imagery of deceased individuals in certain circumstances. Immediate family members and other authorized individuals may request the removal of images or video of deceased individuals, from when critical injury occurs to the moments before or after death, by sending a request to Twitter Inc. via our [privacy form](#). When reviewing such media removal requests, Twitter considers public interest factors such as the newsworthiness of the content and may not be able to honor every request. Note: Find information about contacting Twitter to [deactivate the account of a deceased or incapacitated user](#).

How to mark media as sensitive links to Media policy

Twitter trends FAQs

Where can I find trends?

Trends are available on the Twitter for iOS app, the Twitter for Android app, and [twitter.com](#).

On Twitter's mobile apps, you can find trends listed under the trends section of the **Explore** tab. When signed in to [twitter.com](#) on a desktop or laptop computer, trends are listed in many places, including the Home timeline, Notifications, search results, and profile pages.

How are trends determined?

Trends are determined by an algorithm and, by default, are tailored for you based on who you follow, your interests, and your location. This algorithm identifies topics that are popular now, rather than topics that have been popular for a while or on a daily basis, to help you discover the hottest emerging topics of discussion on Twitter.

You can choose to see trends that are not tailored for you by selecting a specific trends location on [twitter.com](#), iOS, or Android

(instructions below). Location trends identify popular topics among people in a specific geographic location.

Note: The number of Tweets that are related to the trends is just one of the factors the algorithm looks at when ranking and determining trends. Algorithmically, trends and hashtags are grouped together if they are related to the same topic. For instance, #MondayMotivation and #MotivationMonday may both be represented by #MondayMotivation.

What does the # sign mean?

You may notice that some trends have # sign before the word or phrase. This is called a **hashtag** and is included specifically in Tweets to mark them as relating to a topic, so that people can follow the conversation in search.

Will I see world and local events in my trends?

Yes, when signed into your account on the twitter.com, iOS, or Android, trends will be tailored for you based on your location and who you follow. There will be many world and local news events and conversations that will appear in your trends regardless of your personalization.

Can I see trends for a specific location?

Yes, you can view trends for a specific location on twitter.com, iOS, or Android.

View trends via twitter.com:

To view trends for a specific location, click **Change** in your trends box. You can choose to **Keep tailored trends** based on your location and who you follow, or you can select **Change** to choose a nearby location or search locations.

Adjust your trends settings via Twitter for iOS:

In the top menu, tap your profile icon.

17. Tap **Settings and privacy**, then tap **Content preferences**.

18. Under **Explore**, tap **Trends**.

T. Drag the slider next to **Trends for you** to turn on and receive personalized trends.

V. When the feature is disabled, you can change your location by tapping **Change location**. In the search box, type your desired trends location or choose from the list.

Adjust your trends settings via Twitter for Android:

O. In the top menu, you will either see a navigation menu icon or your **profile**

icon. Tap whichever icon you have, and select **Settings and privacy**.

17. Tap **Content preferences**.

18. Under **Explore**, tap **Trends**.

T. Drag the slider next to **Trends for you** to turn on and receive personalized trends.

V. When the feature is disabled, you can change your location by tapping **Change location**. In the search box, type your desired trends location or choose from the list.

Note: If you don't find your preferred country or city, it means we're not receiving enough Tweets from that geographical area to create a list. You can look up local Tweets on any topic by using [advanced Twitter search operators](#).

What happens when I click or tap on a trend?

Clicking or tapping any of the trends takes you to the Twitter search results for that trend. You'll see all Tweets including that phrase or hashtag. To see what people are saying about a previous trend, perform a [search for that keyword](#).

How can I participate in a trend?

Simply post a Tweet including the exact word or phrase as it appears in the trends list (with the hashtag, if you see one). Due to the large number of people Tweeting about these specific trends, you may not always be able to find your particular Tweet in search, but your followers will always see your Tweets.

Note: Twitter also filters searches for quality. To understand what types of behavior could cause your Tweets to be filtered out, read about our [search rules and restrictions](#).

Are there rules for trends?

Yes. We want trends to promote healthy discussions on Twitter. This means that at times, we may prevent certain content from trending. These include trends that:

Contain profanity or adult/graphic references.

Incite hate on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or disease.

Violate the [Twitter Rules](#).

In some cases, we may also consider the newsworthiness of the content, or if it is in the public interest when evaluating potential violations. In these cases, the content might continue to trend on our platform.

Note that even if we prevent the hashtag or the content from appearing on the list of trending topics, you may still be able to access conversations around that content on Twitter.

Do trends have descriptions?

Below some trends, you may see the approximate number of Tweets associated with the trend, or personalized context such as who in your network is Tweeting about it.

Additionally, when a Moment is found to be relevant to a trend, we may use the title of the Moment as the trend's description.

Username squatting policy

Username squatting is prohibited by [The Twitter Rules](#).

Please note that if an account has had no updates, no profile image, and there is no intent to mislead, it typically means there's no name-squatting or impersonation. Note that we will not release inactive or squatted usernames except in cases of trademark infringement. If your report involves [trademark infringement](#), please consult those policies for instructions for reporting these accounts.

Attempts to sell, buy, or solicit other forms of payment in exchange for usernames are also violations and may result in permanent account suspension.

Child sexual exploitation policy

Twitter rules

Twitter does not tolerate any material that features or promotes child sexual exploitation. This may include media, text, illustrated, or computer generated images.

Rationale

Regardless of the intent, viewing, sharing, or linking to child sexual exploitation media contributes to the re-victimization of the depicted minor and, therefore, is not allowed on Twitter. This also applies to content that may further contribute to victimization of children through promotion or glorification of child sexual exploitation.

When this applies

For the purposes of this policy, a minor means any person under the age of 18 years. Examples of content that depicts or promotes child sexual exploitation include, but are not limited to:

Visual depiction of a minor engaging in sexually explicit or sexually suggestive act.

Illustrated, computer-generated or other forms of realistic depictions of a human minor in a sexually explicit context, or engaged in a sexually explicit act.

Links to third-party sites that host child sexual exploitation material.

The following behaviors are also not allowed on Twitter: Fantasizing about or promoting engagement in child sexual exploitation.

Expressing a desire to obtain materials that feature child sexual exploitation.

Recruiting, advertising or expressing interest in: (1) a commercial sex act involving a child, or (2) in harboring and/or transporting a child for sexual purposes.

Sending sexually explicit images to a child. Engaging a minor in a sexually explicit conversation.

Trying to obtain sexually explicit images from a minor or trying to recruit a minor for sexual services through incentive or blackmail.

Identifying alleged victims of childhood sexual exploitation by name or photo.

Consequences

In the majority of cases, the consequence for violating our child sexual exploitation policy is immediate and permanent suspension. In addition, violators will not be allowed to create new accounts in the future.

In a limited number of situations when no malicious intent has been identified, users may instead be asked to remove the violating Tweet and be temporarily locked out of their account before they can Tweet again.

When we are made aware of links depicting or promoting child sexual exploitation they will be removed from the site without further notice and reported to [The National Center for Missing & Exploited Children](#) (NCMEC).

What you can do

If you think you've found a Twitter account distributing or promoting child sexual exploitation, please report it through our designated [form](#). Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. For help on how to find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with NCMEC [here](#) or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the International Association of Internet Hotlines [website](#).

Our partners

We partner with multiple organisations around the world whose work is dedicated to fighting child sexual exploitation. You can find a full list of our safety partners [here](#).

Hateful conduct policy

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple

underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you".

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., "I hope you get cancer and die."

Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”

Saying that a group of individuals deserve serious physical injury, e.g., “If this group of protesters don’t shut up, they deserve to be shot.”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct,

as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can

[submit an appeal](#).

Report private information posted on Twitter

Posting another person's private and confidential information is a violation of the [Twitter Rules](#). Read about our [policy regarding private information posted on Twitter](#).

Do I need to have a Twitter account to make a report?

No, you do not need to have a Twitter account to file a report that your private and confidential information has been posted on Twitter.

Can I report unauthorized posting of Copyrighted material?

If your copyrighted material has been Tweeted without your permission, please review our [Copyright and DMCA policy](#).

What can I do if my information is posted on other websites?

We cannot take down any information that is posted on a website other than Twitter. If your information is posted on another website, please contact the help center or support staff for the website where your information is posted.

How do I report that my private information is posted on Twitter?

You can report that your private information has been posted on Twitter directly from the offending Tweet(s). Learn [how to report individual Tweets](#).

File a [report that someone has posted your private information](#).

What happens when Twitter receives a valid report?

When we receive a complete and valid report that private information has been posted on Twitter, we'll investigate the account and Tweets reported. We will review where, if anywhere, the information has been made publicly available before taking action on the account or Tweets. If the information you reported was previously posted elsewhere on the Internet, it is not a violation of our policy and we will not take action.

Learn more about [protecting your private information on Twitter and other websites](#).

Read our [policy regarding private information posted on Twitter](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

The user shares your name but has no other commonalities, or

The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

Following rules and best practices

To learn more about following, please read "[FAQs about following](#)."

What are the limits?

We don't limit the number of followers you can have. However, **we do monitor how aggressively users follow other users**. We try to make sure that none of our limits restrain reasonable usage, and will not affect most Twitter users.

We monitor all accounts for aggressive following and follow churn (repeatedly following and unfollowing large numbers of other users). You can read more about these below, but if you don't follow or unfollow hundreds of users in a single day, and you aren't using automated methods of following users, you should be fine.

Please note that Twitter does not permit any automated or bulk following or unfollowing behavior. Please review our [Automation rules and best practices](#) for more information on automating your account.

What is aggressive following?

Aggressive following is defined as indiscriminately following hundreds of accounts just to garner attention. However, following a few users if their accounts seem interesting is normal and is not considered aggressive.

What is aggressive follow churn?

If you decide to follow someone and then change your mind later, you can just visit the person's profile page and unfollow them. Aggressive follow churn is when an account *repeatedly* follows and then unfollows a large number of users. This may be done to get lots of people to notice them, to circumvent a Twitter limit, or to change their follower-to-following ratio. These behaviors negatively impact the Twitter experience for other users, are common spam tactics, and may lead to account suspension.

Are there additional limits if you are following 5000+ accounts?

The rules about aggressive following and follow churn still apply. In addition, **every user can follow 5000 people total**. Once you've followed 5000 users, there are limits to the number of additional users you can follow: this limit is different for every user and is based on your ratio of followers to following.

When you hit this limit, we'll tell you by showing an error message in your browser. You'll need to wait until you have more followers in order to follow more users—for example, you can't follow 10,000 people if only 100 people follow you. When you reach a daily or total limit and we show you an error message, you've hit a technical limit imposed to limit egregious behavior by spam accounts and to prevent strain on the site. These are just the technical limits for your account; in addition, you are prohibited from aggressive following behaviors. These behaviors may result in account suspension, regardless of your account's technical ratio.

Limits improve site performance by ensuring that when we send a person's message to all of their followers, the sending of that message is meaningful. Follow limits cannot be lifted by Twitter, and everyone is subject to them, including verified and developer accounts. Based on current behavior in the Twitter community, we've concluded that this is both fair and reasonable.

Why does Twitter monitor following behavior at all? Isn't this a social network?

Twitter works quite differently from social networks: when you accept friend requests on other social networks like Facebook, it usually means you appear in that person's network and they appear in yours. Following on Twitter is different because **following is not mutual**. Twitter allows people to opt-in to (or opt-out of) receiving a person's updates without requiring a mutual relationship.

Limits on Twitter alleviate some of the strain on the invisible part of Twitter, which prevents error pages and downtime on the visible part. For the sake of reliability, we've placed limits on actions like following, API requests per hour, and number of updates per day ([see this page for more information on updating and API limits](#)).

Finally, follower violations are one tactic that spammers often use to abuse Twitter. Monitoring for abuse is one way to reduce spam on Twitter.

What are third party applications to “Get More Followers Fast!”?

You may encounter websites or applications claiming they can help you get lots of followers quickly. These programs may ask for payment for followers, or ask you to follow a list of other users in order to participate. **Using these is not allowed according to the Twitter Rules.**

- [The Twitter Rules](#) addresses the use of these programs by disallowing: using or promoting third-party sites that claim to get you more followers (such as follower trains, sites promising ‘more followers fast,’ or any other site that offers to automatically add followers to your account).

When you give out your username and password to another site or application, you are giving control of your account to someone else. They may then post duplicated, spam, or malicious updates and links, send unwanted direct messages, aggressively follow, or violate other Twitter rules with your account. When these applications do add followers to your account, they are often abandoned or bot accounts that are not reading your updates. If a third-party application causes your account to violate the Twitter Rules, your account may be suspended.

Some third-party applications have been implicated in spam behavior, fraud, the selling of usernames and passwords, and phishing. **Please do not give your username and password out to any third-party application that you have not thoroughly researched.**

Revoke access for any third-party application that you don't recognize by visiting the [Applications tab](#) in Account Settings.

Are there following limits if I'm whitelisted?

Some API administrators have whitelist status so their applications can function without hitting certain system limits. Whitelisting means that an application can have more API requests per hour; it does not increase the follow limits. All whitelisted accounts are still subject to follow limits.

So how am I supposed to get followers?

Remember, Twitter isn't a race to get the most followers. If you follow users that you're interested in and post meaningful content, it's more likely that legitimate users will find you and read your updates. People follow other users on Twitter to read updates that are interesting to them. Aggressively following and unfollowing accounts is frustrating to other Twitter users and degrades the Twitter experience for everyone.

Additional resources

Users who viewed this article also found these articles helpful:

[Common following issues](#)

[Rules and best practices](#)

[Automation rules and best practices](#)

Automation rules

Updated November 3, 2017

This page is primarily intended for **developers**.

For Twitter users: You are ultimately responsible for the actions taken with your account, or by applications associated with your account. Before authorizing a third-party application to access or use your account, make sure you've thoroughly investigated the application and understand what it will do. If automated activity on your account violates the [Twitter Rules](#) or these automation rules, Twitter may take action on your

account, including [filtering your Tweets from search results](#) or suspending your account.

For more information on third-party applications, please see our article on [connecting and revoking third-party applications](#).

If you are a developer and have questions about these Automation Rules, please reach out via <https://twittercommunity.com/c/rules-and-policies>.

If your application has been suspended or restricted from performing write actions, you may submit an appeal at </forms/platform>.

I. Ground Rules Do!

Build solutions that automatically broadcast helpful information in Tweets.

Run creative campaigns that auto-reply to users who engage with your content.

Build solutions that automatically respond to users in Direct Messages.

Try new things that help people (and comply with our rules).

Make sure your application provides a good user experience and performs well — and confirm that remains the case over time.

Don't!

Violate these or other policies. Be extra mindful of our rules about abuse and user privacy.

Abuse the Twitter API or attempt to circumvent rate limits.

Use non-API-based forms of automation, such as scripting the Twitter website. The use of these techniques may result in the permanent suspension of your account.

Spam or bother users, or otherwise send them unsolicited messages.

A. The Twitter Rules and the Developer Agreement and Policy

As with all activity on Twitter, automated activity is subject to the [Twitter Rules](#) and, if you're a developer using the Twitter API, the [Developer Agreement and Policy](#).

You should carefully review these policies to ensure that your automated activity is compliant. Automated applications or activities that violate these policies, or that facilitate or induce users to violate them, may be subject to enforcement action,

potentially including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these policies.

Although all aspects of the Twitter Rules and the Developer Agreement and Policy apply to automated activity, you should keep the following rules top of mind:

Spamming: You may not send automated Tweets or Direct Messages that are spam, or otherwise engage in spamming activity. Some examples of spammy behavior to avoid with automation include:

Trending topics: You may not automatically post about trending topics on Twitter, or use automation to attempt to influence or manipulate trending topics.

Multiple posts/accounts: You may not post duplicative or substantially similar Tweets on one account or over multiple accounts you operate.

Duplicate accounts: You may not create and/or automate multiple accounts for duplicative or substantially similar use cases.

However, automating multiple accounts for related but non-duplicative use cases is permitted. For example, you may automate separate accounts to Tweet when the Hubble Space Telescope passes over different cities, such as [San Francisco](#) or [Hong Kong](#).

Misleading links: You may not send automated Tweets or Direct Messages containing links that are misleading, including links that maliciously or deceptively redirect through landing pages or ad pages before displaying the final content.

Sensitive media: Automated Tweets and Direct Messages must comply with the [Twitter media policy](#), and you should mark your account as potentially sensitive if you intend to post graphic, pornographic, or potentially sensitive media.

Abusive behavior: You may not engage in any automated activity that encourages, promotes, or incites abuse, violence, hateful conduct, or harassment, on or off Twitter.

Private information: You may not post private or confidential information about a person without their prior express authorization.

B. Other Ground Rules for Automated Activity

In addition to the policies above, the following ground rules apply to all automated activity on Twitter:

Don't surprise or mislead users: Automated activity should honor users' expectations. Ask for the user's permission before taking an action if you aren't sure.

Mature content or profanity: Don't Direct Message, mention, or reply to users with potentially sensitive content (including profanity), unless they've clearly indicated an intent to receive it in advance.

Be thoughtful about the information you request or exchange on Twitter

Tweets: Don't ask users to send you personal or private information via a public Tweet. If you need additional personal or private information from a user to provide them with customer service (or other similar use cases), you should ask the user to share such information by Direct Message or another private channel. You might even consider adding a [Direct Message deep link](#) to your Tweet.

Direct Messages: You should only ask users for the minimum amount of information you need to provide them with service. If you need to request or exchange particularly sensitive information (such as credit card information), you should consider directing users to your website or other appropriate channel to do so.

II. Activity-Specific Rules

The activity-specific rules in this section apply to taking specific automated actions on Twitter. Please read these rules carefully, as they outline both permitted and prohibited use cases of automation.

Automated applications or activities that violate these rules, or that facilitate or induce users to violate them, may be subject to enforcement action, including suspension of associated Twitter accounts. We may also rate limit, suspend, or terminate developers' access to the Twitter API based on violations of these rules. As a reminder, you should also carefully review the spam guidelines in the [Twitter Rules](#) to avoid having activities performed by you, your app, or other users through your app or service flagged as spam.

A. Automated Actions Through Another User's Account

Twitter users may authorize your app or service to [access their Twitter account through OAuth](#). A user authorizing your app or service to access their Twitter account through OAuth does not by itself constitute sufficient consent to take automated actions through that user's account.

You may only take automated actions through another Twitter user's account if you:

clearly describe to the user the types of automated actions that will occur; receive express consent from the user to take those automated actions; and immediately honor a user's request to opt-out of further automated actions.

If you substantially change the purpose or functionality of your app or service, you must re-obtain express consent from the user to take automated action through their account before doing so.

These requirements apply to any automated action taken through another Twitter user's account, including posting Tweets, sending Direct Messages, deleting Tweets or Direct Messages, or following/unfollowing other accounts. For applications that offer users the ability to delete Tweets in a bulk or automated manner, you must also clearly state that Tweets are not recoverable once deleted.

B. Automated Tweets

1. Posting automated Tweets

Automated Tweets that cross-post outside information:

You may post automated Tweets based on sources of outside information — such as an RSS feed, weather data, etc. — as long as you are sufficiently authorized to publish such information.

Other automated Tweets (excluding mentions or replies):

Provided you comply with all other rules, you may post automated Tweets for entertainment, informational, or novelty

purposes. As a reminder, accounts posting duplicative, spammy, or otherwise prohibited content may be subject to suspension.

2. Posting automated mentions and replies

The reply and mention functions are intended to make communication between Twitter users easier. Automating these actions to reach many users on an unsolicited basis is an abuse of the feature, and is not permitted. For example, sending automated replies to Tweets based on keyword searches alone is not permitted. Spamy or duplicative use of mentions and replies may result in enforcement action, such as the removal of your Tweets from Search or the suspension of your app or account.

However, you may send automated replies or mentions to Twitter users so long as:

in advance of sending the automated reply, the recipient or mentioned user(s) have requested or have clearly indicated an intent on Twitter to be contacted by you (i.e. opted in), for example by replying to a Tweet from your account, or by sending you a Direct Message;

you provide a clear and easy way for such users to opt-out of receiving automated replies and mentions, and promptly honor all such opt-out requests;

you only send one automated reply or mention per user interaction; and

the automated reply or mention is a reply to the user's original Tweet (if your campaign is based on users posting a reply to your Tweet).

Opt-in techniques and indications of user intent take many different forms, depending on the specifics of your use case and implementation. Some examples include:

A Tweet from your account that clearly indicates that a user taking a specific action on that Tweet (such as Retweeting it) will opt the user into receiving an automated response.

A mention of your account by the user in a manner suggesting the user clearly wishes or intends to receive a response. If you want to run an auto-reply campaign with a campaign- or use-case-specific hashtag, users should also mention you in their Tweets.

Note that a user following your account is not on its own a sufficient indication of user intent to receive an automated response.

: Advertisers, publishers, and brands using auto-response campaigns must request approval from Twitter and may be subject to additional rules. Please reach out to your account or partner manager for assistance. If you do not have a partner manager but are interested in managed client services, you can learn more [here](#).

Additionally, we recommend that any accounts that will communicate with

users via automated mentions or replies:

Note

1. Appropriately filter responses based on potentially sensitive language in user handles,

display names, and Tweet text, as well as potentially sensitive media;

2. Check that the Tweet you are mentioning or replying to still exists (for example, using the statuses/lookup endpoint on the Twitter API).

C. Automated Direct Messages

Sending automated Direct Messages to users

You may not send unsolicited Direct Messages in a bulk or automated manner, and should be thoughtful about the frequency with which you contact users via Direct Message.

You may send automated Direct Messages to users so long as:

in advance of sending the Direct Message, the recipient(s) have requested or have clearly indicated an intent on Twitter to be contacted by you via Direct Message, for example by sending you a Direct Message; **and**

you provide a clear and easy way for such users to opt-out of receiving automated Direct Messages, and promptly honor all such opt-out requests.

The fact that a user is technically able to receive a Direct Message from you (e.g. because the user follows you, has enabled the ability to receive Direct Messages from any

account, or because the user is in a pre-existing Direct Message conversation with you) does not necessarily mean they have requested or expect to receive automated Direct Messages from you.

Interacting with users via Direct Message

Per the Ground Rules, remember to be thoughtful about the amount and type of information you request or exchange with users via Direct Messages. If you will be asking a user to provide personal or private information via an automated Direct Message, you must clearly explain how you will use the information you're collecting. Consider including a link to your privacy policy in your Direct Message to the user, as well as in your Twitter profile bio.

Don't publicly share information received in a Direct Message conversation with a user without first obtaining explicit consent from the user. For example, if a user asks you via Direct Message about a purchase they made from you, you may not mention the user in a public Tweet that includes information about their purchase unless you have the user's explicit consent to do so.

After a user-initiated interaction ends, don't send additional follow-up Direct Messages or mention users in a Tweet unless you get permission from the user.

D. Automated actions you take on Tweets or accounts

Automated likes: You may not like Tweets in an automated manner.

Automated Retweets: Provided you comply with all other rules, you may Retweet or Quote Tweet in an automated manner for entertainment, informational, or novelty purposes. Automated Retweets often lead to negative user experiences, and bulk, aggressive, or spammy Retweeting is a violation of the [Twitter Rules](#).

Automated following/unfollowing: You may not follow or unfollow Twitter accounts in a bulk, aggressive, or indiscriminate manner. Aggressive following is a violation of the [Twitter Rules](#). Please also review our [following rules and best practices](#) to ensure you are in compliance. Note that applications that claim to get users more followers are also prohibited under the [Twitter Rules](#).

Automated adding to lists or collections: You may not add Twitter users to lists or add Tweets to collections in a bulk or indiscriminate manner. Adding a large number of unrelated users to lists is a violation of the [Twitter Rules](#).

About search rules and restrictions

This article details the rules for using Twitter search, and explains reasons why content may not show up in Twitter search.

Why doesn't all content show up in Twitter search?

In order to keep your search results relevant, Twitter filters search results for quality Tweets and accounts. Material that jeopardizes search quality or creates a bad search experience for other people may be automatically removed from Twitter search. Read more about this [here](#).

General rules for using Twitter search

The [Twitter Rules](#) explain what behaviors are permitted on Twitter. In addition to these rules, we've included some tips below to keep your content relevant (and your Tweets in Twitter search).

Please do not:

- Repeatedly post duplicate or near-duplicate content (links or Tweets).
- Abuse trending topics or hashtags (topic words with a # sign).
- Send automated Tweets or replies.
- Use bots or applications to post similar messages based on keywords.
- Post similar messages over multiple accounts.
- Aggressively follow and unfollow people.

Twitter may automatically remove accounts engaging in these behaviors from search (or even suspend in some cases) in order to ensure the best experience for everyone.

Search results filtering

To learn how to filter your search results, including how to disable or enable safe search, read our [Using Twitter search](#) article.

Search results filtering in third-party applications

Accounts may be filtered from search for posting similar messages over several accounts. Your account may be filtered from search if you are logged into a third-party application that is updating many accounts with similar or duplicate content.

Rules for participating in contests on Twitter

Businesses sometimes host contests and offer prizes to Retweet an update, following a particular account, or posting Tweets with a specific hashtag topic or reply. If you're running a contest, check out this [article](#) for guidelines for contests on Twitter.

If you repeatedly post near-duplicate updates or duplicate links in order to enter contests, you may be filtered from search. Please keep in mind that you are responsible for the updates that you post to your Twitter account.

What to do if you are missing from search

Please see this [article](#) for more information related to this issue.

If you are being filtered from search, rest assured that your followers will still receive your updates. If you post a reply or mention to a specific person, it will be delivered to that person.

Having trouble?

You can learn more about searching by keyword, location, and trending topics in our [Using Twitter search](#) help article.

The Twitter Rules

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;

- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Terrorism and violent extremism policy

Overview

March 2019

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for terrorist organizations or violent extremist groups and individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments in this context are informed by national and international terrorism designations. We also assess organizations under our violent extremist group criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

We examine a group's activities both on and off Twitter to determine whether they engage in and/or promote violence against civilians to advance a political, religious and/or social cause.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a terrorist organization or violent extremist group;
- recruiting for a terrorist organization or violent extremist group;
- providing or distributing services (e.g., financial, media/propaganda) to further a terrorist organization's or violent extremist group's stated goals; and
- using the insignia or symbols of terrorist organizations or violent extremist groups to promote them.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to military or government entities.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wish for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy

Overview

March 2019

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives; and
- identifying alleged victims of childhood sexual exploitation by name or image.

What is not a violation of this policy?

Discussions related to child sexual exploitation as a phenomenon or attraction towards minors are permitted, provided they don't promote or glorify child sexual exploitation **in any way**. Artistic depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios e.g., works by internationally renowned artists that feature minors.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

What you can do

If you think you've found a Twitter account distributing or promoting child sexual exploitation, please report it through our designated [form](#). Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. For help on how to find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with NCMEC [here](#) or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the International Association of Internet Hotlines [website](#).

Our partners

We partner with multiple organisations around the world whose work is dedicated to fighting child sexual exploitation. You can find a full list of our safety partners [here](#).

Abusive behavior

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing or hoping serious harm on a person or group of people

We do not tolerate content that wishes, hopes or expresses a desire for death, serious bodily harm or fatal disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone's body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not act against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives. We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted. We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you".

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease, e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident, e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserve serious physical injury, e.g., "If this group of protesters don't shut up, they deserve to be shot."

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

- media that depicts victims of the Holocaust;
- media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., "all [religious group] are terrorists".

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

Glorifying self-harm and suicide

Twitter Rules: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Rationale

While we want people to feel safe sharing their thoughts and feelings, we draw the line at encouraging or promoting self-harm and suicidal behavior, which can pose safety risks for others. With that in mind, we apply a two-pronged approach to the issue: supporting people who are undergoing experiences with self-harm or suicidal thoughts, but prohibiting the promotion or encouragement of self-harming behaviors.

[Learn more](#) about how you can help someone dealing with thoughts of self-harm or suicide.

When this applies

Some examples of encouraging or promoting self-harm include (but may not be limited to) encouraging or glorifying:

- Acts of injuring oneself, such as cutting
- Eating disorders such as bulimia, anorexia, and serious illnesses that cause severe disturbances to a person's eating behaviors.
- The act of taking one's own life, including sharing strategies for suicide and participating in group suicides or suicide games

Some examples of behavior that are not considered glorifying or encouraging:

- Discussions or organizations focused on research, advocacy, and education of self-harm or suicide prevention
- Sharing coping mechanisms and resources for thoughts of self-harm and suicide

Consequences

Our approach towards suicide and self-harm depends on the type of content, whether or not the account is encouraging or promoting self-harm, and the account's previous record of violations.

The first time someone violates this policy, they will be required to remove the violating Tweet and be temporarily locked out of their account before they can Tweet again.

People who repeatedly violate this policy and/or accounts dedicated to promoting or encouraging self-harm and suicide may be suspended. If someone believes their account was suspended in error, they can [submit an appeal](#).

After we assess a report of self-harm or suicide, Twitter will contact the reported individual and let them know that someone who cares about them identified that they might be at risk. We will provide the reported individual with available resources and encourage them to seek help.

In some cases where content is not glorifying or encouraging suicide or self-harm, we may still interstitial Tweets as sensitive. Learn more about our [sensitive media policy](#).

Learn more about dealing with [self-harm and suicide](#).

Visit our Safety Center for a list of local [mental health resources](#).

Sensitive media policy

Overview

March 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive image**.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive image**.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;

- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;
 - descriptions of physical appearance;
 - gossip, rumours, accusations, and allegations; and
 - screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.

3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy

Overview

March 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy

Overview

March 2019

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are; and
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting.

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate; and
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in “follow trains,” “decks,” and “Retweet for Retweet” behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- “follow churn” – following and then unfollowing large numbers of accounts in an effort to inflate one's own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account's followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person's browser (malware) or computer or to compromise a person's privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don't violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Blacklisting URLs

We blacklist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [hacked materials policy](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Election integrity policy

Overview

April 2019

You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote.

The public conversation occurring on Twitter is never more important than during elections. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt elections, including through the distribution of false or misleading information about the electoral process or when or how to vote.

What is in violation of this policy?

We prohibit 3 categories of manipulative behavior and content under this policy:

Misleading information about how to participate

You can't share false or misleading information about how to participate in an election or other civic event. This includes but is not limited to:

- misleading information about how to vote or register to vote (for example, that you can vote by Tweet, text message, email, or phone call);
- misleading information about requirements for voting, including identification requirements; and
- misleading statements or information about the official, announced date or time of an election.

Voter suppression and intimidation

You can't share false or misleading information intended to intimidate or dissuade voters from participating in an election. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to polling places or elections;
- misleading claims about long lines or equipment problems at voting locations;
- misleading claims about voting procedures or techniques which could dissuade voters from participating in an election; and
- threats regarding voting locations (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account](#) policy.

What is not a violation of this policy?

Not all false or untrue information about politics or political events constitutes manipulation or interference in an election. In the absence of other violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election. Therefore, we enable reporting of false or misleading information about voting or registering to vote for people accessing Twitter in advance of major elections. We also work with select government and civil society partners in these countries to provide additional options for reporting and expedited review.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about voting**.
3. Select the option that best tells us how the Tweet is misleading about voting.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about voting**.
3. Select the option that best tells us how the Tweet is misleading about voting.

4. Submit your report.

What happens if you violate this policy?

The consequences for violating our election integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

The user shares your name but has no other commonalities, or

The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

Trademark policy

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Copyright policy

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see

our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and

4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Our approach to policy development and enforcement philosophy

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (20190607*)

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;

- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Terrorism and violent extremism policy

(20190607*)

Overview

March 2019

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for terrorist organizations or violent extremist groups and individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments in this context are informed by national and international terrorism designations. We also assess organizations under our violent extremist group criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

We examine a group's activities both on and off Twitter to determine whether they engage in and/or promote violence against civilians to advance a political, religious and/or social cause.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a terrorist organization or violent extremist group;
- recruiting for a terrorist organization or violent extremist group;
- providing or distributing services (e.g., financial, media/propaganda) to further a terrorist organization's or violent extremist group's stated goals; and
- using the insignia or symbols of terrorist organizations or violent extremist groups to promote them.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to military or government entities.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wish for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy **(20190607*)**

Overview

March 2019

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives; and
- identifying alleged victims of childhood sexual exploitation by name or image.

What is not a violation of this policy?

Discussions related to child sexual exploitation as a phenomenon or attraction towards minors are permitted, provided they don't promote or glorify child sexual exploitation **in any way**. Artistic depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios e.g., works by internationally renowned artists that feature minors.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third

party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

What you can do

If you think you've found a Twitter account distributing or promoting child sexual exploitation, please report it through our designated [form](#). Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. For help on how to find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with NCMEC [here](#) or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the International Association of Internet Hotlines [website](#).

Our partners

We partner with multiple organisations around the world whose work is dedicated to fighting child sexual exploitation. You can find a full list of our safety partners [here](#).

Abusive behavior (20190607*)

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing or hoping serious harm on a person or group of people

We do not tolerate content that wishes, hopes or expresses a desire for death, serious bodily harm or fatal disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone's body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy (20190727*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically

marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you".

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., "I hope you get cancer and die."

Wishing for someone to fall victim to a serious accident, e.g., "I wish that you would get run over by a car next time you run your mouth."

Saying that a group of individuals deserve serious physical injury, e.g., "If this group of protesters don't shut up, they deserve to be shot."

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected

category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can

[submit an appeal](#).

Glorifying self-harm and suicide (20190607*)

[Twitter Rules](#): You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Rationale

While we want people to feel safe sharing their thoughts and feelings, we draw the line at encouraging or promoting self-harm and suicidal behavior, which can pose safety risks for others. With that in mind, we apply a two-pronged approach to the issue: supporting people who are undergoing experiences with self-harm or suicidal thoughts, but prohibiting the promotion or encouragement of self-harming behaviors.

[Learn more](#) about how you can help someone dealing with thoughts of self-harm or suicide.

When this applies

Some examples of encouraging or promoting self-harm include (but may not be limited to) encouraging or glorifying:

- Acts of injuring oneself, such as cutting
- Eating disorders such as bulimia, anorexia, and serious illnesses that cause severe disturbances to a person's eating behaviors.
- The act of taking one's own life, including sharing strategies for suicide and participating in group suicides or suicide games

Some examples of behavior that are not considered glorifying or encouraging:

- Discussions or organizations focused on research, advocacy, and education of self-harm or suicide prevention
- Sharing coping mechanisms and resources for thoughts of self-harm and suicide

Consequences

Our approach towards suicide and self-harm depends on the type of content, whether or not the account is encouraging or promoting self-harm, and the account's previous record of violations.

The first time someone violates this policy, they will be required to remove the violating Tweet and be temporarily locked out of their account before they can Tweet again.

People who repeatedly violate this policy and/or accounts dedicated to promoting or encouraging self-harm and suicide may be suspended. If someone believes their account was suspended in error, they can [submit an appeal](#).

After we assess a report of self-harm or suicide, Twitter will contact the reported individual and let them know that someone who cares about them identified that they might be at risk. We will provide the reported individual with available resources and encourage them to seek help.

In some cases where content is not glorifying or encouraging suicide or self-harm, we may still interstitial Tweets as sensitive. Learn more about our [sensitive media policy](#).

Learn more about dealing with [self-harm and suicide](#).

Visit our Safety Center for a list of local [mental health resources](#).

Sensitive media policy (20190607*)

Overview

March 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to

promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive image**.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive image**.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20190607*)

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;

- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;
 - descriptions of physical appearance;
 - gossip, rumours, accusations, and allegations; and
 - screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy

(20190607*)

Overview

March 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;

- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy **(20190607*)**

Overview

March 2019

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are; and

- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting.

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate; and
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in "follow trains," "decks," and "Retweet for Retweet" behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- “follow churn” – following and then unfollowing large numbers of accounts in an effort to inflate one’s own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account’s followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person’s browser (malware) or computer or to compromise a person’s privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don’t violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Blacklisting URLs

We blacklist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [hacked materials policy](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Election integrity policy (20190607*)

Overview

April 2019

You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote.

The public conversation occurring on Twitter is never more important than during elections. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt elections, including through the distribution of false or misleading information about the electoral process or when or how to vote.

What is in violation of this policy?

We prohibit 3 categories of manipulative behavior and content under this policy:

Misleading information about how to participate

You can't share false or misleading information about how to participate in an election or other civic event. This includes but is not limited to:

- misleading information about how to vote or register to vote (for example, that you can vote by Tweet, text message, email, or phone call);
- misleading information about requirements for voting, including identification requirements; and
- misleading statements or information about the official, announced date or time of an election.

Voter suppression and intimidation

You can't share false or misleading information intended to intimidate or dissuade voters from participating in an election. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to polling places or elections;
- misleading claims about long lines or equipment problems at voting locations;
- misleading claims about voting procedures or techniques which could dissuade voters from participating in an election; and
- threats regarding voting locations (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account](#) policy.

What is not a violation of this policy?

Not all false or untrue information about politics or political events constitutes manipulation or interference in an election. In the absence of other violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election. Therefore, we enable reporting of false or misleading information about voting or registering to vote for people accessing Twitter in advance of major elections. We also work with select government and civil society partners in these countries to provide additional options for reporting and expedited review.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about voting**.
3. Select the option that best tells us how the Tweet is misleading about voting.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about voting**.
3. Select the option that best tells us how the Tweet is misleading about voting.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our election integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Impersonation policy (20190130*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

The user shares your name but has no other commonalities, or

The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody](#), [newsfeed](#), [commentary](#), and [fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be

permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Copyright policy \(20181206*\)](#)

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity’s name, please review Twitter’s [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

<https://web.archive.org/web/20180601194918/https://help.twitter.com/en/rules-and-policies/copyright-policy> 1/3

27.11.2020 Copyright policy

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

<https://web.archive.org/web/20180601194918/https://help.twitter.com/en/rules-and-policies/copyright-policy> 2/3

27.11.2020 Copyright policy

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Our approach to policy development and enforcement philosophy **(20190520*)**

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (20190607*)

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and

permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.

4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Terrorism and violent extremism policy

(20190607*)

Overview

March 2019

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for terrorist organizations or violent extremist groups and individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments in this context are informed by national and international terrorism designations. We also assess organizations under our violent extremist group criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

We examine a group's activities both on and off Twitter to determine whether they engage in and/or promote violence against civilians to advance a political, religious and/or social cause.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a terrorist organization or violent extremist group;
- recruiting for a terrorist organization or violent extremist group;
- providing or distributing services (e.g., financial, media/propaganda) to further a terrorist organization's or violent extremist group's stated goals; and
- using the insignia or symbols of terrorist organizations or violent extremist groups to promote them.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to military or government entities.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wish for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy

(20190607*)

Overview

March 2019

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives; and
- identifying alleged victims of childhood sexual exploitation by name or image.

What is not a violation of this policy?

Discussions related to child sexual exploitation as a phenomenon or attraction towards minors are permitted, provided they don't promote or glorify child sexual exploitation **in any way**. Artistic depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios e.g., works by internationally renowned artists that feature minors.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

What you can do

If you think you've found a Twitter account distributing or promoting child sexual exploitation, please report it through our designated [form](#). Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. For help on how to find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with NCMEC [here](#) or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the International Association of Internet Hotlines [website](#).

Our partners

We partner with multiple organisations around the world whose work is dedicated to fighting child sexual exploitation. You can find a full list of our safety partners [here](#).

Abusive behavior (20190607*)

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing or hoping serious harm on a person or group of people

We do not tolerate content that wishes, hopes or expresses a desire for death, serious bodily harm or fatal disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone's body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy (20190829*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., “I will kill you”.

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., “I hope you get cancer and die.”

Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”

Saying that a group of individuals deserve serious physical injury, e.g., “If this group of protesters don’t shut up, they deserve to be shot.”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can

[submit an appeal](#).

Glorifying self-harm and suicide (20190607*)

[Twitter Rules](#): You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Rationale

While we want people to feel safe sharing their thoughts and feelings, we draw the line at encouraging or promoting self-harm and suicidal behavior, which can pose safety risks for others. With that in mind, we apply a two-pronged approach to the issue: supporting people who are undergoing experiences with self-harm or suicidal thoughts, but prohibiting the promotion or encouragement of self-harming behaviors.

[Learn more](#) about how you can help someone dealing with thoughts of self-harm or suicide.

When this applies

Some examples of encouraging or promoting self-harm include (but may not be limited to) encouraging or glorifying:

- Acts of injuring oneself, such as cutting
- Eating disorders such as bulimia, anorexia, and serious illnesses that cause severe disturbances to a person's eating behaviors.
- The act of taking one's own life, including sharing strategies for suicide and participating in group suicides or suicide games

Some examples of behavior that are not considered glorifying or encouraging:

- Discussions or organizations focused on research, advocacy, and education of self-harm or suicide prevention
- Sharing coping mechanisms and resources for thoughts of self-harm and suicide

Consequences

Our approach towards suicide and self-harm depends on the type of content, whether or not the account is encouraging or promoting self-harm, and the account's previous record of violations.

The first time someone violates this policy, they will be required to remove the violating Tweet and be temporarily locked out of their account before they can Tweet again.

People who repeatedly violate this policy and/or accounts dedicated to promoting or encouraging self-harm and suicide may be suspended. If someone believes their account was suspended in error, they can [submit an appeal](#).

After we assess a report of self-harm or suicide, Twitter will contact the reported individual and let them know that someone who cares about them identified that they might be at risk. We will provide the reported individual with available resources and encourage them to seek help.

In some cases where content is not glorifying or encouraging suicide or self-harm, we may still interstitial Tweets as sensitive. Learn more about our [sensitive media policy](#).

Learn more about dealing with [self-harm and suicide](#).

Visit our Safety Center for a list of local [mental health resources](#).

Sensitive media policy (20190607*)

Overview

March 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to

promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive image**.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive image**.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20190607*)

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;

- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;
 - descriptions of physical appearance;
 - gossip, rumours, accusations, and allegations; and
 - screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy

(20190607*)

Overview

March 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;

- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy (20190607*)

Overview

March 2019

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are; and

- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting.

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate; and
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in "follow trains," "decks," and "Retweet for Retweet" behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- “follow churn” – following and then unfollowing large numbers of accounts in an effort to inflate one’s own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account’s followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person’s browser (malware) or computer or to compromise a person’s privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don’t violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Blacklisting URLs

We blacklist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [hacked materials policy](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Election integrity policy (20190607*)

Overview

April 2019

You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote.

The public conversation occurring on Twitter is never more important than during elections. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt elections, including through the distribution of false or misleading information about the electoral process or when or how to vote.

What is in violation of this policy?

We prohibit 3 categories of manipulative behavior and content under this policy:

Misleading information about how to participate

You can't share false or misleading information about how to participate in an election or other civic event. This includes but is not limited to:

- misleading information about how to vote or register to vote (for example, that you can vote by Tweet, text message, email, or phone call);
- misleading information about requirements for voting, including identification requirements; and
- misleading statements or information about the official, announced date or time of an election.

Voter suppression and intimidation

You can't share false or misleading information intended to intimidate or dissuade voters from participating in an election. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to polling places or elections;
- misleading claims about long lines or equipment problems at voting locations;
- misleading claims about voting procedures or techniques which could dissuade voters from participating in an election; and
- threats regarding voting locations (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account](#) policy.

What is not a violation of this policy?

Not all false or untrue information about politics or political events constitutes manipulation or interference in an election. In the absence of other violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election. Therefore, we enable reporting of false or misleading information about voting or registering to vote for people accessing Twitter in advance of major elections. We also work with select government and civil society partners in these countries to provide additional options for reporting and expedited review.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about voting**.
3. Select the option that best tells us how the Tweet is misleading about voting.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about voting**.
3. Select the option that best tells us how the Tweet is misleading about voting.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our election integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Impersonation policy (20190130*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

The user shares your name but has no other commonalities, or

The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody](#), [newsfeed](#), [commentary](#), and [fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be

permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Copyright policy \(20181206*\)](#)

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity’s name, please review Twitter’s [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

<https://web.archive.org/web/20180601194918/https://help.twitter.com/en/rules-and-policies/copyright-policy> 1/3

27.11.2020 Copyright policy

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

<https://web.archive.org/web/20180601194918/https://help.twitter.com/en/rules-and-policies/copyright-policy> 2/3

27.11.2020 Copyright policy

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Our approach to policy development and enforcement philosophy **(20190520*)**

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (20190607*)

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and

permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.

4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Terrorism and violent extremism policy

(20190607*)

Overview

March 2019

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for terrorist organizations or violent extremist groups and individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments in this context are informed by national and international terrorism designations. We also assess organizations under our violent extremist group criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

We examine a group's activities both on and off Twitter to determine whether they engage in and/or promote violence against civilians to advance a political, religious and/or social cause.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a terrorist organization or violent extremist group;
- recruiting for a terrorist organization or violent extremist group;
- providing or distributing services (e.g., financial, media/propaganda) to further a terrorist organization's or violent extremist group's stated goals; and
- using the insignia or symbols of terrorist organizations or violent extremist groups to promote them.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to military or government entities.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wish for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy

(20190607*)

Overview

March 2019

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives; and
- identifying alleged victims of childhood sexual exploitation by name or image.

What is not a violation of this policy?

Discussions related to child sexual exploitation as a phenomenon or attraction towards minors are permitted, provided they don't promote or glorify child sexual exploitation **in any way**. Artistic depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios e.g., works by internationally renowned artists that feature minors.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

What you can do

If you think you've found a Twitter account distributing or promoting child sexual exploitation, please report it through our designated [form](#). Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. For help on how to find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with NCMEC [here](#) or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the International Association of Internet Hotlines [website](#).

Our partners

We partner with multiple organisations around the world whose work is dedicated to fighting child sexual exploitation. You can find a full list of our safety partners [here](#).

Abusive behavior (20190607*)

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing or hoping serious harm on a person or group of people

We do not tolerate content that wishes, hopes or expresses a desire for death, serious bodily harm or fatal disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone's body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy (20190829*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., “I will kill you”.

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., “I hope you get cancer and die.”

Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”

Saying that a group of individuals deserve serious physical injury, e.g., “If this group of protesters don’t shut up, they deserve to be shot.”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can

[submit an appeal](#).

Glorifying self-harm and suicide (20190607*)

[Twitter Rules](#): You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Rationale

While we want people to feel safe sharing their thoughts and feelings, we draw the line at encouraging or promoting self-harm and suicidal behavior, which can pose safety risks for others. With that in mind, we apply a two-pronged approach to the issue: supporting people who are undergoing experiences with self-harm or suicidal thoughts, but prohibiting the promotion or encouragement of self-harming behaviors.

[Learn more](#) about how you can help someone dealing with thoughts of self-harm or suicide.

When this applies

Some examples of encouraging or promoting self-harm include (but may not be limited to) encouraging or glorifying:

- Acts of injuring oneself, such as cutting
- Eating disorders such as bulimia, anorexia, and serious illnesses that cause severe disturbances to a person's eating behaviors.
- The act of taking one's own life, including sharing strategies for suicide and participating in group suicides or suicide games

Some examples of behavior that are not considered glorifying or encouraging:

- Discussions or organizations focused on research, advocacy, and education of self-harm or suicide prevention
- Sharing coping mechanisms and resources for thoughts of self-harm and suicide

Consequences

Our approach towards suicide and self-harm depends on the type of content, whether or not the account is encouraging or promoting self-harm, and the account's previous record of violations.

The first time someone violates this policy, they will be required to remove the violating Tweet and be temporarily locked out of their account before they can Tweet again.

People who repeatedly violate this policy and/or accounts dedicated to promoting or encouraging self-harm and suicide may be suspended. If someone believes their account was suspended in error, they can [submit an appeal](#).

After we assess a report of self-harm or suicide, Twitter will contact the reported individual and let them know that someone who cares about them identified that they might be at risk. We will provide the reported individual with available resources and encourage them to seek help.

In some cases where content is not glorifying or encouraging suicide or self-harm, we may still interstitial Tweets as sensitive. Learn more about our [sensitive media policy](#).

Learn more about dealing with [self-harm and suicide](#).

Visit our Safety Center for a list of local [mental health resources](#).

Sensitive media policy (20190607*)

Overview

March 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to

promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive image**.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive image**.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20190607*)

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;

- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;
 - descriptions of physical appearance;
 - gossip, rumours, accusations, and allegations; and
 - screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy

(20190607*)

Overview

March 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;

- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy **(20190925*)**

Overview

September 2019

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are; and

- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting.

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate; and
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in "follow trains," "decks," and "Retweet for Retweet" behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- “follow churn” – following and then unfollowing large numbers of accounts in an effort to inflate one’s own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account’s followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person’s browser (malware) or computer or to compromise a person’s privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don’t violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Blacklisting URLs

We blacklist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [financial scam policy](#), our [hacked materials policy](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Election integrity policy (20190607*)

Overview

April 2019

You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote.

The public conversation occurring on Twitter is never more important than during elections. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt elections, including through the distribution of false or misleading information about the electoral process or when or how to vote.

What is in violation of this policy?

We prohibit 3 categories of manipulative behavior and content under this policy:

Misleading information about how to participate

You can't share false or misleading information about how to participate in an election or other civic event. This includes but is not limited to:

- misleading information about how to vote or register to vote (for example, that you can vote by Tweet, text message, email, or phone call);
- misleading information about requirements for voting, including identification requirements; and
- misleading statements or information about the official, announced date or time of an election.

Voter suppression and intimidation

You can't share false or misleading information intended to intimidate or dissuade voters from participating in an election. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to polling places or elections;
- misleading claims about long lines or equipment problems at voting locations;
- misleading claims about voting procedures or techniques which could dissuade voters from participating in an election; and
- threats regarding voting locations (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account](#) policy.

What is not a violation of this policy?

Not all false or untrue information about politics or political events constitutes manipulation or interference in an election. In the absence of other violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election. Therefore, we enable reporting of false or misleading information about voting or registering to vote for people accessing Twitter in advance of major elections. We also work with select government and civil society partners in these countries to provide additional options for reporting and expedited review.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about voting**.
3. Select the option that best tells us how the Tweet is misleading about voting.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about voting**.
3. Select the option that best tells us how the Tweet is misleading about voting.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our election integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Impersonation policy (20190130*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

The user shares your name but has no other commonalities, or

The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody](#), [newsfeed](#), [commentary](#), and [fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be

permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Copyright policy \(20181206*\)](#)

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity’s name, please review Twitter’s [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

<https://web.archive.org/web/20180601194918/https://help.twitter.com/en/rules-and-policies/copyright-policy> 1/3

27.11.2020 Copyright policy

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

<https://web.archive.org/web/20180601194918/https://help.twitter.com/en/rules-and-policies/copyright-policy> 2/3

27.11.2020 Copyright policy

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Our approach to policy development and enforcement philosophy **(20190520*)**

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (20190607*)

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and

permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.

4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Terrorism and violent extremism policy

(20190607*)

Overview

March 2019

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for terrorist organizations or violent extremist groups and individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments in this context are informed by national and international terrorism designations. We also assess organizations under our violent extremist group criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

We examine a group's activities both on and off Twitter to determine whether they engage in and/or promote violence against civilians to advance a political, religious and/or social cause.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a terrorist organization or violent extremist group;
- recruiting for a terrorist organization or violent extremist group;
- providing or distributing services (e.g., financial, media/propaganda) to further a terrorist organization's or violent extremist group's stated goals; and
- using the insignia or symbols of terrorist organizations or violent extremist groups to promote them.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to military or government entities.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wish for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy

(20190607*)

Overview

March 2019

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives; and
- identifying alleged victims of childhood sexual exploitation by name or image.

What is not a violation of this policy?

Discussions related to child sexual exploitation as a phenomenon or attraction towards minors are permitted, provided they don't promote or glorify child sexual exploitation **in any way**. Artistic depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios e.g., works by internationally renowned artists that feature minors.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

What you can do

If you think you've found a Twitter account distributing or promoting child sexual exploitation, please report it through our designated [form](#). Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. For help on how to find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with NCMEC [here](#) or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the International Association of Internet Hotlines [website](#).

Our partners

We partner with multiple organisations around the world whose work is dedicated to fighting child sexual exploitation. You can find a full list of our safety partners [here](#).

Abusive behavior (20190607*)

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing or hoping serious harm on a person or group of people

We do not tolerate content that wishes, hopes or expresses a desire for death, serious bodily harm or fatal disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone's body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy (20190829*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., “I will kill you”.

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., “I hope you get cancer and die.”

Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”

Saying that a group of individuals deserve serious physical injury, e.g., “If this group of protesters don’t shut up, they deserve to be shot.”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can

[submit an appeal](#).

Glorifying self-harm and suicide (20190607*)

[Twitter Rules](#): You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Rationale

While we want people to feel safe sharing their thoughts and feelings, we draw the line at encouraging or promoting self-harm and suicidal behavior, which can pose safety risks for others. With that in mind, we apply a two-pronged approach to the issue: supporting people who are undergoing experiences with self-harm or suicidal thoughts, but prohibiting the promotion or encouragement of self-harming behaviors.

[Learn more](#) about how you can help someone dealing with thoughts of self-harm or suicide.

When this applies

Some examples of encouraging or promoting self-harm include (but may not be limited to) encouraging or glorifying:

- Acts of injuring oneself, such as cutting
- Eating disorders such as bulimia, anorexia, and serious illnesses that cause severe disturbances to a person's eating behaviors.
- The act of taking one's own life, including sharing strategies for suicide and participating in group suicides or suicide games

Some examples of behavior that are not considered glorifying or encouraging:

- Discussions or organizations focused on research, advocacy, and education of self-harm or suicide prevention
- Sharing coping mechanisms and resources for thoughts of self-harm and suicide

Consequences

Our approach towards suicide and self-harm depends on the type of content, whether or not the account is encouraging or promoting self-harm, and the account's previous record of violations.

The first time someone violates this policy, they will be required to remove the violating Tweet and be temporarily locked out of their account before they can Tweet again.

People who repeatedly violate this policy and/or accounts dedicated to promoting or encouraging self-harm and suicide may be suspended. If someone believes their account was suspended in error, they can [submit an appeal](#).

After we assess a report of self-harm or suicide, Twitter will contact the reported individual and let them know that someone who cares about them identified that they might be at risk. We will provide the reported individual with available resources and encourage them to seek help.

In some cases where content is not glorifying or encouraging suicide or self-harm, we may still interstitial Tweets as sensitive. Learn more about our [sensitive media policy](#).

Learn more about dealing with [self-harm and suicide](#).

Visit our Safety Center for a list of local [mental health resources](#).

Sensitive media policy (20190607*)

Overview

March 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to

promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive image**.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive image**.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20190607*)

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;

- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;
 - descriptions of physical appearance;
 - gossip, rumours, accusations, and allegations; and
 - screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy

(20190607*)

Overview

March 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;

- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy **(20190925*)**

Overview

September 2019

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are; and

- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting.

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate; and
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in "follow trains," "decks," and "Retweet for Retweet" behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- “follow churn” – following and then unfollowing large numbers of accounts in an effort to inflate one’s own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account’s followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person’s browser (malware) or computer or to compromise a person’s privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don’t violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Blacklisting URLs

We blacklist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [financial scam policy](#), our [hacked materials policy](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Election integrity policy (20190607*)

Overview

April 2019

You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote.

The public conversation occurring on Twitter is never more important than during elections. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt elections, including through the distribution of false or misleading information about the electoral process or when or how to vote.

What is in violation of this policy?

We prohibit 3 categories of manipulative behavior and content under this policy:

Misleading information about how to participate

You can't share false or misleading information about how to participate in an election or other civic event. This includes but is not limited to:

- misleading information about how to vote or register to vote (for example, that you can vote by Tweet, text message, email, or phone call);
- misleading information about requirements for voting, including identification requirements; and
- misleading statements or information about the official, announced date or time of an election.

Voter suppression and intimidation

You can't share false or misleading information intended to intimidate or dissuade voters from participating in an election. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to polling places or elections;
- misleading claims about long lines or equipment problems at voting locations;
- misleading claims about voting procedures or techniques which could dissuade voters from participating in an election; and
- threats regarding voting locations (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account](#) policy.

What is not a violation of this policy?

Not all false or untrue information about politics or political events constitutes manipulation or interference in an election. In the absence of other violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election. Therefore, we enable reporting of false or misleading information about voting or registering to vote for people accessing Twitter in advance of major elections. We also work with select government and civil society partners in these countries to provide additional options for reporting and expedited review.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about voting**.
3. Select the option that best tells us how the Tweet is misleading about voting.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about voting**.
3. Select the option that best tells us how the Tweet is misleading about voting.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our election integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Impersonation policy (20191001*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and

- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Copyright policy \(20181206*\)](#)

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has

satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

<https://web.archive.org/web/20180601194918/https://help.twitter.com/en/rules-and-policies/copyright-policy> 1/3

27.11.2020 Copyright policy

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

<https://web.archive.org/web/20180601194918/https://help.twitter.com/en/rules-and-policies/copyright-policy> 2/3

27.11.2020 Copyright policy

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);

3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Our approach to policy development and enforcement philosophy

(20190520*)

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (20190607*)

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and

permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.

4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Terrorism and violent extremism policy

(20190607*)

Overview

March 2019

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for terrorist organizations or violent extremist groups and individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments in this context are informed by national and international terrorism designations. We also assess organizations under our violent extremist group criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

We examine a group's activities both on and off Twitter to determine whether they engage in and/or promote violence against civilians to advance a political, religious and/or social cause.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a terrorist organization or violent extremist group;
- recruiting for a terrorist organization or violent extremist group;
- providing or distributing services (e.g., financial, media/propaganda) to further a terrorist organization's or violent extremist group's stated goals; and
- using the insignia or symbols of terrorist organizations or violent extremist groups to promote them.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to military or government entities.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wish for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy

(20190607*)

Overview

March 2019

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives; and
- identifying alleged victims of childhood sexual exploitation by name or image.

What is not a violation of this policy?

Discussions related to child sexual exploitation as a phenomenon or attraction towards minors are permitted, provided they don't promote or glorify child sexual exploitation **in any way**. Artistic depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios e.g., works by internationally renowned artists that feature minors.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

What you can do

If you think you've found a Twitter account distributing or promoting child sexual exploitation, please report it through our designated [form](#). Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. For help on how to find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with NCMEC [here](#) or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the International Association of Internet Hotlines [website](#).

Our partners

We partner with multiple organisations around the world whose work is dedicated to fighting child sexual exploitation. You can find a full list of our safety partners [here](#).

Abusive behavior (20190607*)

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing or hoping serious harm on a person or group of people

We do not tolerate content that wishes, hopes or expresses a desire for death, serious bodily harm or fatal disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone's body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy (20190829*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., “I will kill you”.

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., “I hope you get cancer and die.”

Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”

Saying that a group of individuals deserve serious physical injury, e.g., “If this group of protesters don’t shut up, they deserve to be shot.”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can

[submit an appeal](#).

Glorifying self-harm and suicide (20190607*)

[Twitter Rules](#): You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Rationale

While we want people to feel safe sharing their thoughts and feelings, we draw the line at encouraging or promoting self-harm and suicidal behavior, which can pose safety risks for others. With that in mind, we apply a two-pronged approach to the issue: supporting people who are undergoing experiences with self-harm or suicidal thoughts, but prohibiting the promotion or encouragement of self-harming behaviors.

[Learn more](#) about how you can help someone dealing with thoughts of self-harm or suicide.

When this applies

Some examples of encouraging or promoting self-harm include (but may not be limited to) encouraging or glorifying:

- Acts of injuring oneself, such as cutting
- Eating disorders such as bulimia, anorexia, and serious illnesses that cause severe disturbances to a person's eating behaviors.
- The act of taking one's own life, including sharing strategies for suicide and participating in group suicides or suicide games

Some examples of behavior that are not considered glorifying or encouraging:

- Discussions or organizations focused on research, advocacy, and education of self-harm or suicide prevention
- Sharing coping mechanisms and resources for thoughts of self-harm and suicide

Consequences

Our approach towards suicide and self-harm depends on the type of content, whether or not the account is encouraging or promoting self-harm, and the account's previous record of violations.

The first time someone violates this policy, they will be required to remove the violating Tweet and be temporarily locked out of their account before they can Tweet again.

People who repeatedly violate this policy and/or accounts dedicated to promoting or encouraging self-harm and suicide may be suspended. If someone believes their account was suspended in error, they can [submit an appeal](#).

After we assess a report of self-harm or suicide, Twitter will contact the reported individual and let them know that someone who cares about them identified that they might be at risk. We will provide the reported individual with available resources and encourage them to seek help.

In some cases where content is not glorifying or encouraging suicide or self-harm, we may still interstitial Tweets as sensitive. Learn more about our [sensitive media policy](#).

Learn more about dealing with [self-harm and suicide](#).

Visit our Safety Center for a list of local [mental health resources](#).

Sensitive media policy (20191204*)

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20190607*)

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;

- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;
 - descriptions of physical appearance;
 - gossip, rumours, accusations, and allegations; and
 - screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy

(20190607*)

Overview

March 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;

- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy **(20190925*)**

Overview

September 2019

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are; and

- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting.

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate; and
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in "follow trains," "decks," and "Retweet for Retweet" behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- “follow churn” – following and then unfollowing large numbers of accounts in an effort to inflate one’s own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account’s followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person’s browser (malware) or computer or to compromise a person’s privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don’t violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Blacklisting URLs

We blacklist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [financial scam policy](#), our [hacked materials policy](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Election integrity policy (20190607*)

Overview

April 2019

You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote.

The public conversation occurring on Twitter is never more important than during elections. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt elections, including through the distribution of false or misleading information about the electoral process or when or how to vote.

What is in violation of this policy?

We prohibit 3 categories of manipulative behavior and content under this policy:

Misleading information about how to participate

You can't share false or misleading information about how to participate in an election or other civic event. This includes but is not limited to:

- misleading information about how to vote or register to vote (for example, that you can vote by Tweet, text message, email, or phone call);
- misleading information about requirements for voting, including identification requirements; and
- misleading statements or information about the official, announced date or time of an election.

Voter suppression and intimidation

You can't share false or misleading information intended to intimidate or dissuade voters from participating in an election. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to polling places or elections;
- misleading claims about long lines or equipment problems at voting locations;
- misleading claims about voting procedures or techniques which could dissuade voters from participating in an election; and
- threats regarding voting locations (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account](#) policy.

What is not a violation of this policy?

Not all false or untrue information about politics or political events constitutes manipulation or interference in an election. In the absence of other violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election. Therefore, we enable reporting of false or misleading information about voting or registering to vote for people accessing Twitter in advance of major elections. We also work with select government and civil society partners in these countries to provide additional options for reporting and expedited review.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about voting**.
3. Select the option that best tells us how the Tweet is misleading about voting.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about voting**.
3. Select the option that best tells us how the Tweet is misleading about voting.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our election integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Impersonation policy (20191001*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and

- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Copyright policy \(20181206*\)](#)

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has

satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

<https://web.archive.org/web/20180601194918/https://help.twitter.com/en/rules-and-policies/copyright-policy> 1/3

27.11.2020 Copyright policy

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

<https://web.archive.org/web/20180601194918/https://help.twitter.com/en/rules-and-policies/copyright-policy> 2/3

27.11.2020 Copyright policy

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);

3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Our approach to policy development and enforcement philosophy

(20190520*)

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (20190607*)

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and

permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.

4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Terrorism and violent extremism policy

(20190607*)

Overview

March 2019

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for terrorist organizations or violent extremist groups and individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments in this context are informed by national and international terrorism designations. We also assess organizations under our violent extremist group criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

We examine a group's activities both on and off Twitter to determine whether they engage in and/or promote violence against civilians to advance a political, religious and/or social cause.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a terrorist organization or violent extremist group;
- recruiting for a terrorist organization or violent extremist group;
- providing or distributing services (e.g., financial, media/propaganda) to further a terrorist organization's or violent extremist group's stated goals; and
- using the insignia or symbols of terrorist organizations or violent extremist groups to promote them.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to military or government entities.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wish for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy

(20190607*)

Overview

March 2019

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives; and
- identifying alleged victims of childhood sexual exploitation by name or image.

What is not a violation of this policy?

Discussions related to child sexual exploitation as a phenomenon or attraction towards minors are permitted, provided they don't promote or glorify child sexual exploitation **in any way**. Artistic depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios e.g., works by internationally renowned artists that feature minors.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

What you can do

If you think you've found a Twitter account distributing or promoting child sexual exploitation, please report it through our designated [form](#). Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. For help on how to find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with NCMEC [here](#) or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the International Association of Internet Hotlines [website](#).

Our partners

We partner with multiple organisations around the world whose work is dedicated to fighting child sexual exploitation. You can find a full list of our safety partners [here](#).

Abusive behavior (20190607*)

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing or hoping serious harm on a person or group of people

We do not tolerate content that wishes, hopes or expresses a desire for death, serious bodily harm or fatal disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone's body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy (20190829*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., “I will kill you”.

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., “I hope you get cancer and die.”

Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”

Saying that a group of individuals deserve serious physical injury, e.g., “If this group of protesters don’t shut up, they deserve to be shot.”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can

[submit an appeal](#).

Glorifying self-harm and suicide (20190607*)

[Twitter Rules](#): You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Rationale

While we want people to feel safe sharing their thoughts and feelings, we draw the line at encouraging or promoting self-harm and suicidal behavior, which can pose safety risks for others. With that in mind, we apply a two-pronged approach to the issue: supporting people who are undergoing experiences with self-harm or suicidal thoughts, but prohibiting the promotion or encouragement of self-harming behaviors.

[Learn more](#) about how you can help someone dealing with thoughts of self-harm or suicide.

When this applies

Some examples of encouraging or promoting self-harm include (but may not be limited to) encouraging or glorifying:

- Acts of injuring oneself, such as cutting
- Eating disorders such as bulimia, anorexia, and serious illnesses that cause severe disturbances to a person's eating behaviors.
- The act of taking one's own life, including sharing strategies for suicide and participating in group suicides or suicide games

Some examples of behavior that are not considered glorifying or encouraging:

- Discussions or organizations focused on research, advocacy, and education of self-harm or suicide prevention
- Sharing coping mechanisms and resources for thoughts of self-harm and suicide

Consequences

Our approach towards suicide and self-harm depends on the type of content, whether or not the account is encouraging or promoting self-harm, and the account's previous record of violations.

The first time someone violates this policy, they will be required to remove the violating Tweet and be temporarily locked out of their account before they can Tweet again.

People who repeatedly violate this policy and/or accounts dedicated to promoting or encouraging self-harm and suicide may be suspended. If someone believes their account was suspended in error, they can [submit an appeal](#).

After we assess a report of self-harm or suicide, Twitter will contact the reported individual and let them know that someone who cares about them identified that they might be at risk. We will provide the reported individual with available resources and encourage them to seek help.

In some cases where content is not glorifying or encouraging suicide or self-harm, we may still interstitial Tweets as sensitive. Learn more about our [sensitive media policy](#).

Learn more about dealing with [self-harm and suicide](#).

Visit our Safety Center for a list of local [mental health resources](#).

Sensitive media policy (20191204*)

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20190607*)

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;

- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;
 - descriptions of physical appearance;
 - gossip, rumours, accusations, and allegations; and
 - screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy **(20191217*)**

Overview

November 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;

- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy **(20190925*)**

Overview

September 2019

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are; and

- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting.

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate; and
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in "follow trains," "decks," and "Retweet for Retweet" behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- “follow churn” – following and then unfollowing large numbers of accounts in an effort to inflate one’s own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account’s followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person’s browser (malware) or computer or to compromise a person’s privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don’t violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Blacklisting URLs

We blacklist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [financial scam policy](#), our [hacked materials policy](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Election integrity policy (20190607*)

Overview

April 2019

You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote.

The public conversation occurring on Twitter is never more important than during elections. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt elections, including through the distribution of false or misleading information about the electoral process or when or how to vote.

What is in violation of this policy?

We prohibit 3 categories of manipulative behavior and content under this policy:

Misleading information about how to participate

You can't share false or misleading information about how to participate in an election or other civic event. This includes but is not limited to:

- misleading information about how to vote or register to vote (for example, that you can vote by Tweet, text message, email, or phone call);
- misleading information about requirements for voting, including identification requirements; and
- misleading statements or information about the official, announced date or time of an election.

Voter suppression and intimidation

You can't share false or misleading information intended to intimidate or dissuade voters from participating in an election. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to polling places or elections;
- misleading claims about long lines or equipment problems at voting locations;
- misleading claims about voting procedures or techniques which could dissuade voters from participating in an election; and
- threats regarding voting locations (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account](#) policy.

What is not a violation of this policy?

Not all false or untrue information about politics or political events constitutes manipulation or interference in an election. In the absence of other violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election. Therefore, we enable reporting of false or misleading information about voting or registering to vote for people accessing Twitter in advance of major elections. We also work with select government and civil society partners in these countries to provide additional options for reporting and expedited review.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about voting**.
3. Select the option that best tells us how the Tweet is misleading about voting.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about voting**.
3. Select the option that best tells us how the Tweet is misleading about voting.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our election integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Impersonation policy (20191001*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and

- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Copyright policy \(20181206*\)](#)

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has

satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

<https://web.archive.org/web/20180601194918/https://help.twitter.com/en/rules-and-policies/copyright-policy> 1/3

27.11.2020 Copyright policy

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

<https://web.archive.org/web/20180601194918/https://help.twitter.com/en/rules-and-policies/copyright-policy> 2/3

27.11.2020 Copyright policy

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);

3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Our approach to policy development and enforcement philosophy

(20190520*)

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (2020113*)

The Twitter Rules

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;

- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Terrorism and violent extremism policy

(20190607*)

Overview

March 2019

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for terrorist organizations or violent extremist groups and individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments in this context are informed by national and international terrorism designations. We also assess organizations under our violent extremist group criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

We examine a group's activities both on and off Twitter to determine whether they engage in and/or promote violence against civilians to advance a political, religious and/or social cause.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a terrorist organization or violent extremist group;
- recruiting for a terrorist organization or violent extremist group;
- providing or distributing services (e.g., financial, media/propaganda) to further a terrorist organization's or violent extremist group's stated goals; and
- using the insignia or symbols of terrorist organizations or violent extremist groups to promote them.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to military or government entities.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wish for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy **(20190607*)**

Overview

March 2019

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives; and
- identifying alleged victims of childhood sexual exploitation by name or image.

What is not a violation of this policy?

Discussions related to child sexual exploitation as a phenomenon or attraction towards minors are permitted, provided they don't promote or glorify child sexual exploitation **in any way**. Artistic depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios e.g., works by internationally renowned artists that feature minors.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third

party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

What you can do

If you think you've found a Twitter account distributing or promoting child sexual exploitation, please report it through our designated [form](#). Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. For help on how to find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with NCMEC [here](#) or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the International Association of Internet Hotlines [website](#).

Our partners

We partner with multiple organisations around the world whose work is dedicated to fighting child sexual exploitation. You can find a full list of our safety partners [here](#).

Abusive behavior (20190607*)

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing or hoping serious harm on a person or group of people

We do not tolerate content that wishes, hopes or expresses a desire for death, serious bodily harm or fatal disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone's body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy (20190829*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically

marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you".

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., "I hope you get cancer and die."

Wishing for someone to fall victim to a serious accident, e.g., "I wish that you would get run over by a car next time you run your mouth."

Saying that a group of individuals deserve serious physical injury, e.g., "If this group of protesters don't shut up, they deserve to be shot."

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected

category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can

[submit an appeal](#).

Glorifying self-harm and suicide (20190607*)

[Twitter Rules](#): You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Rationale

While we want people to feel safe sharing their thoughts and feelings, we draw the line at encouraging or promoting self-harm and suicidal behavior, which can pose safety risks for others. With that in mind, we apply a two-pronged approach to the issue: supporting people who are undergoing experiences with self-harm or suicidal thoughts, but prohibiting the promotion or encouragement of self-harming behaviors.

[Learn more](#) about how you can help someone dealing with thoughts of self-harm or suicide.

When this applies

Some examples of encouraging or promoting self-harm include (but may not be limited to) encouraging or glorifying:

- Acts of injuring oneself, such as cutting
- Eating disorders such as bulimia, anorexia, and serious illnesses that cause severe disturbances to a person's eating behaviors.
- The act of taking one's own life, including sharing strategies for suicide and participating in group suicides or suicide games

Some examples of behavior that are not considered glorifying or encouraging:

- Discussions or organizations focused on research, advocacy, and education of self-harm or suicide prevention
- Sharing coping mechanisms and resources for thoughts of self-harm and suicide

Consequences

Our approach towards suicide and self-harm depends on the type of content, whether or not the account is encouraging or promoting self-harm, and the account's previous record of violations.

The first time someone violates this policy, they will be required to remove the violating Tweet and be temporarily locked out of their account before they can Tweet again.

People who repeatedly violate this policy and/or accounts dedicated to promoting or encouraging self-harm and suicide may be suspended. If someone believes their account was suspended in error, they can [submit an appeal](#).

After we assess a report of self-harm or suicide, Twitter will contact the reported individual and let them know that someone who cares about them identified that they might be at risk. We will provide the reported individual with available resources and encourage them to seek help.

In some cases where content is not glorifying or encouraging suicide or self-harm, we may still interstitial Tweets as sensitive. Learn more about our [sensitive media policy](#).

Learn more about dealing with [self-harm and suicide](#).

Visit our Safety Center for a list of local [mental health resources](#).

Sensitive media policy (20191204*)

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20190607*)

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;

- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;
 - descriptions of physical appearance;
 - gossip, rumours, accusations, and allegations; and
 - screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy **(20191217*)**

Overview

November 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;

- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy **(20190925*)**

Overview

September 2019

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are; and

- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting.

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate; and
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in "follow trains," "decks," and "Retweet for Retweet" behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- “follow churn” – following and then unfollowing large numbers of accounts in an effort to inflate one’s own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account’s followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person’s browser (malware) or computer or to compromise a person’s privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don’t violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Blacklisting URLs

We blacklist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [financial scam policy](#), our [hacked materials policy](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Election integrity policy (20190607*)

Overview

April 2019

You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote.

The public conversation occurring on Twitter is never more important than during elections. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt elections, including through the distribution of false or misleading information about the electoral process or when or how to vote.

What is in violation of this policy?

We prohibit 3 categories of manipulative behavior and content under this policy:

Misleading information about how to participate

You can't share false or misleading information about how to participate in an election or other civic event. This includes but is not limited to:

- misleading information about how to vote or register to vote (for example, that you can vote by Tweet, text message, email, or phone call);
- misleading information about requirements for voting, including identification requirements; and
- misleading statements or information about the official, announced date or time of an election.

Voter suppression and intimidation

You can't share false or misleading information intended to intimidate or dissuade voters from participating in an election. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to polling places or elections;
- misleading claims about long lines or equipment problems at voting locations;
- misleading claims about voting procedures or techniques which could dissuade voters from participating in an election; and
- threats regarding voting locations (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account](#) policy.

What is not a violation of this policy?

Not all false or untrue information about politics or political events constitutes manipulation or interference in an election. In the absence of other violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election. Therefore, we enable reporting of false or misleading information about voting or registering to vote for people accessing Twitter in advance of major elections. We also work with select government and civil society partners in these countries to provide additional options for reporting and expedited review.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about voting**.
3. Select the option that best tells us how the Tweet is misleading about voting.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about voting**.
3. Select the option that best tells us how the Tweet is misleading about voting.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our election integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Impersonation policy (20191001*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and

- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Copyright policy \(20181206*\)](#)

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has

satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

<https://web.archive.org/web/20180601194918/https://help.twitter.com/en/rules-and-policies/copyright-policy> 1/3

27.11.2020 Copyright policy

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

<https://web.archive.org/web/20180601194918/https://help.twitter.com/en/rules-and-policies/copyright-policy> 2/3

27.11.2020 Copyright policy

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);

3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Our approach to policy development and enforcement philosophy

(20190520*)

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (2020113*)

The Twitter Rules

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;

- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Terrorism and violent extremism policy

(20190607*)

Overview

March 2019

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for terrorist organizations or violent extremist groups and individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments in this context are informed by national and international terrorism designations. We also assess organizations under our violent extremist group criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

We examine a group's activities both on and off Twitter to determine whether they engage in and/or promote violence against civilians to advance a political, religious and/or social cause.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a terrorist organization or violent extremist group;
- recruiting for a terrorist organization or violent extremist group;
- providing or distributing services (e.g., financial, media/propaganda) to further a terrorist organization's or violent extremist group's stated goals; and
- using the insignia or symbols of terrorist organizations or violent extremist groups to promote them.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to military or government entities.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wish for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy **(20190607*)**

Overview

March 2019

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives; and
- identifying alleged victims of childhood sexual exploitation by name or image.

What is not a violation of this policy?

Discussions related to child sexual exploitation as a phenomenon or attraction towards minors are permitted, provided they don't promote or glorify child sexual exploitation **in any way**. Artistic depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios e.g., works by internationally renowned artists that feature minors.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third

party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

What you can do

If you think you've found a Twitter account distributing or promoting child sexual exploitation, please report it through our designated [form](#). Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. For help on how to find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with NCMEC [here](#) or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the International Association of Internet Hotlines [website](#).

Our partners

We partner with multiple organisations around the world whose work is dedicated to fighting child sexual exploitation. You can find a full list of our safety partners [here](#).

Abusive behavior (20190607*)

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing or hoping serious harm on a person or group of people

We do not tolerate content that wishes, hopes or expresses a desire for death, serious bodily harm or fatal disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone's body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy (2020114)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically

marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you".

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., "I hope you get cancer and die."

Wishing for someone to fall victim to a serious accident, e.g., "I wish that you would get run over by a car next time you run your mouth."

Saying that a group of individuals deserve serious physical injury, e.g., "If this group of protesters don't shut up, they deserve to be shot."

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected

category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can

[submit an appeal](#).

Glorifying self-harm and suicide (20190607*)

[Twitter Rules](#): You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Rationale

While we want people to feel safe sharing their thoughts and feelings, we draw the line at encouraging or promoting self-harm and suicidal behavior, which can pose safety risks for others. With that in mind, we apply a two-pronged approach to the issue: supporting people who are undergoing experiences with self-harm or suicidal thoughts, but prohibiting the promotion or encouragement of self-harming behaviors.

[Learn more](#) about how you can help someone dealing with thoughts of self-harm or suicide.

When this applies

Some examples of encouraging or promoting self-harm include (but may not be limited to) encouraging or glorifying:

- Acts of injuring oneself, such as cutting
- Eating disorders such as bulimia, anorexia, and serious illnesses that cause severe disturbances to a person's eating behaviors.
- The act of taking one's own life, including sharing strategies for suicide and participating in group suicides or suicide games

Some examples of behavior that are not considered glorifying or encouraging:

- Discussions or organizations focused on research, advocacy, and education of self-harm or suicide prevention
- Sharing coping mechanisms and resources for thoughts of self-harm and suicide

Consequences

Our approach towards suicide and self-harm depends on the type of content, whether or not the account is encouraging or promoting self-harm, and the account's previous record of violations.

The first time someone violates this policy, they will be required to remove the violating Tweet and be temporarily locked out of their account before they can Tweet again.

People who repeatedly violate this policy and/or accounts dedicated to promoting or encouraging self-harm and suicide may be suspended. If someone believes their account was suspended in error, they can [submit an appeal](#).

After we assess a report of self-harm or suicide, Twitter will contact the reported individual and let them know that someone who cares about them identified that they might be at risk. We will provide the reported individual with available resources and encourage them to seek help.

In some cases where content is not glorifying or encouraging suicide or self-harm, we may still interstitial Tweets as sensitive. Learn more about our [sensitive media policy](#).

Learn more about dealing with [self-harm and suicide](#).

Visit our Safety Center for a list of local [mental health resources](#).

Sensitive media policy (20191204*)

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20190607*)

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;

- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;
 - descriptions of physical appearance;
 - gossip, rumours, accusations, and allegations; and
 - screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy **(20191217*)**

Overview

November 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;

- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy **(20190925*)**

Overview

September 2019

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are; and

- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting.

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate; and
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in "follow trains," "decks," and "Retweet for Retweet" behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- “follow churn” – following and then unfollowing large numbers of accounts in an effort to inflate one’s own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account’s followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person’s browser (malware) or computer or to compromise a person’s privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don’t violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Blacklisting URLs

We blacklist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [financial scam policy](#), our [hacked materials policy](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Election integrity policy (20190607*)

Overview

April 2019

You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote.

The public conversation occurring on Twitter is never more important than during elections. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt elections, including through the distribution of false or misleading information about the electoral process or when or how to vote.

What is in violation of this policy?

We prohibit 3 categories of manipulative behavior and content under this policy:

Misleading information about how to participate

You can't share false or misleading information about how to participate in an election or other civic event. This includes but is not limited to:

- misleading information about how to vote or register to vote (for example, that you can vote by Tweet, text message, email, or phone call);
- misleading information about requirements for voting, including identification requirements; and
- misleading statements or information about the official, announced date or time of an election.

Voter suppression and intimidation

You can't share false or misleading information intended to intimidate or dissuade voters from participating in an election. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to polling places or elections;
- misleading claims about long lines or equipment problems at voting locations;
- misleading claims about voting procedures or techniques which could dissuade voters from participating in an election; and
- threats regarding voting locations (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account](#) policy.

What is not a violation of this policy?

Not all false or untrue information about politics or political events constitutes manipulation or interference in an election. In the absence of other violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election. Therefore, we enable reporting of false or misleading information about voting or registering to vote for people accessing Twitter in advance of major elections. We also work with select government and civil society partners in these countries to provide additional options for reporting and expedited review.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about voting**.
3. Select the option that best tells us how the Tweet is misleading about voting.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about voting**.
3. Select the option that best tells us how the Tweet is misleading about voting.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our election integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Impersonation policy (20191001*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and

- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Copyright policy \(20181206*\)](#)

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has

satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

<https://web.archive.org/web/20180601194918/https://help.twitter.com/en/rules-and-policies/copyright-policy> 1/3

27.11.2020 Copyright policy

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

<https://web.archive.org/web/20180601194918/https://help.twitter.com/en/rules-and-policies/copyright-policy> 2/3

27.11.2020 Copyright policy

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);

3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Our approach to policy development and enforcement philosophy

(20190520*)

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

[Violence](#): You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

[Terrorism/violent extremism](#): You may not threaten or promote terrorism or violent extremism. [Learn more](#).

[Child sexual exploitation](#): We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

[Abuse/harassment](#): You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

[Hateful conduct](#): You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

[Suicide or self-harm](#): You may not promote or encourage suicide or self-harm. [Learn more](#).

[Sensitive media, including graphic violence and adult content](#): You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

[Illegal or certain regulated goods or services](#): You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

[Private information](#): You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

[Non-consensual nudity](#): You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

[Platform manipulation and spam](#): You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

[Election integrity](#): You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

[Impersonation](#): You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Synthetic and manipulated media: You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Terrorism and violent extremism policy

Overview

March 2019

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for terrorist organizations or violent extremist groups and individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments in this context are informed by national and international terrorism designations. We also assess organizations under our violent extremist group criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;

- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

We examine a group's activities both on and off Twitter to determine whether they engage in and/or promote violence against civilians to advance a political, religious and/or social cause.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a terrorist organization or violent extremist group;
- recruiting for a terrorist organization or violent extremist group;
- providing or distributing services (e.g., financial, media/propaganda) to further a terrorist organization's or violent extremist group's stated goals; and
- using the insignia or symbols of terrorist organizations or violent extremist groups to promote them.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to military or government entities.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wish for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy

Overview

March 2019

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives; and

- identifying alleged victims of childhood sexual exploitation by name or image.

What is not a violation of this policy?

Discussions related to child sexual exploitation as a phenomenon or attraction towards minors are permitted, provided they don't promote or glorify child sexual exploitation **in any way**. Artistic depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios e.g., works by internationally renowned artists that feature minors.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

What you can do

If you think you've found a Twitter account distributing or promoting child sexual exploitation, please report it through our designated [form](#). Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. For help on how to find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with NCMEC [here](#) or with your local

law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the International Association of Internet Hotlines [website](#).

Our partners

We partner with multiple organisations around the world whose work is dedicated to fighting child sexual exploitation. You can find a full list of our safety partners [here](#).

Abusive behavior

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing or hoping serious harm on a person or group of people

We do not tolerate content that wishes, hopes or expresses a desire for death, serious bodily harm or fatal disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone's body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name,

or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you".

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., “I hope you get cancer and die.”

Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”

Saying that a group of individuals deserve serious physical injury, e.g., “If this group of protesters don’t shut up, they deserve to be shot.”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).
If someone believes their account was suspended in error, they can
[submit an appeal](#).

Glorifying self-harm and suicide

Twitter Rules: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Rationale

While we want people to feel safe sharing their thoughts and feelings, we draw the line at encouraging or promoting self-harm and suicidal behavior, which can pose safety risks for others. With that in mind, we apply a two-pronged approach to the issue: supporting people who are undergoing experiences with self-harm or suicidal thoughts, but prohibiting the promotion or encouragement of self-harming behaviors.

[Learn more](#) about how you can help someone dealing with thoughts of self-harm or suicide.

When this applies

Some examples of encouraging or promoting self-harm include (but may not be limited to) encouraging or glorifying:

- Acts of injuring oneself, such as cutting
- Eating disorders such as bulimia, anorexia, and serious illnesses that cause severe disturbances to a person's eating behaviors.
- The act of taking one's own life, including sharing strategies for suicide and participating in group suicides or suicide games

Some examples of behavior that are not considered glorifying or encouraging:

- Discussions or organizations focused on research, advocacy, and education of self-harm or suicide prevention
- Sharing coping mechanisms and resources for thoughts of self-harm and suicide

Consequences

Our approach towards suicide and self-harm depends on the type of content, whether or not the account is encouraging or promoting self-harm, and the account's previous record of violations.

The first time someone violates this policy, they will be required to remove the violating Tweet and be temporarily locked out of their account before they can Tweet again.

People who repeatedly violate this policy and/or accounts dedicated to promoting or encouraging self-harm and suicide may be suspended. If someone believes their account was suspended in error, they can [submit an appeal](#).

After we assess a report of self-harm or suicide, Twitter will contact the reported individual and let them know that someone who cares about them identified that they might be at risk. We will provide the reported individual with available resources and encourage them to seek help.

In some cases where content is not glorifying or encouraging suicide or self-harm, we may still interstitial Tweets as sensitive. Learn more about our [sensitive media policy](#).

Learn more about dealing with [self-harm and suicide](#).

Visit our Safety Center for a list of local [mental health resources](#).

Sensitive media policy

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;

- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing

their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;
 - descriptions of physical appearance;
 - gossip, rumours, accusations, and allegations; and
 - screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy

Overview

March 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy

Overview

March 2019

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are; and
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting.

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate; and
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in “follow trains,” “decks,” and “Retweet for Retweet” behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- "follow churn" – following and then unfollowing large numbers of accounts in an effort to inflate one's own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account's followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person's browser (malware) or computer or to compromise a person's privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don't violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Blacklisting URLs

We blacklist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;

- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [hacked materials policy](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Election integrity policy

Overview

April 2019

You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote.

The public conversation occurring on Twitter is never more important than during elections. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt elections, including through the distribution of false or misleading information about the electoral process or when or how to vote.

What is in violation of this policy?

We prohibit 3 categories of manipulative behavior and content under this policy:

Misleading information about how to participate

You can't share false or misleading information about how to participate in an election or other civic event. This includes but is not limited to:

- misleading information about how to vote or register to vote (for example, that you can vote by Tweet, text message, email, or phone call);
- misleading information about requirements for voting, including identification requirements; and
- misleading statements or information about the official, announced date or time of an election.

Voter suppression and intimidation

You can't share false or misleading information intended to intimidate or dissuade voters from participating in an election. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to polling places or elections;

- misleading claims about long lines or equipment problems at voting locations;
- misleading claims about voting procedures or techniques which could dissuade voters from participating in an election; and
- threats regarding voting locations (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account](#) policy.

What is not a violation of this policy?

Not all false or untrue information about politics or political events constitutes manipulation or interference in an election. In the absence of other violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election. Therefore, we enable reporting of false or misleading information about voting or registering to vote for people accessing Twitter in advance of major elections. We also work with select government and civil society partners in these countries to provide additional options for reporting and expedited review.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about voting**.
3. Select the option that best tells us how the Tweet is misleading about voting.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about voting**.
3. Select the option that best tells us how the Tweet is misleading about voting.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our election integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Synthetic and manipulated media policy

Overview

You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context.

You should be able to find reliable information on Twitter. That means understanding whether the content you see is real or fabricated and having the ability to find more context about what you see on Twitter. Therefore, we may label Tweets that include media (videos, audio, and images) that have been deceptively altered or fabricated. In addition, you may not share deceptively altered media on Twitter in ways that mislead or deceive people about the media's authenticity where threats to physical safety or other serious harm may result.

We use the following criteria as we consider Tweets and media for labeling or removal under this policy as part of our ongoing work to enforce our rules and ensure healthy and safe conversation on Twitter (additional information is available below):

Is the content significantly and deceptively altered or fabricated?	Is the content shared in a deceptive manner?	Is the content likely to impact public safety or cause serious harm?	
✓	✗	✗	Content may be labeled.
✗	✓	✗	Content may be labeled.
✓	✗	✓	Content is likely to be labeled, or may be removed.*
✓	✓	✗	Content is likely to be labeled.
✓	✓	✓	Content is likely to be removed.

1. Is the content synthetic or manipulated?

In order for content to be labeled or removed under this policy, we must have reason to believe that media, or the context in which media are presented, are significantly and deceptively altered or manipulated. Synthetic and manipulated media take many different forms and people can employ a wide range of technologies to produce these media. In assessing whether media have been significantly and deceptively altered or fabricated, some of the factors we consider include:

- whether the content has been substantially edited in a manner that fundamentally alters its composition, sequence, timing, or framing;
- any visual or auditory information (such as new video frames, overdubbed audio, or modified subtitles) that has been added or removed; and
- whether media depicting a real person have been fabricated or simulated

We are most likely to take action (either labeling or removal, as described below) on more significant forms of alteration, such as wholly synthetic audio or video or content that has been doctored (spliced and reordered, slowed down) to change its meaning. Subtler forms of manipulated media, such as isolative editing, omission of context, or presentation with false context, may be labeled or removed on a case-by-case basis.

We will not take action to label or remove media that have been edited in ways that do not fundamentally alter their meaning, such as retouched photos or color-corrected videos.

In order to determine if media have been significantly and deceptively altered or fabricated, we may use our own technology or receive reports through partnerships with third parties. In situations where we are unable to reliably determine if media have been altered or fabricated, we may not take action to label or remove them.

2. Is the content shared in a deceptive manner?

We also consider whether the context in which media are shared could result in confusion or misunderstanding or suggests a deliberate intent to deceive people about the nature or origin of the content, for example by falsely claiming that it depicts reality. We assess the context provided alongside media to see whether it makes clear that the media have been altered or fabricated. Some of the types of context we assess in order to make this determination include:

- The text of the Tweet accompanying or within media
- Metadata associated with media
- Information on the profile of the account sharing media
- Websites linked in the Tweet, or in the profile of the account sharing media

3. Is the content likely to impact public safety or cause serious harm?

Tweets that share synthetic and manipulated media are subject to removal under this policy if they are likely to cause serious harm. Some specific harms we consider include:

- Threats to the physical safety of a person or group
- Risk of mass violence or widespread civil unrest
- Threats to the privacy or ability of a person or group to freely express themselves or participate in civic events, such as:
 - Stalking or unwanted and obsessive attention
 - Targeted content that includes tropes, epithets, or material that aims to silence someone
 - Voter suppression or intimidation

While we have other rules also intended to address these forms of harm, including our policies on violent threats, election integrity, and hateful conduct, we will err toward removal in borderline cases that might otherwise not violate existing rules for Tweets that include synthetic or manipulated media.

We also consider the time frame within which the content may be likely to impact public safety or cause serious harm, and are more likely to remove content under this policy if we find that immediate harms are likely to result from the content's presence on Twitter.

Note: We may also take action on synthetic and manipulated content under our [non-consensual nudity policy](#) (such as pornographic media altered to insert the faces of people not actually involved) or other parts of the Twitter Rules.

Labeling and removal

In most cases, if we have reason to believe that media shared in a Tweet have been significantly and deceptively altered or fabricated, we will provide additional context on Tweets sharing the media where they appear on Twitter. This means we may:

- Apply a label to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended; and/or
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or landing page.

In most cases, we will take all of the above actions on Tweets we label.

Media that meet **all three of the criteria** defined above—i.e. that are synthetic or manipulated, shared in a deceptive manner, and is likely to cause harm—may not be shared on Twitter and are subject to removal. Accounts engaging in repeated or severe violations of this policy may be permanently suspended.

* Other parts of the Twitter Rules apply and may lead to the removal of the content, particularly where there is high likelihood of severe harm, such as a threat to someone's life or physical safety.

Trademark policy

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

[**Copyright policy**](#)

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity’s name, please review Twitter’s [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and

6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to [twitter.com](#), you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information

concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Our approach to policy development and enforcement philosophy

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a

set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (20200204*)

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Synthetic and manipulated media: You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.

5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Terrorism and violent extremism policy (20190607*)

Overview

March 2019

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for terrorist organizations or violent extremist groups and individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments in this context are informed by national and international terrorism designations. We also assess organizations under our violent extremist group criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

We examine a group's activities both on and off Twitter to determine whether they engage in and/or promote violence against civilians to advance a political, religious and/or social cause.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a terrorist organization or violent extremist group;
- recruiting for a terrorist organization or violent extremist group;
- providing or distributing services (e.g., financial, media/propaganda) to further a terrorist organization's or violent extremist group's stated goals; and
- using the insignia or symbols of terrorist organizations or violent extremist groups to promote them.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to military or government entities.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wish for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy

(20190607*)

Overview

March 2019

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;

- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives; and
- identifying alleged victims of childhood sexual exploitation by name or image.

What is not a violation of this policy?

Discussions related to child sexual exploitation as a phenomenon or attraction towards minors are permitted, provided they don't promote or glorify child sexual exploitation **in any way**. Artistic depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios e.g., works by internationally renowned artists that feature minors.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

What you can do

If you think you've found a Twitter account distributing or promoting child sexual exploitation, please report it through our designated [form](#). Please provide the username

and links to the relevant tweets that lead you to believe the account should be investigated. For help on how to find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with NCMEC [here](#) or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the International Association of Internet Hotlines [website](#).

Our partners

We partner with multiple organisations around the world whose work is dedicated to fighting child sexual exploitation. You can find a full list of our safety partners [here](#).

Abusive behavior (20190607*)

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing or hoping serious harm on a person or group of people

We do not tolerate content that wishes, hopes or expresses a desire for death, serious bodily harm or fatal disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone's body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy (20200331*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you".

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., “I hope you get cancer and die.”

Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”

Saying that a group of individuals deserve serious physical injury, e.g., “If this group of protesters don’t shut up, they deserve to be shot.”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion, age, disability, or serious disease.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can

[submit an appeal](#).

Glorifying self-harm and suicide (20190607*)

Twitter Rules: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Rationale

While we want people to feel safe sharing their thoughts and feelings, we draw the line at encouraging or promoting self-harm and suicidal behavior, which can pose safety risks for others. With that in mind, we apply a two-pronged approach to the issue: supporting people who are undergoing experiences with self-harm or suicidal thoughts, but prohibiting the promotion or encouragement of self-harming behaviors.

[Learn more](#) about how you can help someone dealing with thoughts of self-harm or suicide.

When this applies

Some examples of encouraging or promoting self-harm include (but may not be limited to) encouraging or glorifying:

- Acts of injuring oneself, such as cutting
- Eating disorders such as bulimia, anorexia, and serious illnesses that cause severe disturbances to a person's eating behaviors.
- The act of taking one's own life, including sharing strategies for suicide and participating in group suicides or suicide games

Some examples of behavior that are not considered glorifying or encouraging:

- Discussions or organizations focused on research, advocacy, and education of self-harm or suicide prevention
- Sharing coping mechanisms and resources for thoughts of self-harm and suicide

Consequences

Our approach towards suicide and self-harm depends on the type of content, whether or not the account is encouraging or promoting self-harm, and the account's previous record of violations.

The first time someone violates this policy, they will be required to remove the violating Tweet and be temporarily locked out of their account before they can Tweet again.

People who repeatedly violate this policy and/or accounts dedicated to promoting or encouraging self-harm and suicide may be suspended. If someone believes their account was suspended in error, they can [submit an appeal](#).

After we assess a report of self-harm or suicide, Twitter will contact the reported individual and let them know that someone who cares about them identified that they might be at risk. We will provide the reported individual with available resources and encourage them to seek help.

In some cases where content is not glorifying or encouraging suicide or self-harm, we may still interstitial Tweets as sensitive. Learn more about our [sensitive media policy](#).

Learn more about dealing with [self-harm and suicide](#).

Visit our Safety Center for a list of local [mental health resources](#).

Sensitive media policy (20191204*)

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.

3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20190607*)

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their

information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;

- location information related to commercial property or places of business, where this information is publicly available;
- descriptions of physical appearance;
- gossip, rumours, accusations, and allegations; and
- screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy

(20191217*)

Overview

November 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy (20190925*)

Overview

September 2019

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are; and
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting.

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:

- posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
- engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate; and
- coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in “follow trains,” “decks,” and “Retweet for Retweet” behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- “follow churn” – following and then unfollowing large numbers of accounts in an effort to inflate one's own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account's followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person's browser (malware) or computer or to compromise a person's privacy (phishing); and

- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don't violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

- Select **Report Tweet** from the icon.
- Select **It's suspicious or spam**.
- Select the option that best tells us how the Tweet is suspicious or spreading spam.
- Submit your report.

Desktop

You can report this content via desktop as follows:

- Select **Report Tweet** from the icon.
- Select **It's suspicious or spam**.
- Select the option that best tells us how the Tweet is suspicious or spreading spam.
- Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Blacklisting URLs

We blacklist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [financial scam policy](#), our [hacked materials policy](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Election integrity policy (20190607*)

Overview

April 2019

You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote.

The public conversation occurring on Twitter is never more important than during elections. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt elections, including through the distribution of false or misleading information about the electoral process or when or how to vote.

What is in violation of this policy?

We prohibit 3 categories of manipulative behavior and content under this policy:

Misleading information about how to participate

You can't share false or misleading information about how to participate in an election or other civic event. This includes but is not limited to:

- misleading information about how to vote or register to vote (for example, that you can vote by Tweet, text message, email, or phone call);
- misleading information about requirements for voting, including identification requirements; and
- misleading statements or information about the official, announced date or time of an election.

Voter suppression and intimidation

You can't share false or misleading information intended to intimidate or dissuade voters from participating in an election. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to polling places or elections;
- misleading claims about long lines or equipment problems at voting locations;
- misleading claims about voting procedures or techniques which could dissuade voters from participating in an election; and
- threats regarding voting locations (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account](#) policy.

What is not a violation of this policy?

Not all false or untrue information about politics or political events constitutes manipulation or interference in an election. In the absence of other violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election. Therefore, we enable reporting of false or misleading information about voting or registering to vote for people accessing Twitter in advance of major elections. We also work with select government and civil society partners in these countries to provide additional options for reporting and expedited review.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about voting**.
3. Select the option that best tells us how the Tweet is misleading about voting.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about voting**.
3. Select the option that best tells us how the Tweet is misleading about voting.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our election integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Impersonation policy (20191001*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Synthetic and manipulated media policy (20200204*)

Overview

You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context.

You should be able to find reliable information on Twitter. That means understanding whether the content you see is real or fabricated and having the ability to find more context about what you see on Twitter. Therefore, we may label Tweets that include media (videos, audio, and images) that have been deceptively altered or fabricated. In addition, you may not share deceptively altered media on Twitter in ways that mislead or deceive people about the media's authenticity where threats to physical safety or other serious harm may result.

We use the following criteria as we consider Tweets and media for labeling or removal under this policy as part of our ongoing work to enforce our rules and ensure healthy and safe conversation on Twitter (additional information is available below):

Is the content significantly and deceptively altered or fabricated?	Is the content shared in a deceptive manner?	Is the content likely to impact public safety or cause serious harm?	
✓	✗	✗	Content may be labeled.
✗	✓	✗	Content may be labeled.
✓	✗	✓	Content is likely to be labeled, or may be removed.*
✓	✓	✗	Content is likely to be labeled.
✓	✓	✓	Content is likely to be removed.

1. Is the content synthetic or manipulated?

In order for content to be labeled or removed under this policy, we must have reason to believe that media, or the context in which media are presented, are significantly and deceptively altered or manipulated. Synthetic and manipulated media take many different forms and people can employ a wide range of technologies to produce these media. In assessing whether media have been significantly and deceptively altered or fabricated, some of the factors we consider include:

- whether the content has been substantially edited in a manner that fundamentally alters its composition, sequence, timing, or framing;

- any visual or auditory information (such as new video frames, overdubbed audio, or modified subtitles) that has been added or removed; and
- whether media depicting a real person have been fabricated or simulated

We are most likely to take action (either labeling or removal, as described below) on more significant forms of alteration, such as wholly synthetic audio or video or content that has been doctored (spliced and reordered, slowed down) to change its meaning. Subtler forms of manipulated media, such as isolative editing, omission of context, or presentation with false context, may be labeled or removed on a case-by-case basis.

We will not take action to label or remove media that have been edited in ways that do not fundamentally alter their meaning, such as retouched photos or color-corrected videos.

In order to determine if media have been significantly and deceptively altered or fabricated, we may use our own technology or receive reports through partnerships with third parties. In situations where we are unable to reliably determine if media have been altered or fabricated, we may not take action to label or remove them.

2. Is the content shared in a deceptive manner?

We also consider whether the context in which media are shared could result in confusion or misunderstanding or suggests a deliberate intent to deceive people about the nature or origin of the content, for example by falsely claiming that it depicts reality. We assess the context provided alongside media to see whether it makes clear that the media have been altered or fabricated. Some of the types of context we assess in order to make this determination include:

- The text of the Tweet accompanying or within media
- Metadata associated with media
- Information on the profile of the account sharing media
- Websites linked in the Tweet, or in the profile of the account sharing media

3. Is the content likely to impact public safety or cause serious harm?

Tweets that share synthetic and manipulated media are subject to removal under this policy if they are likely to cause serious harm. Some specific harms we consider include:

- Threats to the physical safety of a person or group
- Risk of mass violence or widespread civil unrest
- Threats to the privacy or ability of a person or group to freely express themselves or participate in civic events, such as:
 - Stalking or unwanted and obsessive attention
 - Targeted content that includes tropes, epithets, or material that aims to silence someone
 - Voter suppression or intimidation

While we have other rules also intended to address these forms of harm, including our policies on violent threats, election integrity, and hateful conduct, we will err toward removal in borderline cases that might otherwise not violate existing rules for Tweets that include synthetic or manipulated media.

We also consider the time frame within which the content may be likely to impact public safety or cause serious harm, and are more likely to remove content under this policy if we find that immediate harms are likely to result from the content's presence on Twitter.

Note: We may also take action on synthetic and manipulated content under our [non-consensual nudity policy](#) (such as pornographic media altered to insert the faces of people not actually involved) or other parts of the Twitter Rules.

Labeling and removal

In most cases, if we have reason to believe that media shared in a Tweet have been significantly and deceptively altered or fabricated, we will provide additional context on Tweets sharing the media where they appear on Twitter. This means we may:

- Apply a label to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended; and/or
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or landing page.

In most cases, we will take all of the above actions on Tweets we label.

Media that meet **all three of the criteria** defined above—i.e. that are synthetic or manipulated, shared in a deceptive manner, and is likely to cause harm—may not be shared on Twitter and are subject to removal. Accounts engaging in repeated or severe violations of this policy may be permanently suspended.

** Other parts of the Twitter Rules apply and may lead to the removal of the content, particularly where there is high likelihood of severe harm, such as a threat to someone's life or physical safety.*

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Copyright policy (20181206*)

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

<https://web.archive.org/web/20180601194918/https://help.twitter.com/en/rules-and-policies/copyright-policy> 1/3

27.11.2020 Copyright policy

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

<https://web.archive.org/web/20180601194918/https://help.twitter.com/en/rules-and-policies/copyright-policy> 2/3

27.11.2020 Copyright policy

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept

service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Our approach to policy development and enforcement philosophy **(20190520*)**

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online

speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

[Violence](#): You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

[Terrorism/violent extremism](#): You may not threaten or promote terrorism or violent extremism. [Learn more](#).

[Child sexual exploitation](#): We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

[Abuse/harassment](#): You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

[Hateful conduct](#): You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

[Suicide or self-harm](#): You may not promote or encourage suicide or self-harm. [Learn more](#).

[Sensitive media, including graphic violence and adult content](#): You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

[Illegal or certain regulated goods or services](#): You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

[Private information](#): You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

[Non-consensual nudity](#): You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

[Platform manipulation and spam](#): You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

[Election integrity](#): You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

[Impersonation](#): You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Synthetic and manipulated media: You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Terrorism and violent extremism policy

Overview

March 2019

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for terrorist organizations or violent extremist groups and individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments in this context are informed by national and international terrorism designations. We also assess organizations under our violent extremist group criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;

- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

We examine a group's activities both on and off Twitter to determine whether they engage in and/or promote violence against civilians to advance a political, religious and/or social cause.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a terrorist organization or violent extremist group;
- recruiting for a terrorist organization or violent extremist group;
- providing or distributing services (e.g., financial, media/propaganda) to further a terrorist organization's or violent extremist group's stated goals; and
- using the insignia or symbols of terrorist organizations or violent extremist groups to promote them.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to military or government entities.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wish for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy

Overview

March 2019

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives; and

- identifying alleged victims of childhood sexual exploitation by name or image.

What is not a violation of this policy?

Discussions related to child sexual exploitation as a phenomenon or attraction towards minors are permitted, provided they don't promote or glorify child sexual exploitation **in any way**. Artistic depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios e.g., works by internationally renowned artists that feature minors.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

What you can do

If you think you've found a Twitter account distributing or promoting child sexual exploitation, please report it through our designated [form](#). Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. For help on how to find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with NCMEC [here](#) or with your local

law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the International Association of Internet Hotlines [website](#).

Our partners

We partner with multiple organisations around the world whose work is dedicated to fighting child sexual exploitation. You can find a full list of our safety partners [here](#).

Abusive behavior

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing or hoping serious harm on a person or group of people

We do not tolerate content that wishes, hopes or expresses a desire for death, serious bodily harm or fatal disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone's body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name,

or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you".

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., “I hope you get cancer and die.”

Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”

Saying that a group of individuals deserve serious physical injury, e.g., “If this group of protesters don’t shut up, they deserve to be shot.”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion, age, disability, or serious disease.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).
If someone believes their account was suspended in error, they can
[submit an appeal](#).

Glorifying self-harm and suicide

Twitter Rules: You may not promote or encourage suicide or self-harm. When we receive reports that a person is threatening suicide or self-harm, we may take a number of steps to assist them, such as reaching out to that person and providing resources such as contact information for our mental health partners.

Rationale

While we want people to feel safe sharing their thoughts and feelings, we draw the line at encouraging or promoting self-harm and suicidal behavior, which can pose safety risks for others. With that in mind, we apply a two-pronged approach to the issue: supporting people who are undergoing experiences with self-harm or suicidal thoughts, but prohibiting the promotion or encouragement of self-harming behaviors.

[Learn more](#) about how you can help someone dealing with thoughts of self-harm or suicide.

When this applies

Some examples of encouraging or promoting self-harm include (but may not be limited to) encouraging or glorifying:

- Acts of injuring oneself, such as cutting
- Eating disorders such as bulimia, anorexia, and serious illnesses that cause severe disturbances to a person's eating behaviors.
- The act of taking one's own life, including sharing strategies for suicide and participating in group suicides or suicide games

Some examples of behavior that are not considered glorifying or encouraging:

- Discussions or organizations focused on research, advocacy, and education of self-harm or suicide prevention
- Sharing coping mechanisms and resources for thoughts of self-harm and suicide

Consequences

Our approach towards suicide and self-harm depends on the type of content, whether or not the account is encouraging or promoting self-harm, and the account's previous record of violations.

The first time someone violates this policy, they will be required to remove the violating Tweet and be temporarily locked out of their account before they can Tweet again.

People who repeatedly violate this policy and/or accounts dedicated to promoting or encouraging self-harm and suicide may be suspended. If someone believes their account was suspended in error, they can [submit an appeal](#).

After we assess a report of self-harm or suicide, Twitter will contact the reported individual and let them know that someone who cares about them identified that they might be at risk. We will provide the reported individual with available resources and encourage them to seek help.

In some cases where content is not glorifying or encouraging suicide or self-harm, we may still interstitial Tweets as sensitive. Learn more about our [sensitive media policy](#).

Learn more about dealing with [self-harm and suicide](#).

Visit our Safety Center for a list of local [mental health resources](#).

Sensitive media policy

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;

- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing

their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;
 - descriptions of physical appearance;
 - gossip, rumours, accusations, and allegations; and
 - screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy

Overview

March 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy

Overview

March 2019

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are; and
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting.

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate; and
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in “follow trains,” “decks,” and “Retweet for Retweet” behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- "follow churn" – following and then unfollowing large numbers of accounts in an effort to inflate one's own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account's followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person's browser (malware) or computer or to compromise a person's privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don't violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Blacklisting URLs

We blacklist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;

- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [hacked materials policy](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Civic integrity policy

Overview

May 2020

You may not use Twitter's services for the purpose of manipulating or interfering in elections or other civic processes. This includes posting or sharing content that may suppress participation or mislead people about when, where, or how to participate in a civic process.

The public conversation occurring on Twitter is never more important than during elections and other civic events. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt civic processes, including through the distribution of false or misleading information about the procedures or circumstances around participation in a civic process.

What is a civic process?

Twitter considers civic processes to be events or procedures mandated, organized, and conducted by the governing and/or electoral body of a country, state, region, district, or municipality to address a matter of common concern through public participation. Some examples of civic processes may include:

- Political elections
- Censuses
- Major referenda and ballot initiatives

What is in violation of this policy?

We prohibit 3 categories of manipulative behavior and content under this policy:

Misleading information about how to participate

You can't share false or misleading information about how to participate in an election or other civic process. This includes but is not limited to:

- misleading information about procedures to participate in a civic process (for example, that you can vote by Tweet, text message, email, or phone call in jurisdictions where these are not a possibility);
- misleading information about requirements for participation, including identification or citizenship requirements; and
- misleading statements or information about the official, announced date or time of a civic process.

Suppression and intimidation

You can't share false or misleading information intended to intimidate or dissuade people from participating in an election or other civic process. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to voting in an election, polling places, or collecting census information;
- misleading claims about long lines, equipment problems, or other disruptions at voting locations during election periods;
- misleading claims about process procedures or techniques which could dissuade people from participating; and
- threats regarding voting locations or other key places or events (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account policy](#).

What is not a violation of this policy?

Not all false or untrue information about politics or civic processes constitutes manipulation or interference. In the absence of other policy violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected or appointed official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information;
- voting and audience participation for competitions, game shows, or other entertainment purposes;
- broad, non-specific statements about the integrity of elections or civic processes (such as unsubstantiated claims that an election is "rigged"); and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election or civic process. Therefore, we enable reporting of false or misleading information about civic processes in advance of major events, for people located in the relevant countries and locations. We also work with select government and civil society partners in these countries to provide additional channels for reporting and expedited review.

For civic processes with multiple stages or parts, such as primary elections or lengthy campaigns, reporting will be enabled leading up to the first officially-sanctioned event associated with the civic process.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country at the relevant time, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in civic processes.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in a civic process.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our civic integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order

to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our range of enforcement options and our approach to policy development and enforcement.

Synthetic and manipulated media policy

Overview

You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context.

You should be able to find reliable information on Twitter. That means understanding whether the content you see is real or fabricated and having the ability to find more context about what you see on Twitter. Therefore, we may label Tweets that include media (videos, audio, and images) that have been deceptively altered or fabricated. In addition, you may not share deceptively altered media on Twitter in ways that mislead or deceive people about the media's authenticity where threats to physical safety or other serious harm may result.

We use the following criteria as we consider Tweets and media for labeling or removal under this policy as part of our ongoing work to enforce our rules and ensure healthy and safe conversation on Twitter (additional information is available below):

Is the content significantly and deceptively altered or fabricated?	Is the content shared in a deceptive manner?	Is the content likely to impact public safety or cause serious harm?	
✓	✗	✗	Content may be labeled.
✗	✓	✗	Content may be labeled.
✓	✗	✓	Content is likely to be labeled, or may be removed.*
✓	✓	✗	Content is likely to be labeled.
✓	✓	✓	Content is likely to be removed.

1. Is the content synthetic or manipulated?

In order for content to be labeled or removed under this policy, we must have reason to believe that media, or the context in which media are presented, are significantly and deceptively altered or manipulated. Synthetic and manipulated media take many different forms and people can employ a wide range of technologies to produce these media. In assessing whether media have been significantly and deceptively altered or fabricated, some of the factors we consider include:

- whether the content has been substantially edited in a manner that fundamentally alters its composition, sequence, timing, or framing;
- any visual or auditory information (such as new video frames, overdubbed audio, or modified subtitles) that has been added or removed; and
- whether media depicting a real person have been fabricated or simulated

We are most likely to take action (either labeling or removal, as described below) on more significant forms of alteration, such as wholly synthetic audio or video or content that has been doctored (spliced and reordered,

slowed down) to change its meaning. Subtler forms of manipulated media, such as isolative editing, omission of context, or presentation with false context, may be labeled or removed on a case-by-case basis.

We will not take action to label or remove media that have been edited in ways that do not fundamentally alter their meaning, such as retouched photos or color-corrected videos.

In order to determine if media have been significantly and deceptively altered or fabricated, we may use our own technology or receive reports through partnerships with third parties. In situations where we are unable to reliably determine if media have been altered or fabricated, we may not take action to label or remove them.

2. Is the content shared in a deceptive manner?

We also consider whether the context in which media are shared could result in confusion or misunderstanding or suggests a deliberate intent to deceive people about the nature or origin of the content, for example by falsely claiming that it depicts reality. We assess the context provided alongside media to see whether it makes clear that the media have been altered or fabricated. Some of the types of context we assess in order to make this determination include:

- The text of the Tweet accompanying or within media
- Metadata associated with media
- Information on the profile of the account sharing media
- Websites linked in the Tweet, or in the profile of the account sharing media

3. Is the content likely to impact public safety or cause serious harm?

Tweets that share synthetic and manipulated media are subject to removal under this policy if they are likely to cause serious harm. Some specific harms we consider include:

- Threats to the physical safety of a person or group
- Risk of mass violence or widespread civil unrest
- Threats to the privacy or ability of a person or group to freely express themselves or participate in civic events, such as:
 - Stalking or unwanted and obsessive attention
 - Targeted content that includes tropes, epithets, or material that aims to silence someone
 - Voter suppression or intimidation

While we have other rules also intended to address these forms of harm, including our policies on violent threats, election integrity, and hateful conduct, we will err toward removal in borderline cases that might otherwise not violate existing rules for Tweets that include synthetic or manipulated media.

We also consider the time frame within which the content may be likely to impact public safety or cause serious harm, and are more likely to remove content under this policy if we find that immediate harms are likely to result from the content's presence on Twitter.

Note: We may also take action on synthetic and manipulated content under our [non-consensual nudity policy](#) (such as pornographic media altered to insert the faces of people not actually involved) or other parts of the Twitter Rules.

Labeling and removal

In most cases, if we have reason to believe that media shared in a Tweet have been significantly and deceptively altered or fabricated, we will provide additional context on Tweets sharing the media where they appear on Twitter. This means we may:

- Apply a label to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended; and/or
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or landing page.

In most cases, we will take all of the above actions on Tweets we label.

Media that meet **all three of the criteria** defined above—i.e. that are synthetic or manipulated, shared in a deceptive manner, and is likely to cause harm—may not be shared on Twitter and are subject to removal. Accounts engaging in repeated or severe violations of this policy may be permanently suspended.

* Other parts of the Twitter Rules apply and may lead to the removal of the content, particularly where there is high likelihood of severe harm, such as a threat to someone's life or physical safety.

Trademark policy

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Copyright policy

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity’s name, please review Twitter’s [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to [twitter.com](#), you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

Tweet withheld

18h

This Tweet from **@Username** has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the

material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Our approach to policy development and enforcement philosophy

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;

- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people.

This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

[Public impact of the content:](#) A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks

to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

[Violence](#): You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

[Terrorism/violent extremism](#): You may not threaten or promote terrorism or violent extremism. [Learn more](#).

[Child sexual exploitation](#): We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

[Abuse/harassment](#): You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

[Hateful conduct](#): You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

[Suicide or self-harm](#): You may not promote or encourage suicide or self-harm. [Learn more](#).

[Sensitive media, including graphic violence and adult content](#): You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

[Illegal or certain regulated goods or services](#): You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

[Private information](#): You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

[Non-consensual nudity](#): You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

[Platform manipulation and spam](#): You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

[Election integrity](#): You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

[Impersonation](#): You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Synthetic and manipulated media: You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Terrorism and violent extremism policy

Overview

March 2019

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for terrorist organizations or violent extremist groups and individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments in this context are informed by national and international terrorism designations. We also assess organizations under our violent extremist group criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;

- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

We examine a group's activities both on and off Twitter to determine whether they engage in and/or promote violence against civilians to advance a political, religious and/or social cause.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a terrorist organization or violent extremist group;
- recruiting for a terrorist organization or violent extremist group;
- providing or distributing services (e.g., financial, media/propaganda) to further a terrorist organization's or violent extremist group's stated goals; and
- using the insignia or symbols of terrorist organizations or violent extremist groups to promote them.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to military or government entities.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wish for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy

Overview

March 2019

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives; and

- identifying alleged victims of childhood sexual exploitation by name or image.

What is not a violation of this policy?

Discussions related to child sexual exploitation as a phenomenon or attraction towards minors are permitted, provided they don't promote or glorify child sexual exploitation **in any way**. Artistic depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios e.g., works by internationally renowned artists that feature minors.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

What you can do

If you think you've found a Twitter account distributing or promoting child sexual exploitation, please report it through our designated [form](#). Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. For help on how to find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with NCMEC [here](#) or with your local

law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the International Association of Internet Hotlines [website](#).

Our partners

We partner with multiple organisations around the world whose work is dedicated to fighting child sexual exploitation. You can find a full list of our safety partners [here](#).

Abusive behavior

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing or hoping serious harm on a person or group of people

We do not tolerate content that wishes, hopes or expresses a desire for death, serious bodily harm or fatal disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone's body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name,

or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you".

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., “I hope you get cancer and die.”

Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”

Saying that a group of individuals deserve serious physical injury, e.g., “If this group of protesters don’t shut up, they deserve to be shot.”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion, age, disability, or serious disease.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).
If someone believes their account was suspended in error, they can
[submit an appeal](#).

Suicide and Self-harm Policy

Overview

You may not promote or encourage suicide or self-harm.

At Twitter, we recognize that suicide and self-harm are significant social & public health challenges that require collaboration between all stakeholders – public, private, and civil society – and that we have a role and responsibility to help people access and receive support when they need it.

When developing this policy, we consulted extensively with experts to ensure that people who have engaged in self-harm or experienced suicidal thoughts can share their personal experiences. We also recognized the need to protect people from the potential harm caused by exposure to content that could promote or encourage self-harm – intentionally or inadvertently. That's why our policy prohibits content that promotes or encourages self-harming behaviors and provides support to those undergoing experiences with self-harm or suicidal thoughts.

What is in violation of this policy?

Under this policy, you can't promote, or otherwise encourage, suicide or self-harm. We define promotion and encouragement to include statements such as "the most effective", "the easiest", "the best", "the most successful", "you should", "why don't you". Violations of this policy can occur via Tweets, images or videos, including live video.

We define suicide to be the act of taking one's own life. We define self-harm to include:

- self-inflicted physical injuries e.g., cutting; and
- eating disorders e.g., bulimia, anorexia.

Violations of this policy include, but are not limited to:

- encouraging someone to physically harm or kill themselves;
- asking others for encouragement to engage in self-harm or suicide, including seeking partners for group suicides or suicide games; and
- sharing information, strategies, methods or instructions that would assist people to engage in self-harm and suicide.

What is not a violation of this policy?

Some examples of behavior that are not considered a violation of this policy include:

- sharing personal stories and experiences related to self-harm or suicide;
- sharing coping mechanisms and resources for addressing self-harm or suicidal thoughts; and
- discussions that are focused on research, advocacy, and education related to self-harm or suicide prevention.

Note: people can share their personal experiences, but should avoid sharing detailed information about specific strategies or methods related to self-harm, as this could inadvertently encourage this behavior.

Who can report violations of this policy?

Anyone can report content that may encourage or promote suicide or self-harm via our in-app reporting or our [specialized reporting form](#). These reports are routed to a dedicated team who evaluate each case individually.

Note: if we receive a report that someone has expressed an intention to engage in self-harm or suicide, we will contact them directly, encourage them to seek support, and provide information about dedicated online and hotline resources. We may also work with law enforcement officials where appropriate, for example, if we receive a valid emergency disclosure request as defined in our [Law Enforcement Guidelines](#).

How to report violations of this policy

To ensure that we handle reports sensitively, our in-app reporting provides separate options for people who may be expressing an intention to harm themselves, and content that is encouraging or promoting self-harm or suicide.

Expressing intentions of self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Reporting form

You can also report this content for review via [our dedicated report form](#).

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Encouraging self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

What happens if you violate this policy?

Our enforcement approach depends on the type of content being shared, whether or not the reported account is encouraging or promoting self-harm or suicide, and the account's previous history of violations.

If you violate this policy by sharing content that intentionally encourages others to harm themselves, ask others to encourage you to harm yourself, or share detailed information or instructions related to self-harm or suicide methods, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy, or if your account is dedicated to promoting or encouraging self-harm or suicide, your account will be permanently suspended. If cases include images or videos related to self-harm or suicide, we will also evaluate this content under [our sensitive media policy](#). If you believe that your account was suspended in error, [you can submit an appeal](#).

We may also take steps to prevent the spread of instructional material hosted on third-party websites by [marking such links as unsafe](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Learn more](#) about how you can support someone experiencing thoughts of self-harm or suicide.

Visit our Safety Center for a list of local [mental health resources](#) and read [our blog on suicide prevention](#) to learn more about our work.

Sensitive media policy

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.

3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their

information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;

- location information related to commercial property or places of business, where this information is publicly available;
- descriptions of physical appearance;
- gossip, rumours, accusations, and allegations; and
- screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy

Overview

March 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:

- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy

Overview

March 2019

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are; and
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting.

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate; and
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in “follow trains,” “decks,” and “Retweet for Retweet” behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- “follow churn” – following and then unfollowing large numbers of accounts in an effort to inflate one's own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account's followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person's browser (malware) or computer or to compromise a person's privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);

- posting links without commentary occasionally; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don't violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Blacklisting URLs

We blacklist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [hacked materials policy](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Civic integrity policy

Overview

May 2020

You may not use Twitter's services for the purpose of manipulating or interfering in elections or other civic processes. This includes posting or sharing content that may suppress participation or mislead people about when, where, or how to participate in a civic process.

The public conversation occurring on Twitter is never more important than during elections and other civic events. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt civic processes, including through the distribution of false or misleading information about the procedures or circumstances around participation in a civic process.

What is a civic process?

Twitter considers civic processes to be events or procedures mandated, organized, and conducted by the governing and/or electoral body of a country, state, region, district, or municipality to address a matter of common concern through public participation. Some examples of civic processes may include:

- Political elections
- Censuses
- Major referenda and ballot initiatives

What is in violation of this policy?

We prohibit 3 categories of manipulative behavior and content under this policy:

Misleading information about how to participate

You can't share false or misleading information about how to participate in an election or other civic process. This includes but is not limited to:

- misleading information about procedures to participate in a civic process (for example, that you can vote by Tweet, text message, email, or phone call in jurisdictions where these are not a possibility);
- misleading information about requirements for participation, including identification or citizenship requirements; and
- misleading statements or information about the official, announced date or time of a civic process.

Suppression and intimidation

You can't share false or misleading information intended to intimidate or dissuade people from participating in an election or other civic process. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to voting in an election, polling places, or collecting census information;
- misleading claims about long lines, equipment problems, or other disruptions at voting locations during election periods;
- misleading claims about process procedures or techniques which could dissuade people from participating; and
- threats regarding voting locations or other key places or events (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account policy](#).

What is not a violation of this policy?

Not all false or untrue information about politics or civic processes constitutes manipulation or interference. In the absence of other policy violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected or appointed official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information;
- voting and audience participation for competitions, game shows, or other entertainment purposes;
- broad, non-specific statements about the integrity of elections or civic processes (such as unsubstantiated claims that an election is "rigged"); and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election or civic process. Therefore, we enable reporting of false or misleading information about civic processes in advance of major events, for people located in the relevant countries and locations. We also work with select government and civil society partners in these countries to provide additional channels for reporting and expedited review.

For civic processes with multiple stages or parts, such as primary elections or lengthy campaigns, reporting will be enabled leading up to the first officially-sanctioned event associated with the civic process.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country at the relevant time, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in civic processes.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in a civic process.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our civic integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Synthetic and manipulated media policy

Overview

You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context.

You should be able to find reliable information on Twitter. That means understanding whether the content you see is real or fabricated and having the ability to find more context about what you see on Twitter. Therefore, we may label Tweets that include media (videos, audio, and images) that have been deceptively altered or fabricated. In addition, you may not share deceptively altered media on Twitter in ways that mislead or deceive people about the media's authenticity where threats to physical safety or other serious harm may result.

We use the following criteria as we consider Tweets and media for labeling or removal under this policy as part of our ongoing work to enforce our rules and ensure healthy and safe conversation on Twitter (additional information is available below):

Is the content significantly and deceptively altered or fabricated?	Is the content shared in a deceptive manner?	Is the content likely to impact public safety or cause serious harm?	
✓	✗	✗	Content may be labeled.
✗	✓	✗	Content may be labeled.
✓	✗	✓	Content is likely to be labeled, or may be removed.*
✓	✓	✗	Content is likely to be labeled.
✓	✓	✓	Content is likely to be removed.

1. Is the content synthetic or manipulated?

In order for content to be labeled or removed under this policy, we must have reason to believe that media, or the context in which media are presented, are significantly and deceptively altered or manipulated. Synthetic and manipulated media take many different forms and people can employ a wide range of technologies to produce these media. In assessing whether media have been significantly and deceptively altered or fabricated, some of the factors we consider include:

- whether the content has been substantially edited in a manner that fundamentally alters its composition, sequence, timing, or framing;
- any visual or auditory information (such as new video frames, overdubbed audio, or modified subtitles) that has been added or removed; and
- whether media depicting a real person have been fabricated or simulated

We are most likely to take action (either labeling or removal, as described below) on more significant forms of alteration, such as wholly synthetic audio or video or content that has been doctored (spliced and reordered, slowed down) to change its meaning. Subtler forms of manipulated media, such as isolative editing, omission of context, or presentation with false context, may be labeled or removed on a case-by-case basis.

We will not take action to label or remove media that have been edited in ways that do not fundamentally alter their meaning, such as retouched photos or color-corrected videos.

In order to determine if media have been significantly and deceptively altered or fabricated, we may use our own technology or receive reports through partnerships with third parties. In situations where we are unable to reliably determine if media have been altered or fabricated, we may not take action to label or remove them.

2. Is the content shared in a deceptive manner?

We also consider whether the context in which media are shared could result in confusion or misunderstanding or suggests a deliberate intent to deceive people about the nature or origin of the content, for example by falsely claiming that it depicts reality. We assess the context provided alongside media to see whether it makes clear that the media have been altered or fabricated. Some of the types of context we assess in order to make this determination include:

- The text of the Tweet accompanying or within media
- Metadata associated with media
- Information on the profile of the account sharing media
- Websites linked in the Tweet, or in the profile of the account sharing media

3. Is the content likely to impact public safety or cause serious harm?

Tweets that share synthetic and manipulated media are subject to removal under this policy if they are likely to cause serious harm. Some specific harms we consider include:

- Threats to the physical safety of a person or group
- Risk of mass violence or widespread civil unrest
- Threats to the privacy or ability of a person or group to freely express themselves or participate in civic events, such as:
 - Stalking or unwanted and obsessive attention
 - Targeted content that includes tropes, epithets, or material that aims to silence someone
 - Voter suppression or intimidation

While we have other rules also intended to address these forms of harm, including our policies on violent threats, election integrity, and hateful conduct, we will err toward removal in borderline cases that might otherwise not violate existing rules for Tweets that include synthetic or manipulated media.

We also consider the time frame within which the content may be likely to impact public safety or cause serious harm, and are more likely to remove content under this policy if we find that immediate harms are likely to result from the content's presence on Twitter.

Note: We may also take action on synthetic and manipulated content under our [non-consensual nudity policy](#) (such as pornographic media altered to insert the faces of people not actually involved) or other parts of the Twitter Rules.

Labeling and removal

In most cases, if we have reason to believe that media shared in a Tweet have been significantly and deceptively altered or fabricated, we will provide additional context on Tweets sharing the media where they appear on Twitter. This means we may:

- Apply a label to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended; and/or
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or landing page.

In most cases, we will take all of the above actions on Tweets we label.

Media that meet **all three of the criteria** defined above—i.e. that are synthetic or manipulated, shared in a deceptive manner, and is likely to cause harm—may not be shared on Twitter and are subject to removal. Accounts engaging in repeated or severe violations of this policy may be permanently suspended.

** Other parts of the Twitter Rules apply and may lead to the removal of the content, particularly where there is high likelihood of severe harm, such as a threat to someone's life or physical safety.*

Trademark policy

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Copyright policy](#)

What Types of Copyright Complaints Does Twitter Respond To?

Twitter responds to copyright notifications submitted under the Digital Millennium Copyright Act (“DMCA”). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly

infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review Twitter's [trademark policy](#). If you are concerned about the use of a fictional character, please see our [parody, commentary, and fan accounts policy](#). These are generally not copyright issues.

Am I a Copyright Holder? How Do I Know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

Tip: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph.

What Information Do You Need To Process a Copyright Complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the notification is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How Do I File a Copyright Complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and filing a [copyright report](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a **full copy of your notice** (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing, as was the case in [Lenz v. Universal](#) and [OPG v. Diebold](#). If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

How Are Claims Processed?

We process reports in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. Please note, submitting duplicate copyright notices may result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) after removing or disabling access to the material, provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice, and forward a redacted copy of the complaint to [Lumen](#).

What Happens Next?

Twitter's response to notices of alleged copyright infringement may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a notice of alleged infringement, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the takedown notice, along with instructions for filing a counter-notification.

If you've not yet received a copy of the copyright notification regarding the content removed from your account, please respond to the support ticket we sent you.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Under appropriate circumstances, Twitter may suspend and warn repeat violators, and in more serious cases, permanently terminate user accounts.

What Happens If I Receive a Copyright Notification?

If you receive a copyright notification, it means that the content described in the notification has been removed from Twitter or access to the content on Twitter has been restricted. Please take the time to read through our notice to you, which includes information on the notification we received as well as instructions on how to file a counter-notice.

When Should I File a Counter-notice?

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

If you believe that the material reported in the copyright notification you received was misidentified or removed in error, you should file a counter-notice as per the instructions below.

Re-posting material removed in response to a copyright notification may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notification rather than simply re-posting the material.

What Information Do You Need to Process a Counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What Happens After I Submit a Counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a Copyright Notice or Counter-notice is Serious Business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

Our approach to policy development and enforcement philosophy

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

[Violence](#): You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

[Terrorism/violent extremism](#): You may not threaten or promote terrorism or violent extremism. [Learn more](#).

[Child sexual exploitation](#): We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

[Abuse/harassment](#): You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

[Hateful conduct](#): You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

[Suicide or self-harm](#): You may not promote or encourage suicide or self-harm. [Learn more](#).

[Sensitive media, including graphic violence and adult content](#): You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

[Illegal or certain regulated goods or services](#): You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

[Private information](#): You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

[Non-consensual nudity](#): You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

[Platform manipulation and spam](#): You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

[Election integrity](#): You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

[Impersonation](#): You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Synthetic and manipulated media: You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Terrorism and violent extremism policy

Overview

March 2019

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for terrorist organizations or violent extremist groups and individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments in this context are informed by national and international terrorism designations. We also assess organizations under our violent extremist group criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;

- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

We examine a group's activities both on and off Twitter to determine whether they engage in and/or promote violence against civilians to advance a political, religious and/or social cause.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a terrorist organization or violent extremist group;
- recruiting for a terrorist organization or violent extremist group;
- providing or distributing services (e.g., financial, media/propaganda) to further a terrorist organization's or violent extremist group's stated goals; and
- using the insignia or symbols of terrorist organizations or violent extremist groups to promote them.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to military or government entities.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wish for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy

Overview

March 2019

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives; and

- identifying alleged victims of childhood sexual exploitation by name or image.

What is not a violation of this policy?

Discussions related to child sexual exploitation as a phenomenon or attraction towards minors are permitted, provided they don't promote or glorify child sexual exploitation **in any way**. Artistic depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios e.g., works by internationally renowned artists that feature minors.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

What you can do

If you think you've found a Twitter account distributing or promoting child sexual exploitation, please report it through our designated [form](#). Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. For help on how to find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with NCMEC [here](#) or with your local

law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the International Association of Internet Hotlines [website](#).

Our partners

We partner with multiple organisations around the world whose work is dedicated to fighting child sexual exploitation. You can find a full list of our safety partners [here](#).

Abusive behavior

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing or hoping serious harm on a person or group of people

We do not tolerate content that wishes, hopes or expresses a desire for death, serious bodily harm or fatal disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone's body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name,

or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you".

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., “I hope you get cancer and die.”

Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”

Saying that a group of individuals deserve serious physical injury, e.g., “If this group of protesters don’t shut up, they deserve to be shot.”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion, age, disability, or serious disease.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).
If someone believes their account was suspended in error, they can
[submit an appeal](#).

Suicide and Self-harm Policy

Overview

You may not promote or encourage suicide or self-harm.

At Twitter, we recognize that suicide and self-harm are significant social & public health challenges that require collaboration between all stakeholders – public, private, and civil society – and that we have a role and responsibility to help people access and receive support when they need it.

When developing this policy, we consulted extensively with experts to ensure that people who have engaged in self-harm or experienced suicidal thoughts can share their personal experiences. We also recognized the need to protect people from the potential harm caused by exposure to content that could promote or encourage self-harm – intentionally or inadvertently. That's why our policy prohibits content that promotes or encourages self-harming behaviors and provides support to those undergoing experiences with self-harm or suicidal thoughts.

What is in violation of this policy?

Under this policy, you can't promote, or otherwise encourage, suicide or self-harm. We define promotion and encouragement to include statements such as "the most effective", "the easiest", "the best", "the most successful", "you should", "why don't you". Violations of this policy can occur via Tweets, images or videos, including live video.

We define suicide to be the act of taking one's own life. We define self-harm to include:

- self-inflicted physical injuries e.g., cutting; and
- eating disorders e.g., bulimia, anorexia.

Violations of this policy include, but are not limited to:

- encouraging someone to physically harm or kill themselves;
- asking others for encouragement to engage in self-harm or suicide, including seeking partners for group suicides or suicide games; and
- sharing information, strategies, methods or instructions that would assist people to engage in self-harm and suicide.

What is not a violation of this policy?

Some examples of behavior that are not considered a violation of this policy include:

- sharing personal stories and experiences related to self-harm or suicide;
- sharing coping mechanisms and resources for addressing self-harm or suicidal thoughts; and
- discussions that are focused on research, advocacy, and education related to self-harm or suicide prevention.

Note: people can share their personal experiences, but should avoid sharing detailed information about specific strategies or methods related to self-harm, as this could inadvertently encourage this behavior.

Who can report violations of this policy?

Anyone can report content that may encourage or promote suicide or self-harm via our in-app reporting or our [specialized reporting form](#). These reports are routed to a dedicated team who evaluate each case individually.

Note: if we receive a report that someone has expressed an intention to engage in self-harm or suicide, we will contact them directly, encourage them to seek support, and provide information about dedicated online and hotline resources. We may also work with law enforcement officials where appropriate, for example, if we receive a valid emergency disclosure request as defined in our [Law Enforcement Guidelines](#).

How to report violations of this policy

To ensure that we handle reports sensitively, our in-app reporting provides separate options for people who may be expressing an intention to harm themselves, and content that is encouraging or promoting self-harm or suicide.

Expressing intentions of self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Reporting form

You can also report this content for review via [our dedicated report form](#).

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Encouraging self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

What happens if you violate this policy?

Our enforcement approach depends on the type of content being shared, whether or not the reported account is encouraging or promoting self-harm or suicide, and the account's previous history of violations.

If you violate this policy by sharing content that intentionally encourages others to harm themselves, ask others to encourage you to harm yourself, or share detailed information or instructions related to self-harm or suicide methods, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy, or if your account is dedicated to promoting or encouraging self-harm or suicide, your account will be permanently suspended. If cases include images or videos related to self-harm or suicide, we will also evaluate this content under [our sensitive media policy](#). If you believe that your account was suspended in error, [you can submit an appeal](#).

We may also take steps to prevent the spread of instructional material hosted on third-party websites by [marking such links as unsafe](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Learn more](#) about how you can support someone experiencing thoughts of self-harm or suicide.

Visit our Safety Center for a list of local [mental health resources](#) and read [our blog on suicide prevention](#) to learn more about our work.

Sensitive media policy

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.

3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their

information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;

- location information related to commercial property or places of business, where this information is publicly available;
- descriptions of physical appearance;
- gossip, rumours, accusations, and allegations; and
- screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy

Overview

March 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:

- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy

Overview

March 2019

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are; and
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting.

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate; and
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in “follow trains,” “decks,” and “Retweet for Retweet” behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- “follow churn” – following and then unfollowing large numbers of accounts in an effort to inflate one's own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account's followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person's browser (malware) or computer or to compromise a person's privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);

- posting links without commentary occasionally; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don't violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Blacklisting URLs

We blacklist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [hacked materials policy](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Civic integrity policy

Overview

May 2020

You may not use Twitter's services for the purpose of manipulating or interfering in elections or other civic processes. This includes posting or sharing content that may suppress participation or mislead people about when, where, or how to participate in a civic process.

The public conversation occurring on Twitter is never more important than during elections and other civic events. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt civic processes, including through the distribution of false or misleading information about the procedures or circumstances around participation in a civic process.

What is a civic process?

Twitter considers civic processes to be events or procedures mandated, organized, and conducted by the governing and/or electoral body of a country, state, region, district, or municipality to address a matter of common concern through public participation. Some examples of civic processes may include:

- Political elections
- Censuses
- Major referenda and ballot initiatives

What is in violation of this policy?

We prohibit 3 categories of manipulative behavior and content under this policy:

Misleading information about how to participate

You can't share false or misleading information about how to participate in an election or other civic process. This includes but is not limited to:

- misleading information about procedures to participate in a civic process (for example, that you can vote by Tweet, text message, email, or phone call in jurisdictions where these are not a possibility);
- misleading information about requirements for participation, including identification or citizenship requirements; and
- misleading statements or information about the official, announced date or time of a civic process.

Suppression and intimidation

You can't share false or misleading information intended to intimidate or dissuade people from participating in an election or other civic process. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to voting in an election, polling places, or collecting census information;
- misleading claims about long lines, equipment problems, or other disruptions at voting locations during election periods;
- misleading claims about process procedures or techniques which could dissuade people from participating; and
- threats regarding voting locations or other key places or events (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account policy](#).

What is not a violation of this policy?

Not all false or untrue information about politics or civic processes constitutes manipulation or interference. In the absence of other policy violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected or appointed official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information;
- voting and audience participation for competitions, game shows, or other entertainment purposes;
- broad, non-specific statements about the integrity of elections or civic processes (such as unsubstantiated claims that an election is "rigged"); and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election or civic process. Therefore, we enable reporting of false or misleading information about civic processes in advance of major events, for people located in the relevant countries and locations. We also work with select government and civil society partners in these countries to provide additional channels for reporting and expedited review.

For civic processes with multiple stages or parts, such as primary elections or lengthy campaigns, reporting will be enabled leading up to the first officially-sanctioned event associated with the civic process.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country at the relevant time, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in civic processes.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in a civic process.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our civic integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Synthetic and manipulated media policy

Overview

You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context.

You should be able to find reliable information on Twitter. That means understanding whether the content you see is real or fabricated and having the ability to find more context about what you see on Twitter. Therefore, we may label Tweets that include media (videos, audio, and images) that have been deceptively altered or fabricated. In addition, you may not share deceptively altered media on Twitter in ways that mislead or deceive people about the media's authenticity where threats to physical safety or other serious harm may result.

We use the following criteria as we consider Tweets and media for labeling or removal under this policy as part of our ongoing work to enforce our rules and ensure healthy and safe conversation on Twitter (additional information is available below):

Is the content significantly and deceptively altered or fabricated?	Is the content shared in a deceptive manner?	Is the content likely to impact public safety or cause serious harm?	
✓	✗	✗	Content may be labeled.
✗	✓	✗	Content may be labeled.
✓	✗	✓	Content is likely to be labeled, or may be removed.*
✓	✓	✗	Content is likely to be labeled.
✓	✓	✓	Content is likely to be removed.

1. Is the content synthetic or manipulated?

In order for content to be labeled or removed under this policy, we must have reason to believe that media, or the context in which media are presented, are significantly and deceptively altered or manipulated. Synthetic and manipulated media take many different forms and people can employ a wide range of technologies to produce these media. In assessing whether media have been significantly and deceptively altered or fabricated, some of the factors we consider include:

- whether the content has been substantially edited in a manner that fundamentally alters its composition, sequence, timing, or framing;
- any visual or auditory information (such as new video frames, overdubbed audio, or modified subtitles) that has been added or removed; and
- whether media depicting a real person have been fabricated or simulated

We are most likely to take action (either labeling or removal, as described below) on more significant forms of alteration, such as wholly synthetic audio or video or content that has been doctored (spliced and reordered, slowed down) to change its meaning. Subtler forms of manipulated media, such as isolative editing, omission of context, or presentation with false context, may be labeled or removed on a case-by-case basis.

We will not take action to label or remove media that have been edited in ways that do not fundamentally alter their meaning, such as retouched photos or color-corrected videos.

In order to determine if media have been significantly and deceptively altered or fabricated, we may use our own technology or receive reports through partnerships with third parties. In situations where we are unable to reliably determine if media have been altered or fabricated, we may not take action to label or remove them.

2. Is the content shared in a deceptive manner?

We also consider whether the context in which media are shared could result in confusion or misunderstanding or suggests a deliberate intent to deceive people about the nature or origin of the content, for example by falsely claiming that it depicts reality. We assess the context provided alongside media to see whether it makes clear that the media have been altered or fabricated. Some of the types of context we assess in order to make this determination include:

- The text of the Tweet accompanying or within media
- Metadata associated with media
- Information on the profile of the account sharing media
- Websites linked in the Tweet, or in the profile of the account sharing media

3. Is the content likely to impact public safety or cause serious harm?

Tweets that share synthetic and manipulated media are subject to removal under this policy if they are likely to cause serious harm. Some specific harms we consider include:

- Threats to the physical safety of a person or group
- Risk of mass violence or widespread civil unrest
- Threats to the privacy or ability of a person or group to freely express themselves or participate in civic events, such as:
 - Stalking or unwanted and obsessive attention
 - Targeted content that includes tropes, epithets, or material that aims to silence someone
 - Voter suppression or intimidation

While we have other rules also intended to address these forms of harm, including our policies on violent threats, election integrity, and hateful conduct, we will err toward removal in borderline cases that might otherwise not violate existing rules for Tweets that include synthetic or manipulated media.

We also consider the time frame within which the content may be likely to impact public safety or cause serious harm, and are more likely to remove content under this policy if we find that immediate harms are likely to result from the content's presence on Twitter.

Note: We may also take action on synthetic and manipulated content under our [non-consensual nudity policy](#) (such as pornographic media altered to insert the faces of people not actually involved) or other parts of the Twitter Rules.

Labeling and removal

In most cases, if we have reason to believe that media shared in a Tweet have been significantly and deceptively altered or fabricated, we will provide additional context on Tweets sharing the media where they appear on Twitter. This means we may:

- Apply a label to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended; and/or
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or landing page.

In most cases, we will take all of the above actions on Tweets we label.

Media that meet **all three of the criteria** defined above—i.e. that are synthetic or manipulated, shared in a deceptive manner, and is likely to cause harm—may not be shared on Twitter and are subject to removal. Accounts engaging in repeated or severe violations of this policy may be permanently suspended.

** Other parts of the Twitter Rules apply and may lead to the removal of the content, particularly where there is high likelihood of severe harm, such as a threat to someone's life or physical safety.*

Trademark policy

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Copyright policy

What types of copyright complaints does Twitter respond to?

Twitter responds to copyright complaints submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly

infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review [Twitter's trademark policy](#). If you are concerned about a parody, newsfeed, commentary, or fan account, please see the relevant policy [here](#). These are generally not copyright issues.

Am I a copyright holder? How do I know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

What to consider before submitting a copyright complaint

Before submitting a copyright complaint to us, please consider whether or not the use could be considered [fair use](#).

If you have considered fair use, and you still wish to continue with a copyright complaint, you may want to first reach out to the user in question to see if you can resolve the matter directly with the user. You can reply to the user's Tweet or send the user a Direct Message and ask for them to remove your copyrighted content without having to contact Twitter.

Prior to submitting a formal complaint with Twitter, please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

Note: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph. If you're unsure whether or not you own the copyrights to a work, or if you're infringing upon someone else's work, please consult an attorney or another advisor.

What information do you need to process a copyright complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the complaint is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How do I file a copyright complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and [filing a copyright complaint](#). If you are logged in to [twitter.com](#), you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a full copy of your notice (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a copyright complaint.

How are claims processed?

We process copyright complaints in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. If you do not receive a ticket confirmation that means we did not receive your complaint and you should re-submit your complaint. However, please note, submitting duplicate copyright complaints will result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) and provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice. We will also forward a redacted copy of the complaint to [Lumen](#), with your personal information removed.

What information gets forwarded to the reported user(s)?

If we remove or disable access to the materials reported in a copyright complaint, the reported user(s) will receive a copy of the complaint, including the reporter's full name, email, street address, and any other information included in the complaint.

If you are uncomfortable sharing your contact information with the reported user(s), you may wish to consider appointing an agent to submit your DMCA notice on your behalf. Your agent will be required to submit the DMCA notice with valid contact information, and identify you as the content owner that they are representing.

What happens next?

Twitter's response to copyright complaints may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a copyright complaint, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the complaint, along with instructions for filing a counter-notice.

If you've not yet received a copy of the copyright complaint regarding the content removed from your account, please respond to the support ticket we sent you. If your complaint involves removal of content from Periscope, please see this help center [article](#) regarding how to request a copy of your complaint.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)

**Media not displayed**

This image has been removed in response to a report from the copyright holder.

My content was removed from Twitter*Why did I receive a copyright complaint?*

If you receive a copyright complaint, it means that access to the content described in the complaint has been restricted. Please take the time to read through our correspondence to you, which includes information on the complaint we received as well as instructions on how to file a counter-notice. Please ensure that you are monitoring the email address associated with your Twitter account.

Tip: Removing the material reported in a copyright complaint will not resolve that complaint.

What if I want to contest the takedown?

If you believe that the materials reported in the copyright complaints were misidentified or removed in error, you may send us a counter-notification(s). A counter-notice is a request for Twitter to reinstate the removed material, and it has legal consequences. Alternatively, you may be able to seek a retraction of the copyright complaint from the reporter.

How do I seek a retraction?

The DMCA complaint you received includes the contact information of the reporter. You may want to reach out and ask them to retract their notice. The reporter can send retractions to copyright@twitter.com, and should include: (1) identification of the material that was disabled, and (2) a statement that the reporter would like to retract their DMCA notice. This is the fastest and most efficient means of resolving an unresolved copyright complaint. A retraction is at the sole discretion of the original reporter.

When should I file a counter-notice?

A counter-notice is a request for Twitter to reinstate the removed material, and is the start of a legal process that has legal consequences. For example, submitting a counter notice indicates that you consent to the jurisdiction of a U.S. Federal court and that you consent to the disclosure of your personal information to the reporter and Lumen website.

With these considerations in mind, you may file a counter-notice if you believe that this material was misidentified, or you have a good faith belief that the material should not have been removed. If you're unsure whether or not you should file a counter-notice, you may want to consult with an attorney.

Tip: Re-posting material removed in response to a copyright complaint may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notice rather than re-posting the material.

What information do you need to process a counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What happens after I submit a counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. This means that the contact information that is submitted in your counter-notice will be shared to the person who filed the original notice.

If the copyright owner disagrees that the content was removed in error or misidentification, they may pursue legal action against you. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a copyright complaint or counter-notice is serious business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

What happens if my account receives multiple copyright complaints?

If multiple copyright complaints are received Twitter may lock accounts or take other actions to warn repeat violators. These warnings may vary across Twitter's services. Under appropriate circumstances we may suspend user accounts under our repeat infringer policy. However, we may take retractions and counter-notices into account when applying our repeat infringer policy.

Our approach to policy development and enforcement philosophy

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (20200204*)

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Synthetic and manipulated media: You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.

5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Terrorism and violent extremism policy (20190607*)

Overview

March 2019

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for terrorist organizations or violent extremist groups and individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments in this context are informed by national and international terrorism designations. We also assess organizations under our violent extremist group criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

We examine a group's activities both on and off Twitter to determine whether they engage in and/or promote violence against civilians to advance a political, religious and/or social cause.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a terrorist organization or violent extremist group;
- recruiting for a terrorist organization or violent extremist group;
- providing or distributing services (e.g., financial, media/propaganda) to further a terrorist organization's or violent extremist group's stated goals; and
- using the insignia or symbols of terrorist organizations or violent extremist groups to promote them.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to military or government entities.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wish for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy

(20190607*)

Overview

March 2019

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;

- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives; and
- identifying alleged victims of childhood sexual exploitation by name or image.

What is not a violation of this policy?

Discussions related to child sexual exploitation as a phenomenon or attraction towards minors are permitted, provided they don't promote or glorify child sexual exploitation **in any way**. Artistic depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios e.g., works by internationally renowned artists that feature minors.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

What you can do

If you think you've found a Twitter account distributing or promoting child sexual exploitation, please report it through our designated [form](#). Please provide the username

and links to the relevant tweets that lead you to believe the account should be investigated. For help on how to find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with NCMEC [here](#) or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the International Association of Internet Hotlines [website](#).

Our partners

We partner with multiple organisations around the world whose work is dedicated to fighting child sexual exploitation. You can find a full list of our safety partners [here](#).

Abusive behavior (20190607*)

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing or hoping serious harm on a person or group of people

We do not tolerate content that wishes, hopes or expresses a desire for death, serious bodily harm or fatal disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone's body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy (20200331*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you".

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., “I hope you get cancer and die.”

Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”

Saying that a group of individuals deserve serious physical injury, e.g., “If this group of protesters don’t shut up, they deserve to be shot.”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion, age, disability, or serious disease.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can

[submit an appeal.](#)

Suicide and Self-harm Policy

(20200805*)

Overview

You may not promote or encourage suicide or self-harm.

At Twitter, we recognize that suicide and self-harm are significant social & public health challenges that require collaboration between all stakeholders – public, private, and civil society – and that we have a role and responsibility to help people access and receive support when they need it.

When developing this policy, we consulted extensively with experts to ensure that people who have engaged in self-harm or experienced suicidal thoughts can share their personal experiences. We also recognized the need to protect people from the potential harm caused by exposure to content that could promote or encourage self-harm – intentionally or inadvertently. That's why our policy prohibits content that promotes or encourages self-harming behaviors and provides support to those undergoing experiences with self-harm or suicidal thoughts.

What is in violation of this policy?

Under this policy, you can't promote, or otherwise encourage, suicide or self-harm. We define promotion and encouragement to include statements such as "the most effective", "the easiest", "the best", "the most successful", "you should", "why don't you". Violations of this policy can occur via Tweets, images or videos, including live video.

We define suicide to be the act of taking one's own life. We define self-harm to include:

- self-inflicted physical injuries e.g., cutting; and
- eating disorders e.g., bulimia, anorexia.

Violations of this policy include, but are not limited to:

- encouraging someone to physically harm or kill themselves;
- asking others for encouragement to engage in self-harm or suicide, including seeking partners for group suicides or suicide games; and
- sharing information, strategies, methods or instructions that would assist people to engage in self-harm and suicide.

What is not a violation of this policy?

Some examples of behavior that are not considered a violation of this policy include:

- sharing personal stories and experiences related to self-harm or suicide;
- sharing coping mechanisms and resources for addressing self-harm or suicidal thoughts; and
- discussions that are focused on research, advocacy, and education related to self-harm or suicide prevention.

Note: people can share their personal experiences, but should avoid sharing detailed information about specific strategies or methods related to self-harm, as this could inadvertently encourage this behavior.

Who can report violations of this policy?

Anyone can report content that may encourage or promote suicide or self-harm via our in-app reporting or our [specialized reporting form](#). These reports are routed to a dedicated team who evaluate each case individually.

Note: if we receive a report that someone has expressed an intention to engage in self-harm or suicide, we will contact them directly, encourage them to seek support, and provide information about dedicated online and hotline resources. We may also work with law enforcement officials where appropriate, for example, if we receive a valid emergency disclosure request as defined in our [Law Enforcement Guidelines](#).

How to report violations of this policy

To ensure that we handle reports sensitively, our in-app reporting provides separate options for people who may be expressing an intention to harm themselves, and content that is encouraging or promoting self-harm or suicide.

Expressing intentions of self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Reporting form

You can also report this content for review via [our dedicated report form](#).

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Encouraging self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"

4. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

What happens if you violate this policy?

Our enforcement approach depends on the type of content being shared, whether or not the reported account is encouraging or promoting self-harm or suicide, and the account's previous history of violations.

If you violate this policy by sharing content that intentionally encourages others to harm themselves, ask others to encourage you to harm yourself, or share detailed information or instructions related to self-harm or suicide methods, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy, or if your account is dedicated to promoting or encouraging self-harm or suicide, your account will be permanently suspended. If cases include images or videos related to self-harm or suicide, we will also evaluate this content under [our sensitive media policy](#). If you believe that your account was suspended in error, [you can submit an appeal](#).

We may also take steps to prevent the spread of instructional material hosted on third-party websites by [marking such links as unsafe](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Learn more](#) about how you can support someone experiencing thoughts of self-harm or suicide.

Visit our Safety Center for a list of local [mental health resources](#) and read [our blog on suicide prevention](#) to learn more about our work.

Sensitive media policy (20191204*)

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and

- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20190607*)

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;
 - descriptions of physical appearance;
 - gossip, rumours, accusations, and allegations; and
 - screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy **(20191217*)**

Overview

November 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy (20200918*)

Overview

September 2020

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are;
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting; and
- coordinated harmful activity that encourages or promotes behavior which violates the [Twitter Rules](#).

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;

- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts or by coordinating with others to violate the Twitter Rules. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate;
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account; and
 - coordinating with others to engage in or promote violations of the Twitter Rules, including violations of our [abusive behavior](#) policy.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in "follow trains," "decks," and "Retweet for Retweet" behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- "follow churn" – following and then unfollowing large numbers of accounts in an effort to inflate one's own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account's followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person's browser (malware) or computer or to compromise a person's privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally;
- coordinating with others to express ideas, viewpoints, support, or opposition towards a cause, provided such behavior does not result in violations of the Twitter Rules; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don't violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Denylisting URLs

We denylist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [financial scam policy](#), our [hacked materials policy](#), our approach to [coordinated harmful activity](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Civic integrity policy (20200605*)

Overview

May 2020

You may not use Twitter's services for the purpose of manipulating or interfering in elections or other civic processes. This includes posting or sharing content that may suppress participation or mislead people about when, where, or how to participate in a civic process.

The public conversation occurring on Twitter is never more important than during elections and other civic events. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt civic processes, including through the distribution of false or misleading information about the procedures or circumstances around participation in a civic process.

What is a civic process?

Twitter considers civic processes to be events or procedures mandated, organized, and conducted by the governing and/or electoral body of a country, state, region, district, or municipality to address a matter of common concern through public participation. Some examples of civic processes may include:

- Political elections
- Censuses
- Major referenda and ballot initiatives

What is in violation of this policy?

We prohibit 3 categories of manipulative behavior and content under this policy:

Misleading information about how to participate

You can't share false or misleading information about how to participate in an election or other civic process. This includes but is not limited to:

- misleading information about procedures to participate in a civic process (for example, that you can vote by Tweet, text message, email, or phone call in jurisdictions where these are not a possibility);
- misleading information about requirements for participation, including identification or citizenship requirements; and
- misleading statements or information about the official, announced date or time of a civic process.

Suppression and intimidation

You can't share false or misleading information intended to intimidate or dissuade people from participating in an election or other civic process. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to voting in an election, polling places, or collecting census information;
- misleading claims about long lines, equipment problems, or other disruptions at voting locations during election periods;
- misleading claims about process procedures or techniques which could dissuade people from participating; and

- threats regarding voting locations or other key places or events (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account policy](#).

What is not a violation of this policy?

Not all false or untrue information about politics or civic processes constitutes manipulation or interference. In the absence of other policy violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected or appointed official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information;
- voting and audience participation for competitions, game shows, or other entertainment purposes;
- broad, non-specific statements about the integrity of elections or civic processes (such as unsubstantiated claims that an election is "rigged"); and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election or civic process. Therefore, we enable reporting of false or misleading information about civic processes in advance of major events, for people located in the relevant countries and locations. We also work with select government and civil society partners in these countries to provide additional channels for reporting and expedited review.

For civic processes with multiple stages or parts, such as primary elections or lengthy campaigns, reporting will be enabled leading up to the first officially-sanctioned event associated with the civic process.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country at the relevant time, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in civic processes.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.

3. Select the option that best tells us how the Tweet is misleading about voting or participation in a civic process.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our civic integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Impersonation policy (20191001*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody](#), [newsfeed](#), [commentary](#), and [fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Synthetic and manipulated media policy **(20200204*)**

Overview

You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context.

You should be able to find reliable information on Twitter. That means understanding whether the content you see is real or fabricated and having the ability to find more context about what you see on Twitter. Therefore, we may label Tweets that include media (videos, audio, and images) that have been deceptively altered or fabricated. In addition, you may not share deceptively altered media on Twitter in ways that mislead or deceive people about the media's authenticity where threats to physical safety or other serious harm may result.

We use the following criteria as we consider Tweets and media for labeling or removal under this policy as part of our ongoing work to enforce our rules and ensure healthy and safe conversation on Twitter (additional information is available below):

Is the content significantly and deceptively altered or fabricated?	Is the content shared in a deceptive manner?	Is the content likely to impact public safety or cause serious harm?	
✓	✗	✗	Content may be labeled.
✗	✓	✗	Content may be labeled.
✓	✗	✓	Content is likely to be labeled, or may be removed.*
✓	✓	✗	Content is likely to be labeled.
✓	✓	✓	Content is likely to be removed.

1. Is the content synthetic or manipulated?

In order for content to be labeled or removed under this policy, we must have reason to believe that media, or the context in which media are presented, are significantly and deceptively altered or manipulated. Synthetic and manipulated media take many different forms and people can employ a wide range of technologies to produce these media. In assessing whether media have been significantly and deceptively altered or fabricated, some of the factors we consider include:

- whether the content has been substantially edited in a manner that fundamentally alters its composition, sequence, timing, or framing;
- any visual or auditory information (such as new video frames, overdubbed audio, or modified subtitles) that has been added or removed; and
- whether media depicting a real person have been fabricated or simulated

We are most likely to take action (either labeling or removal, as described below) on more significant forms of alteration, such as wholly synthetic audio or video or content that has been doctored (spliced and reordered, slowed down) to change its meaning. Subtler forms of manipulated media, such as isolative editing, omission of context, or presentation with false context, may be labeled or removed on a case-by-case basis.

We will not take action to label or remove media that have been edited in ways that do not fundamentally alter their meaning, such as retouched photos or color-corrected videos.

In order to determine if media have been significantly and deceptively altered or fabricated, we may use our own technology or receive reports through partnerships with third parties. In situations where we are unable to reliably determine if media have been altered or fabricated, we may not take action to label or remove them.

2. Is the content shared in a deceptive manner?

We also consider whether the context in which media are shared could result in confusion or misunderstanding or suggests a deliberate intent to deceive people about the nature or origin of the content, for example by falsely claiming that it depicts reality. We assess the context provided alongside media to see whether it makes clear that the media have been altered or fabricated. Some of the types of context we assess in order to make this determination include:

- The text of the Tweet accompanying or within media
- Metadata associated with media
- Information on the profile of the account sharing media
- Websites linked in the Tweet, or in the profile of the account sharing media

3. Is the content likely to impact public safety or cause serious harm?

Tweets that share synthetic and manipulated media are subject to removal under this policy if they are likely to cause serious harm. Some specific harms we consider include:

- Threats to the physical safety of a person or group
- Risk of mass violence or widespread civil unrest
- Threats to the privacy or ability of a person or group to freely express themselves or participate in civic events, such as:
 - Stalking or unwanted and obsessive attention
 - Targeted content that includes tropes, epithets, or material that aims to silence someone
 - Voter suppression or intimidation

While we have other rules also intended to address these forms of harm, including our policies on violent threats, election integrity, and hateful conduct, we will err toward removal in borderline cases that might otherwise not violate existing rules for Tweets that include synthetic or manipulated media.

We also consider the time frame within which the content may be likely to impact public safety or cause serious harm, and are more likely to remove content under this policy if we find that immediate harms are likely to result from the content's presence on Twitter.

Note: We may also take action on synthetic and manipulated content under our [non-consensual nudity policy](#) (such as pornographic media altered to insert the faces of people not actually involved) or other parts of the Twitter Rules.

Labeling and removal

In most cases, if we have reason to believe that media shared in a Tweet have been significantly and deceptively altered or fabricated, we will provide additional context on Tweets sharing the media where they appear on Twitter. This means we may:

- Apply a label to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended; and/or
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or landing page.

In most cases, we will take all of the above actions on Tweets we label.

Media that meet **all three of the criteria** defined above—i.e. that are synthetic or manipulated, shared in a deceptive manner, and is likely to cause harm—may not be shared on Twitter and are subject to removal. Accounts engaging in repeated or severe violations of this policy may be permanently suspended.

* Other parts of the Twitter Rules apply and may lead to the removal of the content, particularly where there is high likelihood of severe harm, such as a threat to someone's life or physical safety.

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Copyright policy (20200917*)

What types of copyright complaints does Twitter respond to?

Twitter responds to copyright complaints submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review [Twitter's trademark policy](#). If you are concerned about a parody, newsfeed, commentary, or fan account, please see the relevant policy [here](#). These are generally not copyright issues.

Am I a copyright holder? How do I know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

What to consider before submitting a copyright complaint

Before submitting a copyright complaint to us, please consider whether or not the use could be considered [fair use](#).

If you have considered fair use, and you still wish to continue with a copyright complaint, you may want to first reach out to the user in question to see if you can resolve the matter directly with the user. You can reply to the user's Tweet or send the user a Direct Message and ask for them to remove your copyrighted content without having to contact Twitter.

Prior to submitting a formal complaint with Twitter, please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

Note: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph. If you're unsure whether or not you own the copyrights to a work, or if you're infringing upon someone else's work, please consult an attorney or another advisor.

What information do you need to process a copyright complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the complaint is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How do I file a copyright complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and [filing a copyright complaint](#). If you are logged in to [twitter.com](#), you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a full copy of your notice (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a copyright complaint.

How are claims processed?

We process copyright complaints in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. If you do not receive a ticket confirmation that means we did not receive your complaint and you should re-submit your complaint. However, please note, submitting duplicate copyright complaints will result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) and provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice. We will also forward a redacted copy of the complaint to [Lumen](#) with your personal information removed.

What information gets forwarded to the reported user(s)?

If we remove or disable access to the materials reported in a copyright complaint, the reported user(s) will receive a copy of the complaint, including the reporter's full name, email, street address, and any other information included in the complaint.

If you are uncomfortable sharing your contact information with the reported user(s), you may wish to consider appointing an agent to submit your DMCA notice on your behalf. Your agent will be required to submit the DMCA notice with valid contact information, and identify you as the content owner that they are representing.

What happens next?

Twitter's response to copyright complaints may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a copyright complaint, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the complaint, along with instructions for filing a counter-notice.

If you've not yet received a copy of the copyright complaint regarding the content removed from your account, please respond to the support ticket we sent you. If your complaint involves removal of content from Periscope, please see this help center [article](#) regarding how to request a copy of your complaint.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

My content was removed from Twitter

Why did I receive a copyright complaint?

If you receive a copyright complaint, it means that access to the content described in the complaint has been restricted. Please take the time to read through our correspondence to you, which includes information on the complaint we received as well as instructions on how to file a counter-notice. Please ensure that you are monitoring the email address associated with your Twitter account.

Tip: Removing the material reported in a copyright complaint will not resolve that complaint.

What if I want to contest the takedown?

If you believe that the materials reported in the copyright complaints were misidentified or removed in error, you may send us a counter-notification(s). A counter-notice is a request for Twitter to reinstate the removed material, and it has legal consequences. Alternatively, you may be able to seek a retraction of the copyright complaint from the reporter.

How do I seek a retraction?

The DMCA complaint you received includes the contact information of the reporter. You may want to reach out and ask them to retract their notice. The reporter can send retractions to copyright@twitter.com, and should include: (1) identification of the material that was disabled, and (2) a statement that the reporter would like to retract their DMCA notice. This is the fastest and most efficient means of resolving an unresolved copyright complaint. A retraction is at the sole discretion of the original reporter.

When should I file a counter-notice?

A counter-notice is a request for Twitter to reinstate the removed material, and is the start of a legal process that has legal consequences. For example, submitting a counter notice indicates that you consent to the jurisdiction of a U.S. Federal court and that you consent to the disclosure of your personal information to the reporter and Lumen website.

With these considerations in mind, you may file a counter-notice if you believe that this material was misidentified, or you have a good faith belief that the material should not have been removed. If you're unsure whether or not you should file a counter-notice, you may want to consult with an attorney.

Tip: Re-posting material removed in response to a copyright complaint may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notice rather than re-posting the material.

What information do you need to process a counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What happens after I submit a counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. This means that the contact information that is submitted in your counter-notice will be shared to the person who filed the original notice.

If the copyright owner disagrees that the content was removed in error or misidentification, they may pursue legal action against you. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a copyright complaint or counter-notice is serious business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

What happens if my account receives multiple copyright complaints?

If multiple copyright complaints are received Twitter may lock accounts or take other actions to warn repeat violators. These warnings may vary across Twitter's services. Under appropriate circumstances we may suspend user accounts under our repeat infringer policy. However, we may take retractions and counter-notices into account when applying our repeat infringer policy.

Our approach to policy development and enforcement philosophy

(20190520*)

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a

larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (20200204*)

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Synthetic and manipulated media: You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.

5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Terrorism and violent extremism policy (20190607*)

Overview

March 2019

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for terrorist organizations or violent extremist groups and individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments in this context are informed by national and international terrorism designations. We also assess organizations under our violent extremist group criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

We examine a group's activities both on and off Twitter to determine whether they engage in and/or promote violence against civilians to advance a political, religious and/or social cause.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a terrorist organization or violent extremist group;
- recruiting for a terrorist organization or violent extremist group;
- providing or distributing services (e.g., financial, media/propaganda) to further a terrorist organization's or violent extremist group's stated goals; and
- using the insignia or symbols of terrorist organizations or violent extremist groups to promote them.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to military or government entities.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wish for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy (20201012*)

Overview

October 2020

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;

- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives;
- identifying alleged victims of childhood sexual exploitation by name or image; and
- promoting or normalizing sexual attraction to minors as a form of identity or sexual orientation.

What is not a violation of this policy?

- Discussions related to child sexual exploitation are permitted, provided they don't normalise, promote or glorify child sexual exploitation **in any way**;
- Conversations about help-seeking behaviour of individuals who may struggle with an attraction to minors;
- Advocacy against illegal or harmful activity involving minors provided there is no sharing of, or linking to any material featuring child sexual exploitation; and
- Depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios (e.g., art works by internationally renowned artists that feature minors, news media reportage, media created and shared for scientific or educational purposes).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

Abusive behavior (20190607*)

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing or hoping serious harm on a person or group of people

We do not tolerate content that wishes, hopes or expresses a desire for death, serious bodily harm or fatal disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone's body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy (20200331*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., “I will kill you”.

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., “I hope you get cancer and die.”

Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”

Saying that a group of individuals deserve serious physical injury, e.g., “If this group of protesters don’t shut up, they deserve to be shot.”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or

violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion, age, disability, or serious disease.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs.

When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can

[submit an appeal](#).

Suicide and Self-harm Policy (20200805*)

Overview

You may not promote or encourage suicide or self-harm.

At Twitter, we recognize that suicide and self-harm are significant social & public health challenges that require collaboration between all stakeholders – public, private, and civil society – and that we have a role and responsibility to help people access and receive support when they need it.

When developing this policy, we consulted extensively with experts to ensure that people who have engaged in self-harm or experienced suicidal thoughts can share their personal experiences. We also recognized the need to

protect people from the potential harm caused by exposure to content that could promote or encourage self-harm – intentionally or inadvertently. That's why our policy prohibits content that promotes or encourages self-harming behaviors and provides support to those undergoing experiences with self-harm or suicidal thoughts.

What is in violation of this policy?

Under this policy, you can't promote, or otherwise encourage, suicide or self-harm. We define promotion and encouragement to include statements such as "the most effective", "the easiest", "the best", "the most successful", "you should", "why don't you". Violations of this policy can occur via Tweets, images or videos, including live video.

We define suicide to be the act of taking one's own life. We define self-harm to include:

- self-inflicted physical injuries e.g., cutting; and
- eating disorders e.g., bulimia, anorexia.

Violations of this policy include, but are not limited to:

- encouraging someone to physically harm or kill themselves;
- asking others for encouragement to engage in self-harm or suicide, including seeking partners for group suicides or suicide games; and
- sharing information, strategies, methods or instructions that would assist people to engage in self-harm and suicide.

What is not a violation of this policy?

Some examples of behavior that are not considered a violation of this policy include:

- sharing personal stories and experiences related to self-harm or suicide;
- sharing coping mechanisms and resources for addressing self-harm or suicidal thoughts; and
- discussions that are focused on research, advocacy, and education related to self-harm or suicide prevention.

Note: people can share their personal experiences, but should avoid sharing detailed information about specific strategies or methods related to self-harm, as this could inadvertently encourage this behavior.

Who can report violations of this policy?

Anyone can report content that may encourage or promote suicide or self-harm via our in-app reporting or our [specialized reporting form](#). These reports are routed to a dedicated team who evaluate each case individually.

Note: if we receive a report that someone has expressed an intention to engage in self-harm or suicide, we will contact them directly, encourage them to seek support, and provide information about dedicated online and hotline resources. We may also work with law enforcement officials where appropriate, for example, if we receive a valid emergency disclosure request as defined in our [Law Enforcement Guidelines](#).

How to report violations of this policy

To ensure that we handle reports sensitively, our in-app reporting provides separate options for people who may be expressing an intention to harm themselves, and content that is encouraging or promoting self-harm or suicide.

Expressing intentions of self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Reporting form

You can also report this content for review via [our dedicated report form](#).

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Encouraging self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

What happens if you violate this policy?

Our enforcement approach depends on the type of content being shared, whether or not the reported account is encouraging or promoting self-harm or suicide, and the account's previous history of violations.

If you violate this policy by sharing content that intentionally encourages others to harm themselves, ask others to encourage you to harm yourself, or share detailed information or instructions related to self-harm or suicide methods, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy, or if your account is dedicated to promoting or encouraging self-harm or suicide, your account will be permanently suspended. If cases include images or videos

related to self-harm or suicide, we will also evaluate this content under [our sensitive media policy](#). If you believe that your account was suspended in error, [you can submit an appeal](#).

We may also take steps to prevent the spread of instructional material hosted on third-party websites by [marking such links as unsafe](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Learn more](#) about how you can support someone experiencing thoughts of self-harm or suicide.

Visit our Safety Center for a list of local [mental health resources](#) and read [our blog on suicide prevention](#) to learn more about our work.

Sensitive media policy (20191204*)

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;

- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20190607*)

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing

their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;
 - descriptions of physical appearance;
 - gossip, rumours, accusations, and allegations; and
 - screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy (20191217*)

Overview

November 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy **(20200918*)**

Overview

September 2020

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are;
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting; and
- coordinated harmful activity that encourages or promotes behavior which violates the [Twitter Rules](#).

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts or by coordinating with others to violate the Twitter Rules. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate;
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account; and
 - coordinating with others to engage in or promote violations of the Twitter Rules, including violations of our [abusive behavior](#) policy.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);

- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in “follow trains,” “decks,” and “Retweet for Retweet” behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- “follow churn” – following and then unfollowing large numbers of accounts in an effort to inflate one’s own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account’s followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person’s browser (malware) or computer or to compromise a person’s privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally;
- coordinating with others to express ideas, viewpoints, support, or opposition towards a cause, provided such behavior does not result in violations of the Twitter Rules; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don’t violate other rules. Some examples include:

- organizations with related but separate chapters or branches, such as a business with multiple locations;
- operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
- hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Denylisting URLs

We denylist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [financial scam policy](#), our [hacked materials policy](#), our approach to [coordinated harmful activity](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Civic integrity policy (20200605*)

Overview

May 2020

You may not use Twitter's services for the purpose of manipulating or interfering in elections or other civic processes. This includes posting or sharing content that may suppress participation or mislead people about when, where, or how to participate in a civic process.

The public conversation occurring on Twitter is never more important than during elections and other civic events. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt civic processes, including through the distribution of false or misleading information about the procedures or circumstances around participation in a civic process.

What is a civic process?

Twitter considers civic processes to be events or procedures mandated, organized, and conducted by the governing and/or electoral body of a country, state, region, district, or municipality to address a matter of common concern through public participation. Some examples of civic processes may include:

- Political elections
- Censuses
- Major referenda and ballot initiatives

What is in violation of this policy?

We prohibit 3 categories of manipulative behavior and content under this policy:

Misleading information about how to participate

You can't share false or misleading information about how to participate in an election or other civic process. This includes but is not limited to:

- misleading information about procedures to participate in a civic process (for example, that you can vote by Tweet, text message, email, or phone call in jurisdictions where these are not a possibility);
- misleading information about requirements for participation, including identification or citizenship requirements; and
- misleading statements or information about the official, announced date or time of a civic process.

Suppression and intimidation

You can't share false or misleading information intended to intimidate or dissuade people from participating in an election or other civic process. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to voting in an election, polling places, or collecting census information;
- misleading claims about long lines, equipment problems, or other disruptions at voting locations during election periods;
- misleading claims about process procedures or techniques which could dissuade people from participating; and
- threats regarding voting locations or other key places or events (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account policy](#).

What is not a violation of this policy?

Not all false or untrue information about politics or civic processes constitutes manipulation or interference. In the absence of other policy violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected or appointed official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information;
- voting and audience participation for competitions, game shows, or other entertainment purposes;
- broad, non-specific statements about the integrity of elections or civic processes (such as unsubstantiated claims that an election is "rigged"); and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election or civic process. Therefore, we enable reporting of false or misleading information about civic processes in advance of major events, for people located in the relevant countries and locations. We also work with select government and civil society partners in these countries to provide additional channels for reporting and expedited review.

For civic processes with multiple stages or parts, such as primary elections or lengthy campaigns, reporting will be enabled leading up to the first officially-sanctioned event associated with the civic process.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country at the relevant time, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in civic processes.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in a civic process.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our civic integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Impersonation policy (20191001*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Synthetic and manipulated media policy **(20200204*)**

Overview

You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context.

You should be able to find reliable information on Twitter. That means understanding whether the content you see is real or fabricated and having the ability to find more context about what you see on Twitter. Therefore, we may label Tweets that include media (videos, audio, and images) that have been deceptively altered or fabricated. In addition, you may not share deceptively altered media on Twitter in ways that mislead or deceive people about the media's authenticity where threats to physical safety or other serious harm may result.

We use the following criteria as we consider Tweets and media for labeling or removal under this policy as part of our ongoing work to enforce our rules and ensure healthy and safe conversation on Twitter (additional information is available below):

Is the content significantly and deceptively altered or fabricated?	Is the content shared in a deceptive manner?	Is the content likely to impact public safety or cause serious harm?	
✓	✗	✗	Content may be labeled.
✗	✓	✗	Content may be labeled.
✓	✗	✓	Content is likely to be labeled, or may be removed.*
✓	✓	✗	Content is likely to be labeled.
✓	✓	✓	Content is likely to be removed.

1. Is the content synthetic or manipulated?

In order for content to be labeled or removed under this policy, we must have reason to believe that media, or the context in which media are presented, are significantly and deceptively altered or manipulated. Synthetic and manipulated media take many different forms and people can employ a wide range of technologies to produce these media. In assessing whether media have been significantly and deceptively altered or fabricated, some of the factors we consider include:

- whether the content has been substantially edited in a manner that fundamentally alters its composition, sequence, timing, or framing;
- any visual or auditory information (such as new video frames, overdubbed audio, or modified subtitles) that has been added or removed; and
- whether media depicting a real person have been fabricated or simulated

We are most likely to take action (either labeling or removal, as described below) on more significant forms of alteration, such as wholly synthetic audio or video or content that has been doctored (spliced and reordered, slowed down) to change its meaning. Subtler forms of manipulated media, such as isolative editing, omission of context, or presentation with false context, may be labeled or removed on a case-by-case basis.

We will not take action to label or remove media that have been edited in ways that do not fundamentally alter their meaning, such as retouched photos or color-corrected videos.

In order to determine if media have been significantly and deceptively altered or fabricated, we may use our own technology or receive reports through partnerships with third parties. In situations where we are unable to reliably determine if media have been altered or fabricated, we may not take action to label or remove them.

2. Is the content shared in a deceptive manner?

We also consider whether the context in which media are shared could result in confusion or misunderstanding or suggests a deliberate intent to deceive people about the nature or origin of the content, for example by falsely claiming that it depicts reality. We assess the context provided alongside media to see whether it makes clear that the media have been altered or fabricated. Some of the types of context we assess in order to make this determination include:

- The text of the Tweet accompanying or within media
- Metadata associated with media
- Information on the profile of the account sharing media
- Websites linked in the Tweet, or in the profile of the account sharing media

3. Is the content likely to impact public safety or cause serious harm?

Tweets that share synthetic and manipulated media are subject to removal under this policy if they are likely to cause serious harm. Some specific harms we consider include:

- Threats to the physical safety of a person or group
- Risk of mass violence or widespread civil unrest
- Threats to the privacy or ability of a person or group to freely express themselves or participate in civic events, such as:
 - Stalking or unwanted and obsessive attention
 - Targeted content that includes tropes, epithets, or material that aims to silence someone
 - Voter suppression or intimidation

While we have other rules also intended to address these forms of harm, including our policies on violent threats, election integrity, and hateful conduct, we will err toward removal in borderline cases that might otherwise not violate existing rules for Tweets that include synthetic or manipulated media.

We also consider the time frame within which the content may be likely to impact public safety or cause serious harm, and are more likely to remove content under this policy if we find that immediate harms are likely to result from the content's presence on Twitter.

Note: We may also take action on synthetic and manipulated content under our [non-consensual nudity policy](#) (such as pornographic media altered to insert the faces of people not actually involved) or other parts of the Twitter Rules.

Labeling and removal

In most cases, if we have reason to believe that media shared in a Tweet have been significantly and deceptively altered or fabricated, we will provide additional context on Tweets sharing the media where they appear on Twitter. This means we may:

- Apply a label to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended; and/or
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or landing page.

In most cases, we will take all of the above actions on Tweets we label.

Media that meet **all three of the criteria** defined above—i.e. that are synthetic or manipulated, shared in a deceptive manner, and is likely to cause harm—may not be shared on Twitter and are subject to removal. Accounts engaging in repeated or severe violations of this policy may be permanently suspended.

* Other parts of the Twitter Rules apply and may lead to the removal of the content, particularly where there is high likelihood of severe harm, such as a threat to someone's life or physical safety.

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Copyright policy (20200917*)

What types of copyright complaints does Twitter respond to?

Twitter responds to copyright complaints submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review [Twitter's trademark policy](#). If you are concerned about a parody, newsfeed, commentary, or fan account, please see the relevant policy [here](#). These are generally not copyright issues.

Am I a copyright holder? How do I know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

What to consider before submitting a copyright complaint

Before submitting a copyright complaint to us, please consider whether or not the use could be considered [fair use](#).

If you have considered fair use, and you still wish to continue with a copyright complaint, you may want to first reach out to the user in question to see if you can resolve the matter directly with the user. You can reply to the user's Tweet or send the user a Direct Message and ask for them to remove your copyrighted content without having to contact Twitter.

Prior to submitting a formal complaint with Twitter, please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

Note: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph. If you're unsure whether or not you own the copyrights to a work, or if you're infringing upon someone else's work, please consult an attorney or another advisor.

What information do you need to process a copyright complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;

4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the complaint is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How do I file a copyright complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and [filing a copyright complaint](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a full copy of your notice (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a copyright complaint.

How are claims processed?

We process copyright complaints in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. If you do not receive a ticket confirmation that means we did not receive your complaint and you should re-submit your complaint. However, please note, submitting duplicate copyright complaints will result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) and provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice. We will also forward a redacted copy of the complaint to [Lumen](#), with your personal information removed.

What information gets forwarded to the reported user(s)?

If we remove or disable access to the materials reported in a copyright complaint, the reported user(s) will receive a copy of the complaint, including the reporter's full name, email, street address, and any other information included in the complaint.

If you are uncomfortable sharing your contact information with the reported user(s), you may wish to consider appointing an agent to submit your DMCA notice on your behalf. Your agent will be required to submit the DMCA notice with valid contact information, and identify you as the content owner that they are representing.

What happens next?

Twitter's response to copyright complaints may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a copyright complaint, Twitter will make a

good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the complaint, along with instructions for filing a counter-notice.

If you've not yet received a copy of the copyright complaint regarding the content removed from your account, please respond to the support ticket we sent you. If your complaint involves removal of content from Periscope, please see this help center [article](#) regarding how to request a copy of your complaint.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from **@Username** has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

My content was removed from Twitter

Why did I receive a copyright complaint?

If you receive a copyright complaint, it means that access to the content described in the complaint has been restricted. Please take the time to read through our correspondence to you, which includes information on the complaint we received as well as instructions on how to file a counter-notice. Please ensure that you are monitoring the email address associated with your Twitter account.

Tip: Removing the material reported in a copyright complaint will not resolve that complaint.

What if I want to contest the takedown?

If you believe that the materials reported in the copyright complaints were misidentified or removed in error, you may send us a counter-notification(s). A counter-notice is a request for Twitter to reinstate the removed material, and it has legal consequences. Alternatively, you may be able to seek a retraction of the copyright complaint from the reporter.

How do I seek a retraction?

The DMCA complaint you received includes the contact information of the reporter. You may want to reach out and ask them to retract their notice. The reporter can send retractions to copyright@twitter.com, and should include: (1) identification of the material that was disabled, and (2) a statement that the reporter would like to retract their DMCA notice. This is the fastest and most efficient means of resolving an unresolved copyright complaint. A retraction is at the sole discretion of the original reporter.

When should I file a counter-notice?

A counter-notice is a request for Twitter to reinstate the removed material, and is the start of a legal process that has legal consequences. For example, submitting a counter notice indicates that you consent to the jurisdiction of

a U.S. Federal court and that you consent to the disclosure of your personal information to the reporter and Lumen website.

With these considerations in mind, you may file a counter-notice if you believe that this material was misidentified, or you have a good faith belief that the material should not have been removed. If you're unsure whether or not you should file a counter-notice, you may want to consult with an attorney.

Tip: Re-posting material removed in response to a copyright complaint may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notice rather than re-posting the material.

What information do you need to process a counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What happens after I submit a counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. This means that the contact information that is submitted in your counter-notice will be shared to the person who filed the original notice.

If the copyright owner disagrees that the content was removed in error or misidentification, they may pursue legal action against you. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a copyright complaint or counter-notice is serious business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

What happens if my account receives multiple copyright complaints?

If multiple copyright complaints are received Twitter may lock accounts or take other actions to warn repeat violators. These warnings may vary across Twitter's services. Under appropriate circumstances we may suspend user accounts under our repeat infringer policy. However, we may take retractions and counter-notices into account when applying our repeat infringer policy.

Our approach to policy development and enforcement philosophy

(20190520*)

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or

country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

[Violence](#): You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

[Terrorism/violent extremism](#): You may not threaten or promote terrorism or violent extremism. [Learn more](#).

[Child sexual exploitation](#): We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

[Abuse/harassment](#): You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

[Hateful conduct](#): You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

[Suicide or self-harm](#): You may not promote or encourage suicide or self-harm. [Learn more](#).

[Sensitive media, including graphic violence and adult content](#): You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

[Illegal or certain regulated goods or services](#): You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

[Private information](#): You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

[Non-consensual nudity](#): You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

[Platform manipulation and spam](#): You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

[Election integrity](#): You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

[Impersonation](#): You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Synthetic and manipulated media: You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Violent organizations policy

Overview

October 2020

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are

informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

Other violent organizations

Other violent organizations are those that meet all of the below criteria:

- a collection of individuals with a shared purpose; and
- have systematically targeted civilians with violence.

We examine a group's activities both on and off Twitter to determine whether they satisfy the above criteria.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a violent organization;
- recruiting for a violent organization;
- providing or distributing services (e.g., financial, media/propaganda) to further a violent organization's stated goals; and
- using the insignia or symbol of violent organizations to promote them or indicate affiliation or support.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to state or governmental organizations.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wishing for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy

Overview

March 2019

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives; and
- identifying alleged victims of childhood sexual exploitation by name or image.

What is not a violation of this policy?

Discussions related to child sexual exploitation as a phenomenon or attraction towards minors are permitted, provided they don't promote or glorify child sexual exploitation **in any way**. Artistic depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios e.g., works by internationally renowned artists that feature minors.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

What you can do

If you think you've found a Twitter account distributing or promoting child sexual exploitation, please report it through our designated [form](#). Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. For help on how to find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with NCMEC [here](#) or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the International Association of Internet Hotlines [website](#).

Our partners

We partner with multiple organisations around the world whose work is dedicated to fighting child sexual exploitation. You can find a full list of our safety partners [here](#).

Abusive behavior

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing or hoping serious harm on a person or group of people

We do not tolerate content that wishes, hopes or expresses a desire for death, serious bodily harm or fatal disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone's body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you".

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., “I hope you get cancer and die.”

Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”

Saying that a group of individuals deserve serious physical injury, e.g., “If this group of protesters don’t shut up, they deserve to be shot.”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion, age, disability, or serious disease.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can

[submit an appeal.](#)

Suicide and Self-harm Policy

Overview

You may not promote or encourage suicide or self-harm.

At Twitter, we recognize that suicide and self-harm are significant social & public health challenges that require collaboration between all stakeholders – public, private, and civil society – and that we have a role and responsibility to help people access and receive support when they need it.

When developing this policy, we consulted extensively with experts to ensure that people who have engaged in self-harm or experienced suicidal thoughts can share their personal experiences. We also recognized the need to protect people from the potential harm caused by exposure to content that could promote or encourage self-harm – intentionally or inadvertently. That's why our policy prohibits content that promotes or encourages self-harming behaviors and provides support to those undergoing experiences with self-harm or suicidal thoughts.

What is in violation of this policy?

Under this policy, you can't promote, or otherwise encourage, suicide or self-harm. We define promotion and encouragement to include statements such as "the most effective", "the easiest", "the best", "the most successful", "you should", "why don't you". Violations of this policy can occur via Tweets, images or videos, including live video.

We define suicide to be the act of taking one's own life. We define self-harm to include:

- self-inflicted physical injuries e.g., cutting; and
- eating disorders e.g., bulimia, anorexia.

Violations of this policy include, but are not limited to:

- encouraging someone to physically harm or kill themselves;
- asking others for encouragement to engage in self-harm or suicide, including seeking partners for group suicides or suicide games; and
- sharing information, strategies, methods or instructions that would assist people to engage in self-harm and suicide.

What is not a violation of this policy?

Some examples of behavior that are not considered a violation of this policy include:

- sharing personal stories and experiences related to self-harm or suicide;
- sharing coping mechanisms and resources for addressing self-harm or suicidal thoughts; and
- discussions that are focused on research, advocacy, and education related to self-harm or suicide prevention.

Note: people can share their personal experiences, but should avoid sharing detailed information about specific strategies or methods related to self-harm, as this could inadvertently encourage this behavior.

Who can report violations of this policy?

Anyone can report content that may encourage or promote suicide or self-harm via our in-app reporting or our [specialized reporting form](#). These reports are routed to a dedicated team who evaluate each case individually.

Note: if we receive a report that someone has expressed an intention to engage in self-harm or suicide, we will contact them directly, encourage them to seek support, and provide information about dedicated online and hotline resources. We may also work with law enforcement officials where appropriate, for example, if we receive a valid emergency disclosure request as defined in our [Law Enforcement Guidelines](#).

How to report violations of this policy

To ensure that we handle reports sensitively, our in-app reporting provides separate options for people who may be expressing an intention to harm themselves, and content that is encouraging or promoting self-harm or suicide.

Expressing intentions of self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Reporting form

You can also report this content for review via [our dedicated report form](#).

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Encouraging self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

What happens if you violate this policy?

Our enforcement approach depends on the type of content being shared, whether or not the reported account is encouraging or promoting self-harm or suicide, and the account's previous history of violations.

If you violate this policy by sharing content that intentionally encourages others to harm themselves, ask others to encourage you to harm yourself, or share detailed information or instructions related to self-harm or suicide methods, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy, or if your account is dedicated to promoting or encouraging self-harm or suicide, your account will be permanently suspended. If cases include images or videos related to self-harm or suicide, we will also evaluate this content under [our sensitive media policy](#). If you believe that your account was suspended in error, [you can submit an appeal](#).

We may also take steps to prevent the spread of instructional material hosted on third-party websites by [marking such links as unsafe](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Learn more](#) about how you can support someone experiencing thoughts of self-harm or suicide.

Visit our Safety Center for a list of local [mental health resources](#) and read [our blog on suicide prevention](#) to learn more about our work.

Sensitive media policy

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to

be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and

- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;
 - descriptions of physical appearance;
 - gossip, rumours, accusations, and allegations; and
 - screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy

Overview

March 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you

violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy

Overview

March 2019

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are; and
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting.

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:

- posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
- engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate; and
- coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in “follow trains,” “decks,” and “Retweet for Retweet” behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- “follow churn” – following and then unfollowing large numbers of accounts in an effort to inflate one's own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account's followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person's browser (malware) or computer or to compromise a person's privacy (phishing); and

- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don't violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

- Select **Report Tweet** from the icon.
- Select **It's suspicious or spam**.
- Select the option that best tells us how the Tweet is suspicious or spreading spam.
- Submit your report.

Desktop

You can report this content via desktop as follows:

- Select **Report Tweet** from the icon.
- Select **It's suspicious or spam**.
- Select the option that best tells us how the Tweet is suspicious or spreading spam.
- Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Blacklisting URLs

We blacklist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [hacked materials policy](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Civic integrity policy

Overview

May 2020

You may not use Twitter's services for the purpose of manipulating or interfering in elections or other civic processes. This includes posting or sharing content that may suppress participation or mislead people about when, where, or how to participate in a civic process.

The public conversation occurring on Twitter is never more important than during elections and other civic events. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt civic processes,

including through the distribution of false or misleading information about the procedures or circumstances around participation in a civic process.

What is a civic process?

Twitter considers civic processes to be events or procedures mandated, organized, and conducted by the governing and/or electoral body of a country, state, region, district, or municipality to address a matter of common concern through public participation. Some examples of civic processes may include:

- Political elections
- Censuses
- Major referenda and ballot initiatives

What is in violation of this policy?

We prohibit 3 categories of manipulative behavior and content under this policy:

Misleading information about how to participate

You can't share false or misleading information about how to participate in an election or other civic process. This includes but is not limited to:

- misleading information about procedures to participate in a civic process (for example, that you can vote by Tweet, text message, email, or phone call in jurisdictions where these are not a possibility);
- misleading information about requirements for participation, including identification or citizenship requirements; and
- misleading statements or information about the official, announced date or time of a civic process.

Suppression and intimidation

You can't share false or misleading information intended to intimidate or dissuade people from participating in an election or other civic process. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to voting in an election, polling places, or collecting census information;
- misleading claims about long lines, equipment problems, or other disruptions at voting locations during election periods;
- misleading claims about process procedures or techniques which could dissuade people from participating; and
- threats regarding voting locations or other key places or events (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account policy](#).

What is not a violation of this policy?

Not all false or untrue information about politics or civic processes constitutes manipulation or interference. In the absence of other policy violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected or appointed official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information;

- voting and audience participation for competitions, game shows, or other entertainment purposes;
- broad, non-specific statements about the integrity of elections or civic processes (such as unsubstantiated claims that an election is "rigged"); and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election or civic process. Therefore, we enable reporting of false or misleading information about civic processes in advance of major events, for people located in the relevant countries and locations. We also work with select government and civil society partners in these countries to provide additional channels for reporting and expedited review.

For civic processes with multiple stages or parts, such as primary elections or lengthy campaigns, reporting will be enabled leading up to the first officially-sanctioned event associated with the civic process.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country at the relevant time, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in civic processes.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in a civic process.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our civic integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Synthetic and manipulated media policy

Overview

You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context.

You should be able to find reliable information on Twitter. That means understanding whether the content you see is real or fabricated and having the ability to find more context about what you see on Twitter. Therefore, we may label Tweets that include media (videos, audio, and images) that have been deceptively altered or fabricated. In addition, you may not share deceptively altered media on Twitter in ways that mislead or deceive people about the media's authenticity where threats to physical safety or other serious harm may result.

We use the following criteria as we consider Tweets and media for labeling or removal under this policy as part of our ongoing work to enforce our rules and ensure healthy and safe conversation on Twitter (additional information is available below):

Is the content significantly and deceptively altered or fabricated?	Is the content shared in a deceptive manner?	Is the content likely to impact public safety or cause serious harm?	
✓	✗	✗	Content may be labeled.
✗	✓	✗	Content may be labeled.
✓	✗	✓	Content is likely to be labeled, or may be removed.*
✓	✓	✗	Content is likely to be labeled.
✓	✓	✓	Content is likely to be removed.

1. Is the content synthetic or manipulated?

In order for content to be labeled or removed under this policy, we must have reason to believe that media, or the context in which media are presented, are significantly and deceptively altered or manipulated. Synthetic and manipulated media take many different forms and people can employ a wide range of technologies to produce these media. In assessing whether media have been significantly and deceptively altered or fabricated, some of the factors we consider include:

- whether the content has been substantially edited in a manner that fundamentally alters its composition, sequence, timing, or framing;
- any visual or auditory information (such as new video frames, overdubbed audio, or modified subtitles) that has been added or removed; and
- whether media depicting a real person have been fabricated or simulated

We are most likely to take action (either labeling or removal, as described below) on more significant forms of alteration, such as wholly synthetic audio or video or content that has been doctored (spliced and reordered, slowed down) to change its meaning. Subtler forms of manipulated media, such as isolative editing, omission of context, or presentation with false context, may be labeled or removed on a case-by-case basis.

We will not take action to label or remove media that have been edited in ways that do not fundamentally alter their meaning, such as retouched photos or color-corrected videos.

In order to determine if media have been significantly and deceptively altered or fabricated, we may use our own technology or receive reports through partnerships with third parties. In situations where we are unable to reliably determine if media have been altered or fabricated, we may not take action to label or remove them.

2. Is the content shared in a deceptive manner?

We also consider whether the context in which media are shared could result in confusion or misunderstanding or suggests a deliberate intent to deceive people about the nature or origin of the content, for example by falsely claiming that it depicts reality. We assess the context provided alongside media to see whether it makes clear that the media have been altered or fabricated. Some of the types of context we assess in order to make this determination include:

- The text of the Tweet accompanying or within media
- Metadata associated with media
- Information on the profile of the account sharing media
- Websites linked in the Tweet, or in the profile of the account sharing media

3. Is the content likely to impact public safety or cause serious harm?

Tweets that share synthetic and manipulated media are subject to removal under this policy if they are likely to cause serious harm. Some specific harms we consider include:

- Threats to the physical safety of a person or group
- Risk of mass violence or widespread civil unrest
- Threats to the privacy or ability of a person or group to freely express themselves or participate in civic events, such as:
 - Stalking or unwanted and obsessive attention
 - Targeted content that includes tropes, epithets, or material that aims to silence someone
 - Voter suppression or intimidation

While we have other rules also intended to address these forms of harm, including our policies on violent threats, election integrity, and hateful conduct, we will err toward removal in borderline cases that might otherwise not violate existing rules for Tweets that include synthetic or manipulated media.

We also consider the time frame within which the content may be likely to impact public safety or cause serious harm, and are more likely to remove content under this policy if we find that immediate harms are likely to result from the content's presence on Twitter.

Note: We may also take action on synthetic and manipulated content under our [non-consensual nudity policy](#) (such as pornographic media altered to insert the faces of people not actually involved) or other parts of the Twitter Rules.

Labeling and removal

In most cases, if we have reason to believe that media shared in a Tweet have been significantly and deceptively altered or fabricated, we will provide additional context on Tweets sharing the media where they appear on Twitter. This means we may:

- Apply a label to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended; and/or
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or landing page.

In most cases, we will take all of the above actions on Tweets we label.

Media that meet **all three of the criteria** defined above—i.e. that are synthetic or manipulated, shared in a deceptive manner, and is likely to cause harm—may not be shared on Twitter and are subject to removal. Accounts engaging in repeated or severe violations of this policy may be permanently suspended.

* Other parts of the Twitter Rules apply and may lead to the removal of the content, particularly where there is high likelihood of severe harm, such as a threat to someone's life or physical safety.

Trademark policy

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Copyright policy

What types of copyright complaints does Twitter respond to?

Twitter responds to copyright complaints submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review [Twitter's trademark policy](#). If you are concerned about a parody, newsfeed, commentary, or fan account, please see the relevant policy [here](#). These are generally not copyright issues.

Am I a copyright holder? How do I know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

What to consider before submitting a copyright complaint

Before submitting a copyright complaint to us, please consider whether or not the use could be considered [fair use](#).

If you have considered fair use, and you still wish to continue with a copyright complaint, you may want to first reach out to the user in question to see if you can resolve the matter directly with the user. You can reply to the user's Tweet or send the user a Direct Message and ask for them to remove your copyrighted content without having to contact Twitter.

Prior to submitting a formal complaint with Twitter, please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

Note: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph. If you're unsure whether or not you own the copyrights to a work, or if you're infringing upon someone else's work, please consult an attorney or another advisor.

What information do you need to process a copyright complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;

4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the complaint is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How do I file a copyright complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and [filing a copyright complaint](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a full copy of your notice (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a copyright complaint.

How are claims processed?

We process copyright complaints in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. If you do not receive a ticket confirmation that means we did not receive your complaint and you should re-submit your complaint. However, please note, submitting duplicate copyright complaints will result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) and provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice. We will also forward a redacted copy of the complaint to [Lumen](#), with your personal information removed.

What information gets forwarded to the reported user(s)?

If we remove or disable access to the materials reported in a copyright complaint, the reported user(s) will receive a copy of the complaint, including the reporter's full name, email, street address, and any other information included in the complaint.

If you are uncomfortable sharing your contact information with the reported user(s), you may wish to consider appointing an agent to submit your DMCA notice on your behalf. Your agent will be required to submit the DMCA notice with valid contact information, and identify you as the content owner that they are representing.

What happens next?

Twitter's response to copyright complaints may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a copyright complaint, Twitter will make a

good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the complaint, along with instructions for filing a counter-notice.

If you've not yet received a copy of the copyright complaint regarding the content removed from your account, please respond to the support ticket we sent you. If your complaint involves removal of content from Periscope, please see this help center [article](#) regarding how to request a copy of your complaint.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from **@Username** has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

My content was removed from Twitter

Why did I receive a copyright complaint?

If you receive a copyright complaint, it means that access to the content described in the complaint has been restricted. Please take the time to read through our correspondence to you, which includes information on the complaint we received as well as instructions on how to file a counter-notice. Please ensure that you are monitoring the email address associated with your Twitter account.

Tip: Removing the material reported in a copyright complaint will not resolve that complaint.

What if I want to contest the takedown?

If you believe that the materials reported in the copyright complaints were misidentified or removed in error, you may send us a counter-notification(s). A counter-notice is a request for Twitter to reinstate the removed material, and it has legal consequences. Alternatively, you may be able to seek a retraction of the copyright complaint from the reporter.

How do I seek a retraction?

The DMCA complaint you received includes the contact information of the reporter. You may want to reach out and ask them to retract their notice. The reporter can send retractions to copyright@twitter.com, and should include: (1) identification of the material that was disabled, and (2) a statement that the reporter would like to retract their DMCA notice. This is the fastest and most efficient means of resolving an unresolved copyright complaint. A retraction is at the sole discretion of the original reporter.

When should I file a counter-notice?

A counter-notice is a request for Twitter to reinstate the removed material, and is the start of a legal process that has legal consequences. For example, submitting a counter notice indicates that you consent to the jurisdiction of

a U.S. Federal court and that you consent to the disclosure of your personal information to the reporter and Lumen website.

With these considerations in mind, you may file a counter-notice if you believe that this material was misidentified, or you have a good faith belief that the material should not have been removed. If you're unsure whether or not you should file a counter-notice, you may want to consult with an attorney.

Tip: Re-posting material removed in response to a copyright complaint may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notice rather than re-posting the material.

What information do you need to process a counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What happens after I submit a counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. This means that the contact information that is submitted in your counter-notice will be shared to the person who filed the original notice.

If the copyright owner disagrees that the content was removed in error or misidentification, they may pursue legal action against you. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a copyright complaint or counter-notice is serious business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

What happens if my account receives multiple copyright complaints?

If multiple copyright complaints are received Twitter may lock accounts or take other actions to warn repeat violators. These warnings may vary across Twitter's services. Under appropriate circumstances we may suspend user accounts under our repeat infringer policy. However, we may take retractions and counter-notices into account when applying our repeat infringer policy.

Our approach to policy development and enforcement philosophy

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

[Violence](#): You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

[Terrorism/violent extremism](#): You may not threaten or promote terrorism or violent extremism. [Learn more](#).

[Child sexual exploitation](#): We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

[Abuse/harassment](#): You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

[Hateful conduct](#): You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

[Suicide or self-harm](#): You may not promote or encourage suicide or self-harm. [Learn more](#).

[Sensitive media, including graphic violence and adult content](#): You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

[Illegal or certain regulated goods or services](#): You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

[Private information](#): You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

[Non-consensual nudity](#): You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

[Platform manipulation and spam](#): You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

[Election integrity](#): You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

[Impersonation](#): You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Synthetic and manipulated media: You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Violent organizations policy

Overview

October 2020

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are

informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

Other violent organizations

Other violent organizations are those that meet all of the below criteria:

- a collection of individuals with a shared purpose; and
- have systematically targeted civilians with violence.

We examine a group's activities both on and off Twitter to determine whether they satisfy the above criteria.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a violent organization;
- recruiting for a violent organization;
- providing or distributing services (e.g., financial, media/propaganda) to further a violent organization's stated goals; and
- using the insignia or symbol of violent organizations to promote them or indicate affiliation or support.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to state or governmental organizations.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wishing for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy

Overview

March 2019

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives; and
- identifying alleged victims of childhood sexual exploitation by name or image.

What is not a violation of this policy?

Discussions related to child sexual exploitation as a phenomenon or attraction towards minors are permitted, provided they don't promote or glorify child sexual exploitation **in any way**. Artistic depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios e.g., works by internationally renowned artists that feature minors.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

What you can do

If you think you've found a Twitter account distributing or promoting child sexual exploitation, please report it through our designated [form](#). Please provide the username and links to the relevant tweets that lead you to believe the account should be investigated. For help on how to find the direct status links to individual Tweets, see our [help page](#).

If you have found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should file a report with NCMEC [here](#) or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the International Association of Internet Hotlines [website](#).

Our partners

We partner with multiple organisations around the world whose work is dedicated to fighting child sexual exploitation. You can find a full list of our safety partners [here](#).

Abusive behavior

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing or hoping serious harm on a person or group of people

We do not tolerate content that wishes, hopes or expresses a desire for death, serious bodily harm or fatal disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone's body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy (20200331*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you".

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

Hoping that someone dies as a result of a serious disease, e.g., “I hope you get cancer and die.”

Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”

Saying that a group of individuals deserve serious physical injury, e.g., “If this group of protesters don’t shut up, they deserve to be shot.”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

media that depicts victims of the Holocaust; media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists”.

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion, age, disability, or serious disease.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

symbols historically associated with hate groups, e.g., the Nazi swastika;

images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can

[submit an appeal.](#)

Suicide and Self-harm Policy

Overview

You may not promote or encourage suicide or self-harm.

At Twitter, we recognize that suicide and self-harm are significant social & public health challenges that require collaboration between all stakeholders – public, private, and civil society – and that we have a role and responsibility to help people access and receive support when they need it.

When developing this policy, we consulted extensively with experts to ensure that people who have engaged in self-harm or experienced suicidal thoughts can share their personal experiences. We also recognized the need to protect people from the potential harm caused by exposure to content that could promote or encourage self-harm – intentionally or inadvertently. That's why our policy prohibits content that promotes or encourages self-harming behaviors and provides support to those undergoing experiences with self-harm or suicidal thoughts.

What is in violation of this policy?

Under this policy, you can't promote, or otherwise encourage, suicide or self-harm. We define promotion and encouragement to include statements such as "the most effective", "the easiest", "the best", "the most successful", "you should", "why don't you". Violations of this policy can occur via Tweets, images or videos, including live video.

We define suicide to be the act of taking one's own life. We define self-harm to include:

- self-inflicted physical injuries e.g., cutting; and
- eating disorders e.g., bulimia, anorexia.

Violations of this policy include, but are not limited to:

- encouraging someone to physically harm or kill themselves;
- asking others for encouragement to engage in self-harm or suicide, including seeking partners for group suicides or suicide games; and
- sharing information, strategies, methods or instructions that would assist people to engage in self-harm and suicide.

What is not a violation of this policy?

Some examples of behavior that are not considered a violation of this policy include:

- sharing personal stories and experiences related to self-harm or suicide;
- sharing coping mechanisms and resources for addressing self-harm or suicidal thoughts; and
- discussions that are focused on research, advocacy, and education related to self-harm or suicide prevention.

Note: people can share their personal experiences, but should avoid sharing detailed information about specific strategies or methods related to self-harm, as this could inadvertently encourage this behavior.

Who can report violations of this policy?

Anyone can report content that may encourage or promote suicide or self-harm via our in-app reporting or our [specialized reporting form](#). These reports are routed to a dedicated team who evaluate each case individually.

Note: if we receive a report that someone has expressed an intention to engage in self-harm or suicide, we will contact them directly, encourage them to seek support, and provide information about dedicated online and hotline resources. We may also work with law enforcement officials where appropriate, for example, if we receive a valid emergency disclosure request as defined in our [Law Enforcement Guidelines](#).

How to report violations of this policy

To ensure that we handle reports sensitively, our in-app reporting provides separate options for people who may be expressing an intention to harm themselves, and content that is encouraging or promoting self-harm or suicide.

Expressing intentions of self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Reporting form

You can also report this content for review via [our dedicated report form](#).

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Encouraging self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

What happens if you violate this policy?

Our enforcement approach depends on the type of content being shared, whether or not the reported account is encouraging or promoting self-harm or suicide, and the account's previous history of violations.

If you violate this policy by sharing content that intentionally encourages others to harm themselves, ask others to encourage you to harm yourself, or share detailed information or instructions related to self-harm or suicide methods, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy, or if your account is dedicated to promoting or encouraging self-harm or suicide, your account will be permanently suspended. If cases include images or videos related to self-harm or suicide, we will also evaluate this content under [our sensitive media policy](#). If you believe that your account was suspended in error, [you can submit an appeal](#).

We may also take steps to prevent the spread of instructional material hosted on third-party websites by [marking such links as unsafe](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Learn more](#) about how you can support someone experiencing thoughts of self-harm or suicide.

Visit our Safety Center for a list of local [mental health resources](#) and read [our blog on suicide prevention](#) to learn more about our work.

Sensitive media policy

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to

be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and

- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;
 - descriptions of physical appearance;
 - gossip, rumours, accusations, and allegations; and
 - screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy

Overview

March 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select **Other**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you

violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy

Overview

March 2019

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are; and
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting.

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:

- posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
- engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate; and
- coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in “follow trains,” “decks,” and “Retweet for Retweet” behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- “follow churn” – following and then unfollowing large numbers of accounts in an effort to inflate one's own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account's followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person's browser (malware) or computer or to compromise a person's privacy (phishing); and

- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don't violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

- Select **Report Tweet** from the icon.
- Select **It's suspicious or spam**.
- Select the option that best tells us how the Tweet is suspicious or spreading spam.
- Submit your report.

Desktop

You can report this content via desktop as follows:

- Select **Report Tweet** from the icon.
- Select **It's suspicious or spam**.
- Select the option that best tells us how the Tweet is suspicious or spreading spam.
- Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Blacklisting URLs

We blacklist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [hacked materials policy](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Civic integrity policy

Overview

October 2020

You may not use Twitter's services for the purpose of manipulating or interfering in elections or other civic processes. This includes posting or sharing content that may suppress participation or mislead people about when, where, or how to participate in a civic process. In addition, we may label and reduce the visibility of Tweets containing false or misleading information about civic processes in order to provide additional context.

The public conversation occurring on Twitter is never more important than during elections and other civic events. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt civic processes, including through the distribution of false or misleading information about the procedures or circumstances around participation in a civic process. In instances where misleading information does not seek to directly manipulate or disrupt civic processes, but leads to confusion on our service, we may label the Tweets to give additional context.

What is a civic process?

Twitter considers civic processes to be events or procedures mandated, organized, and conducted by the governing and/or electoral body of a country, state, region, district, or municipality to address a matter of common concern through public participation. Some examples of civic processes may include:

- Political elections
- Censuses
- Major referenda and ballot initiatives

What is in violation of this policy?

This policy addresses 4 categories of misleading behavior and content:

Misleading information about how to participate

We will label or remove false or misleading information about how to participate in an election or other civic process. This includes but is not limited to:

- misleading information about procedures to participate in a civic process (for example, that you can vote by Tweet, text message, email, or phone call in jurisdictions where these are not a possibility);
- misleading information about requirements for participation, including identification or citizenship requirements;
- misleading claims that cause confusion about the established laws, regulations, procedures, and methods of a civic process, or about the actions of officials or entities executing those civic processes; and
- misleading statements or information about the official, announced date or time of a civic process.

Suppression and intimidation

We will label or remove false or misleading information intended to intimidate or dissuade people from participating in an election or other civic process. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to voting in an election, polling places, or collecting census information;
- misleading claims about long lines, equipment problems, or other disruptions at voting locations during election periods;
- misleading claims about process procedures or techniques which could dissuade people from participating; and
- threats regarding voting locations or other key places or events (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

Misleading information about outcomes

We will label or remove false or misleading information intended to undermine public confidence in an election or other civic process. This includes but is not limited to:

- disputed claims that could undermine faith in the process itself, such as unverified information about election rigging, ballot tampering, vote tallying, or certification of election results; and
- misleading claims about the results or outcome of a civic process which calls for or could lead to interference with the implementation of the results of the process, e.g. claiming victory before election

results have been certified, inciting unlawful conduct to prevent the procedural or practical implementation of election results (note that our violent threats policy may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account policy](#).

What is not a violation of this policy?

Not all false or untrue information about politics or civic processes constitutes manipulation or interference. In the absence of other policy violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected or appointed official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information;
- voting and audience participation for competitions, game shows, or other entertainment purposes; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election or civic process. Therefore, we enable reporting of false or misleading information about civic processes in advance of major events, for people located in the relevant countries and locations. We also work with select government and civil society partners in these countries to provide additional channels for reporting and expedited review.

For civic processes with multiple stages or parts, such as primary elections or lengthy campaigns, reporting will be enabled leading up to the first officially-sanctioned event associated with the civic process.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country at the relevant time, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in civic processes.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in a civic process.

4. Submit your report.

What happens if you violate this policy?

The consequences for violating our civic integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Labeling

In circumstances where we do not remove content which violates this policy, we may provide additional context on Tweets sharing the content where they appear on Twitter. This means we may:

- Apply a label and/or warning message to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended; and/or
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or relevant Twitter policies.

In most cases, we will take all of the above actions on Tweets we label. In some instances, we'll also turn off your ability to reply, Retweet, or like the Tweet. We prioritize producing Twitter Moments in cases where misleading content on Twitter is gaining significant attention and has caused public confusion on our service.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Impersonation policy

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Synthetic and manipulated media policy

Overview

You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context.

You should be able to find reliable information on Twitter. That means understanding whether the content you see is real or fabricated and having the ability to find more context about what you see on Twitter. Therefore, we may label Tweets that include media (videos, audio, and images) that have been deceptively altered or fabricated. In addition, you may not share deceptively altered media on Twitter in ways that mislead or deceive people about the media's authenticity where threats to physical safety or other serious harm may result.

We use the following criteria as we consider Tweets and media for labeling or removal under this policy as part of our ongoing work to enforce our rules and ensure healthy and safe conversation on Twitter (additional information is available below):

Is the content significantly and deceptively altered or fabricated?	Is the content shared in a deceptive manner?	Is the content likely to impact public safety or cause serious harm?	
✓	✗	✗	Content may be labeled.
✗	✓	✗	Content may be labeled.
✓	✗	✓	Content is likely to be labeled, or may be removed.*
✓	✓	✗	Content is likely to be labeled.
✓	✓	✓	Content is likely to be removed.

1. Is the content synthetic or manipulated?

In order for content to be labeled or removed under this policy, we must have reason to believe that media, or the context in which media are presented, are significantly and deceptively altered or manipulated. Synthetic and manipulated media take many different forms and people can employ a wide range of technologies to produce these media. In assessing whether media have been significantly and deceptively altered or fabricated, some of the factors we consider include:

- whether the content has been substantially edited in a manner that fundamentally alters its composition, sequence, timing, or framing;
- any visual or auditory information (such as new video frames, overdubbed audio, or modified subtitles) that has been added or removed; and
- whether media depicting a real person have been fabricated or simulated

We are most likely to take action (either labeling or removal, as described below) on more significant forms of alteration, such as wholly synthetic audio or video or content that has been doctored (spliced and reordered, slowed down) to change its meaning. Subtler forms of manipulated media, such as isolative editing, omission of context, or presentation with false context, may be labeled or removed on a case-by-case basis.

We will not take action to label or remove media that have been edited in ways that do not fundamentally alter their meaning, such as retouched photos or color-corrected videos.

In order to determine if media have been significantly and deceptively altered or fabricated, we may use our own technology or receive reports through partnerships with third parties. In situations where we are unable to reliably determine if media have been altered or fabricated, we may not take action to label or remove them.

2. Is the content shared in a deceptive manner?

We also consider whether the context in which media are shared could result in confusion or misunderstanding or suggests a deliberate intent to deceive people about the nature or origin of the content, for example by falsely claiming that it depicts reality. We assess the context provided alongside media to see whether it makes clear that the media have been altered or fabricated. Some of the types of context we assess in order to make this determination include:

- The text of the Tweet accompanying or within media
- Metadata associated with media
- Information on the profile of the account sharing media
- Websites linked in the Tweet, or in the profile of the account sharing media

3. Is the content likely to impact public safety or cause serious harm?

Tweets that share synthetic and manipulated media are subject to removal under this policy if they are likely to cause serious harm. Some specific harms we consider include:

- Threats to the physical safety of a person or group
- Risk of mass violence or widespread civil unrest
- Threats to the privacy or ability of a person or group to freely express themselves or participate in civic events, such as:
 - Stalking or unwanted and obsessive attention
 - Targeted content that includes tropes, epithets, or material that aims to silence someone
 - Voter suppression or intimidation

While we have other rules also intended to address these forms of harm, including our policies on violent threats, election integrity, and hateful conduct, we will err toward removal in borderline cases that might otherwise not violate existing rules for Tweets that include synthetic or manipulated media.

We also consider the time frame within which the content may be likely to impact public safety or cause serious harm, and are more likely to remove content under this policy if we find that immediate harms are likely to result from the content's presence on Twitter.

Note: We may also take action on synthetic and manipulated content under our [non-consensual nudity policy](#) (such as pornographic media altered to insert the faces of people not actually involved) or other parts of the Twitter Rules.

Labeling and removal

In most cases, if we have reason to believe that media shared in a Tweet have been significantly and deceptively altered or fabricated, we will provide additional context on Tweets sharing the media where they appear on Twitter. This means we may:

- Apply a label to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended; and/or
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or landing page.

In most cases, we will take all of the above actions on Tweets we label.

Media that meet **all three of the criteria** defined above—i.e. that are synthetic or manipulated, shared in a deceptive manner, and is likely to cause harm—may not be shared on Twitter and are subject to removal. Accounts engaging in repeated or severe violations of this policy may be permanently suspended.

** Other parts of the Twitter Rules apply and may lead to the removal of the content, particularly where there is high likelihood of severe harm, such as a threat to someone's life or physical safety.*

Trademark policy

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Copyright policy

What types of copyright complaints does Twitter respond to?

Twitter responds to copyright complaints submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review [Twitter's trademark policy](#). If you are concerned about a parody, newsfeed, commentary, or fan account, please see the relevant policy [here](#). These are generally not copyright issues.

Am I a copyright holder? How do I know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

What to consider before submitting a copyright complaint

Before submitting a copyright complaint to us, please consider whether or not the use could be considered [fair use](#).

If you have considered fair use, and you still wish to continue with a copyright complaint, you may want to first reach out to the user in question to see if you can resolve the matter directly with the user. You can reply to the user's Tweet or send the user a Direct Message and ask for them to remove your copyrighted content without having to contact Twitter.

Prior to submitting a formal complaint with Twitter, please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

Note: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph. If you're unsure whether or not you own the copyrights to a work, or if you're infringing upon someone else's work, please consult an attorney or another advisor.

What information do you need to process a copyright complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the complaint is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How do I file a copyright complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and [filing a copyright complaint](#). If you are logged in to [twitter.com](#), you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a full copy of your notice (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a copyright complaint.

How are claims processed?

We process copyright complaints in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. If you do not receive a ticket confirmation that means we did not receive your

complaint and you should re-submit your complaint. However, please note, submitting duplicate copyright complaints will result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) and provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice. We will also forward a redacted copy of the complaint to [Lumen](#), with your personal information removed.

What information gets forwarded to the reported user(s)?

If we remove or disable access to the materials reported in a copyright complaint, the reported user(s) will receive a copy of the complaint, including the reporter's full name, email, street address, and any other information included in the complaint.

If you are uncomfortable sharing your contact information with the reported user(s), you may wish to consider appointing an agent to submit your DMCA notice on your behalf. Your agent will be required to submit the DMCA notice with valid contact information, and identify you as the content owner that they are representing.

What happens next?

Twitter's response to copyright complaints may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a copyright complaint, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the complaint, along with instructions for filing a counter-notice.

If you've not yet received a copy of the copyright complaint regarding the content removed from your account, please respond to the support ticket we sent you. If your complaint involves removal of content from Periscope, please see this help center [article](#) regarding how to request a copy of your complaint.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

My content was removed from Twitter

Why did I receive a copyright complaint?

If you receive a copyright complaint, it means that access to the content described in the complaint has been restricted. Please take the time to read through our correspondence to you, which includes information on the complaint we received as well as instructions on how to file a counter-notice. Please ensure that you are monitoring the email address associated with your Twitter account.

Tip: Removing the material reported in a copyright complaint will not resolve that complaint.

What if I want to contest the takedown?

If you believe that the materials reported in the copyright complaints were misidentified or removed in error, you may send us a counter-notification(s). A counter-notice is a request for Twitter to reinstate the removed material, and it has legal consequences. Alternatively, you may be able to seek a retraction of the copyright complaint from the reporter.

How do I seek a retraction?

The DMCA complaint you received includes the contact information of the reporter. You may want to reach out and ask them to retract their notice. The reporter can send retractions to copyright@twitter.com, and should include: (1) identification of the material that was disabled, and (2) a statement that the reporter would like to retract their DMCA notice. This is the fastest and most efficient means of resolving an unresolved copyright complaint. A retraction is at the sole discretion of the original reporter.

When should I file a counter-notice?

A counter-notice is a request for Twitter to reinstate the removed material, and is the start of a legal process that has legal consequences. For example, submitting a counter notice indicates that you consent to the jurisdiction of a U.S. Federal court and that you consent to the disclosure of your personal information to the reporter and Lumen website.

With these considerations in mind, you may file a counter-notice if you believe that this material was misidentified, or you have a good faith belief that the material should not have been removed. If you're unsure whether or not you should file a counter-notice, you may want to consult with an attorney.

Tip: Re-posting material removed in response to a copyright complaint may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notice rather than re-posting the material.

What information do you need to process a counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What happens after I submit a counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. This means that the contact information that is submitted in your counter-notice will be shared to the person who filed the original notice.

If the copyright owner disagrees that the content was removed in error or misidentification, they may pursue legal action against you. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a copyright complaint or counter-notice is serious business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

What happens if my account receives multiple copyright complaints?

If multiple copyright complaints are received Twitter may lock accounts or take other actions to warn repeat violators. These warnings may vary across Twitter's services. Under appropriate circumstances we may suspend user accounts under our repeat infringer policy. However, we may take retractions and counter-notices into account when applying our repeat infringer policy.

Our approach to policy development and enforcement philosophy

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

[Public impact of the content](#): A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet

does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (20200204*)

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Synthetic and manipulated media: You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.

5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Violent organizations policy (20201026*)

Overview

October 2020

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

Other violent organizations

Other violent organizations are those that meet all of the below criteria:

- a collection of individuals with a shared purpose; and
- have systematically targeted civilians with violence.

We examine a group's activities both on and off Twitter to determine whether they satisfy the above criteria.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a violent organization;
- recruiting for a violent organization;
- providing or distributing services (e.g., financial, media/propaganda) to further a violent organization's stated goals; and
- using the insignia or symbol of violent organizations to promote them or indicate affiliation or support.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to state or governmental organizations.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.

6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wishing for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy (20201012*)

Overview

October 2020

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives;
- identifying alleged victims of childhood sexual exploitation by name or image; and
- promoting or normalizing sexual attraction to minors as a form of identity or sexual orientation.

What is not a violation of this policy?

- Discussions related to child sexual exploitation are permitted, provided they don't normalise, promote or glorify child sexual exploitation **in any way**;
- Conversations about help-seeking behaviour of individuals who may struggle with an attraction to minors;
- Advocacy against illegal or harmful activity involving minors provided there is no sharing of, or linking to any material featuring child sexual exploitation; and
- Depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios (e.g., art works by internationally renowned artists that feature minors, news media reportage, media created and shared for scientific or educational purposes).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

Abusive behavior (20190607*)

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing or hoping serious harm on a person or group of people

We do not tolerate content that wishes, hopes or expresses a desire for death, serious bodily harm or fatal disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone's body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy (20201203)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and have a higher impact on those targeted.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you".

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease, e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident, e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserve serious physical injury, e.g., "If this group of protesters don't shut up, they deserve to be shot."

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

- media that depicts victims of the Holocaust;
- media that depicts lynchings.

Inciting fear about a protected category

We prohibit targeting individuals and groups with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., "all [religious group] are terrorists".

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion, caste, age, disability, serious disease, national origin, race, or ethnicity.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

Suicide and Self-harm Policy **(20200805*)**

Overview

You may not promote or encourage suicide or self-harm.

At Twitter, we recognize that suicide and self-harm are significant social & public health challenges that require collaboration between all stakeholders – public, private, and civil society – and that we have a role and responsibility to help people access and receive support when they need it.

When developing this policy, we consulted extensively with experts to ensure that people who have engaged in self-harm or experienced suicidal thoughts can share their personal experiences. We also recognized the need to protect people from the potential harm caused by exposure to content that could promote or encourage self-harm – intentionally or inadvertently. That's why our policy prohibits content that promotes or encourages self-harming behaviors and provides support to those undergoing experiences with self-harm or suicidal thoughts.

What is in violation of this policy?

Under this policy, you can't promote, or otherwise encourage, suicide or self-harm. We define promotion and encouragement to include statements such as "the most effective", "the easiest", "the best", "the most successful", "you should", "why don't you". Violations of this policy can occur via Tweets, images or videos, including live video.

We define suicide to be the act of taking one's own life. We define self-harm to include:

- self-inflicted physical injuries e.g., cutting; and
- eating disorders e.g., bulimia, anorexia.

Violations of this policy include, but are not limited to:

- encouraging someone to physically harm or kill themselves;
- asking others for encouragement to engage in self-harm or suicide, including seeking partners for group suicides or suicide games; and
- sharing information, strategies, methods or instructions that would assist people to engage in self-harm and suicide.

What is not a violation of this policy?

Some examples of behavior that are not considered a violation of this policy include:

- sharing personal stories and experiences related to self-harm or suicide;
- sharing coping mechanisms and resources for addressing self-harm or suicidal thoughts; and
- discussions that are focused on research, advocacy, and education related to self-harm or suicide prevention.

Note: people can share their personal experiences, but should avoid sharing detailed information about specific strategies or methods related to self-harm, as this could inadvertently encourage this behavior.

Who can report violations of this policy?

Anyone can report content that may encourage or promote suicide or self-harm via our in-app reporting or our [specialized reporting form](#). These reports are routed to a dedicated team who evaluate each case individually.

Note: if we receive a report that someone has expressed an intention to engage in self-harm or suicide, we will contact them directly, encourage them to seek support, and provide information about dedicated online and hotline resources. We may also work with law enforcement officials where appropriate, for example, if we receive a valid emergency disclosure request as defined in our [Law Enforcement Guidelines](#).

How to report violations of this policy

To ensure that we handle reports sensitively, our in-app reporting provides separate options for people who may be expressing an intention to harm themselves, and content that is encouraging or promoting self-harm or suicide.

Expressing intentions of self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Reporting form

You can also report this content for review via [our dedicated report form](#).

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Encouraging self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"

4. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

What happens if you violate this policy?

Our enforcement approach depends on the type of content being shared, whether or not the reported account is encouraging or promoting self-harm or suicide, and the account's previous history of violations.

If you violate this policy by sharing content that intentionally encourages others to harm themselves, ask others to encourage you to harm yourself, or share detailed information or instructions related to self-harm or suicide methods, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy, or if your account is dedicated to promoting or encouraging self-harm or suicide, your account will be permanently suspended. If cases include images or videos related to self-harm or suicide, we will also evaluate this content under [our sensitive media policy](#). If you believe that your account was suspended in error, [you can submit an appeal](#).

We may also take steps to prevent the spread of instructional material hosted on third-party websites by [marking such links as unsafe](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Learn more](#) about how you can support someone experiencing thoughts of self-harm or suicide.

Visit our Safety Center for a list of local [mental health resources](#) and read [our blog on suicide prevention](#) to learn more about our work.

Sensitive media policy (20191204*)

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and

- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20190607*)

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;
 - descriptions of physical appearance;
 - gossip, rumours, accusations, and allegations; and
 - screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy **(20191217*)**

Overview

November 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy (20200918*)

Overview

September 2020

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are;
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting; and
- coordinated harmful activity that encourages or promotes behavior which violates the [Twitter Rules](#).

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;

- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts or by coordinating with others to violate the Twitter Rules. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate;
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account; and
 - coordinating with others to engage in or promote violations of the Twitter Rules, including violations of our [abusive behavior](#) policy.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in "follow trains," "decks," and "Retweet for Retweet" behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- "follow churn" – following and then unfollowing large numbers of accounts in an effort to inflate one's own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account's followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person's browser (malware) or computer or to compromise a person's privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally;
- coordinating with others to express ideas, viewpoints, support, or opposition towards a cause, provided such behavior does not result in violations of the Twitter Rules; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don't violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Denylisting URLs

We denylist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [financial scam policy](#), our [hacked materials policy](#), our approach to [coordinated harmful activity](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Civic integrity policy (20201028*)

Overview

October 2020

You may not use Twitter's services for the purpose of manipulating or interfering in elections or other civic processes. This includes posting or sharing content that may suppress participation or mislead people about when, where, or how to participate in a civic process. In addition, we may label and reduce the visibility of Tweets containing false or misleading information about civic processes in order to provide additional context.

The public conversation occurring on Twitter is never more important than during elections and other civic events. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt civic processes, including through the distribution of false or misleading information about the procedures or circumstances around participation in a civic process. In instances where misleading information does not seek to directly manipulate or disrupt civic processes, but leads to confusion on our service, we may label the Tweets to give additional context.

What is a civic process?

Twitter considers civic processes to be events or procedures mandated, organized, and conducted by the governing and/or electoral body of a country, state, region, district, or municipality to address a matter of common concern through public participation. Some examples of civic processes may include:

- Political elections
- Censuses
- Major referenda and ballot initiatives

What is in violation of this policy?

This policy addresses 4 categories of misleading behavior and content:

Misleading information about how to participate

We will label or remove false or misleading information about how to participate in an election or other civic process. This includes but is not limited to:

- misleading information about procedures to participate in a civic process (for example, that you can vote by Tweet, text message, email, or phone call in jurisdictions where these are not a possibility);
- misleading information about requirements for participation, including identification or citizenship requirements;
- misleading claims that cause confusion about the established laws, regulations, procedures, and methods of a civic process, or about the actions of officials or entities executing those civic processes; and
- misleading statements or information about the official, announced date or time of a civic process.

Suppression and intimidation

We will label or remove false or misleading information intended to intimidate or dissuade people from participating in an election or other civic process. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;

- misleading claims about police or law enforcement activity related to voting in an election, polling places, or collecting census information;
- misleading claims about long lines, equipment problems, or other disruptions at voting locations during election periods;
- misleading claims about process procedures or techniques which could dissuade people from participating; and
- threats regarding voting locations or other key places or events (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

Misleading information about outcomes

We will label or remove false or misleading information intended to undermine public confidence in an election or other civic process. This includes but is not limited to:

- disputed claims that could undermine faith in the process itself, such as unverified information about election rigging, ballot tampering, vote tallying, or certification of election results; and
- misleading claims about the results or outcome of a civic process which calls for or could lead to interference with the implementation of the results of the process, e.g. claiming victory before election results have been certified, inciting unlawful conduct to prevent the procedural or practical implementation of election results (note that our violent threats policy may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account policy](#).

What is not a violation of this policy?

Not all false or untrue information about politics or civic processes constitutes manipulation or interference. In the absence of other policy violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected or appointed official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information;
- voting and audience participation for competitions, game shows, or other entertainment purposes; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election or civic process. Therefore, we enable reporting of false or misleading information about civic processes in advance of major events, for people located in the relevant countries and locations. We also work with select government and civil society partners in these countries to provide additional channels for reporting and expedited review.

For civic processes with multiple stages or parts, such as primary elections or lengthy campaigns, reporting will be enabled leading up to the first officially-sanctioned event associated with the civic process.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country at the relevant time, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in civic processes.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in a civic process.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our civic integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Labeling

In circumstances where we do not remove content which violates this policy, we may provide additional context on Tweets sharing the content where they appear on Twitter. This means we may:

- Apply a label and/or warning message to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended; and/or
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or relevant Twitter policies.

In most cases, we will take all of the above actions on Tweets we label. In some instances, we'll also turn off your ability to reply, Retweet, or like the Tweet. We prioritize producing Twitter Moments in cases where misleading content on Twitter is gaining significant attention and has caused public confusion on our service.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Impersonation policy (20191001*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Synthetic and manipulated media policy **(20200207*)**

Overview

You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context.

You should be able to find reliable information on Twitter. That means understanding whether the content you see is real or fabricated and having the ability to find more context about what you see on Twitter. Therefore, we may label Tweets that include media (videos, audio, and images) that have been deceptively altered or fabricated. In addition, you may not share deceptively altered media on Twitter in ways that mislead or deceive people about the media's authenticity where threats to physical safety or other serious harm may result.

We use the following criteria as we consider Tweets and media for labeling or removal under this policy as part of our ongoing work to enforce our rules and ensure healthy and safe conversation on Twitter (additional information is available below):

Is the content significantly and deceptively altered or fabricated?	Is the content shared in a deceptive manner?	Is the content likely to impact public safety or cause serious harm?	
✓	✗	✗	Content may be labeled.
✗	✓	✗	Content may be labeled.
✓	✗	✓	Content is likely to be labeled, or may be removed.*
✓	✓	✗	Content is likely to be labeled.
✓	✓	✓	Content is likely to be removed.

1. Is the content synthetic or manipulated?

In order for content to be labeled or removed under this policy, we must have reason to believe that media, or the context in which media are presented, are significantly and deceptively altered or manipulated. Synthetic and manipulated media take many different forms and people can employ a wide range of technologies to produce these media. In assessing whether media have been significantly and deceptively altered or fabricated, some of the factors we consider include:

- whether the content has been substantially edited in a manner that fundamentally alters its composition, sequence, timing, or framing;
- any visual or auditory information (such as new video frames, overdubbed audio, or modified subtitles) that has been added or removed; and
- whether media depicting a real person have been fabricated or simulated

We are most likely to take action (either labeling or removal, as described below) on more significant forms of alteration, such as wholly synthetic audio or video or content that has been doctored (spliced and reordered, slowed down) to change its meaning. Subtler forms of manipulated media, such as isolative editing, omission of context, or presentation with false context, may be labeled or removed on a case-by-case basis.

We will not take action to label or remove media that have been edited in ways that do not fundamentally alter their meaning, such as retouched photos or color-corrected videos.

In order to determine if media have been significantly and deceptively altered or fabricated, we may use our own technology or receive reports through partnerships with third parties. In situations where we are unable to reliably determine if media have been altered or fabricated, we may not take action to label or remove them.

2. Is the content shared in a deceptive manner?

We also consider whether the context in which media are shared could result in confusion or misunderstanding or suggests a deliberate intent to deceive people about the nature or origin of the content, for example by falsely claiming that it depicts reality. We assess the context provided alongside media to see whether it makes clear that the media have been altered or fabricated. Some of the types of context we assess in order to make this determination include:

- The text of the Tweet accompanying or within media
- Metadata associated with media
- Information on the profile of the account sharing media
- Websites linked in the Tweet, or in the profile of the account sharing media

3. Is the content likely to impact public safety or cause serious harm?

Tweets that share synthetic and manipulated media are subject to removal under this policy if they are likely to cause serious harm. Some specific harms we consider include:

- Threats to the physical safety of a person or group
- Risk of mass violence or widespread civil unrest
- Threats to the privacy or ability of a person or group to freely express themselves or participate in civic events, such as:
 - Stalking or unwanted and obsessive attention
 - Targeted content that includes tropes, epithets, or material that aims to silence someone
 - Voter suppression or intimidation

While we have other rules also intended to address these forms of harm, including our policies on violent threats, election integrity, and hateful conduct, we will err toward removal in borderline cases that might otherwise not violate existing rules for Tweets that include synthetic or manipulated media.

We also consider the time frame within which the content may be likely to impact public safety or cause serious harm, and are more likely to remove content under this policy if we find that immediate harms are likely to result from the content's presence on Twitter.

Note: We may also take action on synthetic and manipulated content under our [non-consensual nudity policy](#) (such as pornographic media altered to insert the faces of people not actually involved) or other parts of the Twitter Rules.

Labeling and removal

In most cases, if we have reason to believe that media shared in a Tweet have been significantly and deceptively altered or fabricated, we will provide additional context on Tweets sharing the media where they appear on Twitter. This means we may:

- Apply a label to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended; and/or
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or landing page.

In most cases, we will take all of the above actions on Tweets we label.

Media that meet **all three of the criteria** defined above—i.e. that are synthetic or manipulated, shared in a deceptive manner, and is likely to cause harm—may not be shared on Twitter and are subject to removal. Accounts engaging in repeated or severe violations of this policy may be permanently suspended.

* Other parts of the Twitter Rules apply and may lead to the removal of the content, particularly where there is high likelihood of severe harm, such as a threat to someone's life or physical safety.

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Copyright policy (20200917*)

What types of copyright complaints does Twitter respond to?

Twitter responds to copyright complaints submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review [Twitter's trademark policy](#). If you are concerned about a parody, newsfeed, commentary, or fan account, please see the relevant policy [here](#). These are generally not copyright issues.

Am I a copyright holder? How do I know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

What to consider before submitting a copyright complaint

Before submitting a copyright complaint to us, please consider whether or not the use could be considered [fair use](#).

If you have considered fair use, and you still wish to continue with a copyright complaint, you may want to first reach out to the user in question to see if you can resolve the matter directly with the user. You can reply to the user's Tweet or send the user a Direct Message and ask for them to remove your copyrighted content without having to contact Twitter.

Prior to submitting a formal complaint with Twitter, please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

Note: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph. If you're unsure whether or not you own the copyrights to a work, or if you're infringing upon someone else's work, please consult an attorney or another advisor.

What information do you need to process a copyright complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;

4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the complaint is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How do I file a copyright complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and [filing a copyright complaint](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a full copy of your notice (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a copyright complaint.

How are claims processed?

We process copyright complaints in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. If you do not receive a ticket confirmation that means we did not receive your complaint and you should re-submit your complaint. However, please note, submitting duplicate copyright complaints will result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) and provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice. We will also forward a redacted copy of the complaint to [Lumen](#), with your personal information removed.

What information gets forwarded to the reported user(s)?

If we remove or disable access to the materials reported in a copyright complaint, the reported user(s) will receive a copy of the complaint, including the reporter's full name, email, street address, and any other information included in the complaint.

If you are uncomfortable sharing your contact information with the reported user(s), you may wish to consider appointing an agent to submit your DMCA notice on your behalf. Your agent will be required to submit the DMCA notice with valid contact information, and identify you as the content owner that they are representing.

What happens next?

Twitter's response to copyright complaints may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a copyright complaint, Twitter will make a

good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the complaint, along with instructions for filing a counter-notice.

If you've not yet received a copy of the copyright complaint regarding the content removed from your account, please respond to the support ticket we sent you. If your complaint involves removal of content from Periscope, please see this help center [article](#) regarding how to request a copy of your complaint.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from **@Username** has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

My content was removed from Twitter

Why did I receive a copyright complaint?

If you receive a copyright complaint, it means that access to the content described in the complaint has been restricted. Please take the time to read through our correspondence to you, which includes information on the complaint we received as well as instructions on how to file a counter-notice. Please ensure that you are monitoring the email address associated with your Twitter account.

Tip: Removing the material reported in a copyright complaint will not resolve that complaint.

What if I want to contest the takedown?

If you believe that the materials reported in the copyright complaints were misidentified or removed in error, you may send us a counter-notification(s). A counter-notice is a request for Twitter to reinstate the removed material, and it has legal consequences. Alternatively, you may be able to seek a retraction of the copyright complaint from the reporter.

How do I seek a retraction?

The DMCA complaint you received includes the contact information of the reporter. You may want to reach out and ask them to retract their notice. The reporter can send retractions to copyright@twitter.com, and should include: (1) identification of the material that was disabled, and (2) a statement that the reporter would like to retract their DMCA notice. This is the fastest and most efficient means of resolving an unresolved copyright complaint. A retraction is at the sole discretion of the original reporter.

When should I file a counter-notice?

A counter-notice is a request for Twitter to reinstate the removed material, and is the start of a legal process that has legal consequences. For example, submitting a counter notice indicates that you consent to the jurisdiction of

a U.S. Federal court and that you consent to the disclosure of your personal information to the reporter and Lumen website.

With these considerations in mind, you may file a counter-notice if you believe that this material was misidentified, or you have a good faith belief that the material should not have been removed. If you're unsure whether or not you should file a counter-notice, you may want to consult with an attorney.

Tip: Re-posting material removed in response to a copyright complaint may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notice rather than re-posting the material.

What information do you need to process a counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What happens after I submit a counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. This means that the contact information that is submitted in your counter-notice will be shared to the person who filed the original notice.

If the copyright owner disagrees that the content was removed in error or misidentification, they may pursue legal action against you. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a copyright complaint or counter-notice is serious business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

What happens if my account receives multiple copyright complaints?

If multiple copyright complaints are received Twitter may lock accounts or take other actions to warn repeat violators. These warnings may vary across Twitter's services. Under appropriate circumstances we may suspend user accounts under our repeat infringer policy. However, we may take retractions and counter-notices into account when applying our repeat infringer policy.

Our approach to policy development and enforcement philosophy

(20190520*)

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or

country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (20200204*)

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Synthetic and manipulated media: You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.

5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Violent organizations policy (20201026*)

Overview

October 2020

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

Other violent organizations

Other violent organizations are those that meet all of the below criteria:

- a collection of individuals with a shared purpose; and
- have systematically targeted civilians with violence.

We examine a group's activities both on and off Twitter to determine whether they satisfy the above criteria.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a violent organization;
- recruiting for a violent organization;
- providing or distributing services (e.g., financial, media/propaganda) to further a violent organization's stated goals; and
- using the insignia or symbol of violent organizations to promote them or indicate affiliation or support.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to state or governmental organizations.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.

6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wishing for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy (20201012*)

Overview

October 2020

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives;
- identifying alleged victims of childhood sexual exploitation by name or image; and
- promoting or normalizing sexual attraction to minors as a form of identity or sexual orientation.

What is not a violation of this policy?

- Discussions related to child sexual exploitation are permitted, provided they don't normalise, promote or glorify child sexual exploitation **in any way**;
- Conversations about help-seeking behaviour of individuals who may struggle with an attraction to minors;
- Advocacy against illegal or harmful activity involving minors provided there is no sharing of, or linking to any material featuring child sexual exploitation; and
- Depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios (e.g., art works by internationally renowned artists that feature minors, news media reportage, media created and shared for scientific or educational purposes).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

Abusive behavior (20190607*)

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing or hoping serious harm on a person or group of people

We do not tolerate content that wishes, hopes or expresses a desire for death, serious bodily harm or fatal disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone's body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy (20210120*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognise that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and more harmful.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you".

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, incites, or expresses a desire for death, serious and lasting bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

- Hoping that an entire protected category and/or individuals who may be members of that category dies as a result of a serious disease, e.g., "I hope all [nationality] get COVID and die."
- Wishing for someone to fall victim to a serious accident, e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserve serious physical injury, e.g., "If this group of protesters don't shut up, they deserve to be shot."
- Encouraging others to commit violence against an individual or a group based on their perceived membership in a protected category, e.g., "I'm in the mood to punch a [racial slur], who's with me?"

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

- media that depicts victims of the Holocaust;
- media that depicts lynchings.

Incitement against protected categories

We prohibit inciting behavior that targets individuals or groups of people belonging to protected categories. This includes content intended:

- to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., "all [religious group] are terrorists".
- to incite others to harass members of a protected category on or off platform, e.g., "I'm sick of these [religious group] thinking they are better than us, if any of you see someone wearing a [religious symbol of the religious group], grab it off them and post pics!"
- to incite discrimination in the form of denial of support to the economic enterprise of an individual or group because of their perceived membership in a protected category, e.g., "If you go to a [religious group] store, you are supporting terrorists, let's stop giving our money to these [religious slur]". This may not include content intended as political in nature, such as political commentary or content relating to boycotts or protests.

Note that content intended to incite violence against a protected category is prohibited under **Wishing, hoping, or calling for serious harm on a person or groups of people.**

We prohibit targeting individuals and groups with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., "all [religious group] are terrorists".

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion, caste, age, disability, serious disease, national origin, race, or ethnicity.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

Suicide and Self-harm Policy **(20200805*)**

Overview

You may not promote or encourage suicide or self-harm.

At Twitter, we recognize that suicide and self-harm are significant social & public health challenges that require collaboration between all stakeholders – public, private, and civil society – and that we have a role and responsibility to help people access and receive support when they need it.

When developing this policy, we consulted extensively with experts to ensure that people who have engaged in self-harm or experienced suicidal thoughts can share their personal experiences. We also recognized the need to protect people from the potential harm caused by exposure to content that could promote or encourage self-harm – intentionally or inadvertently. That's why our policy prohibits content that promotes or encourages self-harming behaviors and provides support to those undergoing experiences with self-harm or suicidal thoughts.

What is in violation of this policy?

Under this policy, you can't promote, or otherwise encourage, suicide or self-harm. We define promotion and encouragement to include statements such as "the most effective", "the easiest", "the best", "the most successful", "you should", "why don't you". Violations of this policy can occur via Tweets, images or videos, including live video.

We define suicide to be the act of taking one's own life. We define self-harm to include:

- self-inflicted physical injuries e.g., cutting; and
- eating disorders e.g., bulimia, anorexia.

Violations of this policy include, but are not limited to:

- encouraging someone to physically harm or kill themselves;
- asking others for encouragement to engage in self-harm or suicide, including seeking partners for group suicides or suicide games; and
- sharing information, strategies, methods or instructions that would assist people to engage in self-harm and suicide.

What is not a violation of this policy?

Some examples of behavior that are not considered a violation of this policy include:

- sharing personal stories and experiences related to self-harm or suicide;
- sharing coping mechanisms and resources for addressing self-harm or suicidal thoughts; and
- discussions that are focused on research, advocacy, and education related to self-harm or suicide prevention.

Note: people can share their personal experiences, but should avoid sharing detailed information about specific strategies or methods related to self-harm, as this could inadvertently encourage this behavior.

Who can report violations of this policy?

Anyone can report content that may encourage or promote suicide or self-harm via our in-app reporting or our [specialized reporting form](#). These reports are routed to a dedicated team who evaluate each case individually.

Note: if we receive a report that someone has expressed an intention to engage in self-harm or suicide, we will contact them directly, encourage them to seek support, and provide information about dedicated online and hotline resources. We may also work with law enforcement officials where appropriate, for example, if we receive a valid emergency disclosure request as defined in our [Law Enforcement Guidelines](#).

How to report violations of this policy

To ensure that we handle reports sensitively, our in-app reporting provides separate options for people who may be expressing an intention to harm themselves, and content that is encouraging or promoting self-harm or suicide.

Expressing intentions of self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"

3. Submit your report

Reporting form

You can also report this content for review via [our dedicated report form](#).

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Encouraging self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

What happens if you violate this policy?

Our enforcement approach depends on the type of content being shared, whether or not the reported account is encouraging or promoting self-harm or suicide, and the account's previous history of violations.

If you violate this policy by sharing content that intentionally encourages others to harm themselves, ask others to encourage you to harm yourself, or share detailed information or instructions related to self-harm or suicide methods, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy, or if your account is dedicated to promoting or encouraging self-harm or suicide, your account will be permanently suspended. If cases include images or videos related to self-harm or suicide, we will also evaluate this content under [our sensitive media policy](#). If you believe that your account was suspended in error, [you can submit an appeal](#).

We may also take steps to prevent the spread of instructional material hosted on third-party websites by [marking such links as unsafe](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Learn more](#) about how you can support someone experiencing thoughts of self-harm or suicide.

Visit our Safety Center for a list of local [mental health resources](#) and read [our blog on suicide prevention](#) to learn more about our work.

Sensitive media policy (20191204*)

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20190607*)

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;
 - descriptions of physical appearance;
 - gossip, rumours, accusations, and allegations; and
 - screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy **(20191217*)**

Overview

November 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy (20200918*)

Overview

September 2020

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are;

- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting; and
- coordinated harmful activity that encourages or promotes behavior which violates the [Twitter Rules](#).

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts or by coordinating with others to violate the Twitter Rules. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate;
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account; and
 - coordinating with others to engage in or promote violations of the Twitter Rules, including violations of our [abusive behavior](#) policy.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in "follow trains," "decks," and "Retweet for Retweet" behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;

- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- “follow churn” – following and then unfollowing large numbers of accounts in an effort to inflate one’s own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account’s followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person’s browser (malware) or computer or to compromise a person’s privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally;
- coordinating with others to express ideas, viewpoints, support, or opposition towards a cause, provided such behavior does not result in violations of the Twitter Rules; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don’t violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Denylisting URLs

We denylist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [financial scam policy](#), our [hacked materials policy](#), our approach to [coordinated harmful activity](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Civic integrity policy (20201028*)

Overview

October 2020

You may not use Twitter's services for the purpose of manipulating or interfering in elections or other civic processes. This includes posting or sharing content that may suppress participation or mislead people about when, where, or how to participate in a civic process. In addition, we may label and reduce the visibility of Tweets containing false or misleading information about civic processes in order to provide additional context.

The public conversation occurring on Twitter is never more important than during elections and other civic events. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt civic processes, including through the distribution of false or misleading information about the procedures or circumstances around participation in a civic process. In instances where misleading information does not seek to directly manipulate or disrupt civic processes, but leads to confusion on our service, we may label the Tweets to give additional context.

What is a civic process?

Twitter considers civic processes to be events or procedures mandated, organized, and conducted by the governing and/or electoral body of a country, state, region, district, or municipality to address a matter of common concern through public participation. Some examples of civic processes may include:

- Political elections
- Censuses
- Major referenda and ballot initiatives

What is in violation of this policy?

This policy addresses 4 categories of misleading behavior and content:

Misleading information about how to participate

We will label or remove false or misleading information about how to participate in an election or other civic process. This includes but is not limited to:

- misleading information about procedures to participate in a civic process (for example, that you can vote by Tweet, text message, email, or phone call in jurisdictions where these are not a possibility);
- misleading information about requirements for participation, including identification or citizenship requirements;
- misleading claims that cause confusion about the established laws, regulations, procedures, and methods of a civic process, or about the actions of officials or entities executing those civic processes; and
- misleading statements or information about the official, announced date or time of a civic process.

Suppression and intimidation

We will label or remove false or misleading information intended to intimidate or dissuade people from participating in an election or other civic process. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to voting in an election, polling places, or collecting census information;
- misleading claims about long lines, equipment problems, or other disruptions at voting locations during election periods;
- misleading claims about process procedures or techniques which could dissuade people from participating; and
- threats regarding voting locations or other key places or events (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

Misleading information about outcomes

We will label or remove false or misleading information intended to undermine public confidence in an election or other civic process. This includes but is not limited to:

- disputed claims that could undermine faith in the process itself, such as unverified information about election rigging, ballot tampering, vote tallying, or certification of election results; and
- misleading claims about the results or outcome of a civic process which calls for or could lead to interference with the implementation of the results of the process, e.g. claiming victory before election results have been certified, inciting unlawful conduct to prevent the procedural or practical implementation of election results (note that our violent threats policy may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account policy](#).

What is not a violation of this policy?

Not all false or untrue information about politics or civic processes constitutes manipulation or interference. In the absence of other policy violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected or appointed official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information;
- voting and audience participation for competitions, game shows, or other entertainment purposes; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election or civic process. Therefore, we enable reporting of false or misleading information about civic processes in advance of major events, for people located in the relevant countries and locations. We also work with select government and civil society partners in these countries to provide additional channels for reporting and expedited review.

For civic processes with multiple stages or parts, such as primary elections or lengthy campaigns, reporting will be enabled leading up to the first officially-sanctioned event associated with the civic process.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country at the relevant time, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in civic processes.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in a civic process.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our civic integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Labeling

In circumstances where we do not remove content which violates this policy, we may provide additional context on Tweets sharing the content where they appear on Twitter. This means we may:

- Apply a label and/or warning message to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;

- Reduce the visibility of the content on Twitter and/or prevent it from being recommended; and/or
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or relevant Twitter policies.

In most cases, we will take all of the above actions on Tweets we label. In some instances, we'll also turn off your ability to reply, Retweet, or like the Tweet. We prioritize producing Twitter Moments in cases where misleading content on Twitter is gaining significant attention and has caused public confusion on our service.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Impersonation policy (20191001*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody](#), [newsfeed](#), [commentary](#), and [fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Synthetic and manipulated media policy **(20200207*)**

Overview

You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context.

You should be able to find reliable information on Twitter. That means understanding whether the content you see is real or fabricated and having the ability to find more context about what you see on Twitter. Therefore, we may label Tweets that include media (videos, audio, and images) that have been deceptively altered or fabricated. In addition, you may not share deceptively altered media on Twitter in ways that mislead or deceive people about the media's authenticity where threats to physical safety or other serious harm may result.

We use the following criteria as we consider Tweets and media for labeling or removal under this policy as part of our ongoing work to enforce our rules and ensure healthy and safe conversation on Twitter (additional information is available below):

Is the content significantly and deceptively altered or fabricated?	Is the content shared in a deceptive manner?	Is the content likely to impact public safety or cause serious harm?	
✓	✗	✗	Content may be labeled.
✗	✓	✗	Content may be labeled.
✓	✗	✓	Content is likely to be labeled, or may be removed.*
✓	✓	✗	Content is likely to be labeled.
✓	✓	✓	Content is likely to be removed.

1. Is the content synthetic or manipulated?

In order for content to be labeled or removed under this policy, we must have reason to believe that media, or the context in which media are presented, are significantly and deceptively altered or manipulated. Synthetic and manipulated media take many different forms and people can employ a wide range of technologies to produce these media. In assessing whether media have been significantly and deceptively altered or fabricated, some of the factors we consider include:

- whether the content has been substantially edited in a manner that fundamentally alters its composition, sequence, timing, or framing;
- any visual or auditory information (such as new video frames, overdubbed audio, or modified subtitles) that has been added or removed; and
- whether media depicting a real person have been fabricated or simulated

We are most likely to take action (either labeling or removal, as described below) on more significant forms of alteration, such as wholly synthetic audio or video or content that has been doctored (spliced and reordered, slowed down) to change its meaning. Subtler forms of manipulated media, such as isolative editing, omission of context, or presentation with false context, may be labeled or removed on a case-by-case basis.

We will not take action to label or remove media that have been edited in ways that do not fundamentally alter their meaning, such as retouched photos or color-corrected videos.

In order to determine if media have been significantly and deceptively altered or fabricated, we may use our own technology or receive reports through partnerships with third parties. In situations where we are unable to reliably determine if media have been altered or fabricated, we may not take action to label or remove them.

2. Is the content shared in a deceptive manner?

We also consider whether the context in which media are shared could result in confusion or misunderstanding or suggests a deliberate intent to deceive people about the nature or origin of the content, for example by falsely claiming that it depicts reality. We assess the context provided alongside media to see whether it makes clear that the media have been altered or fabricated. Some of the types of context we assess in order to make this determination include:

- The text of the Tweet accompanying or within media
- Metadata associated with media
- Information on the profile of the account sharing media
- Websites linked in the Tweet, or in the profile of the account sharing media

3. Is the content likely to impact public safety or cause serious harm?

Tweets that share synthetic and manipulated media are subject to removal under this policy if they are likely to cause serious harm. Some specific harms we consider include:

- Threats to the physical safety of a person or group
- Risk of mass violence or widespread civil unrest
- Threats to the privacy or ability of a person or group to freely express themselves or participate in civic events, such as:
 - Stalking or unwanted and obsessive attention
 - Targeted content that includes tropes, epithets, or material that aims to silence someone
 - Voter suppression or intimidation

While we have other rules also intended to address these forms of harm, including our policies on violent threats, election integrity, and hateful conduct, we will err toward removal in borderline cases that might otherwise not violate existing rules for Tweets that include synthetic or manipulated media.

We also consider the time frame within which the content may be likely to impact public safety or cause serious harm, and are more likely to remove content under this policy if we find that immediate harms are likely to result from the content's presence on Twitter.

Note: We may also take action on synthetic and manipulated content under our [non-consensual nudity policy](#) (such as pornographic media altered to insert the faces of people not actually involved) or other parts of the Twitter Rules.

Labeling and removal

In most cases, if we have reason to believe that media shared in a Tweet have been significantly and deceptively altered or fabricated, we will provide additional context on Tweets sharing the media where they appear on Twitter. This means we may:

- Apply a label to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended; and/or
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or landing page.

In most cases, we will take all of the above actions on Tweets we label.

Media that meet **all three of the criteria** defined above—i.e. that are synthetic or manipulated, shared in a deceptive manner, and is likely to cause harm—may not be shared on Twitter and are subject to removal. Accounts engaging in repeated or severe violations of this policy may be permanently suspended.

* Other parts of the Twitter Rules apply and may lead to the removal of the content, particularly where there is high likelihood of severe harm, such as a threat to someone's life or physical safety.

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Copyright policy (20200917*)

What types of copyright complaints does Twitter respond to?

Twitter responds to copyright complaints submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review [Twitter's trademark policy](#). If you are concerned about a parody, newsfeed, commentary, or fan account, please see the relevant policy [here](#). These are generally not copyright issues.

Am I a copyright holder? How do I know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

What to consider before submitting a copyright complaint

Before submitting a copyright complaint to us, please consider whether or not the use could be considered [fair use](#).

If you have considered fair use, and you still wish to continue with a copyright complaint, you may want to first reach out to the user in question to see if you can resolve the matter directly with the user. You can reply to the user's Tweet or send the user a Direct Message and ask for them to remove your copyrighted content without having to contact Twitter.

Prior to submitting a formal complaint with Twitter, please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

Note: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph. If you're unsure whether or not you own the copyrights to a work, or if you're infringing upon someone else's work, please consult an attorney or another advisor.

What information do you need to process a copyright complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;

4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the complaint is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How do I file a copyright complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and [filing a copyright complaint](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a full copy of your notice (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a copyright complaint.

How are claims processed?

We process copyright complaints in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. If you do not receive a ticket confirmation that means we did not receive your complaint and you should re-submit your complaint. However, please note, submitting duplicate copyright complaints will result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) and provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice. We will also forward a redacted copy of the complaint to [Lumen](#), with your personal information removed.

What information gets forwarded to the reported user(s)?

If we remove or disable access to the materials reported in a copyright complaint, the reported user(s) will receive a copy of the complaint, including the reporter's full name, email, street address, and any other information included in the complaint.

If you are uncomfortable sharing your contact information with the reported user(s), you may wish to consider appointing an agent to submit your DMCA notice on your behalf. Your agent will be required to submit the DMCA notice with valid contact information, and identify you as the content owner that they are representing.

What happens next?

Twitter's response to copyright complaints may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a copyright complaint, Twitter will make a

good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the complaint, along with instructions for filing a counter-notice.

If you've not yet received a copy of the copyright complaint regarding the content removed from your account, please respond to the support ticket we sent you. If your complaint involves removal of content from Periscope, please see this help center [article](#) regarding how to request a copy of your complaint.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from **@Username** has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

My content was removed from Twitter

Why did I receive a copyright complaint?

If you receive a copyright complaint, it means that access to the content described in the complaint has been restricted. Please take the time to read through our correspondence to you, which includes information on the complaint we received as well as instructions on how to file a counter-notice. Please ensure that you are monitoring the email address associated with your Twitter account.

Tip: Removing the material reported in a copyright complaint will not resolve that complaint.

What if I want to contest the takedown?

If you believe that the materials reported in the copyright complaints were misidentified or removed in error, you may send us a counter-notification(s). A counter-notice is a request for Twitter to reinstate the removed material, and it has legal consequences. Alternatively, you may be able to seek a retraction of the copyright complaint from the reporter.

How do I seek a retraction?

The DMCA complaint you received includes the contact information of the reporter. You may want to reach out and ask them to retract their notice. The reporter can send retractions to copyright@twitter.com, and should include: (1) identification of the material that was disabled, and (2) a statement that the reporter would like to retract their DMCA notice. This is the fastest and most efficient means of resolving an unresolved copyright complaint. A retraction is at the sole discretion of the original reporter.

When should I file a counter-notice?

A counter-notice is a request for Twitter to reinstate the removed material, and is the start of a legal process that has legal consequences. For example, submitting a counter notice indicates that you consent to the jurisdiction of

a U.S. Federal court and that you consent to the disclosure of your personal information to the reporter and Lumen website.

With these considerations in mind, you may file a counter-notice if you believe that this material was misidentified, or you have a good faith belief that the material should not have been removed. If you're unsure whether or not you should file a counter-notice, you may want to consult with an attorney.

Tip: Re-posting material removed in response to a copyright complaint may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notice rather than re-posting the material.

What information do you need to process a counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What happens after I submit a counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. This means that the contact information that is submitted in your counter-notice will be shared to the person who filed the original notice.

If the copyright owner disagrees that the content was removed in error or misidentification, they may pursue legal action against you. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a copyright complaint or counter-notice is serious business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

What happens if my account receives multiple copyright complaints?

If multiple copyright complaints are received Twitter may lock accounts or take other actions to warn repeat violators. These warnings may vary across Twitter's services. Under appropriate circumstances we may suspend user accounts under our repeat infringer policy. However, we may take retractions and counter-notices into account when applying our repeat infringer policy.

Our approach to policy development and enforcement philosophy

(20190520*)

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or

country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (20200204*)

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Synthetic and manipulated media: You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.

5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Violent organizations policy (20201026*)

Overview

October 2020

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

Other violent organizations

Other violent organizations are those that meet all of the below criteria:

- a collection of individuals with a shared purpose; and
- have systematically targeted civilians with violence.

We examine a group's activities both on and off Twitter to determine whether they satisfy the above criteria.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a violent organization;
- recruiting for a violent organization;
- providing or distributing services (e.g., financial, media/propaganda) to further a violent organization's stated goals; and
- using the insignia or symbol of violent organizations to promote them or indicate affiliation or support.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to state or governmental organizations.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.

6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wishing for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy (20201012*)

Overview

October 2020

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives;
- identifying alleged victims of childhood sexual exploitation by name or image; and
- promoting or normalizing sexual attraction to minors as a form of identity or sexual orientation.

What is not a violation of this policy?

- Discussions related to child sexual exploitation are permitted, provided they don't normalise, promote or glorify child sexual exploitation **in any way**;
- Conversations about help-seeking behaviour of individuals who may struggle with an attraction to minors;
- Advocacy against illegal or harmful activity involving minors provided there is no sharing of, or linking to any material featuring child sexual exploitation; and
- Depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios (e.g., art works by internationally renowned artists that feature minors, news media reportage, media created and shared for scientific or educational purposes).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

Abusive behavior (20190607*)

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing or hoping serious harm on a person or group of people

We do not tolerate content that wishes, hopes or expresses a desire for death, serious bodily harm or fatal disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone's body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy (20210127*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognize that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and more harmful.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you."

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, incites, or expresses a desire for death, serious bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

- Hoping that an entire protected category and/or individuals who may be members of that category dies as a result of a serious disease, e.g., "I hope all [nationality] get COVID and die."
- Wishing for someone to fall victim to a serious accident, e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserve serious physical injury, e.g., "If this group of [slur] don't shut up, they deserve to be shot."
- Encouraging others to commit violence against an individual or a group based on their perceived membership in a protected category, e.g., "I'm in the mood to punch a [racial slur], who's with me?"

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

- media that depicts victims of the Holocaust;
- media that depicts lynchings.

Incitement against protected categories

We prohibit inciting behavior that targets individuals or groups of people belonging to protected categories. This includes content intended:

- to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., "all [religious group] are terrorists."
- to incite others to harass members of a protected category on or off platform, e.g., "I'm sick of these [religious group] thinking they are better than us, if any of you see someone wearing a [religious symbol of the religious group], grab it off them and post pics!"
- to incite others to discriminate in the form of denial of support to the economic enterprise of an individual or group because of their perceived membership in a protected category, e.g., "If you go to a [religious group] store, you are supporting those [slur], let's stop giving our money to these [religious slur]." This may not include content intended as political in nature, such as political commentary or content relating to boycotts or protests.

Note that content intended to incite violence against a protected category is prohibited under **Wishing, hoping, or calling for serious harm on a person or groups of people.**

We prohibit targeting individuals and groups with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., "all [religious group] are terrorists."

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion, caste, age, disability, serious disease, national origin, race, or ethnicity.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

Suicide and Self-harm Policy **(20200805*)**

Overview

You may not promote or encourage suicide or self-harm.

At Twitter, we recognize that suicide and self-harm are significant social & public health challenges that require collaboration between all stakeholders – public, private, and civil society – and that we have a role and responsibility to help people access and receive support when they need it.

When developing this policy, we consulted extensively with experts to ensure that people who have engaged in self-harm or experienced suicidal thoughts can share their personal experiences. We also recognized the need to protect people from the potential harm caused by exposure to content that could promote or encourage self-harm – intentionally or inadvertently. That's why our policy prohibits content that promotes or encourages self-harming behaviors and provides support to those undergoing experiences with self-harm or suicidal thoughts.

What is in violation of this policy?

Under this policy, you can't promote, or otherwise encourage, suicide or self-harm. We define promotion and encouragement to include statements such as "the most effective", "the easiest", "the best", "the most successful", "you should", "why don't you". Violations of this policy can occur via Tweets, images or videos, including live video.

We define suicide to be the act of taking one's own life. We define self-harm to include:

- self-inflicted physical injuries e.g., cutting; and
- eating disorders e.g., bulimia, anorexia.

Violations of this policy include, but are not limited to:

- encouraging someone to physically harm or kill themselves;
- asking others for encouragement to engage in self-harm or suicide, including seeking partners for group suicides or suicide games; and
- sharing information, strategies, methods or instructions that would assist people to engage in self-harm and suicide.

What is not a violation of this policy?

Some examples of behavior that are not considered a violation of this policy include:

- sharing personal stories and experiences related to self-harm or suicide;
- sharing coping mechanisms and resources for addressing self-harm or suicidal thoughts; and
- discussions that are focused on research, advocacy, and education related to self-harm or suicide prevention.

Note: people can share their personal experiences, but should avoid sharing detailed information about specific strategies or methods related to self-harm, as this could inadvertently encourage this behavior.

Who can report violations of this policy?

Anyone can report content that may encourage or promote suicide or self-harm via our in-app reporting or our [specialized reporting form](#). These reports are routed to a dedicated team who evaluate each case individually.

Note: if we receive a report that someone has expressed an intention to engage in self-harm or suicide, we will contact them directly, encourage them to seek support, and provide information about dedicated online and hotline resources. We may also work with law enforcement officials where appropriate, for example, if we receive a valid emergency disclosure request as defined in our [Law Enforcement Guidelines](#).

How to report violations of this policy

To ensure that we handle reports sensitively, our in-app reporting provides separate options for people who may be expressing an intention to harm themselves, and content that is encouraging or promoting self-harm or suicide.

Expressing intentions of self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"

3. Submit your report

Reporting form

You can also report this content for review via [our dedicated report form](#).

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Encouraging self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

What happens if you violate this policy?

Our enforcement approach depends on the type of content being shared, whether or not the reported account is encouraging or promoting self-harm or suicide, and the account's previous history of violations.

If you violate this policy by sharing content that intentionally encourages others to harm themselves, ask others to encourage you to harm yourself, or share detailed information or instructions related to self-harm or suicide methods, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy, or if your account is dedicated to promoting or encouraging self-harm or suicide, your account will be permanently suspended. If cases include images or videos related to self-harm or suicide, we will also evaluate this content under [our sensitive media policy](#). If you believe that your account was suspended in error, [you can submit an appeal](#).

We may also take steps to prevent the spread of instructional material hosted on third-party websites by [marking such links as unsafe](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Learn more](#) about how you can support someone experiencing thoughts of self-harm or suicide.

Visit our Safety Center for a list of local [mental health resources](#) and read [our blog on suicide prevention](#) to learn more about our work.

Sensitive media policy (20191204*)

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20190607*)

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;
 - descriptions of physical appearance;
 - gossip, rumours, accusations, and allegations; and
 - screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy **(20191217*)**

Overview

November 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy (20200918*)

Overview

September 2020

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are;

- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting; and
- coordinated harmful activity that encourages or promotes behavior which violates the [Twitter Rules](#).

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts or by coordinating with others to violate the Twitter Rules. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate;
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account; and
 - coordinating with others to engage in or promote violations of the Twitter Rules, including violations of our [abusive behavior](#) policy.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in "follow trains," "decks," and "Retweet for Retweet" behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;

- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- “follow churn” – following and then unfollowing large numbers of accounts in an effort to inflate one’s own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account’s followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person’s browser (malware) or computer or to compromise a person’s privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally;
- coordinating with others to express ideas, viewpoints, support, or opposition towards a cause, provided such behavior does not result in violations of the Twitter Rules; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don’t violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Denylisting URLs

We denylist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [financial scam policy](#), our [hacked materials policy](#), our approach to [coordinated harmful activity](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Civic integrity policy (20201028*)

Overview

October 2020

You may not use Twitter's services for the purpose of manipulating or interfering in elections or other civic processes. This includes posting or sharing content that may suppress participation or mislead people about when, where, or how to participate in a civic process. In addition, we may label and reduce the visibility of Tweets containing false or misleading information about civic processes in order to provide additional context.

The public conversation occurring on Twitter is never more important than during elections and other civic events. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt civic processes, including through the distribution of false or misleading information about the procedures or circumstances around participation in a civic process. In instances where misleading information does not seek to directly manipulate or disrupt civic processes, but leads to confusion on our service, we may label the Tweets to give additional context.

What is a civic process?

Twitter considers civic processes to be events or procedures mandated, organized, and conducted by the governing and/or electoral body of a country, state, region, district, or municipality to address a matter of common concern through public participation. Some examples of civic processes may include:

- Political elections
- Censuses
- Major referenda and ballot initiatives

What is in violation of this policy?

This policy addresses 4 categories of misleading behavior and content:

Misleading information about how to participate

We will label or remove false or misleading information about how to participate in an election or other civic process. This includes but is not limited to:

- misleading information about procedures to participate in a civic process (for example, that you can vote by Tweet, text message, email, or phone call in jurisdictions where these are not a possibility);
- misleading information about requirements for participation, including identification or citizenship requirements;
- misleading claims that cause confusion about the established laws, regulations, procedures, and methods of a civic process, or about the actions of officials or entities executing those civic processes; and
- misleading statements or information about the official, announced date or time of a civic process.

Suppression and intimidation

We will label or remove false or misleading information intended to intimidate or dissuade people from participating in an election or other civic process. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to voting in an election, polling places, or collecting census information;
- misleading claims about long lines, equipment problems, or other disruptions at voting locations during election periods;
- misleading claims about process procedures or techniques which could dissuade people from participating; and
- threats regarding voting locations or other key places or events (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

Misleading information about outcomes

We will label or remove false or misleading information intended to undermine public confidence in an election or other civic process. This includes but is not limited to:

- disputed claims that could undermine faith in the process itself, such as unverified information about election rigging, ballot tampering, vote tallying, or certification of election results; and
- misleading claims about the results or outcome of a civic process which calls for or could lead to interference with the implementation of the results of the process, e.g. claiming victory before election results have been certified, inciting unlawful conduct to prevent the procedural or practical implementation of election results (note that our violent threats policy may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account policy](#).

What is not a violation of this policy?

Not all false or untrue information about politics or civic processes constitutes manipulation or interference. In the absence of other policy violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected or appointed official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information;
- voting and audience participation for competitions, game shows, or other entertainment purposes; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election or civic process. Therefore, we enable reporting of false or misleading information about civic processes in advance of major events, for people located in the relevant countries and locations. We also work with select government and civil society partners in these countries to provide additional channels for reporting and expedited review.

For civic processes with multiple stages or parts, such as primary elections or lengthy campaigns, reporting will be enabled leading up to the first officially-sanctioned event associated with the civic process.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country at the relevant time, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in civic processes.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in a civic process.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our civic integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Labeling

In circumstances where we do not remove content which violates this policy, we may provide additional context on Tweets sharing the content where they appear on Twitter. This means we may:

- Apply a label and/or warning message to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;

- Reduce the visibility of the content on Twitter and/or prevent it from being recommended; and/or
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or relevant Twitter policies.

In most cases, we will take all of the above actions on Tweets we label. In some instances, we'll also turn off your ability to reply, Retweet, or like the Tweet. We prioritize producing Twitter Moments in cases where misleading content on Twitter is gaining significant attention and has caused public confusion on our service.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Impersonation policy (20191001*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody](#), [newsfeed](#), [commentary](#), and [fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Synthetic and manipulated media policy **(20200207*)**

Overview

You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context.

You should be able to find reliable information on Twitter. That means understanding whether the content you see is real or fabricated and having the ability to find more context about what you see on Twitter. Therefore, we may label Tweets that include media (videos, audio, and images) that have been deceptively altered or fabricated. In addition, you may not share deceptively altered media on Twitter in ways that mislead or deceive people about the media's authenticity where threats to physical safety or other serious harm may result.

We use the following criteria as we consider Tweets and media for labeling or removal under this policy as part of our ongoing work to enforce our rules and ensure healthy and safe conversation on Twitter (additional information is available below):

Is the content significantly and deceptively altered or fabricated?	Is the content shared in a deceptive manner?	Is the content likely to impact public safety or cause serious harm?	
✓	✗	✗	Content may be labeled.
✗	✓	✗	Content may be labeled.
✓	✗	✓	Content is likely to be labeled, or may be removed.*
✓	✓	✗	Content is likely to be labeled.
✓	✓	✓	Content is likely to be removed.

1. Is the content synthetic or manipulated?

In order for content to be labeled or removed under this policy, we must have reason to believe that media, or the context in which media are presented, are significantly and deceptively altered or manipulated. Synthetic and manipulated media take many different forms and people can employ a wide range of technologies to produce these media. In assessing whether media have been significantly and deceptively altered or fabricated, some of the factors we consider include:

- whether the content has been substantially edited in a manner that fundamentally alters its composition, sequence, timing, or framing;
- any visual or auditory information (such as new video frames, overdubbed audio, or modified subtitles) that has been added or removed; and
- whether media depicting a real person have been fabricated or simulated

We are most likely to take action (either labeling or removal, as described below) on more significant forms of alteration, such as wholly synthetic audio or video or content that has been doctored (spliced and reordered, slowed down) to change its meaning. Subtler forms of manipulated media, such as isolative editing, omission of context, or presentation with false context, may be labeled or removed on a case-by-case basis.

We will not take action to label or remove media that have been edited in ways that do not fundamentally alter their meaning, such as retouched photos or color-corrected videos.

In order to determine if media have been significantly and deceptively altered or fabricated, we may use our own technology or receive reports through partnerships with third parties. In situations where we are unable to reliably determine if media have been altered or fabricated, we may not take action to label or remove them.

2. Is the content shared in a deceptive manner?

We also consider whether the context in which media are shared could result in confusion or misunderstanding or suggests a deliberate intent to deceive people about the nature or origin of the content, for example by falsely claiming that it depicts reality. We assess the context provided alongside media to see whether it makes clear that the media have been altered or fabricated. Some of the types of context we assess in order to make this determination include:

- The text of the Tweet accompanying or within media
- Metadata associated with media
- Information on the profile of the account sharing media
- Websites linked in the Tweet, or in the profile of the account sharing media

3. Is the content likely to impact public safety or cause serious harm?

Tweets that share synthetic and manipulated media are subject to removal under this policy if they are likely to cause serious harm. Some specific harms we consider include:

- Threats to the physical safety of a person or group
- Risk of mass violence or widespread civil unrest
- Threats to the privacy or ability of a person or group to freely express themselves or participate in civic events, such as:
 - Stalking or unwanted and obsessive attention
 - Targeted content that includes tropes, epithets, or material that aims to silence someone
 - Voter suppression or intimidation

While we have other rules also intended to address these forms of harm, including our policies on violent threats, election integrity, and hateful conduct, we will err toward removal in borderline cases that might otherwise not violate existing rules for Tweets that include synthetic or manipulated media.

We also consider the time frame within which the content may be likely to impact public safety or cause serious harm, and are more likely to remove content under this policy if we find that immediate harms are likely to result from the content's presence on Twitter.

Note: We may also take action on synthetic and manipulated content under our [non-consensual nudity policy](#) (such as pornographic media altered to insert the faces of people not actually involved) or other parts of the Twitter Rules.

Labeling and removal

In most cases, if we have reason to believe that media shared in a Tweet have been significantly and deceptively altered or fabricated, we will provide additional context on Tweets sharing the media where they appear on Twitter. This means we may:

- Apply a label to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended; and/or
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or landing page.

In most cases, we will take all of the above actions on Tweets we label.

Media that meet **all three of the criteria** defined above—i.e. that are synthetic or manipulated, shared in a deceptive manner, and is likely to cause harm—may not be shared on Twitter and are subject to removal. Accounts engaging in repeated or severe violations of this policy may be permanently suspended.

* Other parts of the Twitter Rules apply and may lead to the removal of the content, particularly where there is high likelihood of severe harm, such as a threat to someone's life or physical safety.

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Copyright policy (20200917*)

What types of copyright complaints does Twitter respond to?

Twitter responds to copyright complaints submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review [Twitter's trademark policy](#). If you are concerned about a parody, newsfeed, commentary, or fan account, please see the relevant policy [here](#). These are generally not copyright issues.

Am I a copyright holder? How do I know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

What to consider before submitting a copyright complaint

Before submitting a copyright complaint to us, please consider whether or not the use could be considered [fair use](#).

If you have considered fair use, and you still wish to continue with a copyright complaint, you may want to first reach out to the user in question to see if you can resolve the matter directly with the user. You can reply to the user's Tweet or send the user a Direct Message and ask for them to remove your copyrighted content without having to contact Twitter.

Prior to submitting a formal complaint with Twitter, please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

Note: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph. If you're unsure whether or not you own the copyrights to a work, or if you're infringing upon someone else's work, please consult an attorney or another advisor.

What information do you need to process a copyright complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;

4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the complaint is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How do I file a copyright complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and [filing a copyright complaint](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a full copy of your notice (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a copyright complaint.

How are claims processed?

We process copyright complaints in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. If you do not receive a ticket confirmation that means we did not receive your complaint and you should re-submit your complaint. However, please note, submitting duplicate copyright complaints will result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) and provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice. We will also forward a redacted copy of the complaint to [Lumen](#), with your personal information removed.

What information gets forwarded to the reported user(s)?

If we remove or disable access to the materials reported in a copyright complaint, the reported user(s) will receive a copy of the complaint, including the reporter's full name, email, street address, and any other information included in the complaint.

If you are uncomfortable sharing your contact information with the reported user(s), you may wish to consider appointing an agent to submit your DMCA notice on your behalf. Your agent will be required to submit the DMCA notice with valid contact information, and identify you as the content owner that they are representing.

What happens next?

Twitter's response to copyright complaints may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a copyright complaint, Twitter will make a

good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the complaint, along with instructions for filing a counter-notice.

If you've not yet received a copy of the copyright complaint regarding the content removed from your account, please respond to the support ticket we sent you. If your complaint involves removal of content from Periscope, please see this help center [article](#) regarding how to request a copy of your complaint.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from **@Username** has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

My content was removed from Twitter

Why did I receive a copyright complaint?

If you receive a copyright complaint, it means that access to the content described in the complaint has been restricted. Please take the time to read through our correspondence to you, which includes information on the complaint we received as well as instructions on how to file a counter-notice. Please ensure that you are monitoring the email address associated with your Twitter account.

Tip: Removing the material reported in a copyright complaint will not resolve that complaint.

What if I want to contest the takedown?

If you believe that the materials reported in the copyright complaints were misidentified or removed in error, you may send us a counter-notification(s). A counter-notice is a request for Twitter to reinstate the removed material, and it has legal consequences. Alternatively, you may be able to seek a retraction of the copyright complaint from the reporter.

How do I seek a retraction?

The DMCA complaint you received includes the contact information of the reporter. You may want to reach out and ask them to retract their notice. The reporter can send retractions to copyright@twitter.com, and should include: (1) identification of the material that was disabled, and (2) a statement that the reporter would like to retract their DMCA notice. This is the fastest and most efficient means of resolving an unresolved copyright complaint. A retraction is at the sole discretion of the original reporter.

When should I file a counter-notice?

A counter-notice is a request for Twitter to reinstate the removed material, and is the start of a legal process that has legal consequences. For example, submitting a counter notice indicates that you consent to the jurisdiction of

a U.S. Federal court and that you consent to the disclosure of your personal information to the reporter and Lumen website.

With these considerations in mind, you may file a counter-notice if you believe that this material was misidentified, or you have a good faith belief that the material should not have been removed. If you're unsure whether or not you should file a counter-notice, you may want to consult with an attorney.

Tip: Re-posting material removed in response to a copyright complaint may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notice rather than re-posting the material.

What information do you need to process a counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What happens after I submit a counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. This means that the contact information that is submitted in your counter-notice will be shared to the person who filed the original notice.

If the copyright owner disagrees that the content was removed in error or misidentification, they may pursue legal action against you. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a copyright complaint or counter-notice is serious business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

What happens if my account receives multiple copyright complaints?

If multiple copyright complaints are received Twitter may lock accounts or take other actions to warn repeat violators. These warnings may vary across Twitter's services. Under appropriate circumstances we may suspend user accounts under our repeat infringer policy. However, we may take retractions and counter-notices into account when applying our repeat infringer policy.

Our approach to policy development and enforcement philosophy

(20190520*)

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or

country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (20200204*)

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Synthetic and manipulated media: You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.

5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Violent organizations policy (20201026*)

Overview

October 2020

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

Other violent organizations

Other violent organizations are those that meet all of the below criteria:

- a collection of individuals with a shared purpose; and
- have systematically targeted civilians with violence.

We examine a group's activities both on and off Twitter to determine whether they satisfy the above criteria.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a violent organization;
- recruiting for a violent organization;
- providing or distributing services (e.g., financial, media/propaganda) to further a violent organization's stated goals; and
- using the insignia or symbol of violent organizations to promote them or indicate affiliation or support.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to state or governmental organizations.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.

6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wishing for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy (20201012*)

Overview

October 2020

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives;
- identifying alleged victims of childhood sexual exploitation by name or image; and
- promoting or normalizing sexual attraction to minors as a form of identity or sexual orientation.

What is not a violation of this policy?

- Discussions related to child sexual exploitation are permitted, provided they don't normalise, promote or glorify child sexual exploitation **in any way**;
- Conversations about help-seeking behaviour of individuals who may struggle with an attraction to minors;
- Advocacy against illegal or harmful activity involving minors provided there is no sharing of, or linking to any material featuring child sexual exploitation; and
- Depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios (e.g., art works by internationally renowned artists that feature minors, news media reportage, media created and shared for scientific or educational purposes).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

Abusive behavior (20210527*)

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing, hoping, or calling for serious harm on a person or group of people

We do not tolerate content that wishes, hopes, promotes, incites, or expresses a desire for death, serious bodily harm or serious disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

About wishes of harm exceptions on Twitter

We recognize that conversations regarding certain individuals credibly accused of severe violence may prompt outrage and associated wishes of harm. In these limited cases, we will request the user to delete the Tweet without any risk of account penalty, strike, or suspension. Examples are, but not limited to:

- “I wish all rapists to die.”
- “Child abusers should be hanged.”

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone’s body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Denying mass casualty events took place

We prohibit content that denies that mass murder or other mass casualty events took place, where we can verify that the event occurred, and when the content is shared with abusive intent. This may include references to such an event as a “hoax” or claims that victims or survivors are fake or “actors.” It includes, but is not limited to, events like the Holocaust, school shootings, terrorist attacks, and natural disasters.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual’s previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy (20210127*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognize that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and more harmful.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you."

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, incites, or expresses a desire for death, serious bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

- Hoping that an entire protected category and/or individuals who may be members of that category dies as a result of a serious disease, e.g., "I hope all [nationality] get COVID and die."
- Wishing for someone to fall victim to a serious accident, e.g., "I wish that you would get run over by a car next time you run your mouth."

- Saying that a group of individuals deserve serious physical injury, e.g., "If this group of [slur] don't shut up, they deserve to be shot."
- Encouraging others to commit violence against an individual or a group based on their perceived membership in a protected category, e.g., "I'm in the mood to punch a [racial slur], who's with me?"

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

- media that depicts victims of the Holocaust;
- media that depicts lynchings.

Incitement against protected categories

We prohibit inciting behavior that targets individuals or groups of people belonging to protected categories. This includes content intended:

- to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., "all [religious group] are terrorists."
- to incite others to harass members of a protected category on or off platform, e.g., "I'm sick of these [religious group] thinking they are better than us, if any of you see someone wearing a [religious symbol of the religious group], grab it off them and post pics!"
- to incite others to discriminate in the form of denial of support to the economic enterprise of an individual or group because of their perceived membership in a protected category, e.g., "If you go to a [religious group] store, you are supporting those [slur], let's stop giving our money to these [religious slur]." This may not include content intended as political in nature, such as political commentary or content relating to boycotts or protests.

Note that content intended to incite violence against a protected category is prohibited under **Wishing, hoping, or calling for serious harm on a person or groups of people.**

We prohibit targeting individuals and groups with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., "all [religious group] are terrorists."

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion, caste, age, disability, serious disease, national origin, race, or ethnicity.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

Suicide and Self-harm Policy **(20200805*)**

Overview

You may not promote or encourage suicide or self-harm.

At Twitter, we recognize that suicide and self-harm are significant social & public health challenges that require collaboration between all stakeholders – public, private, and civil society – and that we have a role and responsibility to help people access and receive support when they need it.

When developing this policy, we consulted extensively with experts to ensure that people who have engaged in self-harm or experienced suicidal thoughts can share their personal experiences. We also recognized the need to protect people from the potential harm caused by exposure to content that could promote or encourage self-harm – intentionally or inadvertently. That's why our policy prohibits content that promotes or encourages self-harming behaviors and provides support to those undergoing experiences with self-harm or suicidal thoughts.

What is in violation of this policy?

Under this policy, you can't promote, or otherwise encourage, suicide or self-harm. We define promotion and encouragement to include statements such as "the most effective", "the easiest", "the best", "the most successful", "you should", "why don't you". Violations of this policy can occur via Tweets, images or videos, including live video.

We define suicide to be the act of taking one's own life. We define self-harm to include:

- self-inflicted physical injuries e.g., cutting; and
- eating disorders e.g., bulimia, anorexia.

Violations of this policy include, but are not limited to:

- encouraging someone to physically harm or kill themselves;
- asking others for encouragement to engage in self-harm or suicide, including seeking partners for group suicides or suicide games; and
- sharing information, strategies, methods or instructions that would assist people to engage in self-harm and suicide.

What is not a violation of this policy?

Some examples of behavior that are not considered a violation of this policy include:

- sharing personal stories and experiences related to self-harm or suicide;
- sharing coping mechanisms and resources for addressing self-harm or suicidal thoughts; and
- discussions that are focused on research, advocacy, and education related to self-harm or suicide prevention.

Note: people can share their personal experiences, but should avoid sharing detailed information about specific strategies or methods related to self-harm, as this could inadvertently encourage this behavior.

Who can report violations of this policy?

Anyone can report content that may encourage or promote suicide or self-harm via our in-app reporting or our [specialized reporting form](#). These reports are routed to a dedicated team who evaluate each case individually.

Note: if we receive a report that someone has expressed an intention to engage in self-harm or suicide, we will contact them directly, encourage them to seek support, and provide information about dedicated online and hotline resources. We may also work with law enforcement officials where appropriate, for example, if we receive a valid emergency disclosure request as defined in our [Law Enforcement Guidelines](#).

How to report violations of this policy

To ensure that we handle reports sensitively, our in-app reporting provides separate options for people who may be expressing an intention to harm themselves, and content that is encouraging or promoting self-harm or suicide.

Expressing intentions of self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Reporting form

You can also report this content for review via [our dedicated report form](#).

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Encouraging self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

What happens if you violate this policy?

Our enforcement approach depends on the type of content being shared, whether or not the reported account is encouraging or promoting self-harm or suicide, and the account's previous history of violations.

If you violate this policy by sharing content that intentionally encourages others to harm themselves, ask others to encourage you to harm yourself, or share detailed information or instructions related to self-harm or suicide methods, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy, or if your account is dedicated to promoting or encouraging self-harm or suicide, your account will be permanently suspended. If cases include images or videos related to self-harm or suicide, we will also evaluate this content under [our sensitive media policy](#). If you believe that your account was suspended in error, [you can submit an appeal](#).

We may also take steps to prevent the spread of instructional material hosted on third-party websites by [marking such links as unsafe](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Learn more](#) about how you can support someone experiencing thoughts of self-harm or suicide.

Visit our Safety Center for a list of local [mental health resources](#) and read [our blog on suicide prevention](#) to learn more about our work.

Sensitive media policy (20191204*)

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.

3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20190607*)

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their

information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;

- location information related to commercial property or places of business, where this information is publicly available;
- descriptions of physical appearance;
- gossip, rumours, accusations, and allegations; and
- screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy

(20191217*)

Overview

November 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy (20200918*)

Overview

September 2020

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are;
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting; and
- coordinated harmful activity that encourages or promotes behavior which violates the [Twitter Rules](#).

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts or by coordinating with others to violate the Twitter Rules. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and

- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate;
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account; and
 - coordinating with others to engage in or promote violations of the Twitter Rules, including violations of our [abusive behavior](#) policy.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in "follow trains," "decks," and "Retweet for Retweet" behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- "follow churn" – following and then unfollowing large numbers of accounts in an effort to inflate one's own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account's followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person's browser (malware) or computer or to compromise a person's privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally;
- coordinating with others to express ideas, viewpoints, support, or opposition towards a cause, provided such behavior does not result in violations of the Twitter Rules; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don't violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Denylisting URLs

We denylist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [financial scam policy](#), our [hacked materials policy](#), our approach to [coordinated harmful activity](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Civic integrity policy (20201028*)

Overview

October 2020

You may not use Twitter's services for the purpose of manipulating or interfering in elections or other civic processes. This includes posting or sharing content that may suppress participation or mislead people about when, where, or how to participate in a civic process. In addition, we may label and reduce the

visibility of Tweets containing false or misleading information about civic processes in order to provide additional context.

The public conversation occurring on Twitter is never more important than during elections and other civic events. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt civic processes, including through the distribution of false or misleading information about the procedures or circumstances around participation in a civic process. In instances where misleading information does not seek to directly manipulate or disrupt civic processes, but leads to confusion on our service, we may label the Tweets to give additional context.

What is a civic process?

Twitter considers civic processes to be events or procedures mandated, organized, and conducted by the governing and/or electoral body of a country, state, region, district, or municipality to address a matter of common concern through public participation. Some examples of civic processes may include:

- Political elections
- Censuses
- Major referenda and ballot initiatives

What is in violation of this policy?

This policy addresses 4 categories of misleading behavior and content:

Misleading information about how to participate

We will label or remove false or misleading information about how to participate in an election or other civic process. This includes but is not limited to:

- misleading information about procedures to participate in a civic process (for example, that you can vote by Tweet, text message, email, or phone call in jurisdictions where these are not a possibility);
- misleading information about requirements for participation, including identification or citizenship requirements;
- misleading claims that cause confusion about the established laws, regulations, procedures, and methods of a civic process, or about the actions of officials or entities executing those civic processes; and
- misleading statements or information about the official, announced date or time of a civic process.

Suppression and intimidation

We will label or remove false or misleading information intended to intimidate or dissuade people from participating in an election or other civic process. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to voting in an election, polling places, or collecting census information;
- misleading claims about long lines, equipment problems, or other disruptions at voting locations during election periods;
- misleading claims about process procedures or techniques which could dissuade people from participating; and
- threats regarding voting locations or other key places or events (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

Misleading information about outcomes

We will label or remove false or misleading information intended to undermine public confidence in an election or other civic process. This includes but is not limited to:

- disputed claims that could undermine faith in the process itself, such as unverified information about election rigging, ballot tampering, vote tallying, or certification of election results; and
- misleading claims about the results or outcome of a civic process which calls for or could lead to interference with the implementation of the results of the process, e.g. claiming victory before election results have been certified, inciting unlawful conduct to prevent the procedural or practical implementation of election results (note that our violent threats policy may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account policy](#).

What is not a violation of this policy?

Not all false or untrue information about politics or civic processes constitutes manipulation or interference. In the absence of other policy violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected or appointed official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information;
- voting and audience participation for competitions, game shows, or other entertainment purposes; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election or civic process. Therefore, we enable reporting of false or misleading information about civic processes in advance of major events, for people located in the relevant countries and locations. We also work with select government and civil society partners in these countries to provide additional channels for reporting and expedited review.

For civic processes with multiple stages or parts, such as primary elections or lengthy campaigns, reporting will be enabled leading up to the first officially-sanctioned event associated with the civic process.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country at the relevant time, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in civic processes.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in a civic process.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our civic integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Labeling

In circumstances where we do not remove content which violates this policy, we may provide additional context on Tweets sharing the content where they appear on Twitter. This means we may:

- Apply a label and/or warning message to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended; and/or
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or relevant Twitter policies.

In most cases, we will take all of the above actions on Tweets we label. In some instances, we'll also turn off your ability to reply, Retweet, or like the Tweet. We prioritize producing Twitter Moments in cases where misleading content on Twitter is gaining significant attention and has caused public confusion on our service.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Impersonation policy (20191001*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our range of enforcement options and our approach to policy development and enforcement.

Synthetic and manipulated media policy (20200207*)

Overview

You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context.

You should be able to find reliable information on Twitter. That means understanding whether the content you see is real or fabricated and having the ability to find more context about what you see on Twitter. Therefore, we may label Tweets that include media (videos, audio, and images) that have been deceptively altered or fabricated. In addition, you may not share deceptively altered media on Twitter in ways that mislead or deceive people about the media's authenticity where threats to physical safety or other serious harm may result.

We use the following criteria as we consider Tweets and media for labeling or removal under this policy as part of our ongoing work to enforce our rules and ensure healthy and safe conversation on Twitter (additional information is available below):

Is the content significantly and deceptively altered or fabricated?	Is the content shared in a deceptive manner?	Is the content likely to impact public safety or cause serious harm?	
✓	✗	✗	Content may be labeled.
✗	✓	✗	Content may be labeled.
✓	✗	✓	Content is likely to be labeled, or may be removed.*
✓	✓	✗	Content is likely to be labeled.
✓	✓	✓	Content is likely to be removed.

1. Is the content synthetic or manipulated?

In order for content to be labeled or removed under this policy, we must have reason to believe that media, or the context in which media are presented, are significantly and deceptively altered or manipulated. Synthetic and manipulated media take many different forms and people can employ a wide range of technologies to produce these media. In assessing whether media have been significantly and deceptively altered or fabricated, some of the factors we consider include:

- whether the content has been substantially edited in a manner that fundamentally alters its composition, sequence, timing, or framing;

- any visual or auditory information (such as new video frames, overdubbed audio, or modified subtitles) that has been added or removed; and
- whether media depicting a real person have been fabricated or simulated

We are most likely to take action (either labeling or removal, as described below) on more significant forms of alteration, such as wholly synthetic audio or video or content that has been doctored (spliced and reordered, slowed down) to change its meaning. Subtler forms of manipulated media, such as isolative editing, omission of context, or presentation with false context, may be labeled or removed on a case-by-case basis.

We will not take action to label or remove media that have been edited in ways that do not fundamentally alter their meaning, such as retouched photos or color-corrected videos.

In order to determine if media have been significantly and deceptively altered or fabricated, we may use our own technology or receive reports through partnerships with third parties. In situations where we are unable to reliably determine if media have been altered or fabricated, we may not take action to label or remove them.

2. Is the content shared in a deceptive manner?

We also consider whether the context in which media are shared could result in confusion or misunderstanding or suggests a deliberate intent to deceive people about the nature or origin of the content, for example by falsely claiming that it depicts reality. We assess the context provided alongside media to see whether it makes clear that the media have been altered or fabricated. Some of the types of context we assess in order to make this determination include:

- The text of the Tweet accompanying or within media
- Metadata associated with media
- Information on the profile of the account sharing media
- Websites linked in the Tweet, or in the profile of the account sharing media

3. Is the content likely to impact public safety or cause serious harm?

Tweets that share synthetic and manipulated media are subject to removal under this policy if they are likely to cause serious harm. Some specific harms we consider include:

- Threats to the physical safety of a person or group
- Risk of mass violence or widespread civil unrest
- Threats to the privacy or ability of a person or group to freely express themselves or participate in civic events, such as:
 - Stalking or unwanted and obsessive attention
 - Targeted content that includes tropes, epithets, or material that aims to silence someone
 - Voter suppression or intimidation

While we have other rules also intended to address these forms of harm, including our policies on violent threats, election integrity, and hateful conduct, we will err toward removal in borderline cases that might otherwise not violate existing rules for Tweets that include synthetic or manipulated media.

We also consider the time frame within which the content may be likely to impact public safety or cause serious harm, and are more likely to remove content under this policy if we find that immediate harms are likely to result from the content's presence on Twitter.

Note: We may also take action on synthetic and manipulated content under our [non-consensual nudity policy](#) (such as pornographic media altered to insert the faces of people not actually involved) or other parts of the Twitter Rules.

Labeling and removal

In most cases, if we have reason to believe that media shared in a Tweet have been significantly and deceptively altered or fabricated, we will provide additional context on Tweets sharing the media where they appear on Twitter. This means we may:

- Apply a label to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended; and/or
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or landing page.

In most cases, we will take all of the above actions on Tweets we label.

Media that meet **all three of the criteria** defined above—i.e. that are synthetic or manipulated, shared in a deceptive manner, and is likely to cause harm—may not be shared on Twitter and are subject to removal. Accounts engaging in repeated or severe violations of this policy may be permanently suspended.

** Other parts of the Twitter Rules apply and may lead to the removal of the content, particularly where there is high likelihood of severe harm, such as a threat to someone's life or physical safety.*

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Copyright policy (20200917*)

What types of copyright complaints does Twitter respond to?

Twitter responds to copyright complaints submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review [Twitter's trademark policy](#). If you are concerned about a parody, newsfeed, commentary, or fan account, please see the relevant policy [here](#). These are generally not copyright issues.

Am I a copyright holder? How do I know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

What to consider before submitting a copyright complaint

Before submitting a copyright complaint to us, please consider whether or not the use could be considered [fair use](#).

If you have considered fair use, and you still wish to continue with a copyright complaint, you may want to first reach out to the user in question to see if you can resolve the matter directly with the user. You can reply to the user's Tweet or send the user a Direct Message and ask for them to remove your copyrighted content without having to contact Twitter.

Prior to submitting a formal complaint with Twitter, please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

Note: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph. If you're unsure whether or not you own the copyrights to a work, or if you're infringing upon someone else's work, please consult an attorney or another advisor.

What information do you need to process a copyright complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the complaint is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How do I file a copyright complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and [filing a copyright complaint](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a full copy of your notice (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a copyright complaint.

How are claims processed?

We process copyright complaints in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. If you do not receive a ticket confirmation that means we did not receive your complaint and you should re-submit your complaint. However, please note, submitting duplicate copyright complaints will result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) and provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice. We will also forward a redacted copy of the complaint to [Lumen](#), with your personal information removed.

What information gets forwarded to the reported user(s)?

If we remove or disable access to the materials reported in a copyright complaint, the reported user(s) will receive a copy of the complaint, including the reporter's full name, email, street address, and any other information included in the complaint.

If you are uncomfortable sharing your contact information with the reported user(s), you may wish to consider appointing an agent to submit your DMCA notice on your behalf. Your agent will be required to submit the DMCA notice with valid contact information, and identify you as the content owner that they are representing.

What happens next?

Twitter's response to copyright complaints may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a copyright complaint, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the complaint, along with instructions for filing a counter-notice.

If you've not yet received a copy of the copyright complaint regarding the content removed from your account, please respond to the support ticket we sent you. If your complaint involves removal of content from Periscope, please see this help center [article](#) regarding how to request a copy of your complaint.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

My content was removed from Twitter

Why did I receive a copyright complaint?

If you receive a copyright complaint, it means that access to the content described in the complaint has been restricted. Please take the time to read through our correspondence to you, which includes information on the complaint we received as well as instructions on how to file a counter-notice. Please ensure that you are monitoring the email address associated with your Twitter account.

Tip: Removing the material reported in a copyright complaint will not resolve that complaint.

What if I want to contest the takedown?

If you believe that the materials reported in the copyright complaints were misidentified or removed in error, you may send us a counter-notification(s). A counter-notice is a request for Twitter to reinstate the removed material, and it has legal consequences. Alternatively, you may be able to seek a retraction of the copyright complaint from the reporter.

How do I seek a retraction?

The DMCA complaint you received includes the contact information of the reporter. You may want to reach out and ask them to retract their notice. The reporter can send retractions to copyright@twitter.com, and should include: (1) identification of the material that was disabled, and (2) a statement that the reporter would like to retract their DMCA notice. This is the fastest and most efficient means of resolving an unresolved copyright complaint. A retraction is at the sole discretion of the original reporter.

When should I file a counter-notice?

A counter-notice is a request for Twitter to reinstate the removed material, and is the start of a legal process that has legal consequences. For example, submitting a counter notice indicates that you consent to the jurisdiction of a U.S. Federal court and that you consent to the disclosure of your personal information to the reporter and Lumen website.

With these considerations in mind, you may file a counter-notice if you believe that this material was misidentified, or you have a good faith belief that the material should not have been removed. If you're unsure whether or not you should file a counter-notice, you may want to consult with an attorney.

Tip: Re-posting material removed in response to a copyright complaint may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notice rather than re-posting the material.

What information do you need to process a counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);

3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What happens after I submit a counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. This means that the contact information that is submitted in your counter-notice will be shared to the person who filed the original notice.

If the copyright owner disagrees that the content was removed in error or misidentification, they may pursue legal action against you. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a copyright complaint or counter-notice is serious business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

What happens if my account receives multiple copyright complaints?

If multiple copyright complaints are received Twitter may lock accounts or take other actions to warn repeat violators. These warnings may vary across Twitter's services. Under appropriate circumstances we may suspend user accounts under our repeat infringer policy. However, we may take retractions and counter-notices into account when applying our repeat infringer policy.

Our approach to policy development and enforcement philosophy **(20190520*)**

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting

violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (20200204*)

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Synthetic and manipulated media: You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.

5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Violent organizations policy (20201026*)

Overview

October 2020

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

Other violent organizations

Other violent organizations are those that meet all of the below criteria:

- a collection of individuals with a shared purpose; and
- have systematically targeted civilians with violence.

We examine a group's activities both on and off Twitter to determine whether they satisfy the above criteria.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a violent organization;
- recruiting for a violent organization;
- providing or distributing services (e.g., financial, media/propaganda) to further a violent organization's stated goals; and
- using the insignia or symbol of violent organizations to promote them or indicate affiliation or support.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to state or governmental organizations.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.

6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wishing for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy (20201012*)

Overview

October 2020

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives;
- identifying alleged victims of childhood sexual exploitation by name or image; and
- promoting or normalizing sexual attraction to minors as a form of identity or sexual orientation.

What is not a violation of this policy?

- Discussions related to child sexual exploitation are permitted, provided they don't normalise, promote or glorify child sexual exploitation **in any way**;
- Conversations about help-seeking behaviour of individuals who may struggle with an attraction to minors;
- Advocacy against illegal or harmful activity involving minors provided there is no sharing of, or linking to any material featuring child sexual exploitation; and
- Depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios (e.g., art works by internationally renowned artists that feature minors, news media reportage, media created and shared for scientific or educational purposes).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

Abusive behavior (20210527*)

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing, hoping, or calling for serious harm on a person or group of people

We do not tolerate content that wishes, hopes, promotes, incites, or expresses a desire for death, serious bodily harm or serious disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

About wishes of harm exceptions on Twitter

We recognize that conversations regarding certain individuals credibly accused of severe violence may prompt outrage and associated wishes of harm. In these limited cases, we will request the user to delete the Tweet without any risk of account penalty, strike, or suspension. Examples are, but not limited to:

- “I wish all rapists to die.”
- “Child abusers should be hanged.”

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone’s body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Denying mass casualty events took place

We prohibit content that denies that mass murder or other mass casualty events took place, where we can verify that the event occurred, and when the content is shared with abusive intent. This may include references to such an event as a “hoax” or claims that victims or survivors are fake or “actors.” It includes, but is not limited to, events like the Holocaust, school shootings, terrorist attacks, and natural disasters.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual’s previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy (20210127*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognize that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and more harmful.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you."

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, incites, or expresses a desire for death, serious bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

- Hoping that an entire protected category and/or individuals who may be members of that category dies as a result of a serious disease, e.g., "I hope all [nationality] get COVID and die."
- Wishing for someone to fall victim to a serious accident, e.g., "I wish that you would get run over by a car next time you run your mouth."

- Saying that a group of individuals deserve serious physical injury, e.g., "If this group of [slur] don't shut up, they deserve to be shot."
- Encouraging others to commit violence against an individual or a group based on their perceived membership in a protected category, e.g., "I'm in the mood to punch a [racial slur], who's with me?"

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

- media that depicts victims of the Holocaust;
- media that depicts lynchings.

Incitement against protected categories

We prohibit inciting behavior that targets individuals or groups of people belonging to protected categories. This includes content intended:

- to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., "all [religious group] are terrorists."
- to incite others to harass members of a protected category on or off platform, e.g., "I'm sick of these [religious group] thinking they are better than us, if any of you see someone wearing a [religious symbol of the religious group], grab it off them and post pics!"
- to incite others to discriminate in the form of denial of support to the economic enterprise of an individual or group because of their perceived membership in a protected category, e.g., "If you go to a [religious group] store, you are supporting those [slur], let's stop giving our money to these [religious slur]." This may not include content intended as political in nature, such as political commentary or content relating to boycotts or protests.

Note that content intended to incite violence against a protected category is prohibited under **Wishing, hoping, or calling for serious harm on a person or groups of people.**

We prohibit targeting individuals and groups with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., "all [religious group] are terrorists."

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion, caste, age, disability, serious disease, national origin, race, or ethnicity.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

Suicide and Self-harm Policy **(20200805*)**

Overview

You may not promote or encourage suicide or self-harm.

At Twitter, we recognize that suicide and self-harm are significant social & public health challenges that require collaboration between all stakeholders – public, private, and civil society – and that we have a role and responsibility to help people access and receive support when they need it.

When developing this policy, we consulted extensively with experts to ensure that people who have engaged in self-harm or experienced suicidal thoughts can share their personal experiences. We also recognized the need to protect people from the potential harm caused by exposure to content that could promote or encourage self-harm – intentionally or inadvertently. That's why our policy prohibits content that promotes or encourages self-harming behaviors and provides support to those undergoing experiences with self-harm or suicidal thoughts.

What is in violation of this policy?

Under this policy, you can't promote, or otherwise encourage, suicide or self-harm. We define promotion and encouragement to include statements such as "the most effective", "the easiest", "the best", "the most successful", "you should", "why don't you". Violations of this policy can occur via Tweets, images or videos, including live video.

We define suicide to be the act of taking one's own life. We define self-harm to include:

- self-inflicted physical injuries e.g., cutting; and
- eating disorders e.g., bulimia, anorexia.

Violations of this policy include, but are not limited to:

- encouraging someone to physically harm or kill themselves;
- asking others for encouragement to engage in self-harm or suicide, including seeking partners for group suicides or suicide games; and
- sharing information, strategies, methods or instructions that would assist people to engage in self-harm and suicide.

What is not a violation of this policy?

Some examples of behavior that are not considered a violation of this policy include:

- sharing personal stories and experiences related to self-harm or suicide;
- sharing coping mechanisms and resources for addressing self-harm or suicidal thoughts; and
- discussions that are focused on research, advocacy, and education related to self-harm or suicide prevention.

Note: people can share their personal experiences, but should avoid sharing detailed information about specific strategies or methods related to self-harm, as this could inadvertently encourage this behavior.

Who can report violations of this policy?

Anyone can report content that may encourage or promote suicide or self-harm via our in-app reporting or our [specialized reporting form](#). These reports are routed to a dedicated team who evaluate each case individually.

Note: if we receive a report that someone has expressed an intention to engage in self-harm or suicide, we will contact them directly, encourage them to seek support, and provide information about dedicated online and hotline resources. We may also work with law enforcement officials where appropriate, for example, if we receive a valid emergency disclosure request as defined in our [Law Enforcement Guidelines](#).

How to report violations of this policy

To ensure that we handle reports sensitively, our in-app reporting provides separate options for people who may be expressing an intention to harm themselves, and content that is encouraging or promoting self-harm or suicide.

Expressing intentions of self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Reporting form

You can also report this content for review via [our dedicated report form](#).

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Encouraging self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

What happens if you violate this policy?

Our enforcement approach depends on the type of content being shared, whether or not the reported account is encouraging or promoting self-harm or suicide, and the account's previous history of violations.

If you violate this policy by sharing content that intentionally encourages others to harm themselves, ask others to encourage you to harm yourself, or share detailed information or instructions related to self-harm or suicide methods, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy, or if your account is dedicated to promoting or encouraging self-harm or suicide, your account will be permanently suspended. If cases include images or videos related to self-harm or suicide, we will also evaluate this content under [our sensitive media policy](#). If you believe that your account was suspended in error, [you can submit an appeal](#).

We may also take steps to prevent the spread of instructional material hosted on third-party websites by [marking such links as unsafe](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Learn more](#) about how you can support someone experiencing thoughts of self-harm or suicide.

Visit our Safety Center for a list of local [mental health resources](#) and read [our blog on suicide prevention](#) to learn more about our work.

Sensitive media policy (20191204*)

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.

3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20190607*)

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their

information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;

- location information related to commercial property or places of business, where this information is publicly available;
- descriptions of physical appearance;
- gossip, rumours, accusations, and allegations; and
- screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy

(20191217*)

Overview

November 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy (20200918*)

Overview

September 2020

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are;
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting; and
- coordinated harmful activity that encourages or promotes behavior which violates the [Twitter Rules](#).

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts or by coordinating with others to violate the Twitter Rules. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and

- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate;
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account; and
 - coordinating with others to engage in or promote violations of the Twitter Rules, including violations of our [abusive behavior](#) policy.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in "follow trains," "decks," and "Retweet for Retweet" behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- "follow churn" – following and then unfollowing large numbers of accounts in an effort to inflate one's own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account's followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person's browser (malware) or computer or to compromise a person's privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally;
- coordinating with others to express ideas, viewpoints, support, or opposition towards a cause, provided such behavior does not result in violations of the Twitter Rules; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don't violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Denylisting URLs

We denylist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [financial scam policy](#), our [hacked materials policy](#), our approach to [coordinated harmful activity](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Civic integrity policy (20201028*)

Overview

October 2020

You may not use Twitter's services for the purpose of manipulating or interfering in elections or other civic processes. This includes posting or sharing content that may suppress participation or mislead people about when, where, or how to participate in a civic process. In addition, we may label and reduce the

visibility of Tweets containing false or misleading information about civic processes in order to provide additional context.

The public conversation occurring on Twitter is never more important than during elections and other civic events. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt civic processes, including through the distribution of false or misleading information about the procedures or circumstances around participation in a civic process. In instances where misleading information does not seek to directly manipulate or disrupt civic processes, but leads to confusion on our service, we may label the Tweets to give additional context.

What is a civic process?

Twitter considers civic processes to be events or procedures mandated, organized, and conducted by the governing and/or electoral body of a country, state, region, district, or municipality to address a matter of common concern through public participation. Some examples of civic processes may include:

- Political elections
- Censuses
- Major referenda and ballot initiatives

What is in violation of this policy?

This policy addresses 4 categories of misleading behavior and content:

Misleading information about how to participate

We will label or remove false or misleading information about how to participate in an election or other civic process. This includes but is not limited to:

- misleading information about procedures to participate in a civic process (for example, that you can vote by Tweet, text message, email, or phone call in jurisdictions where these are not a possibility);
- misleading information about requirements for participation, including identification or citizenship requirements;
- misleading claims that cause confusion about the established laws, regulations, procedures, and methods of a civic process, or about the actions of officials or entities executing those civic processes; and
- misleading statements or information about the official, announced date or time of a civic process.

Suppression and intimidation

We will label or remove false or misleading information intended to intimidate or dissuade people from participating in an election or other civic process. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to voting in an election, polling places, or collecting census information;
- misleading claims about long lines, equipment problems, or other disruptions at voting locations during election periods;
- misleading claims about process procedures or techniques which could dissuade people from participating; and
- threats regarding voting locations or other key places or events (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

Misleading information about outcomes

We will label or remove false or misleading information intended to undermine public confidence in an election or other civic process. This includes but is not limited to:

- disputed claims that could undermine faith in the process itself, such as unverified information about election rigging, ballot tampering, vote tallying, or certification of election results; and
- misleading claims about the results or outcome of a civic process which calls for or could lead to interference with the implementation of the results of the process, e.g. claiming victory before election results have been certified, inciting unlawful conduct to prevent the procedural or practical implementation of election results (note that our violent threats policy may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account policy](#).

What is not a violation of this policy?

Not all false or untrue information about politics or civic processes constitutes manipulation or interference. In the absence of other policy violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected or appointed official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information;
- voting and audience participation for competitions, game shows, or other entertainment purposes; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election or civic process. Therefore, we enable reporting of false or misleading information about civic processes in advance of major events, for people located in the relevant countries and locations. We also work with select government and civil society partners in these countries to provide additional channels for reporting and expedited review.

For civic processes with multiple stages or parts, such as primary elections or lengthy campaigns, reporting will be enabled leading up to the first officially-sanctioned event associated with the civic process.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country at the relevant time, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in civic processes.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in a civic process.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our civic integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Labeling

In circumstances where we do not remove content which violates this policy, we may provide additional context on Tweets sharing the content where they appear on Twitter. This means we may:

- Apply a label and/or warning message to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended; and/or
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or relevant Twitter policies.

In most cases, we will take all of the above actions on Tweets we label. In some instances, we'll also turn off your ability to reply, Retweet, or like the Tweet. We prioritize producing Twitter Moments in cases where misleading content on Twitter is gaining significant attention and has caused public confusion on our service.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Impersonation policy (20191001*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Synthetic and manipulated media policy **(20210626*)**

Overview

You may not deceptively promote synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context.

You should be able to find reliable information on Twitter. That means understanding whether the content you see is real or fabricated and having the ability to find more context about what you see on Twitter. Therefore, we may label Tweets that include media (videos, audio, and images) that have been deceptively altered or fabricated. In addition, you may not share deceptively altered media on Twitter in ways that mislead or deceive people about the media's authenticity where threats to physical safety or other serious harm may result.

We use the following criteria as we consider Tweets and media for labeling or removal under this policy as part of our ongoing work to enforce our rules and ensure healthy and safe conversation on Twitter (additional information is available below):

1. Is the content synthetic or manipulated?

In order for content to be labeled or removed under this policy, we must have reason to believe that media, or the context in which media are presented, are significantly and deceptively altered or manipulated. Synthetic and manipulated media take many different forms and people can employ a wide range of technologies to produce these media. In assessing whether media have been significantly and deceptively altered or fabricated, some of the factors we consider include:

- whether the content has been substantially edited in a manner that fundamentally alters its composition, sequence, timing, or framing;
- any visual or auditory information (such as new video frames, overdubbed audio, or modified subtitles) that has been added or removed; and
- whether media depicting a real person have been fabricated or simulated

We are most likely to take action (either labeling or removal, as described below) on more significant forms of alteration, such as wholly synthetic audio or video or content that has been doctored (spliced and reordered, slowed down) to change its meaning. Subtler forms of manipulated media, such as isolative editing, omission of context, or presentation with false context, may be labeled or removed on a case-by-case basis.

We will not take action to label or remove media that have been edited in ways that do not fundamentally alter their meaning, such as retouched photos or color-corrected videos.

In order to determine if media have been significantly and deceptively altered or fabricated, we may use our own technology or receive reports through partnerships with third parties. In situations where we are unable to reliably determine if media have been altered or fabricated, we may not take action to label or remove them.

2. Is the content shared in a deceptive manner?

We also consider whether the context in which media are shared could result in confusion or misunderstanding or suggests a deliberate intent to deceive people about the nature or origin of the content, for example by falsely claiming that it depicts reality. We assess the context provided alongside media to see whether it makes clear that the media have been altered or fabricated. Some of the types of context we assess in order to make this determination include:

- The text of the Tweet accompanying or within media
- Metadata associated with media
- Information on the profile of the account sharing media
- Websites linked in the Tweet, or in the profile of the account sharing media

3. Is the content likely to impact public safety or cause serious harm?

Tweets that share synthetic and manipulated media are subject to removal under this policy if they are likely to cause serious harm. Some specific harms we consider include:

- Threats to the physical safety of a person or group
- Risk of mass violence or widespread civil unrest
- Threats to the privacy or ability of a person or group to freely express themselves or participate in civic events, such as:
 - Stalking or unwanted and obsessive attention
 - Targeted content that includes tropes, epithets, or material that aims to silence someone
 - Voter suppression or intimidation

While we have other rules also intended to address these forms of harm, including our policies on violent threats, election integrity, and hateful conduct, we will err toward removal in borderline cases that might otherwise not violate existing rules for Tweets that include synthetic or manipulated media.

We also consider the time frame within which the content may be likely to impact public safety or cause serious harm, and are more likely to remove content under this policy if we find that immediate harms are likely to result from the content's presence on Twitter.

Note: We may also take action on synthetic and manipulated content under our [non-consensual nudity policy](#) (such as pornographic media altered to insert the faces of people not actually involved) or other parts of the Twitter Rules.

Labeling and removal

In most cases, if we have reason to believe that media shared in a Tweet have been significantly and deceptively altered or fabricated, we will provide additional context on Tweets sharing the media where they appear on Twitter. This means we may:

- Apply a label to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended;
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or landing page; and/or
- Turn off likes, replies, and Retweets.

In most cases, we will take all of the above actions on Tweets we label.

Media that meet **all three of the criteria** defined above—i.e. that are synthetic or manipulated, shared in a deceptive manner, and is likely to cause harm—may not be shared on Twitter and are subject to removal. Accounts engaging in repeated or severe violations of this policy may be permanently suspended.

** Other parts of the Twitter Rules apply and may lead to the removal of the content, particularly where there is high likelihood of severe harm, such as a threat to someone's life or physical safety.*

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Copyright policy (20200917*)

What types of copyright complaints does Twitter respond to?

Twitter responds to copyright complaints submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review [Twitter's trademark policy](#). If you are concerned about a parody, newsfeed, commentary, or fan account, please see the relevant policy [here](#). These are generally not copyright issues.

Am I a copyright holder? How do I know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

What to consider before submitting a copyright complaint

Before submitting a copyright complaint to us, please consider whether or not the use could be considered [fair use](#).

If you have considered fair use, and you still wish to continue with a copyright complaint, you may want to first reach out to the user in question to see if you can resolve the matter directly with the user. You can reply to the

user's Tweet or send the user a Direct Message and ask for them to remove your copyrighted content without having to contact Twitter.

Prior to submitting a formal complaint with Twitter, please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

Note: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph. If you're unsure whether or not you own the copyrights to a work, or if you're infringing upon someone else's work, please consult an attorney or another advisor.

What information do you need to process a copyright complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the complaint is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How do I file a copyright complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and [filing a copyright complaint](#). If you are logged in to [twitter.com](#), you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a full copy of your notice (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a copyright complaint.

How are claims processed?

We process copyright complaints in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. If you do not receive a ticket confirmation that means we did not receive your complaint and you should re-submit your complaint. However, please note, submitting duplicate copyright complaints will result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) and provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice. We will also forward a redacted copy of the complaint to [Lumen](#), with your personal information removed.

What information gets forwarded to the reported user(s)?

If we remove or disable access to the materials reported in a copyright complaint, the reported user(s) will receive a copy of the complaint, including the reporter's full name, email, street address, and any other information included in the complaint.

If you are uncomfortable sharing your contact information with the reported user(s), you may wish to consider appointing an agent to submit your DMCA notice on your behalf. Your agent will be required to submit the DMCA notice with valid contact information, and identify you as the content owner that they are representing.

What happens next?

Twitter's response to copyright complaints may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a copyright complaint, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the complaint, along with instructions for filing a counter-notice.

If you've not yet received a copy of the copyright complaint regarding the content removed from your account, please respond to the support ticket we sent you. If your complaint involves removal of content from Periscope, please see this help center [article](#) regarding how to request a copy of your complaint.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

My content was removed from Twitter

Why did I receive a copyright complaint?

If you receive a copyright complaint, it means that access to the content described in the complaint has been restricted. Please take the time to read through our correspondence to you, which includes information on the

complaint we received as well as instructions on how to file a counter-notice. Please ensure that you are monitoring the email address associated with your Twitter account.

Tip: Removing the material reported in a copyright complaint will not resolve that complaint.

What if I want to contest the takedown?

If you believe that the materials reported in the copyright complaints were misidentified or removed in error, you may send us a counter-notification(s). A counter-notice is a request for Twitter to reinstate the removed material, and it has legal consequences. Alternatively, you may be able to seek a retraction of the copyright complaint from the reporter.

How do I seek a retraction?

The DMCA complaint you received includes the contact information of the reporter. You may want to reach out and ask them to retract their notice. The reporter can send retractions to copyright@twitter.com, and should include: (1) identification of the material that was disabled, and (2) a statement that the reporter would like to retract their DMCA notice. This is the fastest and most efficient means of resolving an unresolved copyright complaint. A retraction is at the sole discretion of the original reporter.

When should I file a counter-notice?

A counter-notice is a request for Twitter to reinstate the removed material, and is the start of a legal process that has legal consequences. For example, submitting a counter notice indicates that you consent to the jurisdiction of a U.S. Federal court and that you consent to the disclosure of your personal information to the reporter and Lumen website.

With these considerations in mind, you may file a counter-notice if you believe that this material was misidentified, or you have a good faith belief that the material should not have been removed. If you're unsure whether or not you should file a counter-notice, you may want to consult with an attorney.

Tip: Re-posting material removed in response to a copyright complaint may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notice rather than re-posting the material.

What information do you need to process a counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What happens after I submit a counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. This means that the contact information that is submitted in your counter-notice will be shared to the person who filed the original notice.

If the copyright owner disagrees that the content was removed in error or misidentification, they may pursue legal action against you. If we do not receive notice within 10 business days that the original reporter is seeking a court

order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a copyright complaint or counter-notice is serious business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

What happens if my account receives multiple copyright complaints?

If multiple copyright complaints are received Twitter may lock accounts or take other actions to warn repeat violators. These warnings may vary across Twitter's services. Under appropriate circumstances we may suspend user accounts under our repeat infringer policy. However, we may take retractions and counter-notices into account when applying our repeat infringer policy.

Our approach to policy development and enforcement philosophy **(20190520*)**

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet

does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (20200204*)

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Synthetic and manipulated media: You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.

5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Violent organizations policy (20201026*)

Overview

October 2020

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

Other violent organizations

Other violent organizations are those that meet all of the below criteria:

- a collection of individuals with a shared purpose; and
- have systematically targeted civilians with violence.

We examine a group's activities both on and off Twitter to determine whether they satisfy the above criteria.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a violent organization;
- recruiting for a violent organization;
- providing or distributing services (e.g., financial, media/propaganda) to further a violent organization's stated goals; and
- using the insignia or symbol of violent organizations to promote them or indicate affiliation or support.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to state or governmental organizations.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.

6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wishing for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy (20201012*)

Overview

October 2020

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives;
- identifying alleged victims of childhood sexual exploitation by name or image; and
- promoting or normalizing sexual attraction to minors as a form of identity or sexual orientation.

What is not a violation of this policy?

- Discussions related to child sexual exploitation are permitted, provided they don't normalise, promote or glorify child sexual exploitation **in any way**;
- Conversations about help-seeking behaviour of individuals who may struggle with an attraction to minors;
- Advocacy against illegal or harmful activity involving minors provided there is no sharing of, or linking to any material featuring child sexual exploitation; and
- Depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios (e.g., art works by internationally renowned artists that feature minors, news media reportage, media created and shared for scientific or educational purposes).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

Abusive behavior (20210527*)

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing, hoping, or calling for serious harm on a person or group of people

We do not tolerate content that wishes, hopes, promotes, incites, or expresses a desire for death, serious bodily harm or serious disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

About wishes of harm exceptions on Twitter

We recognize that conversations regarding certain individuals credibly accused of severe violence may prompt outrage and associated wishes of harm. In these limited cases, we will request the user to delete the Tweet without any risk of account penalty, strike, or suspension. Examples are, but not limited to:

- “I wish all rapists to die.”
- “Child abusers should be hanged.”

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone’s body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Denying mass casualty events took place

We prohibit content that denies that mass murder or other mass casualty events took place, where we can verify that the event occurred, and when the content is shared with abusive intent. This may include references to such an event as a “hoax” or claims that victims or survivors are fake or “actors.” It includes, but is not limited to, events like the Holocaust, school shootings, terrorist attacks, and natural disasters.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual’s previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy (20210127*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognize that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and more harmful.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you."

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, incites, or expresses a desire for death, serious bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

- Hoping that an entire protected category and/or individuals who may be members of that category dies as a result of a serious disease, e.g., "I hope all [nationality] get COVID and die."
- Wishing for someone to fall victim to a serious accident, e.g., "I wish that you would get run over by a car next time you run your mouth."

- Saying that a group of individuals deserve serious physical injury, e.g., "If this group of [slur] don't shut up, they deserve to be shot."
- Encouraging others to commit violence against an individual or a group based on their perceived membership in a protected category, e.g., "I'm in the mood to punch a [racial slur], who's with me?"

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

- media that depicts victims of the Holocaust;
- media that depicts lynchings.

Incitement against protected categories

We prohibit inciting behavior that targets individuals or groups of people belonging to protected categories. This includes content intended:

- to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., "all [religious group] are terrorists."
- to incite others to harass members of a protected category on or off platform, e.g., "I'm sick of these [religious group] thinking they are better than us, if any of you see someone wearing a [religious symbol of the religious group], grab it off them and post pics!"
- to incite others to discriminate in the form of denial of support to the economic enterprise of an individual or group because of their perceived membership in a protected category, e.g., "If you go to a [religious group] store, you are supporting those [slur], let's stop giving our money to these [religious slur]." This may not include content intended as political in nature, such as political commentary or content relating to boycotts or protests.

Note that content intended to incite violence against a protected category is prohibited under **Wishing, hoping, or calling for serious harm on a person or groups of people.**

We prohibit targeting individuals and groups with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., "all [religious group] are terrorists."

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion, caste, age, disability, serious disease, national origin, race, or ethnicity.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

Suicide and Self-harm Policy **(20200805*)**

Overview

You may not promote or encourage suicide or self-harm.

At Twitter, we recognize that suicide and self-harm are significant social & public health challenges that require collaboration between all stakeholders – public, private, and civil society – and that we have a role and responsibility to help people access and receive support when they need it.

When developing this policy, we consulted extensively with experts to ensure that people who have engaged in self-harm or experienced suicidal thoughts can share their personal experiences. We also recognized the need to protect people from the potential harm caused by exposure to content that could promote or encourage self-harm – intentionally or inadvertently. That's why our policy prohibits content that promotes or encourages self-harming behaviors and provides support to those undergoing experiences with self-harm or suicidal thoughts.

What is in violation of this policy?

Under this policy, you can't promote, or otherwise encourage, suicide or self-harm. We define promotion and encouragement to include statements such as "the most effective", "the easiest", "the best", "the most successful", "you should", "why don't you". Violations of this policy can occur via Tweets, images or videos, including live video.

We define suicide to be the act of taking one's own life. We define self-harm to include:

- self-inflicted physical injuries e.g., cutting; and
- eating disorders e.g., bulimia, anorexia.

Violations of this policy include, but are not limited to:

- encouraging someone to physically harm or kill themselves;
- asking others for encouragement to engage in self-harm or suicide, including seeking partners for group suicides or suicide games; and
- sharing information, strategies, methods or instructions that would assist people to engage in self-harm and suicide.

What is not a violation of this policy?

Some examples of behavior that are not considered a violation of this policy include:

- sharing personal stories and experiences related to self-harm or suicide;
- sharing coping mechanisms and resources for addressing self-harm or suicidal thoughts; and
- discussions that are focused on research, advocacy, and education related to self-harm or suicide prevention.

Note: people can share their personal experiences, but should avoid sharing detailed information about specific strategies or methods related to self-harm, as this could inadvertently encourage this behavior.

Who can report violations of this policy?

Anyone can report content that may encourage or promote suicide or self-harm via our in-app reporting or our [specialized reporting form](#). These reports are routed to a dedicated team who evaluate each case individually.

Note: if we receive a report that someone has expressed an intention to engage in self-harm or suicide, we will contact them directly, encourage them to seek support, and provide information about dedicated online and hotline resources. We may also work with law enforcement officials where appropriate, for example, if we receive a valid emergency disclosure request as defined in our [Law Enforcement Guidelines](#).

How to report violations of this policy

To ensure that we handle reports sensitively, our in-app reporting provides separate options for people who may be expressing an intention to harm themselves, and content that is encouraging or promoting self-harm or suicide.

Expressing intentions of self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Reporting form

You can also report this content for review via [our dedicated report form](#).

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Encouraging self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

What happens if you violate this policy?

Our enforcement approach depends on the type of content being shared, whether or not the reported account is encouraging or promoting self-harm or suicide, and the account's previous history of violations.

If you violate this policy by sharing content that intentionally encourages others to harm themselves, ask others to encourage you to harm yourself, or share detailed information or instructions related to self-harm or suicide methods, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy, or if your account is dedicated to promoting or encouraging self-harm or suicide, your account will be permanently suspended. If cases include images or videos related to self-harm or suicide, we will also evaluate this content under [our sensitive media policy](#). If you believe that your account was suspended in error, [you can submit an appeal](#).

We may also take steps to prevent the spread of instructional material hosted on third-party websites by [marking such links as unsafe](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Learn more](#) about how you can support someone experiencing thoughts of self-harm or suicide.

Visit our Safety Center for a list of local [mental health resources](#) and read [our blog on suicide prevention](#) to learn more about our work.

Sensitive media policy (20191204*)

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.

3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20210914*)

Overview

September 2021

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

In addition, you may not share private media, such as images or videos of private individuals, without their consent. However, we recognise that there are instances where users may share images or videos of private individuals, who are not public

figures, as part of a newsworthy event or to further public discourse on issues or events of public interest. In such cases, we may allow the media to remain on the platform.

Sharing private information:

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We take this into consideration because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from physical harm as a result of their information being shared, so we consider information such as physical location and phone numbers to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly

available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

Sharing private media:

Posting images is an important part of our users' experience on Twitter. Where individuals have a reasonable expectation of privacy in an individual piece of media, we believe they should be able to determine whether or not it is shared. Sharing such media could potentially violate users' privacy and may lead to emotional or physical harm. When we are notified by individuals depicted, or their authorized representative, that they did not consent to having image shared, we will remove the media. This policy is not applicable to public figures.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;

- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details;
- other private information, including biometric data or medical records;
- media of private individuals without the permission of the person(s) depicted.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;

- descriptions of physical appearance;
- gossip, rumours, accusations, and allegations; and
- screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

For media, the following are not in violation of our policy:

- the media is publicly available or is being covered by mainstream media;
- the media and the accompanying tweet text add value to the public discourse or are shared in public interest;
- contains eyewitness accounts or on the ground reports from developing events;
- the subject of the media is a public figure.

Who can report violations of this policy?

Sharing private information:

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

Sharing private media:

When reporting private media, we need a first person report in order to make the determination that the image or video has been shared without their permission. We accept reports from:

1. Individuals depicted in the reported media or
2. Their authorized representative such as:
 - a. parents/legal guardians,

- b. lawyers, or
- c. legal representative

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

Sharing private information:

The first time you violate this policy by sharing private information (such as home address, identity documents etc.), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy by posting private information again after your first warning, your account will be permanently suspended.

Sharing private media:

If you violate this policy by sharing private media of an individual without their consent, we will require you to remove the content and temporarily lock your account while you remove the media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy **(20191217*)**

Overview

November 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepsshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy (20200918*)

Overview

September 2020

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are;
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting; and
- coordinated harmful activity that encourages or promotes behavior which violates the [Twitter Rules](#).

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts or by coordinating with others to violate the Twitter Rules. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and

- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate;
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account; and
 - coordinating with others to engage in or promote violations of the Twitter Rules, including violations of our [abusive behavior](#) policy.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in "follow trains," "decks," and "Retweet for Retweet" behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- "follow churn" – following and then unfollowing large numbers of accounts in an effort to inflate one's own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account's followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person's browser (malware) or computer or to compromise a person's privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally;
- coordinating with others to express ideas, viewpoints, support, or opposition towards a cause, provided such behavior does not result in violations of the Twitter Rules; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don't violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Denylisting URLs

We denylist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [financial scam policy](#), our [hacked materials policy](#), our approach to [coordinated harmful activity](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Civic integrity policy (20201028*)

Overview

October 2020

You may not use Twitter's services for the purpose of manipulating or interfering in elections or other civic processes. This includes posting or sharing content that may suppress participation or mislead people about when, where, or how to participate in a civic process. In addition, we may label and reduce the

visibility of Tweets containing false or misleading information about civic processes in order to provide additional context.

The public conversation occurring on Twitter is never more important than during elections and other civic events. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt civic processes, including through the distribution of false or misleading information about the procedures or circumstances around participation in a civic process. In instances where misleading information does not seek to directly manipulate or disrupt civic processes, but leads to confusion on our service, we may label the Tweets to give additional context.

What is a civic process?

Twitter considers civic processes to be events or procedures mandated, organized, and conducted by the governing and/or electoral body of a country, state, region, district, or municipality to address a matter of common concern through public participation. Some examples of civic processes may include:

- Political elections
- Censuses
- Major referenda and ballot initiatives

What is in violation of this policy?

This policy addresses 4 categories of misleading behavior and content:

Misleading information about how to participate

We will label or remove false or misleading information about how to participate in an election or other civic process. This includes but is not limited to:

- misleading information about procedures to participate in a civic process (for example, that you can vote by Tweet, text message, email, or phone call in jurisdictions where these are not a possibility);
- misleading information about requirements for participation, including identification or citizenship requirements;
- misleading claims that cause confusion about the established laws, regulations, procedures, and methods of a civic process, or about the actions of officials or entities executing those civic processes; and
- misleading statements or information about the official, announced date or time of a civic process.

Suppression and intimidation

We will label or remove false or misleading information intended to intimidate or dissuade people from participating in an election or other civic process. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to voting in an election, polling places, or collecting census information;
- misleading claims about long lines, equipment problems, or other disruptions at voting locations during election periods;
- misleading claims about process procedures or techniques which could dissuade people from participating; and
- threats regarding voting locations or other key places or events (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

Misleading information about outcomes

We will label or remove false or misleading information intended to undermine public confidence in an election or other civic process. This includes but is not limited to:

- disputed claims that could undermine faith in the process itself, such as unverified information about election rigging, ballot tampering, vote tallying, or certification of election results; and
- misleading claims about the results or outcome of a civic process which calls for or could lead to interference with the implementation of the results of the process, e.g. claiming victory before election results have been certified, inciting unlawful conduct to prevent the procedural or practical implementation of election results (note that our violent threats policy may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account policy](#).

What is not a violation of this policy?

Not all false or untrue information about politics or civic processes constitutes manipulation or interference. In the absence of other policy violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected or appointed official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information;
- voting and audience participation for competitions, game shows, or other entertainment purposes; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election or civic process. Therefore, we enable reporting of false or misleading information about civic processes in advance of major events, for people located in the relevant countries and locations. We also work with select government and civil society partners in these countries to provide additional channels for reporting and expedited review.

For civic processes with multiple stages or parts, such as primary elections or lengthy campaigns, reporting will be enabled leading up to the first officially-sanctioned event associated with the civic process.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country at the relevant time, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in civic processes.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in a civic process.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our civic integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Labeling

In circumstances where we do not remove content which violates this policy, we may provide additional context on Tweets sharing the content where they appear on Twitter. This means we may:

- Apply a label and/or warning message to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended; and/or
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or relevant Twitter policies.

In most cases, we will take all of the above actions on Tweets we label. In some instances, we'll also turn off your ability to reply, Retweet, or like the Tweet. We prioritize producing Twitter Moments in cases where misleading content on Twitter is gaining significant attention and has caused public confusion on our service.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Impersonation policy (20191001*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Synthetic and manipulated media policy

(20210626*)

Overview

You may not deceptively promote synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context.

You should be able to find reliable information on Twitter. That means understanding whether the content you see is real or fabricated and having the ability to find more context about what you see on Twitter. Therefore, we may label Tweets that include media (videos, audio, and images) that have been deceptively altered or fabricated. In addition, you may not share deceptively altered media on Twitter in ways that mislead or deceive people about the media's authenticity where threats to physical safety or other serious harm may result.

We use the following criteria as we consider Tweets and media for labeling or removal under this policy as part of our ongoing work to enforce our rules and ensure healthy and safe conversation on Twitter (additional information is available below):

1. Is the content synthetic or manipulated?

In order for content to be labeled or removed under this policy, we must have reason to believe that media, or the context in which media are presented, are significantly and deceptively altered or manipulated. Synthetic and manipulated media take many different forms and people can employ a wide range of technologies to produce these media. In assessing whether media have been significantly and deceptively altered or fabricated, some of the factors we consider include:

- whether the content has been substantially edited in a manner that fundamentally alters its composition, sequence, timing, or framing;
- any visual or auditory information (such as new video frames, overdubbed audio, or modified subtitles) that has been added or removed; and
- whether media depicting a real person have been fabricated or simulated

We are most likely to take action (either labeling or removal, as described below) on more significant forms of alteration, such as wholly synthetic audio or video or content that has been doctored (spliced and reordered, slowed down) to change its meaning. Subtler forms of manipulated media, such as isolative editing, omission of context, or presentation with false context, may be labeled or removed on a case-by-case basis.

We will not take action to label or remove media that have been edited in ways that do not fundamentally alter their meaning, such as retouched photos or color-corrected videos.

In order to determine if media have been significantly and deceptively altered or fabricated, we may use our own technology or receive reports through partnerships with third parties. In situations where we are unable to reliably determine if media have been altered or fabricated, we may not take action to label or remove them.

2. Is the content shared in a deceptive manner?

We also consider whether the context in which media are shared could result in confusion or misunderstanding or suggests a deliberate intent to deceive people about the nature or origin of the content, for example by falsely claiming that it depicts reality. We assess the context provided alongside media to see whether it makes clear that the media have been altered or fabricated. Some of the types of context we assess in order to make this determination include:

- The text of the Tweet accompanying or within media
- Metadata associated with media
- Information on the profile of the account sharing media
- Websites linked in the Tweet, or in the profile of the account sharing media

3. Is the content likely to impact public safety or cause serious harm?

Tweets that share synthetic and manipulated media are subject to removal under this policy if they are likely to cause serious harm. Some specific harms we consider include:

- Threats to the physical safety of a person or group
- Risk of mass violence or widespread civil unrest
- Threats to the privacy or ability of a person or group to freely express themselves or participate in civic events, such as:
 - Stalking or unwanted and obsessive attention
 - Targeted content that includes tropes, epithets, or material that aims to silence someone
 - Voter suppression or intimidation

While we have other rules also intended to address these forms of harm, including our policies on violent threats, election integrity, and hateful conduct, we will err toward removal in borderline cases that might otherwise not violate existing rules for Tweets that include synthetic or manipulated media.

We also consider the time frame within which the content may be likely to impact public safety or cause serious harm, and are more likely to remove content under this policy if we find that immediate harms are likely to result from the content's presence on Twitter.

Note: We may also take action on synthetic and manipulated content under our [non-consensual nudity policy](#) (such as pornographic media altered to insert the faces of people not actually involved) or other parts of the Twitter Rules.

Labeling and removal

In most cases, if we have reason to believe that media shared in a Tweet have been significantly and deceptively altered or fabricated, we will provide additional context on Tweets sharing the media where they appear on Twitter. This means we may:

- Apply a label to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended;
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or landing page; and/or
- Turn off likes, replies, and Retweets.

In most cases, we will take all of the above actions on Tweets we label.

Media that meet **all three of the criteria** defined above—i.e. that are synthetic or manipulated, shared in a deceptive manner, and is likely to cause harm—may not be shared on Twitter and are subject to removal. Accounts engaging in repeated or severe violations of this policy may be permanently suspended.

** Other parts of the Twitter Rules apply and may lead to the removal of the content, particularly where there is high likelihood of severe harm, such as a threat to someone's life or physical safety.*

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Copyright policy (20200917*)

What types of copyright complaints does Twitter respond to?

Twitter responds to copyright complaints submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review [Twitter's trademark policy](#). If you are concerned about a parody, newsfeed, commentary, or fan account, please see the relevant policy [here](#). These are generally not copyright issues.

Am I a copyright holder? How do I know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

What to consider before submitting a copyright complaint

Before submitting a copyright complaint to us, please consider whether or not the use could be considered [fair use](#).

If you have considered fair use, and you still wish to continue with a copyright complaint, you may want to first reach out to the user in question to see if you can resolve the matter directly with the user. You can reply to the

user's Tweet or send the user a Direct Message and ask for them to remove your copyrighted content without having to contact Twitter.

Prior to submitting a formal complaint with Twitter, please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

Note: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph. If you're unsure whether or not you own the copyrights to a work, or if you're infringing upon someone else's work, please consult an attorney or another advisor.

What information do you need to process a copyright complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the complaint is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How do I file a copyright complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and [filing a copyright complaint](#). If you are logged in to [twitter.com](#), you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a full copy of your notice (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a copyright complaint.

How are claims processed?

We process copyright complaints in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. If you do not receive a ticket confirmation that means we did not receive your complaint and you should re-submit your complaint. However, please note, submitting duplicate copyright complaints will result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) and provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice. We will also forward a redacted copy of the complaint to [Lumen](#), with your personal information removed.

What information gets forwarded to the reported user(s)?

If we remove or disable access to the materials reported in a copyright complaint, the reported user(s) will receive a copy of the complaint, including the reporter's full name, email, street address, and any other information included in the complaint.

If you are uncomfortable sharing your contact information with the reported user(s), you may wish to consider appointing an agent to submit your DMCA notice on your behalf. Your agent will be required to submit the DMCA notice with valid contact information, and identify you as the content owner that they are representing.

What happens next?

Twitter's response to copyright complaints may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a copyright complaint, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the complaint, along with instructions for filing a counter-notice.

If you've not yet received a copy of the copyright complaint regarding the content removed from your account, please respond to the support ticket we sent you. If your complaint involves removal of content from Periscope, please see this help center [article](#) regarding how to request a copy of your complaint.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

My content was removed from Twitter

Why did I receive a copyright complaint?

If you receive a copyright complaint, it means that access to the content described in the complaint has been restricted. Please take the time to read through our correspondence to you, which includes information on the

complaint we received as well as instructions on how to file a counter-notice. Please ensure that you are monitoring the email address associated with your Twitter account.

Tip: Removing the material reported in a copyright complaint will not resolve that complaint.

What if I want to contest the takedown?

If you believe that the materials reported in the copyright complaints were misidentified or removed in error, you may send us a counter-notification(s). A counter-notice is a request for Twitter to reinstate the removed material, and it has legal consequences. Alternatively, you may be able to seek a retraction of the copyright complaint from the reporter.

How do I seek a retraction?

The DMCA complaint you received includes the contact information of the reporter. You may want to reach out and ask them to retract their notice. The reporter can send retractions to copyright@twitter.com, and should include: (1) identification of the material that was disabled, and (2) a statement that the reporter would like to retract their DMCA notice. This is the fastest and most efficient means of resolving an unresolved copyright complaint. A retraction is at the sole discretion of the original reporter.

When should I file a counter-notice?

A counter-notice is a request for Twitter to reinstate the removed material, and is the start of a legal process that has legal consequences. For example, submitting a counter notice indicates that you consent to the jurisdiction of a U.S. Federal court and that you consent to the disclosure of your personal information to the reporter and Lumen website.

With these considerations in mind, you may file a counter-notice if you believe that this material was misidentified, or you have a good faith belief that the material should not have been removed. If you're unsure whether or not you should file a counter-notice, you may want to consult with an attorney.

Tip: Re-posting material removed in response to a copyright complaint may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notice rather than re-posting the material.

What information do you need to process a counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What happens after I submit a counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. This means that the contact information that is submitted in your counter-notice will be shared to the person who filed the original notice.

If the copyright owner disagrees that the content was removed in error or misidentification, they may pursue legal action against you. If we do not receive notice within 10 business days that the original reporter is seeking a court

order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a copyright complaint or counter-notice is serious business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

What happens if my account receives multiple copyright complaints?

If multiple copyright complaints are received Twitter may lock accounts or take other actions to warn repeat violators. These warnings may vary across Twitter's services. Under appropriate circumstances we may suspend user accounts under our repeat infringer policy. However, we may take retractions and counter-notices into account when applying our repeat infringer policy.

Our approach to policy development and enforcement philosophy **(20190520*)**

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

[Public impact of the content](#): A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet

does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (20200204*)

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Synthetic and manipulated media: You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.

5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Violent organizations policy (20201026*)

Overview

October 2020

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

Other violent organizations

Other violent organizations are those that meet all of the below criteria:

- a collection of individuals with a shared purpose; and
- have systematically targeted civilians with violence.

We examine a group's activities both on and off Twitter to determine whether they satisfy the above criteria.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a violent organization;
- recruiting for a violent organization;
- providing or distributing services (e.g., financial, media/propaganda) to further a violent organization's stated goals; and
- using the insignia or symbol of violent organizations to promote them or indicate affiliation or support.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to state or governmental organizations.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.

6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wishing for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy (20201012*)

Overview

October 2020

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives;
- identifying alleged victims of childhood sexual exploitation by name or image; and
- promoting or normalizing sexual attraction to minors as a form of identity or sexual orientation.

What is not a violation of this policy?

- Discussions related to child sexual exploitation are permitted, provided they don't normalise, promote or glorify child sexual exploitation **in any way**;
- Conversations about help-seeking behaviour of individuals who may struggle with an attraction to minors;
- Advocacy against illegal or harmful activity involving minors provided there is no sharing of, or linking to any material featuring child sexual exploitation; and
- Depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios (e.g., art works by internationally renowned artists that feature minors, news media reportage, media created and shared for scientific or educational purposes).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

Abusive behavior (20210527*)

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing, hoping, or calling for serious harm on a person or group of people

We do not tolerate content that wishes, hopes, promotes, incites, or expresses a desire for death, serious bodily harm or serious disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

About wishes of harm exceptions on Twitter

We recognize that conversations regarding certain individuals credibly accused of severe violence may prompt outrage and associated wishes of harm. In these limited cases, we will request the user to delete the Tweet without any risk of account penalty, strike, or suspension. Examples are, but not limited to:

- “I wish all rapists to die.”
- “Child abusers should be hanged.”

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone’s body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Denying mass casualty events took place

We prohibit content that denies that mass murder or other mass casualty events took place, where we can verify that the event occurred, and when the content is shared with abusive intent. This may include references to such an event as a “hoax” or claims that victims or survivors are fake or “actors.” It includes, but is not limited to, events like the Holocaust, school shootings, terrorist attacks, and natural disasters.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual’s previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy (20210127*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognize that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and more harmful.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you."

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, incites, or expresses a desire for death, serious bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

- Hoping that an entire protected category and/or individuals who may be members of that category dies as a result of a serious disease, e.g., "I hope all [nationality] get COVID and die."
- Wishing for someone to fall victim to a serious accident, e.g., "I wish that you would get run over by a car next time you run your mouth."

- Saying that a group of individuals deserve serious physical injury, e.g., "If this group of [slur] don't shut up, they deserve to be shot."
- Encouraging others to commit violence against an individual or a group based on their perceived membership in a protected category, e.g., "I'm in the mood to punch a [racial slur], who's with me?"

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

- media that depicts victims of the Holocaust;
- media that depicts lynchings.

Incitement against protected categories

We prohibit inciting behavior that targets individuals or groups of people belonging to protected categories. This includes content intended:

- to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., "all [religious group] are terrorists."
- to incite others to harass members of a protected category on or off platform, e.g., "I'm sick of these [religious group] thinking they are better than us, if any of you see someone wearing a [religious symbol of the religious group], grab it off them and post pics!"
- to incite others to discriminate in the form of denial of support to the economic enterprise of an individual or group because of their perceived membership in a protected category, e.g., "If you go to a [religious group] store, you are supporting those [slur], let's stop giving our money to these [religious slur]." This may not include content intended as political in nature, such as political commentary or content relating to boycotts or protests.

Note that content intended to incite violence against a protected category is prohibited under **Wishing, hoping, or calling for serious harm on a person or groups of people.**

We prohibit targeting individuals and groups with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., "all [religious group] are terrorists."

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion, caste, age, disability, serious disease, national origin, race, or ethnicity.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

Suicide and Self-harm Policy **(20200805*)**

Overview

You may not promote or encourage suicide or self-harm.

At Twitter, we recognize that suicide and self-harm are significant social & public health challenges that require collaboration between all stakeholders – public, private, and civil society – and that we have a role and responsibility to help people access and receive support when they need it.

When developing this policy, we consulted extensively with experts to ensure that people who have engaged in self-harm or experienced suicidal thoughts can share their personal experiences. We also recognized the need to protect people from the potential harm caused by exposure to content that could promote or encourage self-harm – intentionally or inadvertently. That's why our policy prohibits content that promotes or encourages self-harming behaviors and provides support to those undergoing experiences with self-harm or suicidal thoughts.

What is in violation of this policy?

Under this policy, you can't promote, or otherwise encourage, suicide or self-harm. We define promotion and encouragement to include statements such as "the most effective", "the easiest", "the best", "the most successful", "you should", "why don't you". Violations of this policy can occur via Tweets, images or videos, including live video.

We define suicide to be the act of taking one's own life. We define self-harm to include:

- self-inflicted physical injuries e.g., cutting; and
- eating disorders e.g., bulimia, anorexia.

Violations of this policy include, but are not limited to:

- encouraging someone to physically harm or kill themselves;
- asking others for encouragement to engage in self-harm or suicide, including seeking partners for group suicides or suicide games; and
- sharing information, strategies, methods or instructions that would assist people to engage in self-harm and suicide.

What is not a violation of this policy?

Some examples of behavior that are not considered a violation of this policy include:

- sharing personal stories and experiences related to self-harm or suicide;
- sharing coping mechanisms and resources for addressing self-harm or suicidal thoughts; and
- discussions that are focused on research, advocacy, and education related to self-harm or suicide prevention.

Note: people can share their personal experiences, but should avoid sharing detailed information about specific strategies or methods related to self-harm, as this could inadvertently encourage this behavior.

Who can report violations of this policy?

Anyone can report content that may encourage or promote suicide or self-harm via our in-app reporting or our [specialized reporting form](#). These reports are routed to a dedicated team who evaluate each case individually.

Note: if we receive a report that someone has expressed an intention to engage in self-harm or suicide, we will contact them directly, encourage them to seek support, and provide information about dedicated online and hotline resources. We may also work with law enforcement officials where appropriate, for example, if we receive a valid emergency disclosure request as defined in our [Law Enforcement Guidelines](#).

How to report violations of this policy

To ensure that we handle reports sensitively, our in-app reporting provides separate options for people who may be expressing an intention to harm themselves, and content that is encouraging or promoting self-harm or suicide.

Expressing intentions of self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Reporting form

You can also report this content for review via [our dedicated report form](#).

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Encouraging self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

What happens if you violate this policy?

Our enforcement approach depends on the type of content being shared, whether or not the reported account is encouraging or promoting self-harm or suicide, and the account's previous history of violations.

If you violate this policy by sharing content that intentionally encourages others to harm themselves, ask others to encourage you to harm yourself, or share detailed information or instructions related to self-harm or suicide methods, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy, or if your account is dedicated to promoting or encouraging self-harm or suicide, your account will be permanently suspended. If cases include images or videos related to self-harm or suicide, we will also evaluate this content under [our sensitive media policy](#). If you believe that your account was suspended in error, [you can submit an appeal](#).

We may also take steps to prevent the spread of instructional material hosted on third-party websites by [marking such links as unsafe](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Learn more](#) about how you can support someone experiencing thoughts of self-harm or suicide.

Visit our Safety Center for a list of local [mental health resources](#) and read [our blog on suicide prevention](#) to learn more about our work.

Sensitive media policy (20191204*)

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.

3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20210914*)

Overview

September 2021

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

In addition, you may not share private media, such as images or videos of private individuals, without their consent. However, we recognise that there are instances where users may share images or videos of private individuals, who are not public

figures, as part of a newsworthy event or to further public discourse on issues or events of public interest. In such cases, we may allow the media to remain on the platform.

Sharing private information:

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We take this into consideration because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from physical harm as a result of their information being shared, so we consider information such as physical location and phone numbers to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly

available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

Sharing private media:

Posting images is an important part of our users' experience on Twitter. Where individuals have a reasonable expectation of privacy in an individual piece of media, we believe they should be able to determine whether or not it is shared. Sharing such media could potentially violate users' privacy and may lead to emotional or physical harm. When we are notified by individuals depicted, or their authorized representative, that they did not consent to having image shared, we will remove the media. This policy is not applicable to public figures.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;

- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details;
- other private information, including biometric data or medical records;
- media of private individuals without the permission of the person(s) depicted.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;

- descriptions of physical appearance;
- gossip, rumours, accusations, and allegations; and
- screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

For media, the following are not in violation of our policy:

- the media is publicly available or is being covered by mainstream media;
- the media and the accompanying tweet text add value to the public discourse or are shared in public interest;
- contains eyewitness accounts or on the ground reports from developing events;
- the subject of the media is a public figure.

Who can report violations of this policy?

Sharing private information:

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

Sharing private media:

When reporting private media, we need a first person report in order to make the determination that the image or video has been shared without their permission. We accept reports from:

1. Individuals depicted in the reported media or
2. Their authorized representative such as:
 - a. parents/legal guardians,

- b. lawyers, or
- c. legal representative

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

Sharing private information:

The first time you violate this policy by sharing private information (such as home address, identity documents etc.), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy by posting private information again after your first warning, your account will be permanently suspended.

Sharing private media:

If you violate this policy by sharing private media of an individual without their consent, we will require you to remove the content and temporarily lock your account while you remove the media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy **(20191217*)**

Overview

November 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepsshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy (20200918*)

Overview

September 2020

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are;
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting; and
- coordinated harmful activity that encourages or promotes behavior which violates the [Twitter Rules](#).

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts or by coordinating with others to violate the Twitter Rules. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and

- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate;
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account; and
 - coordinating with others to engage in or promote violations of the Twitter Rules, including violations of our [abusive behavior](#) policy.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in "follow trains," "decks," and "Retweet for Retweet" behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- "follow churn" – following and then unfollowing large numbers of accounts in an effort to inflate one's own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account's followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person's browser (malware) or computer or to compromise a person's privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally;
- coordinating with others to express ideas, viewpoints, support, or opposition towards a cause, provided such behavior does not result in violations of the Twitter Rules; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don't violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Denylisting URLs

We denylist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [financial scam policy](#), our [hacked materials policy](#), our approach to [coordinated harmful activity](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Civic integrity policy (20201028*)

Overview

October 2020

You may not use Twitter's services for the purpose of manipulating or interfering in elections or other civic processes. This includes posting or sharing content that may suppress participation or mislead people about when, where, or how to participate in a civic process. In addition, we may label and reduce the

visibility of Tweets containing false or misleading information about civic processes in order to provide additional context.

The public conversation occurring on Twitter is never more important than during elections and other civic events. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt civic processes, including through the distribution of false or misleading information about the procedures or circumstances around participation in a civic process. In instances where misleading information does not seek to directly manipulate or disrupt civic processes, but leads to confusion on our service, we may label the Tweets to give additional context.

What is a civic process?

Twitter considers civic processes to be events or procedures mandated, organized, and conducted by the governing and/or electoral body of a country, state, region, district, or municipality to address a matter of common concern through public participation. Some examples of civic processes may include:

- Political elections
- Censuses
- Major referenda and ballot initiatives

What is in violation of this policy?

This policy addresses 4 categories of misleading behavior and content:

Misleading information about how to participate

We will label or remove false or misleading information about how to participate in an election or other civic process. This includes but is not limited to:

- misleading information about procedures to participate in a civic process (for example, that you can vote by Tweet, text message, email, or phone call in jurisdictions where these are not a possibility);
- misleading information about requirements for participation, including identification or citizenship requirements;
- misleading claims that cause confusion about the established laws, regulations, procedures, and methods of a civic process, or about the actions of officials or entities executing those civic processes; and
- misleading statements or information about the official, announced date or time of a civic process.

Suppression and intimidation

We will label or remove false or misleading information intended to intimidate or dissuade people from participating in an election or other civic process. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to voting in an election, polling places, or collecting census information;
- misleading claims about long lines, equipment problems, or other disruptions at voting locations during election periods;
- misleading claims about process procedures or techniques which could dissuade people from participating; and
- threats regarding voting locations or other key places or events (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

Misleading information about outcomes

We will label or remove false or misleading information intended to undermine public confidence in an election or other civic process. This includes but is not limited to:

- disputed claims that could undermine faith in the process itself, such as unverified information about election rigging, ballot tampering, vote tallying, or certification of election results; and
- misleading claims about the results or outcome of a civic process which calls for or could lead to interference with the implementation of the results of the process, e.g. claiming victory before election results have been certified, inciting unlawful conduct to prevent the procedural or practical implementation of election results (note that our violent threats policy may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account policy](#).

What is not a violation of this policy?

Not all false or untrue information about politics or civic processes constitutes manipulation or interference. In the absence of other policy violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected or appointed official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information;
- voting and audience participation for competitions, game shows, or other entertainment purposes; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election or civic process. Therefore, we enable reporting of false or misleading information about civic processes in advance of major events, for people located in the relevant countries and locations. We also work with select government and civil society partners in these countries to provide additional channels for reporting and expedited review.

For civic processes with multiple stages or parts, such as primary elections or lengthy campaigns, reporting will be enabled leading up to the first officially-sanctioned event associated with the civic process.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country at the relevant time, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in civic processes.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in a civic process.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our civic integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. The actions we take may include the following:

Tweet deletion

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Labeling

In circumstances where we do not remove content which violates this policy, we may provide additional context on Tweets sharing the content where they appear on Twitter. This means we may:

- Apply a label and/or warning message to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended; and/or
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or relevant Twitter policies.

In most cases, we will take all of the above actions on Tweets we label. In some instances, we'll also turn off your ability to reply, Retweet, or like the Tweet. We prioritize producing Twitter Moments in cases where misleading content on Twitter is gaining significant attention and has caused public confusion on our service.

Permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Impersonation policy (20191001*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Synthetic and manipulated media policy

(20210626*)

Overview

You may not deceptively promote synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context.

You should be able to find reliable information on Twitter. That means understanding whether the content you see is real or fabricated and having the ability to find more context about what you see on Twitter. Therefore, we may label Tweets that include media (videos, audio, and images) that have been deceptively altered or fabricated. In addition, you may not share deceptively altered media on Twitter in ways that mislead or deceive people about the media's authenticity where threats to physical safety or other serious harm may result.

We use the following criteria as we consider Tweets and media for labeling or removal under this policy as part of our ongoing work to enforce our rules and ensure healthy and safe conversation on Twitter (additional information is available below):

1. Is the content synthetic or manipulated?

In order for content to be labeled or removed under this policy, we must have reason to believe that media, or the context in which media are presented, are significantly and deceptively altered or manipulated. Synthetic and manipulated media take many different forms and people can employ a wide range of technologies to produce these media. In assessing whether media have been significantly and deceptively altered or fabricated, some of the factors we consider include:

- whether the content has been substantially edited in a manner that fundamentally alters its composition, sequence, timing, or framing;
- any visual or auditory information (such as new video frames, overdubbed audio, or modified subtitles) that has been added or removed; and
- whether media depicting a real person have been fabricated or simulated

We are most likely to take action (either labeling or removal, as described below) on more significant forms of alteration, such as wholly synthetic audio or video or content that has been doctored (spliced and reordered, slowed down) to change its meaning. Subtler forms of manipulated media, such as isolative editing, omission of context, or presentation with false context, may be labeled or removed on a case-by-case basis.

We will not take action to label or remove media that have been edited in ways that do not fundamentally alter their meaning, such as retouched photos or color-corrected videos.

In order to determine if media have been significantly and deceptively altered or fabricated, we may use our own technology or receive reports through partnerships with third parties. In situations where we are unable to reliably determine if media have been altered or fabricated, we may not take action to label or remove them.

2. Is the content shared in a deceptive manner?

We also consider whether the context in which media are shared could result in confusion or misunderstanding or suggests a deliberate intent to deceive people about the nature or origin of the content, for example by falsely claiming that it depicts reality. We assess the context provided alongside media to see whether it makes clear that the media have been altered or fabricated. Some of the types of context we assess in order to make this determination include:

- The text of the Tweet accompanying or within media
- Metadata associated with media
- Information on the profile of the account sharing media
- Websites linked in the Tweet, or in the profile of the account sharing media

3. Is the content likely to impact public safety or cause serious harm?

Tweets that share synthetic and manipulated media are subject to removal under this policy if they are likely to cause serious harm. Some specific harms we consider include:

- Threats to the physical safety of a person or group
- Risk of mass violence or widespread civil unrest
- Threats to the privacy or ability of a person or group to freely express themselves or participate in civic events, such as:
 - Stalking or unwanted and obsessive attention
 - Targeted content that includes tropes, epithets, or material that aims to silence someone
 - Voter suppression or intimidation

While we have other rules also intended to address these forms of harm, including our policies on violent threats, election integrity, and hateful conduct, we will err toward removal in borderline cases that might otherwise not violate existing rules for Tweets that include synthetic or manipulated media.

We also consider the time frame within which the content may be likely to impact public safety or cause serious harm, and are more likely to remove content under this policy if we find that immediate harms are likely to result from the content's presence on Twitter.

Note: We may also take action on synthetic and manipulated content under our [non-consensual nudity policy](#) (such as pornographic media altered to insert the faces of people not actually involved) or other parts of the Twitter Rules.

Labeling and removal

In most cases, if we have reason to believe that media shared in a Tweet have been significantly and deceptively altered or fabricated, we will provide additional context on Tweets sharing the media where they appear on Twitter. This means we may:

- Apply a label to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended;
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or landing page; and/or
- Turn off likes, replies, and Retweets.

In most cases, we will take all of the above actions on Tweets we label.

Media that meet **all three of the criteria** defined above—i.e. that are synthetic or manipulated, shared in a deceptive manner, and is likely to cause harm—may not be shared on Twitter and are subject to removal. Accounts engaging in repeated or severe violations of this policy may be permanently suspended.

** Other parts of the Twitter Rules apply and may lead to the removal of the content, particularly where there is high likelihood of severe harm, such as a threat to someone's life or physical safety.*

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Copyright policy (20200917*)

What types of copyright complaints does Twitter respond to?

Twitter responds to copyright complaints submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review [Twitter's trademark policy](#). If you are concerned about a parody, newsfeed, commentary, or fan account, please see the relevant policy [here](#). These are generally not copyright issues.

Am I a copyright holder? How do I know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

What to consider before submitting a copyright complaint

Before submitting a copyright complaint to us, please consider whether or not the use could be considered [fair use](#).

If you have considered fair use, and you still wish to continue with a copyright complaint, you may want to first reach out to the user in question to see if you can resolve the matter directly with the user. You can reply to the

user's Tweet or send the user a Direct Message and ask for them to remove your copyrighted content without having to contact Twitter.

Prior to submitting a formal complaint with Twitter, please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

Note: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph. If you're unsure whether or not you own the copyrights to a work, or if you're infringing upon someone else's work, please consult an attorney or another advisor.

What information do you need to process a copyright complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the complaint is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How do I file a copyright complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and [filing a copyright complaint](#). If you are logged in to [twitter.com](#), you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a full copy of your notice (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a copyright complaint.

How are claims processed?

We process copyright complaints in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. If you do not receive a ticket confirmation that means we did not receive your complaint and you should re-submit your complaint. However, please note, submitting duplicate copyright complaints will result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) and provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice. We will also forward a redacted copy of the complaint to [Lumen](#), with your personal information removed.

What information gets forwarded to the reported user(s)?

If we remove or disable access to the materials reported in a copyright complaint, the reported user(s) will receive a copy of the complaint, including the reporter's full name, email, street address, and any other information included in the complaint.

If you are uncomfortable sharing your contact information with the reported user(s), you may wish to consider appointing an agent to submit your DMCA notice on your behalf. Your agent will be required to submit the DMCA notice with valid contact information, and identify you as the content owner that they are representing.

What happens next?

Twitter's response to copyright complaints may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a copyright complaint, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the complaint, along with instructions for filing a counter-notice.

If you've not yet received a copy of the copyright complaint regarding the content removed from your account, please respond to the support ticket we sent you. If your complaint involves removal of content from Periscope, please see this help center [article](#) regarding how to request a copy of your complaint.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

My content was removed from Twitter

Why did I receive a copyright complaint?

If you receive a copyright complaint, it means that access to the content described in the complaint has been restricted. Please take the time to read through our correspondence to you, which includes information on the

complaint we received as well as instructions on how to file a counter-notice. Please ensure that you are monitoring the email address associated with your Twitter account.

Tip: Removing the material reported in a copyright complaint will not resolve that complaint.

What if I want to contest the takedown?

If you believe that the materials reported in the copyright complaints were misidentified or removed in error, you may send us a counter-notification(s). A counter-notice is a request for Twitter to reinstate the removed material, and it has legal consequences. Alternatively, you may be able to seek a retraction of the copyright complaint from the reporter.

How do I seek a retraction?

The DMCA complaint you received includes the contact information of the reporter. You may want to reach out and ask them to retract their notice. The reporter can send retractions to copyright@twitter.com, and should include: (1) identification of the material that was disabled, and (2) a statement that the reporter would like to retract their DMCA notice. This is the fastest and most efficient means of resolving an unresolved copyright complaint. A retraction is at the sole discretion of the original reporter.

When should I file a counter-notice?

A counter-notice is a request for Twitter to reinstate the removed material, and is the start of a legal process that has legal consequences. For example, submitting a counter notice indicates that you consent to the jurisdiction of a U.S. Federal court and that you consent to the disclosure of your personal information to the reporter and Lumen website.

With these considerations in mind, you may file a counter-notice if you believe that this material was misidentified, or you have a good faith belief that the material should not have been removed. If you're unsure whether or not you should file a counter-notice, you may want to consult with an attorney.

Tip: Re-posting material removed in response to a copyright complaint may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notice rather than re-posting the material.

What information do you need to process a counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What happens after I submit a counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. This means that the contact information that is submitted in your counter-notice will be shared to the person who filed the original notice.

If the copyright owner disagrees that the content was removed in error or misidentification, they may pursue legal action against you. If we do not receive notice within 10 business days that the original reporter is seeking a court

order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a copyright complaint or counter-notice is serious business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

What happens if my account receives multiple copyright complaints?

If multiple copyright complaints are received Twitter may lock accounts or take other actions to warn repeat violators. These warnings may vary across Twitter's services. Under appropriate circumstances we may suspend user accounts under our repeat infringer policy. However, we may take retractions and counter-notices into account when applying our repeat infringer policy.

Our approach to policy development and enforcement philosophy **(20190520*)**

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet

does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (20200204*)

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Synthetic and manipulated media: You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.

5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Violent organizations policy (20201026*)

Overview

October 2020

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

Other violent organizations

Other violent organizations are those that meet all of the below criteria:

- a collection of individuals with a shared purpose; and
- have systematically targeted civilians with violence.

We examine a group's activities both on and off Twitter to determine whether they satisfy the above criteria.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a violent organization;
- recruiting for a violent organization;
- providing or distributing services (e.g., financial, media/propaganda) to further a violent organization's stated goals; and
- using the insignia or symbol of violent organizations to promote them or indicate affiliation or support.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to state or governmental organizations.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.

6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wishing for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy (20201012*)

Overview

October 2020

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives;
- identifying alleged victims of childhood sexual exploitation by name or image; and
- promoting or normalizing sexual attraction to minors as a form of identity or sexual orientation.

What is not a violation of this policy?

- Discussions related to child sexual exploitation are permitted, provided they don't normalise, promote or glorify child sexual exploitation **in any way**;
- Conversations about help-seeking behaviour of individuals who may struggle with an attraction to minors;
- Advocacy against illegal or harmful activity involving minors provided there is no sharing of, or linking to any material featuring child sexual exploitation; and
- Depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios (e.g., art works by internationally renowned artists that feature minors, news media reportage, media created and shared for scientific or educational purposes).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

Abusive behavior (20210527*)

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing, hoping, or calling for serious harm on a person or group of people

We do not tolerate content that wishes, hopes, promotes, incites, or expresses a desire for death, serious bodily harm or serious disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

About wishes of harm exceptions on Twitter

We recognize that conversations regarding certain individuals credibly accused of severe violence may prompt outrage and associated wishes of harm. In these limited cases, we will request the user to delete the Tweet without any risk of account penalty, strike, or suspension. Examples are, but not limited to:

- “I wish all rapists to die.”
- “Child abusers should be hanged.”

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone’s body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Denying mass casualty events took place

We prohibit content that denies that mass murder or other mass casualty events took place, where we can verify that the event occurred, and when the content is shared with abusive intent. This may include references to such an event as a “hoax” or claims that victims or survivors are fake or “actors.” It includes, but is not limited to, events like the Holocaust, school shootings, terrorist attacks, and natural disasters.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual’s previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy (20210127*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognize that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and more harmful.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals with abuse based on protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., "I will kill you."

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, incites, or expresses a desire for death, serious bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

- Hoping that an entire protected category and/or individuals who may be members of that category dies as a result of a serious disease, e.g., "I hope all [nationality] get COVID and die."
- Wishing for someone to fall victim to a serious accident, e.g., "I wish that you would get run over by a car next time you run your mouth."

- Saying that a group of individuals deserve serious physical injury, e.g., "If this group of [slur] don't shut up, they deserve to be shot."
- Encouraging others to commit violence against an individual or a group based on their perceived membership in a protected category, e.g., "I'm in the mood to punch a [racial slur], who's with me?"

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to sending someone:

- media that depicts victims of the Holocaust;
- media that depicts lynchings.

Incitement against protected categories

We prohibit inciting behavior that targets individuals or groups of people belonging to protected categories. This includes content intended:

- to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., "all [religious group] are terrorists."
- to incite others to harass members of a protected category on or off platform, e.g., "I'm sick of these [religious group] thinking they are better than us, if any of you see someone wearing a [religious symbol of the religious group], grab it off them and post pics!"
- to incite others to discriminate in the form of denial of support to the economic enterprise of an individual or group because of their perceived membership in a protected category, e.g., "If you go to a [religious group] store, you are supporting those [slur], let's stop giving our money to these [religious slur]." This may not include content intended as political in nature, such as political commentary or content relating to boycotts or protests.

Note that content intended to incite violence against a protected category is prohibited under **Wishing, hoping, or calling for serious harm on a person or groups of people.**

We prohibit targeting individuals and groups with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., "all [religious group] are terrorists."

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting individuals with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals.

We also prohibit the dehumanization of a group of people based on their religion, caste, age, disability, serious disease, national origin, race, or ethnicity.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent account suspension. If an account is engaging primarily in abusive behavior, or is deemed to have shared a violent threat, we will permanently suspend the account upon initial review.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

Suicide and Self-harm Policy **(20200805*)**

Overview

You may not promote or encourage suicide or self-harm.

At Twitter, we recognize that suicide and self-harm are significant social & public health challenges that require collaboration between all stakeholders – public, private, and civil society – and that we have a role and responsibility to help people access and receive support when they need it.

When developing this policy, we consulted extensively with experts to ensure that people who have engaged in self-harm or experienced suicidal thoughts can share their personal experiences. We also recognized the need to protect people from the potential harm caused by exposure to content that could promote or encourage self-harm – intentionally or inadvertently. That's why our policy prohibits content that promotes or encourages self-harming behaviors and provides support to those undergoing experiences with self-harm or suicidal thoughts.

What is in violation of this policy?

Under this policy, you can't promote, or otherwise encourage, suicide or self-harm. We define promotion and encouragement to include statements such as "the most effective", "the easiest", "the best", "the most successful", "you should", "why don't you". Violations of this policy can occur via Tweets, images or videos, including live video.

We define suicide to be the act of taking one's own life. We define self-harm to include:

- self-inflicted physical injuries e.g., cutting; and
- eating disorders e.g., bulimia, anorexia.

Violations of this policy include, but are not limited to:

- encouraging someone to physically harm or kill themselves;
- asking others for encouragement to engage in self-harm or suicide, including seeking partners for group suicides or suicide games; and
- sharing information, strategies, methods or instructions that would assist people to engage in self-harm and suicide.

What is not a violation of this policy?

Some examples of behavior that are not considered a violation of this policy include:

- sharing personal stories and experiences related to self-harm or suicide;
- sharing coping mechanisms and resources for addressing self-harm or suicidal thoughts; and
- discussions that are focused on research, advocacy, and education related to self-harm or suicide prevention.

Note: people can share their personal experiences, but should avoid sharing detailed information about specific strategies or methods related to self-harm, as this could inadvertently encourage this behavior.

Who can report violations of this policy?

Anyone can report content that may encourage or promote suicide or self-harm via our in-app reporting or our [specialized reporting form](#). These reports are routed to a dedicated team who evaluate each case individually.

Note: if we receive a report that someone has expressed an intention to engage in self-harm or suicide, we will contact them directly, encourage them to seek support, and provide information about dedicated online and hotline resources. We may also work with law enforcement officials where appropriate, for example, if we receive a valid emergency disclosure request as defined in our [Law Enforcement Guidelines](#).

How to report violations of this policy

To ensure that we handle reports sensitively, our in-app reporting provides separate options for people who may be expressing an intention to harm themselves, and content that is encouraging or promoting self-harm or suicide.

Expressing intentions of self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Reporting form

You can also report this content for review via [our dedicated report form](#).

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Encouraging self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

What happens if you violate this policy?

Our enforcement approach depends on the type of content being shared, whether or not the reported account is encouraging or promoting self-harm or suicide, and the account's previous history of violations.

If you violate this policy by sharing content that intentionally encourages others to harm themselves, ask others to encourage you to harm yourself, or share detailed information or instructions related to self-harm or suicide methods, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy, or if your account is dedicated to promoting or encouraging self-harm or suicide, your account will be permanently suspended. If cases include images or videos related to self-harm or suicide, we will also evaluate this content under [our sensitive media policy](#). If you believe that your account was suspended in error, [you can submit an appeal](#).

We may also take steps to prevent the spread of instructional material hosted on third-party websites by [marking such links as unsafe](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Learn more](#) about how you can support someone experiencing thoughts of self-harm or suicide.

Visit our Safety Center for a list of local [mental health resources](#) and read [our blog on suicide prevention](#) to learn more about our work.

Sensitive media policy (20191204*)

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or

- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.

3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20210922*)

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private

information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;

- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;
 - descriptions of physical appearance;
 - gossip, rumours, accusations, and allegations; and
 - screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking

enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy **(20191217*)**

Overview

November 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy **(20200918*)**

Overview

September 2020

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are;
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting; and
- coordinated harmful activity that encourages or promotes behavior which violates the [Twitter Rules](#).

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts or by coordinating with others to violate the Twitter Rules. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate;
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account; and
 - coordinating with others to engage in or promote violations of the Twitter Rules, including violations of our [abusive behavior](#) policy.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in "follow trains," "decks," and "Retweet for Retweet" behavior); and

- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- "follow churn" – following and then unfollowing large numbers of accounts in an effort to inflate one's own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account's followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person's browser (malware) or computer or to compromise a person's privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally;
- coordinating with others to express ideas, viewpoints, support, or opposition towards a cause, provided such behavior does not result in violations of the Twitter Rules; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don't violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Denylisting URLs

We denylist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.

- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [financial scam policy](#), our [hacked materials policy](#), our approach to [coordinated harmful activity](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Civic integrity policy (20211006*)

Overview

October 2021

You may not use Twitter's services for the purpose of manipulating or interfering in elections or other civic processes. This includes posting or sharing content that may suppress participation or mislead people about when, where, or how to participate in a civic process. In addition, we may label and reduce the visibility of Tweets containing false or misleading information about civic processes in order to provide additional context.

The public conversation occurring on Twitter is never more important than during elections and other civic events. Any attempts to undermine the integrity of our service

is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt civic processes, including through the distribution of false or misleading information about the procedures or circumstances around participation in a civic process. In instances where misleading information does not seek to directly manipulate or disrupt civic processes, but leads to confusion on our service, we may label the Tweets to give additional context. Given the significant risks of confusion about key election information, we may take these actions even if Tweets contain (or attempt to contain) satirical or humorous elements.

What is a civic process?

Twitter considers civic processes to be events or procedures mandated, organized, and conducted by the governing and/or electoral body of a country, state, region, district, or municipality to address a matter of common concern through public participation. Some examples of civic processes may include:

- Political elections
- Censuses
- Major referenda and ballot initiatives

What is in violation of this policy?

This policy addresses 4 categories of misleading behavior and content:

Misleading information about how to participate

We will label or remove false or misleading information about how to participate in an election or other civic process. This includes but is not limited to:

- misleading information about procedures to participate in a civic process (for example, that you can vote by Tweet, text message, email, or phone call in jurisdictions where these are not a possibility);
- misleading information about requirements for participation, including identification or citizenship requirements;
- misleading claims that cause confusion about the established laws, regulations, procedures, and methods of a civic process, or about the actions of officials or entities executing those civic processes; and
- misleading statements or information about the official, announced date or time of a civic process.

Suppression and intimidation

We will label or remove false or misleading information intended to intimidate or dissuade people from participating in an election or other civic process. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to voting in an election, polling places, or collecting census information;
- misleading claims about long lines, equipment problems, or other disruptions at voting locations during election periods;
- misleading claims about process procedures or techniques which could dissuade people from participating; and
- threats regarding voting locations or other key places or events (note that our [violent threats policy](#) may also be relevant for threats not covered by this

policy).

Misleading information about outcomes

We will label or remove false or misleading information intended to undermine public confidence in an election or other civic process. This includes but is not limited to:

- disputed claims that could undermine faith in the process itself, such as unverified information about election rigging, ballot tampering, vote tallying, or certification of election results; and
- misleading claims about the results or outcome of a civic process which calls for or could lead to interference with the implementation of the results of the process, e.g. claiming victory before election results have been certified, inciting unlawful conduct to prevent the procedural or practical implementation of election results (note that our violent threats policy may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account policy.](#)

What is not a violation of this policy?

Not all false or untrue information about politics or civic processes constitutes manipulation or interference. In the absence of other policy violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected or appointed official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information;
- voting and audience participation for competitions, game shows, or other entertainment purposes; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election or civic process. Therefore, we enable reporting of false or misleading information about civic processes in advance of major events, for people located in the relevant countries and locations. We also work with select government and civil society partners in these countries to provide additional channels for reporting and expedited review.

For civic processes with multiple stages or parts, such as primary elections or lengthy campaigns, reporting will be enabled leading up to the first officially-sanctioned event associated with the civic process.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country at the relevant time, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in civic processes.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in a civic process.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our civic integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. In instances where accounts repeatedly violate this policy, we will use a strike system to determine if further enforcement actions should be applied. We believe this system further helps to reduce the spread of potentially harmful and misleading information on Twitter, particularly for high-severity violations of our rules.

The actions we take may include the following:

Tweet deletion

For high-severity violations of this policy, including (1) misleading information about how to participate, and (2) suppression and intimidation, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Tweet deletions accrue 2 strikes.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Labeling

In circumstances where we do not remove content which violates this policy, we may provide additional context on Tweets sharing the content where they appear on Twitter. This means we may:

- Apply a label and/or warning message to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Turn off people's ability to reply, Retweet, or like the Tweet;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended;
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or relevant Twitter policies; and/or
- Turn off likes, replies, and Retweets.

In most cases, we will take all of the above actions on Tweets we label. In some instances, we'll also turn off your ability to reply, Retweet, or like the Tweet. We

prioritize producing Twitter Moments in cases where misleading content on Twitter is gaining significant attention and has caused public confusion on our service. Labels applied to Tweets accrue 1 strike.

Account locks and permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

Repeated violations of this policy are enforced against on the basis of the number of strikes an account has accrued for violations of this policy:

- 1 strike: No account-level action
- 2 strikes: 12-hour account lock
- 3 strikes: 12-hour account lock
- 4 strikes: 7-day account lock
- 5 or more strikes: Permanent suspension

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Impersonation policy (20191001*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Synthetic and manipulated media policy

(20210626*)

Overview

You may not deceptively promote synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context.

You should be able to find reliable information on Twitter. That means understanding whether the content you see is real or fabricated and having the ability to find more context about what you see on Twitter. Therefore, we may label Tweets that include media (videos, audio, and images) that have been deceptively altered or fabricated. In addition, you may not share deceptively altered media on Twitter in ways that mislead or deceive people about the media's authenticity where threats to physical safety or other serious harm may result.

We use the following criteria as we consider Tweets and media for labeling or removal under this policy as part of our ongoing work to enforce our rules and ensure healthy and safe conversation on Twitter (additional information is available below):

1. Is the content synthetic or manipulated?

In order for content to be labeled or removed under this policy, we must have reason to believe that media, or the context in which media are presented, are significantly and

deceptively altered or manipulated. Synthetic and manipulated media take many different forms and people can employ a wide range of technologies to produce these media. In assessing whether media have been significantly and deceptively altered or fabricated, some of the factors we consider include:

- whether the content has been substantially edited in a manner that fundamentally alters its composition, sequence, timing, or framing;
- any visual or auditory information (such as new video frames, overdubbed audio, or modified subtitles) that has been added or removed; and
- whether media depicting a real person have been fabricated or simulated

We are most likely to take action (either labeling or removal, as described below) on more significant forms of alteration, such as wholly synthetic audio or video or content that has been doctored (spliced and reordered, slowed down) to change its meaning. Subtler forms of manipulated media, such as isolative editing, omission of context, or presentation with false context, may be labeled or removed on a case-by-case basis.

We will not take action to label or remove media that have been edited in ways that do not fundamentally alter their meaning, such as retouched photos or color-corrected videos.

In order to determine if media have been significantly and deceptively altered or fabricated, we may use our own technology or receive reports through partnerships with third parties. In situations where we are unable to reliably determine if media have been altered or fabricated, we may not take action to label or remove them.

2. Is the content shared in a deceptive manner?

We also consider whether the context in which media are shared could result in confusion or misunderstanding or suggests a deliberate intent to deceive people about the nature or origin of the content, for example by falsely claiming that it depicts reality. We assess the context provided alongside media to see whether it makes clear

that the media have been altered or fabricated. Some of the types of context we assess in order to make this determination include:

- The text of the Tweet accompanying or within media
- Metadata associated with media
- Information on the profile of the account sharing media
- Websites linked in the Tweet, or in the profile of the account sharing media

3. Is the content likely to impact public safety or cause serious harm?

Tweets that share synthetic and manipulated media are subject to removal under this policy if they are likely to cause serious harm. Some specific harms we consider include:

- Threats to the physical safety of a person or group
- Risk of mass violence or widespread civil unrest
- Threats to the privacy or ability of a person or group to freely express themselves or participate in civic events, such as:
 - Stalking or unwanted and obsessive attention
 - Targeted content that includes tropes, epithets, or material that aims to silence someone
 - Voter suppression or intimidation

While we have other rules also intended to address these forms of harm, including our policies on violent threats, election integrity, and hateful conduct, we will err toward removal in borderline cases that might otherwise not violate existing rules for Tweets that include synthetic or manipulated media.

We also consider the time frame within which the content may be likely to impact public safety or cause serious harm, and are more likely to remove content under this policy if we find that immediate harms are likely to result from the content's presence on Twitter.

Note: We may also take action on synthetic and manipulated content under our [non-consensual nudity policy](#) (such as pornographic media altered to insert the faces of people not actually involved) or other parts of the Twitter Rules.

Labeling and removal

In most cases, if we have reason to believe that media shared in a Tweet have been significantly and deceptively altered or fabricated, we will provide additional context on Tweets sharing the media where they appear on Twitter. This means we may:

- Apply a label to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended;
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or landing page; and/or
- Turn off likes, replies, and Retweets.

In most cases, we will take all of the above actions on Tweets we label.

Media that meet **all three of the criteria** defined above—i.e. that are synthetic or manipulated, shared in a deceptive manner, and is likely to cause harm—may not be shared on Twitter and are subject to removal. Accounts engaging in repeated or severe violations of this policy may be permanently suspended.

** Other parts of the Twitter Rules apply and may lead to the removal of the content, particularly where there is high likelihood of severe harm, such as a threat to someone's life or physical safety.*

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Copyright policy (20200917*)

What types of copyright complaints does Twitter respond to?

Twitter responds to copyright complaints submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review [Twitter's trademark policy](#). If you are concerned about a parody, newsfeed, commentary, or fan account, please see the relevant policy [here](#). These are generally not copyright issues.

Am I a copyright holder? How do I know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

What to consider before submitting a copyright complaint

Before submitting a copyright complaint to us, please consider whether or not the use could be considered [fair use](#).

If you have considered fair use, and you still wish to continue with a copyright complaint, you may want to first reach out to the user in question to see if you can resolve the matter directly with the user. You can reply to the user's Tweet or send the user a Direct Message and ask for them to remove your copyrighted content without having to contact Twitter.

Prior to submitting a formal complaint with Twitter, please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

Note: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph. If you're unsure whether or not you own the copyrights to a work, or if you're infringing upon someone else's work, please consult an attorney or another advisor.

What information do you need to process a copyright complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the complaint is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How do I file a copyright complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and [filing a copyright complaint](#). If you are logged in to [twitter.com](#), you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a full copy of your notice (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a copyright complaint.

How are claims processed?

We process copyright complaints in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. If you do not receive a ticket confirmation that means we did not receive your complaint and you should re-submit your complaint. However, please note, submitting duplicate copyright complaints will result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) and provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice. We will also forward a redacted copy of the complaint to [Lumen](#) with your personal information removed.

What information gets forwarded to the reported user(s)?

If we remove or disable access to the materials reported in a copyright complaint, the reported user(s) will receive a copy of the complaint, including the reporter's full name, email, street address, and any other information included in the complaint.

If you are uncomfortable sharing your contact information with the reported user(s), you may wish to consider appointing an agent to submit your DMCA notice on your behalf. Your agent will be required to submit the DMCA notice with valid contact information, and identify you as the content owner that they are representing.

What happens next?

Twitter's response to copyright complaints may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a copyright complaint, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the complaint, along with instructions for filing a counter-notice.

If you've not yet received a copy of the copyright complaint regarding the content removed from your account, please respond to the support ticket we sent you. If your complaint involves removal of content from Periscope, please see this help center [article](#) regarding how to request a copy of your complaint.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

My content was removed from Twitter

Why did I receive a copyright complaint?

If you receive a copyright complaint, it means that access to the content described in the complaint has been restricted. Please take the time to read through our correspondence to you, which includes information on the complaint we received as well as instructions on how to file a counter-notice. Please ensure that you are monitoring the email address associated with your Twitter account.

Tip: Removing the material reported in a copyright complaint will not resolve that complaint.

What if I want to contest the takedown?

If you believe that the materials reported in the copyright complaints were misidentified or removed in error, you may send us a counter-notification(s). A counter-notice is a request for Twitter to reinstate the removed material, and it has legal consequences. Alternatively, you may be able to seek a retraction of the copyright complaint from the reporter.

How do I seek a retraction?

The DMCA complaint you received includes the contact information of the reporter. You may want to reach out and ask them to retract their notice. The reporter can send retractions to copyright@twitter.com, and should include: (1) identification of the material that was disabled, and (2) a statement that the reporter would like to retract their DMCA notice. This is the fastest and most efficient means of resolving an unresolved copyright complaint. A retraction is at the sole discretion of the original reporter.

When should I file a counter-notice?

A counter-notice is a request for Twitter to reinstate the removed material, and is the start of a legal process that has legal consequences. For example, submitting a counter notice indicates that you consent to the jurisdiction of a U.S. Federal court and that you consent to the disclosure of your personal information to the reporter and Lumen website.

With these considerations in mind, you may file a counter-notice if you believe that this material was misidentified, or you have a good faith belief that the material should not have been removed. If you're unsure whether or not you should file a counter-notice, you may want to consult with an attorney.

Tip: Re-posting material removed in response to a copyright complaint may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notice rather than re-posting the material.

What information do you need to process a counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What happens after I submit a counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. This means that the contact information that is submitted in your counter-notice will be shared to the person who filed the original notice.

If the copyright owner disagrees that the content was removed in error or misidentification, they may pursue legal action against you. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a copyright complaint or counter-notice is serious business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

What happens if my account receives multiple copyright complaints?

If multiple copyright complaints are received Twitter may lock accounts or take other actions to warn repeat violators. These warnings may vary across Twitter's services. Under appropriate circumstances we may suspend user accounts under our repeat infringer policy. However, we may take retractions and counter-notices into account when applying our repeat infringer policy.

Our approach to policy development and enforcement philosophy

(20190520*)

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a

larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (20200204*)

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Synthetic and manipulated media: You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.

5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Violent organizations policy (20201026*)

Overview

October 2020

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

Other violent organizations

Other violent organizations are those that meet all of the below criteria:

- a collection of individuals with a shared purpose; and
- have systematically targeted civilians with violence.

We examine a group's activities both on and off Twitter to determine whether they satisfy the above criteria.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a violent organization;
- recruiting for a violent organization;
- providing or distributing services (e.g., financial, media/propaganda) to further a violent organization's stated goals; and
- using the insignia or symbol of violent organizations to promote them or indicate affiliation or support.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to state or governmental organizations.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.

6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wishing for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy (20201012*)

Overview

October 2020

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives;
- identifying alleged victims of childhood sexual exploitation by name or image; and
- promoting or normalizing sexual attraction to minors as a form of identity or sexual orientation.

What is not a violation of this policy?

- Discussions related to child sexual exploitation are permitted, provided they don't normalise, promote or glorify child sexual exploitation **in any way**;
- Conversations about help-seeking behaviour of individuals who may struggle with an attraction to minors;
- Advocacy against illegal or harmful activity involving minors provided there is no sharing of, or linking to any material featuring child sexual exploitation; and
- Depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios (e.g., art works by internationally renowned artists that feature minors, news media reportage, media created and shared for scientific or educational purposes).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

Abusive behavior (20210527*)

Twitter Rules: You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Wishing, hoping, or calling for serious harm on a person or group of people

We do not tolerate content that wishes, hopes, promotes, incites, or expresses a desire for death, serious bodily harm or serious disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., "I hope you get cancer and die."
- Wishing for someone to fall victim to a serious accident e.g., "I wish that you would get run over by a car next time you run your mouth."
- Saying that a group of individuals deserves serious physical injury e.g., "If this group of protesters don't shut up, they deserve to be shot."

About wishes of harm exceptions on Twitter

We recognize that conversations regarding certain individuals credibly accused of severe violence may prompt outrage and associated wishes of harm. In these limited cases, we will request the user to delete the Tweet without any risk of account penalty, strike, or suspension. Examples are, but not limited to:

- “I wish all rapists to die.”
- “Child abusers should be hanged.”

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone’s body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using aggressive insults with the purpose of harassing or intimidating others

We take action against excessively aggressive insults that target an individual, including content that contains slurs or similar language. Please also note that while some individuals may find certain terms to be offensive, we will not action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Denying mass casualty events took place

We prohibit content that denies that mass murder or other mass casualty events took place, where we can verify that the event occurred, and when the content is shared with abusive intent. This may include references to such an event as a “hoax” or claims that victims or survivors are fake or “actors.” It includes, but is not limited to, events like the Holocaust, school shootings, terrorist attacks, and natural disasters.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual’s previous record of rule violations.

For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension. If an account is engaging primarily in abusive behavior, we may permanently suspend the account upon initial review.

Learn more about [our range of enforcement options](#).

Hateful conduct policy (20211010*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognize that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and more harmful.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically

marginalized. For this reason, we prohibit behavior that targets individuals or groups with abuse based on their perceived membership in a protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., “I will kill you.”

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, incites, or expresses a desire for death, serious bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

- Hoping that an entire protected category and/or individuals who may be members of that category dies as a result of a serious disease, e.g., “I hope all [nationality] get COVID and die.”

- Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”
- Saying that a group of individuals deserve serious physical injury, e.g., “If this group of [slur] don’t shut up, they deserve to be shot.”
- Encouraging others to commit violence against an individual or a group based on their perceived membership in a protected category, e.g., “I’m in the mood to punch a [racial slur], who’s with me?”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals or groups with content that references forms of violence or violent events where a protected category was the primary target or victims, where the intent is to harass. This includes, but is not limited to media or text that refers to or depicts:

- genocides, (e.g., the Holocaust);
- lynchings.

Incitement against protected categories

We prohibit inciting behavior that targets individuals or groups of people belonging to protected categories. This includes content intended:

- to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists.”
- to incite others to harass members of a protected category on or off platform, e.g., “I’m sick of these [religious group] thinking they are better than us, if any of you see someone wearing a [religious symbol of the religious group], grab it off them and post pics!”

- to incite others to discriminate in the form of denial of support to the economic enterprise of an individual or group because of their perceived membership in a protected category, e.g., “If you go to a [religious group] store, you are supporting those [slur], let’s stop giving our money to these [religious slur].” This may not include content intended as political in nature, such as political commentary or content relating to boycotts or protests.

Note that content intended to incite violence against a protected category is prohibited under **Wishing, hoping, or calling for serious harm on a person or groups of people**.

We prohibit targeting individuals and groups with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists.”

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting others with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals. We also prohibit the dehumanization of a group of people based on their religion, caste, age, disability, serious disease, national origin, race, or ethnicity. In some cases, such as (but not limited to) severe, repetitive usage of slurs, epithets, or racist/sexy tropes where the primary intent is to harass or intimidate others, we may require Tweet removal. In other cases, such as (but not limited to) moderate, isolated usage where the primary intent is to harass or intimidate others, we may limit Tweet visibility as further described below.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected

category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous record of rule violations. The following is a list of potential enforcement options for content that violates this policy:

- Downranking Tweets in replies, except when the user follows the Tweet author.
- Making Tweets ineligible for amplification in Top search results and/or on timelines for users who don't follow the Tweet author.
- Excluding Tweets and/or accounts in email or in-product recommendations.
- Requiring Tweet removal.
 - For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension.
- Suspending accounts whose primary use we've determined is to engage in hateful conduct as defined in this policy, or who have shared violent threats.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

Suicide and Self-harm Policy

(20200805*)

Overview

You may not promote or encourage suicide or self-harm.

At Twitter, we recognize that suicide and self-harm are significant social & public health challenges that require collaboration between all stakeholders – public, private, and civil society – and that we have a role and responsibility to help people access and receive support when they need it.

When developing this policy, we consulted extensively with experts to ensure that people who have engaged in self-harm or experienced suicidal thoughts can share their personal experiences. We also recognized the need to protect people from the potential harm caused by exposure to content that could promote or encourage self-harm – intentionally or inadvertently. That's why our policy prohibits content that promotes or encourages self-harming behaviors and provides support to those undergoing experiences with self-harm or suicidal thoughts.

What is in violation of this policy?

Under this policy, you can't promote, or otherwise encourage, suicide or self-harm. We define promotion and encouragement to include statements such as "the most effective", "the easiest", "the best", "the most successful", "you should", "why don't you". Violations of this policy can occur via Tweets, images or videos, including live video.

We define suicide to be the act of taking one's own life. We define self-harm to include:

- self-inflicted physical injuries e.g., cutting; and
- eating disorders e.g., bulimia, anorexia.

Violations of this policy include, but are not limited to:

- encouraging someone to physically harm or kill themselves;
- asking others for encouragement to engage in self-harm or suicide, including seeking partners for group suicides or suicide games; and
- sharing information, strategies, methods or instructions that would assist people to engage in self-harm and suicide.

What is not a violation of this policy?

Some examples of behavior that are not considered a violation of this policy include:

- sharing personal stories and experiences related to self-harm or suicide;
- sharing coping mechanisms and resources for addressing self-harm or suicidal thoughts; and
- discussions that are focused on research, advocacy, and education related to self-harm or suicide prevention.

Note: people can share their personal experiences, but should avoid sharing detailed information about specific strategies or methods related to self-harm, as this could inadvertently encourage this behavior.

Who can report violations of this policy?

Anyone can report content that may encourage or promote suicide or self-harm via our in-app reporting or our [specialized reporting form](#). These reports are routed to a dedicated team who evaluate each case individually.

Note: if we receive a report that someone has expressed an intention to engage in self-harm or suicide, we will contact them directly, encourage them to seek support, and provide information about dedicated online and hotline resources. We may also work with law enforcement officials where appropriate, for example, if we receive a valid emergency disclosure request as defined in our [Law Enforcement Guidelines](#).

How to report violations of this policy

To ensure that we handle reports sensitively, our in-app reporting provides separate options for people who may be expressing an intention to harm themselves, and content that is encouraging or promoting self-harm or suicide.

Expressing intentions of self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Reporting form

You can also report this content for review via [our dedicated report form](#).

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Encouraging self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

What happens if you violate this policy?

Our enforcement approach depends on the type of content being shared, whether or not the reported account is encouraging or promoting self-harm or suicide, and the account's previous history of violations.

If you violate this policy by sharing content that intentionally encourages others to harm themselves, ask others to encourage you to harm yourself, or share detailed information or instructions related to self-harm or suicide methods, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy, or if your account is dedicated to promoting or encouraging self-harm or suicide, your account will be permanently suspended. If cases include images or videos related to self-harm or suicide, we will also evaluate this content under [our sensitive media policy](#). If you believe that your account was suspended in error, [you can submit an appeal](#).

We may also take steps to prevent the spread of instructional material hosted on third-party websites by [marking such links as unsafe](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Learn more](#) about how you can support someone experiencing thoughts of self-harm or suicide.

Visit our Safety Center for a list of local [mental health resources](#) and read [our blog on suicide prevention](#) to learn more about our work.

Sensitive media policy (20191204*)

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and
- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and
- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement

request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20210922*)

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;

- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;
 - descriptions of physical appearance;
 - gossip, rumours, accusations, and allegations; and
 - screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking

enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy **(20191217*)**

Overview

November 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy **(20200918*)**

Overview

September 2020

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are;
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting; and
- coordinated harmful activity that encourages or promotes behavior which violates the [Twitter Rules](#).

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts or by coordinating with others to violate the Twitter Rules. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate;
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account; and
 - coordinating with others to engage in or promote violations of the Twitter Rules, including violations of our [abusive behavior](#) policy.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in "follow trains," "decks," and "Retweet for Retweet" behavior); and

- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- "follow churn" – following and then unfollowing large numbers of accounts in an effort to inflate one's own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account's followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person's browser (malware) or computer or to compromise a person's privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally;
- coordinating with others to express ideas, viewpoints, support, or opposition towards a cause, provided such behavior does not result in violations of the Twitter Rules; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don't violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Denylisting URLs

We denylist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.

- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [financial scam policy](#), our [hacked materials policy](#), our approach to [coordinated harmful activity](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Civic integrity policy (20211006*)

Overview

October 2021

You may not use Twitter's services for the purpose of manipulating or interfering in elections or other civic processes. This includes posting or sharing content that may suppress participation or mislead people about when, where, or how to participate in a civic process. In addition, we may label and reduce the visibility of Tweets containing false or misleading information about civic processes in order to provide additional context.

The public conversation occurring on Twitter is never more important than during elections and other civic events. Any attempts to undermine the integrity of our service

is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt civic processes, including through the distribution of false or misleading information about the procedures or circumstances around participation in a civic process. In instances where misleading information does not seek to directly manipulate or disrupt civic processes, but leads to confusion on our service, we may label the Tweets to give additional context. Given the significant risks of confusion about key election information, we may take these actions even if Tweets contain (or attempt to contain) satirical or humorous elements.

What is a civic process?

Twitter considers civic processes to be events or procedures mandated, organized, and conducted by the governing and/or electoral body of a country, state, region, district, or municipality to address a matter of common concern through public participation. Some examples of civic processes may include:

- Political elections
- Censuses
- Major referenda and ballot initiatives

What is in violation of this policy?

This policy addresses 4 categories of misleading behavior and content:

Misleading information about how to participate

We will label or remove false or misleading information about how to participate in an election or other civic process. This includes but is not limited to:

- misleading information about procedures to participate in a civic process (for example, that you can vote by Tweet, text message, email, or phone call in jurisdictions where these are not a possibility);
- misleading information about requirements for participation, including identification or citizenship requirements;
- misleading claims that cause confusion about the established laws, regulations, procedures, and methods of a civic process, or about the actions of officials or entities executing those civic processes; and
- misleading statements or information about the official, announced date or time of a civic process.

Suppression and intimidation

We will label or remove false or misleading information intended to intimidate or dissuade people from participating in an election or other civic process. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to voting in an election, polling places, or collecting census information;
- misleading claims about long lines, equipment problems, or other disruptions at voting locations during election periods;
- misleading claims about process procedures or techniques which could dissuade people from participating; and
- threats regarding voting locations or other key places or events (note that our [violent threats policy](#) may also be relevant for threats not covered by this

policy).

Misleading information about outcomes

We will label or remove false or misleading information intended to undermine public confidence in an election or other civic process. This includes but is not limited to:

- disputed claims that could undermine faith in the process itself, such as unverified information about election rigging, ballot tampering, vote tallying, or certification of election results; and
- misleading claims about the results or outcome of a civic process which calls for or could lead to interference with the implementation of the results of the process, e.g. claiming victory before election results have been certified, inciting unlawful conduct to prevent the procedural or practical implementation of election results (note that our violent threats policy may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account policy.](#)

What is not a violation of this policy?

Not all false or untrue information about politics or civic processes constitutes manipulation or interference. In the absence of other policy violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected or appointed official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information;
- voting and audience participation for competitions, game shows, or other entertainment purposes; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election or civic process. Therefore, we enable reporting of false or misleading information about civic processes in advance of major events, for people located in the relevant countries and locations. We also work with select government and civil society partners in these countries to provide additional channels for reporting and expedited review.

For civic processes with multiple stages or parts, such as primary elections or lengthy campaigns, reporting will be enabled leading up to the first officially-sanctioned event associated with the civic process.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country at the relevant time, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in civic processes.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in a civic process.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our civic integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. In instances where accounts repeatedly violate this policy, we will use a strike system to determine if further enforcement actions should be applied. We believe this system further helps to reduce the spread of potentially harmful and misleading information on Twitter, particularly for high-severity violations of our rules.

The actions we take may include the following:

Tweet deletion

For high-severity violations of this policy, including (1) misleading information about how to participate, and (2) suppression and intimidation, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Tweet deletions accrue 2 strikes.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Labeling

In circumstances where we do not remove content which violates this policy, we may provide additional context on Tweets sharing the content where they appear on Twitter. This means we may:

- Apply a label and/or warning message to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Turn off people's ability to reply, Retweet, or like the Tweet;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended;
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or relevant Twitter policies; and/or
- Turn off likes, replies, and Retweets.

In most cases, we will take all of the above actions on Tweets we label. In some instances, we'll also turn off your ability to reply, Retweet, or like the Tweet. We

prioritize producing Twitter Moments in cases where misleading content on Twitter is gaining significant attention and has caused public confusion on our service. Labels applied to Tweets accrue 1 strike.

Account locks and permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

Repeated violations of this policy are enforced against on the basis of the number of strikes an account has accrued for violations of this policy:

- 1 strike: No account-level action
- 2 strikes: 12-hour account lock
- 3 strikes: 12-hour account lock
- 4 strikes: 7-day account lock
- 5 or more strikes: Permanent suspension

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Impersonation policy (20191001*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Synthetic and manipulated media policy

(20210626*)

Overview

You may not deceptively promote synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context.

You should be able to find reliable information on Twitter. That means understanding whether the content you see is real or fabricated and having the ability to find more context about what you see on Twitter. Therefore, we may label Tweets that include media (videos, audio, and images) that have been deceptively altered or fabricated. In addition, you may not share deceptively altered media on Twitter in ways that mislead or deceive people about the media's authenticity where threats to physical safety or other serious harm may result.

We use the following criteria as we consider Tweets and media for labeling or removal under this policy as part of our ongoing work to enforce our rules and ensure healthy and safe conversation on Twitter (additional information is available below):

1. Is the content synthetic or manipulated?

In order for content to be labeled or removed under this policy, we must have reason to believe that media, or the context in which media are presented, are significantly and

deceptively altered or manipulated. Synthetic and manipulated media take many different forms and people can employ a wide range of technologies to produce these media. In assessing whether media have been significantly and deceptively altered or fabricated, some of the factors we consider include:

- whether the content has been substantially edited in a manner that fundamentally alters its composition, sequence, timing, or framing;
- any visual or auditory information (such as new video frames, overdubbed audio, or modified subtitles) that has been added or removed; and
- whether media depicting a real person have been fabricated or simulated

We are most likely to take action (either labeling or removal, as described below) on more significant forms of alteration, such as wholly synthetic audio or video or content that has been doctored (spliced and reordered, slowed down) to change its meaning. Subtler forms of manipulated media, such as isolative editing, omission of context, or presentation with false context, may be labeled or removed on a case-by-case basis.

We will not take action to label or remove media that have been edited in ways that do not fundamentally alter their meaning, such as retouched photos or color-corrected videos.

In order to determine if media have been significantly and deceptively altered or fabricated, we may use our own technology or receive reports through partnerships with third parties. In situations where we are unable to reliably determine if media have been altered or fabricated, we may not take action to label or remove them.

2. Is the content shared in a deceptive manner?

We also consider whether the context in which media are shared could result in confusion or misunderstanding or suggests a deliberate intent to deceive people about the nature or origin of the content, for example by falsely claiming that it depicts reality. We assess the context provided alongside media to see whether it makes clear

that the media have been altered or fabricated. Some of the types of context we assess in order to make this determination include:

- The text of the Tweet accompanying or within media
- Metadata associated with media
- Information on the profile of the account sharing media
- Websites linked in the Tweet, or in the profile of the account sharing media

3. Is the content likely to impact public safety or cause serious harm?

Tweets that share synthetic and manipulated media are subject to removal under this policy if they are likely to cause serious harm. Some specific harms we consider include:

- Threats to the physical safety of a person or group
- Risk of mass violence or widespread civil unrest
- Threats to the privacy or ability of a person or group to freely express themselves or participate in civic events, such as:
 - Stalking or unwanted and obsessive attention
 - Targeted content that includes tropes, epithets, or material that aims to silence someone
 - Voter suppression or intimidation

While we have other rules also intended to address these forms of harm, including our policies on violent threats, election integrity, and hateful conduct, we will err toward removal in borderline cases that might otherwise not violate existing rules for Tweets that include synthetic or manipulated media.

We also consider the time frame within which the content may be likely to impact public safety or cause serious harm, and are more likely to remove content under this policy if we find that immediate harms are likely to result from the content's presence on Twitter.

Note: We may also take action on synthetic and manipulated content under our [non-consensual nudity policy](#) (such as pornographic media altered to insert the faces of people not actually involved) or other parts of the Twitter Rules.

Labeling and removal

In most cases, if we have reason to believe that media shared in a Tweet have been significantly and deceptively altered or fabricated, we will provide additional context on Tweets sharing the media where they appear on Twitter. This means we may:

- Apply a label to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended;
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or landing page; and/or
- Turn off likes, replies, and Retweets.

In most cases, we will take all of the above actions on Tweets we label.

Media that meet **all three of the criteria** defined above—i.e. that are synthetic or manipulated, shared in a deceptive manner, and is likely to cause harm—may not be shared on Twitter and are subject to removal. Accounts engaging in repeated or severe violations of this policy may be permanently suspended.

** Other parts of the Twitter Rules apply and may lead to the removal of the content, particularly where there is high likelihood of severe harm, such as a threat to someone's life or physical safety.*

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Copyright policy (20200917*)

What types of copyright complaints does Twitter respond to?

Twitter responds to copyright complaints submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review [Twitter's trademark policy](#). If you are concerned about a parody, newsfeed, commentary, or fan account, please see the relevant policy [here](#). These are generally not copyright issues.

Am I a copyright holder? How do I know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

What to consider before submitting a copyright complaint

Before submitting a copyright complaint to us, please consider whether or not the use could be considered [fair use](#).

If you have considered fair use, and you still wish to continue with a copyright complaint, you may want to first reach out to the user in question to see if you can resolve the matter directly with the user. You can reply to the user's Tweet or send the user a Direct Message and ask for them to remove your copyrighted content without having to contact Twitter.

Prior to submitting a formal complaint with Twitter, please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

Note: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph. If you're unsure whether or not you own the copyrights to a work, or if you're infringing upon someone else's work, please consult an attorney or another advisor.

What information do you need to process a copyright complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the complaint is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How do I file a copyright complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and [filing a copyright complaint](#). If you are logged in to [twitter.com](#), you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a full copy of your notice (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a copyright complaint.

How are claims processed?

We process copyright complaints in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. If you do not receive a ticket confirmation that means we did not receive your complaint and you should re-submit your complaint. However, please note, submitting duplicate copyright complaints will result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) and provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice. We will also forward a redacted copy of the complaint to [Lumen](#) with your personal information removed.

What information gets forwarded to the reported user(s)?

If we remove or disable access to the materials reported in a copyright complaint, the reported user(s) will receive a copy of the complaint, including the reporter's full name, email, street address, and any other information included in the complaint.

If you are uncomfortable sharing your contact information with the reported user(s), you may wish to consider appointing an agent to submit your DMCA notice on your behalf. Your agent will be required to submit the DMCA notice with valid contact information, and identify you as the content owner that they are representing.

What happens next?

Twitter's response to copyright complaints may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a copyright complaint, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the complaint, along with instructions for filing a counter-notice.

If you've not yet received a copy of the copyright complaint regarding the content removed from your account, please respond to the support ticket we sent you. If your complaint involves removal of content from Periscope, please see this help center [article](#) regarding how to request a copy of your complaint.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

My content was removed from Twitter

Why did I receive a copyright complaint?

If you receive a copyright complaint, it means that access to the content described in the complaint has been restricted. Please take the time to read through our correspondence to you, which includes information on the complaint we received as well as instructions on how to file a counter-notice. Please ensure that you are monitoring the email address associated with your Twitter account.

Tip: Removing the material reported in a copyright complaint will not resolve that complaint.

What if I want to contest the takedown?

If you believe that the materials reported in the copyright complaints were misidentified or removed in error, you may send us a counter-notification(s). A counter-notice is a request for Twitter to reinstate the removed material, and it has legal consequences. Alternatively, you may be able to seek a retraction of the copyright complaint from the reporter.

How do I seek a retraction?

The DMCA complaint you received includes the contact information of the reporter. You may want to reach out and ask them to retract their notice. The reporter can send retractions to copyright@twitter.com, and should include: (1) identification of the material that was disabled, and (2) a statement that the reporter would like to retract their DMCA notice. This is the fastest and most efficient means of resolving an unresolved copyright complaint. A retraction is at the sole discretion of the original reporter.

When should I file a counter-notice?

A counter-notice is a request for Twitter to reinstate the removed material, and is the start of a legal process that has legal consequences. For example, submitting a counter notice indicates that you consent to the jurisdiction of a U.S. Federal court and that you consent to the disclosure of your personal information to the reporter and Lumen website.

With these considerations in mind, you may file a counter-notice if you believe that this material was misidentified, or you have a good faith belief that the material should not have been removed. If you're unsure whether or not you should file a counter-notice, you may want to consult with an attorney.

Tip: Re-posting material removed in response to a copyright complaint may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notice rather than re-posting the material.

What information do you need to process a counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What happens after I submit a counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. This means that the contact information that is submitted in your counter-notice will be shared to the person who filed the original notice.

If the copyright owner disagrees that the content was removed in error or misidentification, they may pursue legal action against you. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a copyright complaint or counter-notice is serious business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

What happens if my account receives multiple copyright complaints?

If multiple copyright complaints are received Twitter may lock accounts or take other actions to warn repeat violators. These warnings may vary across Twitter's services. Under appropriate circumstances we may suspend user accounts under our repeat infringer policy. However, we may take retractions and counter-notices into account when applying our repeat infringer policy.

Our approach to policy development and enforcement philosophy

(20190520*)

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a

larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (20200204*)

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Synthetic and manipulated media: You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.

5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Violent organizations policy (20201026*)

Overview

October 2020

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

Other violent organizations

Other violent organizations are those that meet all of the below criteria:

- a collection of individuals with a shared purpose; and
- have systematically targeted civilians with violence.

We examine a group's activities both on and off Twitter to determine whether they satisfy the above criteria.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a violent organization;
- recruiting for a violent organization;
- providing or distributing services (e.g., financial, media/propaganda) to further a violent organization's stated goals; and
- using the insignia or symbol of violent organizations to promote them or indicate affiliation or support.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to state or governmental organizations.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.

6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wishing for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy (20201012*)

Overview

October 2020

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives;
- identifying alleged victims of childhood sexual exploitation by name or image; and
- promoting or normalizing sexual attraction to minors as a form of identity or sexual orientation.

What is not a violation of this policy?

- Discussions related to child sexual exploitation are permitted, provided they don't normalise, promote or glorify child sexual exploitation **in any way**;
- Conversations about help-seeking behaviour of individuals who may struggle with an attraction to minors;
- Advocacy against illegal or harmful activity involving minors provided there is no sharing of, or linking to any material featuring child sexual exploitation; and
- Depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios (e.g., art works by internationally renowned artists that feature minors, news media reportage, media created and shared for scientific or educational purposes).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

Abusive behavior (20211017*)

[Twitter Rules](#): You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it

may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., “I will kill you.”

Note: We have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account. Wishing, hoping, or calling for serious harm on a person or group of people.

Wishing, hoping, or calling for serious harm on a person or group of people

We do not tolerate content that wishes, hopes, promotes, incites, or expresses a desire for death, serious bodily harm or serious disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., “I hope you get cancer and die.”
- Wishing for someone to fall victim to a serious accident e.g., “I wish that you would get run over by a car next time you run your mouth.”

- Saying that a group of individuals deserves serious physical injury e.g., “If this group of protesters don’t shut up, they deserve to be shot.”

About wishes of harm exceptions on Twitter

We recognize that conversations regarding certain individuals credibly accused of severe violence may prompt outrage and associated wishes of harm. In these limited cases, we will request the user to delete the Tweet without any risk of account penalty, strike, or suspension. Examples are, but not limited to:

- “I wish all rapists to die.”
- “Child abusers should be hanged.”

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone’s body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using insults, profanity, or slurs with the purpose of harassing or intimidating others

We take action against the use of insults, profanity, or slurs to target others. In some cases, such as (but not limited to) severe, repetitive usage of insults or slurs where the primary intent is to harass or intimidate others, we may require Tweet removal. In other cases, such as (but not limited to) moderate, isolated usage of insults and profanity

where the primary intent is to harass or intimidate others, we may limit Tweet visibility as further described below. Please also note that while some individuals may find certain terms to be offensive, we will not take action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Denying mass casualty events took place

We prohibit content that denies that mass murder or other mass casualty events took place, where we can verify that the event occurred, and when the content is shared with abusive intent. This may include references to such an event as a “hoax” or claims that victims or survivors are fake or “actors.” It includes, but is not limited to, events like the Holocaust, school shootings, terrorist attacks, and natural disasters.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual’s previous record of rule violations. The following is a list of potential enforcement options for content that violates this policy:

- Downranking Tweets in replies, except when the user follows the Tweet author.
- Making Tweets ineligible for amplification in Top search results and/or on timelines for users who don't follow the Tweet author.
- Excluding Tweets and/or accounts in email or in-product recommendations.
- Requiring Tweet removal.
 - For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension.
- Suspending accounts whose primary use we've determined is to engage in abusive behavior as defined in this policy, or who have shared violent threats.

Learn more about [our range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

Hateful conduct policy (20211010*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognize that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and more harmful.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals or groups with abuse based on their perceived membership in a protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., “I will kill you.”

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, incites, or expresses a desire for death, serious bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

- Hoping that an entire protected category and/or individuals who may be members of that category dies as a result of a serious disease, e.g., “I hope all [nationality] get COVID and die.”
- Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”
- Saying that a group of individuals deserve serious physical injury, e.g., “If this group of [slur] don’t shut up, they deserve to be shot.”
- Encouraging others to commit violence against an individual or a group based on their perceived membership in a protected category, e.g., “I’m in the mood to punch a [racial slur], who’s with me?”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals or groups with content that references forms of violence or violent events where a protected category was the primary target or victims,

where the intent is to harass. This includes, but is not limited to media or text that refers to or depicts:

- genocides, (e.g., the Holocaust);
- lynchings.

Incitement against protected categories

We prohibit inciting behavior that targets individuals or groups of people belonging to protected categories. This includes content intended:

- to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists.”
- to incite others to harass members of a protected category on or off platform, e.g., “I’m sick of these [religious group] thinking they are better than us, if any of you see someone wearing a [religious symbol of the religious group], grab it off them and post pics!”
- to incite others to discriminate in the form of denial of support to the economic enterprise of an individual or group because of their perceived membership in a protected category, e.g., “If you go to a [religious group] store, you are supporting those [slur], let’s stop giving our money to these [religious slur].”
This may not include content intended as political in nature, such as political commentary or content relating to boycotts or protests.

Note that content intended to incite violence against a protected category is prohibited under **Wishing, hoping, or calling for serious harm on a person or groups of people**.

We prohibit targeting individuals and groups with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g.,

“all [religious group] are terrorists.”

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting others with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals. We also prohibit the dehumanization of a group of people based on their religion, caste, age, disability, serious disease, national origin, race, or ethnicity. In some cases, such as (but not limited to) severe, repetitive usage of slurs, epithets, or racist/sexy tropes where the primary intent is to harass or intimidate others, we may require Tweet removal. In other cases, such as (but not limited to) moderate, isolated usage where the primary intent is to harass or intimidate others, we may limit Tweet visibility as further described below.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous

record of rule violations. The following is a list of potential enforcement options for content that violates this policy:

- Downranking Tweets in replies, except when the user follows the Tweet author.
- Making Tweets ineligible for amplification in Top search results and/or on timelines for users who don't follow the Tweet author.
- Excluding Tweets and/or accounts in email or in-product recommendations.
- Requiring Tweet removal.
 - For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension.
- Suspending accounts whose primary use we've determined is to engage in hateful conduct as defined in this policy, or who have shared violent threats.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

Suicide and Self-harm Policy **(20200805*)**

Overview

You may not promote or encourage suicide or self-harm.

At Twitter, we recognize that suicide and self-harm are significant social & public health challenges that require collaboration between all stakeholders – public, private, and civil society – and that we have a role and responsibility to help people access and receive support when they need it.

When developing this policy, we consulted extensively with experts to ensure that people who have engaged in self-harm or experienced suicidal thoughts can share their personal experiences. We also recognized the need to protect people from the potential harm caused by exposure to content that could promote or encourage self-harm – intentionally or inadvertently. That's why our policy prohibits content that promotes or encourages self-harming behaviors and provides support to those undergoing experiences with self-harm or suicidal thoughts.

What is in violation of this policy?

Under this policy, you can't promote, or otherwise encourage, suicide or self-harm. We define promotion and encouragement to include statements such as "the most effective", "the easiest", "the best", "the most successful", "you should", "why don't you". Violations of this policy can occur via Tweets, images or videos, including live video.

We define suicide to be the act of taking one's own life. We define self-harm to include:

- self-inflicted physical injuries e.g., cutting; and
- eating disorders e.g., bulimia, anorexia.

Violations of this policy include, but are not limited to:

- encouraging someone to physically harm or kill themselves;
- asking others for encouragement to engage in self-harm or suicide, including seeking partners for group suicides or suicide games; and
- sharing information, strategies, methods or instructions that would assist people to engage in self-harm and suicide.

What is not a violation of this policy?

Some examples of behavior that are not considered a violation of this policy include:

- sharing personal stories and experiences related to self-harm or suicide;
- sharing coping mechanisms and resources for addressing self-harm or suicidal thoughts; and
- discussions that are focused on research, advocacy, and education related to self-harm or suicide prevention.

Note: people can share their personal experiences, but should avoid sharing detailed information about specific strategies or methods related to self-harm, as this could inadvertently encourage this behavior.

Who can report violations of this policy?

Anyone can report content that may encourage or promote suicide or self-harm via our in-app reporting or our [specialized reporting form](#). These reports are routed to a dedicated team who evaluate each case individually.

Note: if we receive a report that someone has expressed an intention to engage in self-harm or suicide, we will contact them directly, encourage them to seek support, and provide information about dedicated online and hotline resources. We may also work with law enforcement officials where appropriate, for example, if we receive a valid emergency disclosure request as defined in our [Law Enforcement Guidelines](#).

How to report violations of this policy

To ensure that we handle reports sensitively, our in-app reporting provides separate options for people who may be expressing an intention to harm themselves, and content that is encouraging or promoting self-harm or suicide.

Expressing intentions of self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu

2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Reporting form

You can also report this content for review via [our dedicated report form](#).

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Encouraging self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

What happens if you violate this policy?

Our enforcement approach depends on the type of content being shared, whether or not the reported account is encouraging or promoting self-harm or suicide, and the account's previous history of violations.

If you violate this policy by sharing content that intentionally encourages others to harm themselves, ask others to encourage you to harm yourself, or share detailed information or instructions related to self-harm or suicide methods, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy, or if your account is dedicated to promoting or encouraging self-harm or suicide, your account will be permanently suspended. If cases include images or videos related to self-harm or suicide, we will also evaluate this content under [our sensitive media policy](#). If you believe that your account was suspended in error, [you can submit an appeal](#).

We may also take steps to prevent the spread of instructional material hosted on third-party websites by [marking such links as unsafe](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Learn more](#) about how you can support someone experiencing thoughts of self-harm or suicide.

Visit our Safety Center for a list of local [mental health resources](#) and read [our blog on suicide prevention](#) to learn more about our work.

Sensitive media policy (20191204*)

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and

- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and

- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20211130)

Overview

November 2021

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

In addition, you may not share private media, such as images or videos of private individuals, without their consent. However, we recognise that there are instances where users may share images or videos of private individuals, who are not public figures, as part of a newsworthy event or to further public discourse on issues or events of public interest. In such cases, we may allow the media to remain on the platform.

Sharing someone's private information:

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private

information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We take this into consideration because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from potential physical harm as a result of their information being shared, so we consider information such as physical location and phone numbers to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

Sharing private media:

Posting images is an important part of our users' experience on Twitter. Where individuals have a reasonable expectation of privacy in an individual piece of media, we believe they should be able to determine whether or not it is shared. Sharing such media could potentially violate users' privacy and may lead to emotional or physical harm. When we are notified by individuals depicted, or their authorized representative, that they did not consent to having media shared, we will remove the media. This policy is not applicable to public figures.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;

- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.
- media of private individuals without the permission of the person(s) depicted.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;

- descriptions of physical appearance;
- gossip, rumours, accusations, and allegations; and
- screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

For media, the following are not in violation of our policy:

- the media is publicly available or is being covered by mainstream media;
- the media and the accompanying tweet text add value to the public discourse or are shared in public interest;
- contains eyewitness accounts or on the ground reports from developing events;
- the subject of the media is a public figure.

Who can report violations of this policy?

Sharing private information:

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

Sharing private media:

When reporting private media, we need a first person report in order to make the determination that the image or video has been shared without their permission. We accept reports from:

1. Individuals depicted in the reported media or

2. Their authorized representative such as:
 - a. parents/legal guardians,
 - b. lawyers, or
 - c. legal representative.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

Sharing private information:

- The first time you violate this policy by sharing private information (such as home address, identity documents etc.), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy by posting private information again after your first warning, your account will be permanently suspended.

Sharing private media:

- If you violate this policy by sharing private media of an individual without their consent, we will require you to remove the content and temporarily lock your account while you remove the media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy

(20191217*)

Overview

November 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you

violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy **(20200918*)**

Overview

September 2020

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are;
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting; and
- coordinated harmful activity that encourages or promotes behavior which violates the [Twitter Rules](#).

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts or by coordinating with others to violate the Twitter Rules. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate;
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account; and
 - coordinating with others to engage in or promote violations of the Twitter Rules, including violations of our [abusive behavior](#) policy.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in "follow trains," "decks," and "Retweet for Retweet" behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- "follow churn" – following and then unfollowing large numbers of accounts in an effort to inflate one's own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account's followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person's browser (malware) or computer or to compromise a person's privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally;
- coordinating with others to express ideas, viewpoints, support, or opposition towards a cause, provided such behavior does not result in violations of the Twitter Rules; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don't violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Denylisting URLs

We denylist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [financial scam policy](#), our [hacked materials policy](#), our approach to [coordinated harmful activity](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Civic integrity policy (20211006*)

Overview

October 2021

You may not use Twitter's services for the purpose of manipulating or interfering in elections or other civic processes. This includes posting or sharing content that may suppress participation or mislead people about when, where, or how to participate in a civic process. In addition, we may label and reduce the visibility of Tweets containing false or misleading information about civic processes in order to provide additional context.

The public conversation occurring on Twitter is never more important than during elections and other civic events. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt civic processes, including through the distribution of false or misleading information about the procedures or circumstances around participation in a civic process. In instances where misleading information does not seek to directly manipulate or disrupt civic processes, but leads to confusion on our service, we may label the Tweets to give additional context. Given the significant risks of confusion about key election information, we may take these actions even if Tweets contain (or attempt to contain) satirical or humorous elements.

What is a civic process?

Twitter considers civic processes to be events or procedures mandated, organized, and conducted by the governing and/or electoral body of a country, state, region, district, or municipality to address a matter of common concern through public participation. Some examples of civic processes may include:

- Political elections
- Censuses
- Major referenda and ballot initiatives

What is in violation of this policy?

This policy addresses 4 categories of misleading behavior and content:

Misleading information about how to participate

We will label or remove false or misleading information about how to participate in an election or other civic process. This includes but is not limited to:

- misleading information about procedures to participate in a civic process (for example, that you can vote by Tweet, text message, email, or phone call in jurisdictions where these are not a possibility);
- misleading information about requirements for participation, including identification or citizenship requirements;
- misleading claims that cause confusion about the established laws, regulations, procedures, and methods of a civic process, or about the actions of officials or entities executing those civic processes; and
- misleading statements or information about the official, announced date or time of a civic process.

Suppression and intimidation

We will label or remove false or misleading information intended to intimidate or dissuade people from participating in an election or other civic process. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to voting in an election, polling places, or collecting census information;
- misleading claims about long lines, equipment problems, or other disruptions at voting locations during election periods;
- misleading claims about process procedures or techniques which could dissuade people from participating; and
- threats regarding voting locations or other key places or events (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

Misleading information about outcomes

We will label or remove false or misleading information intended to undermine public confidence in an election or other civic process. This includes but is not limited to:

- disputed claims that could undermine faith in the process itself, such as unverified information about election rigging, ballot tampering, vote tallying, or certification of election results; and
- misleading claims about the results or outcome of a civic process which calls for or could lead to interference with the implementation of the results of the process, e.g. claiming victory before election results have been certified, inciting unlawful conduct to prevent the procedural or practical implementation of election results (note that our violent threats policy may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral

authority, or government entity. Read more about our [parody, commentary, and fan account policy](#).

What is not a violation of this policy?

Not all false or untrue information about politics or civic processes constitutes manipulation or interference. In the absence of other policy violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected or appointed official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information;
- voting and audience participation for competitions, game shows, or other entertainment purposes; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election or civic process. Therefore, we enable reporting of false or misleading information about civic processes in advance of major events, for people located in the relevant countries and locations. We also work with select government and civil society partners in these countries to provide additional channels for reporting and expedited review.

For civic processes with multiple stages or parts, such as primary elections or lengthy campaigns, reporting will be enabled leading up to the first officially-sanctioned event associated with the civic process.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country at the relevant time, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in civic processes.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in a civic process.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our civic integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. In instances where accounts repeatedly violate this policy, we will use a strike system to determine if further enforcement actions should be applied. We believe this system further helps to reduce the spread of potentially harmful and misleading information on Twitter, particularly for high-severity violations of our rules.

The actions we take may include the following:

Tweet deletion

For high-severity violations of this policy, including (1) misleading information about how to participate, and (2) suppression and intimidation, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Tweet deletions accrue 2 strikes.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Labeling

In circumstances where we do not remove content which violates this policy, we may provide additional context on Tweets sharing the content where they appear on Twitter. This means we may:

- Apply a label and/or warning message to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;

- Turn off people's ability to reply, Retweet, or like the Tweet;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended;
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or relevant Twitter policies; and/or
- Turn off likes, replies, and Retweets.

In most cases, we will take all of the above actions on Tweets we label. In some instances, we'll also turn off your ability to reply, Retweet, or like the Tweet. We prioritize producing Twitter Moments in cases where misleading content on Twitter is gaining significant attention and has caused public confusion on our service. Labels applied to Tweets accrue 1 strike.

Account locks and permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

Repeated violations of this policy are enforced against on the basis of the number of strikes an account has accrued for violations of this policy:

- 1 strike: No account-level action
- 2 strikes: 12-hour account lock
- 3 strikes: 12-hour account lock
- 4 strikes: 7-day account lock
- 5 or more strikes: Permanent suspension

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Impersonation policy (20191001*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Synthetic and manipulated media policy **(20211207*)**

Overview

You may not share synthetic, manipulated, or out-of-context media that may deceive or confuse people and lead to harm (“misleading media”). In addition, we may label Tweets containing misleading media to help people understand their authenticity and to provide additional context.

What is in violation of this policy

In order for content with **misleading media** (including images, videos, audios, gifs, and URLs hosting relevant content) to be labeled or removed under this policy, it must:

- Include media that is significantly and deceptively altered, manipulated, or fabricated, or
- Include media that is shared in a deceptive manner or with false context, and

- Include media likely to result in widespread confusion on public issues, impact public safety, or cause serious harm

We use the following criteria as we consider Tweets and media for labeling or removal under this policy as part of our ongoing work to enforce our rules and ensure healthy and safe conversations on Twitter:

1. Is the content significantly and deceptively altered, manipulated, or fabricated?

In order for content to be labeled or removed under this policy, we must have reason to believe that media are significantly and deceptively altered, manipulated, or fabricated. Synthetic and manipulated media take many different forms and people can employ a wide range of technologies to produce these media. Some of the factors we consider include:

- whether media have been substantially edited or post-processed in a manner that fundamentally alters their composition, sequence, timing, or framing and distorts their meaning;
- whether there are any visual or auditory information (such as new video frames, overdubbed audio, or modified subtitles) that has been added, edited, or removed that fundamentally changes the understanding, meaning, or context of the media;
- whether media have been created, edited, or post-processed with enhancements or use of filters that fundamentally changes the understanding, meaning, or context of the content; and
- whether media depicting a real person have been fabricated or simulated, especially through use of artificial intelligence algorithms

We will not take action to label or remove media that have been edited in ways that do not fundamentally alter their meaning, such as retouched photos or color-corrected videos.

In order to determine if media have been significantly and deceptively altered or fabricated, we may use our own technology or receive reports through partnerships

with third parties. In situations where we are unable to reliably determine if media have been altered or fabricated, we may not take action to label or remove them.

2. Is the content shared in a deceptive manner or with false context?

We also consider whether the context in which media are shared could result in confusion or suggests a deliberate intent to deceive people about the nature or origin of the content, for example, by falsely claiming that it depicts reality. We assess the context provided alongside media to see whether it provides true and factual information. Some of the types of context we assess in order to make this determination include:

- whether misleading, inauthentic, fictional, or produced media are presented or being endorsed as fact or reality, including produced or staged works, reenactments, or exhibitions portrayed as actual events;
- whether media are presented with false or misleading context surrounding the source, location, time, or authenticity of the media;
- whether media are presented with false or misleading context surrounding the identity of the individuals visually depicted in the media

We will not take action to label or remove media that have been shared with commentary or opinions that do not advance or present a misleading claim on the context of the media as listed above.

In order to determine if media have been shared in a deceptive manner or with false context, we may use our own technology or receive reports through partnerships with third parties. In situations where we are unable to reliably determine if media have been shared with false context, we will not label or remove the content.

3. Is the content likely to result in widespread confusion on public issues, impact public safety, or cause serious harm?

Tweets that share misleading media are subject to removal under this policy if they are likely to cause serious harm. Some specific harms we consider include:

- Threats to physical safety of a person or group
- Incitement of abusive behavior to a person or group
- Risk of mass violence or widespread civil unrest
- Risk of impeding or complicating provision of public services, protection efforts, or emergency response
- Threats to the privacy or to the ability of a person or group to freely express themselves or participate in civic events, such as:
 - Stalking or unwanted and obsessive attention
 - Targeted content that includes tropes, epithets, or material that aims to silence someone
 - Voter suppression or intimidation

We also consider the time frame within which the content may be likely to impact public safety or cause serious harm, and are more likely to remove content under this policy if immediate harm is likely to result.

Tweets with misleading media that are not likely to result in immediate harm but still have a potential to impact public safety, result in harm, or cause widespread confusion towards a public issue (health, environment, safety, civil rights and equality, immigration, and social and political stability) may be labeled to reduce their spread and to provide additional context.

While we have other rules also intended to address these forms of harm, including our policies on violent threats, civic integrity, COVID-19 misleading information, and hateful conduct, we will err toward removal in borderline cases that might otherwise not violate existing rules for Tweets that include misleading media.

What is not a violation of this policy

We seek to protect public conversation surrounding various issues. Media often accompany these conversations and encourage further discourse. In the absence of other policy violations, the following are generally not in violation of this policy:

- **Memes or satire**, provided these do not cause significant confusion about the authenticity of the media;
- **Animations, illustrations, and cartoons**, provided these do not cause significant confusion about the authenticity of the media.
- **Commentary, reviews, opinions, and/or reactions**. Sharing media with edits that only add commentary, reviews, opinions, or reactions allows for further debate and discourse relating to various issues and are not in violation of this policy.
- **Counterspeech**. We allow for direct responses to misleading information which seek to undermine its impact by correcting the record, amplifying credible information, and educating the wider community about the prevalence and dynamics of misleading information.
- **Doctored or fake Tweets, social media posts, or chat messages**. Due to the challenges associated with conclusively verifying whether an alleged Tweet, post, or message existed, we do not enforce on doctored or fake Tweets, social media posts, or chat messages under this policy.

Who can report violations of this policy?

We enforce this policy in close coordination with trusted partners, including [our partnership with AP and Reuters](#), other news agencies, public health authorities, and governments. Our team has open lines of communication with various partners to consult and get various media and claims reviewed.

In Australia, South Korea, and the US, Twitter has [begun testing](#) a new reporting feature that will allow users to report Tweets that seem misleading. As part of the experiment, the phrase “It’s misleading” will appear as an option when you select **Report an issue**.

What happens if you violate this policy?

The consequences for violating our synthetic and manipulated media policy depends on the severity of the violation.

Tweet Deletion

For high-severity violations of the policy, including misleading media that have a

serious risk of harm to individuals or communities, we will require you to remove this content.

Labeling

In circumstances where we do not remove content which violates this policy, we may provide additional context on Tweets sharing the misleading media where they appear on Twitter. This means we may:

- Apply a label and/or warning message to the Tweet
- Show a warning to people before they share or like the Tweet;
- Reduce the visibility of the Tweet on Twitter and/or prevent it from being recommended;
- Turn off likes, replies, and Retweets; and/or
- Provide a link to additional explanations or clarifications, such as in a curated landing page (Twitter Moments) or relevant Twitter policies.

In most cases, we will take a combination of the above actions on Tweets we label. We prioritize producing Twitter Moments in cases where misleading content on Twitter is gaining significant attention and has caused public confusion on our service.

Account locks

If we determine that an account has advanced harmful misleading narratives that violate the synthetic and manipulated media policy, we may temporarily lock or suspend the account.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our work and how we build rules to fight misleading media [here](#).

Learn more about [our range of enforcement options](#) and [our approach to policy development and enforcement](#).

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Copyright policy (20200917*)

What types of copyright complaints does Twitter respond to?

Twitter responds to copyright complaints submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review [Twitter's trademark policy](#). If you are concerned about a parody, newsfeed, commentary, or fan account, please see the relevant policy [here](#). These are generally not copyright issues.

Am I a copyright holder? How do I know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

What to consider before submitting a copyright complaint

Before submitting a copyright complaint to us, please consider whether or not the use could be considered [fair use](#).

If you have considered fair use, and you still wish to continue with a copyright complaint, you may want to first reach out to the user in question to see if you can resolve the matter directly with the user. You can reply to the user's Tweet or send the user a Direct Message and ask for them to remove your copyrighted content without having to contact Twitter.

Prior to submitting a formal complaint with Twitter, please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

Note: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph. If you're unsure whether or not you own the copyrights to a work, or if you're infringing upon someone else's work, please consult an attorney or another advisor.

What information do you need to process a copyright complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the complaint is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How do I file a copyright complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and [filing a copyright complaint](#). If you are logged in to [twitter.com](#), you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a full copy of your notice (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a copyright complaint.

How are claims processed?

We process copyright complaints in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. If you do not receive a ticket confirmation that means we did not receive your complaint and you should re-submit your complaint. However, please note, submitting duplicate copyright complaints will result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) and provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice. We will also forward a redacted copy of the complaint to [Lumen](#), with your personal information removed.

What information gets forwarded to the reported user(s)?

If we remove or disable access to the materials reported in a copyright complaint, the reported user(s) will receive a copy of the complaint, including the reporter's full name, email, street address, and any other information included in the complaint.

If you are uncomfortable sharing your contact information with the reported user(s), you may wish to consider appointing an agent to submit your DMCA notice on your behalf. Your agent will be required to submit the DMCA notice with valid contact information, and identify you as the content owner that they are representing.

What happens next?

Twitter's response to copyright complaints may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a copyright complaint, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the complaint, along with instructions for filing a counter-notice.

If you've not yet received a copy of the copyright complaint regarding the content removed from your account, please respond to the support ticket we sent you. If your complaint involves removal of content from Periscope, please see this help center [article](#) regarding how to request a copy of your complaint.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from **@Username** has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

My content was removed from Twitter

Why did I receive a copyright complaint?

If you receive a copyright complaint, it means that access to the content described in the complaint has been restricted. Please take the time to read through our correspondence to you, which includes information on the complaint we received as well as instructions on how to file a counter-notice. Please ensure that you are monitoring the email address associated with your Twitter account.

Tip: Removing the material reported in a copyright complaint will not resolve that complaint.

What if I want to contest the takedown?

If you believe that the materials reported in the copyright complaints were misidentified or removed in error, you may send us a counter-notification(s). A counter-notice is a request for Twitter to reinstate the removed material, and it has legal consequences. Alternatively, you may be able to seek a retraction of the copyright complaint from the reporter.

How do I seek a retraction?

The DMCA complaint you received includes the contact information of the reporter. You may want to reach out and ask them to retract their notice. The reporter can send retractions to copyright@twitter.com, and should

include: (1) identification of the material that was disabled, and (2) a statement that the reporter would like to retract their DMCA notice. This is the fastest and most efficient means of resolving an unresolved copyright complaint. A retraction is at the sole discretion of the original reporter.

When should I file a counter-notice?

A counter-notice is a request for Twitter to reinstate the removed material, and is the start of a legal process that has legal consequences. For example, submitting a counter notice indicates that you consent to the jurisdiction of a U.S. Federal court and that you consent to the disclosure of your personal information to the reporter and Lumen website.

With these considerations in mind, you may file a counter-notice if you believe that this material was misidentified, or you have a good faith belief that the material should not have been removed. If you're unsure whether or not you should file a counter-notice, you may want to consult with an attorney.

Tip: Re-posting material removed in response to a copyright complaint may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notice rather than re-posting the material.

What information do you need to process a counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What happens after I submit a counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. This means that the contact information that is submitted in your counter-notice will be shared to the person who filed the original notice.

If the copyright owner disagrees that the content was removed in error or misidentification, they may pursue legal action against you. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a copyright complaint or counter-notice is serious business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

What happens if my account receives multiple copyright complaints?

If multiple copyright complaints are received Twitter may lock accounts or take other actions to warn repeat violators. These warnings may vary across Twitter's services. Under appropriate circumstances we may suspend

user accounts under our repeat infringer policy. However, we may take retractions and counter-notices into account when applying our repeat infringer policy.

Our approach to policy development and enforcement philosophy

(20190520*)

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating

Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (20200204*)

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Synthetic and manipulated media: You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.

5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Violent organizations policy (20201026*)

Overview

October 2020

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

Other violent organizations

Other violent organizations are those that meet all of the below criteria:

- a collection of individuals with a shared purpose; and
- have systematically targeted civilians with violence.

We examine a group's activities both on and off Twitter to determine whether they satisfy the above criteria.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a violent organization;
- recruiting for a violent organization;
- providing or distributing services (e.g., financial, media/propaganda) to further a violent organization's stated goals; and
- using the insignia or symbol of violent organizations to promote them or indicate affiliation or support.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to state or governmental organizations.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.

6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wishing for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy (20201012*)

Overview

October 2020

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives;
- identifying alleged victims of childhood sexual exploitation by name or image; and
- promoting or normalizing sexual attraction to minors as a form of identity or sexual orientation.

What is not a violation of this policy?

- Discussions related to child sexual exploitation are permitted, provided they don't normalise, promote or glorify child sexual exploitation **in any way**;
- Conversations about help-seeking behaviour of individuals who may struggle with an attraction to minors;
- Advocacy against illegal or harmful activity involving minors provided there is no sharing of, or linking to any material featuring child sexual exploitation; and
- Depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios (e.g., art works by internationally renowned artists that feature minors, news media reportage, media created and shared for scientific or educational purposes).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

Abusive behavior (20211017*)

[Twitter Rules](#): You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it

may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., “I will kill you.”

Note: We have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account. Wishing, hoping, or calling for serious harm on a person or group of people.

Wishing, hoping, or calling for serious harm on a person or group of people

We do not tolerate content that wishes, hopes, promotes, incites, or expresses a desire for death, serious bodily harm or serious disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., “I hope you get cancer and die.”
- Wishing for someone to fall victim to a serious accident e.g., “I wish that you would get run over by a car next time you run your mouth.”

- Saying that a group of individuals deserves serious physical injury e.g., “If this group of protesters don’t shut up, they deserve to be shot.”

About wishes of harm exceptions on Twitter

We recognize that conversations regarding certain individuals credibly accused of severe violence may prompt outrage and associated wishes of harm. In these limited cases, we will request the user to delete the Tweet without any risk of account penalty, strike, or suspension. Examples are, but not limited to:

- “I wish all rapists to die.”
- “Child abusers should be hanged.”

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone’s body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using insults, profanity, or slurs with the purpose of harassing or intimidating others

We take action against the use of insults, profanity, or slurs to target others. In some cases, such as (but not limited to) severe, repetitive usage of insults or slurs where the primary intent is to harass or intimidate others, we may require Tweet removal. In other cases, such as (but not limited to) moderate, isolated usage of insults and profanity

where the primary intent is to harass or intimidate others, we may limit Tweet visibility as further described below. Please also note that while some individuals may find certain terms to be offensive, we will not take action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Denying mass casualty events took place

We prohibit content that denies that mass murder or other mass casualty events took place, where we can verify that the event occurred, and when the content is shared with abusive intent. This may include references to such an event as a “hoax” or claims that victims or survivors are fake or “actors.” It includes, but is not limited to, events like the Holocaust, school shootings, terrorist attacks, and natural disasters.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual’s previous record of rule violations. The following is a list of potential enforcement options for content that violates this policy:

- Downranking Tweets in replies, except when the user follows the Tweet author.
- Making Tweets ineligible for amplification in Top search results and/or on timelines for users who don't follow the Tweet author.
- Excluding Tweets and/or accounts in email or in-product recommendations.
- Requiring Tweet removal.
 - For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension.
- Suspending accounts whose primary use we've determined is to engage in abusive behavior as defined in this policy, or who have shared violent threats.

Learn more about [our range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

Hateful conduct policy (20211010*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognize that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and more harmful.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals or groups with abuse based on their perceived membership in a protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., “I will kill you.”

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, incites, or expresses a desire for death, serious bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

- Hoping that an entire protected category and/or individuals who may be members of that category dies as a result of a serious disease, e.g., “I hope all [nationality] get COVID and die.”
- Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”
- Saying that a group of individuals deserve serious physical injury, e.g., “If this group of [slur] don’t shut up, they deserve to be shot.”
- Encouraging others to commit violence against an individual or a group based on their perceived membership in a protected category, e.g., “I’m in the mood to punch a [racial slur], who’s with me?”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals or groups with content that references forms of violence or violent events where a protected category was the primary target or victims,

where the intent is to harass. This includes, but is not limited to media or text that refers to or depicts:

- genocides, (e.g., the Holocaust);
- lynchings.

Incitement against protected categories

We prohibit inciting behavior that targets individuals or groups of people belonging to protected categories. This includes content intended:

- to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists.”
- to incite others to harass members of a protected category on or off platform, e.g., “I’m sick of these [religious group] thinking they are better than us, if any of you see someone wearing a [religious symbol of the religious group], grab it off them and post pics!”
- to incite others to discriminate in the form of denial of support to the economic enterprise of an individual or group because of their perceived membership in a protected category, e.g., “If you go to a [religious group] store, you are supporting those [slur], let’s stop giving our money to these [religious slur].” This may not include content intended as political in nature, such as political commentary or content relating to boycotts or protests.

Note that content intended to incite violence against a protected category is prohibited under **Wishing, hoping, or calling for serious harm on a person or groups of people**.

We prohibit targeting individuals and groups with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g.,

“all [religious group] are terrorists.”

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting others with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals. We also prohibit the dehumanization of a group of people based on their religion, caste, age, disability, serious disease, national origin, race, or ethnicity. In some cases, such as (but not limited to) severe, repetitive usage of slurs, epithets, or racist/sexy tropes where the primary intent is to harass or intimidate others, we may require Tweet removal. In other cases, such as (but not limited to) moderate, isolated usage where the primary intent is to harass or intimidate others, we may limit Tweet visibility as further described below.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous

record of rule violations. The following is a list of potential enforcement options for content that violates this policy:

- Downranking Tweets in replies, except when the user follows the Tweet author.
- Making Tweets ineligible for amplification in Top search results and/or on timelines for users who don't follow the Tweet author.
- Excluding Tweets and/or accounts in email or in-product recommendations.
- Requiring Tweet removal.
 - For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension.
- Suspending accounts whose primary use we've determined is to engage in hateful conduct as defined in this policy, or who have shared violent threats.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

Suicide and Self-harm Policy **(20200805*)**

Overview

You may not promote or encourage suicide or self-harm.

At Twitter, we recognize that suicide and self-harm are significant social & public health challenges that require collaboration between all stakeholders – public, private, and civil society – and that we have a role and responsibility to help people access and receive support when they need it.

When developing this policy, we consulted extensively with experts to ensure that people who have engaged in self-harm or experienced suicidal thoughts can share their personal experiences. We also recognized the need to protect people from the potential harm caused by exposure to content that could promote or encourage self-harm – intentionally or inadvertently. That's why our policy prohibits content that promotes or encourages self-harming behaviors and provides support to those undergoing experiences with self-harm or suicidal thoughts.

What is in violation of this policy?

Under this policy, you can't promote, or otherwise encourage, suicide or self-harm. We define promotion and encouragement to include statements such as "the most effective", "the easiest", "the best", "the most successful", "you should", "why don't you". Violations of this policy can occur via Tweets, images or videos, including live video.

We define suicide to be the act of taking one's own life. We define self-harm to include:

- self-inflicted physical injuries e.g., cutting; and
- eating disorders e.g., bulimia, anorexia.

Violations of this policy include, but are not limited to:

- encouraging someone to physically harm or kill themselves;
- asking others for encouragement to engage in self-harm or suicide, including seeking partners for group suicides or suicide games; and
- sharing information, strategies, methods or instructions that would assist people to engage in self-harm and suicide.

What is not a violation of this policy?

Some examples of behavior that are not considered a violation of this policy include:

- sharing personal stories and experiences related to self-harm or suicide;
- sharing coping mechanisms and resources for addressing self-harm or suicidal thoughts; and
- discussions that are focused on research, advocacy, and education related to self-harm or suicide prevention.

Note: people can share their personal experiences, but should avoid sharing detailed information about specific strategies or methods related to self-harm, as this could inadvertently encourage this behavior.

Who can report violations of this policy?

Anyone can report content that may encourage or promote suicide or self-harm via our in-app reporting or our [specialized reporting form](#). These reports are routed to a dedicated team who evaluate each case individually.

Note: if we receive a report that someone has expressed an intention to engage in self-harm or suicide, we will contact them directly, encourage them to seek support, and provide information about dedicated online and hotline resources. We may also work with law enforcement officials where appropriate, for example, if we receive a valid emergency disclosure request as defined in our [Law Enforcement Guidelines](#).

How to report violations of this policy

To ensure that we handle reports sensitively, our in-app reporting provides separate options for people who may be expressing an intention to harm themselves, and content that is encouraging or promoting self-harm or suicide.

Expressing intentions of self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu

2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Reporting form

You can also report this content for review via [our dedicated report form](#).

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Encouraging self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

What happens if you violate this policy?

Our enforcement approach depends on the type of content being shared, whether or not the reported account is encouraging or promoting self-harm or suicide, and the account's previous history of violations.

If you violate this policy by sharing content that intentionally encourages others to harm themselves, ask others to encourage you to harm yourself, or share detailed information or instructions related to self-harm or suicide methods, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy, or if your account is dedicated to promoting or encouraging self-harm or suicide, your account will be permanently suspended. If cases include images or videos related to self-harm or suicide, we will also evaluate this content under [our sensitive media policy](#). If you believe that your account was suspended in error, [you can submit an appeal](#).

We may also take steps to prevent the spread of instructional material hosted on third-party websites by [marking such links as unsafe](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Learn more](#) about how you can support someone experiencing thoughts of self-harm or suicide.

Visit our Safety Center for a list of local [mental health resources](#) and read [our blog on suicide prevention](#) to learn more about our work.

Sensitive media policy (20191204*)

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and

- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and

- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20211130)

Overview

November 2021

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

In addition, you may not share private media, such as images or videos of private individuals, without their consent. However, we recognise that there are instances where users may share images or videos of private individuals, who are not public figures, as part of a newsworthy event or to further public discourse on issues or events of public interest. In such cases, we may allow the media to remain on the platform.

Sharing someone's private information:

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private

information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We take this into consideration because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from potential physical harm as a result of their information being shared, so we consider information such as physical location and phone numbers to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

Sharing private media:

Posting images is an important part of our users' experience on Twitter. Where individuals have a reasonable expectation of privacy in an individual piece of media, we believe they should be able to determine whether or not it is shared. Sharing such media could potentially violate users' privacy and may lead to emotional or physical harm. When we are notified by individuals depicted, or their authorized representative, that they did not consent to having media shared, we will remove the media. This policy is not applicable to public figures.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;

- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.
- media of private individuals without the permission of the person(s) depicted.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;

- descriptions of physical appearance;
- gossip, rumours, accusations, and allegations; and
- screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

For media, the following are not in violation of our policy:

- the media is publicly available or is being covered by mainstream media;
- the media and the accompanying tweet text add value to the public discourse or are shared in public interest;
- contains eyewitness accounts or on the ground reports from developing events;
- the subject of the media is a public figure.

Who can report violations of this policy?

Sharing private information:

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

Sharing private media:

When reporting private media, we need a first person report in order to make the determination that the image or video has been shared without their permission. We accept reports from:

1. Individuals depicted in the reported media or

2. Their authorized representative such as:
 - a. parents/legal guardians,
 - b. lawyers, or
 - c. legal representative.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

Sharing private information:

- The first time you violate this policy by sharing private information (such as home address, identity documents etc.), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy by posting private information again after your first warning, your account will be permanently suspended.

Sharing private media:

- If you violate this policy by sharing private media of an individual without their consent, we will require you to remove the content and temporarily lock your account while you remove the media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy

(20191217*)

Overview

November 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you

violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy **(20200918*)**

Overview

September 2020

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are;
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting; and
- coordinated harmful activity that encourages or promotes behavior which violates the [Twitter Rules](#).

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts or by coordinating with others to violate the Twitter Rules. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate;
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account; and
 - coordinating with others to engage in or promote violations of the Twitter Rules, including violations of our [abusive behavior](#) policy.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in "follow trains," "decks," and "Retweet for Retweet" behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- "follow churn" – following and then unfollowing large numbers of accounts in an effort to inflate one's own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account's followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person's browser (malware) or computer or to compromise a person's privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally;
- coordinating with others to express ideas, viewpoints, support, or opposition towards a cause, provided such behavior does not result in violations of the Twitter Rules; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don't violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Denylisting URLs

We denylist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [financial scam policy](#), our [hacked materials policy](#), our approach to [coordinated harmful activity](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Civic integrity policy (20211006*)

Overview

October 2021

You may not use Twitter's services for the purpose of manipulating or interfering in elections or other civic processes. This includes posting or sharing content that may suppress participation or mislead people about when, where, or how to participate in a civic process. In addition, we may label and reduce the visibility of Tweets containing false or misleading information about civic processes in order to provide additional context.

The public conversation occurring on Twitter is never more important than during elections and other civic events. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt civic processes, including through the distribution of false or misleading information about the procedures or circumstances around participation in a civic process. In instances where misleading information does not seek to directly manipulate or disrupt civic processes, but leads to confusion on our service, we may label the Tweets to give additional context. Given the significant risks of confusion about key election information, we may take these actions even if Tweets contain (or attempt to contain) satirical or humorous elements.

What is a civic process?

Twitter considers civic processes to be events or procedures mandated, organized, and conducted by the governing and/or electoral body of a country, state, region, district, or municipality to address a matter of common concern through public participation. Some examples of civic processes may include:

- Political elections
- Censuses
- Major referenda and ballot initiatives

What is in violation of this policy?

This policy addresses 4 categories of misleading behavior and content:

Misleading information about how to participate

We will label or remove false or misleading information about how to participate in an election or other civic process. This includes but is not limited to:

- misleading information about procedures to participate in a civic process (for example, that you can vote by Tweet, text message, email, or phone call in jurisdictions where these are not a possibility);
- misleading information about requirements for participation, including identification or citizenship requirements;
- misleading claims that cause confusion about the established laws, regulations, procedures, and methods of a civic process, or about the actions of officials or entities executing those civic processes; and
- misleading statements or information about the official, announced date or time of a civic process.

Suppression and intimidation

We will label or remove false or misleading information intended to intimidate or dissuade people from participating in an election or other civic process. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to voting in an election, polling places, or collecting census information;
- misleading claims about long lines, equipment problems, or other disruptions at voting locations during election periods;
- misleading claims about process procedures or techniques which could dissuade people from participating; and
- threats regarding voting locations or other key places or events (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

Misleading information about outcomes

We will label or remove false or misleading information intended to undermine public confidence in an election or other civic process. This includes but is not limited to:

- disputed claims that could undermine faith in the process itself, such as unverified information about election rigging, ballot tampering, vote tallying, or certification of election results; and
- misleading claims about the results or outcome of a civic process which calls for or could lead to interference with the implementation of the results of the process, e.g. claiming victory before election results have been certified, inciting unlawful conduct to prevent the procedural or practical implementation of election results (note that our violent threats policy may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral

authority, or government entity. Read more about our [parody, commentary, and fan account policy](#).

What is not a violation of this policy?

Not all false or untrue information about politics or civic processes constitutes manipulation or interference. In the absence of other policy violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected or appointed official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information;
- voting and audience participation for competitions, game shows, or other entertainment purposes; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election or civic process. Therefore, we enable reporting of false or misleading information about civic processes in advance of major events, for people located in the relevant countries and locations. We also work with select government and civil society partners in these countries to provide additional channels for reporting and expedited review.

For civic processes with multiple stages or parts, such as primary elections or lengthy campaigns, reporting will be enabled leading up to the first officially-sanctioned event associated with the civic process.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country at the relevant time, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in civic processes.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in a civic process.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our civic integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. In instances where accounts repeatedly violate this policy, we will use a strike system to determine if further enforcement actions should be applied. We believe this system further helps to reduce the spread of potentially harmful and misleading information on Twitter, particularly for high-severity violations of our rules.

The actions we take may include the following:

Tweet deletion

For high-severity violations of this policy, including (1) misleading information about how to participate, and (2) suppression and intimidation, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Tweet deletions accrue 2 strikes.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Labeling

In circumstances where we do not remove content which violates this policy, we may provide additional context on Tweets sharing the content where they appear on Twitter. This means we may:

- Apply a label and/or warning message to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;

- Turn off people's ability to reply, Retweet, or like the Tweet;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended;
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or relevant Twitter policies; and/or
- Turn off likes, replies, and Retweets.

In most cases, we will take all of the above actions on Tweets we label. In some instances, we'll also turn off your ability to reply, Retweet, or like the Tweet. We prioritize producing Twitter Moments in cases where misleading content on Twitter is gaining significant attention and has caused public confusion on our service. Labels applied to Tweets accrue 1 strike.

Account locks and permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

Repeated violations of this policy are enforced against on the basis of the number of strikes an account has accrued for violations of this policy:

- 1 strike: No account-level action
- 2 strikes: 12-hour account lock
- 3 strikes: 12-hour account lock
- 4 strikes: 7-day account lock
- 5 or more strikes: Permanent suspension

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Impersonation policy (20191001*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Synthetic and manipulated media policy **(20210626*)**

Overview

You may not deceptively promote synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context.

You should be able to find reliable information on Twitter. That means understanding whether the content you see is real or fabricated and having the ability to find more context about what you see on Twitter. Therefore, we may label Tweets that include media (videos, audio, and images) that have been deceptively altered or fabricated. In addition, you may not share deceptively altered media on Twitter in ways that mislead or deceive people about the media's authenticity where threats to physical safety or other serious harm may result.

We use the following criteria as we consider Tweets and media for labeling or removal under this policy as part of our ongoing work to enforce our rules and ensure healthy and safe conversation on Twitter (additional information is available below):

1. Is the content synthetic or manipulated?

In order for content to be labeled or removed under this policy, we must have reason to believe that media, or the context in which media are presented, are significantly and deceptively altered or manipulated. Synthetic and manipulated media take many different forms and people can employ a wide range of technologies to produce these media. In assessing whether media have been significantly and deceptively altered or fabricated, some of the factors we consider include:

- whether the content has been substantially edited in a manner that fundamentally alters its composition, sequence, timing, or framing;
- any visual or auditory information (such as new video frames, overdubbed audio, or modified subtitles) that has been added or removed; and
- whether media depicting a real person have been fabricated or simulated

We are most likely to take action (either labeling or removal, as described below) on more significant forms of alteration, such as wholly synthetic audio or video or content that has been doctored (spliced and reordered, slowed down) to change its meaning. Subtler forms of manipulated media, such as isolative editing, omission of context, or presentation with false context, may be labeled or removed on a case-by-case basis.

We will not take action to label or remove media that have been edited in ways that do not fundamentally alter their meaning, such as retouched photos or color-corrected videos.

In order to determine if media have been significantly and deceptively altered or fabricated, we may use our own technology or receive reports through partnerships with third parties. In situations where we are unable to reliably determine if media have been altered or fabricated, we may not take action to label or remove them.

2. Is the content shared in a deceptive manner?

We also consider whether the context in which media are shared could result in confusion or misunderstanding or suggests a deliberate intent to deceive people about the nature or origin of the content, for example by falsely claiming that it depicts reality. We assess the context provided alongside media to see whether it makes clear that the media have been altered or fabricated. Some of the types of context we assess in order to make this determination include:

- The text of the Tweet accompanying or within media
- Metadata associated with media
- Information on the profile of the account sharing media
- Websites linked in the Tweet, or in the profile of the account sharing media

3. Is the content likely to impact public safety or cause serious harm?

Tweets that share synthetic and manipulated media are subject to removal under this policy if they are likely to cause serious harm. Some specific harms we consider include:

- Threats to the physical safety of a person or group
- Risk of mass violence or widespread civil unrest

- Threats to the privacy or ability of a person or group to freely express themselves or participate in civic events, such as:
 - Stalking or unwanted and obsessive attention
 - Targeted content that includes tropes, epithets, or material that aims to silence someone
 - Voter suppression or intimidation

While we have other rules also intended to address these forms of harm, including our policies on violent threats, election integrity, and hateful conduct, we will err toward removal in borderline cases that might otherwise not violate existing rules for Tweets that include synthetic or manipulated media.

We also consider the time frame within which the content may be likely to impact public safety or cause serious harm, and are more likely to remove content under this policy if we find that immediate harms are likely to result from the content's presence on Twitter.

Note: We may also take action on synthetic and manipulated content under our [non-consensual nudity policy](#) (such as pornographic media altered to insert the faces of people not actually involved) or other parts of the Twitter Rules.

Labeling and removal

In most cases, if we have reason to believe that media shared in a Tweet have been significantly and deceptively altered or fabricated, we will provide additional context on Tweets sharing the media where they appear on Twitter. This means we may:

- Apply a label to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended;
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or landing page; and/or

- Turn off likes, replies, and Retweets.

In most cases, we will take all of the above actions on Tweets we label.

Media that meet **all three of the criteria** defined above—i.e. that are synthetic or manipulated, shared in a deceptive manner, and is likely to cause harm—may not be shared on Twitter and are subject to removal. Accounts engaging in repeated or severe violations of this policy may be permanently suspended.

** Other parts of the Twitter Rules apply and may lead to the removal of the content, particularly where there is high likelihood of severe harm, such as a threat to someone's life or physical safety.*

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and

- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Copyright policy (20200917*)

What types of copyright complaints does Twitter respond to?

Twitter responds to copyright complaints submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review [Twitter's trademark policy](#). If you are concerned about a parody, newsfeed, commentary, or fan account, please see the relevant policy [here](#). These are generally not copyright issues.

Am I a copyright holder? How do I know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

What to consider before submitting a copyright complaint

Before submitting a copyright complaint to us, please consider whether or not the use could be considered [fair use](#).

If you have considered fair use, and you still wish to continue with a copyright complaint, you may want to first reach out to the user in question to see if you can resolve the matter directly with the user. You can reply to the user's Tweet or send the user a Direct Message and ask for them to remove your copyrighted content without having to contact Twitter.

Prior to submitting a formal complaint with Twitter, please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

Note: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph. If you're unsure whether or not you own the copyrights to a work, or if you're infringing upon someone else's work, please consult an attorney or another advisor.

What information do you need to process a copyright complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the complaint is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How do I file a copyright complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and [filing a copyright complaint](#). If you are logged in to twitter.com, you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a full copy of your notice (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a copyright complaint.

How are claims processed?

We process copyright complaints in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. If you do not receive a ticket confirmation that means we did not receive your complaint and you should re-submit your complaint. However, please note, submitting duplicate copyright complaints will result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) and provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice. We will also forward a redacted copy of the complaint to [Lumen](#), with your personal information removed.

What information gets forwarded to the reported user(s)?

If we remove or disable access to the materials reported in a copyright complaint, the reported user(s) will receive a copy of the complaint, including the reporter's full name, email, street address, and any other information included in the complaint.

If you are uncomfortable sharing your contact information with the reported user(s), you may wish to consider appointing an agent to submit your DMCA notice on your behalf. Your agent will be required to submit the DMCA notice with valid contact information, and identify you as the content owner that they are representing.

What happens next?

Twitter's response to copyright complaints may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a copyright complaint, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the complaint, along with instructions for filing a counter-notice.

If you've not yet received a copy of the copyright complaint regarding the content removed from your account, please respond to the support ticket we sent you. If your complaint involves removal of content from Periscope, please see this help center [article](#) regarding how to request a copy of your complaint.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from [@Username](#) has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

My content was removed from Twitter

Why did I receive a copyright complaint?

If you receive a copyright complaint, it means that access to the content described in the complaint has been restricted. Please take the time to read through our correspondence to you, which includes information on the complaint we received as well as instructions on how to file a counter-notice. Please ensure that you are monitoring the email address associated with your Twitter account.

Tip: Removing the material reported in a copyright complaint will not resolve that complaint.

What if I want to contest the takedown?

If you believe that the materials reported in the copyright complaints were misidentified or removed in error, you may send us a counter-notification(s). A counter-notice is a request for Twitter to reinstate the removed material, and it has legal consequences. Alternatively, you may be able to seek a retraction of the copyright complaint from the reporter.

How do I seek a retraction?

The DMCA complaint you received includes the contact information of the reporter. You may want to reach out and ask them to retract their notice. The reporter can send retractions to copyright@twitter.com, and should include: (1) identification of the material that was disabled, and (2) a statement that the reporter would like to retract their DMCA notice. This is the fastest and most efficient means of resolving an unresolved copyright complaint. A retraction is at the sole discretion of the original reporter.

When should I file a counter-notice?

A counter-notice is a request for Twitter to reinstate the removed material, and is the start of a legal process that has legal consequences. For example, submitting a counter notice indicates that you consent to the jurisdiction of a U.S. Federal court and that you consent to the disclosure of your personal information to the reporter and Lumen website.

With these considerations in mind, you may file a counter-notice if you believe that this material was misidentified, or you have a good faith belief that the material should not have been removed. If you're unsure whether or not you should file a counter-notice, you may want to consult with an attorney.

Tip: Re-posting material removed in response to a copyright complaint may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notice rather than re-posting the material.

What information do you need to process a counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);

3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What happens after I submit a counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. This means that the contact information that is submitted in your counter-notice will be shared to the person who filed the original notice.

If the copyright owner disagrees that the content was removed in error or misidentification, they may pursue legal action against you. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a copyright complaint or counter-notice is serious business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

What happens if my account receives multiple copyright complaints?

If multiple copyright complaints are received Twitter may lock accounts or take other actions to warn repeat violators. These warnings may vary across Twitter's services. Under appropriate circumstances we may suspend user accounts under our repeat infringer policy. However, we may take retractions and counter-notices into account when applying our repeat infringer policy.

Our approach to policy development and enforcement philosophy (20190520*)

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting

violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (20200204*)

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Synthetic and manipulated media: You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.

5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Violent organizations policy (20201026*)

Overview

October 2020

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

Other violent organizations

Other violent organizations are those that meet all of the below criteria:

- a collection of individuals with a shared purpose; and
- have systematically targeted civilians with violence.

We examine a group's activities both on and off Twitter to determine whether they satisfy the above criteria.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a violent organization;
- recruiting for a violent organization;
- providing or distributing services (e.g., financial, media/propaganda) to further a violent organization's stated goals; and
- using the insignia or symbol of violent organizations to promote them or indicate affiliation or support.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to state or governmental organizations.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.

6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wishing for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy (20201012*)

Overview

October 2020

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives;
- identifying alleged victims of childhood sexual exploitation by name or image; and
- promoting or normalizing sexual attraction to minors as a form of identity or sexual orientation.

What is not a violation of this policy?

- Discussions related to child sexual exploitation are permitted, provided they don't normalise, promote or glorify child sexual exploitation **in any way**;
- Conversations about help-seeking behaviour of individuals who may struggle with an attraction to minors;
- Advocacy against illegal or harmful activity involving minors provided there is no sharing of, or linking to any material featuring child sexual exploitation; and
- Depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios (e.g., art works by internationally renowned artists that feature minors, news media reportage, media created and shared for scientific or educational purposes).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

Abusive behavior (20211017*)

[Twitter Rules](#): You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it

may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., “I will kill you.”

Note: We have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account. Wishing, hoping, or calling for serious harm on a person or group of people.

Wishing, hoping, or calling for serious harm on a person or group of people

We do not tolerate content that wishes, hopes, promotes, incites, or expresses a desire for death, serious bodily harm or serious disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., “I hope you get cancer and die.”
- Wishing for someone to fall victim to a serious accident e.g., “I wish that you would get run over by a car next time you run your mouth.”

- Saying that a group of individuals deserves serious physical injury e.g., “If this group of protesters don’t shut up, they deserve to be shot.”

About wishes of harm exceptions on Twitter

We recognize that conversations regarding certain individuals credibly accused of severe violence may prompt outrage and associated wishes of harm. In these limited cases, we will request the user to delete the Tweet without any risk of account penalty, strike, or suspension. Examples are, but not limited to:

- “I wish all rapists to die.”
- “Child abusers should be hanged.”

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone’s body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using insults, profanity, or slurs with the purpose of harassing or intimidating others

We take action against the use of insults, profanity, or slurs to target others. In some cases, such as (but not limited to) severe, repetitive usage of insults or slurs where the primary intent is to harass or intimidate others, we may require Tweet removal. In other cases, such as (but not limited to) moderate, isolated usage of insults and profanity

where the primary intent is to harass or intimidate others, we may limit Tweet visibility as further described below. Please also note that while some individuals may find certain terms to be offensive, we will not take action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Denying mass casualty events took place

We prohibit content that denies that mass murder or other mass casualty events took place, where we can verify that the event occurred, and when the content is shared with abusive intent. This may include references to such an event as a “hoax” or claims that victims or survivors are fake or “actors.” It includes, but is not limited to, events like the Holocaust, school shootings, terrorist attacks, and natural disasters.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual’s previous record of rule violations. The following is a list of potential enforcement options for content that violates this policy:

- Downranking Tweets in replies, except when the user follows the Tweet author.
- Making Tweets ineligible for amplification in Top search results and/or on timelines for users who don't follow the Tweet author.
- Excluding Tweets and/or accounts in email or in-product recommendations.
- Requiring Tweet removal.
 - For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension.
- Suspending accounts whose primary use we've determined is to engage in abusive behavior as defined in this policy, or who have shared violent threats.

Learn more about [our range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

Hateful conduct policy (20211010*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognize that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and more harmful.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals or groups with abuse based on their perceived membership in a protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., “I will kill you.”

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, incites, or expresses a desire for death, serious bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

- Hoping that an entire protected category and/or individuals who may be members of that category dies as a result of a serious disease, e.g., “I hope all [nationality] get COVID and die.”
- Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”
- Saying that a group of individuals deserve serious physical injury, e.g., “If this group of [slur] don’t shut up, they deserve to be shot.”
- Encouraging others to commit violence against an individual or a group based on their perceived membership in a protected category, e.g., “I’m in the mood to punch a [racial slur], who’s with me?”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals or groups with content that references forms of violence or violent events where a protected category was the primary target or victims,

where the intent is to harass. This includes, but is not limited to media or text that refers to or depicts:

- genocides, (e.g., the Holocaust);
- lynchings.

Incitement against protected categories

We prohibit inciting behavior that targets individuals or groups of people belonging to protected categories. This includes content intended:

- to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists.”
- to incite others to harass members of a protected category on or off platform, e.g., “I’m sick of these [religious group] thinking they are better than us, if any of you see someone wearing a [religious symbol of the religious group], grab it off them and post pics!”
- to incite others to discriminate in the form of denial of support to the economic enterprise of an individual or group because of their perceived membership in a protected category, e.g., “If you go to a [religious group] store, you are supporting those [slur], let’s stop giving our money to these [religious slur].”
This may not include content intended as political in nature, such as political commentary or content relating to boycotts or protests.

Note that content intended to incite violence against a protected category is prohibited under **Wishing, hoping, or calling for serious harm on a person or groups of people**.

We prohibit targeting individuals and groups with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g.,

“all [religious group] are terrorists.”

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting others with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals. We also prohibit the dehumanization of a group of people based on their religion, caste, age, disability, serious disease, national origin, race, or ethnicity. In some cases, such as (but not limited to) severe, repetitive usage of slurs, epithets, or racist/sexy tropes where the primary intent is to harass or intimidate others, we may require Tweet removal. In other cases, such as (but not limited to) moderate, isolated usage where the primary intent is to harass or intimidate others, we may limit Tweet visibility as further described below.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous

record of rule violations. The following is a list of potential enforcement options for content that violates this policy:

- Downranking Tweets in replies, except when the user follows the Tweet author.
- Making Tweets ineligible for amplification in Top search results and/or on timelines for users who don't follow the Tweet author.
- Excluding Tweets and/or accounts in email or in-product recommendations.
- Requiring Tweet removal.
 - For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension.
- Suspending accounts whose primary use we've determined is to engage in hateful conduct as defined in this policy, or who have shared violent threats.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

Suicide and Self-harm Policy **(20200805*)**

Overview

You may not promote or encourage suicide or self-harm.

At Twitter, we recognize that suicide and self-harm are significant social & public health challenges that require collaboration between all stakeholders – public, private, and civil society – and that we have a role and responsibility to help people access and receive support when they need it.

When developing this policy, we consulted extensively with experts to ensure that people who have engaged in self-harm or experienced suicidal thoughts can share their personal experiences. We also recognized the need to protect people from the potential harm caused by exposure to content that could promote or encourage self-harm – intentionally or inadvertently. That's why our policy prohibits content that promotes or encourages self-harming behaviors and provides support to those undergoing experiences with self-harm or suicidal thoughts.

What is in violation of this policy?

Under this policy, you can't promote, or otherwise encourage, suicide or self-harm. We define promotion and encouragement to include statements such as "the most effective", "the easiest", "the best", "the most successful", "you should", "why don't you". Violations of this policy can occur via Tweets, images or videos, including live video.

We define suicide to be the act of taking one's own life. We define self-harm to include:

- self-inflicted physical injuries e.g., cutting; and
- eating disorders e.g., bulimia, anorexia.

Violations of this policy include, but are not limited to:

- encouraging someone to physically harm or kill themselves;
- asking others for encouragement to engage in self-harm or suicide, including seeking partners for group suicides or suicide games; and
- sharing information, strategies, methods or instructions that would assist people to engage in self-harm and suicide.

What is not a violation of this policy?

Some examples of behavior that are not considered a violation of this policy include:

- sharing personal stories and experiences related to self-harm or suicide;
- sharing coping mechanisms and resources for addressing self-harm or suicidal thoughts; and
- discussions that are focused on research, advocacy, and education related to self-harm or suicide prevention.

Note: people can share their personal experiences, but should avoid sharing detailed information about specific strategies or methods related to self-harm, as this could inadvertently encourage this behavior.

Who can report violations of this policy?

Anyone can report content that may encourage or promote suicide or self-harm via our in-app reporting or our [specialized reporting form](#). These reports are routed to a dedicated team who evaluate each case individually.

Note: if we receive a report that someone has expressed an intention to engage in self-harm or suicide, we will contact them directly, encourage them to seek support, and provide information about dedicated online and hotline resources. We may also work with law enforcement officials where appropriate, for example, if we receive a valid emergency disclosure request as defined in our [Law Enforcement Guidelines](#).

How to report violations of this policy

To ensure that we handle reports sensitively, our in-app reporting provides separate options for people who may be expressing an intention to harm themselves, and content that is encouraging or promoting self-harm or suicide.

Expressing intentions of self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu

2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Reporting form

You can also report this content for review via [our dedicated report form](#).

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Encouraging self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

What happens if you violate this policy?

Our enforcement approach depends on the type of content being shared, whether or not the reported account is encouraging or promoting self-harm or suicide, and the account's previous history of violations.

If you violate this policy by sharing content that intentionally encourages others to harm themselves, ask others to encourage you to harm yourself, or share detailed information or instructions related to self-harm or suicide methods, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy, or if your account is dedicated to promoting or encouraging self-harm or suicide, your account will be permanently suspended. If cases include images or videos related to self-harm or suicide, we will also evaluate this content under [our sensitive media policy](#). If you believe that your account was suspended in error, [you can submit an appeal](#).

We may also take steps to prevent the spread of instructional material hosted on third-party websites by [marking such links as unsafe](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Learn more](#) about how you can support someone experiencing thoughts of self-harm or suicide.

Visit our Safety Center for a list of local [mental health resources](#) and read [our blog on suicide prevention](#) to learn more about our work.

Sensitive media policy (20191204*)

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and

- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and

- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20210922*)

Overview

March 2019

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We consider this because certain types of private information carry higher risks than

others, if they're shared without permission. Our primary aim is to protect individuals from coming to physical harm as a result of their information being shared, so we consider information like physical location to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;
- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;
 - descriptions of physical appearance;
 - gossip, rumours, accusations, and allegations; and
 - screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

Who can report violations of this policy?

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

The consequences for violating our private information policy depends on the severity of the violation and the accounts' previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you

violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy **(20191217*)**

Overview

November 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and

- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

- Select **Report Tweet** from the icon.
- Select **It displays a sensitive photo or video**.
- Select **An unauthorized photo or video**.
- Select **It includes unauthorized, intimate content of me or someone else**.
- Select the relevant option depending on who you are reporting on behalf of.
- Select up to 5 Tweets to report for review.
- Submit your report.

Desktop

You can report this content for review via desktop as follows:

- Select **Report Tweet** from the icon.
- Select **It displays a sensitive photo or video**.
- Select **An unauthorized photo or video**.
- Select **It includes unauthorized, intimate content of me or someone else**.
- Select the relevant option depending on who you are reporting on behalf of.

6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy **(20200918*)**

Overview

September 2020

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are;
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting; and
- coordinated harmful activity that encourages or promotes behavior which violates the [Twitter Rules](#).

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts or by coordinating with others to violate the Twitter Rules. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate;
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account; and
 - coordinating with others to engage in or promote violations of the Twitter Rules, including violations of our [abusive behavior](#) policy.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in “follow trains,” “decks,” and “Retweet for Retweet” behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- “follow churn” – following and then unfollowing large numbers of accounts in an effort to inflate one’s own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account’s followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person’s browser (malware) or computer or to compromise a person’s privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally;
- coordinating with others to express ideas, viewpoints, support, or opposition towards a cause, provided such behavior does not result in violations of the Twitter Rules; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don’t violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Denylisting URLs

We denylist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [financial scam policy](#), our [hacked materials policy](#), our approach to [coordinated harmful activity](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Civic integrity policy (20211006*)

Overview

October 2021

You may not use Twitter's services for the purpose of manipulating or interfering in elections or other civic processes. This includes posting or sharing content that may suppress participation or mislead people about when, where, or how to participate in a civic process. In addition, we may label and reduce the visibility of Tweets containing false or misleading information about civic processes in order to provide additional context.

The public conversation occurring on Twitter is never more important than during elections and other civic events. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt civic processes, including through the distribution of false or misleading information about the procedures or circumstances around participation in a civic process. In instances where misleading information does not seek to directly manipulate or disrupt civic processes, but leads to confusion on our service, we may label the Tweets to give additional context. Given the significant risks of confusion about key election information, we may take these actions even if Tweets contain (or attempt to contain) satirical or humorous elements.

What is a civic process?

Twitter considers civic processes to be events or procedures mandated, organized, and conducted by the governing and/or electoral body of a country, state, region, district, or municipality to address a matter of common concern through public participation. Some examples of civic processes may include:

- Political elections
- Censuses
- Major referenda and ballot initiatives

What is in violation of this policy?

This policy addresses 4 categories of misleading behavior and content:

Misleading information about how to participate

We will label or remove false or misleading information about how to participate in an election or other civic process. This includes but is not limited to:

- misleading information about procedures to participate in a civic process (for example, that you can vote by Tweet, text message, email, or phone call in jurisdictions where these are not a possibility);
- misleading information about requirements for participation, including identification or citizenship requirements;
- misleading claims that cause confusion about the established laws, regulations, procedures, and methods of a civic process, or about the actions of officials or entities executing those civic processes; and

- misleading statements or information about the official, announced date or time of a civic process.

Suppression and intimidation

We will label or remove false or misleading information intended to intimidate or dissuade people from participating in an election or other civic process. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to voting in an election, polling places, or collecting census information;
- misleading claims about long lines, equipment problems, or other disruptions at voting locations during election periods;
- misleading claims about process procedures or techniques which could dissuade people from participating; and
- threats regarding voting locations or other key places or events (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

Misleading information about outcomes

We will label or remove false or misleading information intended to undermine public confidence in an election or other civic process. This includes but is not limited to:

- disputed claims that could undermine faith in the process itself, such as unverified information about election rigging, ballot tampering, vote tallying, or certification of election results; and
- misleading claims about the results or outcome of a civic process which calls for or could lead to interference with the implementation of the results of the

process, e.g. claiming victory before election results have been certified, inciting unlawful conduct to prevent the procedural or practical implementation of election results (note that our violent threats policy may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral authority, or government entity. Read more about our [parody, commentary, and fan account policy.](#)

What is not a violation of this policy?

Not all false or untrue information about politics or civic processes constitutes manipulation or interference. In the absence of other policy violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected or appointed official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information;
- voting and audience participation for competitions, game shows, or other entertainment purposes; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election or civic process. Therefore, we enable reporting of false or misleading information about civic processes in advance of major events, for people located in the relevant countries and locations. We also work with select government and civil society partners in these countries to provide additional channels for reporting and expedited review.

For civic processes with multiple stages or parts, such as primary elections or lengthy campaigns, reporting will be enabled leading up to the first officially-sanctioned event associated with the civic process.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country at the relevant time, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in civic processes.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in a civic process.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our civic integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. In instances where accounts repeatedly violate this policy, we will use a strike system to determine if further enforcement actions should be applied. We believe this system further helps to reduce the spread of potentially harmful and misleading information on Twitter, particularly for high-severity violations of our rules.

The actions we take may include the following:

Tweet deletion

For high-severity violations of this policy, including (1) misleading information about how to participate, and (2) suppression and intimidation, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Tweet deletions accrue 2 strikes.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Labeling

In circumstances where we do not remove content which violates this policy, we may provide additional context on Tweets sharing the content where they appear on Twitter. This means we may:

- Apply a label and/or warning message to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Turn off people's ability to reply, Retweet, or like the Tweet;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended;
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or relevant Twitter policies; and/or
- Turn off likes, replies, and Retweets.

In most cases, we will take all of the above actions on Tweets we label. In some instances, we'll also turn off your ability to reply, Retweet, or like the Tweet. We prioritize producing Twitter Moments in cases where misleading content on Twitter is gaining significant attention and has caused public confusion on our service. Labels applied to Tweets accrue 1 strike.

Account locks and permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

Repeated violations of this policy are enforced against on the basis of the number of strikes an account has accrued for violations of this policy:

- 1 strike: No account-level action

- 2 strikes: 12-hour account lock
- 3 strikes: 12-hour account lock
- 4 strikes: 7-day account lock
- 5 or more strikes: Permanent suspension

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Impersonation policy (20191001*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody](#), [newsfeed](#), [commentary](#), and [fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Synthetic and manipulated media policy **(20210626*)**

Overview

You may not deceptively promote synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context.

You should be able to find reliable information on Twitter. That means understanding whether the content you see is real or fabricated and having the ability to find more context about what you see on Twitter. Therefore, we may label Tweets that include media (videos, audio, and images) that have been deceptively altered or fabricated. In addition, you may not share deceptively altered media on Twitter in ways that mislead or deceive people about the media's authenticity where threats to physical safety or other serious harm may result.

We use the following criteria as we consider Tweets and media for labeling or removal under this policy as part of our ongoing work to enforce our rules and ensure healthy and safe conversation on Twitter (additional information is available below):

1. Is the content synthetic or manipulated?

In order for content to be labeled or removed under this policy, we must have reason to believe that media, or the context in which media are presented, are significantly and deceptively altered or manipulated. Synthetic and manipulated media take many different forms and people can employ a wide range of technologies to produce these media. In assessing whether media have been significantly and deceptively altered or fabricated, some of the factors we consider include:

- whether the content has been substantially edited in a manner that fundamentally alters its composition, sequence, timing, or framing;
- any visual or auditory information (such as new video frames, overdubbed audio, or modified subtitles) that has been added or removed; and

- whether media depicting a real person have been fabricated or simulated

We are most likely to take action (either labeling or removal, as described below) on more significant forms of alteration, such as wholly synthetic audio or video or content that has been doctored (spliced and reordered, slowed down) to change its meaning. Subtler forms of manipulated media, such as isolative editing, omission of context, or presentation with false context, may be labeled or removed on a case-by-case basis.

We will not take action to label or remove media that have been edited in ways that do not fundamentally alter their meaning, such as retouched photos or color-corrected videos.

In order to determine if media have been significantly and deceptively altered or fabricated, we may use our own technology or receive reports through partnerships with third parties. In situations where we are unable to reliably determine if media have been altered or fabricated, we may not take action to label or remove them.

2. Is the content shared in a deceptive manner?

We also consider whether the context in which media are shared could result in confusion or misunderstanding or suggests a deliberate intent to deceive people about the nature or origin of the content, for example by falsely claiming that it depicts reality. We assess the context provided alongside media to see whether it makes clear that the media have been altered or fabricated. Some of the types of context we assess in order to make this determination include:

- The text of the Tweet accompanying or within media
- Metadata associated with media
- Information on the profile of the account sharing media
- Websites linked in the Tweet, or in the profile of the account sharing media

3. Is the content likely to impact public safety or cause serious harm?

Tweets that share synthetic and manipulated media are subject to removal under this policy if they are likely to cause serious harm. Some specific harms we consider include:

- Threats to the physical safety of a person or group
- Risk of mass violence or widespread civil unrest
- Threats to the privacy or ability of a person or group to freely express themselves or participate in civic events, such as:
 - Stalking or unwanted and obsessive attention
 - Targeted content that includes tropes, epithets, or material that aims to silence someone
 - Voter suppression or intimidation

While we have other rules also intended to address these forms of harm, including our policies on violent threats, election integrity, and hateful conduct, we will err toward removal in borderline cases that might otherwise not violate existing rules for Tweets that include synthetic or manipulated media.

We also consider the time frame within which the content may be likely to impact public safety or cause serious harm, and are more likely to remove content under this policy if we find that immediate harms are likely to result from the content's presence on Twitter.

Note: We may also take action on synthetic and manipulated content under our [non-consensual nudity policy](#) (such as pornographic media altered to insert the faces of people not actually involved) or other parts of the Twitter Rules.

Labeling and removal

In most cases, if we have reason to believe that media shared in a Tweet have been significantly and deceptively altered or fabricated, we will provide additional context on Tweets sharing the media where they appear on Twitter. This means we may:

- Apply a label to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended;
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or landing page; and/or
- Turn off likes, replies, and Retweets.

In most cases, we will take all of the above actions on Tweets we label.

Media that meet **all three of the criteria** defined above—i.e. that are synthetic or manipulated, shared in a deceptive manner, and is likely to cause harm—may not be shared on Twitter and are subject to removal. Accounts engaging in repeated or severe violations of this policy may be permanently suspended.

** Other parts of the Twitter Rules apply and may lead to the removal of the content, particularly where there is high likelihood of severe harm, such as a threat to someone's life or physical safety.*

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Copyright policy (20200917*)

What types of copyright complaints does Twitter respond to?

Twitter responds to copyright complaints submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review [Twitter's trademark policy](#). If you are concerned about a parody, newsfeed, commentary, or fan account, please see the relevant policy [here](#). These are generally not copyright issues.

Am I a copyright holder? How do I know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

What to consider before submitting a copyright complaint

Before submitting a copyright complaint to us, please consider whether or not the use could be considered [fair use](#).

If you have considered fair use, and you still wish to continue with a copyright complaint, you may want to first reach out to the user in question to see if you can resolve the matter directly with the user. You can reply to the user's Tweet or send the user a Direct Message and ask for them to remove your copyrighted content without having to contact Twitter.

Prior to submitting a formal complaint with Twitter, please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

Note: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph. If you're unsure whether or not you own the copyrights to a work, or if you're infringing upon someone else's work, please consult an attorney or another advisor.

What information do you need to process a copyright complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the complaint is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How do I file a copyright complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and [filing a copyright complaint](#). If you are logged in to [twitter.com](#), you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a full copy of your notice (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a copyright complaint.

How are claims processed?

We process copyright complaints in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. If you do not receive a ticket confirmation that means we did not receive your complaint and you should re-submit your complaint. However, please note, submitting duplicate copyright complaints will result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) and provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice. We will also forward a redacted copy of the complaint to [Lumen](#), with your personal information removed.

What information gets forwarded to the reported user(s)?

If we remove or disable access to the materials reported in a copyright complaint, the reported user(s) will receive a copy of the complaint, including the reporter's full name, email, street address, and any other information included in the complaint.

If you are uncomfortable sharing your contact information with the reported user(s), you may wish to consider appointing an agent to submit your DMCA notice on your behalf. Your agent will be required to submit the DMCA notice with valid contact information, and identify you as the content owner that they are representing.

What happens next?

Twitter's response to copyright complaints may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a copyright complaint, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the complaint, along with instructions for filing a counter-notice.

If you've not yet received a copy of the copyright complaint regarding the content removed from your account, please respond to the support ticket we sent you. If your complaint involves removal of content from Periscope, please see this help center [article](#) regarding how to request a copy of your complaint.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from **@Username** has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

My content was removed from Twitter

Why did I receive a copyright complaint?

If you receive a copyright complaint, it means that access to the content described in the complaint has been restricted. Please take the time to read through our correspondence to you, which includes information on the complaint we received as well as instructions on how to file a counter-notice. Please ensure that you are monitoring the email address associated with your Twitter account.

Tip: Removing the material reported in a copyright complaint will not resolve that complaint.

What if I want to contest the takedown?

If you believe that the materials reported in the copyright complaints were misidentified or removed in error, you may send us a counter-notification(s). A counter-notice is a request for Twitter to reinstate the removed material, and it has legal consequences. Alternatively, you may be able to seek a retraction of the copyright complaint from the reporter.

How do I seek a retraction?

The DMCA complaint you received includes the contact information of the reporter. You may want to reach out and ask them to retract their notice. The reporter can send retractions to copyright@twitter.com, and should include: (1) identification of the material that was disabled, and (2) a statement that the reporter would like to retract their DMCA notice. This is the fastest and most efficient means of resolving an unresolved copyright complaint. A retraction is at the sole discretion of the original reporter.

When should I file a counter-notice?

A counter-notice is a request for Twitter to reinstate the removed material, and is the start of a legal process that has legal consequences. For example, submitting a counter notice indicates that you consent to the jurisdiction of a U.S. Federal court and that you consent to the disclosure of your personal information to the reporter and Lumen website.

With these considerations in mind, you may file a counter-notice if you believe that this material was misidentified, or you have a good faith belief that the material should not have been removed. If you're unsure whether or not you should file a counter-notice, you may want to consult with an attorney.

Tip: Re-posting material removed in response to a copyright complaint may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notice rather than re-posting the material.

What information do you need to process a counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What happens after I submit a counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. This means that the contact information that is submitted in your counter-notice will be shared to the person who filed the original notice.

If the copyright owner disagrees that the content was removed in error or misidentification, they may pursue legal action against you. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a copyright complaint or counter-notice is serious business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

What happens if my account receives multiple copyright complaints?

If multiple copyright complaints are received Twitter may lock accounts or take other actions to warn repeat violators. These warnings may vary across Twitter's services. Under appropriate circumstances we may suspend user accounts under our repeat infringer policy. However, we may take retractions and counter-notices into account when applying our repeat infringer policy.

Our approach to policy development and enforcement philosophy (20190520*)

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was

broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.

The Twitter Rules (20200204*)

Twitter's purpose is to serve the public conversation. Violence, harassment and other similar types of behavior discourage people from expressing themselves, and ultimately diminish the value of global public conversation. Our rules are to ensure all people can participate in the public conversation freely and safely.

Safety

Violence: You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence. Learn more about our [violent threat](#) and [glorification of violence](#) policies.

Terrorism/violent extremism: You may not threaten or promote terrorism or violent extremism. [Learn more](#).

Child sexual exploitation: We have zero tolerance for child sexual exploitation on Twitter. [Learn more](#).

Abuse/harassment: You may not engage in the targeted harassment of someone, or incite other people to do so. This includes wishing or hoping that someone experiences physical harm. [Learn more](#).

Hateful conduct: You may not promote violence against, threaten, or harass other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. [Learn more](#).

Suicide or self-harm: You may not promote or encourage suicide or self-harm. [Learn more](#).

Sensitive media, including graphic violence and adult content: You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted. [Learn more](#).

Illegal or certain regulated goods or services: You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. [Learn more](#).

Privacy

Private information: You may not publish or post other people's private information (such as home phone number and address) without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so. [Learn more](#).

Non-consensual nudity: You may not post or share intimate photos or videos of someone that were produced or distributed without their consent. [Learn more](#).

Authenticity

Platform manipulation and spam: You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter. [Learn more](#).

Election integrity: You may not use Twitter's services for the purpose of manipulating or interfering in elections. This includes posting or sharing content that may suppress voter turnout or mislead people about when, where, or how to vote. [Learn more](#).

Impersonation: You may not impersonate individuals, groups, or organizations in a manner that is intended to or does mislead, confuse, or deceive others. [Learn more](#).

Synthetic and manipulated media: You may not deceptively share synthetic or manipulated media that are likely to cause harm. In addition, we may label Tweets containing synthetic and manipulated media to help people understand their authenticity and to provide additional context. [Learn more](#).

Copyright and trademark: You may not violate others' intellectual property rights, including copyright and trademark. Learn more about our [trademark policy](#) and [copyright policy](#).

Enforcement and Appeals

Learn more about [our approach to enforcement](#), including potential consequences for violating these rules or attempting to circumvent enforcement, as well as how to appeal.

Third-party advertising in video content

You may not submit, post, or display any video content on or through our services that includes third-party advertising, such as pre-roll video ads or sponsorship graphics, without our prior consent.

Note: we may need to change these rules from time to time in order to support our goal of promoting a healthy public conversation. The most current version is always available at <https://twitter.com/rules>.

Violent threats policy (20190607*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Healthy conversation is only possible when people feel safe from abuse and don't resort to using violent language. For this reason, we have a policy against threatening violence on Twitter. We define violent threats as statements of an intent to kill or inflict serious physical harm on a specific person or group of people.

What is in violation of this policy?

Under this policy, you can't state an intention to inflict violence on a specific person or group of people. We define intent to include statements like "I will", "I'm going to", or "I plan to", as well as conditional statements like "If you do X, I will". Violations of this policy include, but are not limited to:

- threatening to kill someone;
- threatening to sexually assault someone;
- threatening to seriously hurt someone and/or commit a other violent act that could lead to someone's death or serious physical injury; and
- asking for or offering a financial reward in exchange for inflicting violence on a specific person or group of people.

What is not a violation of this policy?

We recognize that some people use violent language as part of hyperbolic speech or between friends, so we also allow some forms of violent speech where it's clear that there is no abusive or violent intent, e.g., "I'll kill you for sending me that plot spoiler!".

This policy is enforced in tandem with our policies on [abusive behavior](#) and [hateful conduct](#). Statements that express a wish or hope that someone experiences physical harm, making vague or indirect threats, or threatening actions that are unlikely to cause serious or lasting injury are not actionable under this policy, but may be reviewed and actioned under those policies. **Note:** we also have a policy against [encouraging self-harm and suicide](#).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will immediately and permanently suspend any account found to be posting violent threats.

In rare cases, we may not suspend an account immediately. For example, if the reported content is a form of hyperbolic speech. In such cases, we may require you to remove this content. We may also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Glorification of violence policy

(20190407*)

Overview

March 2019

You may not threaten violence against an individual or a group of people. We also prohibit the glorification of violence.

Glorifying violent acts could inspire others to take part in similar acts of violence. Additionally, glorifying violent events where people were targeted on the basis of their protected characteristics (including: race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease) could incite or lead to further violence motivated by hatred and intolerance. For these reasons, we have a policy against content that glorifies acts of violence in a way that may inspire others to replicate those violent acts and cause real offline harm, or events where members of a protected group were the primary targets or victims.

What is in violation of this policy?

Under this policy, you can't glorify, celebrate, praise or condone violent crimes, violent events where people were targeted because of their membership in a protected group, or the perpetrators of such acts. We define glorification to include praising, celebrating, or condoning statements, such as "I'm glad this happened", "This person is my hero", "I wish more people did things like this", or "I hope this inspires others to act".

Violations of this policy include, but are not limited to, glorifying, praising, condoning, or celebrating:

- violent acts committed by civilians that resulted in death or serious physical injury, e.g., murders, mass shootings;
- attacks carried out by terrorist organizations or violent extremist groups (as defined by [our terrorism and violent extremism policy](#)); and
- violent events that targeted protected groups, e.g., the Holocaust, Rwandan genocide.

What is not a violation of this policy?

Our focus is on preventing the glorification of violence that could inspire others to replicate violent acts, as well as violent events where protected groups were the primary targets or victims. Exceptions may be made for violent acts by state actors, where violence was not primarily targeting protected groups.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.

5. Select up to 5 Tweets to report for review.
6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via [our abusive behavior reporting form](#), by selecting the **Harassment** option.

What happens if you violate this policy?

The consequences for violating our glorification of violence policy depends on the severity of the violation and the account's previous history of violations.

The first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy after receiving a warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn about the link between glorifying violent acts and offline harm, you can refer to the research of Susan Benesch: [Countering Dangerous Speech: New Ideas for Genocide Prevention](#) and the [Dangerous Speech Project](#).

Violent organizations policy (20201026*)

Overview

October 2020

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

You may not threaten or promote terrorism or violent extremism.

There is no place on Twitter for violent organizations, including terrorist organizations, violent extremist groups, or individuals who affiliate with and promote their illicit activities. The violence that these groups engage in and/or promote jeopardizes the physical safety and well-being of those targeted. Our assessments under this policy are informed by national and international terrorism designations, as well as our violent extremist group and violent organizations criteria.

Violent extremist groups are those that meet all of the below criteria:

- identify through their stated purpose, publications, or actions as an extremist group;
- have engaged in, or currently engage in, violence and/or the promotion of violence as a means to further their cause; and
- target civilians in their acts and/or promotion of violence.

Other violent organizations

Other violent organizations are those that meet all of the below criteria:

- a collection of individuals with a shared purpose; and
- have systematically targeted civilians with violence.

We examine a group's activities both on and off Twitter to determine whether they satisfy the above criteria.

What is in violation of this policy?

Under this policy, you can't affiliate with and promote the illicit activities of a terrorist organization or violent extremist group. Examples of the types of content that violate this policy include, but are not limited to:

- engaging in or promoting acts on behalf of a violent organization;
- recruiting for a violent organization;
- providing or distributing services (e.g., financial, media/propaganda) to further a violent organization's stated goals; and
- using the insignia or symbol of violent organizations to promote them or indicate affiliation or support.

What is not a violation of this policy?

We may make limited exceptions for groups that have reformed or are currently engaging in a peaceful resolution process, as well as groups with representatives who have been elected to public office through democratic elections. We may also make exceptions related to the discussion of terrorism or extremism for clearly educational or documentary purposes. This policy also doesn't apply to state or governmental organizations.

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

In-app

You can report this content for review in-app on iOS as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.

6. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Threatening violence or physical harm**.
4. Select the relevant option depending on who you are reporting on behalf of.
5. Select up to 5 Tweets to report for review.
6. Submit your report.

Report form

You can also report this content for review via our [abusive behavior reporting form](#), by selecting the **Specific violent threats involving physical safety or well-being** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we determine to be in violation of this policy**. If you believe that your account was suspended in error, [you can submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

To learn more about violations related to [specific threats of violence](#) or incidents involving wishing for the serious physical harm, death, or disease of an individual or group of people, please see our [glorification of violence policy](#).

Child sexual exploitation policy (20201012*)

Overview

October 2020

We have a zero-tolerance child sexual exploitation policy on Twitter.

Twitter has **zero tolerance towards any material that features or promotes child sexual exploitation**, one of the most serious violations of the [Twitter Rules](#). This may include media, text, illustrated, or computer-generated images. Regardless of the intent, viewing, sharing, or linking to child sexual exploitation material contributes to the re-victimization of the depicted children. This also applies to content that may further contribute to victimization of children through the promotion or glorification of child sexual exploitation. For the purposes of this policy, a minor is any person under the age of 18.

What is in violation of this policy?

Any content that depicts or promotes child sexual exploitation including, but not limited to:

- visual depictions of a child engaging in sexually explicit or sexually suggestive acts;
- illustrated, computer-generated or other forms of realistic depictions of a human child in a sexually explicit context, or engaging in sexually explicit acts;
- sexualized commentaries about or directed at a known or unknown minor; and
- links to third-party sites that host child sexual exploitation material.

The following behaviors are also not permitted:

- sharing fantasies about or promoting engagement in child sexual exploitation;
- expressing a desire to obtain materials that feature child sexual exploitation;
- recruiting, advertising or expressing an interest in a commercial sex act involving a child, or in harboring and/or transporting a child for sexual purposes;
- sending sexually explicit media to a child;
- engaging or trying to engage a child in a sexually explicit conversation;
- trying to obtain sexually explicit media from a child or trying to engage a child in sexual activity through blackmail or other incentives;
- identifying alleged victims of childhood sexual exploitation by name or image; and
- promoting or normalizing sexual attraction to minors as a form of identity or sexual orientation.

What is not a violation of this policy?

- Discussions related to child sexual exploitation are permitted, provided they don't normalise, promote or glorify child sexual exploitation **in any way**;
- Conversations about help-seeking behaviour of individuals who may struggle with an attraction to minors;
- Advocacy against illegal or harmful activity involving minors provided there is no sharing of, or linking to any material featuring child sexual exploitation; and
- Depictions of nude minors in a non-sexualized context or setting may be permitted in a limited number of scenarios (e.g., art works by internationally renowned artists that feature minors, news media reportage, media created and shared for scientific or educational purposes).

Who can report violations of this policy?

Anyone can report potential violations of this policy, whether they have a Twitter account or not.

How can I report violations of this policy?

If you think you've found a Twitter account distributing or promoting child sexual exploitation, [you can report it via our child sexual exploitation form](#). Provide the username and links to all relevant Tweets that led you to believe the account should be reviewed.

What happens if you violate this policy?

In the majority of cases, the consequence for violating our child sexual exploitation policy is **immediate and permanent suspension**. In addition, violators will be prohibited from creating any new accounts in the future. Note: when we're made aware of content depicting or promoting child sexual exploitation, including links to third party sites where this content can be accessed, they will be removed without further notice and reported to the National Center for Missing & Exploited Children (NCMEC).

In a limited number of situations, where we haven't identified any malicious intent, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Further violations will lead to your account being permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

If you've found content elsewhere on the internet that is distributing or promoting child sexual exploitation, you should [file a report with NCMEC](#), or with your local law enforcement. If you believe the content is hosted on a website outside of the United States, you can report it on the [International Association of Internet Hotlines website](#).

We partner with multiple organisations whose work is dedicated to fighting child sexual exploitation around the world. Please see the full list of [our child protection partners](#).

Abusive behavior (20211017*)

[Twitter Rules](#): You may not engage in the targeted harassment of someone, or incite other people to do so. We consider abusive behavior an attempt to harass, intimidate, or silence someone else's voice.

Rationale

On Twitter, you should feel safe expressing your unique point of view. We believe in freedom of expression and open dialogue, but that means little as an underlying philosophy if voices are silenced because people are afraid to speak up.

In order to facilitate healthy dialogue on the platform, and empower individuals to express diverse opinions and beliefs, we prohibit behavior that harasses or intimidates, or is otherwise intended to shame or degrade others. In addition to posing risks to people's safety, abusive behavior may also lead to physical and emotional hardship for those affected.

Learn more about our approach to [policy development and our enforcement philosophy](#).

When this applies

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation. When we review this type of content, it

may not be clear whether it is intended to harass an individual, or if it is part of a consensual conversation. To help our teams understand the context of a conversation, we may need to hear directly from the person being targeted, to ensure that we have the information needed prior to taking any enforcement action.

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior within Tweets or Direct Messages. For accounts engaging in abusive behavior on their profile, please refer to our [abusive profile policy](#). For behavior targeting people based on their race, ethnicity, national origin, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease, this may be in violation of our [hateful conduct policy](#).

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., “I will kill you.”

Note: We have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account. Wishing, hoping, or calling for serious harm on a person or group of people.

Wishing, hoping, or calling for serious harm on a person or group of people

We do not tolerate content that wishes, hopes, promotes, incites, or expresses a desire for death, serious bodily harm or serious disease against an individual or group of people. This includes, but is not limited to:

- Hoping that someone dies as a result of a serious disease e.g., “I hope you get cancer and die.”
- Wishing for someone to fall victim to a serious accident e.g., “I wish that you would get run over by a car next time you run your mouth.”

- Saying that a group of individuals deserves serious physical injury e.g., “If this group of protesters don’t shut up, they deserve to be shot.”

About wishes of harm exceptions on Twitter

We recognize that conversations regarding certain individuals credibly accused of severe violence may prompt outrage and associated wishes of harm. In these limited cases, we will request the user to delete the Tweet without any risk of account penalty, strike, or suspension. Examples are, but not limited to:

- “I wish all rapists to die.”
- “Child abusers should be hanged.”

Unwanted sexual advances

While some [consensual nudity and adult content is permitted](#) on Twitter, we prohibit unwanted sexual advances and content that sexually objectifies an individual without their consent. This includes, but is not limited to:

- sending someone unsolicited and/or unwanted adult media, including images, videos, and GIFs;
- unwanted sexual discussion of someone’s body;
- solicitation of sexual acts; and
- any other content that otherwise sexualizes an individual without their consent.

Using insults, profanity, or slurs with the purpose of harassing or intimidating others

We take action against the use of insults, profanity, or slurs to target others. In some cases, such as (but not limited to) severe, repetitive usage of insults or slurs where the primary intent is to harass or intimidate others, we may require Tweet removal. In other cases, such as (but not limited to) moderate, isolated usage of insults and profanity

where the primary intent is to harass or intimidate others, we may limit Tweet visibility as further described below. Please also note that while some individuals may find certain terms to be offensive, we will not take action against every instance where insulting terms are used.

Encouraging or calling for others to harass an individual or group of people

We prohibit behavior that encourages others to harass or target specific individuals or groups with abusive behavior. This includes, but is not limited to; calls to target people with abuse or harassment online and behavior that urges offline action such as physical harassment.

Denying mass casualty events took place

We prohibit content that denies that mass murder or other mass casualty events took place, where we can verify that the event occurred, and when the content is shared with abusive intent. This may include references to such an event as a “hoax” or claims that victims or survivors are fake or “actors.” It includes, but is not limited to, events like the Holocaust, school shootings, terrorist attacks, and natural disasters.

Do I need to be the target of this content for it to be reviewed for violating the Twitter Rules?

No, we review both first-person and bystander reports of such content.

Consequences

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual’s previous record of rule violations. The following is a list of potential enforcement options for content that violates this policy:

- Downranking Tweets in replies, except when the user follows the Tweet author.
- Making Tweets ineligible for amplification in Top search results and/or on timelines for users who don't follow the Tweet author.
- Excluding Tweets and/or accounts in email or in-product recommendations.
- Requiring Tweet removal.
 - For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension.
- Suspending accounts whose primary use we've determined is to engage in abusive behavior as defined in this policy, or who have shared violent threats.

Learn more about [our range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

Hateful conduct policy (20211221*)

Hateful conduct: You may not promote violence against or directly attack or threaten other people on the basis of race, ethnicity, national origin, caste, sexual orientation, gender, gender identity, religious affiliation, age, disability, or serious disease. We also do not allow accounts whose primary purpose is inciting harm towards others on the basis of these categories.

Hateful imagery and display names: You may not use hateful images or symbols in your profile image or profile header. You also may not use your username, display name, or profile bio to engage in abusive behavior, such as targeted harassment or expressing hate towards a person, group, or protected category.

Rationale

Twitter's mission is to give everyone the power to create and share ideas and information, and to express their opinions and beliefs without barriers. Free expression is a human right – we believe that everyone has a voice, and the right to use it. Our role is to serve the public conversation, which requires representation of a diverse range of perspectives.

We recognize that if people experience abuse on Twitter, it can jeopardize their ability to express themselves. Research has shown that some groups of people are disproportionately targeted with abuse online. This includes; women, people of color, lesbian, gay, bisexual, transgender, queer, intersex, asexual individuals, marginalized and historically underrepresented communities. For those who identify with multiple underrepresented groups, abuse may be more common, more severe in nature and more harmful.

We are committed to combating abuse motivated by hatred, prejudice or intolerance, particularly abuse that seeks to silence the voices of those who have been historically marginalized. For this reason, we prohibit behavior that targets individuals or groups with abuse based on their perceived membership in a protected category.

If you see something on Twitter that you believe violates our hateful conduct policy, please [report it to us](#).

When this applies

We will review and take action against reports of accounts targeting an individual or group of people with any of the following behavior, whether within Tweets or Direct Messages.

Violent threats

We prohibit content that makes violent threats against an identifiable target. Violent threats are declarative statements of intent to inflict injuries that would result in serious and lasting bodily harm, where an individual could die or be significantly injured, e.g., “I will kill you.”

Note: we have a zero tolerance policy against violent threats. Those deemed to be sharing violent threats will face immediate and permanent suspension of their account.

Wishing, hoping or calling for serious harm on a person or group of people

We prohibit content that wishes, hopes, promotes, incites, or expresses a desire for death, serious bodily harm, or serious disease against an entire protected category and/or individuals who may be members of that category. This includes, but is not limited to:

- Hoping that an entire protected category and/or individuals who may be members of that category dies as a result of a serious disease, e.g., “I hope all [nationality] get COVID and die.”
- Wishing for someone to fall victim to a serious accident, e.g., “I wish that you would get run over by a car next time you run your mouth.”
- Saying that a group of individuals deserve serious physical injury, e.g., “If this group of [slur] don’t shut up, they deserve to be shot.”
- Encouraging others to commit violence against an individual or a group based on their perceived membership in a protected category, e.g., “I’m in the mood to punch a [racial slur], who’s with me?”

References to mass murder, violent events, or specific means of violence where protected groups have been the primary targets or victims

We prohibit targeting individuals or groups with content that references forms of violence or violent events where a protected category was the primary target or victims,

where the intent is to harass. This includes, but is not limited to media or text that refers to or depicts:

- genocides, (e.g., the Holocaust);
- lynchings.

Incitement against protected categories

We prohibit inciting behavior that targets individuals or groups of people belonging to protected categories. This includes content intended:

- to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g., “all [religious group] are terrorists.”
- to incite others to harass members of a protected category on or off platform, e.g., “I’m sick of these [religious group] thinking they are better than us, if any of you see someone wearing a [religious symbol of the religious group], grab it off them and post pics!”
- to incite others to discriminate in the form of denial of support to the economic enterprise of an individual or group because of their perceived membership in a protected category, e.g., “If you go to a [religious group] store, you are supporting those [slur], let’s stop giving our money to these [religious slur].”
This may not include content intended as political in nature, such as political commentary or content relating to boycotts or protests.

Note that content intended to incite violence against a protected category is prohibited under **Wishing, hoping, or calling for serious harm on a person or groups of people**.

We prohibit targeting individuals and groups with content intended to incite fear or spread fearful stereotypes about a protected category, including asserting that members of a protected category are more likely to take part in dangerous or illegal activities, e.g.,

“all [religious group] are terrorists.”

Repeated and/or non-consensual slurs, epithets, racist and sexist tropes, or other content that degrades someone

We prohibit targeting others with repeated slurs, tropes or other content that intends to dehumanize, degrade or reinforce negative or harmful stereotypes about a protected category. This includes targeted misgendering or deadnaming of transgender individuals. We also prohibit the dehumanization of a group of people based on their religion, caste, age, disability, serious disease, national origin, race, ethnicity, gender, gender identity, or sexual orientation. In some cases, such as (but not limited to) severe, repetitive usage of slurs, epithets, or racist/sexy tropes where the primary intent is to harass or intimidate others, we may require Tweet removal. In other cases, such as (but not limited to) moderate, isolated usage where the primary intent is to harass or intimidate others, we may limit Tweet visibility as further described below.

Hateful imagery

We consider hateful imagery to be logos, symbols, or images whose purpose is to promote hostility and malice against others based on their race, religion, disability, sexual orientation, gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

Media depicting hateful imagery is not permitted within live video, account bio, profile or header images. All other instances must be marked as sensitive media. Additionally, sending an individual unsolicited hateful imagery is a violation of our [abusive behavior policy](#).

Do I need to be the target of this content for it to be a violation of the Twitter Rules?

Some Tweets may appear to be hateful when viewed in isolation, but may not be when viewed in the context of a larger conversation. For example, members of a protected category may refer to each other using terms that are typically considered as slurs. When used consensually, the intent behind these terms is not abusive, but a means to reclaim terms that were historically used to demean individuals.

When we review this type of content, it may not be clear whether the intention is to abuse an individual on the basis of their protected status, or if it is part of a consensual conversation. To help our teams understand the context, we sometimes need to hear directly from the person being targeted to ensure that we have the information needed prior to taking any enforcement action.

Note: individuals do not need to be a member of a specific protected category for us to take action. We will never ask people to prove or disprove membership in any protected category and we will not investigate this information.

Consequences

Under this policy, we take action against behavior that targets individuals or an entire protected category with hateful conduct, as described above. Targeting can happen in a number of ways, for example, mentions, including a photo of an individual, referring to someone by their full name, etc.

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to the severity of the violation and an individual's previous

record of rule violations. The following is a list of potential enforcement options for content that violates this policy:

- Downranking Tweets in replies, except when the user follows the Tweet author.
- Making Tweets ineligible for amplification in Top search results and/or on timelines for users who don't follow the Tweet author.
- Excluding Tweets and/or accounts in email or in-product recommendations.
- Requiring Tweet removal.
 - For example, we may ask someone to remove the violating content and serve a period of time in read-only mode before they can Tweet again. Subsequent violations will lead to longer read-only periods and may eventually result in permanent suspension.
- Suspending accounts whose primary use we've determined is to engage in hateful conduct as defined in this policy, or who have shared violent threats.

Learn more about our [range of enforcement options](#).

If someone believes their account was suspended in error, they can [submit an appeal](#).

Suicide and Self-harm Policy **(20200805*)**

Overview

You may not promote or encourage suicide or self-harm.

At Twitter, we recognize that suicide and self-harm are significant social & public health challenges that require collaboration between all stakeholders – public, private, and civil society – and that we have a role and responsibility to help people access and receive support when they need it.

When developing this policy, we consulted extensively with experts to ensure that people who have engaged in self-harm or experienced suicidal thoughts can share their personal experiences. We also recognized the need to protect people from the potential harm caused by exposure to content that could promote or encourage self-harm – intentionally or inadvertently. That's why our policy prohibits content that promotes or encourages self-harming behaviors and provides support to those undergoing experiences with self-harm or suicidal thoughts.

What is in violation of this policy?

Under this policy, you can't promote, or otherwise encourage, suicide or self-harm. We define promotion and encouragement to include statements such as "the most effective", "the easiest", "the best", "the most successful", "you should", "why don't you". Violations of this policy can occur via Tweets, images or videos, including live video.

We define suicide to be the act of taking one's own life. We define self-harm to include:

- self-inflicted physical injuries e.g., cutting; and
- eating disorders e.g., bulimia, anorexia.

Violations of this policy include, but are not limited to:

- encouraging someone to physically harm or kill themselves;
- asking others for encouragement to engage in self-harm or suicide, including seeking partners for group suicides or suicide games; and
- sharing information, strategies, methods or instructions that would assist people to engage in self-harm and suicide.

What is not a violation of this policy?

Some examples of behavior that are not considered a violation of this policy include:

- sharing personal stories and experiences related to self-harm or suicide;
- sharing coping mechanisms and resources for addressing self-harm or suicidal thoughts; and
- discussions that are focused on research, advocacy, and education related to self-harm or suicide prevention.

Note: people can share their personal experiences, but should avoid sharing detailed information about specific strategies or methods related to self-harm, as this could inadvertently encourage this behavior.

Who can report violations of this policy?

Anyone can report content that may encourage or promote suicide or self-harm via our in-app reporting or our [specialized reporting form](#). These reports are routed to a dedicated team who evaluate each case individually.

Note: if we receive a report that someone has expressed an intention to engage in self-harm or suicide, we will contact them directly, encourage them to seek support, and provide information about dedicated online and hotline resources. We may also work with law enforcement officials where appropriate, for example, if we receive a valid emergency disclosure request as defined in our [Law Enforcement Guidelines](#).

How to report violations of this policy

To ensure that we handle reports sensitively, our in-app reporting provides separate options for people who may be expressing an intention to harm themselves, and content that is encouraging or promoting self-harm or suicide.

Expressing intentions of self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu

2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Reporting form

You can also report this content for review via [our dedicated report form](#).

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It expresses intentions of self-harm or suicide"
3. Submit your report

Encouraging self-harm or suicide

In-App

You can report content for review in-app as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

Desktop

You can report this content for review on desktop as follows:

1. Select Report Tweet from drop-down menu
2. Select "It's abusive or harmful"
3. Select "They're encouraging self-harm or suicide"
4. Submit your report

What happens if you violate this policy?

Our enforcement approach depends on the type of content being shared, whether or not the reported account is encouraging or promoting self-harm or suicide, and the account's previous history of violations.

If you violate this policy by sharing content that intentionally encourages others to harm themselves, ask others to encourage you to harm yourself, or share detailed information or instructions related to self-harm or suicide methods, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you continue to violate this policy, or if your account is dedicated to promoting or encouraging self-harm or suicide, your account will be permanently suspended. If cases include images or videos related to self-harm or suicide, we will also evaluate this content under [our sensitive media policy](#). If you believe that your account was suspended in error, [you can submit an appeal](#).

We may also take steps to prevent the spread of instructional material hosted on third-party websites by [marking such links as unsafe](#).

Additional resources

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

[Learn more](#) about how you can support someone experiencing thoughts of self-harm or suicide.

Visit our Safety Center for a list of local [mental health resources](#) and read [our blog on suicide prevention](#) to learn more about our work.

Sensitive media policy (20191204*)

Overview

November 2019

You may not post media that is excessively gory or share violent or adult content within live video or in profile or header images. Media depicting sexual violence and/or assault is also not permitted.

People use Twitter to show what's happening in the world, often sharing images and videos as part of the conversation. Sometimes, this media can depict sensitive topics, including violent and adult content. We recognize that some people may not want to be exposed to sensitive content, which is why we balance allowing people to share this type of media with helping people who want to avoid it to do so.

For this reason, you can't include violent, hateful, or adult content within areas that are highly visible on Twitter, including in live video, profile or header images. If you share this content within Tweets, you need to [mark your account as sensitive](#). Doing so places images and videos behind an interstitial (or warning message), that needs to be acknowledged before your media can be viewed. Using this feature means that people who don't want to see sensitive media can avoid it, or make an informed decision before they choose to view it.

Under this policy, there are also some types of sensitive media content that we don't allow at all, because they have the potential to normalize violence and cause distress to those who view them.

We group sensitive media content into the following categories:

Graphic violence

Graphic violence is any media that depicts death, violence, medical procedures, or serious physical injury in graphic detail. Some examples include, but are not limited to, depictions of:

- violent crimes or accidents;
- physical fights;
- physical child abuse;
- bodily fluids including blood, feces, semen etc.;
- serious physical harm, including visible wounds; and
- severely injured or mutilated animals.

Note: exceptions may be made for documentary or educational content.

Adult content

Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal. Some examples include, but are not limited to, depictions of:

- full or partial nudity, including close-ups of genitals, buttocks, or breasts (excluding content related to breastfeeding);
- simulated sexual acts; and
- sexual intercourse or other sexual acts – this also applies to cartoons, hentai, or anime involving humans or depictions of animals with human-like features.

Note: exceptions may be made for artistic, medical, health, or educational content.

For content that was created or distributed without the consent of those featured, please refer to our [non-consensual nudity policy](#).

Violent sexual conduct

Violent sexual conduct is any media that depicts violence, whether real or simulated, in association with sexual acts. Some examples include, but are not limited to, depictions of:

- rape and other forms of violent sexual assault, or sexual acts that occur without the consent of participants, including a simulated lack of consent; and
- sexualized violence – inflicting physical harm on an individual within an intimate setting, where it is not immediately obvious if those involved have consented to take part.

Gratuitous gore

Gratuitous gore is any media that depicts excessively graphic or gruesome content related to death, violence or severe physical harm, or violent content that is shared for sadistic purposes. Some examples include, but are not limited to, depictions of:

- dismembered or mutilated humans;
- charred or burned human remains;
- exposed internal organs or bones; and
- animal torture or killing. **Note:** exceptions may be made for religious sacrifice, food preparation or processing, and hunting.

Hateful imagery

Hateful imagery is any logo, symbol, or image that has the intention to promote hostility against people on the basis of race, religious affiliation, disability, sexual orientation, gender/gender identity or ethnicity/national origin. Some examples of hateful imagery include, but are not limited to:

- symbols historically associated with hate groups, e.g., the Nazi swastika;
- images depicting others as less than human, or altered to include hateful symbols, e.g., altering images of individuals to include animalistic features; or
- images altered to include hateful symbols or references to a mass murder that targeted a protected category, e.g., manipulating images of individuals to include yellow Star of David badges, in reference to the Holocaust.

What is in violation of this policy?

Our aim is to limit exposure to sensitive images and videos and to prevent the sharing of potentially disturbing types of sensitive media. For this reason, we differentiate our enforcement approach depending on the type of media that has been shared and where it has been shared.

Graphic violence, adult content, and hateful imagery

- you can't target people with unsolicited images or videos that contain graphic violence, adult content, or hateful imagery; and

- you can't include graphic violence, adult content, or hateful imagery within live video, profile, or header images.

Violent sexual conduct and gratuitous gore

We prohibit violent sexual conduct to prevent the normalization of sexual assault and non-consensual violence associated with sexual acts. We prohibit gratuitous gore content because research has shown that repeated exposure to violent content online may negatively impact an individual's wellbeing. For these reasons, you can't share images or videos that depict violent sexual conduct or gratuitous gore on Twitter. **Note:** very limited exceptions may be made for gory media associated with newsworthy events.

What is not a violation of this policy?

You can share graphic violence and consensually produced adult content within your Tweets, provided that you mark this media as sensitive. We may also allow limited sharing of hateful imagery, provided that it is not used to promote a terrorist or violent extremist group, that you mark this content as sensitive and don't target it at an individual (via mentioning someone or including an identifiable individual within such images).

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we will do so manually if your content is reported for review.

Who can report violations of this policy?

Anyone can report potential violations of this policy via our dedicated reporting flows.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select the relevant option depending on what you are reporting.

What happens if you violate this policy?

The enforcement action we take depends on the type of media you have shared, and where you have shared it.

Graphic violence, adult content, and hateful imagery

- live video and profile images – the first time you violate this policy, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended;
- sending someone unsolicited violent or adult content – if you target someone with sensitive media in a clearly abusive or unsolicited way, we'll require you to remove it under our [abusive behavior policy](#); and

- accounts dedicated to posting sensitive media – your account may be permanently suspended if the majority of your activity on Twitter is sharing sensitive media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Violent sexual conduct and gratuitous gore

- if you share media that depicts violent sexual conduct or gratuitous gore anywhere on Twitter, we will require you to remove this content.
- if your account is dedicated to posting this type of content, your account will be immediately permanently suspended.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Adult content that was created or shared without the consent of those depicted is reviewed under our [non-consensual nudity policy](#).

Illegal or certain regulated goods or services (20190607*)

Overview

April 2019

You may not use our service for any unlawful purpose or in furtherance of illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services.

Twitter takes the safety of our users seriously. In an effort to address the potential for real-world harm, we do not allow the use of Twitter for any unlawful behavior or to further illegal activities. This includes selling, buying, or facilitating transactions in illegal goods or services, as well as certain types of regulated goods or services. In some cases, we may ask you to contact a law enforcement agency and have them contact us via our law enforcement request page to ensure we have enough context to enforce this policy. In addition to reports received, we proactively surface activity that may violate this policy for human review.

Goods or services covered under this policy include, but are not limited to:

- counterfeit goods and services*;
- drugs and controlled substances;
- human trafficking;
- products made from endangered or protected species;
- sexual services;
- stolen goods; and
- weapons, including firearms, ammunition, and explosives, and instructions on making weapons (e.g. bombs, 3D printed guns, etc.)

*Learn more about our [counterfeit](#) policy.

What happens if you violate this policy?

The consequences for violating this policy depends on the severity of the violation and the account's previous history of violations.

If you violate this policy more than once and/or if your account is dedicated to the sale of illegal or regulated goods and/or services, your account may be suspended permanently.

Accounts that appear to be using misleading account information in order to engage in spamming, abusive, or disruptive behavior to promote the sale of illegal and regulated goods and/or services may be subject to suspension under our [platform manipulation and spam](#) policy.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Private information policy (20211130)

Overview

November 2021

You may not publish or post other people's private information without their express authorization and permission. We also prohibit threatening to expose private information or incentivizing others to do so.

In addition, you may not share private media, such as images or videos of private individuals, without their consent. However, we recognise that there are instances where users may share images or videos of private individuals, who are not public figures, as part of a newsworthy event or to further public discourse on issues or events of public interest. In such cases, we may allow the media to remain on the platform.

Sharing someone's private information:

Sharing someone's private information online without their permission, sometimes called doxxing, is a breach of their privacy and of the [Twitter Rules](#). Sharing private

information can pose serious safety and security risks for those affected and can lead to physical, emotional, and financial hardship.

When reviewing reports under this policy, we consider a number of things, including:

What type of information is being shared?

We take this into consideration because certain types of private information carry higher risks than others, if they're shared without permission. Our primary aim is to protect individuals from potential physical harm as a result of their information being shared, so we consider information such as physical location and phone numbers to be a higher risk than other types of information.

Who is sharing the information?

We also consider who is sharing the reported information and whether or not they have the consent of the person it belongs to. We do this because we know that there are times when people may want some forms of their personal information to be shared publicly. For example, sharing a personal phone number or email for professional networking or to coordinate social events or publicly sharing someone's home addresses to seek help after a natural disaster.

Is the information available elsewhere online?

If the reported information was shared somewhere else before it was shared on Twitter, e.g., someone sharing their personal phone number on their own publicly accessible website, we may not treat this information as private, as the owner has made it publicly available. Note: we may take action against home addresses being shared, even if they are publicly available, due to the potential for physical harm.

Why is the information being shared?

We also factor in the intent of the person sharing the information. For example, if we believe that someone is sharing information with an abusive intent, or to harass or encourage others to harass another person, we will take action. On the other hand, if someone is sharing information in an effort to help someone involved in a crisis situation like in the aftermath of a violent event, we may not take action.

Sharing private media:

Posting images is an important part of our users' experience on Twitter. Where individuals have a reasonable expectation of privacy in an individual piece of media, we believe they should be able to determine whether or not it is shared. Sharing such media could potentially violate users' privacy and may lead to emotional or physical harm. When we are notified by individuals depicted, or their authorized representative, that they did not consent to having media shared, we will remove the media. This policy is not applicable to public figures.

What is in violation of this policy?

Under this policy, you can't share the following types of private information, without the permission of the person who it belongs to:

- home address or physical location information, including street addresses, GPS coordinates or other identifying information related to locations that are considered private;
- identity documents, including government-issued IDs and social security or other national identity numbers – note: we may make limited exceptions in regions where this information is not considered to be private;
- contact information, including non-public personal phone numbers or email addresses;

- financial account information, including bank account and credit card details; and
- other private information, including biometric data or medical records.
- media of private individuals without the permission of the person(s) depicted.

The following behaviors are also not permitted:

- threatening to publicly expose someone's private information;
- sharing information that would enable individuals to hack or gain access to someone's private information without their consent, e.g., sharing sign-in credentials for online banking services;
- asking for or offering a bounty or financial reward in exchange for posting someone's private information;
- asking for a bounty or financial reward in exchange for not posting someone's private information, sometimes referred to as blackmail.

What is not a violation of this policy?

The following are not in violation of this policy:

- people sharing their own private information;
- sharing information that is publicly available elsewhere, in a non-abusive manner; and
- sharing information that we don't consider to be private, including:
 - name;
 - birthdate or age;
 - place of education or employment;
 - location information related to commercial property or places of business, where this information is publicly available;

- descriptions of physical appearance;
- gossip, rumours, accusations, and allegations; and
- screenshots of text messages or messages from other platforms (unless they contain private information e.g., phone number).

For media, the following are not in violation of our policy:

- the media is publicly available or is being covered by mainstream media;
- the media and the accompanying tweet text add value to the public discourse or are shared in public interest;
- contains eyewitness accounts or on the ground reports from developing events;
- the subject of the media is a public figure.

Who can report violations of this policy?

Sharing private information:

Anyone can report private information that has been shared in a clearly abusive way (whether they have a Twitter account or not). In cases where the information hasn't been shared with a clearly abusive intent, we need to hear directly from the owner of this information (or an authorized representative, such as a lawyer) before taking enforcement action.

Sharing private media:

When reporting private media, we need a first person report in order to make the determination that the image or video has been shared without their permission. We accept reports from:

1. Individuals depicted in the reported media or

2. Their authorized representative such as:
 - a. parents/legal guardians,
 - b. lawyers, or
 - c. legal representative.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's abusive or harmful**.
3. Select **Includes private information**.
4. Select the type of information that you're reporting.
5. Select the relevant option depending on who owns the information you are reporting.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

You can also report this content for review via [our private information report form](#), by selecting the type of private information that you want to report.

What happens if you violate this policy?

When determining the penalty for violating this policy, we consider a number of factors including, but not limited to, the severity of the violation and an individual's previous record of rule violations.

Sharing private information:

- The first time you violate this policy by sharing private information (such as home address, identity documents etc.), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy by posting private information again after your first warning, your account will be permanently suspended.

Sharing private media:

- If you violate this policy by sharing private media of an individual without their consent, we will require you to remove the content and temporarily lock your account while you remove the media.

If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

The distribution of other types of content without the consent of the owner may be actionable under our distribution of hacked material policy.

Non-consensual nudity policy

(20191217*)

Overview

November 2019

You may not post or share intimate photos or videos of someone that were produced or distributed without their consent.

Sharing explicit sexual images or videos of someone online without their consent is a severe violation of their privacy and the [Twitter Rules](#). Sometimes referred to as revenge porn, this content poses serious safety and security risks for people affected and can lead to physical, emotional, and financial hardship.

What is a violation of this policy?

Under this policy, you can't post or share explicit images or videos that were taken, appear to have been taken or that were shared without the consent of the people involved.

Examples of the types of content that violate this policy include, but are not limited to:

- hidden camera content featuring nudity, partial nudity, and/or sexual acts;
- creepshots or upskirts - images or videos taken of people's buttocks, up an individual's skirt/dress or other clothes that allows people to see the person's genitals, buttocks, or breasts;
- images or videos that superimpose or otherwise digitally manipulate an individual's face onto another person's nude body;
- images or videos that are taken in an intimate setting and not intended for public distribution; and
- offering a bounty or financial reward in exchange for intimate images or videos.

What is not a violation of this policy?

Pornography and other forms of consensually produced adult content are allowed on Twitter, provided that this media is marked as sensitive. Doing so provides people who may not want to see this type of content with a warning that they will need to acknowledge before viewing your media.

To mark your media as sensitive, navigate to your [safety settings](#) and select the **Mark media you Tweet as containing material that may be sensitive** option. If you don't mark your media as sensitive, we may do so manually if your content is reported for review.

Who can report violations of this policy?

Because Twitter allows some types of adult content, we need to evaluate context to assess if reported content has been created or shared without the consent of those involved.

Content reportable by anyone

We recognize that it can be difficult for those impacted to report this type of content for review. To reduce the burden on those affected, anyone can report the following types of content:

- creepshots or upskirts;
- content where a bounty or financial reward is offered in exchange for non-consensual nudity media; and
- intimate images or videos that are accompanied by:
- text that wishes/hopes for harm to come to those depicted or otherwise refers to revenge e.g., "I hope you get what you deserve when people see this"; and
- information that could be used to contact those depicted e.g., "You can tell my ex what you think by calling them on 1234567".

Content reportable by featured individual(s)

For other types of content, we may need to hear directly from the individual(s) featured (or an authorized representative, such as a lawyer) to ensure that we have sufficient context before taking any enforcement action.

How can I report violations of this policy?

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Desktop

You can report this content for review via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It displays a sensitive photo or video**.
3. Select **An unauthorized photo or video**.
4. Select **It includes unauthorized, intimate content of me or someone else**.
5. Select the relevant option depending on who you are reporting on behalf of.
6. Select up to 5 Tweets to report for review.
7. Submit your report.

Report form

You can also report this content via [our private information report form](#), by selecting the **An unauthorized photo or video** option.

What happens if you violate this policy?

We will **immediately and permanently suspend any account that we identify as the original poster of intimate media** that was created or shared without consent. We will do the same with any account that posts only this type of content, e.g., accounts dedicated to sharing upskirt images.

In other cases, we may not suspend an account immediately. This is because some people share this content inadvertently, to express shock, disbelief or to denounce this practice. In these cases, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you

violate this policy again after your first warning, your account will be permanently suspended. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Platform manipulation and spam policy **(20200918*)**

Overview

September 2020

You may not use Twitter's services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on Twitter.

We want Twitter to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using Twitter to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on Twitter to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are;
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting; and
- coordinated harmful activity that encourages or promotes behavior which violates the [Twitter Rules](#).

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Accounts and identity

You can't mislead others on Twitter by operating fake accounts. This includes using misleading account information to engage in spamming, abusive, or disruptive behavior. Some of the factors that we take into consideration include:

- use of stock or stolen profile photos, particularly those depicting other people;
- use of stolen or copied profile bios; and
- use of intentionally misleading profile information, including profile location.

You can't artificially amplify or disrupt conversations through the use of multiple accounts or by coordinating with others to violate the Twitter Rules. This includes:

- overlapping accounts – operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;
- mutually interacting accounts – operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and
- coordination – creating multiple accounts to post duplicative content or create fake engagement, including:
 - posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;
 - engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate;
 - coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account; and
 - coordinating with others to engage in or promote violations of the Twitter Rules, including violations of our [abusive behavior](#) policy.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement. This includes:

- selling/purchasing Tweet or account metric inflation – selling or purchasing followers or engagements (Retweets, Likes, mentions, Twitter Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to Tweets;
- reciprocal inflation – trading or coordinating to exchange follows or Tweet engagements (including but not limited to participation in "follow trains," "decks," and "Retweet for Retweet" behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of Twitter accounts, usernames, or temporary access to Twitter accounts.

Misuse of Twitter product features

You can't misuse Twitter product features to disrupt others' experience. This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical Tweets, or repeatedly sending identical Direct Messages; and
- repeatedly posting Tweets or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your Tweet/Direct Message activity.

Following

- "follow churn" – following and then unfollowing large numbers of accounts in an effort to inflate one's own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account's followers, particularly using automation.

Engagement

- aggressively or automatically engaging with Tweets to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- Tweeting with excessive, unrelated hashtags in a single Tweet or across multiple Tweets.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person's browser (malware) or computer or to compromise a person's privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- using Twitter pseudonymously or as a [parody, commentary, or fan account](#);
- posting links without commentary occasionally;
- coordinating with others to express ideas, viewpoints, support, or opposition towards a cause, provided such behavior does not result in violations of the Twitter Rules; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don't violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Tweets via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's suspicious or spam**.
3. Select the option that best tells us how the Tweet is suspicious or spreading spam.
4. Submit your report.

Report form

You can also report this content for review via our [spam reporting form](#) by selecting the **I want to report spam on Twitter** option.

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Denylisting URLs

We denylist or provide warnings about URLs we believe to be unsafe. Read more about [unsafe links](#), including how to appeal if we've falsely identified your URL as unsafe.

Tweet deletion and temporary account locks

- If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one or more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.
- In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.
- If we believe you may be in violation of our fake accounts policy, we may require you provide government-issued identification (such as a driver's license or passport) in order to reinstate your account.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that Twitter is able to reliably attribute to entities known to violate the [Twitter Rules](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our [automation rules](#) for developers, our [election integrity](#) efforts, our [financial scam policy](#), our [hacked materials policy](#), our approach to [coordinated harmful activity](#), and our [guidelines for promotions and contests](#).

Learn more about our [range of enforcement options](#) and our approach to [policy development and enforcement](#).

Civic integrity policy (20211006*)

Overview

October 2021

You may not use Twitter's services for the purpose of manipulating or interfering in elections or other civic processes. This includes posting or sharing content that may suppress participation or mislead people about when, where, or how to participate in a civic process. In addition, we may label and reduce the visibility of Tweets containing false or misleading information about civic processes in order to provide additional context.

The public conversation occurring on Twitter is never more important than during elections and other civic events. Any attempts to undermine the integrity of our service is antithetical to our fundamental rights and undermines the core tenets of freedom of expression, the value upon which our company is based.

We believe we have a responsibility to protect the integrity of those conversations from interference and manipulation. Therefore, we prohibit attempts to use our services to manipulate or disrupt civic processes, including through the distribution of false or misleading information about the procedures or circumstances around participation in a civic process. In instances where misleading information does not seek to directly manipulate or disrupt civic processes, but leads to confusion on our service, we may label the Tweets to give additional context. Given the significant risks of confusion about key election information, we may take these actions even if Tweets contain (or attempt to contain) satirical or humorous elements.

What is a civic process?

Twitter considers civic processes to be events or procedures mandated, organized, and conducted by the governing and/or electoral body of a country, state, region, district, or municipality to address a matter of common concern through public participation. Some examples of civic processes may include:

- Political elections
- Censuses
- Major referenda and ballot initiatives

What is in violation of this policy?

This policy addresses 4 categories of misleading behavior and content:

Misleading information about how to participate

We will label or remove false or misleading information about how to participate in an election or other civic process. This includes but is not limited to:

- misleading information about procedures to participate in a civic process (for example, that you can vote by Tweet, text message, email, or phone call in jurisdictions where these are not a possibility);
- misleading information about requirements for participation, including identification or citizenship requirements;
- misleading claims that cause confusion about the established laws, regulations, procedures, and methods of a civic process, or about the actions of officials or entities executing those civic processes; and
- misleading statements or information about the official, announced date or time of a civic process.

Suppression and intimidation

We will label or remove false or misleading information intended to intimidate or dissuade people from participating in an election or other civic process. This includes but is not limited to:

- misleading claims that polling places are closed, that polling has ended, or other misleading information relating to votes not being counted;
- misleading claims about police or law enforcement activity related to voting in an election, polling places, or collecting census information;
- misleading claims about long lines, equipment problems, or other disruptions at voting locations during election periods;
- misleading claims about process procedures or techniques which could dissuade people from participating; and
- threats regarding voting locations or other key places or events (note that our [violent threats policy](#) may also be relevant for threats not covered by this policy).

Misleading information about outcomes

We will label or remove false or misleading information intended to undermine public confidence in an election or other civic process. This includes but is not limited to:

- disputed claims that could undermine faith in the process itself, such as unverified information about election rigging, ballot tampering, vote tallying, or certification of election results; and
- misleading claims about the results or outcome of a civic process which calls for or could lead to interference with the implementation of the results of the process, e.g. claiming victory before election results have been certified, inciting unlawful conduct to prevent the procedural or practical implementation of election results (note that our violent threats policy may also be relevant for threats not covered by this policy).

False or misleading affiliation

You can't create fake accounts which misrepresent their affiliation, or share content that falsely represents its affiliation, to a candidate, elected official, political party, electoral

authority, or government entity. Read more about our [parody, commentary, and fan account policy](#).

What is not a violation of this policy?

Not all false or untrue information about politics or civic processes constitutes manipulation or interference. In the absence of other policy violations, the following are generally not in violation of this policy:

- inaccurate statements about an elected or appointed official, candidate, or political party;
- organic content that is polarizing, biased, hyperpartisan, or contains controversial viewpoints expressed about elections or politics;
- discussion of public polling information;
- voting and audience participation for competitions, game shows, or other entertainment purposes; and
- using Twitter pseudonymously or as a [parody, commentary, or fan account](#) to discuss elections or politics.

Who can report violations of this policy?

Accurate reporting of suspected violations of this policy requires information and knowledge specific to an election or civic process. Therefore, we enable reporting of false or misleading information about civic processes in advance of major events, for people located in the relevant countries and locations. We also work with select government and civil society partners in these countries to provide additional channels for reporting and expedited review.

For civic processes with multiple stages or parts, such as primary elections or lengthy campaigns, reporting will be enabled leading up to the first officially-sanctioned event associated with the civic process.

How can I report violations of this policy?

If the reporting option for this policy is enabled in your country at the relevant time, you can report this content in-app or on desktop.

In-app

You can report this content for review in-app as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in civic processes.
4. Submit your report.

Desktop

You can report this content for review on desktop as follows:

1. Select **Report Tweet** from the icon.
2. Select **It's misleading about a political election or other civic event**.
3. Select the option that best tells us how the Tweet is misleading about voting or participation in a civic process.
4. Submit your report.

What happens if you violate this policy?

The consequences for violating our civic integrity policy depends on the severity and type of the violation and the accounts' history of previous violations. In instances where accounts repeatedly violate this policy, we will use a strike system to determine if further enforcement actions should be applied. We believe this system further helps to reduce the spread of potentially harmful and misleading information on Twitter, particularly for high-severity violations of our rules.

The actions we take may include the following:

Tweet deletion

For high-severity violations of this policy, including (1) misleading information about how to participate, and (2) suppression and intimidation, we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. Tweet deletions accrue 2 strikes.

Profile modifications

If you violate this policy within your profile information (e.g., your bio), we will require you to remove this content. We will also temporarily lock you out of your account before you can Tweet again. If you violate this policy again after your first warning, your account will be permanently suspended.

Labeling

In circumstances where we do not remove content which violates this policy, we may provide additional context on Tweets sharing the content where they appear on Twitter. This means we may:

- Apply a label and/or warning message to the content where it appears in the Twitter product;
- Show a warning to people before they share or like the content;

- Turn off people's ability to reply, Retweet, or like the Tweet;
- Reduce the visibility of the content on Twitter and/or prevent it from being recommended;
- Provide a link to additional explanations or clarifications, such as in a Twitter Moment or relevant Twitter policies; and/or
- Turn off likes, replies, and Retweets.

In most cases, we will take all of the above actions on Tweets we label. In some instances, we'll also turn off your ability to reply, Retweet, or like the Tweet. We prioritize producing Twitter Moments in cases where misleading content on Twitter is gaining significant attention and has caused public confusion on our service. Labels applied to Tweets accrue 1 strike.

Account locks and permanent suspension

For severe or repeated violations of this policy, accounts will be permanently suspended.

Repeated violations of this policy are enforced against on the basis of the number of strikes an account has accrued for violations of this policy:

- 1 strike: No account-level action
- 2 strikes: 12-hour account lock
- 3 strikes: 12-hour account lock
- 4 strikes: 7-day account lock
- 5 or more strikes: Permanent suspension

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Impersonation policy (20191001*)

Impersonation is a violation of the [Twitter Rules](#). Twitter accounts that pose as another person, brand, or organization in a confusing or deceptive manner may be permanently suspended under Twitter's impersonation policy.

What is not an impersonation policy violation?

Accounts with similar usernames or that are similar in appearance (e.g., the same profile image) are not automatically in violation of the impersonation policy. In order to violate our impersonation policy, the account must portray another entity in a misleading or deceptive manner.

An account will not be removed if:

- The user shares your name but has no other commonalities, or
- The profile clearly states it is not affiliated with or connected to any similarly-named individuals or brands.

Twitter users are allowed to create parody, newsfeed, commentary, or fan accounts. Please refer to Twitter's [parody, newsfeed, commentary, and fan account policy](#) for more information about these types of accounts.

When does Twitter review accounts under this policy?

We review impersonation claims upon receipt of a valid impersonation report. We do not actively monitor users' content.

How do I file an impersonation report?

If you are the impersonated party or an authorized representative, you can [file an impersonation report](#). If you believe an account is impersonating somebody else, you can flag it as a bystander by [reporting](#) directly from the account's profile.

Please note, submitting a report with incomplete information will delay the processing of your report.

Can I request a specific username?

We generally do not accept individual requests for inactive or suspended usernames.

What if my trademark is being misused?

Companies and businesses should report accounts misusing their trademarks by filing a complaint under our [trademark policy](#). Please review the trademark policy page for more information.

What happens if you violate this policy?

If we determine that you violated our impersonation policy, we may suspend your account. In some instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Synthetic and manipulated media policy **(20211207*)**

Overview

You may not share synthetic, manipulated, or out-of-context media that may deceive or confuse people and lead to harm (“misleading media”). In addition, we may label Tweets containing misleading media to help people understand their authenticity and to provide additional context.

What is in violation of this policy

In order for content with **misleading media** (including images, videos, audios, gifs, and URLs hosting relevant content) to be labeled or removed under this policy, it must:

- Include media that is significantly and deceptively altered, manipulated, or fabricated, or
- Include media that is shared in a deceptive manner or with false context, and

- Include media likely to result in widespread confusion on public issues, impact public safety, or cause serious harm

We use the following criteria as we consider Tweets and media for labeling or removal under this policy as part of our ongoing work to enforce our rules and ensure healthy and safe conversations on Twitter:

1. Is the content significantly and deceptively altered, manipulated, or fabricated?

In order for content to be labeled or removed under this policy, we must have reason to believe that media are significantly and deceptively altered, manipulated, or fabricated. Synthetic and manipulated media take many different forms and people can employ a wide range of technologies to produce these media. Some of the factors we consider include:

- whether media have been substantially edited or post-processed in a manner that fundamentally alters their composition, sequence, timing, or framing and distorts their meaning;
- whether there are any visual or auditory information (such as new video frames, overdubbed audio, or modified subtitles) that has been added, edited, or removed that fundamentally changes the understanding, meaning, or context of the media;
- whether media have been created, edited, or post-processed with enhancements or use of filters that fundamentally changes the understanding, meaning, or context of the content; and
- whether media depicting a real person have been fabricated or simulated, especially through use of artificial intelligence algorithms

We will not take action to label or remove media that have been edited in ways that do not fundamentally alter their meaning, such as retouched photos or color-corrected videos.

In order to determine if media have been significantly and deceptively altered or fabricated, we may use our own technology or receive reports through partnerships

with third parties. In situations where we are unable to reliably determine if media have been altered or fabricated, we may not take action to label or remove them.

2. Is the content shared in a deceptive manner or with false context?

We also consider whether the context in which media are shared could result in confusion or suggests a deliberate intent to deceive people about the nature or origin of the content, for example, by falsely claiming that it depicts reality. We assess the context provided alongside media to see whether it provides true and factual information. Some of the types of context we assess in order to make this determination include:

- whether misleading, inauthentic, fictional, or produced media are presented or being endorsed as fact or reality, including produced or staged works, reenactments, or exhibitions portrayed as actual events;
- whether media are presented with false or misleading context surrounding the source, location, time, or authenticity of the media;
- whether media are presented with false or misleading context surrounding the identity of the individuals visually depicted in the media

We will not take action to label or remove media that have been shared with commentary or opinions that do not advance or present a misleading claim on the context of the media as listed above.

In order to determine if media have been shared in a deceptive manner or with false context, we may use our own technology or receive reports through partnerships with third parties. In situations where we are unable to reliably determine if media have been shared with false context, we will not label or remove the content.

3. Is the content likely to result in widespread confusion on public issues, impact public safety, or cause serious harm?

Tweets that share misleading media are subject to removal under this policy if they are likely to cause serious harm. Some specific harms we consider include:

- Threats to physical safety of a person or group
- Incitement of abusive behavior to a person or group
- Risk of mass violence or widespread civil unrest
- Risk of impeding or complicating provision of public services, protection efforts, or emergency response
- Threats to the privacy or to the ability of a person or group to freely express themselves or participate in civic events, such as:
 - Stalking or unwanted and obsessive attention
 - Targeted content that includes tropes, epithets, or material that aims to silence someone
 - Voter suppression or intimidation

We also consider the time frame within which the content may be likely to impact public safety or cause serious harm, and are more likely to remove content under this policy if immediate harm is likely to result.

Tweets with misleading media that are not likely to result in immediate harm but still have a potential to impact public safety, result in harm, or cause widespread confusion towards a public issue (health, environment, safety, civil rights and equality, immigration, and social and political stability) may be labeled to reduce their spread and to provide additional context.

While we have other rules also intended to address these forms of harm, including our policies on violent threats, civic integrity, COVID-19 misleading information, and hateful conduct, we will err toward removal in borderline cases that might otherwise not violate existing rules for Tweets that include misleading media.

What is not a violation of this policy

We seek to protect public conversation surrounding various issues. Media often accompany these conversations and encourage further discourse. In the absence of other policy violations, the following are generally not in violation of this policy:

- **Memes or satire**, provided these do not cause significant confusion about the authenticity of the media;
- **Animations, illustrations, and cartoons**, provided these do not cause significant confusion about the authenticity of the media.
- **Commentary, reviews, opinions, and/or reactions**. Sharing media with edits that only add commentary, reviews, opinions, or reactions allows for further debate and discourse relating to various issues and are not in violation of this policy.
- **Counterspeech**. We allow for direct responses to misleading information which seek to undermine its impact by correcting the record, amplifying credible information, and educating the wider community about the prevalence and dynamics of misleading information.
- **Doctored or fake Tweets, social media posts, or chat messages**. Due to the challenges associated with conclusively verifying whether an alleged Tweet, post, or message existed, we do not enforce on doctored or fake Tweets, social media posts, or chat messages under this policy.

Who can report violations of this policy?

We enforce this policy in close coordination with trusted partners, including [our partnership with AP and Reuters](#), other news agencies, public health authorities, and governments. Our team has open lines of communication with various partners to consult and get various media and claims reviewed.

In Australia, South Korea, and the US, Twitter has [begun testing](#) a new reporting feature that will allow users to report Tweets that seem misleading. As part of the experiment, the phrase “It’s misleading” will appear as an option when you select **Report an issue**.

What happens if you violate this policy?

The consequences for violating our synthetic and manipulated media policy depends on the severity of the violation.

Tweet Deletion

For high-severity violations of the policy, including misleading media that have a

serious risk of harm to individuals or communities, we will require you to remove this content.

Labeling

In circumstances where we do not remove content which violates this policy, we may provide additional context on Tweets sharing the misleading media where they appear on Twitter. This means we may:

- Apply a label and/or warning message to the Tweet
- Show a warning to people before they share or like the Tweet;
- Reduce the visibility of the Tweet on Twitter and/or prevent it from being recommended;
- Turn off likes, replies, and Retweets; and/or
- Provide a link to additional explanations or clarifications, such as in a curated landing page (Twitter Moments) or relevant Twitter policies.

In most cases, we will take a combination of the above actions on Tweets we label. We prioritize producing Twitter Moments in cases where misleading content on Twitter is gaining significant attention and has caused public confusion on our service.

Account locks

If we determine that an account has advanced harmful misleading narratives that violate the synthetic and manipulated media policy, we may temporarily lock or suspend the account.

If you believe that your account was locked or suspended in error, you can [submit an appeal](#).

Additional resources

Learn more about our work and how we build rules to fight misleading media [here](#).

Learn more about [our range of enforcement options](#) and [our approach to policy development and enforcement](#).

Trademark policy (20190607*)

Overview

March 2019

You may not violate others' intellectual property rights, including copyright and trademark.

A trademark is a word, logo, phrase, or device that distinguishes a trademark holder's good or service in the marketplace. Trademark law may prevent others from using a trademark in an unauthorized or confusing manner.

What is in violation of this policy?

Using another's trademark in a way that may mislead or confuse people about your affiliation may be a violation of our trademark policy.

What is not a violation of this policy?

Referencing another's trademark is not automatically a violation of Twitter's trademark policy. Examples of non-violations include:

- using a trademark in a way that is outside the scope of the trademark registration e.g., in a different territory, or a different class of goods or services than that identified in the registration; and
- using a trademark in a nominative or other fair use manner. For more information, see our [parody, newsfeed, commentary, and fan account policy](#).

Who can report violations of this policy?

Twitter only investigates requests that are submitted by the trademark holder or their authorized representative e.g., a legal representative or other representative for a brand.

How can I report violations of this policy?

You can submit a trademark report through our [trademark report form](#). Please provide all the information requested in the form. If you submit an incomplete report, we'll need to follow up about the missing information. Please note that this will result in a delay in processing your report.

Note: We may provide the account holder with your name and other information included in the copy of the report.

What happens if you violate this policy?

If we determine that you violated our trademark policy, we may suspend your account. Depending on the type of violation, we may give you an opportunity to comply with our policies. In other instances, an account may be permanently suspended upon first review. If you believe that your account was suspended in error, you can [submit an appeal](#).

What is our policy about trademark issues within ads?

Read more about [Twitter Ads](#) and our [Trademark Policy for Promoted Ads](#).

Additional resources

Learn more about [our range of enforcement options](#) and our approach to [policy development and enforcement](#).

Copyright policy (20200917*)

What types of copyright complaints does Twitter respond to?

Twitter responds to copyright complaints submitted under the Digital Millennium Copyright Act ("DMCA"). Section 512 of the DMCA outlines the statutory requirements necessary for formally reporting copyright infringement, as well as providing instructions on how an affected party can appeal a removal by submitting a compliant counter-notice.

Twitter will respond to reports of alleged copyright infringement, such as allegations concerning the unauthorized use of a copyrighted image as a profile or header photo, allegations concerning the unauthorized use of a copyrighted video or image uploaded through our media hosting services, or Tweets containing links to allegedly infringing materials. Note that not all unauthorized uses of copyrighted materials are infringements (see our [fair use](#) article for more information).

If you are concerned about the use of your brand or entity's name, please review [Twitter's trademark policy](#). If you are concerned about a parody, newsfeed, commentary, or fan account, please see the relevant policy [here](#). These are generally not copyright issues.

Am I a copyright holder? How do I know?

If you are unsure whether you hold rights to a particular work, please consult an attorney or another adviser as Twitter cannot provide legal advice. There are plenty of resources to learn more about copyright law including <http://copyright.gov>, <https://lumendatabase.org/>, and <http://www.eff.org/issues/bloggers/legal/liability/IP>, to name a few.

What to consider before submitting a copyright complaint

Before submitting a copyright complaint to us, please consider whether or not the use could be considered [fair use](#).

If you have considered fair use, and you still wish to continue with a copyright complaint, you may want to first reach out to the user in question to see if you can resolve the matter directly with the user. You can reply to the user's Tweet or send the user a Direct Message and ask for them to remove your copyrighted content without having to contact Twitter.

Prior to submitting a formal complaint with Twitter, please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a notification with us.

Note: In general, the photographer and NOT the subject of a photograph is the actual rights holder of the resulting photograph. If you're unsure whether or not you own the copyrights to a work, or if you're infringing upon someone else's work, please consult an attorney or another advisor.

What information do you need to process a copyright complaint?

To submit a notice of claimed copyright infringement, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice) of the copyright owner or a person authorized to act on their behalf;
2. Identification of the copyrighted work claimed to have been infringed (e.g., a link to your original work or clear description of the materials allegedly being infringed upon);
3. Identification of the infringing material and information reasonably sufficient to permit Twitter to locate the material on our website or services;
4. Your contact information, including your address, telephone number, and an email address;
5. A statement that you have a good faith belief that the use of the material in the manner asserted is not authorized by the copyright owner, its agent, or the law; and
6. A statement that the information in the complaint is accurate, and, under penalty of perjury, that you are authorized to act on behalf of the copyright owner.

If you are reporting the content of a Tweet, please give us a [direct link to that Tweet](#). Or please specify if the alleged infringement is in the header, avatar, etc. A LINK TO A PROFILE PAGE IS INSUFFICIENT FOR TWITTER TO IDENTIFY INFRINGING MATERIALS.

How do I file a copyright complaint?

You can report alleged copyright infringement by visiting Twitter's Help Center and [filing a copyright complaint](#). If you are logged in to [twitter.com](#), you can visit the Twitter Help Center directly from your Twitter account by clicking the 'Help' link located in the sidebar.

Filing a DMCA complaint is the start of a pre-defined legal process. Your complaint will be reviewed for accuracy, validity, and completeness. If your complaint has satisfied these requirements, we will take action on your request - which includes forwarding a full copy of your notice (including your name, address, phone and email address) to the user(s) who posted the allegedly infringing material in question.

If you are concerned about your contact information being forwarded, you may wish to use an agent to report for you.

Please be aware that under 17 U.S.C. § 512(f), you may be liable for any damages, including costs and attorneys' fees incurred by us or our users, if you knowingly materially misrepresent that material or activity is infringing. If you are unsure whether the material you are reporting is in fact infringing, you may wish to contact an attorney before filing a copyright complaint.

How are claims processed?

We process copyright complaints in the order in which they are received. Once you've submitted your ticket, we will email you a ticket confirmation. If you do not receive a ticket confirmation that means we did not receive your complaint and you should re-submit your complaint. However, please note, submitting duplicate copyright complaints will result in a delay in processing.

If we decide to remove or disable access to the material, we will notify the affected user(s) and provide them with a full copy of the reporter's complaint (including the provided contact information) along with instructions on how to file a counter-notice. We will also forward a redacted copy of the complaint to [Lumen](#), with your personal information removed.

What information gets forwarded to the reported user(s)?

If we remove or disable access to the materials reported in a copyright complaint, the reported user(s) will receive a copy of the complaint, including the reporter's full name, email, street address, and any other information included in the complaint.

If you are uncomfortable sharing your contact information with the reported user(s), you may wish to consider appointing an agent to submit your DMCA notice on your behalf. Your agent will be required to submit the DMCA notice with valid contact information, and identify you as the content owner that they are representing.

What happens next?

Twitter's response to copyright complaints may include the removal or restriction of access to allegedly infringing material. If we remove or restrict access to user content in response to a copyright complaint, Twitter will make a good faith effort to contact the affected account holder with information concerning the removal or restriction of access, including a full copy of the complaint, along with instructions for filing a counter-notice.

If you've not yet received a copy of the copyright complaint regarding the content removed from your account, please respond to the support ticket we sent you. If your complaint involves removal of content from Periscope, please see this help center [article](#) regarding how to request a copy of your complaint.

In an effort to be as transparent as possible regarding the removal or restriction of access to user-posted content, we clearly mark withheld Tweets and media to indicate to viewers when content has been withheld (examples below). We also send a redacted copy of each copyright complaint and counter-notice that we process to Lumen, where they are posted to a public-facing website (with your personal information removed).

Tweet withheld

18h

This Tweet from **@Username** has been withheld in response to a report from the copyright holder. [Learn more](#)



Media not displayed

This image has been removed in response to a report from the copyright holder.

My content was removed from Twitter

Why did I receive a copyright complaint?

If you receive a copyright complaint, it means that access to the content described in the complaint has been restricted. Please take the time to read through our correspondence to you, which includes information on the complaint we received as well as instructions on how to file a counter-notice. Please ensure that you are monitoring the email address associated with your Twitter account.

Tip: Removing the material reported in a copyright complaint will not resolve that complaint.

What if I want to contest the takedown?

If you believe that the materials reported in the copyright complaints were misidentified or removed in error, you may send us a counter-notification(s). A counter-notice is a request for Twitter to reinstate the removed material, and it has legal consequences. Alternatively, you may be able to seek a retraction of the copyright complaint from the reporter.

How do I seek a retraction?

The DMCA complaint you received includes the contact information of the reporter. You may want to reach out and ask them to retract their notice. The reporter can send retractions to copyright@twitter.com, and should

include: (1) identification of the material that was disabled, and (2) a statement that the reporter would like to retract their DMCA notice. This is the fastest and most efficient means of resolving an unresolved copyright complaint. A retraction is at the sole discretion of the original reporter.

When should I file a counter-notice?

A counter-notice is a request for Twitter to reinstate the removed material, and is the start of a legal process that has legal consequences. For example, submitting a counter notice indicates that you consent to the jurisdiction of a U.S. Federal court and that you consent to the disclosure of your personal information to the reporter and Lumen website.

With these considerations in mind, you may file a counter-notice if you believe that this material was misidentified, or you have a good faith belief that the material should not have been removed. If you're unsure whether or not you should file a counter-notice, you may want to consult with an attorney.

Tip: Re-posting material removed in response to a copyright complaint may result in permanent account suspension. If you believe the content was removed in error, please file a counter-notice rather than re-posting the material.

What information do you need to process a counter-notice?

To submit a counter-notice, you will need to provide us with the following information:

1. A physical or electronic signature (typing your full name will suffice);
2. Identification of the material that has been removed or to which access has been disabled and the location at which the material appeared before it was removed or access to it was disabled (the description from the copyright notice will suffice);
3. A statement under penalty of perjury that you have a good faith belief that the material was removed or disabled as a result of mistake or misidentification of the material to be removed or disabled; and
4. Your name, address, and telephone number, and a statement that you consent to the jurisdiction of the Federal District Court for the judicial district in which the address is located, or if your address is outside of the United States, for any judicial district in which Twitter may be found, and that you will accept service of process from the person who provided notification under subsection (c)(1)(C) or an agent of such person.

To submit a counter-notice, please respond to our original email notification of the removal and include the required information in the body of your reply as we discard all attachments for security reasons.

What happens after I submit a counter-notice?

Upon receipt of a valid counter-notice, we will promptly forward a copy to the person who filed the original notice. This means that the contact information that is submitted in your counter-notice will be shared to the person who filed the original notice.

If the copyright owner disagrees that the content was removed in error or misidentification, they may pursue legal action against you. If we do not receive notice within 10 business days that the original reporter is seeking a court order to prevent further infringement of the material at issue, we may replace or cease disabling access to the material that was removed.

We cannot offer any legal advice. Should you have questions, please consult an attorney.

Filing a copyright complaint or counter-notice is serious business!

Please think twice before submitting a claim or counter-notice, especially if you are unsure whether you are the actual rights holder or authorized to act on a rights holder's behalf. There are legal and financial consequences for fraudulent and/or bad faith submissions. Please be sure that you are the actual rights holder, or that you have a good faith belief that the material was removed in error, and that you understand the repercussions of submitting a false claim.

What happens if my account receives multiple copyright complaints?

If multiple copyright complaints are received Twitter may lock accounts or take other actions to warn repeat violators. These warnings may vary across Twitter's services. Under appropriate circumstances we may suspend

user accounts under our repeat infringer policy. However, we may take retractions and counter-notices into account when applying our repeat infringer policy.

Our approach to policy development and enforcement philosophy

(20190520*)

Twitter is reflective of real conversations happening in the world and that sometimes includes perspectives that may be offensive, controversial, and/or bigoted to others. While we welcome everyone to express themselves on our service, we will not tolerate behavior that harasses, threatens, or uses fear to silence the voices of others.

We have the [Twitter Rules](#) in place to help ensure everyone feels safe expressing their beliefs and we strive to enforce them with uniform consistency. Learn more about [different enforcement actions](#).

Our policy development process

Creating a new policy or making a policy change requires in-depth research around trends in online behavior, developing clear external language that sets expectations around what's allowed, and creating enforcement guidance for reviewers that can be scaled across millions of Tweets.

While drafting policy language, we gather feedback from a variety of internal teams as well as our Trust & Safety Council. This is vital to ensure we are considering global perspectives around the changing nature of online speech, including how our rules are applied and interpreted in different cultural and social contexts. Finally, we train our global review teams, update the Twitter Rules, and start enforcing the new policy.

Our enforcement philosophy

We empower people to understand different sides of an issue and encourage dissenting opinions and viewpoints to be discussed openly. This approach allows many forms of speech to exist on our platform and, in particular, promotes counterspeech: speech that presents facts to correct misstatements or misperceptions, points out hypocrisy or contradictions, warns of offline or online consequences, denounces hateful or dangerous speech, or helps change minds and disarm.

Thus, **context matters**. When determining whether to take enforcement action, we may consider a number of factors, including (but not limited to) whether:

- the behavior is directed at an individual, group, or protected category of people;
- the report has been filed by the target of the abuse or a bystander;
- the user has a history of violating our policies;
- the severity of the violation;
- the content may be a topic of legitimate public interest.

Is the behavior directed at an individual or group of people?

To strike a balance between allowing different opinions to be expressed on the platform, and protecting our users, we enforce policies when someone reports abusive behavior that targets a specific person or group of people. This targeting can happen in a number of ways (for example, @mentions, tagging a photo, mentioning them by name, and more).

Has the report been filed by the target of the potential abuse or a bystander?

Some Tweets may seem to be abusive when viewed in isolation, but may not be when viewed in the context of a larger conversation or historical relationship between people on the platform. For example, friendly banter between friends could appear offensive to bystanders, and certain remarks that are acceptable in one culture or country may not be acceptable in another. To help prevent our teams from making a mistake and removing consensual interactions, in certain scenarios we require a [report](#) from the actual target (or their authorized representative) prior to taking any enforcement action.

Does the user have a history of violating our policies?

We start from a position of assuming that people do not intend to violate our Rules. Unless a violation is so egregious that we must immediately suspend an account, we first try to educate people about our Rules and give them a chance to correct their behavior. We show the violator the offending Tweet(s), explain which Rule was broken, and require them to remove the content before they can Tweet again. If someone repeatedly violates our Rules then our enforcement actions become stronger. This includes requiring violators to remove the Tweet(s) and taking additional actions like verifying account ownership and/or temporarily limiting their ability to Tweet for a set period of time. If someone continues to violate Rules beyond that point then their account may be permanently suspended.

What is the severity of the violation?

Certain types of behavior may pose serious safety and security risks and/or result in physical, emotional, and financial hardship for the people involved. These egregious violations of the Twitter Rules — such as posting violent threats, non-consensual intimate media, or content that sexually exploits children — result in the immediate and permanent suspension of an account. Other violations could lead to a range of different steps, like requiring someone to remove the offending Tweet(s) and/or temporarily limiting their ability to post new Tweet(s).

Is the behavior newsworthy and in the legitimate public interest?

Twitter moves at the speed of public consciousness and people come to the service to stay informed about what matters. Exposure to different viewpoints can help people learn from one another, become more tolerant, and make decisions about the type of society we want to live in.

To help ensure people have an opportunity to see every side of an issue, there may be the rare occasion when we allow controversial content or behavior which may otherwise violate our Rules to remain on our service because we believe there is a legitimate public interest in its availability. Each situation is evaluated on a case by case basis and ultimately decided upon by a cross-functional team.

Some of the factors that help inform our decision-making about content are the impact it may have on the public, the source of the content, and the availability of alternative coverage of an event.

Public impact of the content: A topic of legitimate public interest is different from a topic in which the public may be curious. We will consider what the impact is to citizens if they do **not** know about this content. If the Tweet does have the potential to impact the lives of large numbers of people, the running of a country, and/or it speaks to an important societal issue then we may allow the the content to remain on the service. Likewise, if the impact on the public is minimal we will most likely remove content in violation of our policies.

Source of the content: Some people, groups, organizations and the content they post on Twitter may be considered a topic of legitimate public interest by virtue of their being in the public consciousness. This does not mean that their Tweets will always remain on the service. Rather, we will consider if there is a legitimate public interest for a particular Tweet to remain up so it can be openly discussed.

Availability of coverage: Everyday people play a crucial role in providing firsthand accounts of what's happening in the world, counterpoints to establishment views, and, in some cases, exposing the abuse of power by someone in a position of authority. As a situation unfolds, removing access to certain information could inadvertently hide context and/or prevent people from seeing every side of the issue. Thus, before actioning a potentially violating

Tweet, we will take into account the role it plays in showing the larger story and whether that content can be found elsewhere.